Vulkan™
Programming Guide

The Official Guide to Learning Vulkan

Graham Sellers

With contributions from John Kessenich
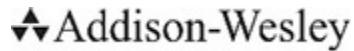
About This E-Book

EPUB is an open, industry-standard format for e-books. However, support for EPUB and its many features varies across reading devices and applications. Use your device or app settings to customize the presentation to your liking. Settings that you can customize often include font, font size, single or double column, landscape or portrait mode, and figures that you can click or tap to enlarge. For additional information about the settings and features on your reading device or app, visit the device manufacturer's Web site.

Many titles include programming code or configuration examples. To optimize the presentation of these elements, view the e-book in single-column, landscape mode and adjust the font size to the smallest setting. In addition to presenting code and configurations in the reflowable text format, we have included images of the code that mimic the presentation found in the print book; therefore, where the reflowable format may compromise the presentation of the code listing, you will see a "Click here to view code image" link. Click the link to view the print-fidelity code image. To return to the previous page viewed, click the Back button on your device or app.

Vulkan™ Programming Guide

The Official Guide to Learning Vulkan

Graham Sellers
With contributions from John Kessenich

✦✦ Addison-Wesley

Boston • Columbus • Indianapolis • New York • San Francisco • Amsterdam • Cape Town
Dubai • London • Madrid • Milan • Munich • Paris • Montreal • Toronto • Delhi • Mexico City
São Paulo • Sydney • Hong Kong • Seoul • Singapore • Taipei • Tokyo

Many of the designations used by manufacturers and sellers to distinguish their products are claimed as trademarks. Where those designations appear in this book, and the publisher was aware of a trademark claim, the designations have been printed with initial capital letters or in all capitals.

The author and publisher have taken care in the preparation of this book, but make no expressed or implied warranty of any kind and assume no responsibility for errors or omissions. No liability is assumed for incidental or consequential damages in connection with or arising out of the use of the information or programs contained herein.

For information about buying this title in bulk quantities, or for special sales opportunities (which may include electronic versions; custom cover designs; and content particular to your business, training goals, marketing focus, or branding interests), please contact our corporate sales department at corpsalespearsoned.com or (800) 382-3419.

For government sales inquiries, please contact governmentsales@pearsoned.com.

For questions about sales outside the U.S., please contact intlcs@pearson.com.

Visit us on the Web: informit.com/aw

Library of Congress Control Number: 2016948832

Vulkan and the Vulkan logo are trademarks of the Khronos Group Inc.

ISBN-13: 978-0-13-446454-1
ISBN-10: 0-13-446454-0
Text printed in the United States.
1   16
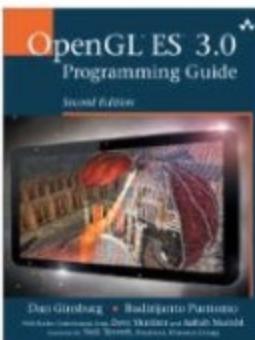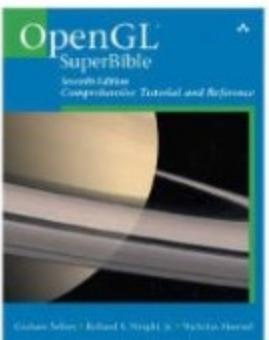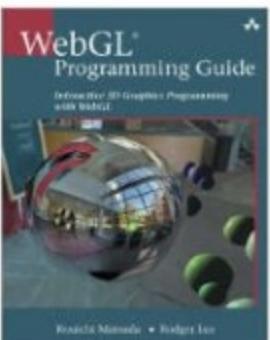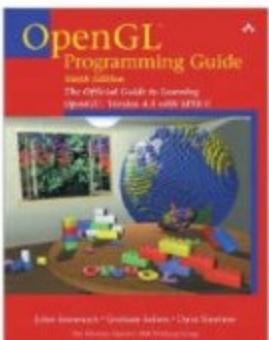
OpenGL Series
from Addison-Wesley

The OpenGL graphics system is a software interface to graphics hardware. ("GL" stands for "Graphics Library".) It allows you to create interactive programs that produce color images of moving, three-dimensional objects. With OpenGL, you can control computer-graphics technology to produce realistic pictures, or ones that depart from reality in imaginative ways.

The **OpenGL Series** from Addison-Wesley Professional comprises tutorial and reference books that help programmers gain a practical understanding of OpenGL standards, along with the insight needed to unlock OpenGL's full potential.

For you, the reader.
—Graham Sellers

Contents

Figures

Tables

Listings

About This Book

This book is about Vulkan. Vulkan is an application programming interface (API) for controlling devices such as graphics processing units (GPUs). Although Vulkan is a logical successor to OpenGL, it is quite different from OpenGL in form. One of the things that experienced practitioners will notice about Vulkan is that it is very verbose. You need to write a lot of application code to get Vulkan to do anything useful, let alone anything remarkable. Many of the things that an OpenGL driver would do are now the responsibility of the Vulkan application writer. These things include synchronization, scheduling, memory management, and so on. As such, you will find a good deal of this book dedicated to such topics, even though they are general topics applicable to more than just Vulkan.

The intended audience for this book is experienced programmers who are already familiar with other graphics and compute APIs. As such, many graphics-related topics are discussed without deep introduction, there are some forward references, and code samples are incomplete or illustrative in scope rather than being complete programs that you can type in. The sample code available from the book's website is complete and tested, however, and should serve as a good reference to follow along with.

Vulkan is intended to be used as the interface between large, complex graphics and compute applications and graphics hardware. Many of the features and responsibilities previously assumed by drivers implementing APIs such as OpenGL now fall to the application. Complex game engines, large rendering packages, and commercial middleware are well-suited to this task; they have more information about their specific behavior than any driver could hope to have. Vulkan is *not* well-suited to simple test applications; neither is it a suitable aid for teaching graphics concepts.

In the first chapters of this book, we introduce Vulkan and some of the fundamental concepts that frame the API. As we progress through the Vulkan system, we cover more advanced topics, eventually producing a more complex rendering system that shows off some of the unique aspects of Vulkan and demonstrates its capabilities.

In Chapter 1, "Overview of Vulkan," we provide a brief introduction to Vulkan and the concepts that form its foundation. We cover the basics of creating Vulkan objects and show the basics of getting started with the Vulkan system.

In Chapter 2, "Memory and Resources," we introduce the memory system of Vulkan, perhaps the most fundamental part of the interface. We show how to allocate memory used by the Vulkan device and by Vulkan drivers and system components running inside your application.

In Chapter 3, "Queues and Commands," we cover *command buffers* and introduce the *queues* to which they are submitted. We show how Vulkan processes work and how your application can build packets of commands to be sent to the device for execution.

In Chapter 4, "Moving Data," we introduce our first few Vulkan commands, all of which are focused on moving data. We use the concepts first discussed in Chapter 3 to build command buffers that can copy and format data stored in the resources and memory introduced in Chapter 2.

In Chapter 5, "Presentation," we show how to get images produced by your application onto the screen. *Presentation* is the term used for interacting with a window system, which is platform-specific, so this chapter delves into some platform-specific topics.

In Chapter 6, "Shaders and Pipelines," we introduce SPIR-V, the binary shading language used by Vulkan. We also introduce the *pipeline* object; show how one is constructed using SPIR-V shaders; and then introduce *compute pipelines*, which can be used to do computation work with Vulkan.

In Chapter 7, "Graphics Pipelines," we build upon what we covered in Chapter 6 and introduce the *graphics pipeline*, which includes all of the configuration necessary to render graphical primitives with Vulkan.

In Chapter 8, "Drawing," we discuss the various drawing commands available in Vulkan, including indexed and nonindexed draws, instancing, and indirect commands. We show how to get data into the graphics pipeline and how to draw more complex geometries than were introduced in Chapter 7.

In Chapter 9, "Geometry Processing," we dig deeper into the first half of the Vulkan graphics pipeline and take another look at the tessellation and geometry shader stages. We show some of the more advanced things that these stages can do and cover the pipeline up to the rasterization stage.

In Chapter 10, "Fragment Processing," we pick up where Chapter 9 left off and cover everything that happens during and after rasterization in order to turn your geometry into a stream of pixels that can be displayed to the user.

In Chapter 11, "Synchronization," we cover the various synchronization primitives available to the Vulkan application, including *fences*, *events*, and *semaphores*. Together, these form the foundation of any application that makes efficient use of the parallel nature of Vulkan.

In Chapter 12, "Getting Data Back," we reverse the direction of communication used in previous chapters and discuss the issues involved in reading data from Vulkan into your application. We show how to time operations executed by a Vulkan device, how to gather statistics about the operation of Vulkan devices, and how to get data produced by Vulkan back into your application.

Finally, in Chapter 13, "Multipass Rendering," we revisit a number of topics covered earlier, tying things together to produce a more advanced application—a deferred rendering application using complex multipass architecture and multiple queues for processing.

The appendix to this book contains a table of the command buffer building functions available to Vulkan applications, providing a quick reference to determine their attributes.

Vulkan is a large, complex, and new system. It is extremely difficult to cover every corner of the API in a book of this scope. The reader is encouraged to thoroughly read the Vulkan specification in addition to this book, as well as to read other books on using heterogeneous compute systems and computer graphics with other APIs. Such material will provide a good foundation in the mathematics and other concepts assumed by this book.

About the Sample Code

The sample code that accompanies this book is available from our website (http://www.vulkanprogrammingguide.com). One thing that seasoned users of other graphics APIs will notice is that Vulkan is very verbose. This is primarily because many of the responsibilities historically assumed by drivers have been delegated to your application. In many cases, however, simple boilerplate code will do the job just fine. Therefore, we have created a simple application framework that deals with much of the functionality that will be common to all samples and real-world applications. This does not mean that this book is a tutorial on how to use our framework. This is simply a practical matter of keeping code samples concise.

Of course, as we discuss specific Vulkan functionality throughout the book, we will include snippets of code, many of which may actually come from the book's sample framework rather than from any

particular example. Some features discussed in the book may not have examples in the code package. This is particularly true of some advanced features that are relevant primarily to large-scale applications. There is no such thing as a short, simple Vulkan example. In many cases, a single example program demonstrates the use of many features. The features that each example uses are listed in that example's read-me file. Again, there is not a 1:1 correspondence between examples and listings in this book and specific examples in the sample code. It shall be assumed that anyone that files a bug asking for a 1:1 list of which samples go with which chapter has not read this paragraph. Such bugs will be summarily closed, quoting this very sentence.

The sample code is designed to link against the latest official Vulkan SDK from LunarG, which is available from http://lunarg.com/vulkan-sdk/. At the time of writing, the latest SDK version is 1.0.22. Newer versions of the SDK are designed to be backward-compatible with older versions, so we recommend that users obtain the latest available version of the SDK before attempting to compile and run the sample applications. The SDK also comes with some samples of its own, and we suggest running those to verify that the SDK and drivers are installed correctly.

In addition to the Vulkan SDK, you will need a working installation of CMake in order to create the build environment for the samples. You will also need an up-to-date compiler. The code samples make use of several C++11 features and rely heavily on the C++ standard libraries for things like threading and synchronization primitives. These features are known to be problematic in early versions of various compiler runtimes, so please make sure that your compilers are up to date. We have tested with Microsoft Visual Studio 2015 on Windows and with GCC 5.3 on Linux. The samples have been tested on 64-bit Windows 7, Windows 10, and Ubuntu 16.10 with recent drivers from AMD, Intel, and NVIDIA.

It should be noted that Vulkan is a cross-platform, cross-vendor, and cross-device system. Many of these samples *should* work on Android and other mobile platforms. We hope to port the samples to as many of these platforms as possible in the future and would very much appreciate help and contributions from you, the reader.

Errata

Vulkan is a new technology. At the time of writing, the specification has been available for only a matter of weeks. Although the author and contributor had a hand in creating the Vulkan specification, it's large and complex and had many contributors. Some of the code in the book is not fully tested, and although it is believed to be correct, it may contain errors. As we were putting the samples together, available Vulkan implementations still had bugs, the validation layers didn't catch as many errors as they could, and the specification itself had gaps and unclear sections. Like the readers, we are still learning Vulkan, so although this text was edited for technical accuracy, we depend on readers to view any updates by visiting this book's website:

http://www.vulkanprogrammingguide.com

Register your copy of *Vulkan^{TM} Programming Guide* at informit.com for convenient access to downloads, updates, and corrections as they become available. To start the registration process, go to informit.com/register and log in or create an account. Enter the product ISBN (9780134464541) and click Submit. Once the process is complete, you will find any available bonus content under "Registered Products."

Acknowledgments

About the Author

Graham Sellers is a software architect at AMD who oversees the development of the OpenGL and Vulkan drivers for AMD's Radeon and FirePro products. His passion for computers and technology began at a young age with a BBC Micro, followed by a long line of 8- and 16-bit home computers that he still enjoys working with. He earned his master's in engineering from the University of Southampton, England, and now lives in Orlando, Florida, with his wife and two children.

Chapter 1. Overview of Vulkan

What You'll Learn in This Chapter

- • What Vulkan is and the fundamentals behind it
- • How to create a minimal Vulkan application
- • The terminology and concepts used in the remainder of this book

---

In this chapter, we introduce Vulkan and explain what it is. We introduce some of the fundamental concepts behind the API, including initialization, object lifetimes, the Vulkan instance, and logical and physical devices. By the end of the chapter, we produce a simple Vulkan application that can initialize the Vulkan system, discover available Vulkan devices and show their properties and capabilities, and finally shut down cleanly.

Introduction

Vulkan is a programming interface for graphics and compute devices. A Vulkan device typically consists of a processor and a number of fixed-function hardware blocks to accelerate operations used in graphics and compute. The processor in the device is usually a very wide multithreaded processor and so the computational model in Vulkan is heavily based on parallel computing. The Vulkan device also has access to memory that may or may not be shared with the main processor upon which your application is running. Vulkan also exposes this memory to you.

Vulkan is an explicit API. That is, almost everything is your responsibility. A driver is a piece of software that takes the commands and data forming the API and translates them into something that hardware can understand. In older APIs such as OpenGL, drivers would track the state of a lot of objects, manage memory and synchronization for you, and check for errors in your application as it ran. This is great for developers but burns valuable CPU time once your application is debugged and running correctly. Vulkan addresses this by placing almost all state tracking, synchronization, and memory management into the hands of the application developer and by delegating correctness checks to *layers* that must be enabled. They do not participate in the execution of your application under normal circumstances.

For these reasons, Vulkan is both very verbose and somewhat fragile. You need to do an awful lot of work to get Vulkan running well, and incorrect usage of the API can often lead to graphical corruption or even program crashes where in older APIs you would have received a helpful error message. In exchange for this, Vulkan provides more control over the device, a clean threading model, and much higher performance than the APIs that it supersedes.

Further, Vulkan has been designed to be more than a *graphics* API. It can be used for heterogeneous devices such as graphics processing units (GPUs), digital signal processors (DSPs), and fixed-function hardware. Functionality is divided into coarse-grained, broadly overlapping categories. The current edition of Vulkan defines the transfer category, which is used for copying data around; the compute category, which is used for running shaders over compute workloads; and the graphics category, which includes rasterization, primitive assembly, blending, depth and stencil tests, and other functionality that will be familiar to graphics programmers.

To the extent that support for each category is optional, it's possible to have a Vulkan device that doesn't support graphics at all. As a consequence, even the APIs to put pictures onto a display device (which is called *presentation*) are not only optional, but are provided as *extensions* to Vulkan rather than being part of the core API.

Instances, Devices, and Queues

Vulkan includes a hierarchy of functionality, starting at the top level with the *instance*, which aggregates all Vulkan-capable *devices* together. Each device then exposes one or more *queues*. It is the queues that perform the work that your application requests.

The Vulkan instance is a software construct that logically separates the state of your application from other applications or libraries running within the context of your application. The physical devices in the sytem are presented as members of the instance, each of which has certain capabilities, including a selection of available queues.

A physical device usually represents a single piece of hardware or a collection of hardware that is interconnected. There is a fixed, finite number of physical devices in any system unless that system supports reconfiguration such as hot-plug. A logical device, which is created by the instance, is the software construct around a physical device and represents a reservation of resources associated with a particular physical device. This includes a possible subset of the available queues on the physical device. It is possible to create multiple logical devices representing a single physical device, and it is the logical device that your application will spend most of its time interacting with.

Figure 1.1 illustrates this hierarchy. In the figure, the application has created two Vulkan instances. There are three physical devices in the system that are available to both instances. After enumeration, the application creates one logical device on the first physical device, two logical devices for the second device, and another for the third. Each logical device enables a different subset of its corresponding physical device's queues. In practice, most Vulkan applications won't be nearly this complex and will simply create a single logical device for one of the physical devices in the system, using a single instance. Figure 1.1 only serves to demonstrate the flexibility of the Vulkan system.
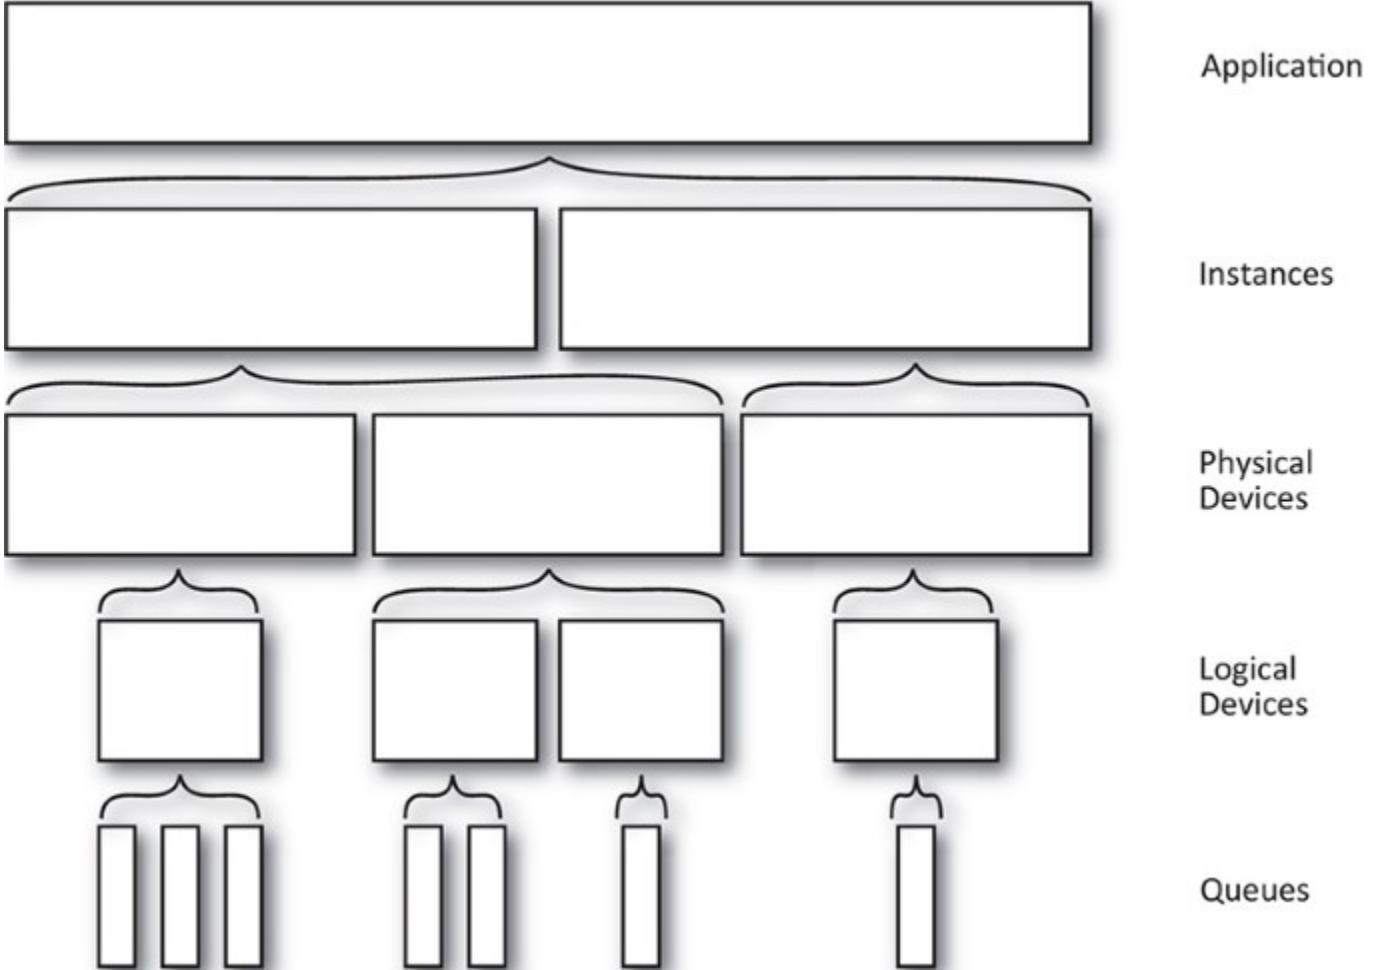
Figure 1.1: Vulkan Hierarchy of Instance, Device, and Queue

The following subsections discuss how to create the Vulkan instance, query the physical devices in the system, attach a logical device corresponding to one of them, and finally retrieve handles to the queues exposed by the device.

The Vulkan Instance

Vulkan can be seen as a subsystem of your application. Once you link your application to the Vulkan libraries and initialize it, it tracks some state. Because Vulkan doesn't introduce any global state into your application, all tracked state must be stored in an object that you provide. This is the *instance* object and is represented by a `VkInstance` object. To construct one, we'll call our first Vulkan function, **`vkCreateInstance()`**, the prototype of which is

Click here to view code image

```
VkResult vkCreateInstance (
    const VkInstanceCreateInfo*        pCreateInfo,
    const VkAllocationCallbacks*       pAllocator,
    VkInstance*                        pInstance);
```

This declaration is typical of a Vulkan function. Where more than a handful of parameters are to be passed to Vulkan, functions often take pointers to structures. Here, `pCreateInfo` is a pointer to an

instance of the `VkInstanceCreateInfo` structure that contains the parameters describing the new Vulkan instance. The definition of `VkInstanceCreateInfo` is

[**Click here to view code image**](#)

```
typedef struct VkInstanceCreateInfo {
    VkStructureType             sType;
    const void*                 pNext;
    VkInstanceCreateFlags       flags;
    const VkApplicationInfo*    pApplicationInfo;
    uint32_t                    enabledLayerCount;
    const char* const*          ppEnabledLayerNames;
    uint32_t                    enabledExtensionCount;
    const char* const*          ppEnabledExtensionNames;
} VkInstanceCreateInfo;
```

The first member in almost every Vulkan structure that is used to pass parameters to the API is the `sType` field, which tells Vulkan what type of structure this is. Each structure in the core API and in any extension has an assigned structure tag. By inspecting this tag, Vulkan tools, layers, and drivers can determine the type of the structure for validation purposes and for use in extensions. Further, the `pNext` field allows a *linked list* of structures to be passed to the function. This allows the set of parameters to be extended without needing to replace the core structure wholesale in an extension. Because we are using the core instance creation structure here, we pass `VK_STRUCTURE_TYPE_INSTANCE_CREATE_INFO` in the `sType` field and simply set `pNext` to `nullptr`.

The `flags` field of `VkInstanceCreateInfo` is reserved for future use and should be set to zero. The next field, `pApplicationInfo`, is an optional pointer to another structure describing your application. You can set this to `nullptr`, but a well-behaved application should fill this in with something useful. `pApplicationInfo` points to an instance of the `VkApplicationInfo` structure, the definition of which is

[**Click here to view code image**](#)

```
typedef struct VkApplicationInfo {
    VkStructureType    sType;
    const void*        pNext;
    const char*        pApplicationName;
    uint32_t           applicationVersion;
    const char*        pEngineName;
    uint32_t           engineVersion;
    uint32_t           apiVersion;
} VkApplicationInfo;
```

Again, we see the `sType` and `pNext` fields in this structure. `sType` should be set to `VK_STRUCTURE_TYPE_APPLICATION_INFO`, and we can leave `pNext` as `nullptr`. `pApplicationName` is a pointer to a `nul`-terminated[1] string containing the name of your application, and `applicationVersion` is the version of the application. This allows tools and drivers to make decisions about how to treat your application without needing to guess[2] which application is running. Likewise, `pEngineName` and `engineVersion` contain the name and version, respectively, of the engine or middleware that your application is based on.

1. Yes, really, `nul`. The ASCII character whose literal value is zero is officially called `NUL`. Now, stop telling me to change it to `NULL`. That's a pointer, not a name of a character.

2. What's best for one application might be different from what's best for another. Also, applications are written by humans, and humans write code with bugs. To optimize fully or work around application bugs, drivers would sometimes use executable names or even application behavior to guess at which application was running and alter behavior appropriately. While not ideal, this new mechanism at least removes the guesswork.

Finally, `apiVersion` contains the version of the Vulkan API that your application is expecting to run on. This should be set to the *absolute minimum* version of Vulkan that your application requires to run—not just to the version of the header that you happen to have installed. This allows the widest possible assortment of devices and platforms to run your application, even if updates to their Vulkan implementations might not be available.

Returning to the `VkInstanceCreateInfo` structure, we see the `enabledLayerCount` and `ppEnabledLayerNames` fields. These are the count of the number of *instance layers* that you wish to enable and their names, respectively. Layers are used to intercept the Vulkan API and provide logging, profiling, debugging, or other additional features. If no layers are needed, simply set `enabledLayerCount` to zero and leave `ppEnabledLayerNames` as `nullptr`. Likewise, `enabledExtensionCount` is the count of the number of extensions you wish to enable,[3] and `ppEnabledExtensionNames` is a list of their names. Again, if we're not using any extensions, we can set these fields to zero and `nullptr`, respectively.

3. As with OpenGL, Vulkan supports extensions as a central part of the API. However, in OpenGL, we would create a context, query the supported extensions, and then start using them. This meant that drivers would need to assume that your application might suddenly start using an extension at any time and be ready for it. Further, it couldn't tell *which* extensions you were looking for, which made the process even more difficult. In Vulkan, applications are required to *opt in* to extensions and explicitly enable them. This allows drivers to disable extensions that aren't in use and makes it harder for applications to accidentally start using functionality that's part of an extension they weren't intending to enable.

Finally, returning to the **`vkCreateInstance()`** function, the `pAllocator` parameter is a pointer to a host memory allocator that your application can supply in order to manage the host memory that the Vulkan system uses. Setting this to `nullptr` causes the Vulkan system to use its own internal allocator, which is what we will do here. Application-managed host memory will be covered in Chapter 2, "Memory and Resources."

Assuming the **`vkCreateInstance()`** function succeeds, it will return `VK_SUCCESS` and place a handle to the new instance in the variable pointed to by the `pInstance` parameter. A *handle* is the value by which objects are referenced. Vulkan handles are always 64 bits wide, regardless of the bitness of the host system. Once we have a handle to our Vulkan instance, we can use it to call other instance functions.

Vulkan Physical Devices

Once we have an instance, we can use it to discover Vulkan-compatible devices installed in the system. Vulkan has two types of devices: physical and logical. Physical devices are normally parts of the system—a graphics card, accelerator, DSP, or other component. There are a fixed number of physical devices in a system, and each has a fixed set of capabilities. A logical device is a software abstraction of a physical device, configured in a way that is specified by the application. The logical device is the one that your application will spend most of its time dealing with, but before we can create a logical device, we must discover the connected physical devices. To do this, we call the **vkEnumeratePhysicalDevices()** function, the prototype of which is

[**Click here to view code image**](#)

```
VkResult vkEnumeratePhysicalDevices (
    VkInstance                          instance,
    uint32_t*                           pPhysicalDeviceCount,
    VkPhysicalDevice*                   pPhysicalDevices);
```

The first parameter to the **vkEnumeratePhysicalDevices()** function, `instance`, is the instance we created earlier. Next, the `pPhysicalDeviceCount` parameter is a pointer to an unsigned integer variable that is both an input and an output. As an output, Vulkan writes the number of physical devices in the system into it. As an input, it should be preinitialized with the maximum number of devices your application can handle. The `pPhysicalDevices` parameter is a pointer to an array of this number of `VkPhysicalDevice` handles.

If you just want to know how many devices there are in the system, set `pPhysicalDevices` to `nullptr`, and Vulkan will ignore the initial value of `pPhysicalDeviceCount`, simply overwriting it with the number of supported devices. You can dynamically adjust the size of your `VkPhysicalDevice` array by calling **vkEnumeratePhysicalDevices()** twice, the first time with only `pPhysicalDevices` set to `nullptr` (although `pPhysicalDeviceCount` must still be a valid pointer) and the second time with `pPhysicalDevices` set to an array that has been appropriately sized for the number of physical devices reported by the first call.

Assuming there are no problems, **vkEnumeratePhysicalDevices()** returns `VK_SUCCESS` and deposits the number of recognized physical devices in `pPhysicalDeviceCount` along with their handles in `pPhysicalDevices`. [Listing 1.1](#) shows an example of constructing the `VkApplicationInfo` and `VkInstanceCreateInfo` structures, creating the Vulkan instance, querying it for the number of supported devices, and finally querying the physical device handles themselves. This is a slightly simplified version of `vkapp::init` from the example framework.

Listing 1.1: Creating a Vulkan Instance

[**Click here to view code image**](#)

```
VkResult vkapp::init()
{
    VkResult result = VK_SUCCESS;
    VkApplicationInfo appInfo = { };
    VkInstanceCreateInfo instanceCreateInfo = { };

    // A generic application info structure
```

```
        appInfo.sType = VK_STRUCTURE_TYPE_APPLICATION_INFO;
        appInfo.pApplicationName = "Application";
        appInfo.applicationVersion = 1;
        appInfo.apiVersion = VK_MAKE_VERSION(1, 0, 0);

        // Create the instance.
        instanceCreateInfo.sType = VK_STRUCTURE_TYPE_INSTANCE_CREATE_INFO;
        instanceCreateInfo.pApplicationInfo = &appInfo;

        result = vkCreateInstance(&instanceCreateInfo, nullptr, &m_instance);

        if (result == VK_SUCCESS)
        {
            // First figure out how many devices are in the system.
            uint32_t physicalDeviceCount = 0;
            vkEnumeratePhysicalDevices(m_instance, &physicalDeviceCount,
    nullptr);

            if (result == VK_SUCCESS)
            {
                // Size the device array appropriately and get the physical
                // device handles.
                m_physicalDevices.resize(physicalDeviceCount);
                vkEnumeratePhysicalDevices(m_instance,
                                           &physicalDeviceCount,
                                           &m_physicalDevices[0]);
            }
        }
        return result;
    }
```

The physical device handle is used to query the device for its capabilities and ultimately to create the logical device. The first query we'll perform is **vkGetPhysicalDeviceProperties()**, which fills in a structure describing all the properties of the physical device. Its prototype is

[Click here to view code image](#)

```
void vkGetPhysicalDeviceProperties (
    VkPhysicalDevice                    physicalDevice,
    VkPhysicalDeviceProperties*         pProperties);
```

When you call **vkGetPhysicalDeviceProperties()**, pass one of the handles returned from **vkEnumeratePhysicalDevices()** in the physicalDevice parameter, and in pProperties, pass a pointer to an instance of the VkPhysicalDeviceProperties structure. This is a large structure that contains a large number of fields describing the properties of the physical device. Its definition is

[Click here to view code image](#)

```
typedef struct VkPhysicalDeviceProperties {
    uint32_t                            apiVersion;
    uint32_t                            driverVersion;
    uint32_t                            vendorID;
    uint32_t                            deviceID;
```

```
    VkPhysicalDeviceType                    deviceType;
    char                                    deviceName
                                                [VK_MAX_PHYSICAL_DEVICE_NAME_SIZE];
    uint8_t                                 pipelineCacheUUID[VK_UUID_SIZE];
    VkPhysicalDeviceLimits                  limits;
    VkPhysicalDeviceSparseProperties        sparseProperties;
} VkPhysicalDeviceProperties;
```

The `apiVersion` field contains the highest version of Vulkan supported by the device, and the `driverVersion` field contains the version of the driver used to control the device. This is vendor-specific, so it doesn't make sense to compare driver versions across vendors. The `vendorID` and `deviceID` fields identify the vendor and the device, and are usually PCI vendor and device identifiers.[4]

4. There is no official central repository of PCI vendor or device identifiers. The PCI SIG (`http://pcisig.com/`) assigns vendor identifiers to its members, and those members assign device identifiers to their products. A fairly comprehensive list in both human- and machine-readable forms is available from `http://pcidatabase.com/`.

The `deviceName` field will contain a human-readable string naming the device. The `pipelineCacheUUID` field is used for pipeline caching, which we will cover in Chapter 6, "Shaders and Pipelines."

In addition to the properties just listed, the `VkPhysicalDeviceProperties` structure embeds `VkPhysicalDeviceLimits` and `VkPhysicalDeviceSparseProperties`, which contain the minimum and maximum limits for the physical device and properties related to sparse textures. There's a lot of information in these structures, and we'll cover the fields separately as the related features are discussed rather than enumerating them all here.

In addition to core features, some of which have optionally higher limits or bounds, Vulkan has a number of optional features that may be supported by a physical device. If a device advertises support for a feature, it must still be enabled (much like an extension), but once enabled, that feature becomes a first-class citizen of the API just like any core feature. To determine which features a physical device supports, call **`vkGetPhysicalDeviceFeatures()`**, the prototype of which is

Click here to view code image

```
void vkGetPhysicalDeviceFeatures (
    VkPhysicalDevice                physicalDevice,
    VkPhysicalDeviceFeatures*       pFeatures);
```

Again, the `VkPhysicalDeviceFeatures` structure is very large and has a Boolean field for each optional feature supported by Vulkan. There are too many fields to list and describe individually here, but the sample application presented at the end of this chapter reads the feature set and prints its content.

Physical Device Memory

In many cases, a Vulkan device is either a separate physical piece of hardware to the main host processor or works sufficiently differently that it will access memory in specialized ways. *Device memory* in Vulkan refers to memory that is accessible to the device and usable as a backing store for textures and other data. Memory is classified into *types*, each of which has a set of properties, such as caching flags and coherency behavior between host and device. Each type of memory is then backed by one of the device's *heaps*, of which there may be several.

To query the configuration of heaps and the memory types supported by the device, call

Click here to view code image

```
void vkGetPhysicalDeviceMemoryProperties (
    VkPhysicalDevice                        physicalDevice,
    VkPhysicalDeviceMemoryProperties*       pMemoryProperties);
```

The resulting memory organization is written into the `VkPhysicalDeviceMemoryProperties` structure, the address of which is passed in `pMemoryProperties`. The `VkPhysicalDeviceMemoryProperties` structure contains the properties of both the device's heaps and its supported memory types. The definition of this structure is

Click here to view code image

```
typedef struct VkPhysicalDeviceMemoryProperties {
    uint32_t        memoryTypeCount;
    VkMemoryType    memoryTypes[VK_MAX_MEMORY_TYPES];
    uint32_t        memoryHeapCount;
    VkMemoryHeap    memoryHeaps[VK_MAX_MEMORY_HEAPS];
} VkPhysicalDeviceMemoryProperties;
```

The number of memory types is reported in the `memoryTypeCount` field. The maximum number of memory types that might be reported is the value of `VK_MAX_MEMORY_TYPES`, which is defined to be 32. The `memoryTypes` array contains `memoryTypeCount` `VkMemoryType` structures describing each of the memory types. The definition of `VkMemoryType` is

Click here to view code image

```
typedef struct VkMemoryType {
    VkMemoryPropertyFlags    propertyFlags;
    uint32_t                 heapIndex;
} VkMemoryType;
```

This is a simple structure containing only a set of flags and the memory type's heap index. The `flags` field describes the type of memory and is made up of a combination of the `VkMemoryPropertyFlagBits` flags. The meanings of the flags are as follows:

- `VK_MEMORY_PROPERTY_DEVICE_LOCAL_BIT` means that the memory is local to (that is, physically connected to) the device. If this bit is not set, then the memory can be assumed to be local to the host.

- `VK_MEMORY_PROPERTY_HOST_VISIBLE_BIT` means that memory allocations made with this type can be mapped and read or written by the host. If this bit is not set then memory of this type cannot be directly accessed by the host and is rather for exclusive use by the device.

- `VK_MEMORY_PROPERTY_HOST_COHERENT_BIT` means that when this type of memory is concurrently accessed by both the host and device, those accesses will be coherent between the two clients. If this bit is not set, then the device or host may not see the results of writes performed by each until caches are explicitly flushed.

- `VK_MEMORY_PROPERTY_HOST_CACHED_BIT` means that data in this type of memory is cached by the host. Read accesses to this type of memory are typically faster than they would be if this bit were not set. However, access by the device may have slightly higher latency, especially if the memory is also coherent.

- `VK_MEMORY_PROPERTY_LAZILY_ALLOCATED_BIT` means that memory allocated with this type doesn't necessarily consume space from the associated heap immediately and that a driver might defer physical memory allocation until the memory object is used to back a resource.

Each memory type reports the heap from which it consumes space in the `heapIndex` field of the `VkMemoryType` structure. This is an index into the `memoryHeaps` array returned in the `VkPhysicalDeviceMemoryProperties` structure from the call to **`vkGetPhysicalDeviceMemoryProperties()`**. Each element of the `memoryHeaps` array describes one of the device's memory heaps. The definition of this structure is

[**Click here to view code image**](#)

```
typedef struct VkMemoryHeap {
    VkDeviceSize        size;
    VkMemoryHeapFlags   flags;
} VkMemoryHeap;
```

Again, this is a simple structure. It contains only the size, in bytes, of the heap and some flags describing the heap. In Vulkan 1.0, the only defined flag is `VK_MEMORY_HEAP_DEVICE_LOCAL_BIT`. If this bit is set, then the heap is local to the device. This corresponds to the similarly named flag describing memory types.

Device Queues

Vulkan devices execute work that is submitted to *queues*. Each device will have one or more queues, and each of those queues will belong to one of the device's queue families. A *queue family* is a group of queues that have identical capabilities but are able to run in parallel. The number of queue families, the capabilities of each family, and the number of queues belonging to each family are all properties of the physical device. To query the device for its queue families, call `vkGetPhysicalDeviceQueueFamilyProperties()`, the prototype of which is

[**Click here to view code image**](#)

```
void vkGetPhysicalDeviceQueueFamilyProperties (
    VkPhysicalDevice                        physicalDevice,
    uint32_t*                               pQueueFamilyPropertyCount,
    VkQueueFamilyProperties*                pQueueFamilyProperties);
```

`vkGetPhysicalDeviceQueueFamilyProperties()` works somewhat like **`vkEnumeratePhysicalDevices()`** in that it is expected that you call it twice. The first time, you pass `nullptr` as `pQueueFamilyProperties`, and in `pQueueFamilyPropertyCount`, you pass a pointer to a variable that will be overwritten with

the number of queue families supported by the device. You can use this number to appropriately size an array of `VkQueueFamilyProperties`. Then, on the second call, pass this array in `pQueueFamilyProperties`, and Vulkan will fill it with the properties of the queues. The definition of `VkQueueFamilyProperties` is

Click here to view code image

```
typedef struct VkQueueFamilyProperties {
    VkQueueFlags      queueFlags;
    uint32_t          queueCount;
    uint32_t          timestampValidBits;
    VkExtent3D        minImageTransferGranularity;
} VkQueueFamilyProperties;
```

The first field in this structure, `queueFlags`, describes the overall capabilities of the queue. This field is made up of a combination of the `VkQueueFlagBits` bits, the meanings of which are as follows:

- If `VK_QUEUE_GRAPHICS_BIT` is set, then queues in this family support graphics operations such as drawing points, lines, and triangles.
- If `VK_QUEUE_COMPUTE_BIT` is set, then queues in this family support compute operations such as dispatching compute shaders.
- If `VK_QUEUE_TRANSFER_BIT` is set, then queues in this family support transfer operations such as copying buffer and image contents.
- If `VK_QUEUE_SPARSE_BINDING_BIT` is set, then queues in this family support memory binding operations used to update sparse resources.

The `queueCount` field indicates the number of queues in the family. This might be set to 1, or it could be substantially higher if the device supports multiple queues with the same basic functionality.

The `timestampValidBits` field indicates how many bits are valid when timestamps are taken from the queue. If this value is zero, then the queue doesn't support timestamps. If it's nonzero, then it's guaranteed to be at least 36 bits. Furthermore, if the `timestampComputeAndGraphics` field of the device's `VkPhysicalDeviceLimits` structure is `VK_TRUE`, then all queues supporting either `VK_QUEUE_GRAPHICS_BIT` or `VK_QUEUE_COMPUTE_BIT` are guaranteed to support timestamps with at least 36 bits of resolution. In this case, there's no need to check each queue individually.

Finally, the `minImageTimestampGranularity` field specifies the units with which the queue supports image transfers (if at all).

Note that it might be the case that a device reports more than one queue family with apparently identical properties. Queues within a family are essentially identical. Queues in different families may have different internal capabilities that can't be expressed easily in the Vulkan API. For this reason, an implementation might choose to report similar queues as members of different families. This places additional restrictions on how resources are shared between those queues, which might allow the implementation to accommodate those differences.

Listing 1.2 illustrates how to query the physical device's memory properties and queue family properties. You will need to retrieve the queue family properties before creating the logical device, as discussed in the next section.

Listing 1.2: Querying Physical Device Properties

```
uint32_t queueFamilyPropertyCount;
std::vector<VkQueueFamilyProperties> queueFamilyProperties;
VkPhysicalDeviceMemoryProperties physicalDeviceMemoryProperties;

// Get the memory properties of the physical device.
vkGetPhysicalDeviceMemoryProperties(m_physicalDevices[deviceIndex],
                                    &physicalDeviceMemoryProperties);

// First determine the number of queue families supported by the physical
// device.
vkGetPhysicalDeviceQueueFamilyProperties(m_physicalDevices[0],
                                         &queueFamilyPropertyCount,
                                         nullptr);

// Allocate enough space for the queue property structures.
queueFamilyProperties.resize(queueFamilyPropertyCount);

// Now query the actual properties of the queue families.
vkGetPhysicalDeviceQueueFamilyProperties(m_physicalDevices[0],
                                         &queueFamilyPropertyCount,
                                         queueFamilyProperties.data());
```

Creating a Logical Device

After enumerating all of the physical devices in the system, your application should choose a device and create a *logical device* corresponding to it. The logical device represents the device in an initialized state. During creation of the logical device, you get to opt in to optional features, turn on extensions you need, and so on. Creating the logical device is performed by calling **vkCreateDevice()**, the prototype of which is

```
VkResult vkCreateDevice (
    VkPhysicalDevice                    physicalDevice,
    const VkDeviceCreateInfo*           pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkDevice*                           pDevice);
```

The physical device to which the new logical device corresponds is passed in physicalDevice. The information about the new logical device is passed in an instance of the VkDeviceCreateInfo structure through the pCreateInfo structure. The definition of VkDeviceCreateInfo is

```
typedef struct VkDeviceCreateInfo {
    VkStructureType                     sType;
    const void*                         pNext;
    VkDeviceCreateFlags                 flags;
    uint32_t                            queueCreateInfoCount;
```

```
    const VkDeviceQueueCreateInfo*    pQueueCreateInfos;
    uint32_t                          enabledLayerCount;
    const char* const*                ppEnabledLayerNames;
    uint32_t                          enabledExtensionCount;
    const char* const*                ppEnabledExtensionNames;
    const VkPhysicalDeviceFeatures*   pEnabledFeatures;
} VkDeviceCreateInfo;
```

The `sType` field of the `VkDeviceCreateInfo` structure should be set to
`VK_STRUCTURE_TYPE_DEVICE_CREATE_INFO`. As usual, unless you're using extensions,
`pNext` should be set to `nullptr`. In the current version of Vulkan, no bits are defined for the
`flags` field of the structure, so set this to zero too.

Next comes the queue creation information. `pQueueCreateInfos` is a pointer to an array of one
or more `VkDeviceQueueCreateInfo` structures, each of which allows the specification of one
or more queues. The number of structures in the array is given in `queueCreateInfoCount`. The
definition of `VkDeviceQueueCreateInfo` is

[Click here to view code image](#)

```
typedef struct VkDeviceQueueCreateInfo {
    VkStructureType            sType;
    const void*                pNext;
    VkDeviceQueueCreateFlags   flags;
    uint32_t                   queueFamilyIndex;
    uint32_t                   queueCount;
    const float*               pQueuePriorities;
} VkDeviceQueueCreateInfo;
```

The `sType` field for `VkDeviceQueueCreateInfo` is
`VK_STRUCTURE_TYPE_DEVICE_QUEUE_CREATE_INFO`. There are currently no flags defined
for use in the `flags` field, so it should be set to zero. The `queueFamilyIndex` field specifies the
family of the queues you want to create. This is an index into the array of queue families returned
from **vkGetPhysicalDeviceQueueFamilyProperties()**. To create queues in this family,
set `queueCount` to the number of queues you want to use. Of course, the device must support at
least this many queues in the family you choose.

The `pQueuePriorities` field is an optional pointer to an array of floating point values
representing the relative priority of work submitted to each of the queues. These numbers are
normalized numbers in the range of 0.0 to 1.0. Queues with higher priority may be allocated more
processing resources or scheduled more aggressively than queues with lower priority. Setting
`pQueuePriorities` to `nullptr` has the effect of leaving the queues at the same, default
priority.

The requested queues are sorted in order of priority and then assigned device-dependent relative
priorities. The number of discrete priorities that a queue may take on is a device-specific parameter.
You can determine this by checking the `discreteQueuePriorities` field of the
`VkPhysicalDeviceLimits` structure returned from a call to
vkGetPhysicalDeviceProperties(). For example, if a device supports only low- and high-
priority workloads, this field will be 2. All devices support at least two discrete priority levels.
However, if a device supports arbitrary priorities, then this field could be much higher. Regardless of

the value of `discreteQueuePriorities`, the relative priorities of the queue are still expressed as floating-point values.

Returning to the `VkDeviceCreateInfo` structure, the `enabledLayerCount`, `ppEnabledLayerNames`, `enabledExtensionCount`, and `ppEnabledExtensionNames` fields are for enabling layers and extensions. We will cover both of these topics later in this chapter. For now, we'll set both `enabledLayerCount` and `enabledExtensionCount` to zero and both `ppEnabledLayerNames` and `ppEnabedExtensionNames` to `nullptr`.

The final field of `VkDeviceCreateInfo`, `pEnabledFeatures`, is a pointer to an instance of the `VkPhysicalDeviceFeatures` structure that specifies which of the optional features that your application wishes to use. If you don't want to use any optional features, you can simply set this to `nullptr`. However, Vulkan in this form is relatively limited, and much of its interesting functionality will be disabled.

To determine which of the optional features the device supports, call **`vkGetPhysicalDeviceFeatures()`** as discussed earlier. **`vkGetPhysicalDeviceFeatures()`** writes the set of features supported by the device into an instance of the `VkPhysicalDeviceFeatures` structure that you pass in. By simply querying the phyiscal device's features and then passing the very same `VkPhysicalDeviceFeatures` structure back to **`vkCreateDevice()`**, you enable every optional feature that the device supports and do not request features that the device does not support.

Simply enabling every supported feature, however, may come with some performance impact. For some features, a Vulkan implementation may need to allocate extra memory, track additional state, configure hardware slightly differently, or perform some other operation that otherwise costs your application. It's not a good idea to enable features that won't be used. In an optimized application, you should query the supported features from the device; then, from the supported features, enable the specific features that your application requires.

[Listing 1.3](#) shows a simple example of querying the device for its supported features, setting up a list of features that the application requires. Support for tessellation and geometry shaders is absolutely required, and support for multidraw indirect is enabled if it is supported by the device. The code then creates a device using a single instance of its first queue.

Listing 1.3: Creating a Logical Device

Click here to view code image

```
VkResult result;
VkPhysicalDeviceFeatures supportedFeatures;
VkPhysicalDeviceFeatures requiredFeatures = {};

vkGetPhysicalDeviceFeatures(m_physicalDevices[0],
                            &supportedFeatures);

requiredFeatures.multiDrawIndirect      =
supportedFeatures.multiDrawIndirect;
requiredFeatures.tessellationShader     = VK_TRUE;
requiredFeatures.geometryShader         = VK_TRUE;

const VkDeviceQueueCreateInfo deviceQueueCreateInfo =
```

```
    {
        VK_STRUCTURE_TYPE_DEVICE_QUEUE_CREATE_INFO,      // sType
        nullptr,                                        // pNext
        0,                                              // flags
        0,                                              // queueFamilyIndex
        1,                                              // queueCount
        nullptr                                         // pQueuePriorities
    };
    const VkDeviceCreateInfo deviceCreateInfo =
    {
        VK_STRUCTURE_TYPE_DEVICE_CREATE_INFO,           // sType
        nullptr,                                        // pNext

        0,                                              // flags
        1,                                              //
    queueCreateInfoCount

        &deviceQueueCreateInfo,                         // pQueueCreateInfos
        0,                                              // enabledLayerCount
        nullptr,                                        // ppEnabledLayerNames
        0,                                              //
    enabledExtensionCount
        nullptr,                                        //
    ppEnabledExtensionNames
        &requiredFeatures                               // pEnabledFeatures
    };

    result = vkCreateDevice(m_physicalDevices[0],
                            &deviceCreateInfo,
                            nullptr,
                            &m_logicalDevice);
```

After the code in Listing 1.3 has run and successfully created the logical device, the set of enabled features is stored in the `requiredFeatures` variable. This may be kept for later so that code that can optionally use a feature can check whether it was successfully enabled and fall back gracefully.

Object Types and Function Conventions

Virtually everything in Vulkan is represented as an object that is referred to by a handle. Handles are divided into two broad categories: dispatchable objects and nondispatchable objects. For the most part, this is not relevant to applications and affects only how the API is structured and how system-level components such as the Vulkan loader and layers interoperate with those objects.

Dispatchable objects are objects that internally contain a dispatch table. This is the table of functions used by various components to determine which parts of code to execute when your application makes calls to Vulkan. These types of objects are generally heavier-weight constructs and currently consist of the instance (`VkInstance`), physical device (`VkPhysicalDevice`), logical device (`VkDevice`), command buffer (`VkCommandBuffer`), and queue (`VkQueue`). All other objects are considered nondispatchable.

The first argument to *any* Vulkan function is always a dispatchable object. The only exceptions to this rule are the functions related to creating and initializing the instance.

Managing Memory

Vulkan provides two types of memory: *host memory* and *device memory*. Objects created by the Vulkan API generally require some amount of host memory. This is where the Vulkan implementation will store the state of the object and any data it needs to implement the Vulkan API. Resource objects such as buffers and images require some amount of device memory. This is the memory where the data stored in the resource is kept.

It is possible for your application to manage host memory for the Vulkan implementation, and it is required that your application manage device memory. To do this, you will need to create a device memory management subsystem. Each resource that you create can be queried for the amount and type of memory it requires for it to be backed. It will be up to your application to allocate the correct amount of memory and attach it to the resource object before it can be used.

In higher-level APIs such as OpenGL, this "magic" is performed by drivers on behalf of your application. However, some applications require a very large number of small resources, and other applications require a smaller number of very large resources. Some applications create and destroy resources over the course of their execution, whereas other applications create all of their resources at startup and do not free them until they are terminated.

The allocation strategies used in these cases might be quite different. There is no one-size-fits-all strategy. An OpenGL driver has no idea how your application will behave and so must adapt allocation strategies to attempt to fit your usage patterns. On the other hand, *you*, the application developer, know exactly how your application will behave. You can partition resources into long-lived and transient groups. You can bucket resources that will be used together into a small number of pooled allocations. You are in the best position to decide the allocation strategies used by your application.

It is important to note that each "live" memory allocation places some cost on the system. Therefore, it is important to keep the number of allocation objects to a minimum. It is recommended that device memory allocators allocate memory in large chunks. Many small resources can be placed inside a much smaller number of device memory blocks. An example of a device memory allocator is discussed in Chapter 2, "Memory and Resources," which discusses memory allocation in much more detail.

Multithreading in Vulkan

Support for multithreaded applications is an integral part of the design of Vulkan. Vulkan generally assumes that the application will ensure that no two threads are mutating the same object at the same time. This is known as *external synchronization*. The vast majority of Vulkan commands in the performance-critical portions of Vulkan (such as building command buffers) provide no synchronization at all.

In order to concretely define the threading requirements of various Vulkan commands, each parameter that must be protected from concurrent access by the host is marked as externally synchronized. In some cases, handles to objects or other data are embdedd in data structures, included in arrays, or otherwise passed to commands through some indirect means. Those paramters must also be externally synchronized.

The intention of this is that a Vulkan implementation never needs to take a mutex or use other synchronization primitives internally to protect data structures. This means that multithreaded programs rarely stall or block across threads.

In addition to requiring the host to synchronize access to shared objects when they are used across threads, Vulkan includes a number of higher-level features that are designed specifically to allow threads to perform work without blocking one another. These include the following:

- Host memory allocations can be handled through a host memory allocation structure passed to object creation functions. By using an allocator per thread, the data structures in that allocator don't need to be protected. Host memory allocators are covered in Chapter 2, "Memory and Resources."

- Command buffers are allocated from pools, and access to the pool is externally synchronized. If an application uses a separate command pool per thread, command buffers can be allocated from those pools without blocking against one another. Command buffers and pools are covered in Chapter 3, "Queues and Commands."

- Descriptors are allocated in sets from descriptor pools. Descriptors are the representation of resources as used by shaders running on the device. They are covered in detail in Chapter 6, "Shaders and Pipelines." If a separate pool is used for each thread, then descriptor sets can be allocated from those pools without the threads blocking one another.

- Second-level command buffers allow the contents of a large renderpass (which must be contained in a single command buffer) to be generated in parallel and then grouped as they're called from the primary command buffer. Secondary command buffers are covered in detail in Chapter 13, "Multipass Rendering."

When you are building a very simple, single-threaded application, the requirement to create pools from which to allocate objects may seem like a lot of unnecessary indirection. However, as applications scale in number of threads, these objects are indispensible to achieving high performance.

Throughout the remainder of this book, any special requirements with respect to threading will be noted as the commands are introduced.

Mathematical Concepts

Computer graphics and most heterogeneous compute applications are fairly heavily math-based. Most Vulkan devices are based on extremely powerful computational processors. At the time of writing, even modest mobile processors are capable of providing many gigaflops of processing power, while higher-end desktop and workstation processors deliver many teraflops of number-crunching ability. As a consequence, really interesting applications will build on math-heavy shaders. Further, several fixed-function sections of Vulkan's processing pipeline are built upon mathematical concepts that are hard-wired into the device and specification.

Vectors and Matrices

One of the fundamental building blocks of any graphics application is the vector. Whether they're the representations of a position, a direction, a color, or some other quantity, vectors are used throughout the graphics literature. One common form of vector is the homogeneous vector, which is a vector in a space one dimension higher than the quantity it's representing. These vectors are used to store projective coordinates. Multiplying a homogeneous vector by any scalar produces a new vector representing the same projective coordinate. To project a point vector, divide through by its last component, producing a vector of the form *x, y, z,* 1.0 (for a four-component vector).

To transform a vector from one coordinate space to another, multiply the vector by a matrix. Just as a point in 3D space is represented as a four-component homogeneous vector, a transformation matrix operating on a 3D homogeneous vector is a 4 × 4 matrix.

A point in 3D space is typically represented as a homogeneous vector of four components conventionally called x, y, z, and w. For a point, the w component generally begins as 1.0 and changes as the vector is transformed through projective matrices. After division through by the w component, the point is projected through whichever transforms it's been subjected to. If none of the transforms is a projective transform, then w remains 1.0, and division by 1.0 has no effect on the vector. If the vector was subjected to a projective transform, then w will not be equal to 1.0, but dividing through by it will project the point and return w to 1.0.

Meanwhile, a direction in 3D space is also represented as a homogeneous vector whose w component is 0.0. Multiplying a direction vector by a properly constructed 4 × 4 projective matrix will leave the w component at 0.0, and it will have no effect on any of the other components. By simply discarding the additional component, you can put a 3D direction vector through the same transformations as a 4D homogeneous 3D point vector and make it undergo rotations, scales, and other transforms consistently.

Coordinate Systems

Vulkan deals with graphics primitives such as lines and triangles by representing their endpoints or corners as points in 3D space. These primitives are known as *vertices*. The inputs to the Vulkan system are vertex coordinates in a 3D coordinate space (represented as homogenous vectors with w components of 1.0) relative to the origin of the object of which they are part. This coordinate space is known as *object space* or sometimes *model space*.

Typically, the first shaders in the pipeline will transform this vertex into *view space*, which is a position relative to the viewer. This transformation is performed by multiplying the vertex's position vector by a transformation matrix. This is often called the *object-to-view matrix* or the *model-view matrix*.

Sometimes, absolute coordinates of a vertex are required, such as when finding the coordinate of a vertex relative to some other object. This global space is known as *world space* and is the position of vertices relative to a global origin.

From view space, the positions of vertices are transformed into *clip space*. This is the final space used by the geometry-processing part of Vulkan and is the space into which vertices are typically transformed when pushing them into the projective space used by typical 3D applications. Clip space is so called because it is the coordinate space in which most implementations perform *clipping*, which removes sections of primitives that lie outside the visible region being rendered.

From clip space, vertex positions are normalized by dividing through by their w components. This yields a coordinate space called *normalized device coordinates*, or NDC, and the process is often called the *perspective divide*. In this space, the visible part of the coordinate system is from −1.0 to 1.0 in the x and y directions and from 0.0 to 1.0 in the z direction. Anything outside this region will be clipped away prior to perspective division.

Finally, a vertex's normalized device coordinate is transformed by the viewport, which describes how NDC maps into a window or image into which the picture is being rendered.

Enhancing Vulkan

Although the core API specification of Vulkan is quite extensive, it's by no means all-encompassing. Some functionality is optional, while yet more is available in the form of layers (which modify or enhance existing behavior) and extensions (which add new functionality to Vulkan). Both enhancement mechanisms are described in the following sections.

Layers

Layers are features of Vulkan that allow its behavior to be modified. Layers generally intercept all or part of Vulkan and add functionality such as logging, tracing, providing diagnostics, profiling, and so on. A layer can be added at the instance level, in which case it affects the whole Vulkan instance and possibly every device created by it. Alternatively, the layer can be added at the device level, in which case it affects only the device for which it is enabled.

To discover the layers available to an instance on a system, call `vkEnumerateInstanceLayerProperties()`, the prototype of which is

[Click here to view code image](#)

```
VkResult vkEnumerateInstanceLayerProperties (
    uint32_t*                          pPropertyCount,
    VkLayerProperties*                 pProperties);
```

If `pProperties` is `nullptr`, then `pPropertyCount` should point to a variable that will be overwritten with the count of the number of layers available to Vulkan. If `pProperties` is not `nullptr`, then it should point to an array of `VkLayerProperties` structures that will be filled with information about the layers registered with the system. In this case, the initial value of the variable pointed to by `pPropertyCount` is the length of the array pointed to by `pProperties`, and this variable will be overwritten with the number of entries in the array overwritten by the command.

Each element of the `pProperties` array is an instance of the `VkLayerProperties` structure, the definition of which is

[Click here to view code image](#)

```
typedef struct VkLayerProperties {
    char        layerName[VK_MAX_EXTENSION_NAME_SIZE];
    uint32_t    specVersion;
    uint32_t    implementationVersion;
    char        description[VK_MAX_DESCRIPTION_SIZE];
} VkLayerProperties;
```

Each layer has a formal name that is stored in the `layerName` member of the `VkLayerProperties` structure. As the specification for each layer might be improved, clarified, or appended to over time, the version of the layer implementation is reported in `specVersion`.

As specifications are improved over time, so too are implementations of those specifications. The implementation version is stored in the `implementationVersion` field of the `VkLayerProperties` structure. This allows implementations to improve performance, fix bugs, implement a wider set of optional features, and so on. An application writer may recognize a

particular implementation of a layer and choose to use it only if the version of that implementation is past a certain version where, for example, a critical bug was known to be fixed.

Finally, a human-readable string describing the layer is stored in `description`. The only purpose of this field is for logging or display in a user interface, and it is for informational purposes only.

illustrates how to query the instance layers supported by the Vulkan system.

Listing 1.4: Querying Instance Layers

Click here to view code image

```
uint32_t numInstanceLayers = 0;
std::vector<VkLayerProperties> instanceLayerProperties;

// Query the instance layers.
vkEnumerateInstanceLayerProperties(&numInstanceExtensions,
                                   nullptr);

// If there are any layers, query their properties.
if (numInstanceLayers != 0)
{
    instanceLayerProperties.resize(numInstanceLayers);
    vkEnumerateInstanceLayerProperties(nullptr,
                                       &numInstanceLayers,
                                       instanceLayerProperties.data());
}
```

As mentioned, it is not only at the instance level that layers can be injected. Layers can also be applied at the device level. To determine which layers are available to devices, call **`vkEnumerateDeviceLayerProperties()`**, the prototype of which is

Click here to view code image

```
VkResult vkEnumerateDeviceLayerProperties (
    VkPhysicalDevice                        physicalDevice,
    uint32_t*                               pPropertyCount,
    VkLayerProperties*                      pProperties);
```

The layers available to each physical device in a system may be different, so each physical device can report a different set of layers. The physical device whose layers to query is passed in `physicalDevice`. The `pPropertyCount` and `pProperties` parameters to **`vkEnumerateDeviceLayerProperties()`** behave similarly to the identically named parameters to **`vkEnumerateInstanceLayerProperties()`**. Device layers are also described by instances of the `VkLayerProperties` structure.

To enable a layer at the instance level, include its name in the `ppEnabledLayerNames` field of the `VkInstanceCreateInfo` structure used to create the instance. Likewise, to enable a layer when creating a logical device corresponding to a physical device in the system, include the layer name in the `ppEnabledLayerNames` member of the `VkDeviceCreateInfo` used to create the device.

Several layers are included in the offical SDK, most of which are related to debugging, parameter validation, and logging. These include the following:

- `VK_LAYER_LUNARG_api_dump` prints Vulkan calls and their parameters and values to the console.
- `VK_LAYER_LUNARG_core_validation` performs validation on parameters and state used in descriptor sets, pipeline state, and dynamic state; validates the interfaces between SPIR-V modules and the graphics pipeline; and tracks and validates usage of GPU memory used to back objects.
- `VK_LAYER_LUNARG_device_limits` ensures that values passed to Vulkan commands as arguments or data structure members fall within the device's supported feature set limits.
- `VK_LAYER_LUNARG_image` validates that image usage is consistent with supported formats.
- `VK_LAYER_LUNARG_object_tracker` performs tracking on Vulkan objects, attempting to catch leaks, use-after-free errors, and other invalid object usage.
- `VK_LAYER_LUNARG_parameter_validation` confirms that all parameter values passed to Vulkan functions are valid.
- `VK_LAYER_LUNARG_swapchain` performs validation on functionality provided by the WSI (Window System Integration) extensions described in [Chapter 5](#), "[Presentation](#)."
- `VK_LAYER_GOOGLE_threading` ensures valid usage of Vulkan commands with respect to threading, ensuring that no two threads access the same object at the same time when they shouldn't.
- `VK_LAYER_GOOGLE_unique_objects` ensures that every object will have a unique handle for easier tracking by the application, avoiding cases where an implementation might de-duplicate handles that represent objects with the same parameters.

In addition to this, a large number of separate layers are grouped into a larger, single layer called `VK_LAYER_LUNARG_standard_validation`, making it easy to turn on. The book's application framework enables this layer when built in debug mode, leaving all layers disabled when built in release mode.

Extensions

Extensions are fundamental to any cross-platform, open API such as Vulkan. They allow implementers to experiment, innovate, and ultimately push technology forward. Eventually, useful features originally introduced as extensions make their way into future versions of the API after having been proved in the field. However, extensions are not without cost. Some may require implementations to track additional state, make additional checks during command buffer build, or come with some performance penalty even when the extension is not in direct use. Therefore, extensions must be explicitly enabled by an application before they can be used. This means that applications that don't use an extension don't pay for it in terms of performance or complexity and that it's almost impossible to accidentally use features from an extension, which improves portability.

Extensions are divided into two categories: instance extensions and device extensions. An *instance* extension is one that generally enhances the entire Vulkan system on a platform. This type of extension is either supplied via a device-independent layer or is simply an extension that is exposed by every device on the system and promoted into an instance. *Device* extensions are extensions that extend the capabilities of one or more devices in the sytem but aren't necessarily available on every device.

Each extension can define new functions, new types, structures, enumerations, and so on. Once enabled, an extension can be considered part of the API that is available to the application. Instance extensions must be enabled when the Vulkan instance is created, and device extensions must be enabled when the device is created. These leaves us with a chicken-and-egg situation: How do we know which extensions are supported before initializing a Vulkan instance?

Querying the supported instance extensions is one of the few pieces of Vulkan functionality that may be used before a Vulkan instance is created. This is performed using the **`vkEnumerateInstanceExtensionProperties()`** function, the prototype of which is

[**Click here to view code image**](#)

```
VkResult vkEnumerateInstanceExtensionProperties (
    const char*                         pLayerName,
    uint32_t*                           pPropertyCount,
    VkExtensionProperties*              pProperties);
```

`pLayerName` is the name of a layer that might provide extensions. For now, set this to `nullptr`. `pPropertyCount` is a pointer to a variable containing the count of the number of instance extensions to query Vulkan about, and `pProperties` is a pointer to an array of `VkExtensionProperties` structures that will be filled with information about the supported extensions. If `pProperties` is `nullptr`, then the initial value of the variable pointed to by `pPropertyCount` is ignored and is overwritten with the number of instance extensions supported.

If `pProperties` is not `nullptr`, then the number of entries in the array is assumed to be in the variable pointed to by `pPropertyCount`, and up to this many entries of the array are populated with information about the supported extensions. The variable pointed to by `pPropertyCount` is then overwritten with the number of entries actually written to `pProperties`.

To correctly query all of the supported instance extensions, call **`vkEnumerateInstanceExtensionProperties()`** twice. The first time, call it with pProperties set to nullptr to retrieve the number of supported instance extensions. Then appropriately size an array to receive the extension properties and call **`vkEnumerateInstanceExtensionProperties()`** again, this time passing the address of the array in pProperties. [Listing 1.5](#) demonstrates how to do this.

Listing 1.5: Querying Instance Extensions

[**Click here to view code image**](#)

```
uint32_t numInstanceExtensions = 0;
std::vector<VkExtensionProperties> instanceExtensionProperties;

// Query the instance extensions.
vkEnumerateInstanceExtensionProperties(nullptr,
                                       &numInstanceExtensions,
                                       nullptr);

// If there are any extensions, query their properties.
if (numInstanceExtensions != 0)
{
    instanceExtensionProperties.resize(numInstanceExtensions);
```

```
        vkEnumerateInstanceExtensionProperties(nullptr,
                                         &numInstanceExtensions,
                                         instanceExtensionProperties.data());
    }
```

After the code in [Listing 1.5](#) has completed execution, `instanceExtensionProperties` will contain a list of the extensions supported by the instance. Each element of the array of `VkExtensionProperties` describes one extension. The definition of `VkExtensionProperties` is

[**Click here to view code image**](#)

```
typedef struct VkExtensionProperties {
    char        extensionName[VK_MAX_EXTENSION_NAME_SIZE];
    uint32_t    specVersion;
} VkExtensionProperties;
```

The `VkExtensionProperties` structure simply contains the name of the extension and the version of that extension. Extensions may add functionality over time as new revisions of the extension are produced. The `specVersion` field allows updates to extensions without the need to create an entirely new extension in order to add minor functionality. The name of the extension is stored in `extensionName`.

As you saw earlier, when creating the Vulkan instance, the `VkInstanceCreateInfo` structure has a member called `ppEnabledExtensionNames`, which is a pointer to an array of strings naming the extensions to enable. If the Vulkan system on a platform supports an extension, that extension will be included in the array returned from **`vkEnumerateInstanceExtensionProperties()`**, and its name can be passed to **`vkCreateInstance()`** via the `ppEnabledExtensionNames` field of the `VkInstanceCreateInfo` structure.

Querying the support for device extensions is a similar process. To do this, call **`vkEnumerateDeviceExtensionProperties()`**, whose prototype is

[**Click here to view code image**](#)

```
VkResult vkEnumerateDeviceExtensionProperties (
    VkPhysicalDevice                          physicalDevice,
    const char*                               pLayerName,
    uint32_t*                                 pPropertyCount,
    VkExtensionProperties*                    pProperties);
```

The prototype of **`vkEnumerateDeviceExtensionProperties()`** is almost identical to that of **`vkEnumerateInstanceExtensionProperties()`**, with the addition of a `physicalDevice` parameter. The `physicalDevice` parameter is the handle to the device whose extensions to query. As with **`vkEnumerateInstanceExtensionProperties()`**, **`vkEnumerateDeviceExtensionProperties()`** overwrites `pPropertyCount` with the number of supported extensions if `pProperties` is `nullptr`, and if `pProprties` is not `nullptr`, it fills that array with information about the supported extensions. The same `VkExtensionProperties` structure is used for device extensions and for instance extensions.

When you are creating the physical device, the `ppEnabledExtensionNames` field of the `VkDeviceCreateInfo` structure may contain a pointer to one of the strings returned from **vkEnumerateDeviceExtensionProperties()**.

Some extensions provide new functionality in the form of additional entry points that you can call. These are exposed as function pointers, the values of which you must query either from the instance or from the device after enabling the extension. Instance functions are functions that are valid for the entire instance. If an extension extends instance-level functionality, you should use an instance-level function pointer to access the new features.

To retrieve an instance-level function pointer, call **vkGetInstanceProcAddr()**, the prototype of which is

[Click here to view code image](#)

```
PFN_vkVoidFunction vkGetInstanceProcAddr (
    VkInstance                              instance,
    const char*                             pName);
```

The `instance` parameter is the handle to the instance for which to retrieve a new function pointer. If your application does use more than one Vulkan instance, then the function pointer returned from this command is valid only for objects owned by the referenced instance. The name of the function is passed in `pName`, which is a `nul`-terminated UTF-8 string. If the function name is recognized and the extension is enabled, the return value from **vkGetInstanceProcAddr()** is a function pointer that you can call from your application.

The `PFN_vkVoidFunction` is a type definition for a pointer to a function of the following declaration:

[Click here to view code image](#)

```
VKAPI_ATTR void VKAPI_CALL vkVoidFunction(void);
```

There are no functions in Vulkan that have this particular signature, and it is unlikely that an extension would introduce such a function. In all likelihood, you will need to cast the resulting function pointer to a pointer of the appropriate signature before you can call it.

Instance-level function pointers are valid for any object owned by the instance, assuming the device that created the object (or the device itself, if the function dispatches on the device) supports the extension and the extension is enabled for that device. Because each device might be implemented inside a different Vulkan driver, instance function pointers must dispatch though a layer of indirection to land in the correct module. Managing this indirection may incur some overhead; to avoid this, you can get a device-specific function pointer that goes directly to the appropriate driver.

To get a device-level function pointer, call **vkGetDeviceProcAddr()**, the prototype of which is

[Click here to view code image](#)

```
PFN_vkVoidFunction vkGetDeviceProcAddr (
    VkDevice                                device,
    const char*                             pName);
```

The device with which the function pointer will be used is passed in `device`. Again, the name of the function you are querying is passed as a `nul`-terminated UTF-8 string in `pName`. The resulting function pointer is valid *only* with the device specified in `device`. `device` must refer to a device

that supports the extension that provides the new function and for which the extension has been enabled.

The function pointer produced by **vkGetDeviceProcAddr()** is specific to device. Even if the same physical device is used to create two or more logical devices with the exact same parameters, you must use the resulting function pointers only with the devices from which they were queried.

Shutting Down Cleanly

Before your program exits, you should clean up after yourself. In many cases, an operating system will clean up resources that you've allocated when your application terminates. However, it will not always be the case that the end of your code is the end of the application. If you are writing a component of a larger application, for example, that application may terminate rendering or compute operations using Vulkan without actually exiting.

When cleaning up, it is generally good practice to do the following:

- Complete or otherwise terminate all work that your application is doing both on the host and the device, in all threads related to Vulkan.

- Destroy objects in the reverse order from the order in which they were created.

The logical device is likely to be the last object (aside from objects used at runtime) that you created during initialization of your application. Before destroying the device, you should ensure that it is not executing any work on behalf of your application. To do this, call **vkDeviceWaitIdle()**, the prototype of which is

[Click here to view code image](#)

```
VkResult vkDeviceWaitIdle (
    VkDevice                            device);
```

The handle to the device is passed in device. When **vkDeviceWaitIdle()** returns, all work submitted to the device on behalf of your application is guaranteed to have completed—unless, of course, you submit more work to the device in the meantime. You should ensure that any other threads that might be submitting work to the device have been terminated.

Once you have ensured that the device is idle, you can safely destroy it. To do this, call **vkDestroyDevice()**, the prototype of which is

[Click here to view code image](#)

```
void vkDestroyDevice (
    VkDevice                            device,
    const VkAllocationCallbacks*        pAllocator);
```

The handle to the device to destroy is passed in the device parameter, access to which must be externally synchronized. Note that access to a device does *not* need to be externally synchronized with respect to any other command. The application should ensure, however, that a device is not destroyed while any other command accessing it is still executing in another thread.

pAllocator should point to an allocation structure that is compatible with the one used to create the device. Once the device object has been destroyed, no more commands can be submitted to it. Further, it is no longer possible to use the device handle as an argument to any function, including

other object-destruction functions that take a device handle as their first argument. This is another reason why you should destroy objects in reverse order from the order in which they were created.

Once all devices associated with a Vulkan instance have been destroyed, it is safe to destroy the instance. This is accomplished by calling the **vkDestroyInstance()** function, whose prototype is

[**Click here to view code image**](#)

```
void vkDestroyInstance (
    VkInstance                          instance,
    const VkAllocationCallbacks*        pAllocator);
```

The handle to the instance to destroy is passed in instance and, as with **vkDestroyDevice()**, a pointer to an allocation structure compatible with the one with which the instance was allocated should be passed in pAllocator. If the pAllocator parameter to **vkCreateInstance()** was nullptr, then the pAllocator parameter to **vkDestroyInstance()** should be too.

Note that it's not necessary to destroy the physical devices. Physical devices are not created with a dedicated creation function as logical devices are. Rather, physical devices are returned from a call to **vkEnumeratePhysicalDevices()** and are considered to be owned by the instance. Therefore, when the instance is destroyed, that instance's resources associated with each physical device are freed as well.

Summary

This chapter introduced you to Vulkan. You have seen how the entirety of Vulkan state is contained within an instance. The instance provides access to physical devices, and each physical device exposes a number of queues that can be used to do work. You have seen how to create a logical device corresponding to the physical device. You have seen how Vulkan can be extended, how to determine the extensions available to the instance and to a device, and how to enable those extensions. You have seen how to cleanly shut down the Vulkan system by waiting for a device to complete work submitted by your application, destroying the device handles, and finally destroying the instance handle.

Chapter 2. Memory and Resources

What You'll Learn in This Chapter

- • How Vulkan manages host and device memory
- • How to manage memory effectively in your application
- • How Vulkan uses images and buffers to consume and produce data

---

Memory is fundamental to the operation of virtually all computing systems, including Vulkan. In Vulkan, there are two fundamental types of memory: host memory and device memory. All resources upon which Vulkan operates must be backed by device memory, and it is the application's responsibility to manage this memory. Further, memory is used to store data structures on the host. Vulkan provides the opportunity for your application to manage this memory too. In this chapter, you'll learn about the mechanisms through which you can manage memory used by Vulkan.

Host Memory Management

Whenever Vulkan creates new objects, it might need memory to store data related to them. For this, it uses *host memory*, which is regular memory accessible to the CPU that might be returned from a call to `malloc` or **new**, for example. However, beyond a normal allocator, Vulkan has particular requirements for some allocations. Most notably, it expects allocations to be aligned correctly. This is because some high-performance CPU instructions work best (or only) on aligned memory addresses. By assuming that allocations storing CPU-side data structures are aligned, Vulkan can use these high-performance instructions unconditionally, providing substantial performance advantages.

Because of these requirements, Vulkan implementations will use advanced allocators to satisfy them. However, it also provides the opportunity for your application to replace the allocators for certain, or even all, operations. This is performed through the `pAllocator` parameter available in most device creation functions. For example, let's revisit the **vkCreateInstance()** function, which is one of the first that your application might call. Its prototype is

[**Click here to view code image**](#)

```
VkResult vkCreateInstance (
    const VkInstanceCreateInfo*      pCreateInfo,
    const VkAllocationCallbacks*     pAllocator,
    VkInstance*                      pInstance);
```

The `pAllocator` parameter is a pointer to a `VkAllocationCallbacks` structure. Until now, we've been setting `pAllocator` to `nullptr`, which tells Vulkan to use its own internal allocator rather than rely on our application. The `VkAllocationCallbacks` structure encapsulates a custom memory allocator that we can provide. The definition of the structure is

[**Click here to view code image**](#)

```
typedef struct VkAllocationCallbacks {
    void*                            pUserData;
    PFN_vkAllocationFunction         pfnAllocation;
```

```
        PFN_vkReallocationFunction                    pfnReallocation;
        PFN_vkFreeFunction                            pfnFree;
        PFN_vkInternalAllocationNotification          pfnInternalAllocation;
        PFN_vkInternalFreeNotification                pfnInternalFree;
    } VkAllocationCallbacks;
```

You can see from the definition of `VkAllocationCallbacks` that the structure is essentially a set of function pointers and an additional void pointer, `pUserData`. The pointer is for your application's use. It can point anywhere; Vulkan will not dereference it. In fact, it doesn't even need to be a pointer. You can put anything in there, so long as it fits into a pointer-size blob. The only thing that Vulkan will do with `pUserData` is pass it back to the callback functions to which the remaining members of `VkAllocationCallbacks` point.

`pfnAllocation`, `pfnReallocation`, and `pfnFree` are used for normal, object-level memory management. They are defined as pointers to functions that match the following declarations:

[Click here to view code image](#)

```
void* VKAPI_CALL Allocation(
    void*                                 pUserData,
    size_t                                size,
    size_t                                alignment,
    VkSystemAllocationScope               allocationScope);

void* VKAPI_CALL Reallocation(
    void*                                 pUserData,
    void*                                 pOriginal
    size_t                                size,
    size_t                                alignment,
    VkSystemAllocationScope               allocationScope);

void VKAPI_CALL Free(
    void*                                 pUserData,
    void*                                 pMemory);
```

Notice that all three functions take a `pUserData` parameter as their first argument. This is the same `pUserData` pointer that's part of the `VkAllocationCallbacks` structure. If your application uses data structures to manage memory, this is a good place to put their addresses. One logical thing to do with this is to implement your memory allocator as a C++ class (assuming you're writing in C++) and then put the class's **this** pointer in `pUserData`.

The `Allocation` function is responsible for making new allocations. The `size` parameter gives the size of the allocation, in bytes. The `alignment` parameter gives the required alignment of the allocation, also in bytes. This is an often-overlooked parameter. It is very tempting to simply hook this function up to a naïve allocator such as `malloc`. If you do this, you will find that it works for a while but that certain functions might mysteriously crash later. *If you provide your own allocator, it must honor the `alignment` parameter.*

The final parameter, `allocationScope`, tells your application what the scope, or lifetime, of the allocation is going to be. It is one of the `VkSystemAllocationScope` values, which have the following meanings:

- `VK_SYSTEM_ALLOCATION_SCOPE_COMMAND` means that the allocation will be live only for the duration of the command that provoked the allocation. Vulkan will likely use this for very short-lived temporary allocations, as it works on a single command.

- `VK_SYSTEM_ALLOCATION_SCOPE_OBJECT` means that the allocation is directly associated with a particular Vulkan object. This allocation will live at least until the object is destroyed. This type of allocation will only ever be made as part of executing a creation command (one beginning with `vkCreate`).

- `VK_SYSTEM_ALLOCATION_SCOPE_CACHE` means that the allocation is associated with some form of internal cache or a `VkPipelineCache` object.

- `VK_SYSTEM_ALLOCATION_SCOPE_DEVICE` means that the allocation is scoped to the device. This type of allocation is made when the Vulkan implementation needs memory associated with the device that is not tied to a single object. For example, if the implementation allocates objects in blocks, this type of allocation might be made in response to a request to create a new object, but because many objects might live in the same block, the allocation can't be tied directly to any specific object.

- `VK_SYSTEM_ALLOCATION_SCOPE_INSTANCE` means that the allocation is scoped to the instance. This is similar to `VK_SYSTEM_ALLOCATION_SCOPE_DEVICE`. This type of allocation is typically made by layers or during early parts of Vulkan startup, such as by **`vkCreateInstance()`** and **`vkEnumeratePhysicalDevices()`**.

The `pfnInternalAllocation` and `pfnInternalFree` function pointers point to alternate allocator functions that are used when Vulkan makes memory allocations using its own allocators. These callbacks have the same signatures as `pfnAllocation` and `pfnFree`, except that `pfnInternalAllocation` doesn't return a value and `pfnInternalFree` shouldn't actually free the memory. These functions are used only for notification so that your application can keep track of how much memory Vulkan is using. The prototypes of these functions should be

[**Click here to view code image**](#)

```
void VKAPI_CALL InternalAllocationNotification(
    void*                                   pUserData,
    size_t                                  size,
    VkInternalAllocationType                allocationType,
    VkSystemAllocationScope                 allocationScope);

void VKAPI_CALL InternalFreeNotification(
    void*                                   pUserData,
    size_t                                  size,
    VkInternalAllocationType                allocationType,
    VkSystemAllocationScope                 allocationScope);
```

There's not much you can do with the information provided through `pfnInternalAllocation` and `pfnInternalFree` besides log it and keep track of the total memory usage made by the application. Specifying these function pointers is optional, but if you supply one, you must supply both. If you don't want to use them, set them both to `nullptr`.

[Listing 2.1](#) shows an example of how to declare a C++ class that can be used as an allocator that maps the Vulkan allocation callback functions. Because the callback functions used by Vulkan are naked C function pointers, the callback functions themselves are declared as static member functions

of the class, whereas the actual implementations of those functions are declared as regular nonstatic member functions.

Listing 2.1: Declaration of a Memory Allocator Class

```
class allocator
{
public:
    // Operator that allows an instance of this class to be used as a
    // VkAllocationCallbacks structure
    inline operator VkAllocationCallbacks() const
    {
        VkAllocationCallbacks result;

        result.pUserData = (void*)this;
        result.pfnAllocation = &Allocation;
        result.pfnReallocation = &Reallocation;
        result.pfnFree = &Free;
        result.pfnInternalAllocation = nullptr;
        result.pfnInternalFree = nullptr;

        return result;
    };

private:
    // Declare the allocator callbacks as static member functions.
    static void* VKAPI_CALL Allocation(
        void*                               pUserData,
        size_t                              size,
        size_t                              alignment,
        VkSystemAllocationScope             allocationScope);

    static void* VKAPI_CALL Reallocation(
        void*                               pUserData,
        void*                               pOriginal,
        size_t                              size,
        size_t                              alignment,
        VkSystemAllocationScope             allocationScope);

    static void VKAPI_CALL Free(
        void*                               pUserData,
        void*                               pMemory);

    // Now declare the nonstatic member functions that will actually perform
    // the allocations.
    void* Allocation(

        size_t                                      size,
        size_t                                      alignment,
        VkSystemAllocationScope                     allocationScope);
```

```
    void* Reallocation(
        void*                                    pOriginal,
        size_t                                   size,
        size_t                                   alignment,
        VkSystemAllocationScope                  allocationScope);

    void Free(
        void*                                    pMemory);
};
```

An example implementation of this class is shown in . It maps the Vulkan allocation functions to the POSIX `aligned_malloc` functions. Note that this allocator is almost certainly *not* better than what most Vulkan implementations use internally and serves only as an example of how to hook the callback functions up to your own code.

Listing 2.2: Implementation of a Memory Allocator Class

Click here to view code image

```
void* allocator::Allocation(
    size_t                                   size,
    size_t                                   alignment,
    VkSystemAllocationScope                  allocationScope)
{
    return aligned_malloc(size, alignment);
}

void* VKAPI_CALL allocator::Allocation(
    void*                                    pUserData,
    size_t                                   size,
    size_t                                   alignment,
    VkSystemAllocationScope                  allocationScope)
{
    return static_cast<allocator*>(pUserData)->Allocation(size,
                                        alignment,
                                        allocationScope);
}

void* allocator::Reallocation(
    void*                                    pOriginal,
    size_t                                   size,
    size_t                                   alignment,
    VkSystemAllocationScope                  allocationScope)
{
    return aligned_realloc(pOriginal, size, alignment);
}
void* VKAPI_CALL allocator::Reallocation(
    void*                                    pUserData,
    void*                                    pOriginal,
    size_t                                   size,
    size_t                                   alignment,
    VkSystemAllocationScope                  allocationScope)
```

```
    {
        return static_cast<allocator*>(pUserData)->Reallocation(pOriginal,
                                                                size,
                                                                alignment,
                                                                allocationScope);
    }
    void allocator::Free(
        void*                                   pMemory)
    {
        aligned_free(pMemory);
    }
    void VKAPI_CALL allocator::Free(
        void*                                   pUserData,
        void*                                   pMemory)
    {
        return static_cast<allocator*>(pUserData)->Free(pMemory);
    }
```

As can be seen in Listing 2.2, the static member functions simply cast the pUserData parameters back to a class instance and call the corresponding nonstatic member function. Because the nonstatic and static member functions are located in the same compilation unit, the nonstatic member function is likely to be inlined into the static one, making the efficiency of this implementation quite high.

Resources

Vulkan operates on data. Everything else is really secondary to this. Data is stored in resources, and resources are backed by memory. There are two fundamental types of resources in Vulkan: buffers and images. A buffer is a simple, linear chunk of data that can be used for almost anything—data structures, raw arrays, and even image data, should you choose to use them that way. Images, on the other hand, are structured and have type and format information, can be multidimensional, form arrays of their own, and support advanced operations for reading and writing data from and to them.

Both types of resources are constructed in two steps: first the resource itself is created, and then the resource needs to be backed by memory. The reason for this is to allow the application to manage memory itself. Memory management is complex, and it is very difficult for a driver to get it right all the time. What works well for one application might not work well for another. Therefore, it is expected that applications can do a better job of managing memory than drivers can. For example, an application that uses a small number of very large resources and keeps them around for a long time might use one strategy in its memory allocator, while another application that continually creates and destroys small resources might implement another.

Although images are more complex structures, the procedure for creating them is similar to buffers. This section looks at buffer creation first and then moves on to discuss images.

Buffers

Buffers are the simplest type of resource but have a wide variety of uses in Vulkan. They are used to store linear structured or unstructured data, which can have a format or be raw bytes in memory. The various uses for buffer objects will be discussed as we introduce those topics. To create a new buffer object, call **vkCreateBuffer()**, the prototype of which is

```
VkResult vkCreateBuffer (
    VkDevice                                device,
    const VkBufferCreateInfo*               pCreateInfo,
    const VkAllocationCallbacks*            pAllocator,
    VkBuffer*                               pBuffer);
```

As with most functions in Vulkan that consume more than a couple of parameters, those parameters are bundled up in a structure and passed to Vulkan via a pointer. Here, the `pCreateInfo` parameter is a pointer to an instance of the `VkBufferCreateInfo` structure, the definition of which is

```
typedef struct VkBufferCreateInfo {
    VkStructureType        sType;
    const void*            pNext;
    VkBufferCreateFlags    flags;
    VkDeviceSize           size;
    VkBufferUsageFlags     usage;
    VkSharingMode          sharingMode;
    uint32_t               queueFamilyIndexCount;
    const uint32_t*        pQueueFamilyIndices;
} VkBufferCreateInfo;
```

The `sType` for `VkBufferCreateInfo` should be set to `VK_STRUCTURE_TYPE_BUFFER_CREATE_INFO`, and the `pNext` member should be set to `nullptr` unless you're using an extension. The `flags` field of the structure gives Vulkan some information about the properties of the new buffer. In the current version of Vulkan, the only bits defined for use in the `flags` field are related to *sparse buffers*, which we will cover later in this chapter. For now, `flags` can be set to zero.

The `size` field of `VkBufferCreateInfo` specifies the size of the buffer, in bytes. The `usage` field tells Vulkan how you're going to use the buffer and is a bitfield made up of a combination of members of the `VkBufferUsageFlagBits` enumeration. On some architectures, the intended usage of the buffer can have an effect on how it's created. The currently defined bits along with the sections where we'll discuss them are as follows:

- `VK_BUFFER_USAGE_TRANSFER_SRC_BIT` and `VK_BUFFER_USAGE_TRANSFER_DST_BIT` mean that the buffer can be the source or destination, respectively, of transfer commands. Transfer operations are operations that copy data from a source to a destination. They are covered in Chapter 4, "Moving Data."

- `VK_BUFFER_USAGE_UNIFORM_TEXEL_BUFFER_BIT` and `VK_BUFFER_USAGE_STORAGE_TEXEL_BUFFER_BIT` mean that the buffer can be used to back a uniform or storage texel buffer, respectively. Texel buffers are formatted arrays of texels that can be used as the source or destination (in the case of storage buffers) of reads and writes by shaders running on the device. Texel buffers are covered in Chapter 6, "Shaders and Pipelines."

- `VK_BUFFER_USAGE_UNIFORM_BUFFER_BIT` and `VK_BUFFER_USAGE_STORAGE_BUFFER_BIT` mean that the buffer can be used to back uniform or storage buffers, respectively. As opposed to texel buffers, regular uniform and

storage buffers have no format associated with them and can therefore be used to store arbitrary data and data structures. They are covered in Chapter 6, "Shaders and Pipelines."

- `VK_BUFFER_USAGE_INDEX_BUFFER_BIT` and `VK_BUFFER_USAGE_VERTEX_BUFFER_BIT` mean that the buffer can be used to store index or vertex data, respectively, used in drawing commands. You'll learn more about drawing commands, including indexed drawing commands, in Chapter 8, "Drawing."

- `VK_BUFFER_USAGE_INDIRECT_BUFFER_BIT` means that the buffer can be used to store parameters used in indirect dispatch and drawing commands, which are commands that take their parameters directly from buffers rather than from your program. These are covered in Chapter 6, "Shaders and Pipelines," and Chapter 8, "Drawing."

The `sharingMode` field of `VkBufferCreateInfo` indicates how the buffer will be used on the multiple command queues supported by the device. Because Vulkan can execute many operations in parallel, some implementations need to know whether the buffer will essentially be used by a single command at a time or potentially by many. Setting `sharingMode` to `VK_SHARING_MODE_EXCLUSIVE` says that the buffer will only be used on a single queue, whereas setting `sharingMode` to `VK_SHARING_MODE_CONCURRENT` indicates that you plan to use the buffer on multiple queues at the same time. Using `VK_SHARING_MODE_CONCURRENT` might result in lower performance on some systems, so unless you need this, set `sharingMode` to `VK_SHARING_MODE_EXCLUSIVE`.

If you do set `sharingMode` to `VK_SHARING_MODE_CONCURRENT`, you need to tell Vulkan which queues you're going to use the buffer on. This is done using the `pQueueFamilyIndices` member of `VkBufferCreateInfo`, which is a pointer to an array of queue families that the resource will be used on. `queueFamilyIndexCount` contains the length of this array—the number of queue families that the buffer will be used with. When `sharingMode` is set to `VK_SHARING_MODE_EXCLUSIVE`, `queueFamilyCount` and `pQueueFamilies` are both ignored.

Listing 2.3 demonstrates how to create a buffer object that is 1MiB in size, usable as the source or destination of transfer operations, and used on only one queue family at a time.

Listing 2.3: Creating a Buffer Object

Click here to view code image

```
static const VkBufferCreateInfo bufferCreateInfo =
{
    VK_STRUCTURE_TYPE_BUFFER_CREATE_INFO, nullptr,
    0,
    1024 * 1024,
    VK_BUFFER_USAGE_TRANSFER_SRC_BIT | VK_BUFFER_USAGE_TRANSFER_DST_BIT,
    VK_SHARING_MODE_EXCLUSIVE,
    0, nullptr
};

VkBuffer buffer = VK_NULL_HANDLE;

vkCreateBuffer(device, &bufferCreateInfo, &buffer);
```

After the code in [Listing 2.3](#) has run, a new `VkBuffer` handle is created and placed in the `buffer` variable. The buffer is not yet fully usable because it first needs to be backed with memory. This operation is covered in "[Device Memory Management](#)" later in this chapter.

Formats and Support

While buffers are relatively simple resources and do not have any notion of the format of the data they contain, images and buffer views (which we will introduce shortly) do include information about their content. Part of that information describes the format of the data in the resource. Some formats have special requirements or restrictions on their use in certain parts of the pipeline. For example, some formats might be readable but not writable, which is common with compressed formats.

In order to determine the properties and level of support for various formats, you can call **`vkGetPhysicalDeviceFormatProperties()`**, the prototype of which is

[**Click here to view code image**](#)

```
void vkGetPhysicalDeviceFormatProperties (
    VkPhysicalDevice                        physicalDevice,
    VkFormat                                format,
    VkFormatProperties*                     pFormatProperties);
```

Because support for particular formats is a property of a physical device rather than a logical one, the physical device handle is specified in `physicalDevice`. If your application absolutely required support for a particular format or set of formats, you could check for support before even creating the logical device and reject particular physical devices from consideration early in application startup, for example. The format for which to check support is specified in `format`. If the device recognizes the format, it will write its level of support into the instance of the `VkFormatProperties` structure pointed to by `pFormatProperties`. The definition of the `VkFormatProperties` structure is

[**Click here to view code image**](#)

```
typedef struct VkFormatProperties {
    VkFormatFeatureFlags    linearTilingFeatures;
    VkFormatFeatureFlags    optimalTilingFeatures;
    VkFormatFeatureFlags    bufferFeatures;
} VkFormatProperties;
```

All three fields in the `VkFormatProperties` structure are bitfields made up from members of the `VkFormatFeatureFlagBits` enumeration. An image can be in one of two primary tiling modes: linear, in which image data is laid out linearly in memory, first by row, then by column, and so on; and optimal, in which image data is laid out in highly optimized patterns that make efficient use of the device's memory subsystem. The `linearTilingFeatures` field indicates the level of support for a format in images in linear tiling, the `optimalTilingFeatures` field indicates the level of support for a format in images in optimal tiling, and the `bufferFeatures` field indicates the level of support for the format when used in a buffer.

The various bits that might be included in these fields are defined as follows:

- `VK_FORMAT_FEATURE_SAMPLED_IMAGE_BIT`: The format may be used in read-only images that will be sampled by shaders.

- `VK_FORMAT_FEATURE_SAMPLED_IMAGE_FILTER_LINEAR_BIT`: Filter modes that include linear filtering may be used when this format is used for a sampled image.
- `VK_FORMAT_FEATURE_STORAGE_IMAGE_BIT`: The format may be used in read-write images that will be read and written by shaders.
- `VK_FORMAT_FEATURE_STORAGE_IMAGE_ATOMIC_BIT`: The format may be used in read-write images that also support atomic operations performed by shaders.
- `VK_FORMAT_FEATURE_UNIFORM_TEXEL_BUFFER_BIT`: The format may be used in a read-only texel buffer that will be read from by shaders.
- `VK_FORMAT_FEATURE_STORAGE_TEXEL_BUFFER_BIT`: The format may be used in read-write texel buffers that may be read from and written to by shaders.
- `VK_FORMAT_FEATURE_STORAGE_TEXEL_BUFFER_ATOMIC_BIT`: The format may be used in read-write texel buffers that also support atomic operations performed by shaders.
- `VK_FORMAT_FEATURE_VERTEX_BUFFER_BIT`: The format may be used as the source of vertex data by the vertex-assembly stage of the graphics pipeline.
- `VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BIT`: The format may be used as a color attachment in the color-blend stage of the graphics pipeline.
- `VK_FORMAT_FEATURE_COLOR_ATTACHMENT_BLEND_BIT`: Images with this format may be used as color attachments when blending is enabled.
- `VK_FORMAT_FEATURE_DEPTH_STENCIL_ATTACHMENT_BIT`: The format may be used as a depth, stencil, or depth-stencil attachment.
- `VK_FORMAT_FEATURE_BLIT_SRC_BIT`: The format may be used as the source of data in an image copy operation.
- `VK_FORMAT_FEATURE_BLIT_DST_BIT`: The format may be used as the destination of an image copy operation.

Many formats will have a number of format support bits turned on. In fact, many formats are compulsory to support. A complete list of the mandatory formats is contained in the Vulkan specification. If a format is on the mandatory list, then it's not strictly necessary to test for support. However, for completeness, implementations are expected to accurately report capabilities for all supported formats, even mandatory ones.

The **`vkGetPhysicalDeviceFormatProperties()`** function really returns only a coarse set of flags indicating whether a format may be used at all under particular scenarios. For images especially, there may be more complex interactions between a specific format and its effect on the level of support within an image. Therefore, to retrieve even more information about the support for a format when used in images, you can call **`vkGetPhysicalDeviceImageFormatProperties()`**, the prototype of which is

[Click here to view code image](#)

```
VkResult vkGetPhysicalDeviceImageFormatProperties (
    VkPhysicalDevice                        physicalDevice,
    VkFormat                                format,
    VkImageType                             type,
    VkImageTiling                           tiling,
    VkImageUsageFlags                       usage,
```

```
    VkImageCreateFlags                      flags,
    VkImageFormatProperties*                pImageFormatProperties);
```

Like **vkGetPhysicalDeviceFormatProperties()**,
vkGetPhysicalDeviceImageFormatProperties() takes a VkPhysicalDevice
handle as its first parameter and reports support for the format for the physical device rather than for
a logical one. The format you're querying support for is passed in format.

The type of image that you want to ask about is specified in type. This should be one of the image
types: VK_IMAGE_TYPE_1D, VK_IMAGE_TYPE_2D, or VK_IMAGE_TYPE_3D. Different image
types might have different restrictions or enhancements. The tiling mode for the image is specified in
tiling and can be either VK_IMAGE_TILING_LINEAR or VK_IMAGE_TILING_OPTIMAL,
indicating linear or optimal tiling, respectively.

The intended use for the image is specified in the usage parameter. This is a bitfield indicating how
the image is to be used. The various uses for an image are discussed later in this chapter. The flags
field should be set to the same value that will be used when creating the image that will use the
format.

If the format is recognized and supported by the Vulkan implementation, then it will write
information about the level of support into the VkImageFormatProperties structure pointed to
by pImageFormatProperties. The definition of VkImageFormatProperties is

[Click here to view code image](#)

```
typedef struct VkImageFormatProperties {
    VkExtent3D             maxExtent;
    uint32_t               maxMipLevels;
    uint32_t               maxArrayLayers;
    VkSampleCountFlags     sampleCounts;
    VkDeviceSize           maxResourceSize;
} VkImageFormatProperties;
```

The extent member of VkImageFormatProperties reports the maximum size of an image
that can be created with the specified format. For example, formats with fewer bits per pixel may
support creating larger images than those with wider pixels. extent is an instance of the
VkExtent3D structure, the definition of which is

[Click here to view code image](#)

```
typedef struct VkExtent3D {
    uint32_t    width;
    uint32_t    height;
    uint32_t    depth;
} VkExtent3D;
```

The maxMipLevels field reports the maximum number of mipmap levels supported for an image
of the requested format along with the other parameters passed to
vkGetPhysicalDeviceImageFormatProperties(). In most cases, maxMipLevels will
either report log_2 (*max* (*extent.x, extent.y, extent.z*)) for the image when mipmaps are supported or 1
when mipmaps are not supported.

The `maxArrayLayers` field reports the maximum number of array layers supported for the image. Again, this is likely to be a fairly high number if arrays are supported or 1 if arrays are not supported.

If the image format supports multisampling, then the supported sample counts are reported through the `sampleCounts` field. This is a bitfield containing one bit for each supported sample count. If bit n is set, then images with 2^n samples are supported in this format. If the format is supported at all, at least one bit of this field will be set. It is very unlikely that you will ever see a format that supports multisampling but does not support a single sample per pixel.

Finally, the `maxResourceSize` field specifies the maximum size, in bytes, that a resource of this format might be. This should not be confused with the maximum extent, which reports the maximum size in each of the dimensions that might be supported. For example, if an implementation reports that it supports images of 16,384 × 16,384 pixels × 2,048 layers with a format containing 128 bits per pixel, then creating an image of the maxium extent in every dimension would produce 8TiB of image data. It's unlikely that an implementation really supports creating an 8TiB image. However, it might well support creating an 8 × 8 × 2,048 array or a 16,384 × 16,284 nonarray image, either of which would fit into a more moderate memory footprint.

Images

Images are more complex than buffers in that they are multidimensional; have specific layouts and format information; and can be used as the source and destination for complex operations such as filtering, blending, depth or stencil testing, and so on. Images are created using the **`vkCreateImage()`** function, the prototype of which is

[Click here to view code image](#)

```
VkResult vkCreateImage (
    VkDevice                        device,
    const VkImageCreateInfo*        pCreateInfo,
    const VkAllocationCallbacks*    pAllocator,
    VkImage*                        pImage);
```

The device that is used to create the image is passed in the `device` parameter. Again, the description of the image is passed through a data structure, the address of which is passed in the `pCreateInfo` parameter. This is a pointer to an instance of the `VkImageCreateInfo` structure, the definition of which is

[Click here to view code image](#)

```
typedef struct VkImageCreateInfo {
    VkStructureType         sType;
    const void*             pNext;
    VkImageCreateFlags      flags;
    VkImageType             imageType;
    VkFormat                format;
    VkExtent3D              extent;
    uint32_t                mipLevels;
    uint32_t                arrayLayers;
    VkSampleCountFlagBits   samples;
    VkImageTiling           tiling;
    VkImageUsageFlags       usage;
    VkSharingMode           sharingMode;
```

```
        uint32_t                    queueFamilyIndexCount;
        const uint32_t*             pQueueFamilyIndices;
        VkImageLayout               initialLayout;
    } VkImageCreateInfo;
```

As you can see, this is a significantly more complex structure than the `VkBufferCreateInfo` structure. The common fields, `sType` and `pNext`, appear at the top, as with most other Vulkan structures. The `sType` field should be set to `VK_STRUCTURE_TYPE_IMAGE_CREATE_INFO`.

The `flags` field of `VkImageCreateInfo` contains flags describing some of the properties of the image. These are a selection of the `VkImageCreateFlagBits` enumeration. The first three—`VK_IMAGE_CREATE_SPARSE_BINDING_BIT`, `VK_IMAGE_CREATE_SPARSE_RESIDENCY_BIT`, and `VK_IMAGE_CREATE_SPARSE_ALIASED_BIT`—are used for controlling sparse images, which are covered later in this chapter.

If `VK_IMAGE_CREATE_MUTABLE_FORMAT_BIT` is set, then you can create *views* of the image with a different format from the parent. Image views are essentially a special type of image that shares data and layout with its parent but can override parameters such as format. This allows data in the image to be interpreted in multiple ways at the same time. Using image views is a way to create two different aliases for the same data. Image views are covered later in this chapter. If `VK_IMAGE_CREATE_CUBE_COMPATIBLE_BIT` is set, then the you will be able to create cube map views of it. Cube maps are covered later in this chapter.

The `imageType` field of the `VkImageCreateInfo` structure specifies the type of image that you want to create. The image type is essentially the dimensionality of the image and can be one of `VK_IMAGE_TYPE_1D`, `VK_IMAGE_TYPE_2D`, or `VK_IMAGE_TYPE_3D` for a 1D, 2D, or 3D image, respectively.

Images also have a format, which describes how texel data is stored in memory and how it is interpreted by Vulkan. The format of the image is specified by the `format` field of the `VkImageCreateInfo` structure and must be one of the image formats represented by a member of the `VkFormat` enumeration. Vulkan supports a large number of formats—too many to list here. We will use some of the formats in the book examples and explain how they work at that time. For the rest, refer to the Vulkan specification.

The *extent* of an image is its size in texels. This is specified in the `extent` field of the `VkImageCreateInfo` structure. This is an instance of the `VkExtent3D` structure, which has three members: `width`, `height`, and `depth`. These should be set to the width, height, and depth of the desired image, respectively. For 1D images, `height` should be set to 1, and for 1D and 2D images, `depth` should be set to 1. Rather than alias the next-higher dimension as an array count, Vulkan uses an explicit array size, which is set in `arrayLayers`.

The maximum size of an image that can be created is device-dependent. To determine the largest image size, call **vkGetPhysicalDeviceFeatures()** and check the `maxImageDimension1D`, `maxImageDimension2D`, and `maxImageDimension3D` fields of the embedded `VkPhysicalDeviceLimits` structure. `maxImageDimension1D` contains the maximum supported width for 1D images, `maxImageDimension2D` the maximum side length for 2D images, and `maxImageDimension3D` the maximum side length for 3D images. Likewise, the maximum number of layers in an array image is contained in the `maxImageArrayLayers` field.

If the image is a cube map, then the maximum side length for the cube is stored in `maxImageDimensionCube`.

`maxImageDimension1D`, `maxImageDimension2D`, and `maxImageDimensionCube` are guaranteed to be at least 4,096 texels, and `maxImageDimensionCube` and `maxImageArrayLayers` are guaranteed to be at least 256. If the image you want to create is smaller than these dimensions, then there's no need to check the device features. Further, it's quite common to find Vulkan implementations that support significantly higher limits. It would be reasonable to make larger image sizes a hard requirement rather than trying to create fallback paths for lower-end devices.

The number of mipmap levels to create in the image is specified in `mipLevels`. Mipmapping is the process of using a set of prefiltered images of successively lower resolution in order to improve image quality when undersampling the image. The images that make up the various mipmap levels are arranged in a pyramid, as shown in Figure 2.1.
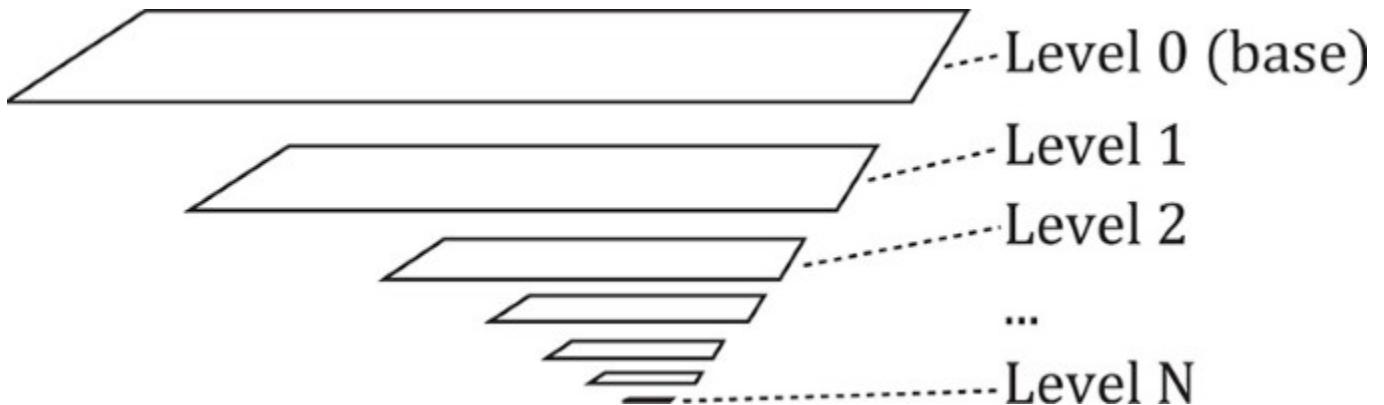
Figure 2.1: Mipmap Image Layout

In a mipmapped texture, the base level is the lowest-numbered level (usually level zero) and has the resolution of the texture. Each successive level is half the size of the level above it until halving the size of the image again in one of the dimensions would result in a single texel in that direction. Sampling from mipmapped textures is covered in some detail in Chapter 6, "Shaders and Pipelines."

Likewise, the number of samples in the image is specified in `samples`. This field is somewhat unlike the others. It must be a member of the `VkSampleCountFlagBits` enumeration, which is actually defined as bits to be used in a bitfield. However, only power-of-two sample counts are currently defined in Vulkan, which means they're "1-hot" values, so single-bit enumerant values work just fine.

The next few fields describe how the image will be used. First is the tiling mode, specified in the `tiling` field. This is a member of the `VkImageTiling` enumeration, which contains only `VK_IMAGE_TILING_LINEAR` or `VK_IMAGE_TILING_OPTIMAL`. *Linear* tiling means that image data is laid out left to right, top to bottom,[1] such that if you map the underlying memory and write it with the CPU, it would form a linear image. Meanwhile, *optimal* tiling is an opaque representation used by Vulkan to lay data out in memory to improve efficiency of the memory subsystem on the device. This is generally what you should choose unless you plan to map and manipulate the image with the CPU. Optimal tiling will likely perform *significantly* better than linear tiling in most operations, and linear tiling might not be supported at all for some operations or formats, depending on the Vulkan implementation.

The `usage` field is a bitfield describing where the image will be used. This is similar to the `usage` field in the `VkBufferCreateInfo` structure. The `usage` field here is made up of members of the `VkImageUsageFlags` enumeration, the members of which are as follows:

- `VK_IMAGE_USAGE_TRANSFER_SRC_BIT` and `VK_IMAGE_USAGE_TRANSFER_DST_BIT` mean that the image will be the source or destination of transfer commands, respectively. Transfer commands operating on images are covered in [Chapter 4](), "[Moving Data]()."

- `VK_IMAGE_USAGE_SAMPLED_BIT` means that the image can be sampled from in a shader.

- `VK_IMAGE_USAGE_STORAGE_BIT` means that the image can be used for general-purpose storage, including writes from a shader.

- `VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT` means that the image can be bound as a color attachment and drawn into using graphics operations. Framebuffers and their attachments are covered in [Chapter 7](), "[Graphics Pipelines]()."

- `VK_IMAGE_USAGE_DEPTH_STENCIL_ATTACHMENT_BIT` means that the image can be bound as a depth or stencil attachment and used for depth or stencil testing (or both). Depth and stencil operations are covered in [Chapter 10](), "[Fragment Processing]()."

- `VK_IMAGE_USAGE_TRANSIENT_ATTACHMENT_BIT` means that the image can be used as a transient attachment, which is a special kind of image used to store intermediate results of a graphics operation. Transient attachments are covered in [Chapter 13](), "[Multipass Rendering]()."

- `VK_IMAGE_USAGE_INPUT_ATTACHMENT_BIT` means that the image can be used as a special input during graphics rendering. Input images differ from regular sampled or storage images in that only fragment shaders can read from them and only at their own pixel location. Input attachments are also covered in detail in [Chapter 13](), "[Multipass Rendering]()."

The `sharingMode` is identical in function to the similarly named field in the `VkBufferCreateInfo` structure described in "[Buffers]()" earlier in this chapter. If it is set to `VK_SHARING_MODE_EXCLUSIVE`, the image will be used with only a single queue family at a time. If it is set to `VK_SHARING_MODE_CONCURRENT`, then the image may be accessed by multiple queues concurrently. Likewise, `queueFamilyIndexCount` and `pQueueFamilyIndices` provide similar function and are used when `sharingMode` is `VK_SHARING_MODE_CONCURRENT`.

Finally, images have a *layout*, which specifies in part how it will be used at any given moment. The `initialLayout` field determines which layout the image will be created in. The available layouts are the members of the `VkImageLayout` enumeration, which are

- `VK_IMAGE_LAYOUT_UNDEFINED`: The state of the image is undefined. The image must be moved into one of the other layouts before it can be used almost for anything.

- `VK_IMAGE_LAYOUT_GENERAL`: This is the "lowest common denominator" layout and is used where no other layout matches the intended use case. Images in `VK_IMAGE_LAYOUT_GENERAL` can be used almost anywhere in the pipeline.

- `VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL`: The image is going to be rendered into using a graphics pipeline.

- `VK_IMAGE_LAYOUT_DEPTH_STENCIL_ATTACHMENT_OPTIMAL`: The image is going to be used as a depth or stencil buffer as part of a graphics pipeline.

- `VK_IMAGE_LAYOUT_DEPTH_STENCIL_READ_ONLY_OPTIMAL`: The image is going to be used for depth testing but will not be written to by the graphics pipeline. In this special state, the image can also be read from in shaders.

- `VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL`: The image will be bound for reading by shaders. This layout is typically used when an image is going to be used as a texture.

- `VK_IMAGE_LAYOUT_TRANSFER_SRC_OPTIMAL`: The image is the source of copy operations.

- `VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL`: The image is the destination of copy operations.

- `VK_IMAGE_LAYOUT_PREINITIALIZED`: The image contains data placed there by an external actor, such as by mapping the underlying memory and writing into it from the host.

- `VK_IMAGE_LAYOUT_PRESENT_SRC_KHR`: The image is used as the source for presentation, which is the act of showing it to the user.

Images can be moved from layout to layout, and we will cover the various layouts as we introduce the topics related to them. However, images must initially be created in either the `VK_IMAGE_LAYOUT_UNDEFINED` or the `VK_IMAGE_LAYOUT_PREINITIALIZED` layout. `VK_IMAGE_LAYOUT_PREINITIALIZED` should be used only when you have data in memory that you will bind to the image resource immediately. `VK_IMAGE_LAYOUT_UNDEFINED` should be used when you plan to move the resource to another layout before use. Images can be moved out of `VK_IMAGE_LAYOUT_UNDEFINED` layout at little or no cost at any time.

The mechanism for changing the layout of an image is known as a *pipeline barrier*, or simply a barrier. A barrier not only serves as a means to change the layout of a resource but can also synchronize access to that resource by different stages in the Vulkan pipeline and even by different queues running concurrently on the same device. As such, a pipeline barrier is fairly complex and quite difficult to get right. Pipeline barriers are discussed in some detail in Chapter 4, "Moving Data," and are further explained in the sections of the book where they are relevant.

Listing 2.4 shows a simple example of creating an image resource.

Listing 2.4: Creating an Image Object

Click here to view code image

```
VkImage image = VK_NULL_HANDLE;
VkResult result = VK_SUCCESS;

static const
VkImageCreateInfo imageCreateInfo =
{
    VK_STRUCTURE_TYPE_IMAGE_CREATE_INFO,        // sType
    nullptr,                                    // pNext
    0,                                          // flags
    VK_IMAGE_TYPE_2D,                           // imageType
    VK_FORMAT_R8G8B8A8_UNORM,                   // format
    { 1024, 1024, 1 },                          // extent
```

```
        10,                                      // mipLevels
        1,                                       // arrayLayers
        VK_SAMPLE_COUNT_1_BIT,                   // samples
        VK_IMAGE_TILING_OPTIMAL,                 // tiling
        VK_IMAGE_USAGE_SAMPLED_BIT,              // usage
        VK_SHARING_MODE_EXCLUSIVE,               // sharingMode
        0,                                       // queueFamilyIndexCount
        nullptr,                                 // pQueueFamilyIndices
        VK_IMAGE_LAYOUT_UNDEFINED                // initialLayout
    };
    result = vkCreateImage(device, &imageCreateInfo, nullptr, &image);
```

The image created by the code in is a 1,024 × 1,024 texel 2D image with a single sample, in `VK_FORMAT_R8G8B8A8_UNORM` format and optimal tiling. The code creates it in the undefined layout, which means that we can move it to another layout later to place data into it. The image is to be used as a texture in one of our shaders, so we set the `VK_IMAGE_USAGE_SAMPLED_BIT` usage flag. In our simple applications, we use only a single queue, so we set the sharing mode to exclusive.

Linear Images

As discussed earlier, two tiling modes are available for use in image resources: `VK_IMAGE_TILING_LINEAR` and `VK_IMAGE_TILING_OPTIMAL`. The `VK_IMAGE_TILING_OPTIMAL` mode represents an opaque, implementation-defined layout that is intended to improve the efficiency of the memory subsystem of the device for read and write operations on the image. However, `VK_IMAGE_TILING_LINEAR` is a transparent layout of the data that is intended to be intuitive. Pixels in the image are laid out left to right, top to bottom. Therefore, it's possible to map the memory used to back the resource to allow the host to read and write to it directly.

In addition to the image's width, height, depth, and pixel format, a few pieces of information are needed to enable host access to the underlying image data. These are the row pitch of the image, which is the distance in bytes between the start of each row of the image; the array pitch, which is the distance between array layers; and the depth pitch, which is the distance between depth slices. Of course, the array pitch and depth pitch apply only to array or 3D images, respectively, and the row pitch applies only to 2D or 3D images.

An image is normally made up of several subresources. Some formats have more than one *aspect*, which is a component of the image such as the depth or stencil component in a depth-stencil image. Mipmap levels and array layers are also considered to be separate subresources. The layout of each subresource within an image may be different and therefore has different layout information. This information can be queried by calling **vkGetImageSubresourceLayout()**, the prototype of which is

[Click here to view code image](#)

```
void vkGetImageSubresourceLayout (
    VkDevice                        device,
    VkImage                         image,
    const VkImageSubresource*       pSubresource,
    VkSubresourceLayout*            pLayout);
```

The device that owns the image that is being queried is passed in `device`, and the image being queried is passed in `image`. A description of the subresource is passed through an instance of the `VkImageSubresource` structure, a pointer to which is passed in the `pSubresource` parameter. The definition of `VkImageSubresource` is

[**Click here to view code image**](#)

```
typedef struct VkImageSubresource {
    VkImageAspectFlags    aspectMask;
    uint32_t              mipLevel;
    uint32_t              arrayLayer;
} VkImageSubresource;
```

The aspect or aspects of the image that you want to query the layout of is specified in `aspectMask`. For color images, this should be `VK_IMAGE_ASPECT_COLOR_BIT`, and for depth, stencil, or depth-stencil images, this should be some combination of `VK_IMAGE_ASPECT_DEPTH_BIT`, and `VK_IMAGE_ASPECT_STENCIL_BIT`. The mipmap level for which the parameters are to be returned is specified in `mipLevel`, and the array layer is specified in `arrayLayer`. You should normally set `arrayLayer` to zero, as the parameters of the image aren't expected to change across layers.

When **vkGetImageSubresourceLayout()** returns, it will have written the layout parameters of the subresource into the `VkSubresourceLayout` structure pointed to by `pLayout`. The definition of `VkSubresourceLayout` is

[**Click here to view code image**](#)

```
typedef struct VkSubresourceLayout {
    VkDeviceSize    offset;
    VkDeviceSize    size;
    VkDeviceSize    rowPitch;
    VkDeviceSize    arrayPitch;
    VkDeviceSize    depthPitch;
} VkSubresourceLayout;
```

The size of the memory region consumed by the requested subresource is returned in `size`, and the offset within the resource where the subresource begins is returned in `offset`. The `rowPitch`, `arrayPitch`, and `depthPitch` fields contain the row, array layer, and depth slice pitches, respectively. The unit of these fields is always bytes, regardless of the pixel format of the images. Pixels within a row are always tightly packed. [Figure 2.2](#) illustrates how these parameters represent memory layout of an image. In the figure, the valid image data is represented by the grey grid, and padding around the image is shown as blank space.
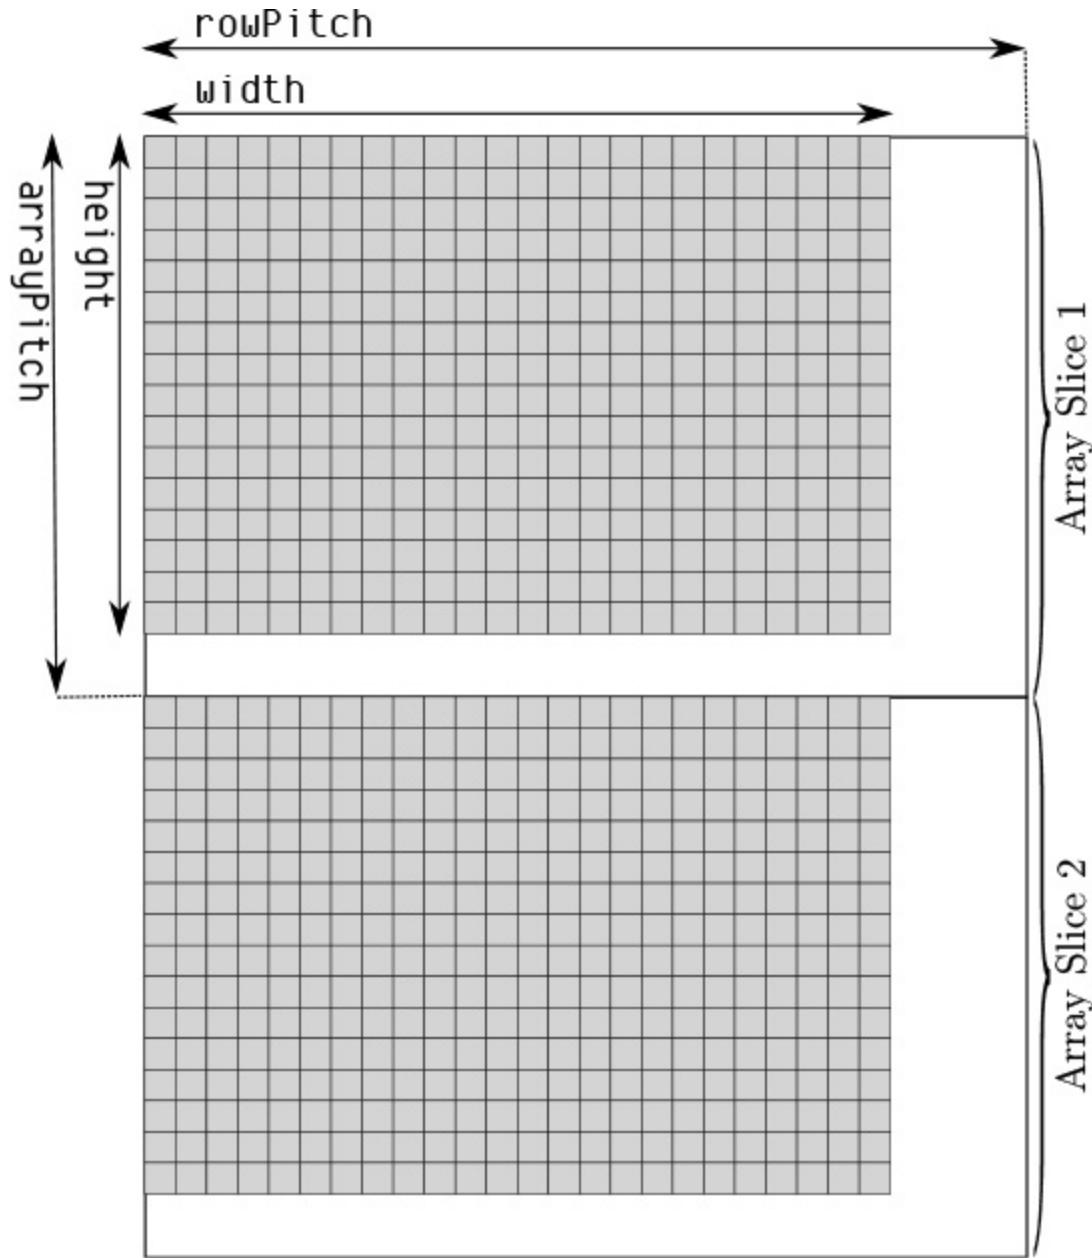
Figure 2.2: Memory Layout of LINEAR Tiled Images

Given the memory layout of an image in LINEAR tiling mode, it is possible to trivially compute the memory address for a single texel within the image. Loading image data into a LINEAR tiled image is then simply a case of loading scanlines from the image into memory at the right location. For many texel formats and image dimensions, it is highly likely that the image's rows are tightly packed in memory—that is, the rowPitch field of the VkSubresourceLayout structure is equal to the subresource's width. In this case, many image-loading libraries will be able to load the image directly into the mapped memory of the image.

Nonlinear Encoding

You may have noticed that some of the Vulkan image formats include `SRGB` in their names. This refers to *sRGB* color encoding, which is a nonlinear encoding that uses a gamma curve approximating that of CRTs. Although CRTs are all but obsolete now, sRGB encoding is still in widespread use for texture and image data.

Because the amount of light energy produced by a CRT is not linear with the amount of electrical energy used to produce the electron beam that excites the phosphor, an inverse mapping must be applied to color signals to make a linear rise in numeric value produce a linear increase in light output. The amount of light output by a CRT is approximately

$$L_{out} = V_{in}{}^{\gamma}$$

The standard value of γ in NTSC television systems (common in North America, parts of South America, and parts of Asia) is 2.2. Meanwhile, the standard value of γ in SECAM and PAL systems (common in Europe, Africa, Australia, and other regions of Asia) is 2.8.

The sRGB curve attempts to compensate for this by applying *gamma correction* to linear data in memory. The standard sRGB transfer function is not a pure gamma curve but is made up of a short linear section followed by a curved, gamma-corrected section. The function applied to data to go from linear to sRGB space is

[Click here to view code image](#)

```
if (cl >= 1.0)
{
    cs = 1.0;
}
else if (cl <= 0.0)
{
    cs = 0.0;
}
else if (cl < 0.0031308)
{
    cs = 12.92 * cl;
}
else
{
    cs = 1.055 * pow(cl, 0.41666) - 0.055;
}
```

To go from sRGB space to linear space, the following transform is made:

[Click here to view code image](#)

```
if (cs >= 1.0)
{
    cl = 1.0;
}
else if (cs <= 0.0)
{
    cl = 0.0;
}
```

```
    else if (cs <= 0.04045)
    {
        cl = cs / 12.92;
    }
    else
    {
        cl = pow((cs + 0.0555) / 1.055), 2.4)
    }
```

In both code snippets, `cs` is the sRGB color space value, and `cl` is the linear value. Figure 2.3 shows a side-by-side comparison of a simple $\gamma = 2.2$ curve and the standard sRGB transfer function. As you can see in the figure, the curves for sRGB correction (shown on the top) and a simple power curve (shown on the bottom) are almost identical. While Vulkan implementations are expected to implement sRGB using the official definition, if you need to perform the transformation manually in your shaders, you may be able to get away with a simple power function without accumulating too much error.
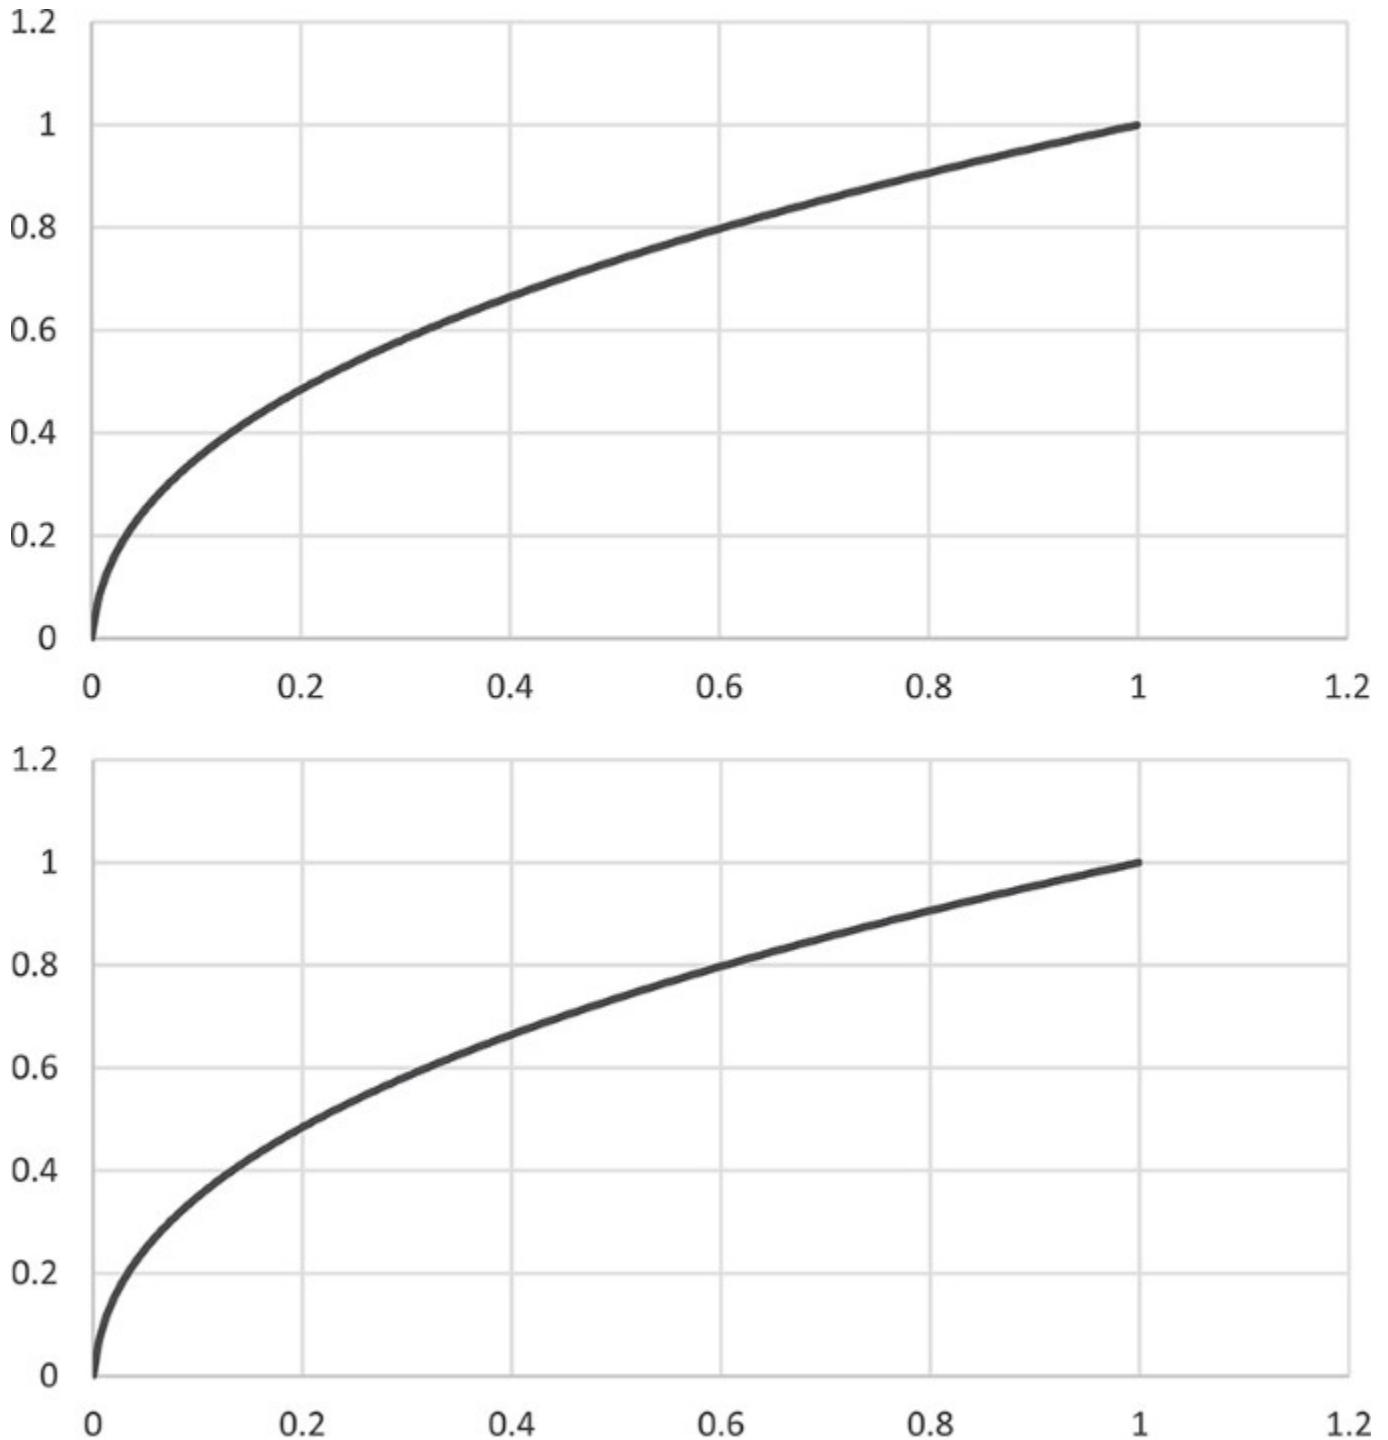
Figure 2.3: Gamma Curves for sRGB (Top) and Simple Powers (Bottom)

When rendering to an image in sRGB format, linear values produced by your shaders are transformed to sRGB encoding before being written into the image. When reading from an image in sRGB format, texels are transformed from sRGB format back to linear space before being returned to your shader.

Blending and interpolation always occurs in linear space such that data read from a framebuffer is first transformed from sRGB to linear space and then blended with the source data in linear space, and the final result is transformed back to sRGB encoding before being written into the framebuffer.

Rendering in sRGB space provides more precision in darker colors and can result in less banding artifacts and richer colors. However, for best image quality, including high-dynamic-range rendering, it's best to choose a floating-point color format and render in a linear space, converting to sRGB as late as possible before display.

Compressed Image Formats

Image resources are likely to be the largest consumers of device memory in your application. For this reason, Vulkan provides the capability for images to be compressed. Image compression provides two significant benefits to an application:

- It reduces the total amount of memory consumed by image resources used by the application.
- It reduces the total *memory bandwidth* consumed while accessing those resources.

All currently defined compressed image formats in Vulkan are what are known as *block compressed* formats. Texels within an image are compressed in small square or rectangular blocks that can be decompressed independently of all others. All formats are lossy, and the compression ratio is not competitive with formats such as JPEG or even PNG. However, decompression is fast and cheap to implement in hardware, and random access to texels is relatively straightforward.

Support for various compressed image formats is optional, but all Vulkan implementations are required to support at least one family of formats. You can determine which family of compressed formats is supported by checking various fields of the device's `VkPhysicalDeviceFeatures` structure as returned from a call to **`vkGetPhysicalDeviceProperties()`**.

If `textureCompressionBC` is `VK_TRUE`, then the device supports the *block compressed* formats, also known as BC formats. The BC family includes

- BC1: Made up of the `VK_FORMAT_BC1_RGB_UNORM_BLOCK`, `VK_FORMAT_BC1_RGB_SRGB_BLOCK`, `VK_FORMAT_BC1_RGBA_UNORM_BLOCK`, and `VK_FORMAT_BC1_RGBA_SRGB_BLOCK` formats, BC1 encodes images in blocks of 4×4 texels, with each block represented as a 64-bit quantity.
- BC2: Consisting of `VK_FORMAT_BC2_UNORM_BLOCK` and `VK_FORMAT_BC2_SRGB_BLOCK`, BC2 encodes images in blocks of 4×4 texels, with each block represented as a 128-bit quantity. BC2 images always have an alpha channel. The encoding for the RGB channels is the same as with BC1 RGB formats, and the alpha is stored as 4 bits per texel in a second 64-bit field before the BC1 encoded RGB data.
- BC3: The `VK_FORMAT_BC3_UNORM_BLOCK` and `VK_FORMAT_BC3_SRGB_BLOCK` formats make up the BC3 family, again encoding texels in 4×4 blocks, with each block consuming 128 bits of storage. The first 64-bit quantity stores compressed alpha values, allowing coherent alpha data to be stored with higher precision than BC2. The second 64-bit quantity stores compressed color data in a similar form to BC1.
- BC4: `VK_FORMAT_BC4_UNORM_BLOCK` and `VK_FORMAT_BC4_SRGB_BLOCK` represent single-channel formats, again encoded as 4×4 blocks of texels, with each block consuming 64 bits of storage. The encoding of the single-channel data is essentially the same as that of the alpha channel of a BC3 image.
- BC5: Made up of `VK_FORMAT_BC5_UNORM_BLOCK` and `VK_FORMAT_BC5_SRGB_BLOCK`, the BC5 family is a two-channel format, with each 4×4 block essentially consisting of two BC4 blocks back-to-back.

- BC6: The `VK_FORMAT_BC6H_SFLOAT_BLOCK` and `VK_FORMAT_BC6H_UFLOAT_BLOCK` formats are signed and unsigned floating-point compressed formats, respectively. Each 4×4 block of RGB texels is stored in 128 bits of data.

- BC7: `VK_FORMAT_BC7_UNORM_BLOCK` and `VK_FORMAT_BC7_SRGB_BLOCK` are four-channel formats with each 4×4 block of RGBA texel data stored in a 128-bit component.

If the `textureCompressionETC2` member of `VkPhysicalDeviceFeatures` is `VK_TRUE`, then the device supports the *ETC* formats, including ETC2 and EAC. The following formats are included in this family:

- `VK_FORMAT_ETC2_R8G8B8_UNORM_BLOCK` and `VK_FORMAT_ETC2_R8G8B8_SRGB_BLOCK`: Unsigned formats where 4×4 blocks of RGB texels are packed into 64 bits of compressed data.

- `VK_FORMAT_ETC2_R8G8B8A1_UNORM_BLOCK` and `VK_FORMAT_ETC2_R8G8B8A1_SRGB_BLOCK`: Unsigned formats where 4×4 blocks of RGB texels plus a one-bit alpha value per texel are packed into 64 bits of compressed data.

- `VK_FORMAT_ETC2_R8G8B8A8_UNORM_BLOCK` and `VK_FORMAT_ETC2_R8G8B8A8_SRGB_BLOCK`: Each 4×4 block of texels is represented as a 128-bit quantity. Each texel has 4 channels.

- `VK_FORMAT_EAC_R11_UNORM_BLOCK` and `VK_FORMAT_EAC_R11_SNORM_BLOCK`: Unsigned and signed single-channel formats with each 4×4 block of texels represented as a 64-bit quantity.

- `VK_FORMAT_EAC_R11G11_UNORM_BLOCK` and `VK_FORMAT_EAC_R11G11_SNORM_BLOCK`: Unsigned and signed two-channel formats with each 4×4 block of texels represented as a 64-bit quantity.

The final family is the ASTC family. If the `textureCompressionASTC_LDR` member of `VkPhysicalDeviceFeatures` is `VK_TRUE`, then the device supports the *ASTC* formats. You may have noticed that for all of the formats in the BC and ETC families, the block size is fixed at 4×4 texels, but depending on format, the texel format and number of bits used to store the compressed data vary.

ASTC is different here in that the number of bits per block is always 128, and all ASTC formats have four channels. However, the block size in texels can vary. The following block sizes are supported: 4×4, 5×4, 5×5, 6×5, 6×6, 8×5, 8×6, 8×8, 10×5, 10×6, 10×8, 10×10, 12×10, and 12×12.

The format of the token name for ASTC formats is formulated as `VK_FORMAT_ASTC_{N}x{M}_{encoding}_BLOCK`, where `{N}` and `{M}` represent the width and height of the block, and `{encoding}` is either `UNORM` or `SRGB`, depending on whether the data is linear or encoded as sRGB nonlinear. For example, `VK_FORMAT_ASTC_8x6_SRGB_BLOCK` is an RGBA ASTC compressed format with 8×6 blocks and sRGB encoded data.

For all formats including SRGB, only the R, G, and B channels use nonlinear encoding. The A channel is always stored with linear encoding.

Resource Views

Buffers and images are the two primary types of resources supported in Vulkan. In addition to creating these two resource types, you can create *views* of existing resources in order to partition them, reinterpret their content, or use them for multiple purposes. Views of buffers, which represent a subrange of a buffer object, are known as *buffer views*, and views of images, which can alias formats or represent a subresource of another image, are known as *image views*.

Before a view of a buffer or image can be created, you need to bind memory to the parent object.

Buffer Views

A buffer view is used to interpret the data in a buffer with a specific format. Because the raw data in the buffer is then treated as a sequence of texels, this is also known as a texel buffer view. A texel buffer view can be accessed directly in shaders, and Vulkan will automatically convert the texels in the buffer into the format expected by the shader. One example use for this functionality is to directly fetch the properties of vertices in a vertex shader by reading from a texel buffer rather than using a vertex buffer. While this is more restrictive, it does allow random access to the data in the buffer.

To create a buffer view, call **vkCreateBufferView()**, the prototype of which is

[Click here to view code image](#)

```
VkResult vkCreateBufferView (
    VkDevice                              device,
    const VkBufferViewCreateInfo*         pCreateInfo,
    const VkAllocationCallbacks*          pAllocator,
    VkBufferView*                         pView);
```

The device that is to create the new view is passed in `device`. This should be the same device that created the buffer of which you are creating a view. The remaining parameters of the new view are passed through a pointer to an instance of the `VkBufferViewCreateInfo` structure, the definition of which is

[Click here to view code image](#)

```
typedef struct VkBufferViewCreateInfo {
    VkStructureType           sType;
    const void*               pNext;
    VkBufferViewCreateFlags   flags;
    VkBuffer                  buffer;
    VkFormat                  format;
    VkDeviceSize              offset;
    VkDeviceSize              range;
} VkBufferViewCreateInfo;
```

The `sType` field of `VkBufferViewCreateInfo` should be set to `VK_STRUCTURE_TYPE_BUFFER_VIEW_CREATE_INFO`, and `pNext` should be set to `nullptr`. The `flags` field is reserved and should be set to 0. The parent buffer is specified in `buffer`. The new view will be a "window" into the parent buffer starting at `offset` bytes and extending for `range` bytes. When bound as a texel buffer, the data in the buffer is interpreted as a sequence of texels with the format as specified in `format`.

The maximum number of texels that can be stored in a texel buffer is determined by inspecting the `maxTexelBufferElements` field of the device's `VkPhysicalDeviceLimits` structure, which can be retrieved by calling **`vkGetPhysicalDeviceProperties()`**. If the buffer is to be used as a texel buffer, then `range` divided by the size of a texel in `format` must be less than or equal to this limit. `maxTexelBufferElements` is guaranteed to be at least 65,536, so if the view you're creating contains fewer texels, there's no need to query this limit.

The parent buffer must have been created with the `VK_BUFFER_USAGE_UNIFORM_TEXEL_BUFFER_BIT` or `VK_BUFFER_USAGE_STORAGE_TEXEL_BUFFER_BIT` flags in the `usage` field of the `VkBufferCreateInfo` used to create the buffer. The specified format must support the `VK_FORMAT_FEATURE_UNIFORM_TEXEL_BUFFER_BIT`, `VK_FORMAT_FEATURE_STORAGE_TEXEL_BUFFER_BIT`, or `VK_FORMAT_FEATURE_STORAGE_TEXEL_BUFFER_ATOMIC_BIT` as reported by **`vkGetPhysicalDeviceFormatProperties()`**.

On success, **`vkCreateBufferView()`** places the handle to the newly created buffer view in the variable pointed to by `pView`. If `pAllocator` is not `nullptr`, then the allocation callbacks specified in the `VkAllocationCallbacks` structure it points to are used to allocate any host memory required by the new object.

Image Views

In many cases, the image resource cannot be used directly, as more information about it is needed than is included in the resource itself. For example, you cannot use an image resource directly as an attachment to a framebuffer or bind an image into a descriptor set in order to sample from it in a shader. To satisfy these additional requirements, you must create an image view, which is essentially a collection of properties and a reference to a parent image resource.

An image view also allows all or part of an existing image to be seen as a different format. The resulting view of the parent image must have the same dimensions as the parent, although a subset of the parent's array layers or mip levels may be included in the view. The format of the parent and child images must also be *compatible*, which usually means that they have the same number of bits per pixel, even if the data formats are completely different and even if there are a different number of channels in the image.

To create a new view of an existing image, call **`vkCreateImageView()`**, the prototype of which is

[**Click here to view code image**](#)

```
VkResult vkCreateImageView (
    VkDevice                        device,
    const VkImageViewCreateInfo*    pCreateInfo,
    const VkAllocationCallbacks*    pAllocator,
    VkImageView*                    pView);
```

The device that will be used to create the new view and that should own the parent image is specified in `device`. The remaining parameters used in the creation of the new view are passed through an instance of the `VkImageViewCreateInfo` structure, a pointer to which is passed in `pCreateInfo`. The definition of `VkImageViewCreateInfo` is

```
typedef struct VkImageViewCreateInfo {
    VkStructureType          sType;
    const void*              pNext;
    VkImageViewCreateFlags   flags;
    VkImage                  image;
    VkImageViewType          viewType;
    VkFormat                 format;
    VkComponentMapping       components;
    VkImageSubresourceRange  subresourceRange;
} VkImageViewCreateInfo;
```

The `sType` field of `VkImageViewCreateInfo` should be set to
`VK_STRUCTURE_TYPE_IMAGE_VIEW_CREATE_INFO`, and `pNext` should be set to `nullptr`.
The `flags` field is reserved for future use and should be set to 0.

The parent image of which to create a new view is specified in `image`. The type of view to create is
specified in `viewType`. The view type must be compatible with the parent's image type and is a
member of the `VkImageViewType` enumeration, which is larger than the `VkImageType`
enumeration used in creating the parent image. The image view types are as follows:

- `VK_IMAGE_VIEW_TYPE_1D`, `VK_IMAGE_VIEW_TYPE_2D`, and
  `VK_IMAGE_VIEW_TYPE_3D` are the "normal" 1D, 2D, and 3D image types.

- `VK_IMAGE_VIEW_TYPE_CUBE` and `VK_IMAGE_VIEW_TYPE_CUBE_ARRAY` are cube
  map and cube map array images.

- `VK_IMAGE_VIEW_TYPE_1D_ARRAY` and `VK_IMAGE_VIEW_TYPE_2D_ARRAY` are 1D
  and 2D array images.

Note that all images are essentially considered array images, even if they only have one layer. It is,
however, possible to create nonarray views of parent images that refer to one of the layers of the
image.

The format of the new view is specified in `format`. This must be a format that is compatible with
that of the parent image. In general, if two formats have the same number of bits per pixel, then they
are considered compatible. If either or both of the formats is a block compressed image format, then
one of two things must be true:

- If both images have compressed formats, then the number of bits per block must match between
  those formats.

- If only one image is compressed and the other is not, then bits per block in the compressed
  image must be the same as the number of bits per texel in the uncompressed image.

By creating an uncompressed view of a compressed image, you give access to the raw, compressed
data, making it possible to do things like write compressed data from a shader into the image or
interpret the compressed data directly in your application. Note that while all block-compressed
formats encode blocks either as 64-bit or 128-bit quantities, there are no uncompressed, single-
channel 64-bit or 128-bit image formats. To alias a compressed image as an uncompressed format,
you need to choose an uncompressed format with the same number of bits per texel and then
aggregate the bits from the different image channels within your shader to extract the individual
fields from the compressed data.

The component ordering in the view may be different from that in the parent. This allows, for example, an RGBA view of a BGRA format image to be created. This remapping is specified using an instance of `VkComponentMapping`, the definition of which is simply

```
typedef struct VkComponentMapping {
    VkComponentSwizzle    r;
    VkComponentSwizzle    g;
    VkComponentSwizzle    b;
    VkComponentSwizzle    a;
} VkComponentMapping;
```

Each member of `VkComponentMapping` specifies the *source* of data in the parent image that will be used to fill the resulting texel fetched from the child view. They are members of the `VkComponentSwizzle` enumeration, the members of which are as follows:

- `VK_COMPONENT_SWIZZLE_R`, `VK_COMPONENT_SWIZZLE_G`, `VK_COMPONENT_SWIZZLE_B`, and `VK_COMPONENT_SWIZZLE_A` indicate that the source data should be read from the R, G, B, or A channels of the parent image, respectively.

- `VK_COMPONENT_SWIZZLE_ZERO` and `VK_COMPONENT_SWIZZLE_ONE` indicate that the data in the child image should be read as zero or one, respectively, regardless of the content of the parent image.

- `VK_COMPONENT_SWIZZLE_IDENTITY` indicates that the data in the child image should be read from the corresponding channel in the parent image. Note that the numeric value of `VK_COMPONENT_SWIZZLE_IDENTITY` is zero, so simply setting the entire `VkComponentMapping` structure to zero will result in an identity mapping between child and parent images.

The child image can be a subset of the parent image. This subset is specified using the embedded `VkImageSubresourceRange` structure in `subresourceRange`. The definition of `VkImageSubresourceRange` is

```
typedef struct VkImageSubresourceRange {
    VkImageAspectFlags    aspectMask;
    uint32_t              baseMipLevel;
    uint32_t              levelCount;
    uint32_t              baseArrayLayer;
    uint32_t              layerCount;
} VkImageSubresourceRange;
```

The `aspectMask` field is a bitfield made up from members of the `VkImageAspectFlagBits` enumeration specifying which aspects of the image are affected by the barrier. Some image types have more than one logical part, even though the data itself might be interleaved or otherwise related. An example of this is depth-stencil images, which have both a depth component and a stencil component. Each of these two components may be viewable as a separate image in its own right, and these subimages are known as *aspects*. The flags that can be included in `aspectMask` are

- `VK_IMAGE_ASPECT_COLOR_BIT`: The color part of an image. There is usually *only* a color aspect in color images.

- `VK_IMAGE_ASPECT_DEPTH_BIT`: The depth aspect of a depth-stencil image.

- `VK_IMAGE_ASPECT_STENCIL_BIT`: The stencil aspect of a depth-stencil image.

- `VK_IMAGE_ASPECT_METADATA_BIT`: Any additional information associated with the image that might track its state and is used, for example, in various compression techniques.

When you create the new view of the parent image, that view can refer to only one aspect of the parent image. Perhaps the most common use case of this is to create a depth- or stencil-only view of a combined depth-stencil format image.

To create a new image view that corresponds only to a subset of the parent's mip chain, use the `baseMipLevel` and `levelCount` to specify where in the mip chain the view begins and how many mip levels it will contain. If the parent image does not have mipmaps, these fields should be set to zero and one, respectively.

Likewise, to create an image view of a subset of a parent's array layers, use the `baseArrayLayer` and `layerCount` fields to specify the starting layer and number of layers, respectively. Again, if the parent image is not an array image, then `baseArrayLayer` should be set to zero and `layerCount` should be set to one.

Image Arrays

The defined image types (`VkImageType`) include only `VK_IMAGE_TYPE_1D`, `VK_IMAGE_TYPE_2D`, or `VK_IMAGE_TYPE_3D`, which are used to create 1D, 2D, and 3D images, respectively. However, in addition to their sizes in each of the *x*, *y*, and *z* dimensions, all images have a layer count, contained in the `arrayLayers` field of their `VkImageCreateInfo` structure.

Images can be aggregated into arrays, and each element of an array image is known as a layer. Array images allow images to be grouped into single objects, and sampling from multiple layers of the same array image is often more performant than sampling from several loose array objects. Because all Vulkan images have a `layerCount` field, they are all technically array images. However, in practice, we only refer to images with a `layerCount` greater than 1 as an array image.

When views are created of images, the view is explicitly marked as either an array or a nonarray. A nonarray view implicitly has only one layer whereas an array view has multiple layers. Sampling from a nonarray view may perform better than sampling from a single layer of an array image, simply because the device needs to perform fewer indirections and parameter lookups.

A 1D array texture is conceptually different from a 2D texture, and a 2D array texture is different from a 3D texture. The primary difference is that linear filtering can be performed in the *y* direction of a 2D texture and in the *z* direction in a 3D texture, whereas filtering cannot be performed across multiple layers in an array image. Notice that there is no 3D array image view type included in `VkImageViewType`, and most Vulkan implementations will not allow you to create a 3D image with an `arrayLayers` field greater than 1.

In addition to image arrays, a *cube map* is a special type of image that allows groups of six layers of an array image to be interpreted as the sides of a cube. Imagine standing in the center of a cube-shaped room. The room has four walls, a floor, and a ceiling. To your left and right are considered the negative and positive *X* directions, behind and in front of you are the negative and positive *Z* directions, and the floor and ceiling are the negative and positive *Y* directions. These faces are often

notated as the -X, +X, -Y, +Y, -Z, and +Z faces. These are the six faces of a cube map, and a group of six consecutive array layers can be interpreted in that order.

A cube map is sampled using a 3D coordinate. This coordinate is interpreted as a vector pointing from the center of the cube map outward, and the point sampled in the cube-map is the point where the vector meets the cube. Again, put yourself back into the cube-map room and imagine you have a laser pointer. As you point the laser in different directions, the spot on the wall or ceiling is the point from which texture data is taken when the cube map is sampled.

Figure 2.4 shows this pictorially. As you can see in the figure, the cube map is constructed from a selection of six consecutive elements from the parent texture. To create a cube-map view, first create a 2D array image with at least six faces. The `imageType` field of the `VkImageCreateInfo` structure should be set to `VK_IMAGE_TYPE_2D` and the `arrayLayers` field should be at least 6. Note that the number of layers in the parent array doesn't have to be a multiple of 6, but it has to be at least 6.
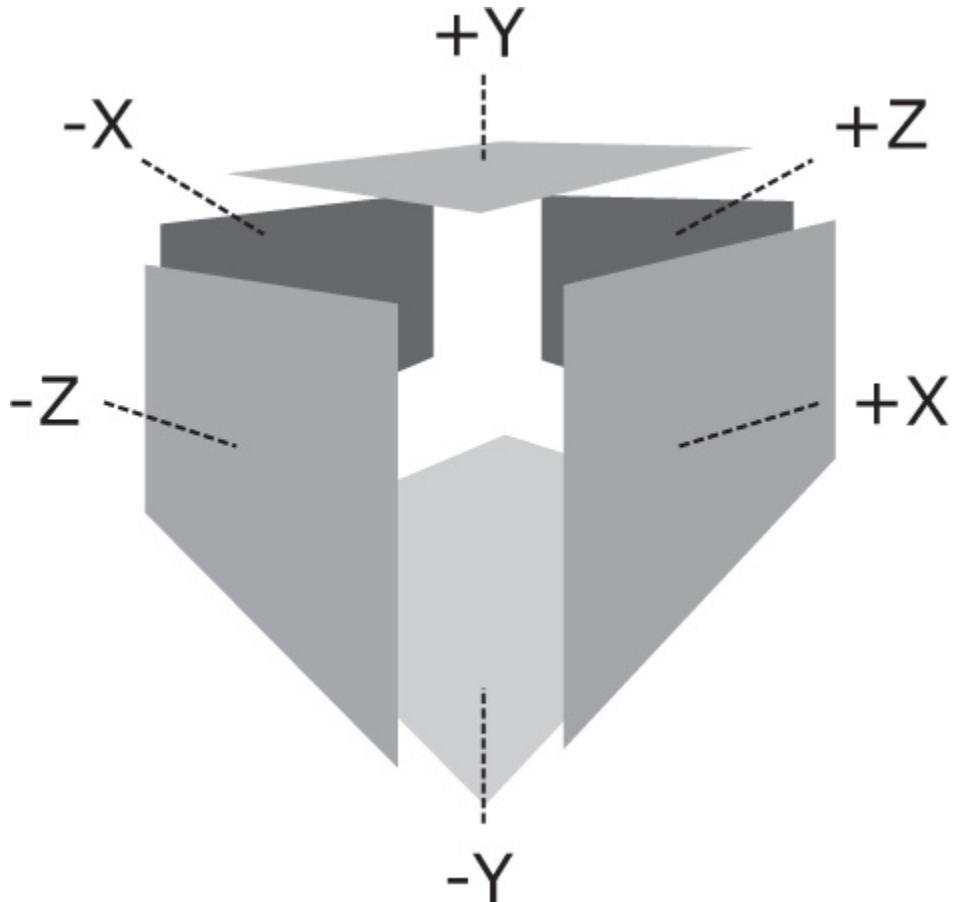
Figure 2.4: Cube Map Construction

The `flags` field of the parent image's `VkImageCreateInfo` structure must have the `VK_IMAGE_CREATE_CUBE_COMPATIBLE_BIT` set, and the image must be square (because the faces of a cube are square).

Next, we create a view of the 2D array parent, but rather than creating a normal 2D (array) view of the image, we create a cube-map view. To do this, set the `viewType` field of the `VkImageViewCreateInfo` structure used to create the view to `VK_IMAGE_VIEW_TYPE_CUBE`. In the embedded `subresourceRange` field, the

`baseArrayLayer` and `layerCount` fields are used to determine where in the array the cube map begins. To create a single cube, `layerCount` should be set to 6.

The first element of the array (at the index specified in the `baseArrayLayer` field) becomes the -X face, and the next five layers become the +X, -Y, +Y, -Z, and +Z faces, in that order.

Cube maps can also form arrays of their own. This is simply a concatenation of an integer multiple of six faces, with each group of six forming a separate cube. To create a cube-map array image, set the `viewType` field of `VkImageViewCreateInfo` to `VK_IMAGE_VIEW_TYPE_CUBE_ARRAY`, and set the `layerCount` to a multiple of 6. The number of cubes in the array is therefore the `layerCount` for the array divided by 6. The number of layers in the parent image must be at least as many layers as are referenced by the cube-map view.

When data is placed in a cube map or cube-map array image, it is treated identically to an array image. Each array layer is laid out consecutively, and commands such as **`vkCmdCopyBufferToImage()`** (which is covered in Chapter 4, "Moving Data") can be used to write into the image. The image can be bound as a color attachment and rendered to. Using layered rendering, you can even write to multiple faces of a cube map in a single drawing command.

Destroying Resources

When you are done with buffers, images, and other resources, it is important to destroy them cleanly. Before destroying a resource, you must make sure that it is not in use and that no work is pending that might access it. Once you are certain that this is the case, you can destroy the resource by calling the appropriate destruction function. To destroy a buffer resource, call **`vkDestroyBuffer()`**, the prototype of which is

Click here to view code image

```
void vkDestroyBuffer (
    VkDevice                                device,
    VkBuffer                                buffer,
    const VkAllocationCallbacks*            pAllocator);
```

The device that owns the buffer object should be specified in `device`, and the handle to the buffer object should be specified in `buffer`. If a host memory allocator was used to create the buffer object, a compatible allocator should be specified in `pAllocator`; otherwise, `pAllocator` should be set to `nullptr`.

Note that destroying a buffer object for which other views exist will also invalidate those views. The view objects themselves must still be destroyed explicitly, but it is not legal to access a view of a buffer that has been destroyed. To destroy a buffer view, call **`vkDestroyBufferView()`**, the prototype of which is

Click here to view code image

```
void vkDestroyBufferView (
    VkDevice                                device,
    VkBufferView                            bufferView,
    const VkAllocationCallbacks*            pAllocator);
```

Again, `device` is a handle to the device that owns the view, and `bufferView` is a handle to the view to be destroyed. `pAllocator` should point to a host memory allocator compatible with that used to create the view or should be set to `nullptr` if no allocator was used to create the view.

Destruction of images is almost identical to that of buffers. To destroy an image object, call **vkDestroyImage()**, the prototype of which is

[**Click here to view code image**](#)

```
void vkDestroyImage (
    VkDevice                             device,
    VkImage                              image,
    const VkAllocationCallbacks*         pAllocator);
```

`device` is the device that owns the image to be destroyed, and `image` is the handle to that image. Again, if a host memory allocator was used to create the original image, then `pAllocator` should point to one compatible with it; otherwise, `pAllocator` should be `nullptr`.

As with buffers, destroying an image invalidates all views of that image. It is not legal to access a view of an image that has already been destroyed. The only thing you can do with such views is to destroy them. Destroying an image view is accomplished by calling **vkDestroyImageView()**, the prototype of which is

[**Click here to view code image**](#)

```
void vkDestroyImageView (
    VkDevice                             device,
    VkImageView                          imageView,
    const VkAllocationCallbacks*         pAllocator);
```

As you might expect, `device` is the device that owns the view being destroyed, and `imageView` is the handle to that view. As with all other destruction functions mentioned so far, `pAllocator` is a pointer to an allocator compatible with the one used to create the view or `nullptr` if no allocator was used.

Device Memory Management

When the Vulkan device operates on data, the data must be stored in *device memory*. This is memory that is accessible to the device. In a Vulkan system there are four classes of memory. Some systems may have only a subset of these, and some may only have two. Given a host (the processor upon which your application is running) and a device (the processor that executes your Vulkan commands), there could be separate memory physically attached to each. In addition, some regions of the physical memory attached to each processor might be accessible to the other processor or processors in the system.

In some cases, the visible region of shared memory might be relatively small, and in other cases, there may actually be only one physical piece of memory, which is shared between the host and the device. [Figure 2.5](#) demonstrates the memory map of a host and device with physically separate memories.
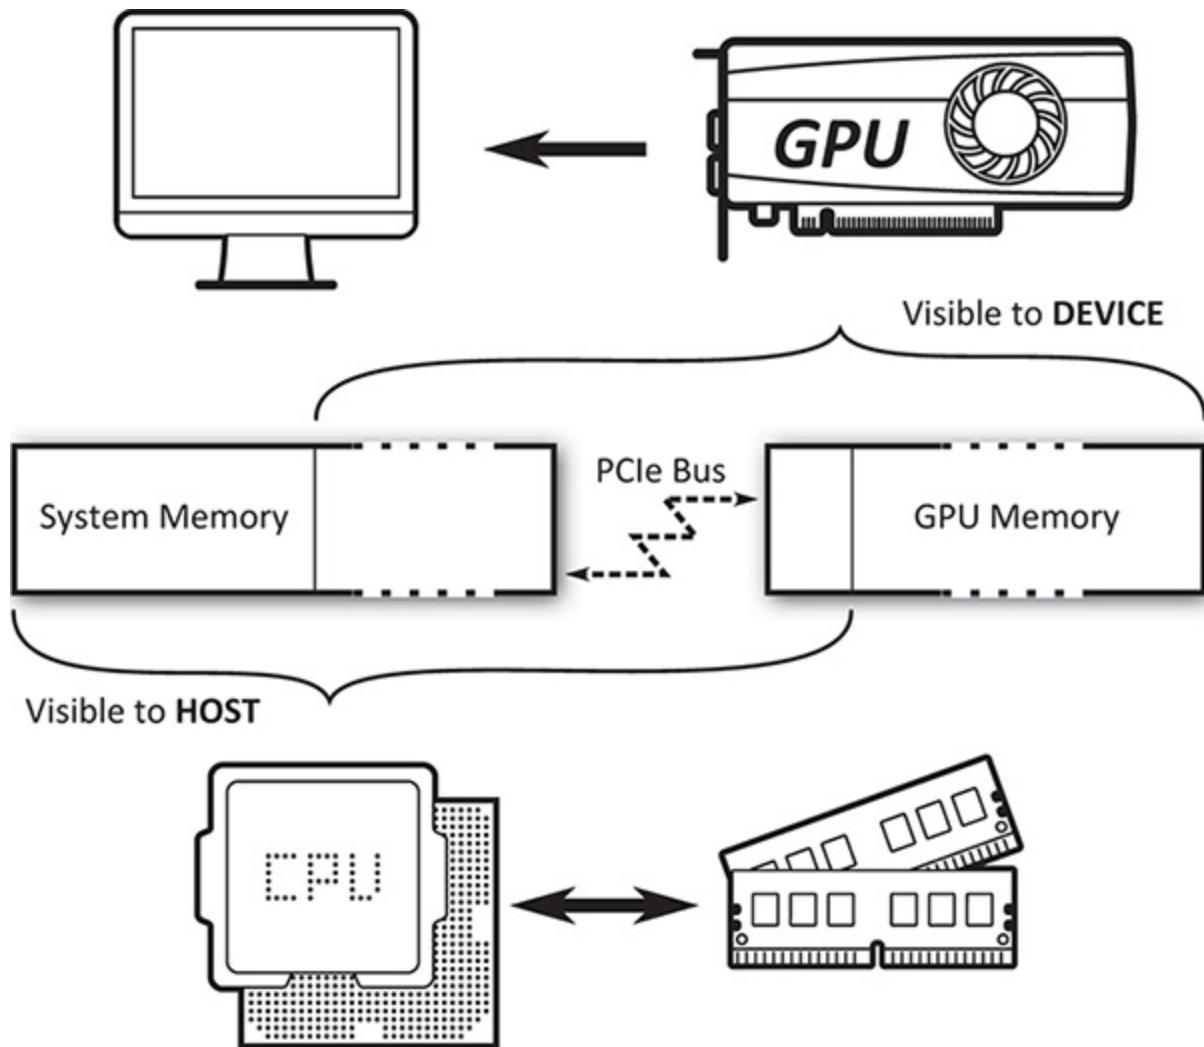
Figure 2.5: Host and Device Memory

Any memory that is accessible to the device is known as *device memory*, even if that memory is physically attached to the host. In this case, it is *host local device memory*. This is distinct from host memory, which might also be known as system memory, which is regular memory allocated with a function such as `malloc` or **`new`**. Device memory may also be accessible to the host through a mapping.

A typical discrete GPU as found on an add-in card plugged into a PCI-Express slot will have an amount of dedicated memory physically attached to its circuit board. Some part of this memory may be accessible *only* to the device, and some part of the memory may be accessible to the host through some form of window. In addition, the GPU will have access to some or all of the host's system memory. All of these pools of memory will appear as a heap to the host, and memory will be mapped into those heaps via the various types of memory.

On the other hand, a typical embedded GPU—such as those found in embedded systems, mobile devices, or even laptop processors—may share memory controller and subsystem with the host processor. In this case, it is likely that access to main system memory is coherent and the device will expose fewer heaps—perhaps only one. This is considered a *unified memory architecture*.

Allocating Device Memory

A device memory allocation is represented as a `VkDeviceMemory` object that is created using the `vkAllocateMemory()` function, the prototype of which is

Click here to view code image

```
VkResult vkAllocateMemory (
    VkDevice                            device,
    const VkMemoryAllocateInfo*         pAllocateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkDeviceMemory*                     pMemory);
```

The device that will use the memory is passed in `device`. `pAllocateInfo` describes the new device memory object which, if the allocation is successful, will be placed in the variable pointed to by `pMemory`. `pAllocateInfo` points to an instance of the `VkMemoryAllocateInfo` structure, the definition of which is

Click here to view code image

```
typedef struct VkMemoryAllocateInfo {
    VkStructureType    sType;
    const void*        pNext;
    VkDeviceSize       allocationSize;
    uint32_t           memoryTypeIndex;
} VkMemoryAllocateInfo;
```

This is a simple structure containing only the size and the memory type to be used for the allocation. `sType` should be set to `VK_STRUCTURE_TYPE_MEMORY_ALLOCATE_INFO`, and `pNext` should be set to `nullptr` unless an extension is in use that requires more information about the allocation. The size of the allocation is passed in `allocationSize` and is measured in bytes. The memory type, passed in `memoryTypeIndex`, is an index into the memory type array returned from a call to `vkGetPhysicalDeviceMemoryProperties()`, as described in "Physical Device Memory" in Chapter 1, "Overview of Vulkan."

Once you have allocated device memory, it can be used to back resources such as buffers and images. Vulkan may use device memory for other purposes, such as other types of device objects, internal allocations and data structures, scratch storage, and so on. These allocations are managed by the Vulkan driver, as the requirements may vary quite widely between implementations.

When you are done with a memory allocation, you need to free it. To do this, call `vkFreeMemory()`, the prototype of which is

Click here to view code image

```
void vkFreeMemory (
    VkDevice                            device,
    VkDeviceMemory                      memory,
    const VkAllocationCallbacks*        pAllocator);
```

`vkFreeMemory()` takes the memory object directly in `memory`. It is your responsibility to ensure that there is no work queued up to a device that might use the memory object before you free it. Vulkan will not track this for you. If a device attempts to access memory after it's been freed, the results can be unpredictable and can easily crash your application.

Further, access to `memory` must be externally synchronized. Attempting to free device memory with a call to **vkFreeMemory()** while another command is executing in another thread will produce undefined behavior and possibly crash your application.

On some platforms, there may be an upper bound to the total number of memory allocations that can exist within a single process. If you try to create more allocations than this limit, allocation could fail. This limit can be determined by calling **vkGetPhysicalDeviceProperties()** and inspecting the `maxMemoryAllocationCount` field of the returned `VkPhysicalDeviceLimits` structure. The limit is guaranteed to be at least 4,096 allocations, though some platforms may report a much higher limit. Although this may seem low, the intention is that you create a small number of large allocations and then suballocate from them to place many resources in the same allocation. There is no upper limit to the total number of resources can be created, memory allowing.

Normally, when you allocate memory from a heap, that memory is permanently assigned to the returned `VkDeviceMemory` object until that object is destroyed by calling **vkFreeMemory()**. In some cases, you (or even the Vulkan implementation) may not know exactly how much memory is required for certain operations, or indeed whether any memory is required at all.

In particular, this is often the case for images that are used for intermediate storage of data during rendering. When the image is created, if the `VK_IMAGE_USAGE_TRANSIENT_ATTACHMENT_BIT` is included in the `VkImageCreateInfo` structure, then Vulkan knows that the data in the image will live for a short time, and therefore, it's possible that it may never need to be written out to device memory.

In this case, you can ask Vulkan to be *lazy* with its allocation of the memory object to defer true allocation until Vulkan can determine that the physical storage for data is really needed. To do this, choose a memory type with the `VK_MEMORY_PROPERTY_LAZILY_ALLOCATED_BIT` set. Choosing an otherwise-appropriate memory type that does not have this bit set will still work correctly but will always allocate the memory up front, even if it never ends up being used.

If you want to know whether a memory allocation is physically backed and how much backing has actually been allocated for a memory object, call **vkGetDeviceMemoryCommitment()**, the prototype of which is

[**Click here to view code image**](#)

```
void vkGetDeviceMemoryCommitment (
    VkDevice                              device,
    VkDeviceMemory                        memory,
    VkDeviceSize*                         pCommittedMemoryInBytes);
```

The device that owns the memory allocation is passed in `device` and the memory allocation to query is passed in `memory`. `pCommittedMemoryInBytes` is a pointer to a variable that will be overwritten with the number of bytes actually allocated for the memory object. That commitment will always come from the heap associated with the memory type used to allocate the memory object.

For memory objects allocated with memory types that don't include `VK_MEMORY_PROPERTY_LAZILY_ALLOCATED_BIT`, or if the memory object ended up fully committed, **vkGetDeviceMemoryCommitment()** will always return the full size of the memory object. The commitment returned from **vkGetDeviceMemoryCommitment()** is informational at best. In many cases, the information could be out of date, and there's not much you can do with the information anyway.

Host Access to Device Memory

As discussed earlier in this chapter, device memory is divided into multiple regions. Pure device memory is accessible only to the device. However, there are regions of memory that are accessible to both the host and the device. The host is the processor upon which your application is running, and it is possible to ask Vulkan to give you a pointer to memory allocated from host-accessible regions. This is known as *mapping* memory.

To map device memory into the host's address space, the memory object to be mapped must have been allocated from a heap that has the VK_MEMORY_PROPERTY_HOST_VISIBLE_BIT flag set in its heap properties. Assuming that this is the case, mapping the memory to obtain a pointer usable by the host is achieved by calling **vkMapMemory()**, the prototype of which is

[Click here to view code image](#)

```
VkResult vkMapMemory (
    VkDevice                                device,
    VkDeviceMemory                          memory,
    VkDeviceSize                            offset,
    VkDeviceSize                            size,
    VkMemoryMapFlags                        flags,
    void**                                  ppData);
```

The device that owns the memory object to be mapped is passed in device, and the handle to the memory object being mapped is passed in memory. Access to the memory object must be externally synchronized. To map a range of a memory object, specify the starting offset in offset and the size of the region in size. If you want to map the entire memory object, set offset to 0 and size to VK_WHOLE_SIZE. Setting offset to a nonzero value and size to VK_WHOLE_SIZE will map the memory object starting from offset to the end. offset and size are both specified in bytes. You should not attempt to map a region of the memory object that extends beyond its bounds.

The flags parameter is reserved for future use and should be set to zero.

If **vkMapMemory()** is successful, a pointer to the mapped region is written into the variable pointed to by ppData. This pointer can then be cast to the appropriate type in your application and dereferenced to directly read and write the device memory. Vulkan guarantees that pointers returned from **vkMapMemory()** are aligned to an integer multiple of the device's minimum memory mapping alignment when offset is subtracted from them.

This value is reported in the minMemoryMapAlignment field of the VkPhysicalDeviceLimits structure returned from a call to **vkGetPhysicalDeviceProperties()**. It is guaranteed to be at least 64 bytes but could be any higher power of two. On some CPU architectures, much higher performance can be achieved by using memory load and store instructions that assume aligned addresses. minMemoryMapAlignment will often match a cache line size or the natural alignment of the machine's widest register, for example, to facilitate this. Some host CPU instructions will fault if passed an unaligned address. Therefore, you can check minMemoryMapAlignment once and decide whether to use optimized functions that assume aligned addressing or fallback functions that can handle unaligned addresses at the expense of performance.

When you're done with the pointer to the mapped memory range, it can be unmapped by calling **vkUnmapMemory()**, the prototype of which is

```
void vkUnmapMemory (
    VkDevice                            device,
    VkDeviceMemory                      memory);
```

The device that owns the memory object is passed in `device`, and the memory object to be unmapped is passed in `memory`. As with **vkMapMemory()**, access to the memory object must be externally synchronized.

It's not possible to map the same memory object more than once at the same time. That is, you can't call **vkMapMemory()** on the same memory object with different memory ranges, whether they overlap or not, without unmapping the memory object in between. The range isn't needed when unmapping the object because Vulkan knows the range that was mapped.

As soon as the memory object is unmapped, any pointer received from a call to **vkMapMemory()** is invalid and should not be used. Also, if you map the same range of the same memory object over and over, you shouldn't assume that the pointer you get back will be the same.

When device memory is mapped into host address space, there are effectively two clients of that memory, which may both perform writes into it. There is likely to be a cache hierarchy on both the host and the device sides of the mapping, and those caches may or may not be coherent. In order to ensure that both the host and the device see a coherent view of data written by the other client, it is necessary to force Vulkan to flush caches that might contain data written by the host but not yet made visible to the device or to invalidate a host cache that might hold stale data that has been overwritten by the device.

Each memory type advertised by the device has a number of properties, one of which might be `VK_MEMORY_PROPERTY_HOST_COHERENT_BIT`. If this is the case, and a mapping is made from a region with this property set, then Vulkan will take care of coherency between caches. In some cases, the caches are automatically coherent because they are either shared between host and device or have some form of coherency protocol to keep them in sync. In other cases, a Vulkan driver might be able to infer when caches need to be flushed or invalidated and then perform these operations behind the scenes.

If `VK_MEMORY_PROPERTY_HOST_COHERENT_BIT` is not set in the memory properties of a mapped memory region, then it is your responsibility to explicitly flush or invalidate caches that might be affected by the mapping. To flush host caches that might contain pending writes, call **vkFlushMappedMemoryRanges()**, the prototype of which is

```
VkResult vkFlushMappedMemoryRanges (
    VkDevice                            device,
    uint32_t                            memoryRangeCount,
    const VkMappedMemoryRange*          pMemoryRanges);
```

The device that owns the mapped memory objects is specified in `device`. The number of ranges to flush is specified in `memoryRangeCount`, and the details of each range are passed in an instance of the `VkMappedMemoryRange` structure. A pointer to an array of `memoryRangeCount` of these structures is passed through the `pMemoryRanges` parameter. The definition of `VkMappedMemoryRange` is

```
typedef struct VkMappedMemoryRange {
    VkStructureType     sType;
    const void*         pNext;
    VkDeviceMemory      memory;
    VkDeviceSize        offset;
    VkDeviceSize        size;
} VkMappedMemoryRange;
```

The `sType` field of `VkMappedMemoryRange` should be set to `VK_STRUCTURE_TYPE_MAPPED_MEMORY_RANGE`, and `pNext` should be set to `nullptr`. Each memory range refers to a mapped memory object specified in the `memory` field and a mapped range within that object, specified by `offset` and `size`. You don't have to flush the entire mapped region of the memory object, so `offset` and `size` don't need to match the parameters used in `vkMapMemory()`. Also, if the memory object is not mapped, or if `offset` and `size` specify a region of the object that isn't mapped, then the flush command has no effect. To just flush any existing mapping on a memory object, set `offset` to zero and `size` to `VK_WHOLE_SIZE`.

A flush is necessary if the host has written to a mapped memory region and needs the device to see the effect of those writes. However, if the device writes to a mapped memory region and you need the host to see the effect of the device's writes, you need to invalidate any caches on the host that might now hold stale data. To do this, call **vkInvalidateMappedMemoryRanges()**, the prototype of which is

```
VkResult vkInvalidateMappedMemoryRanges (
    VkDevice                              device,
    uint32_t                              memoryRangeCount,
    const VkMappedMemoryRange*            pMemoryRanges);
```

As with **vkFlushMappedMemoryRanges()**, `device` is the device that owns the memory objects whose mapped regions are to be invalidated. The number of regions is specified in `memoryRangeCount`, and a pointer to an array of `memoryRangeCount` `VkMappedMemoryRange` structures is passed in `pMemoryRanges`. The fields of the `VkMappedMemoryRange` structures are interpreted exactly as they are in **vkFlushMappedMemoryRanges()**, except that the operation performed is an invalidation rather than a flush.

vkFlushMappedMemoryRanges() and **vkInvalidateMappedMemoryRanges()** affect only caches and coherency of access by the *host* and have no effect on the device. Regardless of whether a memory mapping is coherent or not, access by the device to memory that has been mapped must still be synchronized using *barriers*, which will be discussed later in this chapter.

Binding Memory to Resources

Before a resource such as a buffer or image can be used by Vulkan to store data, memory must be bound to it. Before memory is bound to a resource, you should determine what *type* of memory and how much of it the resource requires. There is a different function for buffers and for textures. They are **vkGetBufferMemoryRequirements()** and **vkGetImageMemoryRequirements()**, and their prototypes are

[**Click here to view code image**](#)

```
void vkGetBufferMemoryRequirements (
    VkDevice                        device,
    VkBuffer                        buffer,
    VkMemoryRequirements*           pMemoryRequirements);
```

and

[**Click here to view code image**](#)

```
void vkGetImageMemoryRequirements (
    VkDevice                        device,
    VkImage                         image,
    VkMemoryRequirements*           pMemoryRequirements);
```

The only difference between these two functions is that **vkGetBufferMemoryRequirements()** takes a handle to a buffer object and **vkGetImageMemoryRequirements()** takes a handle to an image object. Both functions return the memory requirements for the resource in an instance of the VkMemoryRequirements structure, the address of which is passed in the pMemoryRequirements parameter. The definition of VkMemoryRequirements is

[**Click here to view code image**](#)

```
typedef struct VkMemoryRequirements {
    VkDeviceSize   size;
    VkDeviceSize   alignment;
    uint32_t       memoryTypeBits;
} VkMemoryRequirements;
```

The amount of memory needed by the resource is placed in the size field, and the alignment requirements of the object are placed in the alignment field. When you bind memory to the object (which we will get to in a moment), you need to ensure that the offset from the start of the memory object meets the alignment requirements of the resource and that there is sufficient space in the memory object to store the object.

The memoryTypeBits field is populated with all the memory types that the resource can be bound to. One bit is turned on, starting from the least significant bit, for each type that can be used with the resource. If you have no particular requirements for the memory, simply find the lowest-set bit and use its index to choose the memory type, which is then used as the memoryTypeIndex field in the allocation info passed to a call to **vkAllocateMemory()**. If you do have particular requirements or preferences for the memory—if you want to be able to map the memory or prefer that it be host local, for example—look for a type that includes those bits and is supported by the resource.

shows an example of an appropriate algorithm for choosing the memory type for an image resource.

Listing 2.5: Choosing a Memory Type for an Image

Click here to view code image

```cpp
uint32_t application::chooseHeapFromFlags(
    const VkMemoryRequirements& memoryRequirements,
    VkMemoryPropertyFlags requiredFlags,
    VkMemoryPropertyFlags preferredFlags)

{
    VkPhysicalDeviceMemoryProperties deviceMemoryProperties;

    vkGetPhysicalDeviceMemoryProperties(m_physicalDevices[0],
                                        &deviceMemoryProperties);

    uint32_t selectedType = ~0u;
    uint32_t memoryType;

    for (memoryType = 0; memoryType < 32; ++memoryType)
    {
        if (memoryRequirements.memoryTypeBits & (1 << memoryType))
        {
            const VkMemoryType& type =
                deviceMemoryProperties.memoryTypes[memoryType];

            // If it exactly matches my preferred properties, grab it.
            if ((type.propertyFlags & preferredFlags) == preferredFlags)
            {
                selectedType = memoryType;
                break;
            }
        }
    }
    if (selectedType != ~0u)
    {
        for (memoryType = 0; memoryType < 32; ++memoryType)
        {
            if (memoryRequirements.memoryTypeBits & (1 << memoryType))
            {
                const VkMemoryType& type =
                    deviceMemoryProperties.memoryTypes[memoryType];

                // If it has all my required properties, it'll do.
                if ((type.propertyFlags & requiredFlags) == requiredFlags)
                {
                    selectedType = memoryType;
                    break;
                }
            }
        }
    }
```

```
    }

    return selectedType;
}
```

The algorithm shown in [Listing 2.5](#) chooses a memory type given the memory requirements for an object, a set of hard requirements, and a set of preferred requirements. First, it iterates through the device's supported memory types and checks each for the set of preferred flags. If there is a memory type that contains all of the flags that the caller prefers, then it immediately returns that memory type. If none of the device's memory types exactly matches the preferred flags, then it iterates again, this time returning the first memory type that meets all of the requirements.

Once you have chosen the memory type for the resource, you can bind a piece of a memory object to that resource by calling either **vkBindBufferMemory()** for buffer objects or **vkBindImageMemory()** for image objects. Their prototypes are

[**Click here to view code image**](#)

```
VkResult vkBindBufferMemory (
    VkDevice                            device,
    VkBuffer                            buffer,
    VkDeviceMemory                      memory,
    VkDeviceSize                        memoryOffset);
```

and

[**Click here to view code image**](#)

```
VkResult vkBindImageMemory (
    VkDevice                            device,
    VkImage                             image,
    VkDeviceMemory                      memory,
    VkDeviceSize                        memoryOffset);
```

Again, these two functions are identical in declaration except that **vkBindBufferMemory()** takes a VkBuffer handle and **vkBindImageMemory()** takes a VkImage handle. In both cases, device must own both the resource and the memory object, whose handle is passed in memory. This is the handle of a memory allocation created through a call to **vkAllocateMemory()**.

Access to buffer and image from **vkBindBufferMemory()** and **vkBindImageMemory()**, respectively, must be externally synchronized. Once memory has been bound to a resource object, the memory binding cannot be changed again. If two threads attempt to execute **vkBindBufferMemory()** or **vkBindImageMemory()** concurrently, then which thread's binding takes effect and which one is invalid is subject to a race condition. Even resolving the race condition would not produce a legal command sequence, so this should be avoided.

The memoryOffset parameter specifies *where* in the memory object the resource will live. The amount of memory consumed by the object is determined from the size of the object's requirements, as discovered with a call to **vkGetBufferMemoryRequirements()** or **vkGetImageMemoryRequirements()**.

It is very strongly recommended that rather than simply creating a new memory allocation for each resource, you create a pool of a small number of relatively large memory allocations and place multiple resources in each one at different offsets. It is possible for two resources to overlap in

memory. In general, aliasing data like this is not well defined, but if you can be sure that two resources are not used at the same time, this can be a good way to reduce the memory requirements of your application.

An example of a device memory allocator is included with the book's source code.

Sparse Resources

Sparse resources are a special type of resource that can be partially backed by memory and can have their memory backing changed after they have been created and even used in the application. A sparse resource must still be bound to memory before it can be used, although that binding can be changed. Additionally, an image or buffer can support sparse residency, which allows parts of the image to *not* be backed by memory at all.

To create a sparse image, set the `VK_IMAGE_CREATE_SPARSE_BINDING_BIT` in the `flags` field of the `VkImageCreateInfo` structure used to create the image. Likewise, to create a sparse buffer, set the `VK_BUFFER_CREATE_SPARSE_BINDING_BIT` in the `flags` field of the `VkBufferCreateInfo` structure used to create the buffer.

If an image was created with the `VK_IMAGE_CREATE_SPARSE_BINDING_BIT` bit set, your application should call **vkGetImageSparseMemoryRequirements()** to determine the additional requirements that the image needs. The prototype of **vkGetImageSparseMemoryRequirements()** is

[**Click here to view code image**](#)

```
void vkGetImageSparseMemoryRequirements (
    VkDevice                             device,
    VkImage                              image,
    uint32_t*                            pSparseMemoryRequirementCount,
    VkSparseImageMemoryRequirements*     pSparseMemoryRequirements);
```

The device that owns the image should be passed in `device`, and the image whose requirements to query should be passed in `image`. The `pSparseMemoryRequirements` parameter points to an array of `VkSparseImageMemoryRequirements` structures that will be filled with the requirements of the image.

If `pSparseMemoryRequirements` is `nullptr`, then the initial value of the variable pointed to by `pSparseMemoryRequirementCount` is ignored and is overwritten with the number of requirements of the image. If `pSparseMemoryRequirements` is not `nullptr`, then the initial value of the variable pointed to by `pSparseMemoryRequirementCount` is the number of elements in the `pSparseMemoryRequirements` array and is overwritten with the number of requirements actually written to the array.

The definition of `VkSparseImageMemoryRequirements` is

[**Click here to view code image**](#)

```
typedef struct VkSparseImageMemoryRequirements {
    VkSparseImageFormatProperties    formatProperties;
    uint32_t                         imageMipTailFirstLod;
    VkDeviceSize                     imageMipTailSize;
    VkDeviceSize                     imageMipTailOffset;
```

```
    VkDeviceSize                        imageMipTailStride;
} VkSparseImageMemoryRequirements;
```

The first field of `VkSparseImageMemoryRequirements` is an instance of the
`VkSparseImageFormatProperties` structure that provides general information about how the
image is laid out in memory with respect to binding.

[Click here to view code image](#)

```
typedef struct VkSparseImageFormatProperties {
    VkImageAspectFlags          aspectMask;
    VkExtent3D                  imageGranularity;
    VkSparseImageFormatFlags    flags;
} VkSparseImageFormatProperties;
```

The `aspectMask` field of `VkSparseImageFormatProperties` is a bitfield indicating the
image aspects to which the properties apply. This will generally be all of the aspects in the image. For
color images, it will be `VK_IMAGE_ASPECT_COLOR_BIT`, and for depth, stencil, and depth-stencil
images, it will be either or both of `VK_IMAGE_ASPECT_DEPTH_BIT` and
`VK_IMAGE_ASPECT_STENCIL_BIT`.

When memory is bound to a sparse image, it is bound in blocks rather than to the whole resource at
once. Memory has to be bound in implementation-specific sized blocks, and the
`imageGranularity` field of `VkSparseImageFormatProperties` contains this size.

Finally, the `flags` field contains some addional flags describing further behavior of the image. The
flags that may be included are

- `VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT`: If this bit is set and the image is
  an array, then the mip tail shares a binding shared by all array layers. If the bit is not set, then
  each array layer has its own mip tail that can be bound to memory independently of others.

- `VK_SPARSE_IMAGE_FORMAT_ALIGNED_MIP_SIZE_BIT`: If this bit is set, it is an
  indicator that the mip tail begins with the first level that is not a multiple of the image's binding
  granularity. If the bit is not set, then the tail begins at the first level that is smaller than the
  image's binding granularity.

- `VK_SPARSE_IMAGE_FORMAT_NONSTANDARD_BLOCK_SIZE_BIT`: If this bit is set, then
  the image's format does support sparse binding, but not with the standard block sizes. The
  values reported in `imageGranularity` are still correct for the image but don't necessarily
  match the standard block for the format.

Unless `VK_SPARSE_IMAGE_FORMAT_NONSTANDARD_BLOCK_SIZE_BIT` is set in `flags`,
then the values in `imageGranularity` match a set of standard block sizes for the format. The
size, in texels, of various formats is shown in [Table 2.1](#).

The remaining fields of `VkSparseImageMemoryRequirements` describe how the format used
by the image behaves in the mip tail. The mip tail is the region of the mipmap chain beginning from
the first level that cannot be sparsely bound to memory. This is typically the first level that is smaller
than the size of the format's granularity. As memory must be bound to sparse resources in units of the
granularity, the mip tail presents an all-or-nothing binding opportunity. Once any level of the
mipmap's tail is bound to memory, *all* levels within the tail become bound.

Texel Size	2D Block Shape	3D Block Shape
8-bit	256×256	$64 \times 32 \times 32$
16-bit	256×128	$32 \times 32 \times 32$
32-bit	128×128	$32 \times 32 \times 16$
64-bit	128×64	$32 \times 16 \times 16$
128-bit	64×64	$16 \times 16 \times 16$

Table 2.1: Sparse Texture Block Sizes

The mip tail begins at the level reported in the `imageMipTailFirstLod` field of `VkSparseImageMemoryRequirements`. The size of the tail, in bytes, is contained in `imageMipTailSize`, and it begins at `imageMipTailOffset` bytes into the image's memory binding region. If the image does not have a single mip tail binding for all array layers (as indicated by the presence of `VK_SPARSE_IMAGE_FORMAT_SINGLE_MIPTAIL_BIT` in the `aspectMask` field of `VkSparseImageFormatProperties`), then `imageMipTailStride` is the distance, in bytes, between the start of the memory binding for each mip tail level.

The properties of a specific format can also be determined by calling **vkGetPhysicalDeviceSparseImageFormatProperties()**, which, given a specific format, will return a `VkSparseImageFormatProperties` describing that format's sparse image requirements without the need to create an image and query it. The prototype of **vkGetPhysicalDeviceSparseImageFormatProperties()** is

[**Click here to view code image**](#)

```
void vkGetPhysicalDeviceSparseImageFormatProperties (
    VkPhysicalDevice                    physicalDevice,
    VkFormat                            format,
    VkImageType                         type,
    VkSampleCountFlagBits               samples,
    VkImageUsageFlags                   usage,
    VkImageTiling                       tiling,
    uint32_t*                           pPropertyCount,
    VkSparseImageFormatProperties*      pProperties);
```

As you can see, **vkGetPhysicalDeviceSparseImageFormatProperties()** takes as parameters many of the properties that would be used to construct an image. Sparse image properties are a function of a physical device, a handle to which should be passed in `physicalDevice`. The format of the image is passed in `format`, and the type of image (`VK_IMAGE_TYPE_1D`, `VK_IMAGE_TYPE_2D`, or `VK_IMAGE_TYPE_3D`) is passed in `type`. If multisampling is required, the number of samples (represented as one of the members of the `VkSampleCountFlagBits` enumeration) is passed in `samples`.

The intended usage for the image is passed in `usage`. This should be a bitfield containing the flags specifying how an image with this format will be used. Be aware that sparse images may not be supported at all under certain use cases, so it's best to set this field conservatively and accurately rather than just turning on every bit and hoping for the best. Finally, the tiling mode to be used for the image is specified in `tiling`. Again, standard block sizes may be supported only in certain tiling

modes. For example, it's very unlikely that an implementation would support standard (or even reasonable) block sizes when LINEAR tiling is used.

Just as with **vkGetPhysicalDeviceImageFormatProperties()**, **vkGetPhysicalDeviceSparseImageFormatProperties()** can return an array of properties. The pPropertyCount parameter points to a variable that will be overwritten with the number of properties reported for the format. If pProperties is nullptr, then the initial value of the variable pointed to by pPropertyCount is ignored and the total number of properties is written into it. If pProperties is not nullptr, then it should be a pointer to an array of VkSparseImageFormatProperties structures that will receive the properties of the image. In this case, the initial value of the variable pointed to by pPropertyCount is the number of elements in the array, and it is overwritten with the number of items populated in the array.

Because the memory binding used to back sparse images can be changed, even after the image is in use, the update to the binding properties of the image is pipelined along with that work. Unlike **vkBindImageMemory()** and **vkBindBufferMemory()**, which are operations likely carried out by the host, memory is bound to a sparse resource using an operation on the queue, allowing the device to execute them. The command to bind memory to a sparse resource is **vkQueueBindSparse()**, the prototype of which is

Click here to view code image

```
VkResult vkQueueBindSparse (
    VkQueue                       queue,
    uint32_t                      bindInfoCount,
    const VkBindSparseInfo*       pBindInfo,
    VkFence                       fence);
```

The queue that will execute the binding operation is specified in queue. Several binding operations can be performed by a single call to **vkQueueBindSparse()**. The number of operations to perform is passed in bindInfoCount, and pBindInfo is a pointer to an array of bindInfoCount VkBindSparseInfo structures, each describing one of the bindings. The definition of VkBindSparseInfo is

Click here to view code image

```
typedef struct VkBindSparseInfo {
    VkStructureType                          sType;
    const void*                              pNext;
    uint32_t                                 waitSemaphoreCount;
    const VkSemaphore*                       pWaitSemaphores;
    uint32_t                                 bufferBindCount;
    const VkSparseBufferMemoryBindInfo*      pBufferBinds;
    uint32_t                                 imageOpaqueBindCount;
    const VkSparseImageOpaqueMemoryBindInfo* pImageOpaqueBinds;
    uint32_t                                 imageBindCount;
    const VkSparseImageMemoryBindInfo*       pImageBinds;
    uint32_t                                 signalSemaphoreCount;
    const VkSemaphore*                       pSignalSemaphores;
} VkBindSparseInfo;
```

The act of binding memory to sparse resources is actually pipelined with other work performed by the device. As you read in Chapter 1, "Overview of Vulkan," work is performed by submitting it to queues. The binding is then performed along with the execution of commands submitted to the same queue. Because **vkQueueBindSparse()** behaves a lot like a command submission, VkBindSparseInfo contains many fields related to synchronization.

The sType field of VkBindSparseInfo should be set to VK_STRUCTURE_TYPE_BIND_SPARSE_INFO, and pNext should be set to nullptr. As with VkSubmitInfo, each sparse binding operation can optionally wait for one or more semaphores to be signaled before performing the operation and can signal one or more semaphores when it is done. This allows updates to the sparse resource's bindings to be synchronized with other work performed by the device.

The number of semaphores to wait on is specified in waitSemaphoreCount, and the number of semaphores to signal is specified in signalSemaphoreCount. The pWaitSemaphores field is a pointer to an array of waitSemaphoreCount semaphore handles to wait on, and pSignalSemaphores is a pointer to an array of signalSemaphoreCount semaphores to signal. Semaphores are covered in some detail in Chapter 11, "Synchronization."

Each binding operation can include updates to buffers and images. The number of buffer binding updates is specified in bufferBindCount and pBufferBinds is a pointer to an array of bufferBindCount VkSparseBufferMemoryBindInfo structures, each describing one of the buffer memory binding operations. The definition of VkSparseBufferMemoryBindInfo is

Click here to view code image

```
typedef struct VkSparseBufferMemoryBindInfo {
    VkBuffer                    buffer;
    uint32_t                    bindCount;
    const VkSparseMemoryBind*   pBinds;
} VkSparseBufferMemoryBindInfo;
```

Each instance of VkSparseBufferMemoryBindInfo contains the handle of the buffer to which memory will be bound. A number of regions of memory can be bound to the buffer at different offsets. The number of memory regions is specified in bindCount, and each binding is described by an instance of the VkSparseMemoryBind structure. pBinds is a pointer to an array of bindCount VkSparseMemoryBind structures. The definition of VkSparseMemoryBind is

Click here to view code image

```
typedef struct VkSparseMemoryBind {
    VkDeviceSize            resourceOffset;
    VkDeviceSize            size;
    VkDeviceMemory          memory;
    VkDeviceSize            memoryOffset;

    VkSparseMemoryBindFlags flags;
} VkSparseMemoryBind;
```

The size of the block of memory to bind to the resource is contained in size. The offsets of the block in the resource and in the memory object are contained in resourceOffset and memoryOffset, respectively, and are both expressed in units of bytes. The memory object that is

the source of storage for the binding is specified in `memory`. When the binding is executed, the block of memory, `size` bytes long and starting at `memoryOffset` bytes into the memory object specified in `memory`, will be bound into the buffer specified in the `buffer` field of the `VkSparseBufferMemoryBindInfo` structure.

The `flags` field contains additional flags that can be used to further control the binding. No flags are used for buffer resources. However, image resources use the same `VkSparseMemoryBind` structure to affect memory bindings directly to images. This is known as an opaque image memory binding, and the opaque image memory bindings to be performed are also passed through the `VkBindSparseInfo` structure. The `pImageOpaqueBinds` member of `VkBindSparseInfo` points to an array of `imageOpaqueBindCount` `VkSparseImageOpaqueMemoryBindInfo` structures defining the opaque memory bindings. The definition of `VkSparseImageOpaqueMemoryBindInfo` is

[Click here to view code image](#)

```
typedef struct VkSparseImageOpaqueMemoryBindInfo {
    VkImage                     image;
    uint32_t                    bindCount;
    const VkSparseMemoryBind*   pBinds;
} VkSparseImageOpaqueMemoryBindInfo;
```

Just as with `VkSparseBufferMemoryBindInfo`, `VkSparseImageOpaqueMemoryBindInfo` contains a handle to the image to which to bind memory in `image` and a pointer to an array of `VkSparseMemoryBind` structures in `pBinds`, which is `bindCount` elements long. This is the same structure used for buffer memory bindings. However, when this structure is used for images, you can set the `flags` field of each `VkSparseMemoryBind` structure to include the `VK_SPARSE_MEMORY_BIND_METADATA_BIT` flag in order to explicitly bind memory to the metadata associated with the image.

When memory is bound opaquely to a sparse image, the blocks of memory have no defined correlation with texels in the image. Rather, the backing store of the image is treated as a large, opaque region of memory with no information about how texels are laid out in it provided to the application. However, so long as memory is bound to the entire image when it is used, results will still be well-defined and consistent. This allows sparse images to be backed by multiple, smaller memory objects, potentially easing pool allocation strategies, for example.

To bind memory to an explicit region of an image, you can perform a nonopaque image memory binding by passing one or more `VkSparseImageMemoryBindInfo` structures through the `VkBindSparseInfo` structures passed to **`vkQueueBindSparse()`**. The definition of `VkSparseImageMemoryBindInfo` is

[Click here to view code image](#)

```
typedef struct VkSparseImageMemoryBindInfo {
    VkImage                         image;
    uint32_t                        bindCount;
    const VkSparseImageMemoryBind*  pBinds;
} VkSparseImageMemoryBindInfo;
```

Again, the `VkSparseImageMemoryBindInfo` structure contains a handle to the image to which to bind memory in `image`, a count of the number of bindings to perform in `bindCount`, and a pointer to an array of structures describing the bindings in `pBinds`. This time, however, `pBinds` points to an array of `bindCount` `VkSparseImageMemoryBind` structures, the definition of which is

Click here to view code image

```
typedef struct VkSparseImageMemoryBind {
    VkImageSubresource        subresource;
    VkOffset3D                offset;
    VkExtent3D                extent;
    VkDeviceMemory            memory;
    VkDeviceSize              memoryOffset;
    VkSparseMemoryBindFlags   flags;
} VkSparseImageMemoryBind;
```

The `VkSparseImageMemoryBind` structure contains much more information about how the memory is to be bound to the image resource. For each binding, the image subresource to which the memory is to be bound is specified in `subresource`, which is an instance of the `VkImageSubresource`, the definition of which is

Click here to view code image

```
typedef struct VkImageSubresource {
    VkImageAspectFlags    aspectMask;
    uint32_t              mipLevel;
    uint32_t              arrayLayer;
} VkImageSubresource;
```

The `VkImageSubresource` allows you to specify the aspect of the image (`VK_IMAGE_ASPECT_COLOR_BIT`, `VK_IMAGE_ASPECT_DEPTH_BIT`, or `VK_IMAGE_ASPECT_STENCIL_BIT`, for example) in `aspectMask`, the mipmap level to which you want to bind memory in `mipLevel`, and the array layer where the memory should be bound in `arrayLayer`. For nonarray images, `arrayLayer` should be set to zero.

Within the subresource, the `offset` and `extent` fields of the `VkSparseImageMemoryBind` structure define the offset and size of the region of texels to bind the image data to. This must be aligned to the tile-size boundaries, which are either the standard sizes as shown in Table 2.1 or the per-format block size that can be retrieved from `vkGetPhysicalDeviceSparseImageFormatProperties()`.

Again, the memory object from which to bind memory is specified in `memory`, and the offset within the memory where the backing store resides is specified in `memoryOffset`. The same flags are available in the `flags` field of `VkSparseImageMemoryBind`.

Summary

This chapter introduced you to the different types of resources that are used by Vulkan. It described how the memory used to back them is allocated and then associated with them. It also explained how you can manage the application memory used by Vulkan through the use of a custom allocator. You have seen how to move resources from state to state and how to synchronize acess to them through pipeline barriers. This enables efficient, parallel access to resources both from multiple stages of the Vulkan pipeline and from the host.

Chapter 3. Queues and Commands

What You'll Learn in This Chapter

    • What a queue is and how to use it

    • How to create commands and send them to Vulkan

    • How to ensure that a device has finished processing your work

---

Vulkan devices expose multiple queues that perform work. In this chapter, we discuss the various queue types and explain how to submit work to them in the form of command buffers. We also show how to instruct a queue to complete all of the work you've sent it.

Device Queues

Each device in Vulkan has one or more queues. The queue is the part of the device that actually performs work. It can be thought of as a subdevice that exposes a subset of the device's functionality. In some implementations, each queue may even be a physically separate part of the system.

Queues are grouped into one or more *queue families*, each containing one or more queues. Queues within a single family are essentially identical. Their capabilities are the same, their performance level and access to system resources is the same, and there is no cost (beyond synchronization) of transferring work between them. If a device contains multiple cores that have the same capabilities but differ in performance, access to memory, or some other factor that might mean they can't operate identically, it may expose them in separate families that otherwise appear identical.

As discussed in Chapter 1, "Overview of Vulkan," you can query the properties of each of a physical device's queue families by calling **vkGetPhysicalDeviceQueueFamilyProperties()**. This function writes the properties of the queue family into an instance of the `VkQueueFamilyProperties` structure that you hand it.

The number and type of queues that you wish to use must be specified when you create the device. As you saw in Chapter 1, "Overview of Vulkan," the `VkDeviceCreateInfo` structure that you pass to **vkCreateDevice()** contains the `queueCreateInfoCount` and `pQueueCreateInfos` members. Chapter 1, "Overview of Vulkan," glossed over them, but now it's time to fill them in. The `queueCreateInfoCount` member contains the number of `VkDeviceQueueCreateInfo` structures stored in the array pointed to by `pQueueCreateInfos`. The definition of the `VkDeviceQueueCreateInfo` structure is

Click here to view code image

```
typedef struct VkDeviceQueueCreateInfo {
    VkStructureType             sType;
    const void*                 pNext;
    VkDeviceQueueCreateFlags    flags;
    uint32_t                    queueFamilyIndex;
    uint32_t                    queueCount;
    const float*                pQueuePriorities;
} VkDeviceQueueCreateInfo;
```

As with most Vulkan structures, the `sType` field is the structure type, which in this case should be `VK_STRUCTURE_TYPE_QUEUE_CREATE_INFO`, and the `pNext` field is used for extensions and should be set to `nullptr` when none are used. The `flags` field contains flags controlling queue construction, but no flag is defined for use in the current version of Vulkan, so this field should be set to zero.

The fields of interest here are `queueFamilyIndex` and `queueCount`. The `queueFamilyIndex` field specifies the family from which you want to allocate queues, and the `queueCount` field specifies the number of queues to allocate from that family. To allocate queues from multiple families, simply pass an array of more than one `VkDeviceQueueCreateInfo` structure in the `pQueueCreateInfos` member of the `VkDeviceCreateInfo` structure.

The queues are constructed when the device is created. For this reason, we don't *create* queues, but obtain them from the device. To do this, call **`vkGetDeviceQueue()`**:

Click here to view code image

```
void vkGetDeviceQueue (
    VkDevice                            device,
    uint32_t                            queueFamilyIndex,
    uint32_t                            queueIndex,
    VkQueue*                            pQueue);
```

The **`vkGetDeviceQueue()`** function takes as arguments the device from which you want to obtain the queue, the family index, and the index of the queue within that family. These are specified in `device`, `queueFamilyIndex`, and `queueIndex`, respectively. The `pQueue` parameter points to the `VkQueue` handle that is to be filled with the handle to the queue. `queueFamilyIndex` and `queueIndex` must refer to a queue that was initialized when the device was created. If they do, a queue handle is placed into the variable pointed to by `pQueue`; otherwise, this variable is set to `VK_NULL_HANDLE`.

Creating Command Buffers

The primary purpose of a queue is to process work on behalf of your application. Work is represented as a sequence of commands that are recorded into *command buffers*. Your application will create command buffers containing the work it needs to do and *submit* them to one of the queues for execution. Before you can record any commands, you need to create a command buffer. Command buffers themselves are not created directly, but allocated from pools. To create a pool, call **`vkCreateCommandPool()`**, whose prototype is

Click here to view code image

```
VkResult vkCreateCommandPool (
    VkDevice                            device,
    const VkCommandPoolCreateInfo*      pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkCommandPool*                      pCommandPool);
```

As with most Vulkan object creation functions, the first parameter, `device`, is the handle to the device that will own the new pool object, and a description of the pool is passed via a structure, a pointer to which is placed in `pCreateInfo`. This structure is an instance of `VkCommandPoolCreateInfo`, the definition of which is

```
typedef struct VkCommandPoolCreateInfo {
    VkStructureType               sType;
    const void*                   pNext;
    VkCommandPoolCreateFlags      flags;
    uint32_t                      queueFamilyIndex;
} VkCommandPoolCreateInfo;
```

As with most Vulkan structures, the first two fields, `sType` and `pNext`, contain the structure type and a pointer to another structure containing more information about the pool to be created. Here, we'll set `sType` to `VK_STRUCTURE_TYPE_COMMAND_POOL_CREATE_INFO` and, because we're not passing any extra information, set `pNext` to `nullptr`.

The `flags` field contains flags that determine the behavior of the pool and the command buffers that are allocated from it. These are members of the `VkCommandPoolCreateFlagBits` enumeration, and there are currently two flags defined for use here.

- Setting the `VK_COMMAND_POOL_CREATE_TRANSIENT_BIT` indicates that command buffers taken from the pool will be short-lived and returned to the pool shortly after use. Not setting this bit suggests to Vulkan that you might keep the command buffers around for some time.

- Setting the `VK_COMMAND_POOL_CREATE_RESET_COMMAND_BUFFER_BIT` allows individual command buffers to be reused by resetting them or restarting them. (Don't worry, we'll cover that in a moment.) If this bit is not specified, then only the pool itself can be reset, which implicitly recycles all of the command buffers allocated from it.

Each of these bits may add some overhead to the work done by a Vulkan implementation to track the resources or otherwise alter its allocation strategy. For example, setting `VK_COMMAND_POOL_CREATE_TRANSIENT_BIT` may cause a Vulkan implementation to employ a more advanced allocation strategy for the pool in order to avoid fragmentation as command buffers are frequently allocated and then returned to it. Setting `VK_COMMAND_POOL_CREATE_RESET_COMMAND_BUFFER_BIT` may cause the implementation to track the reset status of each command buffer rather than simply track it at the pool level.

In this case, we're actually going to set both bits. This gives us the most flexibility, possibly at the expense of some performance in cases where we could have managed command buffers in bulk.

Finally, the `queueFamilyIndex` field of `VkCommandPoolCreateInfo` specifies the family of queues to which command buffers allocated from this pool will be submitted. This is necessary because even where two queues on a device have the same capabilities and support the same set of commands, issuing a particular command to one queue may work differently from issuing that same command to another queue.

The `pAllocator` parameter is used for application-managed host memory allocations, which is covered in Chapter 2, "Memory and Resources." Assuming successful creation of the command pool, its handle will be written into the variable pointed to by `pCommandPool`, and **vkCreateCommandPool()** will return `VK_SUCCESS`.

Once we have a pool from which to allocate command buffers, we can grab new command buffers by calling **vkAllocateCommandBuffers()**, which is defined as

```
VkResult vkAllocateCommandBuffers (
    VkDevice                                device,
    const VkCommandBufferAllocateInfo*      pAllocateInfo,
    VkCommandBuffer*                        pCommandBuffers);
```

The device used to allocate the command buffers is passed in `device`, and the remaining parameters describing the command buffers to allocate are passed in an instance of the `VkCommandBufferAllocateInfo` structure, the address of which is passed in `pCommandBuffers`. The definition of `VkCommandBufferAllocateInfo` is

```
typedef struct VkCommandBufferAllocateInfo {
    VkStructureType        sType;
    const void*            pNext;
    VkCommandPool          commandPool;
    VkCommandBufferLevel   level;
    uint32_t               commandBufferCount;
} VkCommandBufferAllocateInfo;
```

The `sType` field should be set to `VK_STRUCTURE_TYPE_COMMAND_BUFFER_ALLOCATE_INFO`, and as we're using only the core feature set here, we set the `pNext` parameter to `nullptr`. A handle to the command pool that we created earlier is placed into the `commandPool` parameter.

The `level` parameter specifies the *level* of the command buffers that we want to allocate. It can be set to either `VK_COMMAND_BUFFER_LEVEL_PRIMARY` or `VK_COMMAND_BUFFER_LEVEL_SECONDARY`. Vulkan allows *primary* command buffers to call *secondary* command buffers. For our first few examples, we will use only primary-level command buffers. We'll cover secondary-level command buffers later in the book.

Finally, `commandBufferCount` specifies the number of command buffers that we want to allocate from the pool. Note that we don't tell Vulkan anything about the length or size of the command buffers we're creating. The internal data structures representing device commands will generally vary too greatly for any unit of measurement, such as bytes or commands, to make much sense. Vulkan will manage the command buffer memory for you.

If **vkAllocateCommandBuffers()** is successful, it will return `VK_SUCCESS` and place the handles to the allocated command buffers in the array pointed to by `pCommandBuffers`. This array should be big enough to hold all the handles. Of course, if you want to allocate only a single command buffer, you can point this at a regular `VkCommandBuffer` handle.

To free command buffers, we use the **vkFreeCommandBuffers()** command, which is declared as

```
void vkFreeCommandBuffers (
    VkDevice                    device,
    VkCommandPool               commandPool,
    uint32_t                    commandBufferCount,
    const VkCommandBuffer*      pCommandBuffers);
```

The `device` parameter is the device that owns the pool from which the command buffers were allocated. `commandPool` is a handle to that pool, `commandBufferCount` is the number of command buffers to free, and `pCommandBuffers` is a pointer to an array of `commandBufferCount` handles to the command buffers to free. Note that freeing a command buffer doesn't necessarily free all of the resources associated with it but returns them to the pool from which they were allocated.

To free all of the resources used by a command pool and all of the command buffers allocated from it, call **`vkDestroyCommandPool()`**, the prototype of which is

[**Click here to view code image**](#)

```
void vkDestroyCommandPool (
    VkDevice                            device,
    VkCommandPool                       commandPool,
    const VkAllocationCallbacks*        pAllocator;
```

The device that owns the command pool is passed in the `device` parameter, and a handle to the command pool to destroy is passed in `commandPool`. A pointer to a host memory allocation structure compatible with the one used to create the pool is passed in `pAllocator`. This parameter should be `nullptr` if the `pAllocator` parameter to **`vkCreateCommandPool()`** was also `nullptr`.

There is no need to explicitly free all of the command buffers allocated from a pool before destroying the pool. The command buffers allocated from the pool are all freed as a part of destroying the pool and freeing its resources. Care should be taken, however, that no command buffers allocated from the pool are still executing or queued for execution on the device when **`vkDestroyCommandPool()`** is called.

Recording Commands

Commands are recorded into command buffers using Vulkan command functions, all of which take a command buffer handle as their first parameter. Access to the command buffer must be externally synchronized, meaning that it is the responsibility of your application to ensure that no two threads simultaneously attempt to record commands into the same command buffer at the same time. However, the following is perfectly acceptable:

- One thread can record commands into multiple command buffers by simply calling command buffer functions on different command buffers in succession.

- Two or more threads can participate in building a single command buffer, so long as the application can guarantee that no two of them are ever executing a command buffer building function concurrently.

One of the key design principles of Vulkan is to enable efficient multithreading. To achieve this, it is important that your application's threads do not block each other's execution by, for example, taking a mutex to protect a shared resource. For this reason, it's best to have one or more command buffers for each thread rather than to try sharing one. Further, as command buffers are allocated from pools, you can go further and create a command pool for each thread, allowing command buffers to be allocated by your worker threads from their respective pools without interacting.

Before you can start recording commands into a command buffer, however, you have to begin the command buffer, which resets it to an initial state. To do this, call **vkBeginCommandBuffer()**, the prototype of which is

```
VkResult vkBeginCommandBuffer (
    VkCommandBuffer                         commandBuffer,
    const VkCommandBufferBeginInfo*      pBeginInfo);
```

The command buffer to begin recording is passed in commandBuffer, and the parameters that are used in recording this command buffer are passed through a pointer to a VkCommandBufferBeginInfo structure specified in pBeginInfo. The definition of VkCommandBufferBeginInfo is

```
typedef struct VkCommandBufferBeginInfo {
    VkStructureType                         sType;
    const void*                             pNext;
    VkCommandBufferUsageFlags               flags;
    const VkCommandBufferInheritanceInfo*  pInheritanceInfo;
} VkCommandBufferBeginInfo;
```

The sType field of VkCommandBufferBeginInfo should be set to VK_STRUCTURE_TYPE_COMMAND_BUFFER_BEGIN_INFO, and pNext should be set to nullptr. The flags field is used to tell Vulkan how the command buffer will be used. This should be a bitwise combination of some of the members of the VkCommandBufferUsageFlagBits enumeration, which include the following:

- VK_COMMAND_BUFFER_USAGE_ONE_TIME_SUBMIT_BIT means that the command buffer will be recorded, executed only once, and then destroyed or recycled.

- VK_COMMAND_BUFFER_USAGE_RENDER_PASS_CONTINUE_BIT means that the command buffer will be used inside a *renderpass* and is valid only for secondary command buffers. The flag is ignored when you create a primary command buffer, which is what we will cover in this chapter. Renderpasses and secondary command buffers are covered in more detail in Chapter 13, "Multipass Rendering."

- VK_COMMAND_BUFFER_USAGE_SIMULTANEOUS_USE_BIT means that the command buffer might be executed or pending execution more than once.

For our purposes, it's safe to set flags to zero, which means that we might execute the command buffer more than once but not simultaneously and that we're not creating a secondary command buffer.

The pInheritanceInfo member of VkCommandBufferBeginInfo is used when beginning a secondary command buffer to define which states are inherited from the primary command buffer that will call it. For primary command buffers, this pointer is ignored. We'll cover the content of the VkCommandBufferInheritanceInfo structure when we introduce secondary command buffers in Chapter 13, "Multipass Rendering."

Now it's time to create our first command. Back in Chapter 2, "Memory and Resources," you learned about buffers, images, and memory. The **vkCmdCopyBuffer()** command is used to copy data between two buffer objects. Its prototype is

Click here to view code image

```
void vkCmdCopyBuffer (
    VkCommandBuffer                    commandBuffer,
    VkBuffer                           srcBuffer,
    VkBuffer                           dstBuffer,
    uint32_t                           regionCount,
    const VkBufferCopy*                pRegions);
```

This is the general form of all Vulkan commands. The first parameter, `commandBuffer`, is the command buffer to which the command is appended. The `srcBuffer` and `dstBuffer` parameters specify the buffer objects to be used as the source and destination of the copy, respectively. Finally, an array of regions is passed to the function. The number of regions is specified in `regionCount`, and the address of the array of regions is specified in `pRegions`. Each region is represented as an instance of the `VkBufferCopy` structure, the definition of which is

Click here to view code image

```
typedef struct VkBufferCopy {
    VkDeviceSize    srcOffset;
    VkDeviceSize    dstOffset;
    VkDeviceSize    size;
} VkBufferCopy;
```

Each element of the array simply contains the source and destination offsets and the size of each region to be copied in `srcOffset`, `dstOffset`, and size, respectively. When the command is executed, for each region in `pRegions`, `size` bytes of data will be copied from `srcOffset` in `srcBuffer` to `dstOffset` in `dstBuffer`. The offsets are also measured in bytes.

One thing that is fundamental to the operation of Vulkan is that the commands are not executed as soon as they are called. Rather, they are simply added to the end of the specified command buffer. If you are copying data to or from a region of memory that is visible to the host (i.e., it's mapped), then you need to be sure of several things:

• Ensure that the data is in the source region before the command is executed by the device.

• Ensure that the data in the source region is valid until after the command has been executed on the device.

• Ensure that you don't read the destination data until after the command has been executed on the device.

The first of these is perhaps the most interesting. In particular, it means that you can build the command buffer containing the copy command *before* putting the source data in memory. So long as the source data is in the right place before the command buffer is executed, things will work out.

Listing 3.1 demonstrates how to use **vkCmdCopyBuffer()** to copy a section of data from one buffer to another. The command buffer to perform the copy with is passed in the `cmdBuffer` parameter; the source and destination buffers are passed in `srcBuffer` and `dstBuffer` parameters, respectively; and the offsets of the data within them is passed in the `srcOffset` and

`dstOffset` parameters. The function packs these parameters, along with the size of the copy, into a `VkBufferCopy` structure and calls **vkCmdCopyBuffer()** to perform the copy operation.

Listing 3.1: Example of Using **vkCmdCopyBuffer()**

[Click here to view code image](#)

```
void CopyDataBetweenBuffers(VkCmdBuffer cmdBuffer,
                            VkBuffer srcBuffer, VkDeviceSize srcOffset,
                            VkBuffer dstBuffer, VkDeviceSize dstOffset,
                            VkDeviceSize size)
{
    const VkBufferCopy copyRegion =
    {
        srcOffset, dstOffset, size
    };

    vkCmdCopyBuffer(cmdBuffer, srcBuffer, dstBuffer, 1, &copyRegion);
}
```

Remember that `srcOffset` and `dstOffset` are relative to the start of the source and destination buffers, respectively, but that each of those buffers could be bound to memory at different offsets and could potentially be bound to the same memory object. Therefore, if one of the memory objects is mapped, the offset within the memory object is the offset at which the buffer object is bound to it *plus* the offset you pass to **vkCmdCopyBuffer()**.

Before the command buffer is ready to be sent to the device for execution, we must tell Vulkan that we're done recording commands into it. To do this, we call **vkEndCommandBuffer()**, the prototype of which is

[Click here to view code image](#)

```
VkResult vkEndCommandBuffer (
    VkCommandBuffer                          commandBuffer);
```

The **vkEndCommandBuffer()** function takes only a single parameter, `commandBuffer`, which is the command buffer to end recording. After **vkEndCommandBuffer()** is executed on a command buffer, Vulkan finishes any final work it needs to do to get the command buffer ready for execution.

Recycling Command Buffers

In many applications, a similar sequence of commands is used to render all or part of each frame. Therefore, it is likely that you will record similar command buffers over and over. Using the commands introduced so far, you would call **vkAllocateCommandBuffers()** to grab one or more command buffer handles, record commands into the command buffers, and then call **vkFreeCommandBuffers()** to return the command buffers to their respective pools. This is a relatively heavyweight operation, and if you know that you will reuse a command buffer for similar work many times in a row, it may be more efficient to *reset* the command buffer. This effectively puts the command buffer back into its original state but does not necessarily interact with the pool at all. Therefore, if the command buffer dynamically allocates resources from the pool as it grows, it can hang on to those resources and avoid the cost of reallocation the second and subsequent times it's rebuilt. To reset a command buffer, call **vkResetCommandBuffer()**, the prototype of which is

[**Click here to view code image**](#)

```
VkResult vkResetCommandBuffer (
    VkCommandBuffer                      commandBuffer,
    VkCommandBufferResetFlags        flags);
```

The command buffer to reset is passed in `commandBuffer`. `flags` specifies additional operations to perform while resetting the command buffer. The only flag defined for use here is `VK_COMMAND_BUFFER_RESET_RELEASE_RESOURCES_BIT`. If this bit is set, then resources allocated from the pool by the command buffer *are* returned to the pool. Even with this bit set, it's probably still more efficient to call **vkResetCommandBuffer()** than it is to free and reallocate a new command buffer.

It's also possible to reset all the command buffers allocated from a pool in one shot. To do this, call **vkResetCommandPool()**, the prototype of which is

[**Click here to view code image**](#)

```
VkResult vkResetCommandPool (
    VkDevice                             device,
    VkCommandPool                        commandPool,
    VkCommandPoolResetFlags          flags);
```

The device that owns the command pool is specified in `device`, and the pool to reset is specified in `commandPool`. Just as with **vkResetCommandBuffer()**, the `flags` parameter specifies additional action to be taken as part of resetting the pool. Again, the only flag defined for use here is `VK_COMMAND_POOL_RESET_RELEASE_RESOURCES_BIT`. When this bit is set, any resources dynamically allocated by the pool are freed as part of the reset operation.

Command buffers allocated from the pool are *not* freed by **vkResetCommandPool()**, but all reenter their initial state as if they had been freshly allocated. **vkResetCommandPool()** is typically used at the end of a frame to return a batch of reusable command buffers to their pool rather than individually reset individual command buffers.

Care should be taken to try to keep the complexity of command buffers consistent over their multiple uses if they are reset without returning resources to the pool. As a command buffer grows, it may allocate resources dynamically from the pool, and the command pool may allocate resources from a systemwide pool. The amount of resources that a command buffer may consume is essentially

unbounded, because there is no hard-wired limit to the number of commands you can place in a single command buffer. If your application uses a mix of very small and very large command buffers, it's possible that eventually all command buffers will grow as large as the most complex command buffers.

To avoid this scenario, either periodically specify the `VK_COMMAND_BUFFER_RESET_RELEASE_RESOURCES_BIT` or `VK_COMMAND_POOL_RESET_RELEASE_RESOURCES_BIT` when resetting command buffers or their pools, respectively, or try to ensure that the same command buffers are always used in the same way—either short, simple command buffers or long, complex command buffers. Avoid mixing use cases.

Submission of Commands

To execute the command buffer on the device, we need to submit it to one of the device's queues. To do this, call **vkQueueSubmit()**, the prototype of which is

Click here to view code image

```
VkResult vkQueueSubmit (
    VkQueue                     queue,
    uint32_t                    submitCount,
    const VkSubmitInfo*         pSubmits,
    VkFence                     fence);
```

This command can submit one or more command buffers to the device for execution. The `queue` parameter specifies the device queue to which to send the command buffer. Access to the queue must be externally synchronized. All of the command buffers to submit were allocated from a pool, and that pool must have been created with respect to one of the device's queue families. This is the `queueFamilyIndex` member of the `VkCommandPoolCreateInfo` structure passed to **vkCreateCommandPool()**. queue must be a member of that family.

The number of submissions is specified in `submitCount`, and an array of structures describing each of the submissions is specified in `pSubmits`. Each submission is represented by an instance of the `VkSubmitInfo` structures, the definition of which is

Click here to view code image

```
typedef struct VkSubmitInfo {
    VkStructureType             sType;
    const void*                 pNext;
    uint32_t                    waitSemaphoreCount;
    const VkSemaphore*          pWaitSemaphores;
    const VkPipelineStageFlags* pWaitDstStageMask;
    uint32_t                    commandBufferCount;
    const VkCommandBuffer*      pCommandBuffers;
    uint32_t                    signalSemaphoreCount;
    const VkSemaphore*          pSignalSemaphores;
} VkSubmitInfo;
```

The `sType` field of `VkSubmitInfo` should be set to `VK_STRUCTURE_TYPE_SUBMIT_INFO`, and `pNext` should be set to `nullptr`. Each `VkSubmitInfo` structure can represent multiple command buffers that are to be executed by the device.

Each set of command buffers can be wrapped in a set of semaphores upon which to wait before beginning execution and can signal one or more semaphores when they complete execution. A *semaphore* is a type of synchronization primitive that allows work executed by different queues to be scheduled and coordinated correctly. We will cover semaphores along with other synchronization primitives in Chapter 11, "Synchronization." For now, we're not going to use these fields, so `waitSemaphoreCount` and `signalSemaphoreCount` can be set to zero, and `pWaitSemaphores`, `pWaitDstStageMask`, and `pSignalSemaphores` can all be set to `nullptr`.

The command buffers we want to execute are placed in an array, and its address is passed in `pCommandBuffers`. The number of command buffers to execute (the length of the `pCommandBuffers` array) is specified in `commandBufferCount`. At some time after the **vkQueueSubmit()** command is called, the commands in the command buffers begin executing on the device. Commands submitted to different queues on the same device (or to queues belonging to different devices) may execute in parallel. **vkQueueSubmit()** returns as soon as the specified command buffers have been scheduled, possibly long before they've even begun executing.

The fence parameter to **vkQueueSubmit()** is a handle to a fence object, which can be used to wait for completion of the commands executed by this submission. A *fence* is another type of synchronization primitive that we will cover in Chapter 11, "Synchronization." For now, we'll set `fence` to `VK_NULL_HANDLE`. Until we cover fences, we can wait for *all* work submitted to a queue to complete by calling **vkQueueWaitIdle()**. Its prototype is

[**Click here to view code image**](#)

```
VkResult vkQueueWaitIdle (
    VkQueue                         queue);
```

The only parameter to **vkQueueWaitIdle()**, queue, is the queue upon which to wait. When **vkQueueWaitIdle()** returns, all command buffers submitted to queue are guaranteed to have completed execution. A shortcut to wait for all commands submitted to all queues on a single device to have completed is to call **vkDeviceWaitIdle()**. Its prototype is

[**Click here to view code image**](#)

```
VkResult vkDeviceWaitIdle (
    VkDevice                        device);
```

Calling **vkQueueWaitIdle()** or **vkDeviceWaitIdle()** is really not recommended, as they fully flush any work on the queue or device and are very heavyweight operations. Neither should be called in any performance-critical part of your application. Suitable use cases include just before shutting down the application or when reinitializing application subsystems such as thread management, memory management, and so on, where there is likely to be a substantial pause anyway.

Summary

This chapter introduced you to command buffers, which are the mechanisms by which commands are communicated by your application to the Vulkan device. We introduced our first Vulkan command and showed you how to ask the device to execute work for you.

We discussed how to send the command buffers to the Vulkan device for execution by submitting them to the queue. You saw how to ensure that all work submitted to a queue or to a device has finished executing. Although we glossed over a number of important topics, such as how to call one command buffer from another and how to implement fine-grained synchronization between the host and the device and between queues on the device, these topics will be discussed in upcoming chapters.

Chapter 4. Moving Data

What You'll Learn in This Chapter

- • How to manage the state of resources as they are used by Vulkan
- • How to copy data between resources and fill buffers and images with a known value
- • How to perform blit operations to stretch and scale image data

---

Graphics and compute operations are generally data-intensive. Vulkan includes several objects that provide a means to store and manipulate data. It is often necessary to move data into and out of those objects, and several commands are provided to do exactly that: copy data and fill buffer and image objects. Further, at any given time a resource may be in one of any number of states, and many parts of the Vulkan pipeline may need access to them. This chapter covers data movement commands that can be used to copy data and fill memory—the commands needed to manage the state of resources as they are accessed by your applications.

Chapter 3, "Queues and Commands," showed that commands executed by the device are placed in command buffers and submitted to one of its queues for execution. This is important because it means that commands are not executed as you call them in your application, but as they are encoutered by the device while it makes its way through the command buffers you've submitted. The first command you were introduced to, **vkCmdCopyBuffer()**, copies data between two buffers or between different regions in the same buffer. This is one of may commands that affect buffers, images, and other objects in Vulkan. This chapter covers similar commands for filling, copying, and clearing buffers and images.

Managing Resource State

At any given time in the execution of a program, each resource can be in one of many different states. For example, if the graphics pipeline is drawing to an image or using it as the source of texture data, or if Vulkan is copying data from the host into an image, each of those usage scenarios is different. For some Vulkan implementations, there may be no real difference between some of these states, and for others, accurately knowing the state of a resource at a given point in time can make the difference between your application working or rendering junk.

Because commands in command buffers are responsible for most access to resources, and because command buffers might be built in a different order from the order in which they are submitted for execution, it's not really practical for Vulkan implementations to attempt to track the state of a resource and make sure it's in the right one for each usage scenario. In particular, a resource may begin in one state and move to another due to the execution of a command buffer. While drivers could track the state of resources as they are used in a command buffer, tracking state across command buffers would require significant effort[1] when the command buffers were submitted for execution. Therefore, this responsibility falls to your application. Resource state is perhaps most important for images because they are complex, structured resources.

1. The validation layers do, in fact, attempt to track this state. While this comes with a substantial performance impact, the layer is capable of catching and reporting many resource-state-related issues.

The state of an image is roughly divided into two essentially orthogonal pieces of state: its layout, which determines how the data is laid out in memory and was discussed briefly earlier in the book, and a record of who last wrote to the image, which affects caching and coherency of data on the device. The initial layout of an image is specified when it is created, and then can be changed throughout the image's lifetime, either explicitly using *barriers* or implicitly using renderpass. Barriers also marshal access to resources from different parts of the Vulkan pipeline, and in some cases, transitioning a resource from one layout to another can be accomplished at other midpipeline synchronization work performed by barriers.

The specific use cases for each layout are discussed in some depth later in the book. However, the fundamental act of moving a resource from state to state is known as a barrier, and it is extremely important to get barriers right and to use them effectively in your application.

Pipeline Barriers

A barrier is a synchronization mechanism for memory access mamangement and resource state movement within the stages of the Vulkan pipeline. The primary command for synchronizing access to resources and moving them from state to state is **`vkCmdPipelineBarrier()`**, the prototype of which is

[**Click here to view code image**](#)

```
void vkCmdPipelineBarrier (
    VkCommandBuffer                        commandBuffer,
    VkPipelineStageFlags                   srcStageMask,
    VkPipelineStageFlags                   dstStageMask,
    VkDependencyFlags                      dependencyFlags,
    uint32_t                               memoryBarrierCount,
    const VkMemoryBarrier*                 pMemoryBarriers,
    uint32_t                               bufferMemoryBarrierCount,
    const VkBufferMemoryBarrier*           pBufferMemoryBarriers,
    uint32_t                               imageMemoryBarrierCount,
    const  VkImageMemoryBarrier*           pImageMemoryBarriers);
```

The command buffer that will execute the barrier is passed in `commandBuffer`. The next two parameters, `srcStageMask` and `dstStageMask`, specify which pipeline stages wrote to the resource last and which stages will read from the resource next, respectively. That is, they specify the source and destination for the data flow represented by the barrier. Each is constructed from a number of the members of the `VkPipelineStageFlagBits` enumeration.

- `VK_PIPELINE_STAGE_TOP_OF_PIPE_BIT`: The *top of pipe* is considered to be hit as soon as the device starts processing the command.
- `VK_PIPELINE_STAGE_DRAW_INDIRECT_BIT`: When the pipeline executes an `indirect` command, it fetches some of the parameters for the command from memory. This is the stage that fetches those parameters.
- `VK_PIPELINE_STAGE_VERTEX_INPUT_BIT`: This is the stage where vertex attributes are fetched from their respective buffers. After this, content of vertex buffers can be overwritten, even if the resulting vertex shaders have not yet completed execution.
- `VK_PIPELINE_STAGE_VERTEX_SHADER_BIT`: This stage is passed when all vertex shader work resulting from a drawing command is completed.

- `VK_PIPELINE_STAGE_TESSELLATION_CONTROL_SHADER_BIT`: This stage is passed when all tessellation control shader invocations produced as the result of a drawing command have completed execution.

- `VK_PIPELINE_STAGE_TESSELLATION_EVALUATION_SHADER_BIT`: This stage is passed when all tessellation evaluation shader invocations produced as the result of a drawing command have completed execution.

- `VK_PIPELINE_STAGE_GEOMETRY_SHADER_BIT`: This stage is passed when all geometry shader invocations produced as the result of a drawing command have completed execution.

- `VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT`: This stage is passed when all fragment shader invocations produced as the result of a drawing command have completed execution. Note that there is no way to know that a primitive has been completely rasterized while the resulting fragment shaders have not yet completed. However, rasterization does not access memory, so no information is lost here.

- `VK_PIPELINE_STAGE_EARLY_FRAGMENT_TESTS_BIT`: All per-fragment tests that might occur *before* the fragment shader is launched have completed.

- `VK_PIPELINE_STAGE_LATE_FRAGMENT_TESTS_BIT`: All per-fragment tests that might occur *after* the fragment shader is executed have completed. Note that outputs to the depth and stencil attachments happen as part of the test, so this stage and the early fragment test stage include the depth and stencil outputs.

- `VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT`: Fragments produced by the pipeline have been written to the color attachments.

- `VK_PIPELINE_STAGE_COMPUTE_SHADER_BIT`: Compute shader invocations produced as the result of a dispatch have completed.

- `VK_PIPELINE_STAGE_TRANSFER_BIT`: Any pending transfers triggered as a result of calls to **vkCmdCopyImage()** or **vkCmdCopyBuffer()**, for example, have completed.

- `VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT`: All operations considered to be part of the graphics pipeline have completed.

- `VK_PIPELINE_STAGE_HOST_BIT`: This pipeline stage corresponds to access from the host.

- `VK_PIPELINE_STAGE_ALL_GRAPHICS_BIT`: When used as a destination, this special flag means that any pipeline stage may access memory. As a source, it's effectively equivalent to `VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT`.

- `VK_PIPELINE_STAGE_ALL_COMMANDS_BIT`: This stage is the big hammer. Whenever you just don't know what's going on, use this; it will synchronize everything with everything. Just use it wisely.

Because the flags specified in `srcStageMask` and `dstStageMask` are used to indicate *when* things happen, it's acceptable for Vulkan implementations to move them around or interpret them in various ways. The `srcStageMask` specifies when the source stage has finished reading or writing a resource. As a result, moving the effective position of that stage later in the pipeline doesn't change the fact that those accesses have completed; it may mean only that the implementation waits longer than it really needs to for them to complete.

Likewise, the `dstStageMask` specifies the point at which the pipeline will wait before proceeding. If an implementation moves that wait point earlier, that will still work. The event that it waits on will

still have completed when the logically later parts of the pipeline begin execution. That implementation just misses the opportunity to perform work when it was instead waiting.

The `dependencyFlags` parameter specifies a set of flags that describes how the dependency represented by the barrier affects the resources referenced by the barrier. The only defined flag is `VK_DEPENDENCY_BY_REGION_BIT`, which indicates that the barrier affects only the region modified by the source stages (if it can be determined), which is consumed by the destination stages.

A single call to **vkCmdPipelineBarrier()** can be used to trigger many barrier operations. There are three types of barrier operations: global memory barriers, buffer barriers, and image barriers. Global memory barriers affect things such as synchronized access to mapped memory between the host and the device. Buffer and image barriers primarily affect device access to buffer and image resources, respectively.

Global Memory Barriers

The number of global memory barriers to be triggered by **vkCmdPipelineBarrier()** is specified in `memoryBarrierCount`. If this is nonzero, then `pMemoryBarriers` points to an array of `memoryBarrierCount` `VkMemoryBarrier` structures, each defining a single memory barrier. The definition of `VkMemoryBarrier` is

[Click here to view code image](#)

```
typedef struct VkMemoryBarrier {
    VkStructureType     sType;
    const void*         pNext;
    VkAccessFlags       srcAccessMask;
    VkAccessFlags       dstAccessMask;
} VkMemoryBarrier;
```

The `sType` field of `VkMemoryBarrier` should be set to `VK_STRUCTURE_TYPE_MEMORY_BARRIER`, and `pNext` should be set to `nullptr`. The only other fields in the structure are the source and destination access masks specified in `srcAccessMask` and `dstAccessMask`, respectively. The access masks are bitfields containing members of the `VkAccessFlagBits`. The source access mask specifies how the memory was last written, and the destination access mask specifies how the memory will next be read. The available access flags are

- `VK_ACCESS_INDIRECT_COMMAND_READ_BIT`: The memory referenced will be the source of commands in an indirect drawing or dispatch command such as **vkCmdDrawIndirect()** or **vkCmdDispatchIndirect()**.

- `VK_ACCESS_INDEX_READ_BIT`: The memory referenced will be the source of index data in an indexed drawing command such as **vkCmdDrawIndexed()** or **vkCmdDrawIndexedIndirect()**.

- `VK_ACCESS_VERTEX_ATTRIBUTE_READ_BIT`: The memory referenced will be the source of vertex data fetched by Vulkan's fixed-function vertex assembly stage.

- `VK_ACCESS_UNIFORM_READ_BIT`: The memory referenced is the source of data for a uniform block accessed by a shader.

- `VK_ACCESS_INPUT_ATTACHMENT_READ_BIT`: The memory referenced is used to back an image used as an input attachment.
- `VK_ACCESS_SHADER_READ_BIT`: The memory referenced is used to back an image object that is read from using image loads or texture reads in a shader.
- `VK_ACCESS_SHADER_WRITE_BIT`: The memory referenced is used to back an image object that is written to using image stores in a shader.
- `VK_ACCESS_COLOR_ATTACHMENT_READ_BIT`: The memory referenced is used to back an image used as a color attachment where reads are performed, perhaps because blending is enabled. Note that this is not the same as an input attachment, where data is read explicitly by the fragment shader.
- `VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT`: The memory referenced is used to back an image used as a color attachment that will be written to.
- `VK_ACCESS_DEPTH_STENCIL_ATTACHMENT_READ_BIT`: The memory referenced is used to back an image used as a depth or stencil attachment that will be read from because the relevant test is enabled.
- `VK_ACCESS_DEPTH_STENCIL_ATTACHMENT_WRITE_BIT`: The memory referenced is used to back an image used as a depth or stencil attachment that will be written to because the relevant write mask is enabled.
- `VK_ACCESS_TRANSFER_READ_BIT`: The memory referenced is used as the source of data in a transfer operation such as **`vkCmdCopyImage()`**, **`vkCmdCopyBuffer()`**, or **`vkCmdCopyBufferToImage()`**.
- `VK_ACCESS_TRANSFER_WRITE_BIT`: The memory referenced is used as the destination of a transfer operation.
- `VK_ACCESS_HOST_READ_BIT`: The memory referenced is mapped and will be read from by the host.
- `VK_ACCESS_HOST_WRITE_BIT`: The memory referenced is mapped and will be written to by the host.
- `VK_ACCESS_MEMORY_READ_BIT`: All other memory reads not explicitly covered by the preceding cases should specify this bit.
- `VK_ACCESS_MEMORY_WRITE_BIT`: All other memory writes not explicitly covered by the preceding cases should specify this bit.

Memory barriers provide two important pieces of functionality. First, they help avoid hazards, and second, they help ensure data consistency.

A *hazard* occurs when read and write operations are reordered relative to the order in which the programmer expects them to execute. They can be very hard to diagnose because they are often platform- or timing-dependent. There are three types of hazards:

- A *read-after-write*, or RaW, hazard occurs when the programmer expects to read from a piece of memory that has recently been written to and that those reads will *see* the results of the writes. If the read is rescheduled and ends up executing before the write is complete, the read will see old data.
- A *write-after-read*, or WaR, hazard occurs when a programmer expects to overwrite a piece of memory that had previously been read by another part of the program. If the write operation

ends up being scheduled before the read operation, then the read operation will see the new data, not the older data it was expecting.

- A *write-after-write*, or WaW, hazard occurs when a programmer expects to overwrite the same location in memory multiple times and that only the results of the last write will be visible to subsequent readers. If the writes are rescheduled with respect to one another, then only the result of the write that happened to execute last will be visible to readers.

There is no such thing as a read-after-read hazard because no data is modified.

In the memory barrier, the source isn't necessarily a producer of data but the first operation that is protected by that barrier. For avoiding RaW hazards, the source is actually a read operation.

For example, to ensure that all texture fetches are complete before overwriting an image with a copy operation, we need to specify `VK_ACCESS_SHADER_READ_BIT` in the `srcAccessMask` field and `VK_ACCESS_TRANSFER_WRITE_BIT` in the `dstAccessMask` field. This tells Vulkan that the first stage is reading from an image in a shader and that the second stage may overwrite that image, so we should not reorder the copy into the image before any shaders that may have read from it.

Note that there is some overlap between the bits in `VkAccessFlagBits` and those in `VkPipelineStageFlagBits`. The `VkAccessFlagBits` flags specify *what* operation is being performed, and the `VkPipelineStageFlagBits` describe *where* in the pipeline the action is performed.

The second piece of functionality provided by the memory barrier is to ensure consistency of the views of data from different parts of the pipeline. For example, if an application contains a shader that writes to a buffer from a shader and then needs to read that data back from the buffer by mapping the underlying memory object, it should specify `VK_ACCESS_SHADER_WRITE_BIT` in `srcAccessMask` and `VK_ACCESS_HOST_READ_BIT` in `dstAccessMask`. If there are caches in the device that may buffer writes performed by shaders, those caches may need to be flushed in order for the host to see the results of the write operations.

Buffer Memory Barriers

Buffer memory barriers provide finer-grained control of the memory used to back buffer objects. The number of buffer memory barriers executed by a call to **vkCmdPipelineBarrier()** is specified in the `bufferMemoryBarrierCount` parameter, and the `pBufferMemoryBarriers` field is a pointer to an array of this many `VkBufferMemoryBarrier` structures, each defining a buffer memory barrier. The definition of `VkBufferMemoryBarrier` is

[**Click here to view code image**](#)

```
typedef struct VkBufferMemoryBarrier {
    VkStructureType    sType;
    const void*        pNext;
    VkAccessFlags      srcAccessMask;
    VkAccessFlags      dstAccessMask;
    uint32_t           srcQueueFamilyIndex;
    uint32_t           dstQueueFamilyIndex;
    VkBuffer           buffer;
    VkDeviceSize       offset;
    VkDeviceSize       size;
} VkBufferMemoryBarrier;
```

The `sType` field of each `VkBufferMemoryBarrier` structure should be set to `VK_STRUCTURE_TYPE_BUFFER_MEMORY_BARRIER`, and `pNext` should be set to `nullptr`. The `srcAccessMask` and `dstAccessMask` fields have the same meanings as they do in the `VkMemoryBarrier` structure. Obviously, some of the flags that refer specifically to images, such as color or depth attachments, have little meaning when dealing with buffer memory.

When ownership of the buffer is being transferred from one queue to another and those queues are in different families, the family indices of the source and destination queues must be supplied in `srcQueueFamilyIndex` and `dstQueueFamilyIndex`, respectively. If there is no transfer of ownership, then `srcQueueFamilyIndex` and `dstQueueFamilyIndex` can both be set to `VK_QUEUE_FAMILY_IGNORED`. In this case, the sole ownership is assumed to be the queue family for which the command buffer is being built.

The buffer the access to which is being controlled by the barrier is specified in `buffer`. To synchronize access to a range of a buffer, use the `offset` and `size` fields of the structure to specify that range, in bytes. To control access to the whole buffer, simply set `offset` to zero and `size` to `VK_WHOLE_SIZE`.

If the buffer will be accessed by work executing on more than one queue, and those queues are of different families, additional action must be taken by your application. Because a single device exposing multiple queue families may actually be made up of multiple physical components, and because those components may have their own caches, scheduling architecture, memory controllers, and so on, Vulkan needs to know when a resource is moved from queue to queue. If this is the case, specify the queue family index of the source queue in `srcQueueFamilyIndex` and the family of the destination queue in `dstQueueFamilyIndex`.

Similarly to image memory barriers, if the resource is not being transferred between queues belonging to different families, `srcQueueFamilyIndex` and `dstQueueFamilyIndex` should be set to `VK_QUEUE_FAMILY_IGNORED`.

Image Memory Barriers

Just as with buffers, special attention should be paid to images, and image memory barriers are used to control access to images. The number of image memory barriers to be performed by the call to **vkCmdPipelineBarrier()** is specified in the `imageMemoryBarrierCount` parameter, and `pImageMemoryBarriers` is a pointer to an array of this many `VkImageMemoryBarrier` structures, each describing a single barrier. The definition of `VkImageMemoryBarrier` is

[Click here to view code image](#)

```
typedef struct VkImageMemoryBarrier {
    VkStructureType           sType;
    const void*               pNext;
    VkAccessFlags             srcAccessMask;
    VkAccessFlags             dstAccessMask;
    VkImageLayout             oldLayout;
    VkImageLayout             newLayout;
    uint32_t                  srcQueueFamilyIndex;
    uint32_t                  dstQueueFamilyIndex;
    VkImage                   image;
    VkImageSubresourceRange   subresourceRange;
} VkImageMemoryBarrier;
```

The sType field of each VkImageMemoryBarrier structure should be set to VK_STRUCTURE_TYPE_IMAGE_MEMORY_BARRIER, and pNext should be set to nullptr. Just as with the other memory barriers, the srcAccessMask and dstAccessMask fields specify the source and destination access type. Again, only some of the access types will apply to images. Also, when you are controlling access across queues, the srcQueueFamilyIndex and dstQueueFamilyIndex fields should be set to the family indices of the queues where the source and destination work will take place.

The oldLayout and newLayout fields specify the layouts to be used for the image before and after the barrier. These are the same fields that can be used when creating the image. The image that the barrier is to affect is specified in image, and the parts of the image to be affected by the barrier are specified in subresourceRange, which is an instance of the VkImageSubresourceRange structure, the definition of which is

[Click here to view code image](#)

```
typedef struct VkImageSubresourceRange {
    VkImageAspectFlags    aspectMask;
    uint32_t              baseMipLevel;
    uint32_t              levelCount;
    uint32_t              baseArrayLayer;
    uint32_t              layerCount;
} VkImageSubresourceRange;
```

The image aspect is the part of the image that is to be included in the barrier. Most image formats and types have only a single aspect. A common exception is a depth-stencil image, which may have a separate aspect for each of the depth and stencil components of the image. It is possible, using the aspect flags, to discard stencil data while keeping depth data for later sampling, for example.

For images with mipmaps, a subset of the mipmaps can be included in the barrier by specifying the lowest-numbered (highest-resolution) mipmap level in the baseMipLevel field and the number of levels in the levelCount field. If the image doesn't have a full mipmap chain, baseMipLevel should be set to 0, and levelCount should be set to 1.

Likewise, for array images, a subset of the image layers can be included in the barrier by setting baseArrayLayer to the index of the first layer and layerCount to the number of layers to include. Again, even if the image is not an array image, you should set baseArrayLayer to 0 and layerCount to 1. In short, treat all images as though they have mipmaps (even if it's only one level) and all images as though they are arrays (even if they have only one layer).

Listing 4.1 shows an example of how to perform an image memory barrier.

Listing 4.1: Image Memory Barrier

[Click here to view code image](#)

```
const VkImageMemoryBarrier imageMemoryBarriers =
{
    VK_STRUCTURE_TYPE_IMAGE_MEMORY_BARRIER,     // sType
    nullptr,                                    // pNext
    VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,       // srcAccessMask
    VK_ACCESS_SHADER_READ_BIT,                  // dstAccessMask
    VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL,   // oldLayout
```

```
        VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL,   // newLayout
        VK_QUEUE_FAMILY_IGNORED,                     // srcQueueFamilyIndex
        VK_QUEUE_FAMILY_IGNORED,                     // dstQueueFamilyIndex
        image,                                       // image
        {                                            // subresourceRange
            VK_IMAGE_ASPECT_COLOR_BIT,               // aspectMask
            0,                                       // baseMipLevel
            VK_REMAINING_MIP_LEVELS,                 // levelCount
            0,                                       // baseArrayLayer
            VK_REMAINING_ARRAY_LAYERS                // layerCount
        }
    };

    vkCmdPipelineBarrier(m_currentCommandBuffer,
                        VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,
                        VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,
                        0,
                        0, nullptr,
                        0, nullptr,
                        1, &imageMemoryBarrier);
```

The image memory barrier shown in Listing 4.1 takes an image that was previously in the
`VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL` layout and moves it to the
`VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL` layout. The source of data is the color
output from the pipeline, as specified by `VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT`, and
the destination of the data is sampling by a shader, as specified by
`VK_ACCESS_SHADER_READ_BIT`.

There is no transfer of ownership across queues, so both `srcQueueFamilyIndex` and
`dstQueueFamilyIndex` are set to `VK_QUEUE_FAMILY_IGNORED`. Also, we're performing the
barrier across all mipmap levels and array layers in the image, so the `levelCount` and
`layerCount` members of the `subresourceRange` structure are set to
`VK_REMAINING_MIP_LEVELS` and `VK_REMAINING_ARRAY_LAYERS`, respectively.

This barrier takes an image that previously was written to as a color attachment by a graphics
pipeline and moves it into a state in which it can be read from by a shader.

Clearing and Filling Buffers

You were introduced to buffer objects in Chapter 2, "Memory and Resources." A buffer is a linear
region of data backed by memory. In order for a buffer to be useful, you need to be able to fill it with
data. In some cases, simply clearing the whole buffer to a known value is all you need to do. This
allows you to, for example, initialize a buffer that you will eventually write into using a shader or
some other operation.

To fill a buffer with a fixed value, call **vkCmdFillBuffer()**, the prototype of which is

[Click here to view code image](#)

```
void vkCmdFillBuffer (
    VkCommandBuffer                          commandBuffer,
    VkBuffer                                 dstBuffer,
    VkDeviceSize                             dstOffset,
```

```
VkDeviceSize                                    size,
uint32_t                                        data);
```

The command buffer into which to place the command is specified in `commandBuffer`. The buffer that will be filled with data is specified in `dstBuffer`. To fill a section of the buffer with data, specify the starting offset of the fill operation, in bytes, in `dstOffset` and the size of the region, again in bytes, in `size`. Both `dstOffset` and `size` must be multiples of 4. To fill from `dstOffset` to the end of the buffer, pass the special value, `VK_WHOLE_SIZE`, in the `size` parameter. It follows that to fill an entire buffer, simply set `dstOffset` to 0 and `size` to `VK_WHOLE_SIZE`.

The value that you want to fill the buffer with is passed in `data`. This is a `uint32_t` variable that is simply replicated for the region of the fill operation. It is as though the buffer is interpreted as an array of `uint32_t`, and each element from `dstOffset` to the end of the region is filled with this value. To clear a buffer with a floating-point value, you can reinterpret the floating-point value as a `uint32_t` value and pass that to `vkCmdFillBuffer()`. [Listing 4.2](#) demonstrates this.

Listing 4.2: Filling a Buffer with Floating-Point Data

[Click here to view code image](#)

```
void FillBufferWithFloats(VkCommandBuffercmdBuffer,
                          VkBuffer dstBuffer,
                          VkDeviceSize offset,
                          VkDeviceSize length,
                          const float value)

{
    vkCmdFillBuffer(cmdBuffer,
                    dstBuffer,
                    0,
                    1024,
                    *(const uint32_t*)&value);
}
```

Sometimes, filling a buffer with a fixed value is not enough, and there is a need to place data more explicitly in a buffer object. When a large amount of data is needed to be transfered into or between buffers, either mapping the buffer and writing to it with the host or copying data from another (possibly mapped) buffer with **vkCmdCopyBuffer()** is most appropriate. However, for small updates, such as updating the values of a vector or small data structures, **vkCmdUpdateBuffer()** can be used to place data directly into a buffer object.

The prototype for **vkCmdUpdateBuffer()** is

[Click here to view code image](#)

```
void vkCmdUpdateBuffer (
    VkCommandBuffer                             commandBuffer,
    VkBuffer                                    dstBuffer,
    VkDeviceSize                                dstOffset,
    VkDeviceSize                                dataSize,
    const uint32_t*                             pData);
```

vkCmdUpdateBuffer() copies data directly from host memory into a buffer object. The data is consumed from host memory as soon as **vkCmdUpdateBuffer()** is called, and as such, it's fine to free the host memory data structure or overwrite its content once **vkCmdUpdateBuffer()** returns. Be aware, though, that the data is not written into the buffer until **vkCmdUpdateBuffer()** is executed by the device after the command buffer has been submitted. For this reason, Vulkan must make a *copy* of the data you've supplied and hold it either in some auxiliary data structure associated with the command buffer or directly inside the command buffer itself.

Again, the command buffer that will contain the command is passed in commandBuffer, and the destination buffer object is passed in dstBuffer. The offset at which the data is to be placed is passed in dstOffset, and the size of the data to place into the buffer is passed in dataSize. Both dstOffset and dataSize are in units of bytes, but as with **vkCmdFillBuffer()**, both must be a multiple of 4. The special value VK_WHOLE_SIZE is not accepted for the size parameter to **vkCmdUpdateBuffer()** because it is also used as the size of the host memory region that is the source of the data. The maximum size of data that can be placed in a buffer with **vkCmdUpdateBuffer()** is 65,536 bytes.

pData points to the host memory containing the data that will eventually be placed into the buffer object. Although the type of the variable expected here is a pointer to uint32_t, any data can be in the buffer. Simply typecast a pointer to any memory region readable by the host to `const` uint32_t*, and pass it to pData. Ensure that the data region is at least size bytes long. For example, it's reasonable to construct a C++ data structure matching the layout of a uniform or shader storage block and simply copy its entire content into a buffer that will be used appropriately in a shader.

Again, be cautious when using **vkCmdFillBuffer()**. It is intended for short, immediate updates to buffers. For example, writing a single value into a uniform buffer is probably much more efficiently achieved with **vkCmdFillBuffer()** than it is with a buffer mapping and a call to **vkCmdCopyBuffer()**.

Clearing and Filling Images

Just as with buffers, it it possible to copy data directly between images and to fill images with fixed values. Images are larger, more complex, opaque data structures, so the raw offsets and data are not generally visible to an application.[2]

> [2]. Of course, it's possible to map the memory that is used for backing an image. In particular, when *linear* tiling is used for an image, this is standard practice. However, in general, this is not recommended.

To clear an image to a fixed value, call vkCmdClearColorImage(), the prototype of which is

[Click here to view code image](#)

```
void vkCmdClearColorImage (
    VkCommandBuffer                  commandBuffer,
    VkImage                          image,
    VkImageLayout                    imageLayout,
    const VkClearColorValue*         pColor,
    uint32_t                         rangeCount,
    const VkImageSubresourceRange*   pRanges);
```

The command buffer that will contain the clear command is passed in `commandBuffer`. The image that is to be cleared is passed in `image`, and the layout that the image is expected to be in when the clear command is executed is passed in `imageLayout`.

The accepted layouts for `imageLayout` are `VK_IMAGE_LAYOUT_GENERAL` and `VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL`. To clear images that are in different layouts, it is necessary to move them to one of these two layouts using a pipeline barrier before executing the clear command.

The values to clear the image to are specified in an instance of the `VkClearColorValue` union, the definition of which is

Click here to view code image

```
typedef union VkClearColorValue {
    float       float32[4];
    int32_t     int32[4];
    uint32_t    uint32[4];
} VkClearColorValue;
```

The `VkClearColorValue` is simply a union of three arrays of four values each. One is for floating-point data, one is for signed integer data, and one is for unsigned integer data. Vulkan will read the appropriate member for the format of the image being cleared. Your application can write into the member that matches the source of data. No data conversion is performed by **`vkCmdClearColorImage()`**; it is up to your application to fill the `VkClearColorValue` union correctly.

Any number of regions of the destination image can be cleared with a single call to **`vkCmdClearColorImage()`**, although each will be cleared with the same values. If you need to clear multiple regions of the same image with different colors, you will need to call **`vkCmdClearColorImage()`** multiple times. However, you want to clear all regions with the same color, specify the number of regions in `rangeCount`, and pass a pointer to an array of `rangeCount` `VkImageSubresourceRange` structures in `pRanges`. The definition of `VkImageSubresourceRange` is

Click here to view code image

```
typedef struct VkImageSubresourceRange {
    VkImageAspectFlags    aspectMask;
    uint32_t              baseMipLevel;
    uint32_t              levelCount;
    uint32_t              baseArrayLayer;
    uint32_t              layerCount;
} VkImageSubresourceRange;
```

This structure was first introduced in Chapter 2, "Memory and Resources," when we discussed creation of image view. Here, it is used to define the regions of the image that you want to clear. Because we are clearing a color image, the `aspectMask` must be set to `VK_IMAGE_ASPECT_COLOR_BIT`. The `baseMipLevel` and `levelCount` fields are used to specify the starting mipmap level and number of levels to clear, respectively, and if the image is an array image, the `baseArrayLayer` and `layerCount` fields are used to specify the starting layer

and number of layers to clear. If the image is not an array image, these fields should be set to 0 and 1, respectively.

Clearing a depth-stencil image is similar to clearing a color image, except that a special `VkClearDepthStencilValue` structure is used to specify the clear values. The prototype of **vkCmdClearDepthStencilImage()** is similar to that of **vkCmdClearColorImage()** and is

[**Click here to view code image**](#)

```
void vkCmdClearDepthStencilImage (
    VkCommandBuffer                     commandBuffer,
    VkImage                             image,
    VkImageLayout                       imageLayout,
    const VkClearDepthStencilValue*     pDepthStencil,
    uint32_t                            rangeCount,
    const VkImageSubresourceRange *     pRanges);
```

Again, the command buffer that will perform the clear operation is specified in `commandBuffer`, the image to clear is specified in `image`, and the layout that the image is expected to be in at the time of the clear operation is specified in `imageLayout`. As with **vkCmdClearColorImage()**, `imageLayout` should be either `VK_IMAGE_LAYOUT_GENERAL` or `VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL`. No other layouts are valid for a clear operation.

The values to which to clear the depth-stencil image are passed through an instance of the `VkClearDepthStencilValue` structure, which contains both the depth and stencil clear values. Its definition is

[**Click here to view code image**](#)

```
typedef struct VkClearDepthStencilValue {
    float       depth;
    uint32_t    stencil;
} VkClearDepthStencilValue;
```

As with **vkCmdClearColorImage()**, a number of ranges of the image can be cleared in a single call to **vkCmdClearDepthStencilImage()**. The number of ranges to clear is specified in `rangeCount`, and the `pRanges` parameter should point to an array of `rangeCount` `VkImageSubresourceRange` structures defining the ranges to be cleared.

Because depth-stencil images may contain both a depth and a stencil aspect, the `aspectMask` field of each member of `pRanges` can contain `VK_IMAGE_ASPECT_DEPTH_BIT`, `VK_IMAGE_ASPECT_STENCIL_BIT`, or both. If `aspectMask` contains `VK_IMAGE_ASPECT_DEPTH_BIT`, then the value stored in the `depth` field of the `VkClearDepthStencilValue` structure is used to clear the depth aspect of the specified range. Likewise, if `aspectMask` contains `VK_IMAGE_ASPECT_STENCIL_BIT`, then the stencil aspect of the specified range will be cleared using the `stencil` member of the `VkClearDepthStencilValue` structure.

Note that it's generally much more efficient to specify a single region with both `VK_IMAGE_ASPECT_DEPTH_BIT` and `VK_IMAGE_ASPECT_STENCIL_BIT` set than it is to specify two regions each with only one bit set.

Copying Image Data

In the previous section, we discussed clearing images to a fixed value passed through a simple structure. In many cases, though, you need to upload texture data into images or copy image data between images. Vulkan supports copying image data from a buffer to an image, between images, and from an image to a buffer.

To copy data from a buffer to one or more regions of an image, call `vkCmdCopyBufferToImage()`, the prototype of which is

[**Click here to view code image**](#)

```
void vkCmdCopyBufferToImage (
    VkCommandBuffer                     commandBuffer,
    VkBuffer                            srcBuffer,
    VkImage                             dstImage,
    VkImageLayout                       dstImageLayout,
    uint32_t                            regionCount,
    const VkBufferImageCopy*            pRegions);
```

The command buffer that will execute the command is specified in `commandBuffer`, the source buffer object is specified in `srcBuffer`, and the image into which the data will be copied is specified in `dstImage`. As with the destination image in clears, the layout of the destination image for copies is expected to be either `VK_IMAGE_LAYOUT_GENERAL` or `VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL` and is specified in the `dstImageLayout` parameter.

The number of regions to update is given in `regionCount`, and `pRegions` is a pointer to an array of `regionCount` `VkBufferImageCopy` structures, each definining an area of the image to copy data into. The definition of `VkBufferImageCopy` is

[**Click here to view code image**](#)

```
typedef struct VkBufferImageCopy {
    VkDeviceSize                bufferOffset;
    uint32_t                    bufferRowLength;
    uint32_t                    bufferImageHeight;
    VkImageSubresourceLayers    imageSubresource;
    VkOffset3D                  imageOffset;
    VkExtent3D                  imageExtent;
} VkBufferImageCopy;
```

The `bufferOffset` field contains the offset of the data in the buffer, in bytes. The data in the buffer is laid out left to right, top to bottom, as shown in [Figure 4.1](#). The `bufferRowLength` field specifies the number of texels in the source image, and `bufferImageHeight` specifies the number of rows of data in the image. If `bufferRowLength` is zero, the image is assumed to be tightly packed in the buffer and therefore equal to `imageExtent.width`. Likewise, if `bufferImageHeight` is zero, then the number of rows in the source image is assumed to be equal to the height of the image extent, which is in `imageExtent.height`.
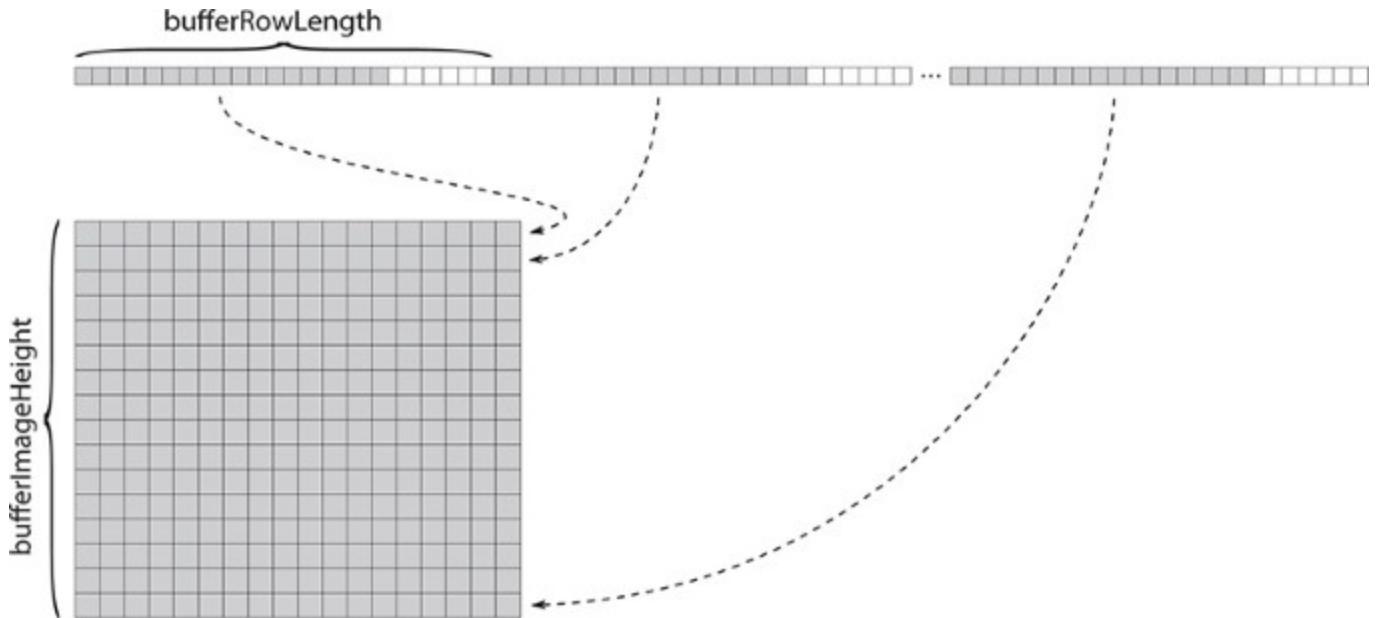
Figure 4.1: Data Layout of Images Stored in Buffers

The subresource into which to copy the image data is specified in an instance of the `VkImageSubresourceLayers` structure, the definition of which is

[**Click here to view code image**](#)

```
typedef struct VkImageSubresourceLayers {
    VkImageAspectFlags    aspectMask;
    uint32_t              mipLevel;
    uint32_t              baseArrayLayer;
    uint32_t              layerCount;
} VkImageSubresourceLayers;
```

The `aspectMask` field of `VkImageSubresourceLayers` contains the aspect or aspects that are the destination of the image copy. Usually, this will be a single bit from the `VkImageAspectFlagBits` enumeration. If the target image is a color image, then this should simply be set to `VK_IMAGE_ASPECT_COLOR_BIT`. If the image is a depth-only image, it should be `VK_IMAGE_ASPECT_DEPTH_BIT`, and if the image is a stencil-only image, it should be `VK_IMAGE_ASPECT_STENCIL_BIT`. If the image is a combined depth-stencil image, then you can copy data into both the depth and stencil aspects simultaneously by specifying both `VK_IMAGE_ASPECT_DEPTH_BIT` and `VK_IMAGE_ASPECT_STENCIL_BIT`.

The target mipmap level is specified in `mipLevel`. You can copy data into only a single mipmap level with each element in the `pRegions` array, although you can of course specify multiple elements, each targeting a different level.

If the target image is an array image, then you can specify the starting layer and number of layers for the image copy in `baseArrayLayer` and `layerCount`, respectively. If the image is not an array image, then these fields should be set to 0 and 1.

Each region can target either an entire mipmap level or a smaller window within each mipmap level. The offset of the window is specified in `imageOffset`, and the size of the window is specified in `imageExtent`. To overwrite an entire mipmap level, set `imageOffset.x` and

`imageOffset.y` to 0, and set `imageExtent.width` and `imageExtent.height` to the size of the mipmap level. It is up to you to calculate this. Vulkan will not do it for you.

It's also possible to perform the copy in the opposite direction—to copy data from an image into a buffer. To do this, call **vkCmdCopyImageToBuffer()**, the prototype of which is

[**Click here to view code image**](#)

```
void vkCmdCopyImageToBuffer (
    VkCommandBuffer                     commandBuffer,
    VkImage                             srcImage,
    VkImageLayout                       srcImageLayout,
    VkBuffer                            dstBuffer,
    uint32_t                            regionCount,
    const VkBufferImageCopy*             pRegions);
```

The command buffer to execute the copy is specified in `commandBufer`, the source image in `srcImage`, and the destination buffer in `dstBuffer`. As with the other copy commands, the `srcImageLayout` parameter specifies the layout that the source image is expected to be in. Because the image is now the `source` of data, the layout should either be `VK_IMAGE_LAYOUT_GENERAL` or `VK_IMAGE_LAYOUT_TRANSFER_SRC_OPTIMAL`.

Again, a number of regions can be copied in a single call to **vkCmdCopyImageToBuffer()**, each represented by an instance of the `VkBufferImageCopy` structure. The number of regions to copy is specified in `regionCount`, and the `pRegions` parameter contains a pointer to an array of `regionCount` `VkBufferImageCopy` structures defining each of these regions. This is the same structure accepted by **vkCmdCopyBufferToImage()**. However, in this use case, `bufferOffset`, `bufferRowLength`, and `bufferImageHeight` contain parameters for the destination of the copy, and `imageSubresource`, `imageOffset`, and `imageExtent` contain parameters for the source of the copy.

Finally, it's also possible to copy data between two images. To do this, use the **vkCmdCopyImage()** command, the prototype of which is

[**Click here to view code image**](#)

```
void vkCmdCopyImage (
    VkCommandBuffer                             commandBuffer,
    VkImage                                     srcImage,
    VkImageLayout                               srcImageLayout,
    VkImage                                     dstImage,
    VkImageLayout                               dstImageLayout,
    uint32_t                                    regionCount,
    const VkImageCopy*                          pRegions);
```

The command buffer that will execute the command is passed in `commandBuffer`, the image containing the source data is passed in `srcImage`, and the image that is the destination for the copy is passed in `dstImage`. Again, the layout for both images must be passed to the copy command. `srcImageLayout` is the expected layout of the source image at the time of the copy and should be either `VK_IMAGE_LAYOUT_GENERAL` or `VK_IMAGE_LAYOUT_TRANSFER_SRC_OPTIMAL` (as this is the source of a transfer operation). Similarly, `dstImageLayout` is the expected layout of the

destination image and should be either `VK_IMAGE_LAYOUT_GENERAL` or
`VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL`.

As with the buffer-to-image and image-to-buffer copy commands, **`vkCmdCopyImage()`** can copy several regions at a time. The number of regions to copy is specified in `regionCount`, and each is represented by an instance of the `VkImageCopy` structure contained in an array, the address of which is passed in `pRegions`. The definition of `VkImageCopy` is

[**Click here to view code image**](#)

```
typedef struct VkImageCopy {
    VkImageSubresourceLayers    srcSubresource;
    VkOffset3D                  srcOffset;
    VkImageSubresourceLayers    dstSubresource;
    VkOffset3D                  dstOffset;
    VkExtent3D                  extent;
} VkImageCopy;
```

Each instance of `VkImageCopy` contains the subresource information and offsets for the source and destination windows. **`vkCmdCopyImage()`** cannot resize image data, so the extent of the source and destination regions is the same and is contained in the `extent` field.

`srcSubresource` contains the subresource definition for the source data and has the same meaning as the `imageSubresource` field in the `VkBufferImageCopy` structure passed to **`vkCmdCopyImageToBuffer()`**. Likewise, the `dstSubresource` field contains the subresource definition for the destination region and has the same meaning as the `imageSubresource` field in the `VkBufferImageCopy` structure passed to **`vkCmdCopyBufferToImage()`**.

The `srcOffset` and `dstOffset` fields contain the coordinates of the source and destination windows, respectively.

Copying Compressed Image Data

As discussed in [Chapter 2](#), "[Memory and Resources](#)," Vulkan supports a number of compressed image formats. All compression formats currently defined are block-based formats with fixed block sizes. For many of these formats, the block size is 4 × 4 texels. For the ASTC formats, the block size varies by image.

When copying data between buffers and images, only an integral number of blocks may be copied. Therefore, the width and height of each image region, in texels, must be integer multiples of the block size used by the image. Further, the origins of copy regions must also be integer multiples of the block size.

It is also possible to copy data between two compressed images or between a compressed and an uncompressed image using **`vkCmdCopyImage()`**. When you do so, the source and destination image formats must have the same compressed block size. That is, if the size of the compressed block is 64 bits, for example, then both the source and destination formats must be compressed images with 64-bit block sizes, or the uncompressed image format must be a 64-bit per-texel format.

When copying from an uncompressed image to a compressed one, each source texel is treated as a single raw value containing the same number of bits as a block in the compressed image. This value is written directly into the compressed image as though it were the compressed data. The texel values

are not compressed by Vulkan. This allows you to create compressed image data in your application or shaders and then copy it into compressed images for later processing. Vulkan does not compress raw image data for you. Further, for uncompressed to compressed copies, the `extent` field of the `VkImageCopy` structure is in units of texels in the source image but must conform to the block size requirements of the destination image.

When copying from a compressed format to an uncompressed format, the opposite is true. Vulkan does not decompress the image data. Rather, it pulls raw 64-bit or 128-bit compressed block values from the source image and deposits them in the destination image. In this case, the destination image should have the same number of bits per texel as bits per block in the source image. For a compressed to uncompressed copy, the `extent` field of the `VkImageCopy` structure is measured in units of texels in the destination image but must conform to the requiremnts imposed by the block size in the source image.

Copying between two block compressed image formats is allowed, so long as both formats have an equal number of bits per block. However, the value of this is debatable, as image data compressed in one format generally does not decode meaningfully when interpreted as another format. Regardless of its worth, when performing this operation, the regions to be copied are still measured in texels, but all offsets and extents must be integer multiples of the common block size.

The only exception to the rule that image copies into, out of, and between compressed images are aligned to multiples of the block size occurs when the source or destination image is not an integer multiple of the block size wide or high, and the region to be copied extends to the edge of the image.

Stretching Images

Of all the image-related commands covered so far, none supports format conversion or resizing of the copied area. To do this, you need to use the **vkCmdBlitImage()** command, which can take images of different formats and stretch or shrink the region to be copied as it is written into the target image. The term *blit* is short for *bl*ock *i*mage *t*ransfer and refers to the operation of not only copying image data, but potentially also processing it along the way.

The prototype of **vkCmdBlitImage()** is

[Click here to view code image](#)

```
void vkCmdBlitImage (
    VkCommandBuffer                         commandBuffer,
    VkImage                                 srcImage,
    VkImageLayout                           srcImageLayout,
    VkImage                                 dstImage,
    VkImageLayout                           dstImageLayout,
    uint32_t                                regionCount,
    const VkImageBlit*                      pRegions,
    VkFilter                                filter);
```

The command buffer that will execute the command is passed in `commandBuffer`. The source and destination images are passed in `srcImage` and `dstImage`, respectively. Again, as with **vkCmdCopyImage()**, the expected layouts of the source and destination images are passed in `srcImageLayout` and `dstImageLayout`. The layout of the source image must be either `VK_IMAGE_LAYOUT_GENERAL` or `VK_IMAGE_LAYOUT_TRANSFER_SRC_OPTIMAL`, and the

layout of the destination image must be either `VK_IMAGE_LAYOUT_GENERAL` or `VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL`.

As with the other copy commands, **`vkCmdBlitImage()`** can copy any number of regions of the source image into the destination image, and each is represented by a data structure. The number of regions to copy is passed in `regionCount`, and `pRegion` points to an array of `regionCount` `VkImageBlit` structures, each defining one of the regions to copy. The definition of `VkImageBlit` is

[**Click here to view code image**](#)

```
typedef struct VkImageBlit {
    VkImageSubresourceLayers    srcSubresource;
    VkOffset3D                  srcOffsets[2];
    VkImageSubresourceLayers    dstSubresource;
    VkOffset3D                  dstOffsets[2];
} VkImageBlit;
```

The `srcSubresource` and `dstSubresource` fields of `VkImageBlit` define the subresource for the source and destination images. Whereas in `VkImageCopy` each region was defined by a `VkOffset3D` structure and shared a `VkExtent3D` structure, in `VkImageBlit` each region is defined by a pair of `VkOffset3D` structures arranged as arrays of two elements.

The first element of the `srcOffsets` and `dstOffsets` arrays defines one corner of the region to be copied, and the second element of these arrays defines the opposite corner of the region. The region defined by `srcOffsets` in the source image is then copied into the region defined by `dstOffsets` in the destination image. If either region is "upside down" with respect to the other, then the copied region will be flipped vertically. Likewise, if one region is "back to front" with respect to the other, then the image will be flipped horizontally. If both of these conditions are met, then the copied region will be rotated 180° with respect to the original.

If the regions are different sizes in the source and destination rectangles, then the image data will be magnified or minified, accordingly. In this case, the filter mode specified in the `filter` parameter to **`vkCmdBlitImage()`** will be used to filter the data. `filter` must be one of `VK_FILTER_NEAREST` or `VK_FILTER_LINEAR` to apply point sampling or linear filtering, respectively.

The format of the source image must be one that supports the `VK_FORMAT_FEATURE_BLIT_SRC_BIT` feature. In most implementations, this will include almost all image formats. Further, the destination format must be one that supports `VK_FORMAT_FEATURE_BLIT_DST_BIT`. In general, this is any format that can be rendered to or written to by the device using image stores in shaders. It is unlikely that any Vulkan device supports blitting to a compressed image format.

Summary

This chapter discussed how to clear images with fixed values and full buffer objects with data. We placed small amounts of data directly into buffer objects using commands embedded inside command buffers and explained how Vulkan is able to copy image data between buffers and images, between images and buffers, and between pairs of images. Finally, we introduced you to the concept of a blit, which is an operation that allows image data to be scaled and to undergo format conversion as it is copied. These operations provide a foundation for getting large amounts of data into and out of the Vulkan device for further processing.

Chapter 5. Presentation

What You'll Learn in This Chapter

  • How to display the results of your program onscreen

  • How to determine the display devices attached to the system

  • How to change display modes and interface with a native window system

---

Vulkan is primarily a graphics API in the sense that the majority of its functionality is dedicated to generating and processing images. Most Vulkan applications will be designed to show their results to the user. This is a process known as *presentation*. However, because the variety of platforms upon which Vulkan runs is large, and because not all applications need to present their outputs to the user visually, presentation is not a core part of the API but is handled by a set of extensions. This chapter discusses how to enable and use those extensions to get pictures on the screen.

Presentation Extension

Presentation in Vulkan is not part of the core API. In fact, a given implementation of Vulkan may not support presentation at all. The reasons for this are

  • Not all Vulkan applications need to present images to the user. Computecentric applications, for example, might produce nonvisual data or produce images that only need to be saved to disk rather than displayed in real time.

  • Presentation is generally handled by the operating system, window system, or other platform-specific library, which can vary quite a bit from platform to platform.

Due to this, presentation is handled by a set of extensions collectively known as the WSI extensions, or Window System Integration systems. Extensions in Vulkan must be enabled explicitly before they can be used, and the extension needed for each platform is slightly different, as some of the functions take platform-specific parameters. Before you can perform any presentation-related operations, therefore, you need to enable the appropriate presentation-related extensions using the mechanisms described in Chapter 1, "Overview of Vulkan."

Presentation in Vulkan is handled by a small suite of extensions. Functionality that is common to almost all platforms that support presenting graphical output to the user is supported by one extension, and functionality that is specific to each platform is supported by a number of smaller, platform-specific extensions.

Presentation Surfaces

The object to which graphics data is rendered in order to be presented is known as a *surface* and is represented by a `VkSurfaceKHR` handle. This special object is introduced by the `VK_KHR_surface` extension. This extension adds general functionality for handling surface objects but is customized on a per-platform basis to provide the platform-specific interfaces to associate a surface with a window. Interfaces are defined for Microsoft Windows, Mir and Wayland, X Windows via either the XCB or Xlib interface, and Android. Further platforms may be added in the future.

The prototypes and data types for the platform-specific parts of the extension are included in the main `vulkan.h` header file but are protected behind platform-specific preprocessor guards. The code for this book supports the Windows platform and the Linux platform via the Xlib and Xcb interfaces. To enable the code for these platforms, before including `vulkan.h`, we must define one of the `VK_USE_PLATFORM_WIN32_KHR`, `VK_USE_PLATFORM_XLIB_KHR`, or `VK_USE_PLATFORM_LIB_XCB_KHR` defines. The build system for the book's source code does this for you using a compiler command-line option.

Vulkan is also supported on a wide range of other operating systems and device types. In particular, many of Vulkan's features are geared to mobile and embedded devices. For example, Vulkan is the API of choice on the Android platform, and in addition to the interfaces covered here, the Android platform has full support through its own platform interfaces. Outside of initialization, though, using Vulkan on Android should be a fairly similar experience to using Vulkan on other platforms.

Presentation on Microsoft Windows

Before we can present, we need to determine whether any queues on a device support presentation operations. Presentation capability is a per-queue-family feature. On Windows platforms, call the **vkGetPhysicalDeviceWin32PresentationSupportKHR()** function to determine whether a particular queue family supports presentation. Its prototype is

[**Click here to view code image**](#)

```
VkBool32 vkGetPhysicalDeviceWin32PresentationSupportKHR(
    VkPhysicalDevice                          physicalDevice,
    uint32_t                                  queueFamilyIndex);
```

The physical device being queried is passed in `physicalDevice`, and the queue family index is passed in `queueFamilyIndex`. If at least one queue family supports presentation, then we can proceed to create presentable surfaces using the device. To create a surface, use the **vkCreateWin32SurfaceKHR()** function, whose prototype is

[**Click here to view code image**](#)

```
VkResult vkCreateWin32SurfaceKHR(
    VkInstance                            instance,
    const VkWin32SurfaceCreateInfoKHR*    pCreateInfo,
    const VkAllocationCallbacks*          pAllocator,
    VkSurfaceKHR*                         pSurface);
```

This function associates a Windows native window handle with a new surface object and returns the object in the variable pointed to by `pSurface`. Only a Vulkan instance is required, and its handle is passed into `instance`. The information describing the new surface is passed through `pCreateInfo`, which is a pointer to an instance of the `VkWin32SurfaceCreateInfoKHR` structure, the definition of which is

[**Click here to view code image**](#)

```
typedef struct VkWin32SurfaceCreateInfoKHR {
    VkStructure Type                sType;
    const void*                     pNext;
    VkWin32SurfaceCreateFlagsKHR    flags;
    HINSTANCE                       hinstance;
```

```
    HWND                                          hwnd;
} VkWin32SurfaceCreateInfoKHR;
```

The `sType` field of `VkWin32SurfaceCreateInfoKHR` should be set to
`VK_STRUCTURE_TYPE_DISPLAY_SURFACE_CREATE_INFO_KHR`, and `pNext` should be set
to `nullptr` unless another extension in use extends the structure. The `flags` field is reserved for
future use and should be set to zero.

The `hinstance` parameter should be set to the `HINSTANCE` of the application or module that was
used to create the the native window. This is typically passed to the application in the first parameter
of `WinMain` or can be obtained by calling the Win32 function `GetModuleHandle` with a null
pointer. The `hwnd` member is the handle to the native window with which to associate the Vulkan
surface. This is the window in which the results of presentation to swap chains created for the surface
will be displayed.

Presentation on Xlib-Based Platforms

The process for creating a surface on an Xlib-based system is similar. First, we need to determine
whether the platform supports presentation to an Xlib surface on an X server. To do this, call
vkGetPhysicalDeviceXlibPresentationSupportKHR(), whose prototype is

[Click here to view code image](#)

```
VkBool32 vkGetPhysicalDeviceXlibPresentationSupportKHR(
    VkPhysicalDevice                              physicalDevice,
    uint32_t                                      queueFamilyIndex,
    Display*                                      dpy,
    VisualID                                      visualID);
```

For the physical device whose handle is specified in `physicalDevice`, and the queue family
index specified in `queueFamilyIndex`,
vkGetPhysicalDeviceXlibPresentationSupportKHR() reports whether queues in that
family support presentation to Xlib surfaces for a given X server. The connection to the X server is
represented by the `dpy` parameter. Presentation is supported on a per-format basis. In Xlib, formats
are represented by visual IDs, and the visual ID for the intended format of the surface is passed in
`visualID`.

Asssuming that at least one queue family on a device supports presentation to a format we'd like to
use, we can then create a surface for an Xlib window by calling the
vkCreateXlibSurfaceKHR() function, the prototype of which is

[Click here to view code image](#)

```
VkResult vkCreateXlibSurfaceKHR(
    VkInstance                                    instance,
    const VkXlibSurfaceCreateInfoKHR*             pCreateInfo,
    const VkAllocationCallbacks*                  pAllocator,
    VkSurfaceKHR*                                 pSurface);
```

vkCreateXlibSurfaceKHR() creates a new surface associated with an Xlib window. The
Vulkan instance should be passed in `instance`, and the remaining parameters controlling the

creation of the surface are passed in `pCreateInfo`, which is a pointer to an instance of the `VkXlibSurfaceCreateInfoKHR` structure, the definition of which is

```
typedef struct VkXlibSurfaceCreateInfoKHR {
    VkStructureType                 sType;
    const void*                     pNext;
    VkXlibSurfaceCreateFlagsKHR     flags;
    Display*                        dpy;
    Window                          window;
} VkXlibSurfaceCreateInfoKHR;
```

The `sType` field of `VkXlibSurfaceCreateInfoKHR` should be set to `VK_STRUCTURE_TYPE_XLIB_SURFACE_CREATE_INFO_KHR`, and `pNext` should be set to `nullptr`. The `flags` field is reserved and should be set to zero.

The `dpy` field is the Xlib `Display` representing the connection to the X server, and `window` is the Xlib `Window` handle to the window with which the new surface will be associated.

If **`vkCreateXlibSurfaceKHR()`** requires any host memory, it will use the host memory allocator passed in `pAllocator`. If `pAllocator` is `nullptr`, then an internal allocator will be used.

If surface creation is successful, the resulting `VkSurface` handle is written into the variable pointed to by `pSurface`.

Presentation with Xcb

Xcb is a slightly lower-level interface to the X protocol than is provided by Xlib and may be a better choice for applications that wish to achieve lower latency. As with Xlib and the other platforms, before creating objects for presentation on an Xcb system, we need to determine whether any of the queues on a physical device support presentation. To do this, call **`vkGetPhysicalDeviceXcbPresentationSupportKHR()`**, the prototype of which is

```
VkBool32 vkGetPhysicalDeviceXcbPresentationSupportKHR(
    VkPhysicalDevice                            physicalDevice,
    uint32_t                                    queueFamilyIndex,
    xcb_connection_t*                           connection,
    xcb_visualid_t                              visual_id);
```

The physical device being queried is passed in `physicalDevice`, and the index of the queue family is passed in `queueFamilyIndex`. The connection to the X server is passed in `connection`. Again, presentation capability is reported on a per-visual ID basis, and the visual ID being queried is passed in `visual_id`.

Once you have determined that at least one queue family on the device supports presentation in the visual ID of your choice, you can create a surface into which to render using **`vkCreateXcbSurfaceKHR()`**, the prototype of which is

```
VkResult vkCreateXcbSurfaceKHR(
    VkInstance                              instance,
    const VkXcbSurfaceCreateInfoKHR*        pCreateInfo,
    const VkAllocationCallbacks*            pAllocator,
    VkSurfaceKHR*                           pSurface);
```

The Vulkan instance is passed in `instance`, and the remaining parameters controlling creation of the surface are passed through an instance of the `VkXcbSurfaceCreateInfoKHR` structure pointed to by `pCreateInfo`. The definition of `VkXcbSurfaceCreateInfoKHR` is

```
typedef struct VkXcbSurfaceCreateInfoKHR {
    VkStructureType             sType;
    const  void*                pNext;
    VkXcbSurfaceCreateFlagsKHR  flags;
    xcb_connection_t*           connection;
    xcb_window_t                window;
} VkXcbSurfaceCreateInfoKHR;
```

The `sType` field for `VkXcbSurfaceCreateInfoKHR` should be set to `VK_STRUCTURE_TYPE_XCB_SURFACE_CREATE_INFO_KHR`, and `pNext` should be set to `nullptr`. The `flags` field is reserved and should be set to zero. The connection to the X server is passed in the `connection` field, and the handle to the window is passed in `window`.

If **vkCreateXcbSurfaceKHR()** is successful, it will write the handle to the new surface into the variable pointed to by `pSurface`. If it needs any host memory to construct the handle and `pAllocator` is not `nullptr`, then it will use your allocator to request that memory.

Swap Chains

Regardless of which platform you're running on, the resulting `VkSurfaceKHR` handle refers to Vulkan's view of a window. In order to actually present anything to that surface, it's necessary to create a special image that can be used to store the data in the window. On most platforms, this type of image is either owned by or tightly integrated with the window system, so rather than creating a normal Vulkan image object, we use a second object called a *swap chain* to manage one or more image objects.

Swap-chain objects are used to ask the native window system to create one or more images that can be used to present into a Vulkan surface. This is exposed using the `VK_KHR_swapchain` extension. Each swap-chain object manages a set of images, usually in some form of ring buffer. The application can ask the swap chain for the next available image, render into it, and then hand the image back to the swap chain ready for display. By managing presentable images in a ring or queue, one image can be presented to the display while another is being drawn to by the application, overlapping the operation of the window system and application.

To create a swap-chain object, call **vkCreateSwapchainKHR()**, the prototype of which is

```
VkResult vkCreateSwapchainKHR(
    VkDevice                            device,
```

```
    const VkSwapchainCreateInfoKHR*         pCreateInfo,
    const VkAllocationCallbacks*            pAllocator,
    VkSwapchainKHR*                         pSwapchain);
```

The device with which the swap chain is to be associated is passed in `device`. The resulting swap chain can be used with any of the queues on `device` that support presentation. The information about the swap chain is passed in an instance of the `VkSwapchainCreateInfoKHR` structure, the address of which is passed in `pCreateInfo`. The definition of `VkSwapchainCreateInfoKHR` is

[Click here to view code image](#)

```
typedef struct VkSwapchainCreateInfoKHR {
    VkStructureType                 sType;
    const void*                     pNext;
    VkSwapchainCreateFlagsKHR       flags;
    VkSurfaceKHR                    surface;
    uint32_t                        minImageCount;
    VkFormat                        imageFormat;
    VkColorSpaceKHR                 imageColorSpace;
    VkExtent2D                      imageExtent;
    uint32_t                        imageArrayLayers;
    VkImageUsageFlags               imageUsage;
    VkSharingMode                   imageSharingMode;
    uint32_t                        queueFamilyIndexCount;
    const uint32_t*                 pQueueFamilyIndices;
    VkSurfaceTransformFlagBitsKHR   preTransform;
    VkCompositeAlphaFlagBitsKHR     compositeAlpha;
    VkPresentModeKHR                presentMode;
    VkBool32                        clipped;
    VkSwapchainKHR                  oldSwapchain;
} VkSwapchainCreateInfoKHR;
```

The `sType` field for `VkSwapchainCreateInfoKHR` should be set to `VK_STRUCTURE_TYPE_SWAPCHAIN_CREATE_INFO_KHR`, and `pNext` should be set to `nullptr`. The `flags` field is reserved for use in future versions of the `VK_KHR_swapchain` extension and should be set to zero.

The surface to which the new swap chain will present is passed in `surface`. This should be a surface created with one of the the platform-specific surface creation functions such as **vkCreateWin32SurfaceKHR()** or **vkCreateXlibSurfaceKHR()**. The number of images in the swap chain is passed in `minImageCount`. For example, to enable double or triple buffering, set `minImageCount` to 2 or 3, respectively. Setting `minImageCount` to 1 represents a request to render to the front buffer or directly to the display. Some platforms don't support this (and may not even support double buffering). To determine the minimum and maximum number of images supported in a swap chain, call **vkGetPhysicalDeviceSurfaceCapabilitiesKHR()**, which is discussed later in this section.

Note that setting `minImageCount` to 2 means that you'll have a single front buffer and a single back buffer. After triggering presentation of a finished back buffer, you won't be able to begin rendering to the other buffer until the presentation has completed. For best performance, possibly at the price of some latency, you should set `minImageCount` to at least 3 if the device supports it.

The format and color space of the presentable images is specified in `imageFormat` and `imageColorSpace`. The format must be a Vulkan format for which the device reports the presentation capability. `imageColorSpace` is a member of the `VkColorSpaceKHR` enumeration, the only member of which is `VK_COLORSPACE_SRGB_NONLINEAR_KHR`, which means that the presentation engine can expect sRGB nonlinear data, if the `imageFormat` member indicates an sRGB format image.

The `imageExtent` field specifies the dimensions of the images in the swap chain, in pixels, and `imageArrayLayers` field specifies the number of layers in each image. This can be used to render to a layered image and then present specific layers of it to the user. The `imageUsage` field is a collection of the standard `VkImageUsageFlags` enumeration specifying how the images will be used (in addition to as present sources). For example, if you want to render to the image as a normal color attachment, you would include `VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT`, and if you want to write directly to it with a shader, you would include `VK_IMAGE_USAGE_STORAGE_BIT`.

The set of usages that are included in `imageUsage` must be selected from the usages supported for swap-chain images. This is determined by calling **`vkGetPhysicalDeviceSurfaceCapabilitiesKHR()`**.

The `sharingMode` field specifies how the images are to be shared across queues. If the image is going to be used by only one queue at a time (which is likely, as presentable images are generally write-only), then set this to `VK_SHARING_MODE_EXCLUSIVE`. If the image is likely to be used across multiple queues, then this can be set to `VK_SHARING_MODE_CONCURRENT`. In this case, `pQueueFamilyIndices` should be a pointer to an array of indices of the queues with which the images will be used, and `queueFamilyIndexCount` is the length of this array, in elements. When `sharingMode` is `VK_SHARING_MODE_EXCLUSIVE`, these two fields are ignored.

The `preTransform` field specifies how the images should be transformed prior to presentation to the user. This allows images to be rotated or flipped (or both) to accommodate things like portrait displays and rear-projection systems. It is a bitwise combination of a selection of members of the `VkSurfaceTransformFlagBitsKHR` enumeration.

The `compositeAlpha` field controls how alpha composition is handled by the window system. This is a member of the `VkCompositeAlphaFlagBitsKHR` enumeration. If this is set to `VK_COMPOSITE_ALPHA_OPAQUE_BIT_KHR`, then the alpha channel of the presentable image (if it exists) is ignored and treated as though it contains constant 1.0 values. Other values of `compositeAlpha` allow partially transparent images to be composited by the native window system.

The `presentMode` field controls synchronization with the window system and the rate at which the images are presented to the surface. The available modes are

- `VK_PRESENT_MODE_IMMEDIATE_KHR`: When presentation is scheduled, the image is presented to the user as soon as possible, without waiting for any external events such as vertical blanking. This provides the highest possible frame rate but can introduce tearing or other artifacts.

- `VK_PRESENT_MODE_MAILBOX_KHR`: When a new image is presented, it is marked as the pending image, and at the next opportunity (probably after the next vertical refresh), the system will display it to the user. If a new image is presented before this happens, that image will be shown, and the previously presented image will be discarded.

- `VK_PRESENT_MODE_FIFO_KHR`: Images to be presented are stored in an internal queue and shown to the user in order. A new image is taken from the queue at regular intervals (usually after each vertical refresh).

- `VK_PRESENT_MODE_FIFO_RELAXED_KHR`: This mode behaves similarly to `VK_PRESENT_MODE_FIFO_KHR`, except that if the queue is empty and vertical refresh occurs, the next image posted to the queue will be displayed immediately, similarly to `VK_PRESENT_MODE_IMMEDIATE_KHR`. This allows an application running faster than the vertical refresh rate to avoid tearing while still going as fast as possible if it can't keep up in places.

As a general rule, if you want to run with vertical sync (vsync) on, select `VK_PRESENT_MODE_FIFO_KHR`, and if you want to run as fast as possible, select `VK_PRESENT_MODE_IMMEDIATE_KHR` or `VK_PRESENT_MODE_MAILBOX_KHR`. `VK_PRESENT_MODE_IMMEDIATE_KHR` will show visible tearing in most cases but provides the lowest possible latency. `VK_PRESENT_MODE_MAILBOX_KHR` continues to flip at regular intervals, producing a maximum latency of one vertical refresh, but will not exhibit tearing.

The `clipped` member of `VkSwapchainCreateInfoKHR` is used to optimize cases where not all of the surface might be visible. For example, if the surface to which the images will be presented represents a window that might be covered or partially off the screen, it may be possible to avoid rendering the parts of it that the user will never see. When `clipped` is `VK_TRUE`, Vulkan can eliminate those parts of the image from rendering operations. When `clipped` is `VK_FALSE`, Vulkan will render the entire image, regardless of whether it's visible or not.

Finally, the `oldSwapchain` field of `VkSwapchainCreateInfoKHR` may be used to pass an existing swap chain associated with the surface back to Vulkan for recycling. This is used when one swap chain is being replaced by another, such as when a window is resized and the swap chain needs to be reallocated with larger images.

The parameters contained in the `VkSwapchainCreateInfoKHR` structure must all conform to the capabilities of the suface, which you can determine by calling **vkGetPhysicalDeviceSurfaceCapabilitiesKHR()**, the prototype of which is

[**Click here to view code image**](#)

```
VkResult vkGetPhysicalDeviceSurfaceCapabilitiesKHR(
    VkPhysicalDevice                            physicalDevice,
    VkSurfaceKHR                                surface,
    VkSurfaceCapabilitiesKHR*                   pSurfaceCapabilities);
```

The physical device that owns the surface is passed in `physicalDevice`, and the surface whose capabilities are being queried is passed in `surface`. **vkGetPhysicalDeviceSurfaceCapabilitiesKHR()** then returns information about the surface in an instance of the `VkSurfaceCapabilitiesKHR` structure, a pointer to which is provided through the `pSurfaceCapabilities` parameter. The definition of `VkSurfaceCapabilitiesKHR` is

[**Click here to view code image**](#)

```
typedef struct VkSurfaceCapabilitiesKHR {
    uint32_t                        minImageCount;
    uint32_t                        maxImageCount;
```

```
    VkExtent2D                          currentExtent;
    VkExtent2D                          minImageExtent;
    VkExtent2D                          maxImageExtent;
    uint32_t                            maxImageArrayLayers;
    VkSurfaceTransformFlagsKHR          supportedTransforms;
    VkSurfaceTransformFlagBitsKHR       currentTransform;
    VkCompositeAlphaFlagsKHR            supportedCompositeAlpha;
    VkImageUsageFlags                   supportedUsageFlags;
} VkSurfaceCapabilitiesKHR;
```

The number of images in the swap chain must fall between the `minImageCount` and `maxImageCount` parameters of the surface's capabilities. The current size of the surface at the time of the query is returned in `currentExtent`. If the surface is resizable (such as a sizeable window on a desktop), then the smallest and largest sizes that the surface can become are returned in `minImageExtent` and `maxImageExtent`. If the surface supports presentation from array images, the maximum number of layers in those images is returned in `minArrayLayers`.

Some surfaces support performing transformations on images as they are presented. For example, an image might be flipped or rotated to accommodate presentation to displays or other devices that are at nonstandard angles. The set of supported transforms is returned in the `supportedTransforms` field of `VkSurfaceCapabilitiesKHR` and is a bitfield made up of a selection of members of the `VkSurfaceTransformFlagBitsKHR` enumeration. One of those bits is set in `currentTransform`, which contains the current transform applied when the query is made.

If the surface supports composition, then the supported composition modes are contained as a combination of flags from the `VkCompositeAlphaFlagBitsKHR` enumeration in the `supportedCompositeAlpha` field.

Finally, the allowed usage for the images created through a swap chain on this surface is returned in `supportedUsageFlags`.

Once you have a swap chain associated with a surface to which you want to present, you need to get handles to the images representing the items in the chain. To do this, call **vkGetSwapchainImagesKHR()**, the prototype of which is

[Click here to view code image](#)

```
VkResult vkGetSwapchainImagesKHR(
    VkDevice                            device,
    VkSwapchainKHR                      swapchain,
    uint32_t*                           pSwapchainImageCount,
    VkImage*                            pSwapchainImages);
```

The device that owns the swap chain should be passed in `device`, and the swap chain from which you are requesting images should be passed in `swapchain`. `pSwapchainImageCount` points to a variable that will contain the number of images received. If `pSwapchainImages` is `nullptr`, then the initial value of `pSwapchainImageCount` will be ignored, and the variable will instead be overwritten with the number of swap-chain images in the swap-chain object. If `pSwapchainImages` is not `nullptr`, then it should be a pointer to an array of `VkImage` handles that will be filled with the images from the swap chain. The intial value of the variable pointed to by `pSwapchainImageCount` is the length of the array, and it will be overwritten with the number of images actually placed in the array.

Because when you create the swap chain, you get to specify only the *minimum* number of images in the swap chain, not the exact number of images, you need to use **vkGetSwapchainImagesKHR()** to determine how many images there really are in a swap chain, even if you just created it. Listing 5.1 demonstrates how to create a swap chain for an existing surface, query the number of images in it, and then query the actual image handles.

Listing 5.1: Creating a Swap Chain

Click here to view code image

```
VkResult result;

// First, we create the swap chain.
VkSwapchainCreateInfoKHR swapChainCreateInfo =
{
    VK_STRUCTURE_TYPE_SWAPCHAIN_CREATE_INFO_KHR, // sType
    nullptr,                                     // pNext
    0,                                           // flags
    m_mainSurface,                               // surface
    2,                                           // minImageCount
    VK_FORMAT_R8G8B8A8_UNORM,                    // imageFormat
    VK_COLORSPACE_SRGB_NONLINEAR_KHR,            // imageColorSpace
    { 1024, 768 },                               // imageExtent
    1,                                           // imageArrayLayers
    VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT,         // imageUsage
    VK_SHARING_MODE_EXCLUSIVE,                   // imageSharingMode
    0,                                           // queueFamilyIndexCount
    nullptr,                                     // pQueueFamilyIndices
    VK_SURFACE_TRANSFORM_INHERIT_BIT_KHR,        // preTransform
    VK_COMPOSITE_ALPHA_OPAQUE_BIT_KHR,           // compositeAlpha
    VK_PRESENT_MODE_FIFO_KHR,                    // presentMode
    VK_TRUE,                                     // clipped
    m_swapChain                                  // oldSwapchain
};

result = vkCreateSwapchainKHR(m_logicalDevice,
                              &swapChainCreateInfo,
                              nullptr,
                              &m_swapChain);

// Next, we query the swap chain for the number of images it actually
contains.
uint32_t swapChainImageCount = 0;

if (result == VK_SUCCESS)
{
    result = vkGetSwapchainImagesKHR(m_logicalDevice,
                                     m_swapChain,
                                     &swapChainImageCount,
                                     nullptr);
}

if (result == VK_SUCCESS)
```

```
{
    // Now we resize our image array and retrieve the image handles from
the
    // swap chain.
    m_swapChainImages.resize(swapChainImageCount);

    result = vkGetSwapchainImagesKHR(m_logicalDevice,
                                     m_swapChain,
                                     &swapChainImageCount,
                                     m_swapChainImages.data());
}

    return result == VK_SUCCESS? m_swapChain : VK_NULL_HANDLE;
```

Note that the code in [Listing 5.1](#) contains many hard-coded values. In a more robust application, you should call **vkGetPhysicalDeviceSurfaceCapabilitiesKHR()** to determine the capabilities of the device with respect to presenting the surface and the capabilities of the surface to support parameters such as the transform mode, number of images in the swap chain, and so on.

In particular, the surface format chosen in the `imageFormat` field of `VkSwapchainCreateInfoKHR` must be a format that is supported by the surface. To determine which formats can be used for swap chains associated with a surface, call **vkGetPhysicalDeviceSurfaceFormatsKHR()**, the prototype of which is

[Click here to view code image](#)

```
VKAPI_ATTR VkResult VKAPI_CALL vkGetPhysicalDeviceSurfaceFormatsKHR(
    VkPhysicalDevice                            physicalDevice,
    VkSurfaceKHR                                surface,
    uint32_t*                                   pSurfaceFormatCount,
    VkSurfaceFormatKHR*                         pSurfaceFormats);
```

The physical device that you are querying is passed in `physicalDevice`, and the surface to which you want to present is passed in `surface`. If `pSurfaceFormats` is `nullptr`, then the variable pointed to by `pSurfaceFormatCount` is overwritten with the number of formats supported by the surface. If `pSurfaceFormats` is not `nullptr`, then it is a pointer to an array of `VkSurfaceFormatKHR` structures large enough to receive the number of formats supported by the surface. In this case, the number of elements in the array is passed as the initial value of the variable pointed to by `pSurfaceFormats`, and this is overwritten with the number of formats actually written to the array.

The definition of `VkSurfaceFormatKHR` is

[Click here to view code image](#)

```
typedef struct VkSurfaceFormatKHR {
    VkFormat           format;

    VkColorSpaceKHR    colorSpace;
} VkSurfaceFormatKHR;
```

The `format` field of `VkSurfaceFormatKHR` is the format of pixels in memory for the surface, and `colorSpace` is the supported color space. At present, the only defined color space is `VK_COLORSPACE_SRGB_NONLINEAR_KHR`.

In some cases, devices will support presentation from almost any format. This is generally true of compositing systems that use the image you've rendered to as an input to some further processing. However, other devices may support presenting from a very limited set of surface formats—perhaps only a single format for a particular surface. This is likely to be the case when you are presenting directly to a display device.

The images you get back from a call to **vkGetSwapchainImagesKHR()** aren't immediately usable. Before you can write any data into them, you need to *acquire* the next available image by using a call to **vkAcquireNextImageKHR()**. This function retrieves the index of the next image in the swap chain that your application should render to. Its prototype is

[**Click here to view code image**](#)

```
VkResult vkAcquireNextImageKHR(
    VkDevice                            device,
    VkSwapchainKHR                      swapchain,
    uint64_t                            timeout,
    VkSemaphore                         semaphore,
    VkFence                             fence,
    uint32_t*                           pImageIndex);
```

The `device` parameter is the device that owns the swap chain, and `swapchain` is the handle to the swap chain to retrieve the next swap-chain image index from.

vkAcquireNextImageKHR() waits for a new image to become available before returning to the application. `timeout` specifies the time, in nanoseconds, that it will wait before returning. If the timeout is exceeded, then **vkAcquireNextImageKHR()** will return `VK_NOT_READY`. By setting `timeout` to 0, you can implement nonblocking behavior whereby **vkAcquireNextImageKHR()** will either return a new image immediately or return `VK_NOT_READY` to indicate that it would block if called with a nonzero timeout.

The index of the next image into which the application should render will be written into the variable pointed to by `pImageIndex`. The presentation engine might still be reading data from the image, so in order to synchronize access to the image, the `semaphore` parameter can be used to pass the handle of a semaphore that will become signaled when the image can be rendered to, or the `fence` parameter can be used to pass the handle to a fence that will become signaled when it is safe to render to the image.

Semaphores and fences are two of the synchronization primitives supported by Vulkan. We will cover synchronization primitives in more detail in [Chapter 11](#), "[Synchronization](#)."

Full-Screen Surfaces

The platform-specific extensions mentioned in the previous section allow a `VkSurface` object to be created that represents a native window owned by the operating system or window system. These extensions are typically used to render into a window that is visible on a desktop. Although it is often possible to create a window with no border that covers an entire display, it is often more efficient to render directly to a display instead.

This functionality is provided by the `VK_KHR_display` and `VK_KHR_display_swapchain` extensions. These extensions provide a platform-independent mechanism for discovering displays attached to a system, determining their properties and supported modes, and so on.

If the Vulkan implementation supports `VK_KHR_display`, you can discover the number of display devices attached to a physical device by calling **`vkGetPhysicalDeviceDisplayPropertiesKHR()`**, the prototype of which is

[**Click here to view code image**](#)

```
VkResult vkGetPhysicalDeviceDisplayPropertiesKHR(
    VkPhysicalDevice                         physicalDevice,
    uint32_t*                                pPropertyCount,
    VkDisplayPropertiesKHR*                  pProperties);
```

Displays are attached to physical devices, and the physical device whose displays you want information about is passed in the `physicalDevice` parameter. `pPropertyCount` is a pointer to a variable that will be overwritten with the number of physical devices attached to the display. If `pProperties` is `nullptr`, then the initial value of the variable pointed to by `pPropertyCount` is ignored, and it is simply overwritten with the total number of displays attached to the device. However, if `pPropertyCount` is not `nullptr`, then it is a pointer to an array of `VkDisplayPropertiesKHR` structures. The length of this array is passed as the initial value of the variable pointed to by `pPropertyCount`. The definition of `VkDisplayPropertiesKHR` is

[**Click here to view code image**](#)

```
typedef struct VkDisplayPropertiesKHR {
    VkDisplayKHR                 display;
    const char*                  displayName;
    VkExtent2D                   physicalDimensions;
    VkExtent2D                   physicalResolution;
    VkSurfaceTransformFlagsKHR   supportedTransforms;
    VkBool32                     planeReorderPossible;
    VkBool32                     persistentContent;
} VkDisplayPropertiesKHR;
```

The `display` member of each of the `VkDisplayPropertiesKHR` structures is a handle to the display that can be used to reference it later. The `displayName` is a human-readable string describing the display. The `physicalDimensions` field gives the dimensions of the display, in millimeters, and the `physicalResolution` field gives the native resolution of the display, in pixels.

Some displays (or display controllers) support flipping or rotating images as they're displayed. If this is the case, those capabilities are reported in the `supportedTransforms` field. This bitfield is made up of the members of the `VkSurfaceTransformsFlagsKHR` enumeration described earlier.

If the display supports more than one plane, then `planeReorderPossible` will be set to `VK_TRUE` if those planes can be reordered with respect to one another. If the planes can be shown only in a fixed order, then `planeReorderPossible` will be set to `VK_FALSE`.

Finally, some displays can accept partial or infrequent updates, which in many cases can improve power efficiency. If the display does support being updated in this manner, `persistentContent` will be set to `VK_TRUE`; otherwise, it will be set to `VK_FALSE`.

All devices will support at least one plane on each connected display. A plane can display images to the user. In some cases, a device will support more than one plane that it can mix in various other planes to produce a final image. These planes are sometimes known as *overlay planes* because each plane can be overlaid on those logically beneath it. When a Vulkan application presents, it presents to one of the planes of the display. It's possible to present to multiple planes from the same application.

The supported plane count is considered to be part of the device, as it is generally the device—not the physical display—that performs composition operations to merge information from the planes into a single image. The physical device can then display a subset of its supported planes on each connected display. To determine the number and type of planes supported by a device, call **`vkGetPhysicalDeviceDisplayPlanePropertiesKHR()`**, the prototype of which is

[**Click here to view code image**](#)

```
VkResult vkGetPhysicalDeviceDisplayPlanePropertiesKHR(
    VkPhysicalDevice                        physicalDevice,
    uint32_t*                               pPropertyCount,
    VkDisplayPlanePropertiesKHR*            pProperties);
```

The physical device whose overlay capabilities to query is passed in `physicalDevice`. If `pProperties` is `nullptr`, then `pPropertyCount` is a pointer to a variable that will be overwritten with the number of display planes supported by the device. If `pProperties` is not `nullptr`, then it must be a pointer to an array of `VkDisplayPlanePropertiesKHR` structures large enough to hold information about the supported display planes. The number of elements in the array is determined from the initial value of the variable pointed to by `pPropertyCount`. The definition of `VkDisplayPlanePropertiesKHR` is

[**Click here to view code image**](#)

```
typedef struct VkDisplayPlanePropertiesKHR {
    VkDisplayKHR    currentDisplay;
    uint32_t        currentStackIndex;
} VkDisplayPlanePropertiesKHR;
```

For each display plane supported by the device, one entry is placed in the `pProperties` array. Each plane appears on a single physical display, which is represented by the `currentDisplay` member, and if the device supports more than one plane on each display, the `currentStackIndex` indicates the order in which the planes are overlaid on one another.

Some of the device's display planes may span multiple physical displays. To determine which display devices a display plane is visible on, you can call **`vkGetDisplayPlaneSupportedDisplaysKHR()`**, which is declared as

[**Click here to view code image**](#)

```
VkResult vkGetDisplayPlaneSupportedDisplaysKHR(
    VkPhysicalDevice                                physicalDevice,
    uint32_t                                        planeIndex,
    uint32_t*                                       pDisplayCount,
    VkDisplayKHR*                                   pDisplays);
```

For a given physical display, specified in `physicalDevice`, and display plane, specified in `planeIndex`, **vkGetDisplayPlaneSupportedDisplaysKHR()** writes the number of displays across which that plane is visible into the variable pointed to by `pDisplayCount`. If `pDisplays` is not `nullptr`, then the handles to those displays are written into the array to which it points.

Each display plane has a set of capabilities such as maximum resolution and whether or not it supports various composition modes, and these capabilities will vary by display mode. To determine these capabilities, call **vkGetDisplayPlaneCapabilitiesKHR()**, the prototype of which is

[**Click here to view code image**](#)

```
VkResult vkGetDisplayPlaneCapabilitiesKHR(
    VkPhysicalDevice                        physicalDevice,
    VkDisplayModeKHR                        mode,
    uint32_t                                planeIndex,
    VkDisplayPlaneCapabilitiesKHR*          pCapabilities);
```

For a given device (passed in `physicalDevice`) and display mode (a handle to which is passed in `mode`), the support for this mode supported by the plane specified in `planeIndex` is written into an instance of the `VkDisplayPlaneCapabilitiesKHR` structure, the address of which is passed in `pCapabilities`. The definition of `VkDisplayPlaneCapabilitiesKHR` is

[**Click here to view code image**](#)

```
typedef struct VkDisplayPlaneCapabilitiesKHR {
    VkDisplayPlaneAlphaFlagsKHR     supportedAlpha;
    VkOffset2D                      minSrcPosition;
    VkOffset2D                      maxSrcPosition;
    VkExtent2D                      minSrcExtent;
    VkExtent2D                      maxSrcExtent;
    VkOffset2D                      minDstPosition;
    VkOffset2D                      maxDstPosition;
    VkExtent2D                      minDstExtent;
    VkExtent2D                      maxDstExtent;
} VkDisplayPlaneCapabilitiesKHR;
```

The supported composition modes for the display plane are reported in `supportedAlpha`. This is a combination of the bits defined in `VkDisplayPlaneAlphaFlagBitsKHR`, which include

- `VK_DISPLAY_PLANE_ALPHA_OPAQUE_BIT_KHR`: The plane does not support blended composition at all, and all surfaces presented on that plane are considered to be fully opaque.

- `VK_DISPLAY_PLANE_ALPHA_GLOBAL_BIT_KHR`: The plane supports a single, global alpha value that is passed through the `globalAlpha` member of the `VkDisplaySurfaceCreateInfoKHR` structure used to create the surface.

- `VK_DISPLAY_PLANE_ALPHA_PER_PIXEL_BIT_KHR`: The plane supports per-pixel transparency that is sourced from the alpha channel of the images presented to the surfaces.

The `minSrcPosition` and `maxSrcPosition` fields specify the minimum and maximum offset of the displayable region within a presentable surface that can be displayed on the plane, and the `minSrcExtent` and `maxSrcExtent` fields specify its minimum and maximum size.

The `minDstPosition` and `maxDstPosition` fields specify the minimum and maximum offset at which the plane may be placed on the corresponding physical display, and `minDstExtent` and `maxDstExtent` indicate its physical size, in pixels, on that display.

Together, these fields allow a subset of a surface to be displayed in a window that may span one or more physical displays. This is considered to be a relatively advanced display capability, and in practice, most devices will report `minSrcPosition`, `minDstPosition`, `maxSrcPosition`, and `maxDstPosition` as the display origin and the maximum extents as the supported resolution of the display.

Each physical display may support multiple display modes. Each mode is represented by a `VkDisplayModeKHR` handle and has a number of properties. Each display can report a list of predefined display modes, which can be retrieved by calling **`vkGetDisplayModePropertiesKHR()`**, the prototype of which is

[Click here to view code image](#)

```
VkResult vkGetDisplayModePropertiesKHR(
    VkPhysicalDevice                        physicalDevice,
    VkDisplayKHR                            display,
    uint32_t*                               pPropertyCount,
    VkDisplayModePropertiesKHR*             pProperties);
```

The physical device to which the display is attached is passed in `physicalDevice`, and the display whose modes you want to query is passed in `display`. Remember that multiple displays may be connected to a single physical device and each may support a different selection of display modes. The `pPropertyCount` parameter points to a variable that will be overwritten with the number of supported display modes. The initial value of this variable is ignored if `pProperties` is `nullptr`. If `pProperties` is not `nullptr`, then it should point to an array of `VkDisplayModePropertiesKHR` structures that will be filled with information about the display modes. The definition of `VkDisplayModePropertiesKHR` is

[Click here to view code image](#)

```
typedef struct VkDisplayModePropertiesKHR {
    VkDisplayModeKHR            displayMode;
    VkDisplayModeParametersKHR    parameters;
} VkDisplayModePropertiesKHR;
```

The first member of `VkDisplayModePropertiesKHR`, `displayMode`, is a `VkDisplayModeKHR` handle to the display mode that can be used to refer to it unambiguously. The second member is an instance of the `VkDisplayModeParametersKHR` structure containing the parameters of the display mode. The definition of this structure is

[Click here to view code image](#)

```
typedef struct VkDisplayModeParametersKHR {
    VkExtent2D    visibleRegion;
    uint32_t      refreshRate;
} VkDisplayModeParametersKHR;
```

The parameters of the display mode are quite simple, containing only the extent of the display, in pixels, represented by the `visibleRegion` member of `VkDisplayModeParametersKHR` and

the refresh rate, measured in thousandths of a Hertz. Generally, an application will enumerate the display modes supported by the device to which it wishes to render and select the most appropriate one. If none of the preexisting display modes are suitable, it's also possible to create new ones by calling **vkCreateDisplayModeKHR()**, the prototype of which is

```
VkResult vkCreateDisplayModeKHR(
    VkPhysicalDevice                        physicalDevice,
    VkDisplayKHR                            display,
    const VkDisplayModeCreateInfoKHR*       pCreateInfo,
    const VkAllocationCallbacks*            pAllocator,
    VkDisplayModeKHR*                       pMode);
```

The physical device that will own the mode is passed in `physicalDevice`, and the display upon which the mode will be used is passed in `display`. If creation of the new mode is successful, a handle to it will be written into the variable pointed to by `pMode`. The parameters of the new mode are passed through a pointer to an instance of the `VkDisplayModeCreateInfoKHR` structure, the definition of which is

```
typedef struct VkDisplayModeCreateInfoKHR {
    VkStructureType              sType;
    const void*                  pNext;
    VkDisplayModeCreateFlagsKHR  flags;
    VkDisplayModeParametersKHR   parameters;
} VkDisplayModeCreateInfoKHR;
```

The `sType` field of the `VkDisplayModeCreateInfoKHR` structure should be set to `VK_STRUCTURE_TYPE_DISPLAY_MODE_CREATE_INFO_KHR`, and `pNext` should be set to `nullptr`. The `flags` field is reserved for future use and should be set to zero. The remaining parameters of the new display mode are contained in an instance of the `VkDisplayModeParametersKHR`.

Once you have determined the topology of the displays connected to the physical devices in the system, their supported planes, and their display modes, you can create a `VkSurfaceKHR` object referencing one of them, which you can use just like a surface referencing a window. To do this, call **vkCreateDisplayPlaneSurfaceKHR()**, the prototype of which is

```
VkResult vkCreateDisplayPlaneSurfaceKHR(
    VkInstance                              instance,
    const VkDisplaySurfaceCreateInfoKHR*    pCreateInfo,
    const VkAllocationCallbacks*            pAllocator,
    VkSurfaceKHR*                           pSurface);
```

vkCreateDisplayPlaneSurfaceKHR() is a function that operates at the instance level because a single display mode might span multiple planes across multiple displays, even being connected to to multiple physical devices. The parameters describing the surface are passed through an instance of the `VkDisplaySurfaceCreateInfoKHR` structure, the address of which is passed in `pCreateInfo`. The definition of `VkDisplaySurfaceCreateInfoKHR` is

```
typedef struct VkDisplaySurfaceCreateInfoKHR {
    VkStructureType                 sType;
    const void*                     pNext;
    VkDisplaySurfaceCreateFlagsKHR  flags;
    VkDisplayModeKHR                displayMode;
    uint32_t                        planeIndex;
    uint32_t                        planeStackIndex;
    VkSurfaceTransformFlagBitsKHR   transform;
    float                           globalAlpha;
    VkDisplayPlaneAlphaFlagBitsKHR  alphaMode;
    VkExtent2D                      imageExtent;
} VkDisplaySurfaceCreateInfoKHR;
```

The `sType` field of `VkDisplaySurfaceCreateInfoKHR` should be set to `VK_STRUCTURE_TYPE_DISPLAY_SURFACE_CREATE_INFO_KHR`, and `pNext` should be set to `nullptr`. The `flags` field is reserved for future use and should be set to zero. The handle to the display mode that is to be used for the new surface is passed through the `displayMode` field. This can be one of the predefined display modes returned from a call to **vkGetDisplayModePropertiesKHR()** or a user-created display mode produced from a call to **vkCreateDisplayModeKHR()**.

The plane to which the surface will be presented is passed in `planeIndex`, and the relative order in which the plane should appear when composited with other planes on the device should be passed in `planeStackIndex`. At time of presentation, the image can be flipped or rotated, assuming that the operation is supported by the display. The operation to be performed is specified in `transform`, which is a single bit selected from the `VkSurfaceTransformFlagBitsKHR` enumeration. This must be a supported transform for the video mode.

The transforms that can be applied to a surface during presentation depend on the device and surface capabilities, which can be retrieved by using a call to **vkGetPhysicalDeviceSurfaceCapabilitiesKHR()**.

If the image is to be composited on top of other planes, it is possible to set the transparency for the surface by using the `globalAlpha` and `alphaMode` fields. If `alphaMode` is `VK_DISPLAY_PLANE_ALPHA_GLOBAL_BIT_KHR`, then `globalAlpha` sets the global alpha value for composition. If `alphaMode` is `VK_DISPLAY_PLANE_ALPHA_PER_PIXEL_BIT_KHR`, then the alpha value for each pixel is taken from the presented image, and the value of `globalAlpha` is ignored. If `alphaMode` is `VK_DISPLAY_PLANE_ALPHA_OPAQUE_BIT_KHR`, then blended composition is disabled.

The `imageExtent` field specifies the size of the presentable surface. In general, for full-screen rendering, this should be the same as the extent of the display mode selected in `displayMode`.

Performing Presentation

Presentation is an operation that occurs in the context of a queue. Generally, commands executed inside command buffers submitted to a queue produce the images that are to be presented, so those images should be shown to the user only when the rendering operations that created them have completed. While a device in a system may support many queues, it is not required that all of them support presentation. Before you can use a queue for presentation to a surface, you must determine whether that queue supports presentation to that surface.

To determine whether a queue supports presentation, pass the physical device, surface and queue family to a call to **vkGetPhysicalDeviceSurfaceSupportKHR()**, the prototype of which is

```
VkResult vkGetPhysicalDeviceSurfaceSupportKHR(
    VkPhysicalDevice                            physicalDevice,
    uint32_t                                    queueFamilyIndex,
    VkSurfaceKHR                                surface,
    VkBool32*                                   pSupported);
```

The physical device to query is passed in `physicalDevice`. All queues are members of a queue family, and all members of a queue family are considered to have identical properties. Therefore, only the family of a queue is needed to determine whether that queue supports presentation. The queue-family index is passed in `queueFamilyIndex`.

The capability of a queue to present is dependent on the surface. For example, some queues may be able to present into windows owned by the operating system but have no direct access to physical hardware that controls full-screen surfaces. Therefore, the surface to which you want to present is passed in `surface`.

If **vkGetPhysicalDeviceSurfaceSupportKHR()** is successful, the ability of queues in the specified family to present to the surface specified in `surface` is written into the variable pointed to by `pSupported`—`VK_TRUE` indicating support and `VK_FALSE` indicating lack of support. If something goes wrong, **vkGetPhysicalDeviceSurfaceSupportKHR()** will return a failure code, and the value of `pSupported` will not be overwritten.

Before an image can be presented, it must be in the correct layout. This state is the `VK_IMAGE_LAYOUT_PRESENT_SRC_KHR` layout. Images are transitioned from layout to layout using image memory barriers, as discussed briefly in Chapter 2, "Memory and Resources." Listing 5.2 shows how to transition an image from `VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL` to `VK_IMAGE_LAYOUT_PRESENT_SRC_KHR` layout using an image memory barrier.

Listing 5.2: Transitioning an Image to Present Source

```
const VkImageMemoryBarrier barrier =
{
    VK_STRUCTURE_TYPE_IMAGE_MEMORY_BARRIER,     // sType
    nullptr,                                    // pNext
    VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,       // srcAccessMask
    VK_ACCESS_MEMORY_READ_BIT,                  // dstAccessMask
```

```
        VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL,     // oldLayout
        VK_IMAGE_LAYOUT_PRESENT_SRC_KHR,              // newLayout
        0,                                            // srcQueueFamilyIndex
        0,                                            // dstQueueFamilyIndex
        sourceImage,                                  // image
        {                                             // subresourceRange
            VK_IMAGE_ASPECT_COLOR_BIT,                // aspectMask
            0,                                        // baseMipLevel
            1,                                        // levelCount
            0,                                        // baseArrayLayer
            1,                                        // layerCount
        }
    };

    vkCmdPipelineBarrier(cmdBuffer,
                         VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,
                         VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT,
                         0,
                         0, nullptr,
                         0, nullptr,
                         1, &barrier);
```

Note that the image memory barrier is executed inside a command buffer and this command buffer should be submitted to a device queue for execution. Once the image is in `VK_IMAGE_LAYOUT_PRESENT_SRC_KHR` layout, it can be presented to the user by calling **vkQueuePresentKHR()**, the prototype of which is

Click here to view code image

```
VkResult vkQueuePresentKHR(
    VkQueue                             queue,
    const VkPresentInfoKHR*             pPresentInfo);
```

The queue to which the image should be submitted for presentation is specified in `queue`. The rest of the parameters to the command are passed through an instance of the `VkPresentInfoKHR` structure, the definition of which is

Click here to view code image

```
typedef struct VkPresentInfoKHR {
    VkStructureType        sType;
    const void*            pNext;
    uint32_t               waitSemaphoreCount;
    const VkSemaphore*     pWaitSemaphores;
    uint32_t               swapchainCount;
    const VkSwapchainKHR*  pSwapchains;
    const uint32_t*        pImageIndices;
    VkResult*              pResults;
} VkPresentInfoKHR;
```

The `sType` field of `VkPresentInfoKHR` should be set to `VK_STRUCTURE_TYPE_PRESENT_INFO_KHR`, and `pNext` should be set to `nullptr`. Before the images are presented, Vulkan will optionally wait on one or more semaphores to enable rendering to the images to be synchronized with the presentation operation. The number of semaphores to wait

on is passed in the `waitSemaphoreCount` member, and the `pWaitSemaphores` member points to an array of this many semaphore handles to wait on.

A single call to **vkQueuePresentKHR()** can actually present multiple images to multiple swap chains at the same time. This is useful, for example, in an application that is rendering to multiple windows at the same time. The number of images to preset is specified in `swapchainCount`. `pSwapchains` is an array of the swap-chain objects to present with.

The images presented to each of the swap chains are not referenced by their `VkImage` handles but by the indices into their arrays of swap-chain images as retrieved from the swap-chain object. For each swap chain that will be presented to, one image index is passed through the corresponding element in the array pointed to by `pImageIndices`.

Each of the separate present operations triggered by the call to **vkQueuePresentKHR()** can produce its own result code. Remember that some values of `VkResult` indicate success. `pResults` is a pointer to an array of `swapchainCount VkResult` variables that will be filled with the results of the present operations.

Cleaning Up

Regardless of the method of presentation you've used in your application, it is important to clean up correctly. First, you should destroy the swap chain to which you are presenting. To do this, call **vkDestroySwapchainKHR()**, the prototype of which is

[**Click here to view code image**](#)

```
void vkDestroySwapchainKHR(
    VkDevice                      device,
    VkSwapchainKHR                swapchain,
    const VkAllocationCallbacks*  pAllocator);
```

The device that owns the swap chain is passed in `device`, and the swap chain to destroy is passed in `swapchain`. If a host memory allocator was used to create the swap chain, then a pointer to a compatible allocator is passed in `pAllocator`.

When the swap chain is destroyed, all of the presentable images associated with the swap chain are also destroyed. Therefore, before you destroy a swap chain, you should ensure that there is no pending work that might write to any of its surfaces and that there are no pending present operations that might read from them. The easiest way to do this is to call **vkDeviceWaitIdle()** on the device. While not normally recommended, destruction of a swap chain usually does not occur in a performance-critical part of an application, so in this case, simple is best.

When images are acquired from a swap chain using **vkAcquireNextImageKHR()** or presented using **vkQueuePresentKHR()**, semaphores are passed to these functions to signal and wait on, resepectively. Care should be taken that the semaphores live long enough that the swap chain can complete any signaling operations on them before they are destroyed. To ensure this, it is best to destroy the swap chain before destroying any semaphores that might have been used with it.

Summary

In this chapter, you learned about the operations supported by Vulkan for getting images onto displays. We covered presenting to various window systems, the mechanism by which you determine which images to render into, and how to enumerate and control the physical display devices attached to a system. We briefly covered synchronization involved in presenting and will dig further into synchronization primitives later in the book. We also discussed methods for configuring display synchronization. With the information in this chapter, you should have a decent understanding of how Vulkan presents images to the user.

Chapter 6. Shaders and Pipelines

What You'll Learn in This Chapter

  • What a shader is and how it's used

  • The basics of *SPIR-V*, the Vulkan shading language

  • How to construct a shader pipeline and use it to do work

Shaders are small programs that are executed directly on the device. These are the fundamental building blocks of any complex Vulkan program. Shaders are perhaps more important to the operation of your program than the Vulkan API itself. This chapter introduces shaders and shader modules, shows how they are constructed from SPIR-V binaries, and illustrates how those binaries are generated from GLSL using standard tools. It discusses the construction of pipelines containing those shaders and other information required to run them, and then shows how to execute the shaders to do work on the device.

Shaders are the fundamental building blocks of work to be executed on the device. Vulkan shaders are represented by SPIR-V, which is a binary intermediate representation of program code. SPIR-V can be generated offline using a compiler, online directly inside your application, or by passing a high-level language to a library at runtime. The sample applications accompanying this book take the first approach: compiling the shaders offline and then loading the resulting SPIR-V binaries from disk.

The original shaders are written in GLSL using the Vulkan profile. This is a modified and enhanced version of the same shading language used with OpenGL. Most of the examples, therefore, discuss Vulkan features in terms of their representation in GLSL. However, it should be clear that Vulkan itself knows nothing of GLSL and doesn't care where SPIR-V shaders come from.

An Overview of GLSL

Although not officially part of the Vulkan specification, Vulkan shares much of its heritage with OpenGL. In OpenGL, the officially supported high-level language is GLSL—the OpenGL Shading Language. Therefore, during the design of SPIR-V, much attention was paid to ensuring that at least one high-level language would be suitable for use in the generation of SPIR-V shaders. Minor modifications were made to the GLSL specification for use with Vulkan. Some features were added to enable GLSL shaders to interact cleanly with the Vulkan system, and legacy features in OpenGL that were not carried forward into Vulkan were removed from the Vulkan profile of GLSL.

The result is a slimline version of GLSL that supports most of the features of Vulkan while enabling a high level of portability between OpenGL and Vulkan. In short, if you stick to the modern features of OpenGL in an OpenGL application, much of what you write in your shaders will compile directly into SPIR-V using the official reference compiler. Of course, you are free to write your own compilers and tools or to use a third-party compiler to produce SPIR-V modules from any language you choose.

The modifications to GLSL to allow it to be used to produce a SPIR-V shader suitable for use with Vulkan are documented in the `GL_KHR_vulkan_glsl` extension.

In this section, we provide a brief overview of GLSL. It is assumed that the reader is somewhat familiar with high-level shading languages in general and is capable of researching GLSL in more depth if needed.

Listing 6.1 shows the simplest possible GLSL shader. It is simply an empty function, returning **void**, that does absolutely nothing. It's actually a valid shader for any stage in the Vulkan pipeline, although executing it in some stages would result in some undefined behavior.

Listing 6.1: Simplest Possible GLSL Shader

```
#version 450 core

void main (void)
{
    // Do nothing!
}
```

All GLSL shaders should begin with a `#version` directive to inform the GLSL compiler which version of GLSL we're using. This allows the compiler to perform appropriate error checks and to allow certain language constructs that have been introduced over time.

When compiled to SPIR-V for use in Vulkan, a compiler should automatically define VULKAN to the version of the GL_KHR_vulkan_glsl extension in use such that you can wrap Vulkan-specific constructs or functionality in your GLSL shader in `#ifdef` VULKAN or `#if` VULKAN > {version} blocks in order to allow the same shaders to be used with OpenGL or Vulkan. Throughout this book, when Vulkan-specific features are discussed in the context of GLSL, it is assumed that the code beging written is either exclusively for Vulkan or is wrapped in the appropriate `#ifdef` preprocessor conditionals to allow it to be compiled for Vulkan.

GLSL is a C-like language, with its syntax and many of its semantics taken from C or C++. If you are a C programmer, you should be comfortable with constructs such as `for` and `while` loops; flow-control keywords such as `break` and `continue`; `switch` statements; relational operators such as ==, <, and >; the ternary operator a ? b : c; and so on. All of these are available in GLSL shaders.

The fundamental data types in GLSL are signed and unsigned integer and floating-point values, denoted as `int`, `uint`, and `float`, respectively. Double-precision floating-point values are also supported using the `double` data type. Inside GLSL, they have no defined bit width, much as in C. GLSL has no stdint analog, so defining a specific bit width for variables is not supported, although the GLSL and SPIR-V specifications do provide some minimum guarantees for the range and precision of numeric representations used by Vulkan. However, bit width and layouts are defined for variables sourced from and written to memory. Integers are stored in memory in twos-component form, and floating-point variables follow IEEE conventions wherever possible, aside from minor differences in precision requirements, handling of denormals and not-a-number (NaN) values, and so on.

In addition to the basic scalar integer and floating-point data types, GLSL represents short vectors of up to four components and small matrices of up to 4 × 4 elements as first-class citizens in the language. Vectors and matrices of the fundamental data types (signed and unsigned integers and floating-point scalars) can be declared. The `vec2`, `vec3`, and `vec4` types, for example, are vectors of two, three, and four floating-point values, respectively. Integer vectors are notated using the i or u

prefixes for signed and unsigned integers, respectively. Thus, `ivec4` is a vector of four signed integers, and `uvec4` is a vector of four unsigned integers. The d prefix is used to denote double-precision floating-point. Thus, `dvec4` is a vector of four double-precision floating-point values.

Matrices are written using the form `matN` or `matNxM`, representing N × N square matrices and N × M rectangular matrices, respectively. The d prefix may also be used with matrix data types to form double-precision matrices. Therefore, `dmat4` is a 4 × 4 matrix of double-precision floating-point values. Matrices of integers, however, are not supported. Matrices are considered to be column-primary and may be treated as arrays of vectors. Thus, writing `m[3]` where m is of type `mat4` yields a four-element vector of floating-point values (a `vec4`) representing the last column of m.

The Boolean type is also a first-class citizen in GLSL and can be formed into vectors (but not matrices), just like floating-point and integer variables. Boolean variables are written using the `bool` type. Relational operators comparing vectors produce vectors of Booleans, with each element representing one of the comparison results. The special built-in functions `any()` and `all()` are used to produce a single Boolean expressing whether any element of the source vector is true and whether all of the elements of the source vector are true, respectively.

Data produced by the system is passed into GLSL shaders through built-in variables. Examples are variables such as `gl_FragCoord`, `gl_VertexIndex`, and so on, which will each be introduced as we reach the relevant parts of Vulkan throughout this book. Built-in variables often have specific semantics, and reading and writing them can change how a shader behaves.

User-specified data is normally passed into shaders through memory. Variables can be bound together into blocks, which can then be bound to device-accessible resources backed by memory that your application can write into. This allows you to pass large amounts of data to shaders. For smaller but more frequently updated data, special variables called *push constants* are available and will be covered later in the book.

GLSL provides a very large number of built-in functions. In contrast to C, however, GLSL has no header files, and it is not necessary to `#include` anything. Rather, the GLSL equivalent to a standard library is automatically provided by the compiler. It includes a large suite of math functions, functions for accessing textures, and special functions such as flow-control functions that control the execution of the shader on the device.

An Overview of SPIR-V

SPIR-V shaders are embedded in *modules*. Each module can contain one or many shaders. Each shader has an entry point that has a name and a shader type, which is used to define which shading *stage* the shader runs in. The entry point is where the shader begins execution when it is run. A SPIR-V module is passed to Vulkan along with creation information, and Vulkan returns an object representing that module. The module can then be used to construct a *pipeline*, which is a fully compiled version of a single shader along with information required to run it on the device.

Representation of SPIR-V

SPIR-V is the only officially supported shading language for Vulkan. It is accepted at the API-level and is ultimately used to construct pipelines, which are objects that configure a Vulkan device to do work for your application.

SPIR-V was designed to be easy for tools and drivers to consume. This improves portability by reducing the variability between implementations. The native representation of a SPIR-V module is a

stream of 32-bit words stored in memory. Unless you are a tool writer or plan to generate SPIR-V yourself, it is unlikely that you will deal with the binary encoding of SPIR-V directly. Rather, you will either look at a textual representation of SPIR-V or generate SPIR-V using a tool such as `glslangvalidator`, the official Khronos GLSL compiler.

Saving the shader in [Listing 6.1](#) as a text file with the `.comp` extension tells `glslangvalidator` that the shader is to be compiled as a compute shader. We can then compile this shader using `glslangvalidator` with the command line

[Click here to view code image](#)

```
glslangvalidator simple.comp -o simple.spv
```

This produces a SPIR-V binary named `simple.spv`. We can disassemble the binary using the SPIR-V disassembler, `spirv-dis`, which outputs a human-readable disassembly of the generated SPIR-V binary. This is shown in [Listing 6.2](#).

Listing 6.2: Simplest SPIR-V

[Click here to view code image](#)

```
; SPIR-V
; Version: 1.0
; Generator: Khronos Glslang Reference Front End; 1
; Bound: 6
; Schema: 0
               OpCapability Shader
          %1 = OpExtInstImport "GLSL.std.450"
               OpMemoryModel Logical GLSL450
               OpEntryPoint GLCompute %4 "main"
               OpExecutionMode %4 LocalSize 1 1 1
               OpSource GLSL 450
               OpName %4 "main"
          %2 = OpTypeVoid
          %3 = OpTypeFunction %2
          %4 = OpFunction %2 None %3
          %5 = OpLabel
               OpReturn
               OpFunctionEnd
```

You can see that the text form of SPIR-V looks like a strange dialect of assembly language. We can step through this disassembly and see how it relates to the original input. Each line of the output assembly represents a single SPIR-V instruction, possibly made up of multiple tokens.

The first instruction in the stream, `OpCapability Shader`, requests that the `Shader` capability be turned on. SPIR-V functionality is roughly divided into related groups of instructions and features. Before your shader can use any of these features, it must declare that it will be using the capability of which the feature is part. The shader in [Listing 6.2](#) is a graphics shader and therefore uses the `Shader` capability. This is the most fundamental capability. Without this, we cannot compile graphics shaders. As we introduce more SPIR-V and Vulkan functionality, we will introduce the various capabilities that each feature depends on.

Next, we see `%1 = OpExtInstImport` **"GLSL.std.450"**. This is essentially importing an additional set of instructions corresponding to the functionality included in GLSL version 450, which is what the original shader was written in. Notice that this instruction is preceeded by `%1 =`. This *names* the result of the instruction by assigning it an ID. The result of `OpExtInstImport` is effectively a library. When we want to call functions in this library, we do so using the `OpExtInst` instruction, which takes both a library (the result of the `OpExtInstImport` instruction) and an instruction index. This allows the SPIR-V instruction set to be arbitrarily extended.

Next, we see some additional declarations. `OpMemoryModel` specifies the working memory model for this module, which in this case is the logical memory model corresponding to GLSL version 450. This means that all memory access is performed through resources rather than a physical memory model, which accesses memory through pointers.

Next is the declaration of an entry point in the module. The `OpEntryPoint GLCompute %4` **"main"** instruction means that there is an available entry point corresponding to an OpenGL compute shader, with ID 4 exported with the function name `main`. This name is used to reference the entry point when we hand the resulting shader module back to Vulkan.

We use this ID in the subsequent instruction, `OpExecutionMode %4 LocalSize 1 1 1`, which defines the execution group size of this shader to be $1 \times 1 \times 1$ work item. This is implicit in GLSL if the local size **layout** qualifier is not present.

The next two instructions are simply informational. `OpSource GLSL 450` indicates that the module was compiled from GLSL version 450, and `OpName 4` **"main"** provides a name for the token with ID 4.

Now we see the real meat of the function. First, `%2 = OpTypeVoid` declares that we want to use ID 2 as the type **void**. Everything in SPIR-V has an ID, even type definitions. Large, aggregate types can be built up by referencing sequentially smaller, simpler types. However, we need to start from somewhere, and assigning a type to **void** is where we're starting.

`%3 = OpTypeFunction %2` means that we're defining ID 3 as a function type taking **void** (previously declared as ID 2) and taking no parameters. We use this in the following line, `%4 = OpFunction %2 None %3`. This means that we're declaring ID 4 (which we previously named **"main"**) to be an instance of the function 3 (declared in the line above), returning **void** (as ID 2), and having no particular decarations. This is indicated by `None` in the instructions and can be used for things like inlining, whether the variable is constant (its constness), and so on.

Finally, we see the declaration of a label (which is unused and only a side effect of the way the compiler operates), the implicit return statement, and eventually the end of the function. This is the end of our SPIR-V module.

The complete binary dump of the shader is 192 bytes long. SPIR-V is quite verbose, as 192 bytes is quite a bit longer than the original shader. However, SPIR-V makes explicit some of the things that are implicit in the original shading language. For example, declaring a memory model is not necessary in GLSL because it supports only a logical memory model. Further, there is some redundancy in the SPIR-V module as compiled here; we don't care about the name of the main function, the label with ID 5 is never actually used, and the shader imports the `GLSL.std.450` library but never actually uses it. It is possible to strip such unneeded instructions from a module. Even after this, because SPIR-V is relatively sparsely encoded, the resulting binaries are fairly easy

to compress with even a generic compressor and probably significantly more compressible with a specialized compression library.

All SPIR-V code is written in *SSA* (*single static assignment*) form, which means that every virtual register (the tokens written as `%n` in the listing above) are written exactly once. Almost every instruction that does work produces a result identifier. As we progress to more complex shaders, you will see that machine-generated SPIR-V is somewhat unwieldy, and because of its verbosity and enforced SSA form, it is quite difficult to write by hand. It is strongly recommended that you use a compiler to generate SPIR-V offline or online by using the compiler as a library that your application can link to.

If you do plan to generate or interpret SPIR-V modules yourself, you can use the defined binary encoder to build tools to parse or create them. However, they have a well-defined binary storage format that is explained later in this chapter.

All SPIR-V modules begin with a *magic number* that can be used to weakly validate that the binary blob is, in fact, a SPIR-V module. This magic number is `0x07230203` when viewed as a native unsigned integer. This number can also be used to deduce the endianness of the module. Because each SPIR-V token is a 32-bit word, if a SPIR-V module is passed by disk or network to a host of a different endianness, the bytes within the word are swapped, and its value is changed. For example, if a SPIR-V module stored in little-endian format is loaded by a big-endian host, the magic number will be read as `0x03022307`, so the host knows to swap the byte order of the words in the module.

Following the magic number are several more words that describe properties of the module. First is the version number of SPIR-V used in the module. This is encoded using the bytes of the 32-bit word where bits 16 through 23 contain the major version and bits 8 through 15 contain the minor version. SPIR-V 1.0 therefore uses the encoding `0x00010000`. The remaining bits in the version number are reserved. Next is a token containing the version of the tool that generated the SPIR-V module. This value is tool-dependent.

Next is the maximum number of IDs used in the module. All variables, functions, and other components of the SPIR-V module are assigned an ID less than this number, so including it up front allows tools consuming SPIR-V to allocate arrays of data structures to hold them in rather than allocating those structures on the fly. The last word in the header is reserved and should be set to zero. Following this is the stream of instructions.

Handing SPIR-V to Vulkan

Vulkan doesn't care too much where the SPIR-V shaders and modules come from. Typically, they will be compiled offline as part of building your application, compiled using an online compiler, or generated directly in your application. Once you have a SPIR-V module, you need to hand it to Vulkan so that a *shader module* object can be created from it. To do this, call **vkCreateShaderModule()**, the prototype of which is

[Click here to view code image](#)

```
VkResult vkCreateShaderModule (
    VkDevice                          device,
    const VkShaderModuleCreateInfo*   pCreateInfo,
    const VkAllocationCallbacks*      pAllocator,
    VkShaderModule*                   pShaderModule);
```

As with all Vulkan object creation functions, **vkCreateShaderModule()** takes a device handle as input along with a pointer to a structure containing a descrtiption of the object being created. In this case, this is a VkShaderModuleCreateInfo structure, the definition of which is

```
typedef struct VkShaderModuleCreateInfo {
    VkStructureType              sType;
    const void*                  pNext;
    VkShaderModuleCreateFlags    flags;
    size_t                       codeSize;
    const uint32_t*              pCode;
} VkShaderModuleCreateInfo;
```

The sType field of VkShaderModuleCreateInfo should be set to VK_STRUCTURE_TYPE_SHADER_MODULE_CREATE_INFO, and pNext should be set to nullptr. The flags field is reserved for future use and should be set to zero. The codeSize field contains the size of the SPIR-V module in bytes, the code for which is passed in pCode.

If the SPIR-V code is valid and understood by Vulkan, then **vkCreateShaderModule()** will return VK_SUCCESS and place a handle to a new shader module in the variable pointed to by pShaderModule. You can then use the shader module to create pipelines, which are the final form of the shader used to do work on the device.

Once you are done with a shader module, you should destroy it to free its resources. This is performed by calling **vkDestroyShaderModule()**, the prototype of which is

```
void vkDestroyShaderModule (
    VkDevice                     device,
    VkShaderModule               shaderModule,
    const VkAllocationCallbacks* pAllocator);
```

A handle to the device that owns the shader module should be passed in device, and the shader module to be destroyed should be passed in shaderModule. Access to the shader module must be externally synchronized. No other access to the shader module needs to be externally synchronized. In particular, it is possible to use the same shader module to create multiple pipelines in parallel, as discussed in the next section. Your application only needs to ensure that the shader module is not destroyed while there are Vulkan commands executing on other threads that might access the same module.

After the module has been destroyed, its handle is immediately invalidated. However, pipelines created using the module remain valid until they are destroyed. If a host memory allocator was used when the shader module was created, then a pointer to a compatible allocator should be passed in pAllocator; otherwise, pAllocator should be set to nullptr.

Pipelines

As you read in the previous sections, Vulkan uses shader modules to represent collections of shaders. Shader modules are created by handing the module code to **vkCreateShaderModule()**, but before they can be used to do useful work on the device, you need to create a pipeline. There are two types of pipeline in Vulkan: compute and graphics. The graphics pipeline is rather complex and contains a lot of state unrelated to shaders. However, a compute pipeline is conceptually much simpler and contains essentially nothing but the shader code itself.

Compute Pipelines

Before we discuss creating a compute pipeline, we should cover the basics of compute shaders in general. The shader and its execution are the core of Vulkan. Vulkan also provides access to various fixed blocks of functionality for performing things such as copying data around and processing pixel data. However, the shader will form the core of any nontrivial application.

The compute shader provides raw access to the compute capabilities of the Vulkan device. The device can be seen as a collection of wide vector processing units that operate on related pieces of data. A compute shader is written as though it were a serial, single track of execution. However, there are hints that many such tracks of execution may run together. This is, in fact, how most Vulkan devices are constructed. Each track of execution is known as an *invocation*.

When a compute shader is executed, many invocations are started at once. The invocations are grouped into local work groups of fixed size, and then one or more such groups are launched together in what is sometimes known as a global work group. Logically, both the local and global work groups are three-dimensional. However, setting the size of any one of the three dimensions to one reduces the dimensionality of the group.

The size of the local work group is set in the compute shader. In GLSL, this is done using a layout qualifier, which is translated to the `LocalSize` decoration on the `OpExecutionMode` declaration in the SPIR-V shader passed to Vulkan. Listing 6.3 shows the size declaration applied to a GLSL shader, and Listing 6.4 shows the resulting SPIR-V disassembly, truncated for clarity.

Listing 6.3: Local Size Declaration in a Compute Shader (GLSL)

Click here to view code image

```
#version 450 core

layout (local_size_x = 4, local_size_y = 5, local_size_z 6) in;

void main(void)
{
    // Do nothing.
}
```

Listing 6.4: Local Size Declaration in a Compute Shader (SPIR-V)

Click here to view code image

```
...
         OpCapability Shader
%1 = OpExtInstImport "GLSL.std.450"
```

```
                  OpMemoryModel Logical GLSL450
                  OpEntryPoint GLCompute %4 "main"
                  OpExecutionMode %4 LocalSize 4 5 6
                  OpSource GLSL 450
    ...
```

As you can see, the `OpExecutionMode` instruction in Listing 6.4 sets the local size of the shader to {4, 5, 6} as specified in Listing 6.3.

The maximum local work group size for a compute shader is generally fairly small and is required only to be at least 128 invocations in the *x* and *y* dimensions and 64 invocations in the *z* dimension. Further, the total "volume" of the work group (the product of the limit in *x*, *y*, and *z* directions) is subject to a further limit, which is required only to be at least 128 invocations. Although many implementations support higher limits, you should always query those limits if you want to exceed the required minimums.

The maximum size of a work group can be determined from the `maxComputeWorkGroupSize` field of the `VkPhysicalDeviceLimits` structure returned from a call to **`vkGetPhysicalDeviceProperties()`**, as explained in Chapter 1, "Overview of Vulkan." Further, the maximum total number of invocations in the local work group is contained in the `maxComputeWorkGroupInvocations` field of the same structure. An implementation will likely reject a SPIR-V shader that exceeds any of those limits, although behavior can technically be undefined in this case.

Creating Pipelines

To create one or more compute pipelines, call **`vkCreateComputePipelines()`**, the prototype of which is

Click here to view code image

```
VkResult vkCreateComputePipelines (
    VkDevice                           device,
    VkPipelineCache                    pipelineCache,
    uint32_t                           createInfoCount,
    const VkComputePipelineCreateInfo*     pCreateInfos,
    const VkAllocationCallbacks*       pAllocator,
    VkPipeline*                        pPipelines);
```

The `device` parameter to **`vkCreateComputePipelines()`** is the device with which the pipelines will be used and is responsible for allocating the pipeline objects. `pipelineCache` is a handle to an object that can be used to accelerate the creation of pipeline objects and is covered later in this chapter. The parameters for the creation of each new pipeline is represented by an instance of the `VkComputePipelineCreateInfo` structure. The number of structures (and therefore the number of pipelines to create) is passed in `createInfoCount`, and the address of an array of these structures is passed in `pCreateInfos`. The definition of `VkComputePipelineCreateInfo` is

Click here to view code image

```
typedef struct VkComputePipelineCreateInfo {
    VkStructureType                    sType;
    const void*                        pNext;
```

```
    VkPipelineCreateFlags                flags;
    VkPipelineShaderStageCreateInfo      stage;
    VkPipelineLayout                     layout;
    VkPipeline                           basePipelineHandle;
    int32_t                              basePipelineIndex;
} VkComputePipelineCreateInfo;
```

The `sType` field of `VkComputePipelineCreateInfo` should be set to
`VK_STRUCTURE_TYPE_COMPUTE_PIPELINE_CREATE_INFO`, and `pNext` should be set to
`nullptr`. The `flags` field is reserved for future use, and in the current version of Vulkan, it should
be set to zero. The `stage` field is an embedded structure containing information about the shader
itself and is an instance of the `VkPipelineShaderStageCreateInfo` structure, the definition
of which is

[**Click here to view code image**](#)

```
typedef struct VkPipelineShaderStageCreateInfo {
    VkStructureType                      sType;
    const void*                          pNext;
    VkPipelineShaderStageCreateFlags     flags;
    VkShaderStageFlagBits                stage;
    VkShaderModule                       module;
    const char*                          pName;
    const VkSpecializationInfo*          pSpecializationInfo;
} VkPipelineShaderStageCreateInfo;
```

The `sType` for `VkPipelineShaderStageCreateInfo` is
`VK_STRUCTURE_TYPE_PIPELINE_SHADER_STAGE_CREATE_INFO`, and `pNext` should be
set to `nullptr`. The `flags` field is reserved in the current version of Vulkan and should be set to
zero.

The `VkPipelineShaderStageCreateInfo` structure is used for all stages of pipeline
creation. Although graphics pipelines have multiple stages (which you'll learn about in [Chapter 7](#),
"[Graphics Pipelines](#)"), compute pipelines have only a single stage, and therefore `stage` should be
set to `VK_SHADER_STAGE_COMPUTE_BIT`.

`module` is the handle to the shader module you created earlier; it contains the shader code for the
compute pipeline you want to create. Because a single shader module can contain multiple entry
points and ultimately many shaders, the entry point that represents this particular pipeline is specified
in the `pName` field of `VkPipelineShaderStageCreateInfo`. This is one of the few instances
where human-readable strings are used in Vulkan.

Specialization Constants

The final field in `VkPipelineShaderStageCreateInfo` is a pointer to an instance of the
`VkSpecializationInfo` structure. This structure contains the information required to *specialize*
a shader, which is the process of building a shader with some of its constants compiled in.

A typical Vulkan implementation will delay final code generation for pipelines until
`vkCreateComputePipelines()` is called. This allows the values of specialization constants to
be considered during the final passes of optimization over the shader. Typical uses and applications
of specialization constants include

- Producing special cases through branching: Including a condition on a Boolean specialization constant will result in the final shader taking only one branch of the **if** statement. The nontaken branch will probably be optimized away. If you have two similar versions of a shader that differ in only a couple of places, this is a good way to merge them into one.

- Special cases through switch statements: Likewise, using an integer specialization constant as the tested variable in a switch statement will result in only one of the cases ever being taken in that particular pipeline. Again, most Vulkan implementations will optimize out all the never-taken cases.

- Unrolling loops: Using an integer specialization constant as the iteration count in a **for** loop may result in the Vulkan implementation making better decisions about how to unroll the loop or whether to unroll it at all. For example, if the loop counter ends up with a value of 1, then the loop goes away and its body becomes straight-line code. A small loop iteration count might result in the compiler unrolling the loop exactly that number of times. A larger iteration count may result in the compiler unrolling the loop by a factor of the count and then looping over that unrolled section a smaller number of times.

- Constant folding: Subexpressions involving specialization constants can be folded just as with any other constant. In particular, expressions involving multiple specialization constants may fold into a single constant.

- Operator simplification: Trivial operations such as adding zero or multiplying by one disappear, multiplying by negative one can be absorbed into additions turning them to subtractions, multiplying by small integers such as two can be turned into additions or even absorbed into other operations, and so on.

In GLSL, a specialization constant is declared as a regular constant that is given an ID in a layout qualifier. Specialization constants in GLSL can be Booleans, integers, floating-point values, or composites such as arrays, structures, vectors, or matrices. When translated to SPIR-V, these become `OpSpecConstant` tokens. Listing 6.5 shows an example GLSL declaration of some specialization constants, and Listing 6.6 shows the resulting SPIR-V produced by the GLSL compiler.

Listing 6.5: Specialization Constants in GLSL

Click here to view code image

```
layout (constant_id = 0) const int numThings = 42;
layout (constant_id = 1) const float thingScale = 4.2f;
layout (constant_id = 2) const bool doThat = false;
```

Listing 6.6: Specialization Constants in SPIR-V

Click here to view code image

```
...
          OpDecorate %7 SpecId 0
          OpDecorate %9 SpecId 1
          OpDecorate %11 SpecId 2
    %6 = OpTypeInt 32 1
    %7 = OpSpecConstant %6 42
    %8 = OpTypeFloat 32
    %9 = OpSpecConstant %8 4.2
```

```
            %10 = OpTypeBool
            %11 = OpSpecConstantFalse %10
    ...
```

Listing 6.6 has been edited to remove portions unrelated to the specialization constants. As you can see, however, `%7` is declared as a specialization constant using `OpSpecConstant` of type `%6` (a 32-bit integer) with a default value of 42. Next, `%9` is declared as a specialization constant of type `%8` (a 32-bit floating-point value) with a default value of 4.2. Finally, `%11` is decalared as a Boolean value (type `%10` in this SPIR-V) with a default value of `false`. Note that Booleans are declared with either `OpSpecConstantTrue` or `OpSpecConstantFalse`, depending on whether their default value is true or false, respectively.

Note that in both the GLSL shader and the resulting SPIR-V shader, the specialization constants are assigned default values. In fact, they *must* be assigned default values. These constants may be used like any other constant in the shader. In particular, they can be used for things like sizing arrays where only compile-time constants are otherwise allowed. If new values are not included in the `VkSpecializationInfo` structure passed to **vkCreateComputePipelines()**, then those default values are used. However, the constants can be overridden by passing new values when the pipeline is created. The definition of `VkSpecializationInfo` is

Click here to view code image

```
typedef struct VkSpecializationInfo {
    uint32_t                          mapEntryCount;
    const VkSpecializationMapEntry*   pMapEntries;
    size_t                            dataSize;
    const void*                       pData;
} VkSpecializationInfo;
```

Inside the `VkSpecializationInfo` structure, `mapEntryCount` contains the number of specialization constants that are to be set, and this is the number of entries in the array of `VkSpecializationMapEntry` structures in the array pointed to by `pMapEntries`. Each of these represents a single specialization constant. The definition of `VkSpecializationMapEntry` is

Click here to view code image

```
typedef struct VkSpecializationMapEntry {
    uint32_t    constantID;
    uint32_t    offset;
    size_t      size;
} VkSpecializationMapEntry;
```

The `constantID` field is the ID of the specialization constant and is used to match the constant ID used in the shader module. This is set using the `constant_id` layout qualifier in GLSL and the `SpecId` decoration in SPIR-V. The `offset` and `size` fields are the offset and size of the raw data containing the values of the specicialization constants. The raw data is pointed to by `pData` field of the `VkSpecializationInfo` structure, and its size is given in `dataSize`. Vulkan uses the data in this blob to initialize the specialization constants. If one or more of the specialization constants in the shader is not specified in the specialization info when the pipeline is constructed, its default value is used.

When you are done with a pipeline and no longer need it, you should destroy it in order to free any resources associated with it. To destroy a pipeline object, call **vkDestroyPipeline()**, the prototype of which is

```
void vkDestroyPipeline (
    VkDevice                            device,
    VkPipeline                          pipeline,
    const VkAllocationCallbacks*        pAllocator);
```

The device that owns the pipeline is specified in `device`, and the pipeline to be destroyed is passed in `pipeline`. If a host memory allocator was used to create the pipeline, a pointer to a compatible allocator should be passed in `pAllocator`; otherwise, `pAllocator` should be set to `nullptr`.

After a pipeline has been destroyed, it should not be used again. This includes any references to it from command buffers that may not yet have completed execution. It is the application's responsibility to ensure that any submitted command buffers referencing the pipeline have completed execution and that any command buffers into which the pipeline is bound are not submitted after the pipeline is destroyed.

Accelerating Pipeline Creation

Creation of pipelines is possibly one of the most expensive operations that your application might perform. Although SPIR-V code is consumed by **vkCreateShaderModule()**, it is not until you call **vkCreateGraphicsPipelines()** or **vkCreateComputePipelines()** that Vulkan sees all of the shader stages and other state associated with the pipeline that might affect the final code that will execute on the device. For this reason, a Vulkan implementation may delay the majority of work involved in creating a ready-to-run pipeline object until the last possible moment. This includes shader compilation and code generation, which are typically fairly intensive operations.

Because an application that runs many times will use the same pipelines over and over, Vulkan provides a mechanism to *cache* the results of pipeline creation across runs of an application. This allows applications that build all of their pipelines at startup to start more quickly. The pipeline cache is represented as an object that is created by calling

```
VkResult vkCreatePipelineCache (
    VkDevice                            device,
    const VkPipelineCacheCreateInfo*    pCreateInfo,
    const VkAllocationCallbacks *       pAllocator,
    VkPipelineCache*                    pPipelineCache);
```

The device that will be used to create the pipeline cache is specified in `device`. The remaining parameters for the creation of the pipeline cache are passed through a pointer to an instance of the `VkPipelineCacheCreateInfo` structure, the definition of which is

```
typedef struct VkPipelineCacheCreateInfo {
    VkStructureType             sType;
    const void *                pNext;
    VkPipelineCacheCreateFlags  flags;
```

```
    size_t                              initialDataSize;
    const void *                        pInitialData;
} VkPipelineCacheCreateInfo;
```

The `sType` field of the `VkPipelineCacheCreateInfo` structure should be set to
`VK_STRUCTURE_TYPE_PIPELINE_CACHE_CREATE_INFO`, and `pNext` should be set to
`nullptr`. The `flags` field is reserved for future use and should be set to zero. If existing data is
available from a previous run of the application, its address can be passed through `pInitialData`.
The size of the data is passed in `initialDataSize`. If no initial data is available,
`initialDataSize` should be set to zero, and `pInitialData` should be set to `nullptr`.

When the cache is created, the initial data (if any) is used to prime the cache. If necessary, Vulkan
makes a copy of the data. The data pointed to by `pInitialData` is not modified. As more
pipelines are created, data describing them may be added to the cache, growing it over time. To
retrieve the data from the cache, call **vkGetPipelineCacheData()**. The prototype of
vkGetPipelineCacheData() is

[**Click here to view code image**](#)

```
VkResult vkGetPipelineCacheData (
    VkDevice                            device,
    VkPipelineCache                     pipelineCache,
    size_t*                             pDataSize,
    void*                               pData);
```

The device that owns the pipeline cache should be specified in `device`, and the handle to the
pipeline cache whose data is being queried should be passed in `pipelineCache`. If `pData` is not
`nullptr`, then it points to a region of memory that will receive the cache data. In this case, the
initial value of the variable pointed to by `pDataSize` is the size, in bytes, of this region of memory.
That variable will be overwritten with the amount of data actually written into memory.

If `pData` is `nullptr`, then the initial value of the variable pointed to by `pDataSize` is ignored,
and the variable is overwritten with the size of the data required to store the cache. In order to store
the entire cache data, call **vkGetPipelineCacheData()** twice; the first time, call it with
`pData` set to `nullptr` and `pDataSize` pointing to a variable that will receive the required size of
the cache data. Then size a buffer appropriately to store the resulting cache data and call
vkGetPipelineCacheData() again, this time passing a pointer to this memory region in
`pData`. [Listing 6.7](#) illustrates how to save the pipeline data to a file.

Listing 6.7: Saving Pipeline Cache Data to a File

[**Click here to view code image**](#)

```
VkResult SaveCacheToFile(VkDevice device, VkPipelineCache cache,
        const char* fileName)
{
    size_t cacheDataSize;
    VkResult result = VK_SUCCESS;

    // Determine the size of the cache data.
    result = vkGetPipelineCacheData(device,
                                    cache,
```

```
                                    &cacheDataSize,
                                    nullptr);

    if (result == VK_SUCCESS && cacheDataSize != 0)
    {
        FILE* pOutputFile;
        void* pData;

        // Allocate a temporary store for the cache data.
        result = VK_ERROR_OUT_OF_HOST_MEMORY;
        pData = malloc(cacheDataSize);

        if (pData != nullptr)
        {
            // Retrieve the actual data from the cache.
            result = vkGetPipelineCacheData(device,
                                            cache,
                                            &cacheDataSize,
                                            pData);

        if (result == VK_SUCCESS)
        {
            // Open the file and write the data to it.
            pOutputFile = fopen(fileName, "wb");

            if (pOutputFile != nullptr)
            {
                fwrite(pData, 1, cacheDataSize, pOutputFile);

                fclose(pOutputFile);
            }
                free(pData);
            }
        }
    }

    return result;
}
```

Once you have received the pipeline data, you can store it to disk or otherwise archive it ready for a future run of your application. There is no defined structure to the content of the cache; it is implementation-dependent. However, the first few words of the cache data always form a header that can be used to verify that a blob of data is a valid cache and which device created it.

The layout of the cache header can be represented as the following C structure:

[Click here to view code image](#)

```
// This structure does not exist in official headers but is included here
// for illustration.
typedef struct VkPipelineCacheHeader {
    uint32_t        length;
    uint32_t        version;
    uint32_t        vendorID;
```

```
    uint32_t          deviceID;
    uint8_t           uuid[16];
} VkPipelineCacheHeader;
```

Although the members of the structure are listed as `uint32_t` typed variables, the data in the cache is not formally of type `uint32_t`. Caches are always stored in *little-endian* byte order, regardless of the byte ordering of the host. This means that if you want to interpret this structure on a big-endian host, you need to reverse the byte order of the `uint32_t` fields.

The `length` field is the size of the header structure, in bytes. In the current revision of the specification, this length should be 32. The `version` field is the version of the structure. The only defined version is 1. The `vendorID` and `deviceID` fields should match the `vendorID` and `deviceID` fields of the `VkPhysicalDeviceProperties` structure returned from a call to **`vkGetPhysicalDeviceProperties()`**. The `uuid` field is an opaque string of bytes that uniquely identifies the device. If there is a mismatch between the `vendorID`, `deviceID`, or `uuid` field and what the Vulkan driver expects, it may reject the cache data and reset the cache to empty. A driver may also embed checksum, encryption, or other data inside the cache to ensure that invalid cache data is not loaded into the device.

If you have two cache objects and wish to merge them, you can do that by calling **`vkMergePipelineCaches()`**, the prototype of which is

[Click here to view code image](#)

```
VkResult vkMergePipelineCaches (
    VkDevice                            device,
    VkPipelineCache                     dstCache,
    uint32_t                            srcCacheCount,
    const VkPipelineCache*              pSrcCaches);
```

The `device` parameter is a handle to the device that owns the caches that are to be merged. `dstCache` is a handle to the destination cache, which will end up as an amalgamation of all of the entries from all source caches. The number of caches that are to be merged is specified in `srcCacheCount`, and `pSrcCaches` is a pointer to an array of `VkPipelineCache` handles to the caches to be merged.

After **`vkMergePipelineCaches()`** has executed, `dstCache` will contain all of the cache entries from all of the source caches specified in `pSrcCaches`. It is then possible to call **`vkGetPipelineCacheData()`** on the destination cache to retrieve a single, large cache data structure representing all entries from all of the caches.

This is particularly useful, for example, when creating pipelines in many threads. Although access to the pipeline cache is thread-safe, implementations may internally take locks to prevent concurrent write access to multiple caches. If you instead create multiple pipeline caches—one for each thread—and use them during initial creation of your pipelines, any per-cache locks taken by the implementation will be uncontested, speeding access to them. Later, when all pipelines are created, you can merge the pipelines in order to save their data in one large resource.

When you are done creating pipelines and no longer need the cache, it's important to destroy it, because it could be quite large. To destroy a pipeline cache object, call **`vkDestroyPipelineCache()`**, the prototype of which is

```
void vkDestroyPipelineCache (
    VkDevice                              device,
    VkPipelineCache                       pipelineCache,
    const VkAllocationCallbacks*          pAllocator);
```

`device` is a handle to the device that owns the pipeline cache, and `pipelineCache` is a handle to the pipeline cache object that is to be destroyed. After the pipeline cache has been destroyed, it should not be used again, although pipelines created using the cache are still valid. Also, any data retrieved from the cache using a call to **vkGetPipelineCacheData()** is valid and can be used to construct a new cache that should result in matches on subsequent pipeline creation requests.

Binding Pipelines

Before you can use a pipeline, it must be *bound* into the a command buffer that will execute drawing or dispatching commands. When such a command is executed, the current pipeline (and all the shaders in it) are used to process the commands. To bind a pipeline to a command buffer, call **vkCmdBindPipeline()**, the prototype of which is

```
void vkCmdBindPipeline (
    VkCommandBuffer                       commandBuffer,
    VkPipelineBindPoint                   pipelineBindPoint,
    VkPipeline                            pipeline);
```

The command buffer to which you are binding the pipeline is specified in `commandBuffer`, and the pipeline you are binding is specified in `pipeline`. There are two binding points for pipelines on each command buffer: the graphics and compute binding points. The compute bind point is where compute pipelines should be bound. Graphics pipelines are covered in the next chapter and should be bound to the graphics pipeline bind point.

To bind a pipeline to the compute binding point, set `pipelineBindPoint` to `VK_PIPELINE_BIND_POINT_COMPUTE`, and to bind the pipeline to the graphics binding point, set `pipelineBindPoint` to `VK_PIPELINE_BIND_POINT_GRAPHICS`.

The current pipeline binding for each of compute and graphics is part of the state of a command buffer. When a new command buffer is begun, this state is undefined. Therefore, you *must* bind a pipeline to the relevant binding point before invoking any work that would use a pipeline.

Executing Work

In the previous section, you saw how to construct a compute pipeline using **vkCreateComputePipelines()** and bind it into a command buffer. Once a pipeline is bound, you can use it to execute work.

Compute shaders running as part of a compute pipeline execute in groups called local work groups. These groups logically execute in lockstep and are of a fixed size specified in the shader. The maximum size of a local work group is typically small but must be at least 128 invocations × 128 invocations × 64 invocations. Further, the maximum number of total invocations in a single local work group may also be smaller than this total volume and is required only to be 128 invocations.

For this reason, local work groups are started in larger groups, sometimes called the global work group or dispatch size. Kicking off work in a compute shader is therefore called *dispatching* work, or a *dispatch*. The local work group is logically a 3D construct, or *volume* of invocations, although one or two of the dimensions can be a single invocation in size, making the work group flat in that direction. Likewise, these local work groups are dispatched together in three dimensions, even if one or more of those dimensions is a single work group deep.

The command to dispatch a global work group using a compute pipeline is **vkCmdDispatch()**, the prototype of which is

```
void vkCmdDispatch (
    VkCommandBuffer                          commandBuffer,
    uint32_t                                 x,
    uint32_t                                 y,
    uint32_t                                 z);
```

The command buffer that will execute the command is passed in `commandBuffer`. The number of local work groups in each of the *x*, *y*, and *z* dimensions is passed in the `x`, `y`, and `z` parameters, respectively. A valid compute pipeline must be bound to the command buffer at the `VK_PIPELINE_BIND_POINT_COMPUTE` binding point. When the command is executed by the device, a global work group of $x \times y \times z$ local work groups begins executing the shader contained in the bound pipeline.

It is perfectly possible to have a local work group that has a different effective dimensionality from that of the global work group. For example, it's fine to have a $32 \times 32 \times 1$ dispatch of $64 \times 1 \times 1$ local work groups.

In addition to being able to specify the number of work groups in the dispatch using parameters to **vkCmdDispatch()**, it's possible to perform an *indirect* dispatch, where the size of the dispatch in work groups is sourced from a buffer object. This allows dispatch sizes to be computed after a command buffer is built by performing an indirect dispatch using a buffer and then overwriting the contents of the buffer using the host. The content of the buffer can even be updated using the device itself to provide a limited means for the device to feed itself work.

The prototype of **vkCmdDispatchIndirect()** is

```
void vkCmdDispatchIndirect (
    VkCommandBuffer                          commandBuffer,
    VkBuffer                                 buffer,
    VkDeviceSize                             offset);
```

Again, the command buffer that will contain the command is passed in `commandBuffer`. Rather than passing the dispatch size as in **vkCmdDispatch()**, the number of workgroups in each dimension are expected to be stored as three consecutive `uint32_t` variables at the offset specified in `offset` (in bytes) in the buffer object specified in `buffer`. The parameters in the buffer essentially represent an instance of the `VkDispatchIndirectCommand` structure, the definition of which is

```
typedef struct VkDispatchIndirectCommand {
    uint32_t   x;
    uint32_t   y;
    uint32_t   z;
} VkDispatchIndirectCommand;
```

Again, the contents of the buffer are not read until the **vkCmdDispatchIndirect()** command is reached during processing of the command buffer by the device.

The maximum number of work groups in each dimension supported by a device can be determined by inspecting the maxComputeWorkGroupCount field of the device's VkPhysicalDeviceLimits structure returned from a call to **vkGetPhysicalDeviceProperties()**, as explained in Chapter 1, "Overview of Vulkan." Exceeding those limits in a call to **vkCmdDispatch()** or placing values outside those limits in the buffer referenced by **vkCmdDispatchIndirect()** will result in undefined (probably bad) behavior.

Resource Access in Shaders

The shaders in your program consume and produce data in one of two ways. The first is through interaction with fixed function hardware, and the second is by directly reading and writing resources. You saw in Chapter 2, "Memory and Resources," how to create buffers and images. In this section, we introduce *descriptor sets*, which are representations of the set of resources that shaders can interact with.

Descriptor Sets

A descriptor set is a set of resources that are bound into the pipeline as a group. Multiple sets can be bound to a pipeline at a time. Each set has a layout, which describes the order and types of resources in the set. Two sets with the same layout are considered to be compatible and interchangeable. The descriptor set layout is represented by an object, and sets are created with respect to this object. Further, the set of sets that are accessible to a pipeline are grouped into another object: the pipeline layout. Pipelines are created with respect to this pipeline layout object.

The relationship between the descriptor set layout and the pipeline layout is illustrated in Figure 6.1. As you can see in the figure, two descriptor sets are defined, the first having a texture, a sampler, and two buffers. The second set contains four textures, two samplers, and three buffers. These descriptor sets' layouts are aggregated into a single pipeline layout. A pipeline can then be created with respect to the pipeline layout, while descriptor sets are created with respect to descriptor set layouts. Those descriptor sets can be bound into command buffers along with compatible pipelines to allow those pipelines to access the resources in them.
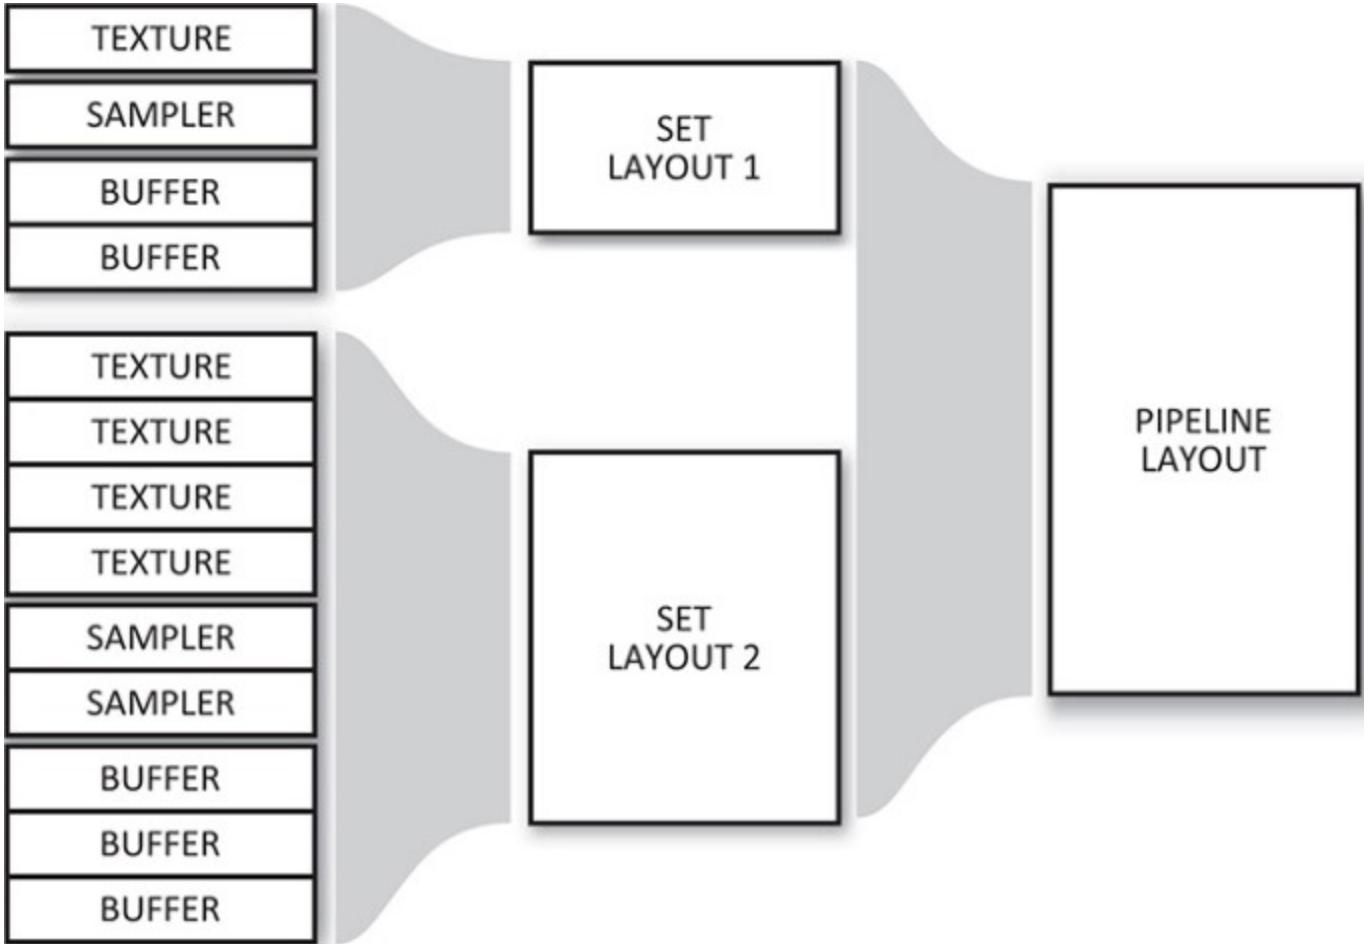
Figure 6.1: Descriptor Sets and Pipeline Sets

At any time, the application can bind a new descriptor set to the command buffer in any point that has an identical layout. The same descriptor set layouts can be used to create multiple pipelines. Therefore, if you have a set of objects that share a common set of resources, but additionally each require some unique resources, you can leave the common set bound and replace the unique resources as your application moves through the objects that need to be rendered.

To create a descriptor set layout object, call **vkCreateDescriptorSetLayout()**, the prototype of which is

Click here to view code image

```
VkResult vkCreateDescriptorSetLayout (
    VkDevice                                 device,
    const VkDescriptorSetLayoutCreateInfo*   pCreateInfo,
    const VkAllocationCallbacks*             pAllocator,
    VkDescriptorSetLayout*                   pSetLayout);
```

As usual, the information required to construct the descriptor set layout object is passed through a pointer to a structure. This is an instance of the VkDescriptorSetLayoutCreateInfo structure, the definition of which is

```
typedef struct VkDescriptorSetLayoutCreateInfo {
    VkStructureType                      sType;
    const void*                          pNext;
    VkDescriptorSetLayoutCreateFlags     flags;
    uint32_t                             bindingCount;
    const VkDescriptorSetLayoutBinding*  pBindings;
} VkDescriptorSetLayoutCreateInfo;
```

The `sType` field of `VkDescriptorSetLayoutCreateInfo` should be set to
`VK_STRUCTURE_TYPE_DESCRIPTOR_SET_LAYOUT_CREATE_INFO`, and `pNext` should be
set to `nullptr`. `flags` is reserved for future use and should be set to zero.

Resources are bound to binding points in the descriptor set. The `bindingCount` and `pBindings`
members of `VkDescriptorSetLayoutCreateInfo` contain the number of binding points that
the set will contain and a pointer to an array containing their descriptions, respectively. Each binding
is described by an instance of the `VkDescriptorSetLayoutBinding` structure, the definition
of which is

```
typedef struct VkDescriptorSetLayoutBinding {
    uint32_t             binding;
    VkDescriptorType     descriptorType;
    uint32_t             descriptorCount;
    VkShaderStageFlags   stageFlags;
    const VkSampler*     pImmutableSamplers;
} VkDescriptorSetLayoutBinding;
```

Each resource accessible to a shader is given a binding number. This binding number is stored in the
`binding` field of `VkDescriptorSetLayoutBinding`. The bindings used in a descriptor set
do not need to be contiguous, and there can be gaps (unused binding numbers) in a set. However, it's
recommended that you don't create sparsely populated sets because this can waste resources in the
device.

The type of descriptor at this binding point is stored in `descriptorType`. This a member of the
`VkDescriptorType` enumeration. We'll discuss the various resource types a little later, but they
include

- `VK_DESCRIPTOR_TYPE_SAMPLER`: A sampler is an object that can be used to perform
  operations such as filtering and sample coordinate transformations when reading data from an
  image.
- `VK_DESCRIPTOR_TYPE_SAMPLED_IMAGE`: A sampled image is an image that can be used
  in conjunction with a sampler to provide filtered data to a shader.
- `VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER`: A combined image-sampler
  object is a sampler and an image paired together. The same sampler is always used to sample
  from the image, which can be more efficient on some architectures.
- `VK_DESCRIPTOR_TYPE_STORAGE_IMAGE`: A storage image is an image that cannot be
  used with a sampler but can be written to. This is in contrast to a sampled image, which cannot
  be written to.

- `VK_DESCRIPTOR_TYPE_UNIFORM_TEXEL_BUFFER`: A uniform texel buffer is a buffer that is filled with homogeneous formatted data that cannot be written by shaders. Knowing that buffer content is constant may allow some Vulkan implementations to optimize access to the buffer better.

- `VK_DESCRIPTOR_TYPE_STORAGE_TEXEL_BUFFER`: A storage texel buffer is a buffer that contains formatted data much like a uniform texel buffer, but a storage buffer can be written to.

- `VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER` and `VK_DESCRIPTOR_TYPE_STORAGE_BUFFER`: These are similar to `VK_DESCRIPTOR_TYPE_UNIFORM_TEXEL_BUFFER` and `VK_DESCRIPTOR_TYPE_STORAGE_TEXEL_BUFFER`, except that the data is unformatted and described by structures declared in the shader.

- `VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER_DYNAMIC` and `VK_DESCRIPTOR_TYPE_STORAGE_BUFFER_DYNAMIC`: These are similar to `VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER` and `VK_DESCRIPTOR_TYPE_STORAGE_BUFFER`, but include an offset and size that are passed when the descriptor set is bound to the pipeline rather than when the descriptor is bound into the set. This allows a single buffer in a single set to be updated at high frequency.

- `VK_DESCRIPTOR_TYPE_INPUT_ATTACHMENT`: An input attachment is a special type of image whose content is generated by earlier operations on the same image in a graphics pipeline.

Listing 6.8 illustrates how a selection of resources is declared inside a GLSL shader.

Listing 6.8 : Declaring Resources in GLSL

Click here to view code image

```
#version 450 core

layout (set = 0, binding = 0) uniform sampler2DmyTexture;
layout (set = 0, binding = 2) uniform sampler3DmyLut;
layout (set = 1, binding = 0) uniform myTransforms
{
    mat4 transform1;
    mat3 transform2;
};

void main(void)
{
    // Do nothing!
}
```

Listing 6.9 shows a condensed form of what the shader in Listing 6.8 translates to when compiled with the GLSL compiler.

Listing 6.9: Declaring Resources in SPIR-V

Click here to view code image

```
; SPIR-V
; Version: 1.0
; Generator: Khronos Glslang Reference Front End; 1
; Bound: 22
; Schema: 0
                OpCapability Shader
           %1 = OpExtInstImport "GLSL.std.450"
                OpMemoryModel Logical GLSL450
                OpEntryPoint GLCompute %4 "main"
                OpExecutionMode %4 LocalSize 1 1 1
                OpSource GLSL 450
                OpName %4 "main"
                OpName %10 "myTexture"
                OpName %14 "myLut"
                OpName %19 "myTransforms"
                OpMemberName %19 0 "transform1"
                OpMemberName %19 1 "transform2"
                OpName %21 ""
                OpDecorate %10 DescriptorSet 0
                OpDecorate %10 Binding 0
                OpDecorate %14 DescriptorSet 0
                OpDecorate %14 Binding 2
                OpMemberDecorate %19 0 ColMajor
                OpMemberDecorate %19 0 Offset 0
                OpMemberDecorate %19 0 MatrixStride 16
                OpMemberDecorate %19 1 ColMajor
                OpMemberDecorate %19 1 Offset 64
                OpMemberDecorate %19 1 MatrixStride 16
                OpDecorate %19 Block
                OpDecorate %21 DescriptorSet 1
                OpDecorate %21 Binding 0
           %2 = OpTypeVoid
           %3 = OpTypeFunction %2
           %6 = OpTypeFloat 32
           %7 = OpTypeImage %6 2D 0 0 0 1 Unknown
           %8 = OpTypeSampledImage %7
           %9 = OpTypePointer UniformConstant %8
          %10 = OpVariable %9 UniformConstant
          %11 = OpTypeImage %6 3D 0 0 0 1 Unknown
          %12 = OpTypeSampledImage %11
          %13 = OpTypePointer UniformConstant %12
          %14 = OpVariable %13 UniformConstant
          %15 = OpTypeVector %6 4
          %16 = OpTypeMatrix %15 4
          %17 = OpTypeVector %6 3
          %18 = OpTypeMatrix %17 3
          %19 = OpTypeStruct %16 %18
          %20 = OpTypePointer Uniform %19
          %21 = OpVariable %20 Uniform
```

```
        %4 = OpFunction %2 None %3
        %5 = OpLabel
             OpReturn
             OpFunctionEnd
```

Multiple descriptor set layouts can be used in a single pipeline. As you can see in Listings 6.8 and 6.9, the resources are placed in two sets, the first containing "myTexture" and "myLut" (both samplers) and the second containing "myTransforms" (a uniform buffer). To group two or more descriptor sets into something the pipeline can use, we need to aggregate them into a VkPipelineLayout object. To do this, call **vkCreatePipelineLayout()**, the prototype of which is

Click here to view code image

```
VkResult vkCreatePipelineLayout (
    VkDevice                            device,
    const VkPipelineLayoutCreateInfo*   pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkPipelineLayout*                   pPipelineLayout);
```

This function uses the device specified in device to create a new VkPipelineLayout object using the information in a VkPipelineLayoutCreateInfo structure passed by address in pCreateInfo. The definition of VkPipelineLayoutCreateInfo is

Click here to view code image

```
typedef struct VkPipelineLayoutCreateInfo {
    VkStructureType                 sType;
    const void*                     pNext;
    VkPipelineLayoutCreateFlags     flags;
    uint32_t                        setLayoutCount;
    const VkDescriptorSetLayout*    pSetLayouts;
    uint32_t                        pushConstantRangeCount;
    const VkPushConstantRange*      pPushConstantRanges;
} VkPipelineLayoutCreateInfo;
```

The sType field for VkPipelineLayoutCreateInfo should be set to VK_STRUCTURE_TYPE_PIPELINE_LAYOUT_CREATE_INFO, and pNext should be set to nullptr. The flags field is reserved in the current version of Vulkan and should be set to zero.

The number of descriptor set layouts (which is the same as the number of sets in the pipeline layout) is given in setLayoutCount, and pSetLayouts is a pointer to an array of VkDescriptorSetLayout handles created previously with calls to **vkCreateDescriptorSetLayout()**. The maximum number of descriptor sets that can be bound at once (and therefore the maximum number of set layouts in a pipeline layout) is at least 4. Some implementations may support a higher limit than this. You can determine the absolute maximum number of layouts supported by inspecting the maxBoundDescriptorSets member of the device's VkPhysicalDeviceLimits structure, which you can retrieve by calling **vkGetPhysicalDeviceProperties()**.

The final two parameters, pushConstantRangeCount and pPushConstantRanges, are used to describe the *push constants* used in the pipeline. Push constants are a special class of resource

that can be used directly as constants in a shader. It is extremely fast to update the values of push constants, requiring no synchronization. We discuss push constants later in this chapter.

When the `VkDescriptorSetLayout` object is created, the resources used by all of the sets within the pipeline layout are aggregated and must fall within a device-dependent limit. Effectively, there is an upper bound on the number and type of resources that can be accessed by a single pipeline.

Further, some devices may not support accessing all of the pipeline's resources from every shader stage simultaneously and therefore have a *per-stage* upper limit on the number of resources accessible from each stage.

Each of the limits can be checked by retrieving the device's `VkPhysicalDeviceLimits` structure through a call to **vkGetPhysicalDeviceProperties()** and checking the relevant members. The members of `VkPhysicalDeviceLimits` associated with pipeline layout maximums are shown in Table 6.1.

Field Limit	Guaranteed Minimum
maxDescriptorSetSamplers Maximum number of samplers in a single pipeline	96
maxDescriptorSetUniformBuffers Maximum number of uniform buffers in a single pipeline	72
maxDescriptorSetUniformBuffersDynamic Maximum number of dynamic uniform buffers in a single pipeline	8
maxDescriptorSetStorageBuffers Maximum number of shader storage buffers in a single pipeline	24
maxDescriptorSetStorageBuffersDynamic Maximum number of dynamic shader storage buffers in a single pipeline	4
maxDescriptorSetSampledImages Maximum number of sampled images in a single pipeline	96
maxDescriptorSetStorageImages Maximum number of storage images in a single pipeline	24
maxDescriptorSetInputAttachments Maximum number of input attachments in a single pipeline	4
maxPerStageDescriptorSamplers Maximum number of samplers in a single stage	16
maxPerStageDescriptorUniformBuffers Maximum number of uniform buffers in a single stage	12
maxPerStageDescriptorStorageBuffers Maximum number of shader storage buffers in a single stage	4
maxPerStageDescriptorSampledImages Maximum number of sampled images in a single stage	16
maxPerStageDescriptorStorageImages Maximum number of storage images in a single stage	4
maxPerStageDescriptorInputAttachments Maximum number of input attachments in a single stage	4
maxPerStageResources Maximum total resources used in a single stage	128

Table 6.1: Pipeline Resource Limits

If your shaders or the resulting pipeline need to use more resources than are guaranteed to be supported as shown in Table 6.1, then you need to check the resource limits and be prepared to fail gracefully if you exceed them. However, if your resource requirements fit comfortably inside this range, there's no reason to query any of it directly as Vulkan guarantees at least this level of support.

Two pipelines can be used with the same set of descriptor sets if their pipeline layouts are considered to be compatible. For two pipeline layouts to be compatible for descriptor sets, they must

- Use the same number of push constant ranges
- Use the same descriptor set layouts (or identical layouts) in the same order

Two pipeline layouts are also considered to be partially compatible if they use the same (or identically defined) set layouts for the first few sets and then differ after that. In that case, the pipelines are compatible up to the point where the descriptor set layouts change.

When a pipeline is bound to a command buffer, it can continue to use any bound descriptor sets that are compatible with the set bindings in the pipeline layout. Thus, switching between two (partially) compatible pipelines doesn't require re-binding any sets up to the point where the pipelines share layouts. If you have a set of resources that you want to be globally available, such as a uniform block containing some per-frame constants, or a set of textures that you want to be available to every shader, put those in the first set(s). Resources that might change at higher frequency can be placed in higher-numbered sets.

Listing 6.10 shows the application-side code to create the descriptor set layouts and pipeline layout describing the sets referenced in Listings 6.8 and 6.9.

Listing 6.10: Creating a Pipeline Layout

Click here to view code image

```
// This describes our combined image-samplers. One set, two disjoint
bindings.
static const VkDescriptorSetLayoutBinding Samplers[] =
{
    {
        0,                                            // Start from binding
0
        VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER,    // Combined image-
sampler
        1,                                            // Create one binding
        VK_SHADER_STAGE_ALL,                          // Usable in all
stages
        nullptr                                       // No static samplers
    },
    {
        2,                                            // Start from binding
2
        VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER,    // Combined image-
sampler
        1,                                            // Create one binding
        VK_SHADER_STAGE_ALL,                          // Usable in all
stages
```

```cpp
        nullptr                                           // No static samplers
    }
};
// This is our uniform block. One set, one binding.
static const VkDescriptorSetLayoutBinding UniformBlock =
{
    0,                                                    // Start from binding
0
    VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER,          // Uniform block
    1,                                          // One binding
    VK_SHADER_STAGE_ALL,                        // All stages
    nullptr                                     // No static samplers
};


// Now create the two descriptor set layouts.
static const VkDescriptorSetLayoutCreateInfo createInfoSamplers =
{
    VK_STRUCTURE_TYPE_DESCRIPTOR_SET_LAYOUT_CREATE_INFO,
    nullptr,
    0,
    2,
    &Samplers[0]
};
static const VkDescriptorSetLayoutCreateInfo createInfoUniforms =
{
    VK_STRUCTURE_TYPE_DESCRIPTOR_SET_LAYOUT_CREATE_INFO,
    nullptr,
    0,
    1,
    &UniformBlock
};
// This array holds the two set layouts.
VkDescriptorSetLayout setLayouts[2];

vkCreateDescriptorSetLayout(device, &createInfoSamplers,
                            nullptr, &setLaouts[0]);
vkCreateDescriptorSetLayout(device, &createInfoUniforms,
                            nullptr, &setLayouts[1]);

// Now create the pipeline layout.
const VkPipelineLayoutCreateInfo pipelineLayoutCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_LAYOUT_CREATE_INFO, nullptr,
    0,
    2, setLayouts,
    0, nullptr
};

VkPipelineLayout pipelineLayout;

vkCreatePipelineLayout(device, &pipelineLayoutCreateInfo,
                       nullptr, pipelineLayout);
```

The pipeline layout we create in [Listing 6.10](#) matches the layout expected by the shader code in [Listings 6.9](#) and [6.10](#). When we create a compute pipeline using that shader, we pass the pipeline layout object created in [Listing 6.10](#) as the `layout` field of the `VkComputePipelineCreateInfo` structure passed to **vkCreateComputePipelines()**.

When the pipeline layout is no longer needed, it should be destroyed by calling **vkDestroyPipelineLayout()**. This frees any resources associated with the pipeline layout object. The prototype of **vkDestroyPipelineLayout()** is

[**Click here to view code image**](#)

```
void vkDestroyPipelineLayout (
    VkDevice                        device,
    VkPipelineLayout                pipelineLayout,
    const VkAllocationCallbacks*    pAllocator);
```

After the pipeline layout object is destroyed, it should not be used again. However, any pipelines created using the pipeline layout object remain valid until they are destroyed. It is therefore not necessary to keep pipeline layout objects around in order to use the pipelines created using them.

To destroy a descriptor set layout object and free its resources, call

[**Click here to view code image**](#)

```
void vkDestroyDescriptorSetLayout (
    VkDevice                        device,
    VkDescriptorSetLayout           descriptorSetLayout,
    const VkAllocationCallbacks*    pAllocator);
```

The device that owns the descriptor set layout should be passed in `device`, and the handle to the descriptor set layout is passed in `descriptorSetLayout`. `pAllocator` should point to a host memory allocation structure that is compatible with the one used to create the descriptor set layout or should be `nullptr` if the `pAllocator` parameter to **vkCreateDescriptorSetLayout()** was also `nullptr`.

After a descriptor set layout is destroyed, its handle is no longer valid, and it should not be used again. However, descriptor sets, pipeline layouts, and other objects created by referencing the set remain valid.

Binding Resources to Descriptor Sets

Resources are represented by *descriptors* and are bound to the pipeline by first binding their descriptors into sets and then binding those descriptor sets to the pipeline. This allows a large number of resources to be bound with very little processing time because the exact set of resources used by a particular drawing command can be determined in advance and the descriptor set holding them created up front.

The descriptors are allocated from pools called *descriptor pools*. Because descriptors for different types of resources are likely to have similar data structures on any given implementation, pooling the allocations used to store descriptors allows drivers to make efficient use of memory. To create a descriptor pool, call **vkCreateDescriptorPool()**, the prototype of which is

```
VkResult vkCreateDescriptorPool (
    VkDevice                              device,
    const VkDescriptorPoolCreateInfo*     pCreateInfo,
    const VkAllocationCallbacks*          pAllocator,
    VkDescriptorPool*                     pDescriptorPool);
```

The device that is to be used to create the descriptor pool is specified in `device`, and the remaining parameters describing the new pool are passed through a pointer to an instance of the `VkDescriptorPoolCreateInfo` structure in `pCreateInfo`. The definition of `VkDescriptorPoolCreateInfo` is

```
typedef struct VkDescriptorPoolCreateInfo {
    VkStructureType             sType;
    const void*                 pNext;
    VkDescriptorPoolCreateFlags flags;
    uint32_t                    maxSets;
    uint32_t                    poolSizeCount;
    const VkDescriptorPoolSize* pPoolSizes;
} VkDescriptorPoolCreateInfo;
```

The `sType` field of `VkDescriptorPoolCreateInfo` should be set to `VK_STRUCTURE_TYPE_DESCRIPTOR_POOL_CREATE_INFO`, and the `pNext` field should be set to `nullptr`. The `flags` field is used to pass additional information about the allocation strategy that should be used to manage the resources consumed by the pool. The only defined flag is `VK_DESCRIPTOR_POOL_CREATE_FREE_DESCRIPTOR_SET_BIT`, which, if set, indicates that the application may free individual descriptors allocated from the pool, so the allocator should be prepared for that. If you don't intend to return individual descriptors to the pool, simply set `flags` to zero.

The `maxSets` field specifies the maximum total number of sets that may be allocated from the pool. Note that this is the total number of sets, regardless of the size of each set or the overall size of the pool. The next two fields, `poolSizeCount` and `pPoolSize`, specify the number of resource descriptors for each type of resoruce that might be stored in the set. `pPoolSize` is a pointer to an array of `poolSizeCount` instances of the `VkDescriptorPoolSize` structure, each one specifying the number of descriptors of a particular type that may be allocated from the pool. The definition of `VkDescriptorPoolSize` is

```
typedef struct VkDescriptorPoolSize {
    VkDescriptorType    type;
    uint32_t            descriptorCount;
} VkDescriptorPoolSize;
```

The first field of `VkDescriptorPoolSize`, `type`, specifies the type of resource, and the second field, `descriptorCount`, specifies the number of descriptors of that type to be stored in the pool. `type` is a member of the `VkDescriptorType` enumeration. If no element of the `pPoolSize` array specifies a particular type of resource, then no descriptors of that type can be allocated from the

resulting pool. If a particular type of resource appears twice in the array, then the sum of all of their `descriptorCount` fields is used to size the pool for that type of resource. The total number of resources in the pool is divided among the sets allocated from the pool.

If creation of the pool is successful, then a handle to the new `VkDescriptorPool` object is written into the variable pointed to by `pDescriptorPool`. To allocate blocks of descriptors from the pool, we create new descriptor set objects by calling **vkAllocateDescriptorSets()**, the declaration of which is

[**Click here to view code image**](#)

```
VkResult vkAllocateDescriptorSets (
    VkDevice                             device,
    const VkDescriptorSetAllocateInfo*   pAllocateInfo,
    VkDescriptorSet*                     pDescriptorSets);
```

The device that owns the descriptor pool from which the sets are to be allocated is passed in `device`. The remaining information describing the sets to be allocated is passed via a pointer to an instance of the `VkDescriptorSetAllocateInfo` structure in `pDescriptorSets`. The definition of `VkDescriptorSetAllocateInfo` is

[**Click here to view code image**](#)

```
typedef struct VkDescriptorSetAllocateInfo {
    VkStructureType                 sType;
    const void*                     pNext;
    VkDescriptorPool                descriptorPool;
    uint32_t                        descriptorSetCount;
    const VkDescriptorSetLayout*    pSetLayouts;
} VkDescriptorSetAllocateInfo;
```

The `sType` field of the `VkDescriptorSetAllocateInfo` structure should be set to `VK_STRUCTURE_TYPE_DESCRIPTOR_SET_ALLOCATE_INFO`, and the `pNext` field should be set to `nullptr`. A handle to the descriptor pool from which to allocate the sets is specified in `descriptorPool`, which should be a handle of a descriptor set created by a call to **vkCreateDescriptorPool()**. Access to `descriptorPool` should be externally synchronized. The number of sets to create is specified in `descriptorSetCount`. The layout of each set is then passed through an array of `VkDescriptorSetLayout` object handles in `pSetLayouts`.

When successful, **vkAllocateDescriptorSets()** consumes sets and descriptors from the specified pool and deposits the new descriptor set handles in the array pointed to by `pDescriptorSets`. The number of descriptors consumed from the pool for each descriptor set is determined from the descriptor set layouts passed through `pSetLayouts`, the creation of which we described earlier.

If the descriptor pool was created with the `flags` member of the `VkDescriptorSetCreateInfo` structure containing `VK_DESCRIPTOR_POOL_CREATE_FREE_DESCRIPTOR_SET_BIT`, then descriptor sets may be returned to the pool by freeing them. To free one or more descriptor sets, call **vkFreeDescriptorSets()**, the prototype of which is

```
VkResult vkFreeDescriptorSets (
    VkDevice                            device,
    VkDescriptorPool                    descriptorPool,
    uint32_t                            descriptorSetCount,
    const VkDescriptorSet*              pDescriptorSets);
```

The device that owns the descriptor pool is specified in `device`, and the pool to which the descriptor sets should be returned is specified in `descriptorPool`. Access to `descriptorPool` must be externally synchronized. The number of descriptor sets to free is passed in `descriptorSetCount`, and `pDescriptorSets` points to an array of `VkDescriptorSet` handles to the objects to free. When the descriptor sets are freed, their resources are returned to the pool from which they came and may be allocated to a new set in the future.

Even if `VK_DESCRIPTOR_POOL_CREATE_FREE_DESCRIPTOR_SET_BIT` was not specified when a descriptor pool was created, it's still possible to recycle all the resources from all sets allocated from the pool. This is accomplished by resetting the pool itself by calling **vkResetDescriptorPool()**. With this command, it's not necessary to explicitly specify every set allocated from the pool. The prototype of **vkResetDescriptorPool()** is

```
VkResult vkResetDescriptorPool (
    VkDevice                            device,
    VkDescriptorPool                    descriptorPool,
    VkDescriptorPoolResetFlags          flags);
```

`device` is a handle to the device that owns the descriptor pool, and `descriptorPool` is a handle to the descriptor pool being reset. Access to the descriptor pool must be externally synchronized. `flags` is reserved for future use and should be set to zero.

Regardless of whether sets are individually freed by calling **vkFreeDescriptorSets()** or freed in bulk by calling **vkResetDescriptorPool()**, care must be taken to ensure that sets are not referenced after they have been freed. In particular, any command buffer containing commands that might reference descriptor sets that are to be freed should either have completed execution or should be discarded without submission.

To completely free the resources associated with a descriptor pool, you should destroy the pool object by calling **vkDestroyDescriptorPool()**, the prototype of which is

```
void vkDestroyDescriptorPool(
    VkDevice                        device,
    VkDescriptorPool                descriptorPool,
    const VkAllocationCallbacks*    pAllocator);
```

A handle to the device that owns the pool should be passed in `device`, and the handle to the pool to destroy is passed in `descriptorPool`. `pAllocator` should point to a host memory allocation structure that is compatible with the one used to create the pool or should be `nullptr` if the `pAllocator` parameter to **vkCreateDescriptorPool()** was also `nullptr`.

When the descriptor pool is destroyed, all of its resources are freed, including any sets allocated from it. There is no need to explicitly free the descriptor sets allocated from the pool before destroying it or to reset the pool with a call to **vkResetDescriptorPool()**. However, just as when descriptor sets are freed explicitly, you must make sure that your application does not access sets allocated from a pool after that pool has been destroyed. This includes any work performed by the device during execution of command buffers submitted but not yet completed.

To bind resources into descriptor sets, we can either write to the descriptor set directly or copy bindings from another descriptor set. In either case, we use the **vkUpdateDescriptorSets()** command, the prototype of which is

[**Click here to view code image**](#)

```
void vkUpdateDescriptorSets (
    VkDevice                        device,
    uint32_t                        descriptorWriteCount,
    const VkWriteDescriptorSet*     pDescriptorWrites,
    uint32_t                        descriptorCopyCount,
    const VkCopyDescriptorSet*      pDescriptorCopies);
```

The device that owns the descriptor set to be updated is passed in `device`. The number of direct writes is passed in `descriptorWriteCount`, and the number of descriptor copies is passed in `descriptorCopyCount`. The parameters for each write are contained in a `VkWriteDescriptorSet` structure, and the parameters for each copy are contained in a `VkCopyDescriptorSet` structure. The `pDescriptorWrites` and `pDescriptorCopies` parameters contain pointers to `descriptorWriteCount` and `descriptorCopyCount` `VkWriteDescriptorSet` and `VkCopyDescriptorSet` structures, respectively. The definition of `VkWriteDescriptorSet` is

[**Click here to view code image**](#)

```
typedef struct VkWriteDescriptorSet {
    VkStructureType                 sType;
    const void*                     pNext;
    VkDescriptorSet                 dstSet;
    uint32_t                        dstBinding;
    uint32_t                        dstArrayElement;
    uint32_t                        descriptorCount;
    VkDescriptorType                descriptorType;
    const VkDescriptorImageInfo*    pImageInfo;
    const VkDescriptorBufferInfo*   pBufferInfo;
    const VkBufferView*             pTexelBufferView;
} VkWriteDescriptorSet;
```

The `sType` field of `VkWriteDescriptorSet` should be set to `VK_STRUCTURE_TYPE_WRITE_DESCRIPTOR_SET`, and `pNext` should be set to `nullptr`. For each write operation, the destination descriptor set is specified in `dstSet`, and the binding index is specified in `dstBinding`. If the binding in the set refers to an array of resources, then `dstArrayElement` is used to specify the starting index of the update, and `descriptorCount` is used to specify the number of consecutive descriptors to update. If the target binding is not an array, then `dstArrayElement` should be set to 0, and `descriptorCount` should be set to 1.

The type of resource being updated is specified in `descriptorType`, which is a member of the `VkDescriptorType` enumeration. The value of this parameter determines which of the next parameters is considered by the function. If the descriptor being updated is an image resource, then `pImageInfo` is a pointer to an instance of the `VkDescriptorImageInfo` structure containing information about the image. The definition of `VkDescriptorImageInfo` is

Click here to view code image

```
typedef struct VkDescriptorImageInfo {
    VkSampler        sampler;
    VkImageView      imageView;
    VkImageLayout    imageLayout;
} VkDescriptorImageInfo;
```

A handle to the image view that is to be bound into the descriptor set is passed in `imageView`. If the resource in the descriptor set is a `VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER`, then a handle to the accompanying sampler is specified in `sampler`. The layout that the image is expected to be in when it is used in the descriptor set is passed in `imageLayout`.

If the resource to be bound into the descriptor set is a buffer, then the parameters describing the binding are stored in an instance of the `VkDescriptorBufferInfo` structure, and a pointer to this structure is stored in the `pBufferInfo` of `VkWriteDescriptorSet`. The definition of `VkDescriptorBufferInfo` is

Click here to view code image

```
typedef struct VkDescriptorBufferInfo {
    VkBuffer        buffer;
    VkDeviceSize    offset;
    VkDeviceSize    range;
} VkDescriptorBufferInfo;
```

The buffer object to bind is specified in `buffer`, and the offset and size of binding, expressed in bytes, are specified in `offset` and `range`, respectively. The bound range must lie entirely inside the buffer object. To bind the whole buffer (inferring the size of the range from the buffer object), then `range` can be set to `VK_WHOLE_SIZE`.

If the referenced buffer binding is a uniform buffer binding, then `range` must be less than or equal to the device's `maxUniformBufferRange` limit, as determined by calling **vkGetPhysicalDeviceProperties()** and inspecting the `VkPhysicalDeviceLimits` structure. Also, the `offset` parameter must be an integer multiple of the device's uniform buffer offset alignment requirement, which is contained in the `minUniformBufferOffsetAlignment` field of the `VkPhysicalDeviceLimits` structure. Likewise, if the buffer binding is a storage buffer, then `range` must be less than or equal to the `maxStorageBufferRange` field of `VkPhysicalDeviceLimits`. For storage buffers, the `offset` parameter must be an integer multiple of the `minStorageBufferOffsetAlignment` field of the `VkPhysicalDeviceLimits` structure.

The `maxUniformBufferRange` and `maxStorageBufferRange` limits are guaranteed to be at least 16,384 and 2^{27}, respectively. If the buffers you're using fit within these limits, there's no reason to query them. Note that the guaranteed maximum size of a storage buffer is *much* larger than

that of a uniform buffer. If you have a very large amount of data, it may be worth considering using storage buffers over uniform buffers, even if access to the buffer's content is uniform in nature.

The `minUniformBufferOffsetAlignment` and `minStorageBufferOffsetAlignment` are guaranteed to be at most 256 bytes. Note that these are maximum minimums, and the values reported for a device may be smaller than this.

In addition to writing directly into descriptor sets, **vkUpdateDescriptorSets()** can copy descriptors from one set to another or between bindings in the same set. These copies are described in the array of `VkCopyDescriptorSet` structures passed through the `pDescriptorCopies` parameter. The definition of `VkCopyDescriptorSet` is

Click here to view code image

```
typedef struct VkCopyDescriptorSet {
    VkStructureType    sType;
    const void*        pNext;
    VkDescriptorSet    srcSet;
    uint32_t           srcBinding;
    uint32_t           srcArrayElement;
    VkDescriptorSet    dstSet;
    uint32_t           dstBinding;
    uint32_t           dstArrayElement;
    uint32_t           descriptorCount;
} VkCopyDescriptorSet;
```

The `sType` field of `VkCopyDescriptorSet` should be set to `VK_STRUCTURE_TYPE_COPY_DESCRIPTOR_SET`, and `pNext` should be set to `nullptr`. The handles of the source and destination descriptor sets are specified in the `srcSet` and `dstSet`, respectively. These can be the same set so long as the range of descriptors to copy does not overlap.

The `srcBinding` and `dstBinding` fields specify the binding indices of the source and destination descriptors, respectively. If the descriptors to be copied form an array of bindings, the indices of the start of the range of descriptors in the source and destination sets are specified in `srcArrayElement` and `dstArrayElement`, respectively. If the descriptors do not form arrays, both of these fields should be set to 0. The length of the array of descriptors to copy is specified in `descriptorCount`. If the copy is not of an array of descriptors, then `descriptorCount` should be set to 1.

When **vkUpdateDescriptorSets()** executes, the updates are performed by the host. Any access by the device to the descriptor sets referenced by the `pDescriptorWrites` or `pDescriptorCopies` arrays must be complete before **vkUpdateDescriptorSets()** is called. This includes work described in command buffers that have already been submitted but may not yet have completed execution.

All of the descriptor writes described by `pWriteDescriptors` are executed first, in the order in which they appear in the array, followed by all of the copies described by `pCopyDescriptors`. This means that if a particular binding is the destination of a write or copy operation more than once, only the last action on that binding as a destination will be visible after all the operations complete.

Binding Descriptor Sets

Just as with pipelines, to access the resources attached to a descriptor set, the descriptor set must be bound to the command buffer which will execute the commands that access those descriptors. There are also two binding points for descriptor sets—one for compute and one for graphics—which are the sets that will be accessed by the pipelines of the appropriate type.

To bind descriptor sets to a command buffer, call **vkCmdBindDescriptorSets()**, the prototype of which is

[Click here to view code image](#)

```
void vkCmdBindDescriptorSets (
    VkCommandBuffer                     commandBuffer,
    VkPipelineBindPoint                 pipelineBindPoint,
    VkPipelineLayout                    layout,
    uint32_t                            firstSet,
    uint32_t                            descriptorSetCount,
    const VkDescriptorSet*              pDescriptorSets,
    uint32_t                            dynamicOffsetCount,
    const uint32_t*                     pDynamicOffsets);
```

The command buffer to which the descriptor sets will be bound is specified in `commandBuffer`. The `pipelineBindPoint` argument specifies whether to bind the descriptor sets into the compute or graphics binding point by setting it to `VK_PIPELINE_BIND_POINT_COMPUTE` or `VK_PIPELINE_BIND_POINT_GRAPHICS`, respectively.

The pipeline layout that will be used by pipelines that will access the descriptors in the set specified in `layout`. This layout needs to be compatible with any pipeline that will use the sets and allows Vulkan to correctly configure the set bindings before a pipeline is bound to the command buffer. This means that the order in which you bind resources and pipelines into a command buffer doesn't matter so long as by the time you issue any drawing or dispatch commands, the layouts match.

To bind a subset of the sets accessible to the pipeline layout, use the `firstSet` and `descriptorSetCount` arguments to specify the index of the first set to bind and the number of sets, respectively. `pDescriptorSets` is a pointer to an array of `VkDescriptorSet` handles to the sets to be bound. These are obtained from calls to **vkAllocateDescriptorSets()** discussed earlier.

vkCmdBindDescriptorSets() is also responsible for setting the offsets used in any dynamic uniform or shader storage bindings. These are passed in the `dynamicOffsetCount` and `pDynamicOffsets` parameters. `dynamicOffsetCount` is the number of dynamic offsets to set, and `pDynamicOffsets` is a pointer to an array of `dynamicOffsetCount` 32-bit offsets. For each dynamic uniform or shader storage buffer in the descriptor set(s) being bound, there should be one offset specified in the `pDynamicOffsets` array. This offset is added to the base of the buffer view bound to the block in the descriptor set. This allows uniform and shader storage blocks to be re-bound to sections of a larger buffer without needing to create a new buffer view each time the offset is updated. Some implementations may need to pass additional information to the shader to account for this offset, but it is still generally faster than creating buffer views on the fly.

Uniform, Texel, and Storage Buffers

Shaders can access the content of buffer memory directly through three types of resources:

- Uniform blocks provide fast access to constant (read-only) data stored in buffer objects. They are declared as though they were structures in a shader and are attached to memory using a buffer resource bound into the descriptor set.
- Shader storage blocks provide read-write access to buffer objects. Declared similarly to uniform blocks, data is arranged as though it were a structure but can be written to. Shader storage blocks also support atomic operations.
- Texel buffers provide access to long, linear arrays of formatted texel data. They are read-only, and a texel buffer binding performs format conversion from the underlying data format into the floating-point representation that the shader expects when the buffer is read.

Which type of resource you use depends on how you want to access it. The maximum size of a uniform block is often limited, while access to it is generally very fast. On the other hand, the maximum size of a shader storage block is very large, but in some implementations, access to it could be slower—especially if write operations are enabled. For access to large arrays of formatted data, a texel buffer is probably the best choice.

Unform and Shader Storage Blocks

To declare a uniform block in GLSL, use the **uniform** keyword, as shown in [Listing 6.11](#). A shader storage block is declared similarly, except that the **uniform** keyword is omitted and the **buffer** keyword is used instead. The listing shows an example of each. The uniform block uses a descriptor set and binding index that is specified using a GLSL **layout** qualifier.

Listing 6.11: Declaring Uniform and Shader Blocks in GLSL

Click here to view code image

```
layout (set = 0, binding = 1) uniform my_uniform_buffer_t
{
    float foo;
    vec4 bar;
    int baz[42];
} my_uniform_buffer;

layout (set = 0, binding = 2) buffer my_storage_buffer_t
{
    int peas;
    float carrots;
    vec3 potatoes[99];
} my_storage_buffer;
```

The layout of variables inside the block is determined by a set of rules, which by default follow the **std140** rules for uniform blocks and **std430** rules for shader storage blocks. These rule sets are named for the version of GLSL in which they were introduced. The set of rules for packing data into memory may be changed by specifying a different layout in GLSL. Vulkan itself, however, does not automatically assign offsets to members of the block. This is the job of the front-end compiler that produces the SPIR-V shader that Vulkan consumes.

The resulting SPIR-V shader must conform to either the **std140** or **std430** layout rules (the latter being more flexible than the former), although those rules are not explicitly part of the SPIR-V specification. When the declarations are translated to SPIR-V by the front-end compiler, the members

of the block are explictly assigned locations. If the shader is generated from something other than GLSL, such as from another high-level language, or programmatically by a component of an application, then so long as the offsets assigned by the SPIR-V generator conform to the appropriate rule set, then the shader will work.

Listing 6.12 shows the shader from Listing 6.11 after it has been translated to SPIR-V by the reference compiler.

Listing 6.12: Declaring Uniform and Shader Blocks in SPIR-V

Click here to view code image

```
        OpCapability Shader
%1 = OpExtInstImport "GLSL.std.450"
        OpMemoryModel Logical GLSL450
        OpEntryPoint Vertex %4 "main"
        OpSource GLSL 450
        OpName %4 "main"
        ;; Name the my_uniform_buffer_t block and its members.
        OpName %12 "my_uniform_buffer_t"
        OpMemberName %12 0 "foo"
        OpMemberName %12 1 "bar"
        OpMemberName %12 2 "baz"
        OpName %14 "my_uniform_buffer"
        ;; Name the my_storage_buffer_t block and its members.
        OpName %18 "my_storage_buffer_t"
        OpMemberName %18 0 "peas"
        OpMemberName %18 1 "carrots"
        OpMemberName %18 2 "potatoes"
        OpName %20 "my_storage_buffer"
        OpDecorate %11 ArrayStride 16
        ;; Assign offsets to the members of my_uniform_buffer_t.
        OpMemberDecorate %12 0 Offset 0
        OpMemberDecorate %12 1 Offset 16
        OpMemberDecorate %12 2 Offset 32
        OpDecorate %12 Block
        OpDecorate %14 DescriptorSet 0
        OpDecorate %14 Binding 1
        OpDecorate %17 ArrayStride 16
        ;; Assign offsets to the members of my_storage_buffer_t.
        OpMemberDecorate %18 0 Offset 0
        OpMemberDecorate %18 1 Offset 4
        OpMemberDecorate %18 2 Offset 16
        OpDecorate %18 BufferBlock
        OpDecorate %20 DescriptorSet 0
        OpDecorate %20 Binding 2
        ...
```

As you can see from Listing 6.12, the compiler has explicitly assigned offsets to each member of the blocks declared in Listing 6.11. Any mention of `std140` and `std430` is absent from the SPIR-V version of the shader.[1]

1. **std140** and **std430** are also absent from the GLSL version of the shader, but they are implied and assumed by the front-end compiler.

Texel Buffers

A *texel buffer* is a special type of buffer binding used in a shader that can perform format conversion when the data is read. Texel buffers are read-only and are declared in GLSL using a **samplerBuffer** typed variable, as shown in Listing 6.13. Sampler buffers can return floating-point or signed or unsigned integer data to the shader. An example of each is shown in Listing 6.13.

Listing 6.13: Declaring Texel Buffers in GLSL

Click here to view code image

```
layout (set = 0, binding = 3) uniform samplerBuffer my_float_texel_buffer;
layout (set = 0, binding = 4) uniform isamplerBuffer
my_signed_texel_buffer;
layout (set = 0, binding = 5) uniform usamplerBuffer
my_unsigned_texel_buffer;
```

When translated to SPIR-V using the reference compiler, the declarations from Listing 6.13 produce the SPIR-V shader shown in Listing 6.14.

Listing 6.14: Declaring Texel Buffers in SPIR-V

Click here to view code image

```
       OpCapability Shader
       OpCapability SampledBuffer
 %1 = OpExtInstImport "GLSL.std.450"
       OpMemoryModel Logical GLSL450
       OpEntryPoint Vertex %4 "main"
       OpSource GLSL 450
       OpName %4 "main"
       ;; Name our texel buffers.
       OpName %10 "my_float_texel_buffer"
       OpName %15 "my_signed_texel_buffer"
       OpName %20 "my_unsigned_texel_buffer"
       ;; Assign set and binding decorations.
       OpDecorate %10 DescriptorSet 0
       OpDecorate %10 Binding 3
       OpDecorate %15 DescriptorSet 0
       OpDecorate %15 Binding 4
       OpDecorate %20 DescriptorSet 0
       OpDecorate %20 Binding 5
 %2 = OpTypeVoid
 %3 = OpTypeFunction %2
       ;; Declare the three texel buffer variables.
 %6 = OpTypeFloat 32
 %7 = OpTypeImage %6 Buffer 0 0 0 1 Unknown
 %8 = OpTypeSampledImage %7
 %9 = OpTypePointer UniformConstant %8
%10 = OpVariable %9 UniformConstant
```

```
%11 = OpTypeInt 32 1
%12 = OpTypeImage %11 Buffer 0 0 0 1 Unknown
%13 = OpTypeSampledImage %12
%14 = OpTypePointer UniformConstant %13
%15 = OpVariable %14 UniformConstant
%16 = OpTypeInt 32 0
%17 = OpTypeImage %16 Buffer 0 0 0 1 Unknown
%18 = OpTypeSampledImage %17
%19 = OpTypePointer UniformConstant %18
%20 = OpVariable %19 UniformConstant
        ...
```

To fetch from a texel buffer in GLSL, the `texelFetch` function is used with the sampler variable to read individual texels. A **`samplerBuffer`** (or the corresponding signed or unsigned integer variants, **`isamplerBuffer`** and **`usamplerBuffer`**) can be thought of as a 1D texture that supports only point sampling. However, the maximum size of a texel buffer attached to one of these variables is generally much, much larger than the maximum size of a 1D texture. For example, the minimum required upper bound for a texel buffer in Vulkan is 65,535 elements, whereas the minimum required size for a 1D texture is only 4,096 texels. In some cases, implementations will support texel buffers that are gigabytes in size.

Push Constants

One special type of resource briefly introduced earlier is the push constant. A *push constant* is a uniform variable in a shader that can be used just like a member of a uniform block, but rather than being backed by memory, it is owned and updated by Vulkan itself.[2] As a consequence, new values for these constants can be *pushed* into the pipeline directly from the command buffer, hence the term.

> [2.] In reality, some implementations may still back push constants with device memory internally. However, it is likely that such implementations will have more optimal paths for updating that memory; some implementations will have fast, dedicated memory or registers for these constants; and a single call to **`vkCmdPushConstants()`** is likely to perform better than memory barriers associated with updating a uniform block in any case.

Push constants are logically considered to be part of the pipeline's resources and are therefore declared along with the other resources in the pipeline layout used to create pipeline objects. In the `VkPipelineLayoutCreateInfo` structure, two fields are used to define how many push constants will be used by the pipeline. Push constants belong to ranges, each defined by a `VkPushConstantRange` structure. The `pushConstantRanges` member of `VkPipelineLayoutCreateInfo` specifies the count of the number of ranges of push constants that are included in the pipeline's resource layout, and the `pPushConstantRanges` of `VkPipelineLayoutCreateInfo` is a pointer to an array of `VkPushConstantRange` structures, each defining a range of push constants used by the pipeline. The definition of `VkPushConstantRange` is

[**Click here to view code image**](#)

```
typedef struct VkPushConstantRange {
    VkShaderStageFlags    stageFlags;
    uint32_t              offset;
    uint32_t              size;
} VkPushConstantRange;
```

The space used for push constants is abstracted as though it were a contiguous region of memory, even if that may not be the case in practice in some implementations. In some implementations, each shader stage has its own space for constant storage, and in such implementations, passing a single constant to multiple shading stages may require broadcasting it and consuming more resources. The stages that will "see" each range of constants are included in the `stageFlags` field of `VkPushConstantRange`. This is a bitwise combination of a selection of flags from `VkShaderStageFlagBits`. The starting offset and size of the region are specified in `offset` and `size`, respectively.

To consume push constants inside the pipeline, variables representing them are declared in the pipeline's shaders. A push constant is declared in a SPIR-V shader using the `PushConstant` storage class on a variable declaration. In GLSL, such a declaration can be produced by declaring a uniform block with the `push_constant` layout qualifier. There is one such block declaration available per pipeline. Logically, it has the same "in-memory" layout as a **std430** block. However, this layout is used only to compute offsets to members and may not be how a Vulkan implementation internally represents the data in the block.

Listing 6.15 shows a GLSL declaration of a push constant block, and Listing 6.16 shows the resulting SPIR-V.

Listing 6.15: Declaring Push Constants in GLSL

Click here to view code image

```
layout (push_constant) uniform my_push_constants_t
{
    int bourbon;
    int scotch;
    int beer;
} my_push_constants;
```

Listing 6.16: Declaring Push Constants in SPIR-V

Click here to view code image

```
    OpCapability Shader
%1 = OpExtInstImport "GLSL.std.450"
    OpMemoryModel Logical GLSL450
    OpEntryPoint GLCompute %4 "main"
    OpExecutionMode %4 LocalSize 1 1 1
    OpSource GLSL 450
    OpName %4 "main"
    ;; Name the push constant block and its members.
    OpName %7 "my_push_constants_t"
    OpMemberName %7 0 "bourbon"
    OpMemberName %7 1 "scotch"
    OpMemberName %7 2 "beer"
    OpName %9 "my_push_constants"
    ;; Assign offsets to the members of the push constant block.
    OpMemberDecorate %7 0 Offset 0
    OpMemberDecorate %7 1 Offset 4
    OpMemberDecorate %7 2 Offset 8
```

```
      OpDecorate %7 Block
%2 = OpTypeVoid
%3 = OpTypeFunction %2
%6 = OpTypeInt 32 1
%7 = OpTypeStruct %6 %6 %6
      ;; Declare the push constant block itself.
%8 = OpTypePointer PushConstant %7
%9 = OpVariable %8 PushConstant
      ...
```

Push constants become part of the layout of the pipeline that will use them. When push constants are included in a pipeline, they may consume some of the resources that Vulkan would otherwise use to track pipeline or descriptor bindings. Therefore, you should treat push constants as relatively precious resources.

To update the content of one or more push constants, call **vkCmdPushConstants()**, the prototype of which is

[Click here to view code image](#)

```
void vkCmdPushConstants (
    VkCommandBuffer                        commandBuffer,
    VkPipelineLayout                       layout,
    VkShaderStageFlags                     stageFlags,
    uint32_t                               offset,
    uint32_t                               size,
    const void*                            pValues);
```

The command buffer that will perform the update is specified in commandBuffer, and the layout that defines the locations of the push constants is specified in layout. This layout must be compatible with any pipeline that is subsequently bound and used in dispatch or drawing commands.

The stages that will need to see the updated constants should be specified in stageFlags. This is a bitwise combination of some of the flags from the VkShaderStageFlagBits enumeration. Although there is only one push constant block available to each pipeline, it may be that in some implementations, push constants are implemented by using per-stage resources. When stageFlags is set accurately, performance may increase by allowing Vulkan to not update the stages that aren't included. Be careful, though: In implementations that do support broadcasting constants across stages at no cost, these flags might be ignored, and your shaders may see the updates anyway.

As push constants are logically represented as backed by memory with a **std430** layout, the content of each push constant "lives" at an offset from the beginning of the block that can be computed using the **std430** rules. The offset of the first constant to update within this virtual block is specified in offset, and the size of the update, in bytes, is specified in size.

A pointer to the data to place in the push constants is passed in pValues. Typically, this will be a pointer to an array of uint32_t or **float** variables. Both offset and size must be a multiple of 4 to align them correctly with respect to the size of these data types. When **vkCmdPushConstants()** is executed, it is as though the contents of the array were copied directly into a **std430** block.

You are free to replace the content of the array or free the memory immediately after calling **vkCmdPushConstants()**. The data in the array is consumed immediately by the command, and

the value of the pointer is not retained. Therefore, it's perfectly fine to set `pValues` to something that lives on the stack or to the address of a local variable.

The total amount of space available for push constants in a single pipeline (or pipeline layout) can be determined by inspecting the `maxPushConstantsSize` field of the device's `VkPhysicalDeviceLimits` structure. This is guaranteed to be at least 128 bytes (enough for a couple of 4 × 4 matrices). It's not particularly large, but if this is sufficient, there's no reason to query the limit. Again, treat push constants as scarce resources. Prefer to use a normal uniform block for larger data structures, and use push constants for single integers or very frequently updated data.

Sampled Images

When shaders read from images, they can do so in one of two ways. The first is to perform a raw load, which directly reads formatted or unformatted data from a specific location in the image, and the second is to *sample* the image using a sampler. Sampling can include operations such as performing basic transformations on the image coordinates or filtering texels to smooth the image data returned to the shader.

The state of a sampler is represented by a sampler object, which is bound into descriptor sets just as images and buffers are. To create a sampler object, call **vkCreateSampler()**, the prototype of which is

[**Click here to view code image**](#)

```
VkResult vkCreateSampler (
    VkDevice                            device,
    const VkSamplerCreateInfo*          pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkSampler*                          pSampler);
```

The device that will create the sampler is passed in `device`, and the remaining parameters of the sampler are passed through a pointer to an instance of the `VkSamplerCreateInfo` structure in `pCreateInfo`. The upper bound on the total number of samplers that can be created by a device is implementation-dependent. It's guaranteed to be at least 4,000. If there's a possibility that your application may create more than this limit, then you need to check the device's level of support for creating large numbers of samplers. The total number of samplers that a device can manage is contained in the `maxSamplerAllocationCount` field of its `VkPhysicalDeviceLimits` structure, which you can obtain from a call to **vkGetPhysicalDeviceProperties()**.

The definition of `VkSamplerCreateInfo` is

[**Click here to view code image**](#)

```
typedef struct VkSamplerCreateInfo {
    VkStructureType         sType;
    const void*             pNext;
    VkSamplerCreateFlags    flags;
    VkFilter                magFilter;
    VkFilter                minFilter;
    VkSamplerMipmapMode     mipmapMode;
    VkSamplerAddressMode    addressModeU;
    VkSamplerAddressMode    addressModeV;
    VkSamplerAddressMode    addressModeW;
```

```
    float                   mipLodBias;
    VkBool32                anisotropyEnable;
    float                   maxAnisotropy;
    VkBool32                compareEnable;
    VkCompareOp             compareOp;
    float                   minLod;
    float                   maxLod;
    VkBorderColor           borderColor;
    VkBool32                unnormalizedCoordinates;
} VkSamplerCreateInfo;
```

The `sType` field of `VkSamplerCreateInfo` should be set to `VK_STRUCTURE_TYPE_SAMPLER_CREATE_INFO`, and `pNext` should be set to `nullptr`. The `flags` field is reserved for future use and should be set to zero.

Image Filtering

The `magFilter` and `minFilter` fields specify the filtering mode to be used when the image is magnified or minified, respectively. Whether an image is magnified or minified is determined by comparing sampling coordinates across adjacent pixels being shaded. If the gradient of the sampling coordinates is greater than one, then the image is minified; otherwise, it is magnified. `magFilter` and `minFilter` are both members of the `VkFilter` enumeration. The members of `VkFilter` are

- `VK_FILTER_NEAREST`: When sampling, the nearest texel in the image is chosen and returned directly to the shader.
- `VK_FILTER_LINEAR`: A 2 × 2 footprint containing the texel coordinates is used to produce a weighted average of four texels, and this average is returned to the shader.

The `VK_FILTER_NEAREST` mode causes Vulkan to simply select the nearest texel to the requested coordinates when sampling from an image. In many cases, this can lead to a blocky or aliased image, causing shimmering artifacts in the rendered picture. The `VK_FILTER_LINEAR` mode tells Vulkan to apply linear filtering to the image when it is sampled.

When you are filtering an image with linear filtering, the requested sample may lie somewhere between two texel centers in 1D, four centers in 2D, and so on. Vulkan will read from the surrounding texels and then combine the results using a weighted sum of the values based on the distance to each center. This is illustrated in Figure 6.2. In the figure, a sample is taken at the ×, which lies between the four texel centers marked A, B, C, and D. Regardless of the *integer* part of the texture coordinate *{u, v}*, the *fractional* part of the texture coordinate is given by *{α, β}*.
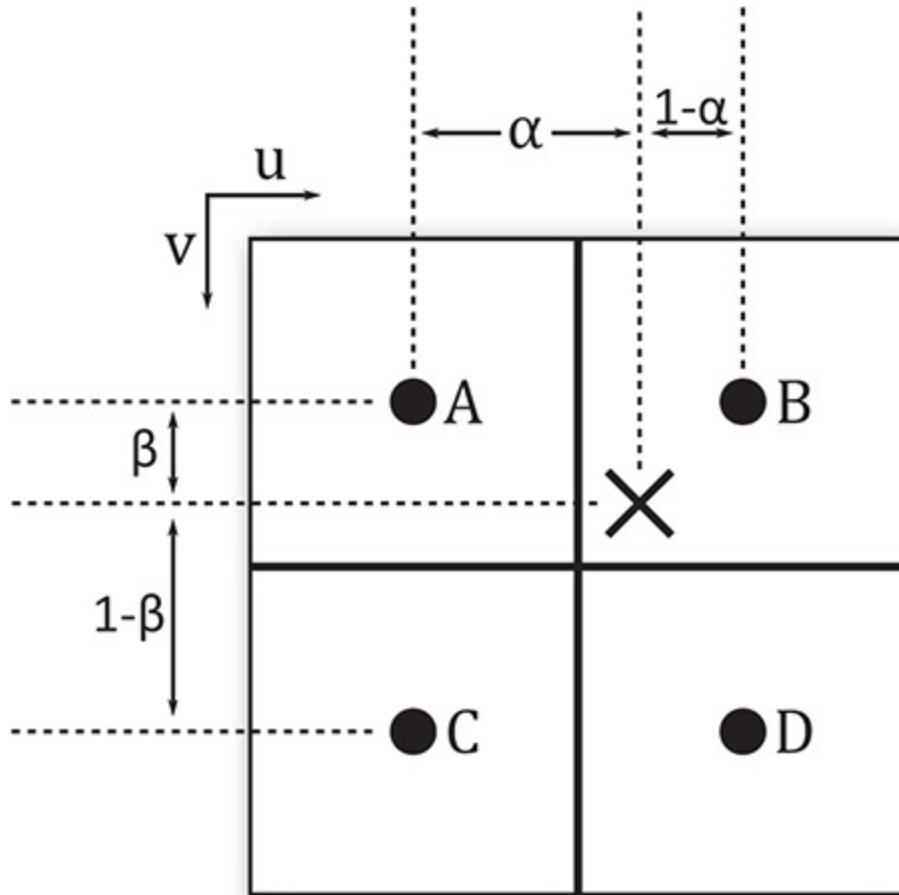
Figure 6.2: Linear Sampling

To form a linear weighted sum of the texels at A and B, their values are simply combined according to the relationship

$T_{u0} = \alpha A + (1 - \alpha)\, B$

This can be writen as

$T_{u0} = B + \alpha\,(B - A)$

Likewise, a weighted sum of C and D is formed as

$T_{u1} = \alpha C + (1 - \alpha)\, D$ or

$T_{u1} = D - \alpha\,(D - C)$

The two temporary values T_{u0} and T_{u1} can then be combined into a single, weighted sum using a similar mechanism, but with β:

$T = \beta T_{u0} + (1 - \beta)\, T_{u1}$ or

$T = T_{u1} + \beta(T_{u1} - T_{u0})$

This can be extended in any number of dimensions, although only texture dimensionalities up to three are defined by Vulkan.

Mipmapping

The `mipmapMode` field specifies how mipmaps are used in the image when it is sampled. This is a member of the `VkSamplerMipmapMode` enumeration, whose members have the following meanings:

- `VK_SAMPLER_MIPMAP_MODE_NEAREST`: The computed level-of-detail is rounded to the nearest integer, and that level is used to select the mipmap level. If sampling from the base level, then the filtering mode specified in `magFilter` is used to sample from that level; otherwise, the `minFilter` filter is used.

- `VK_SAMPLER_MIPMAP_MODE_LINEAR`: The computed level of detail is rounded both up and down, and the two resulting mipmap levels are sampled. The two resulting texel values are then blended and returned to the shader.

To select a mipmap from the image, Vulkan will compute the derivative of the coordinates used to sample from the texture. The exact math is covered in some detail in the Vulkan specification. In short, the level selected is the log_2 of the maximum of the derivatives of each of the texture coordinate dimensions. This level can also be biased using parameters taken from the sampler or supplied by the shader, or it can be entirely specified in the shader. Regardless of its source, the result may not be an exact integer.

When the mipmap mode is `VK_SAMPLER_MIPMAP_MODE_NEAREST`, then the selected mipmap level is simply rounded down to the next-lowest integer, and then that level is sampled as though it were a single-level image. When the mipmap mode is `VK_SAMPLER_MIPMAP_MODE_LINEAR`, a sample is taken from each of the next-lower and next-higher levels using the filtering mode selected by the `minFilter` field, and then those two samples are further combined using a weighted average, similarly to how the samples are combined during linear sampling, as described earlier.

Note that this filtering mode applies only to minification, which is the process of sampling from a mipmap level other than the base level. When the log_2 of the texture coordinate derivatives is less than 1, then the 0^{th} level is selected, so only a single level is available for sampling. This is known as magnification and uses the filtering mode specified in `magFilter` to sample from the base level only.

The next three fields in `VkSamplerCreateInfo`—`addressModeU`, `addressModeV`, and `addressModeW`—are used to select the transform that is applied to texture coordinates that would otherwise sample outside the image. The following modes are available:

- `VK_SAMPLER_ADDRESS_MODE_REPEAT`: As the texture coordinate progresses from 0.0 to 1.0 and beyond, it is wrapped back to 0.0, effectively using only the fractional part of the coordinate to sample from the image. The effect is to tile the image indefinitely.

- `VK_SAMPLER_ADDRESS_MODE_MIRRORED_REPEAT`: The texture coordinate progresses from 0.0 to 1.0 as normal, and then in the range 1.0 to 2.0, the fractional part is subtracted from 1.0 to form a new coordinate moving back toward 0.0. The effect is to alternately tile the normal and mirror-image version of a texture.

- `VK_SAMPLER_ADDRESS_MODE_CLAMP_TO_EDGE`: Texture coordinates beyond 1.0 are clamped to 1.0, and negative coordinates are clamped to 0.0. This clamped coordinate is used to sample from the image. The effect is that the texels along the edge of the image are used to fill any area that would be sampled outside the image.

- `VK_SAMPLER_ADDRESS_MODE_CLAMP_TO_BORDER`: Sampling from the texture outside its bounds will result in texels of the border color, as specified in the `borderColor` field, being returned rather than data from the image.

- `VK_SAMPLER_ADDRESS_MODE_MIRROR_CLAMP_TO_EDGE`: This is a hybrid mode that first applies a single mirroring of the texture coordinate and then behaves like `VK_SAMPLER_ADDRESS_MODE_CLAMP_TO_EDGE`.

The effect of each of the sampler addressing modes applied to an image is shown in [Figure 6.3](#). In the figure, the top-left image shows the result of the `VK_SAMPLER_ADDRESS_MODE_REPEAT` addressing mode. As you can see, the texture is simply repeated across the frame. The top-right image shows the result of `VK_SAMPLER_ADDRESS_MODE_MIRRORED_REPEAT`. Each alternate repetition of the texture is mirrored in the X or Y direction.
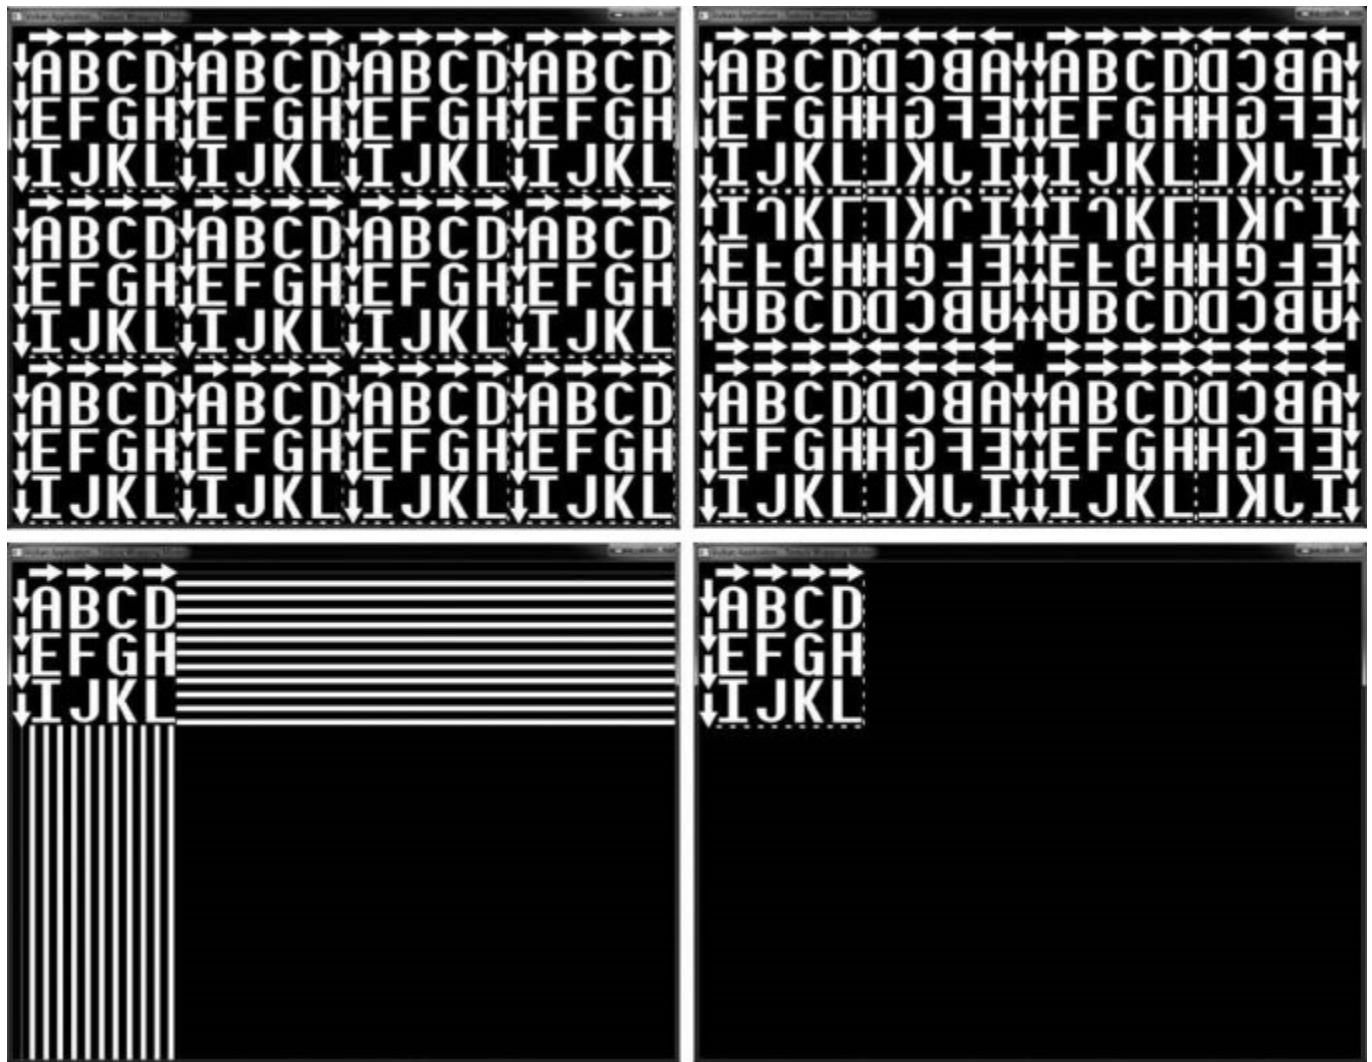
Figure 6.3: Effect of Sampling Modes

The bottom-left image applies the `VK_SAMPLER_ADDRESS_MODE_CLAMP_TO_EDGE` addressing mode to the texture. Here, the last column or row of pixels is repeated indefinitely after the sampling coordinates leave the texture. Finally, the bottom-right image shows the result of the `VK_SAMPLER_ADDRESS_MODE_CLAMP_TO_BORDER` mode. This texture was created with a

black border color, so the area outside the texture appears to be blank, but Vulkan is actually sampling black texels from this region. This allows you to see the original texture.

When the filter mode is `VK_FILTER_LINEAR`, wrapping or clamping the texture coordinate is applied to each of the generated coordinates in the 2×2 footprint used to create the resulting texel. The result is that filtering is applied as though the image really wrapped.

For the `VK_SAMPLER_ADDRESS_MODE_CLAMP_TO_BORDER` filter mode, when texels are sampled from the border (i.e., what would be outside the image), the border color is substituted rather than fetching data from the image. The color that is used depends on the value of the `borderColor` field. This is not a full color specification, but a member of the `VkBorderColor` enumeration, which allows one of a small, predefined set of colors to be selected. These are

- `VK_BORDER_COLOR_FLOAT_TRANSPARENT_BLACK`: Returns floating-point zeros to the shader in all channels
- `VK_BORDER_COLOR_INT_TRANSPARENT_BLACK`: Returns integer zeros to the shader in all channels
- `VK_BORDER_COLOR_FLOAT_OPAQUE_BLACK`: Returns floating-point zeros in the R, G, and B channels, and floating-point one in A
- `VK_BORDER_COLOR_INT_OPAQUE_BLACK`: Returns integer zeros in R, G, and B, and integer one in A
- `VK_BORDER_COLOR_FLOAT_OPAQUE_WHITE`: Returns floating-point ones to the shader in all channels
- `VK_BORDER_COLOR_INT_OPAQUE_WHITE`: Returns integer ones to the shader in all channels

The `mipLodBias` field of `VkSamplerCreateInfo` specifies a floating-point bias that is added to the computed level of detail before mipmap selection is made. This allows you to nudge the level of detail up or down the mipmap chain to make the resulting filtered texture look sharper or blurrier than it might otherwise.

If you want to use anisotropic filtering, set `anisotropyEnable` to `VK_TRUE`. The exact details of anisotropic filtering are implementation-dependent. Anisotropic filtering generally works by considering a projected footprint of the area to be sampled rather than using a fixed 2×2 footprint. An approximation to an area sample is formed by taking many samples within the footprint.

Because the number of samples taken can be quite large, anisotropic filtering can have a negative effect on performance. Also, under extreme cases, the projected footprint can be quite large, and this can result in a large area and correspondingly blurry filter result. To limit these effects, you can clamp the maximum amount of anisotropy by setting `maxAnisotropy` to a value between 1.0 and the maximum value supported by the device. You can determine this by calling **`vkGetPhysicalDeviceProperties()`** and inspecting the `maxSamplerAnisotropy` member of the embedded `VkPhysicalDeviceLimits` structure.

When a sampler is used with a depth image, it can be configured to perform a comparison operation and return the result of the comparison rather than the raw values stored in the image. When this mode is enabled, the comparison is performed on each sample taken from the image, and the resulting value is the fraction of the total samples taken that passed the test. This can be used to implement a technique known as *percentage closer filtering*, or PCF. To enable this mode, set `compareEnable` to `VK_TRUE`, and set the comparison operation in `compareOp`.

`compareOp` is a member of the `VkCompareOp` enumeration, which is used in many places in Vulkan. As you will see in [Chapter 7](#), "[Graphics Pipelines](#)," this is the same enumeration that is used to specify the depth test operation. The available operations and how they are interpreted in the context of shader accesses to depth resources are shown in [Table 6.2](#).

The sampler can be configured to restrict sampling to a subset of the mip levels in an image with mipmaps. The range of mipmaps to sample from is specified in `minLod` and `maxLod`, which contain the lowest (highest-resolution) and highest (lowest-resolution) mipmaps that should be sampled from, respectively. To sample from the entire mipmap chain, set `minLod` to 0.0, and set `maxLod` to a level of detail high enough that the computed level of detail will never be clamped.

Finally, `unnormalizedCoordinates` is a flag that, when set to `VK_TRUE`, indicates that the coordinates used to sample from the image are in units of raw texels, rather than a value that is normalized between 0.0 and 1.0 across each dimension of the texture. This allows texels to be explicitly fetched from the image. However, several restrictions exist in this mode. When `unnormalizedCoordinates` is `VK_TRUE`, `minFilter` and `magFilter` must be the same, `mipmapMode` must be `VK_SAMPLER_MIPMAP_MODE_NEAREST`, and `anisotropyEnable` and `compareEnable` must be `VK_FALSE`.

When you are done with a sampler, you should destroy it by calling **`vkDestroySampler()`**, the prototype of which is

Click here to view code image

```
void vkDestroySampler (
    VkDevice                        device,
    VkSampler                       sampler,
    const VkAllocationCallbacks*    pAllocator);
```

`device` is the device that owns the sampler object, and `sampler` is the sampler object to destroy. If a host memory allocator was used when the sampler was created, a compatible allocator should be passed through the `pAllocator` parameter; otherwise, `pAllocator` should be `nullptr`.

Function VK_COMPARE_OP_...	Meaning
ALWAYS	The comparison always passes; the return value will be 1.0.
NEVER	The comparison never passes; the return value will be 0.0.
LESS	The comparison passes if the shader's reference value is less than the value in the image.
LESS_OR_EQUAL	The comparison passes if the shader's reference value is less than or equal to the value in the image.
EQUAL	The comparison passes if the shader's reference value is equal to the value in the image.
NOT_EQUAL	The comparison passes if the shader's reference value is not equal to the value in the image.
GREATER	The comparison passes if the shader's reference value is greater than the value in the image.
GREATER_OR_EQUAL	The comparison passes if the shader's reference value is greater than or equal to the value in the image.

Table 6.2: Texture Comparison Functions

Summary

This chapter covered the basics of the shading language supported by Vulkan, SPIR-V, including how SPIR-V shader modules are consumed by Vulkan and how pipelines containing those shaders are constructed. You saw how to construct a compute shader and use it to create a compute pipeline, how to dispatch work into that pipeline, and how a pipeline can access resources in order to consume and produce data. In upcoming chapters, we will build on the concept of pipelines to produce pipeline objects with multiple stages and make use of more advanced features.

Chapter 7. Graphics Pipelines

What You'll Learn in This Chapter

 • What the Vulkan graphics pipeline looks like

 • How to create a graphics pipeline object

 • How to draw graphical primitives with Vulkan

Perhaps the most common use of Vulkan is as a graphics API. Graphics are a fundamental part of Vulkan and drive the core of almost any visual application. Graphics processing in Vulkan can be seen as a pipeline that takes graphics commands through the many stages required to produce a picture on a display. This chapter covers the basics of graphics pipelines in Vulkan and introduces our first graphics example.

The Logical Graphics Pipeline

The graphics pipeline in Vulkan can be seen as a production line, where commands enter the front of the pipeline and are processed in stages. Each stage performs some kind of transform, taking the commands and their associated data and turning them into something else. By the end of the pipeline, the commands have been transformed into colorful pixels making up your output picture.

Many parts of the graphics pipeline are optional and can be disabled or might not even be supported by a Vulkan implementation. The only part of the pipeline that an application *must* enable is the *vertex shader*. The full Vulkan graphics pipeline is shown in Figure 7.1. However, don't be alarmed; we'll introduce each stage gently in this chapter and dig into more details in subsequent parts of the book.
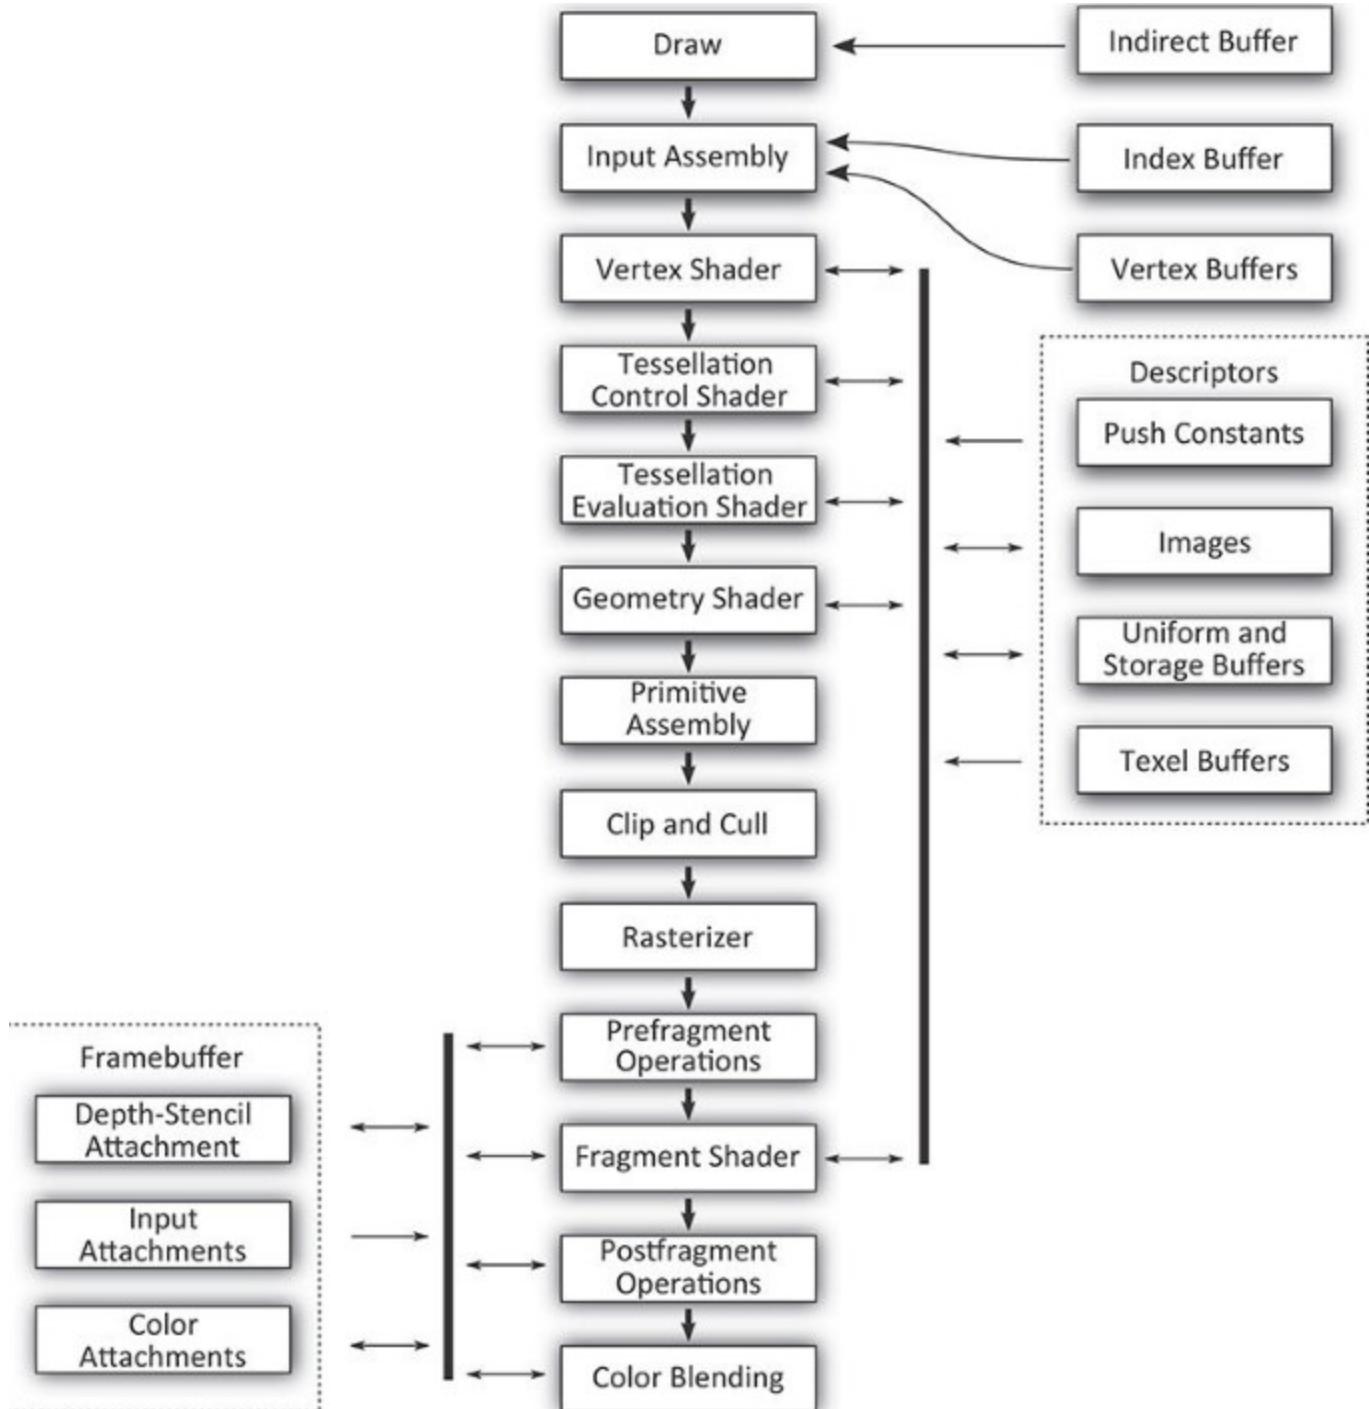
Figure 7.1: The Full Vulkan Graphics Pipeline

The following is a brief description of each stage of the pipeline and what it does.

- Draw: This is where your commands enter the Vulkan graphics pipeline. Typically, a small processor or dedicated piece of hardware inside the Vulkan device interprets the commands in the command buffer and directly interacts with the hardware to induce work.

- Input assembly: This stage reads the *index* and *vertex* buffers that contain information about the vertices making up the draw you've sent.

- Vertex shader: This is where the vertex shader executes. It takes as input the properties of the vertex and prepares transformed and processed vertex data for the next stage.

- Tessellation control shader: This programmable shading stage is responsible for producing tessellation factors and other per-patch data that is used by the fixed-function tessellation engine.

- Tessellation primitive generation: Not shown in [Figure 7.1](#), this fixed function stage uses the tessellation factors produced in the tessellation control shader to break patch primitives into many smaller, simpler primitives ready to be shaded by the tessellation evaluation shader.

- Tessellation evaluation shader: This shading stage runs on each new vertex produced by the tessellation primitive generator. It operates similarly to a vertex shader except that the incoming vertices are generated rather than read from memory.

- Geometry shader: This shading stage operates on full primiitves. The primitives might be points, lines or triangles, or special variations of them that include additional vertices surrounding them. This stage also has the ability to change the primitive type midpipeline.

- Primitive assembly: This stage groups vertices produced by the vertex, tessellation, or geometry stage and groups them into primitives suitable for rasterization. It also culls and clips primitives and transforms them into the appropriate viewport.

- Clip and cull: This fixed-function stage determines which parts of which primitives might contribute to the output image and discards parts of primitives that do not, forwarding potentially visible primitives to the rasterizer.

- Rasterizer: Rasterization is the fundamental core of all graphics in Vulkan. The rasterizer takes assembled primitives that are still represented by a sequence of vertices and turns them into individual fragments, which may become the pixels that make up your image.

- Prefragment operations: Several operations can be performed on fragments once their positions are known but before they are shaded. These prefragment operations include depth and stencil tests when they are enabled.

- Fragment assembly: Not shown in the figure, the fragment assembly stage takes the output of the rasterizer along with any per-fragment data and sends it, as a group, into the fragment shading stage.

- Fragment shader: This stage runs the final shader in the pipeline, which is responsible for computing the data that will be sent on to the final fixed-function processing stages that follow.

- Postfragment operations: In some circumstances, the fragment shader modifies data that would normally be used in prefragment operations. In these cases, those prefragment operations move to the postfragment stage and are executed here.

- Color blending: The color operations take the final results of the fragment shader and postfragment operations and use them to update the framebuffer. The color operations include blending and logic operations.

As you can tell, there are a lot of interrelated stages in the graphics pipeline. Unlike the compute pipeline introduced in [Chapter 6](#), "[Shaders and Pipelines](#)," the graphics pipeline contains not only the configuration of a wide selection of fixed functionality, but also up to five shader stages. Further, depending on the implementation, some of the logically fixed-function stages are actually at least partially implemented in shader code generated by drivers.

The purpose of representing the graphics pipeline as an object in Vulkan is to provide the implementation as much information as needed to move parts of the pipeline between fixed-function hardware and programmable shader cores. If the information were not all available at the same time in the same object, it would mean that some implementations of Vulkan may need to recompile a shader based on configurable state. The set of states contained in the graphics pipeline has been carefully chosen to prevent this, making switching states as fast as possible.

The fundamental unit of drawing in Vulkan is a *vertex*. Vertices are grouped into *primitives* and processed by the Vulkan pipeline. The simplest drawing command in Vulkan is **vkCmdDraw()**, whose prototype is

[**Click here to view code image**](#)

```
void vkCmdDraw (
    VkCommandBuffer                     commandBuffer,
    uint32_t                            vertexCount,
    uint32_t                            instanceCount,
    uint32_t                            firstVertex,
    uint32_t                            firstInstance);
```

Like other Vulkan commands, **vkCmdDraw()** appends a command to a command buffer that will later be executed by the device. The command buffer to append to is specified in `commandBuffer`. The number of vertices to push into the pipeline is specified in `vertexCount`. If you want to draw the same set of vertices over and over with slightly different parameters, you can specify the number of *instances* in `instanceCount`. This is known as *instancing*, and we'll cover that later in this chapter. For now, we can just set `instanceCount` to 1. It's also possible to start drawing from a vertex or instance other than 0. To do this, we can use `firstVertex` and `firstInstance`, respectively. Again, we'll cover that later. For the time being, we'll set both of these parameters to 0.

Before you can draw anything, you must bind a graphics pipeline to the command buffer, and before that, you must create a graphics pipeline. Undefined behavior (generally bad behavior) will occur if you try drawing without binding a pipeline first.

When you call **vkCmdDraw()**, `vertexCount` vertices are generated and pushed into the current Vulkan graphics pipeline. For each vertex, input assembly is executed, followed by your vertex shader. Declaring inputs beyond what is provided for you by Vulkan is optional, but having a vertex shader is not. Thus, the simplest possible graphics pipeline consists only of a vertex shader.

Renderpasses

One of the things that distinguishes a Vulkan graphics pipeline from a compute pipeline is that, usually, you'll be using the graphics pipeline to render pixels into images that you will either further process or display to the user. In complex graphics applications, the picture is built up over many passes where each pass is responsible for producing a different part of the scene, applying full-frame effects such as postprocessing or composition, rendering user interface elements, and so on.

Such passes can be represented in Vulkan using a renderpass object. A single renderpass object encapsulates multiple passes or rendering phases over a single set of output images. Each pass within the renderpass is known as a *subpass*. Renderpass objects can contain many subpasses, but even in simple applications with only a single pass over a single output image, the renderpass object contains information about that output image.

All drawing must be contained inside a renderpass. Further, graphics pipelines need to know where they're rendering to; therefore, it's necessary to create a renderpass object before creating a graphics pipeline so that we can tell the pipeline about the images it'll be producing. Renderpasses are covered in great depth in Chapter 13, "Multipass Rendering." In this chapter, we'll create the simplest possible renderpass object that will allow us to render into an image.

To create a renderpass object, call **vkCreateRenderPass()**, the prototype of which is

```
VkResult vkCreateRenderPass (
    VkDevice                       device,
    const VkRenderPassCreateInfo*  pCreateInfo,
    const VkAllocationCallbacks*   pAllocator,
    VkRenderPass*                  pRenderPass);
```

The `device` parameter to **vkCreateRenderPass()** is the device that will create the renderpass object, and `pCreateInfo` points to a structure defining the renderpass. This is an instance of the `VkRenderPassCreateInfo` structure, whose definition is

```
typedef struct VkRenderPassCreateInfo {
    VkStructureType                sType;
    const void*                    pNext;
    VkRenderPassCreateFlags        flags;
    uint32_t                       attachmentCount;
    const VkAttachmentDescription* pAttachments;
    uint32_t                       subpassCount;
    const VkSubpassDescription*    pSubpasses;
    uint32_t                       dependencyCount;
    const VkSubpassDependency*     pDependencies;
} VkRenderPassCreateInfo;
```

The `sType` field of `VkRenderPassCreateInfo` should be set to `VK_STRUCTURE_TYPE_RENDERPASS_CREATE_INFO`, and `pNext` should be set to `nullptr`. The `flags` field is reserved for future use and should be set to zero.

`pAttachments` is a pointer to an array of `attachmentCount` `VkAttachmentDescription` structures that define the *attachments* associated with the renderpass. Each of these structures defines a single image that is to be used as an input, output, or both within one or more of the subpasses in the renderpass. If there really are no attachments associated with the renderpass, you can set `attachmentCount` to zero and `pAttachments` to `nullptr`. However, outside of some advanced use cases, almost all graphics rendering will use at least one attachment. The definition of `VkAttachmentDescription` is

```
typedef struct VkAttachmentDescription {
    VkAttachmentDescriptionFlags   flags;
    VkFormat                       format;
    VkSampleCountFlagBits          samples;
    VkAttachmentLoadOp             loadOp;
    VkAttachmentStoreOp            storeOp;
```

```
    VkAttachmentLoadOp              stencilLoadOp;
    VkAttachmentStoreOp             stencilStoreOp;
    VkImageLayout                   initialLayout;
    VkImageLayout                   finalLayout;
} VkAttachmentDescription;
```

The `flags` field is used to give Vulkan additional information about the attachment. The only defined bit is `VK_ATTACHMENT_DESCRIPTION_MAY_ALIAS_BIT`, which, if set, indicates that the attachment might be using the same memory as another attachment referenced by the same renderpass. This tells Vulkan not to do anything that might make data in that attachment inconsistent. This bit can be used in some advanced cases where memory is at a premium and you are trying to optimize its usage. In most cases, `flags` can be set to zero.

The `format` field specifies the format of the attachment. This is one of the `VkFormat` enumerations and should match the format of the image used as the attachment. Likewise, `samples` indicates the number of samples in the image and is used for *multisampling*. When multisampling is not in use, set `samples` to `VK_SAMPLE_COUNT_1_BIT`.

The next four fields specify what to do with the attachment at the beginning and end of the renderpass. The load operations tell Vulkan what to do with the attachment when the renderpass begins. This can be set to one of the following values:

- `VK_ATTACHMENT_LOAD_OP_LOAD` indicates that the attachment has data in it already and that you want to keep rendering to it. This causes Vulkan to treat the contents of the attachment as valid when the renderpass begins.

- `VK_ATTACHMENT_LOAD_OP_CLEAR` indicates that you want Vulkan to clear the attachment for you when the renderpass begins. The color to which you want to clear the attachments is specified when the renderpass has begun.

- `VK_ATTACHMENT_LOAD_OP_DONT_CARE` indicates that you don't care about the content of the attachment at the beginning of the renderpass and that Vulkan is free to do whatever it wishes with it. You can use this if you plan to explicitly clear the attachment or if you know that you'll replace the content of the attachment inside the renderpass.

Likewise, the store operations tell Vulkan what you want it to do with the contents of the attachments when the renderpass ends. These can be set to one of the following values:

- `VK_ATTACHMENT_STORE_OP_STORE` indicates that you want Vulkan to keep the contents of the attachment for later use, which usually means that it should write them out into memory. This is usually the case for images you want to display to the user, read from later, or use as an attachment in another renderpass (with the `VK_ATTACHMENT_LOAD_OP_LOAD` load operation).

- `VK_ATTACHMENT_STORE_OP_DONT_CARE` indicates that you don't need the content after the renderpass has ended. This is normally used for intermediate storage or for the depth or stencil buffers.

If the attachment is a combined depth-stencil attachment, then the `stencilLoadOp` and `stencilStoreOp` fields tell Vulkan what to do with the stencil part of the attachment (the regular `loadOp` and `storeOp` fields specify what should happen to the depth part of the attachment), which can be different from the depth part.

The `initialLayout` and `finalLayout` fields tell Vulkan what layout to expect the image to be in when the renderpass begins and what layout to leave it in when the renderpass ends. Note that renderpass objects *do not* automatically move images into the initial layout. This is the layout that the image is expected to be in when the renderpass is used. The renderpass does, however, move the image to the final layout when it's done.

You can use barriers to explicitly move images from layout to layout, but where possible, it's best to try to move images from layout to layout inside renderpasses. This gives Vulkan the best opportunity to choose the right layout for each part of the renderpass and even perform any operations required to move images between layouts in parallel with other rendering. Advanced usage of these fields and renderpasses in general is covered in Chapter 13, "Multipass Rendering."

After you define all of the attachments that are going to be used in the renderpass, you need to define all of the subpasses. Each subpass references a number of attachments (from the array you passed in `pAttachments`) as inputs or outputs. Those descriptions are specified in an array of `VkSubpassDescription` structures, one for each subpass in the renderpass. The definition of `VkSubpassDescription` is

Click here to view code image

```
typedef struct VkSubpassDescription {
    VkSubpassDescriptionFlags       flags;
    VkPipelineBindPoint             pipelineBindPoint;
    uint32_t                        inputAttachmentCount;
    const VkAttachmentReference*    pInputAttachments;
    uint32_t                        colorAttachmentCount;
    const VkAttachmentReference*    pColorAttachments;
    const VkAttachmentReference*    pResolveAttachments;
    const VkAttachmentReference*    pDepthStencilAttachment;
    uint32_t                        preserveAttachmentCount;
    const uint32_t*                 pPreserveAttachments;
} VkSubpassDescription;
```

The `flags` field of `VkSubpassDescription` is reserved for future use and should be set to zero. Also, the current version of Vulkan supports renderpasses only for graphics, so `pipelineBindPoint` should be set to `VK_PIPELINE_BIND_POINT_GRAPHICS`. The remaining fields describe the attachments used by the subpass. Each subpass can have a number of input attachments, which are attachments from which it can read data; color attachments, which are attachments to which its outputs are written; and a depth-stencil attachment, which is used as a depth and stencil buffer. These attachments are specified in the `pInputAttachments`, `pColorAttachments`, and `pDepthStencilAttachment` fields, respectively. The numbers of input and color attachments are specified in `inputAttachmentCount` and `colorAttachmentCount`, respectively. There is only one depth-stencil attachment, so this parameter is not an array and has no associated count.

The maximum number of color attachments that a single subpass can render to can be determined by inspecting the `maxColorAttachments` member of the device's `VkPhysicalDeviceLimits` structure, which you can retrieve by calling **vkGetPhysicalDeviceProperties()**. `maxColorAttachments` is guaranteed to be at least 4, so if you never use more than this many color attachments, you don't need to query the limit. However, many implementations support a

higher limit than this, so you may be able to implement more advanced algorithms in fewer passes by writing to more outputs at once.

Each of these arguments is a pointer to either a single `VkAttachmentReference` structure or an array of them and forms a reference to one of the attachments described in `pAttachments`. The definition of `VkAttachmentReference` is

[**Click here to view code image**](#)

```
typedef struct VkAttachmentReference {
    uint32_t         attachment;
    VkImageLayout    layout;
} VkAttachmentReference;
```

Each attachment reference is a simple structure containing an index into the array of attachments in `attachment` and the image layout that the attachment is expected to be in at this subpass. In addition to the input and output attachment references, two further sets of references are provided to each subpass.

First, the resolve attachments, which are specified through `pResolveAttachments`, are the attachments to which multisample image data is resolved. These attachments correspond to the color attachments specified in `pColorAttachments`, and the number of resolve attachments is assumed to be the same, as specified in `colorAttachmentCount`.

If one of the elements of `pColorAttachments` is a multisample image, but only the final, resolved image is needed after the renderpass is complete, you can ask Vulkan to resolve the image for you as part of the renderpass, and possibly disacard the original multisample data. To do this, set the store operation for the multisample color attachment to `VK_ATTACHMENT_STORE_OP_DONT_CARE`, and set a corresponding single-sample attachment in the matching element of `pResolveAttachments`. The store operation for the resolve attachment should be set to `VK_ATTACHMENT_STORE_OP_STORE`, which will cause Vulkan to keep the single-sample data but throw out the original multisample data.

Second, if there are attachments that you want to live across a subpass but that are not directly referenced by the subpass, you should reference them in the `pPreserveAttachments` array. This reference will prevent Vulkan from making any optimizations that might disturb the contents of those attachments.

When there is more than one subpass in a renderpass, Vulkan can figure out which subpasses are dependent on one another by following the attachment references and looking for inputs and outputs that make subpasses dependent on one another. However, there are cases in which dependencies cannot easily be represented by a simple input-to-output relationship. This generally happens when a subpass writes directly to a resource such as an image or buffer and a subsequent subpass reads that data back. Vulkan cannot figure this out automatically, so you must provide such dependency information explicitly. This is done using the `pDependencies` member of `VkRenderPassCreateInfo`, which is a pointer to an array of `dependencyCount` `VkSubpassDependency` structures. The definition of `VkSubpassDependency` is

[**Click here to view code image**](#)

```
typedef struct VkSubpassDependency {
    uint32_t              srcSubpass;
    uint32_t              dstSubpass;
```

```
    VkPipelineStageFlags      srcStageMask;
    VkPipelineStageFlags      dstStageMask;
    VkAccessFlags             srcAccessMask;
    VkAccessFlags             dstAccessMask;
    VkDependencyFlags         dependencyFlags;
} VkSubpassDependency;
```

Each dependency is a reference from a source subpass (the producer of data) and a destination subpass (the consumer of that data), specified in `srcSubpass` and `dstSubpass`, respectively. Both are indices into the array of subpasses that make up the renderpass. The `srcStageMask` is a bitfield specifying which pipeline stage(s) of the source subpass produced the data. Likewise, `dstStageMask` is a bitfield specifying which stages of the destination subpass will consume the data.

The `srcAccessMask` and `dstAccessMask` fields are also bitfields. They specify how each of the source and destination subpasses access the data. For example, the source stage may perform image stores from its vertex shader or write to a color attachment through regular fragment shader outputs. Meanwhile, the destination subpass may read through an input attachment or an image load.

For the purpose of creating a simple renderpass with a single subpass, with a single output attachment and no external dependencies, the data structures are mostly empty. Listing 7.1 demonstrates how to set up a simple renderpass in this configuration.

Listing 7.1: Creating a Simple Renderpass

Click here to view code image

```
// This is our color attachment. It's an R8G8B8A8_UNORM single sample
image.
// We want to clear it at the start of the renderpass and save the
contents
// when we're done. It starts in UNDEFINED layout, which is a key to
// Vulkan that it's allowed to throw the old content away, and we want to
// leave it in COLOR_ATTACHMENT_OPTIMAL state when we're done.
static const VkAttachmentDescription attachments[] =
{
    {
        0,                                      // flags
        VK_FORMAT_R8G8B8A8_UNORM,               // format
        VK_SAMPLE_COUNT_1_BIT,                  // samples
        VK_ATTACHMENT_LOAD_OP_CLEAR,            // loadOp
        VK_ATTACHMENT_STORE_OP_STORE,           // storeOp
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,        // stencilLoadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,       // stencilStoreOp
        VK_IMAGE_LAYOUT_UNDEFINED,              // initialLayout
        VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL   // finalLayout
    }
};

// This is the single reference to our single attachment.
static const VkAttachmentReference attachmentReferences[] =
{
    {
```

```
        0,                                             // attachment
        VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL    // layout
    }
};

// There is one subpass in this renderpass, with only a reference to the
// single output attachment.
static const VkSubpassDescription subpasses[] =
{
    {
        0,                                             // flags
        VK_PIPELINE_BIND_POINT_GRAPHICS,            // pipelineBindPoint
        0,                                             // inputAttachmentCount
        nullptr,                                      // pInputAttachments
        1,                                             // colorAttachmentCount
        &attachmentReferences[0],                     // pColorAttachments
        nullptr,                                      // pResolveAttachments
        nullptr,                                      //
pDepthStencilAttachment
        0,                                             //
preserveAttachmentCount
        nullptr                                       // pPreserveAttachments
    }
};

// Finally, this is the information that Vulkan needs to create the
renderpass
// object.
static VkRenderPassCreateInfo renderpassCreateInfo =
{
    VK_STRUCTURE_TYPE_RENDER_PASS_CREATE_INFO,      // sType
    nullptr,                                         // pNext
    0,                                               // flags
    1,                                               // attachmentCount
    attachments[0],                                  // pAttachments
    1,                                               // subpassCount
    &subpasses[0],                                   // pSubpasses
    0,                                               // dependencyCount
    nullptr                                          // pDependencies
};

VkRenderPass renderpass = VK_NULL_HANDLE;
// The only code that actually executes is this single call, which creates
// the renderpass object.
vkCreateRenderPass(device,
                   &renderpassCreateInfo,
                   nullptr,
                   &renderpass);
```

In Listing 7.1, we set up a simple renderpass with a single color attachment of format
VK_FORMAT_R8G8B8A8_UNORM, no depth-stencil attachment, and no dependencies. It looks like
a lot of code, but that's because we need to specify full data structures even though we're not using
most of the fields. As your applications grow more complex, the amount of code you need to write

doesn't actually grow correspondingly. Further, because the structures are constant, the amount of code *executed* by Listing 7.1 is minimal.

We'll use the renderpass created in Listing 7.1 to create a graphics pipeline in the next section.

Of course, when we are done using the renderpass object, we should destroy it. To do this, call **vkDestroyRenderPass()**, the prototype of which is

```
void vkDestroyRenderPass (
    VkDevice                          device,
    VkRenderPass                      renderPass,
    const VkAllocationCallbacks*      pAllocator);
```

device is the device that created the renderpass, and renderPass is the handle to the renderpass object to destroy. If a host memory allocator was used to create the renderpass, pAllocator should point to a compatible allocator; otherwise, pAllocator should be nullptr.

The Framebuffer

The *framebuffer* is an object that represents the set of images that graphics pipelines will render into. These affect the last few stages in the pipeline: depth and stencil tests, blending, logic operations, multisampling, and so on. A framebuffer object is created by using a reference to a renderpass and can be used with any renderpass that has a similar arrangement of attachments.

To create a framebuffer object, call **vkCreateFramebuffer()**, the prototype of which is

```
VkResult vkCreateFramebuffer (
    VkDevice                          device,
    const VkFramebufferCreateInfo*    pCreateInfo,
    const VkAllocationCallbacks*      pAllocator,
    VkFramebuffer*                    pFramebuffer);
```

The device that will be used to create the framebuffer object is passed in device, and the remaining parameters describing the new framebuffer object are passed through a pointer to an instance of the VkFramebufferCreateInfo structure in pCreateInfo. The definition of VkFramebufferCreateInfo is

```
typedef struct VkFramebufferCreateInfo {
    VkStructureType           sType;
    const void*               pNext;
    VkFramebufferCreateFlags  flags;
    VkRenderPass              renderPass;
    uint32_t                  attachmentCount;
    const VkImageView*        pAttachments;
    uint32_t                  width;
    uint32_t                  height;
    uint32_t                  layers;
} VkFramebufferCreateInfo;
```

The `sType` field of `VkFramebufferCreateInfo` should be set to `VK_STRUCTURE_TYPE_FRAMEBUFFER_CREATE_INFO`, and `pNext` should be set to `nullptr`. The `flags` field is reserved and should be set to zero.

A handle to a renderpass object that is compatible with the framebuffer being created should be passed in `renderPass`. For the purposes of compatibility with framebuffer objects, two renderpasses are compatible if their attachment references are the same.

The set of images that is to be bound into the framebuffer object is passed through an array of `VkImageView` handles, a pointer to which is passed in `pAttachments`. The number of images in `pAttachments` is specified in `attachmentCount`. The passes comprising the renderpass make references to the image attachments, and those references are specified as indices into the array specified in `pAttachments`. If you know that a particular renderpass doesn't use some of the attachments, but you want the framebuffer to be compatible with several renderpass objects or to keep a consistent layout of images in your application, some of the image handles in `pAttachments` can be `VkNullHandle`.

Although each of the images in the framebuffer has a native width, height, and (in the case of array images) layer count, you must still specify the dimensions of the framebuffer. These dimensions are passed in the `width`, `height`, and `layers` fields of the `VkFramebufferCreateInfo` structure. Rendering to regions of the framebuffer that are outside some of the images results in no rendering to those parts of the attachment images that are outside the image while continuing to render to those parts of the images that are.

The maximum supported size of a framebuffer is device-dependent. To determine the supported dimensions of the framebuffer, check the `maxFramebufferWidth`, `maxFramebufferHeight`, and `maxFramebufferLayers` fields of the device's `VkPhysicalDeviceLimits` structure. These provide the maximum supported width, height, and layer count for framebuffers, respectively. The supported width and height are guaranteed to be at least 4,096 pixels, and the number of supported layers is guaranteed to be at least 256. However, most desktop-class hardware will support limits of 16,384 pixels in width and height and 2,048 layers.

It's also possible to create a framebuffer with *no* attachments at all. This is known as an *attachmentless framebuffer*. In this case, the framebuffer's dimensions are solely defined by the `width`, `height`, and `layers` fields. This type of framebuffer is typically used with fragment shaders that have other side effects, such as performing image stores, or with occlusion queries, which can measure other aspects of rendering but don't necessarily require that the result of rendering be stored anywhere.

If **vkCreateFramebuffer()** is successful, it will write the new `VkFramebuffer` handle into the variable pointed to by `pFramebuffer`. If it requires any host memory, it will use the allocator pointed to by `pAllocator` to allocate it. If `pAllocator` is not `nullptr`, then a compatible allocator should be used when the framebuffer is destroyed.

As you will see in [Chapter 8](#), "[Drawing](#)," we will use the framebuffer object in conjunction with a renderpass in order to draw into the images attached to the framebuffer. When you are done using a framebuffer, you should destroy it by calling **vkDestroyFramebuffer()**, the prototype of which is

```
void vkDestroyFramebuffer (
    VkDevice                            device,
    VkFramebuffer                       framebuffer,
    const VkAllocationCallbacks*        pAllocator);
```

`device` is a handle to the device that created the framebuffer object, and `framebuffer` is a handle to the framebuffer object being destroyed. If a host memory allocator was used to allocate the framebuffer, a compatible allocator should be passed through the `pAllocator` object.

Destroying a framebuffer object does not affect any of the images attached to the framebuffer. Images can be attached to multiple framebuffers at the same time and can be used in multiple ways at the same time as being attached to a framebuffer. However, even if the images are not destroyed, the framebuffer should not be used—including any access in command buffers by the device. You should ensure that any command buffers referencing the framebuffer have completed execution if they have been submitted or have not been submitted after the framebuffer object is destroyed.

Creating a Simple Graphics Pipeline

Creating a graphics pipeline is achieved using a method similar to the one for creating a compute pipeline, as described in Chapter 6, "Shaders and Pipelines." However, as you have seen, the graphics pipeline includes many shading stages and fixed-function processing blocks, so the description of a graphics pipeline is correspondingly that much more complex. Graphics pipelines are created by calling **vkCreateGraphicsPipelines()**, the prototype of which is

```
VkResult vkCreateGraphicsPipelines (
    VkDevice                                device,
    VkPipelineCache                         pipelineCache,
    uint32_t                                createInfoCount,
    const VkGraphicsPipelineCreateInfo*     pCreateInfos,
    const VkAllocationCallbacks*            pAllocator,
    VkPipeline*                             pPipelines);
```

As you can see, the prototype for **vkCreateGraphicsPipelines()** is similar to **vkCreateComputePipelines()**. It takes a device (`device`), a handle to a pipeline cache (`pipelineCache`), and an array of `createInfo` structures along with the count of the number of structures in the array (`pCreateInfos` and `createInfoCount`, respectively). This is where the real guts of the function are. `VkGraphicsPipelineCreateInfo` is a large, complex structure, and it contains pointers to several other structures along with handles to other objects that you need to have created. Take a deep breath: The definition of `VkGraphicsPipelineCreateInfo` is

```
typedef struct VkGraphicsPipelineCreateInfo {
    VkStructureType                             sType;
    const void*                                 pNext;
    VkPipelineCreateFlags                       flags;
    uint32_t                                    stageCount;
    const VkPipelineShaderStageCreateInfo*      pStages;
```

```
    const VkPipelineVertexInputStateCreateInfo*        pVertexInputState;
    const
VkPipelineInputAssemblyStateCreateInfo*        pInputAssemblyState;
    const VkPipelineTessellationStateCreateInfo*        pTessellationState;
    const VkPipelineViewportStateCreateInfo*        pViewportState;
    const
VkPipelineRasterizationStateCreateInfo*        pRasterizationState;
    const VkPipelineMultisampleStateCreateInfo*        pMultisampleState;
    const VkPipelineDepthStencilStateCreateInfo*        pDepthStencilState;
    const VkPipelineColorBlendStateCreateInfo*        pColorBlendState;
    const VkPipelineDynamicStateCreateInfo*        pDynamicState;
    VkPipelineLayout                                layout;
    VkRenderPass                                    renderPass;
    uint32_t                                        subpass;
    VkPipeline                                      basePipelineHandle;
    int32_t                                         basePipelineIndex;
} VkGraphicsPipelineCreateInfo;
```

As you were warned, `VkGraphicsPipelineCreateInfo` is a large structure with many substructures referenced by pointers. However, it's easy enough to break down into blocks, and many of the additional creation info is optional and can be left as `nullptr`. As with all other Vulkan creation info structures, `VkGraphicsPipelineCreateInfo` starts with an `sType` field and a `pNext` field. The `sType` for `VkGraphicsPipelineCreateInfo` is `VK_GRAPHICS_PIPELINE_CREATE_INFO`, and `pNext` can be left as `nullptr` unless extensions are in use.

The `flags` field contains information about how the pipeline will be used. Three flags are defined in the current version of Vulkan, and their meanings are as follows:

- `VK_PIPELINE_CREATE_DISABLE_OPTIMIZATION_BIT` tells Vulkan that this pipeline is not going to be used in performance-critical applications and that you would prefer to receive a ready-to-go pipeline object quickly rather than have Vulkan spend a lot of time optimizing the pipeline. You might use this for things like simple shaders for displaying splash screens or user interface elements that you want to display quickly.

- `VK_PIPELINE_CREATE_ALLOW_DERIVATIVES_BIT` and `VK_PIPELINE_CREATE_DERIVATIVE_BIT` are used with *derivative* pipelines. This is a feature whereby you can group similar pipelines and tell Vulkan that you'll switch rapidly among them. The `VK_PIPELINE_CREATE_ALLOW_DERIVATIVES_BIT` flag tells Vulkan that you will want to create derivatives of the new pipeline, and `VK_PIPELINE_CREATE_DERIVATIVE_BIT` tells Vulkan that this pipeline *is* a pipeline.

Graphics Shader Stages

The next two fields in the `VkGraphicsPipelineCreateInfo` structure, `stageCount` and `pStages`, are where you pass your shaders into the pipeline. `pStages` is a pointer to an array of `stageCount VkPipelineShaderStageCreateInfo` structures, each describing one of the shading stages. These are the same structures that you saw in the definition of `VkComputePipelineCreateInfo`, except now you have an array of them. The definition of `VkPipelineShaderStageCreateInfo` is

```
typedef struct VkPipelineShaderStageCreateInfo {
    VkStructureType                     sType;
    const void*                         pNext;
    VkPipelineShaderStageCreateFlags    flags;
    VkShaderStageFlagBits               stage;
    VkShaderModule                      module;
    const char*                         pName;
    const VkSpecializationInfo*         pSpecializationInfo;
} VkPipelineShaderStageCreateInfo;
```

All graphics pipelines must have at least a vertex shader, and the vertex shader is always the first shading stage in the pipeline. Therefore, the `pStages` of `VkGraphicsPipelineCreateInfo` should point to a `VkPipelineShaderStageCreateInfo` describing a vertex shader. The parameters in the `VkPipelineShaderStageCreateInfo` structure have the same meaning as they did when we created a compute pipeline in [Chapter 6](#), "[Shaders and Pipelines](#)." `module` should be a shader module that contains at least one vertex shader, and `pName` should be the name of a vertex shader entry point in that module.

Because in our simple pipeline we're not using most of the stages of the Vulkan graphics pipeline, we can leave most of the other fields of the `VkGraphicsPipelineCreateInfo` structure as their defaults or as `nullptr` for the pointers. The `layout` field is the same as the `layout` field in the `VkComputePipelineCreateInfo` structure and specifies the pipeline layout used for resources by this pipeline.

We can set the `renderPass` member of our structure to the handle of the renderpass object we created earlier in [Listing 7.1](#). There's only one subpass in this renderpass, so we can set `subpass` to zero.

[Listing 7.2](#) shows a minimal example of creating a graphics pipeline containing only a vertex shader. It looks long, but most of it is setting up default values in structures that are not actually used by the pipeline. These structures will be explained in the following few paragraphs.

Listing 7.2: Creating a Simple Graphics Pipeline

```
VkPipelineShaderStageCreateInfo shaderStageCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_SHADER_STAGE_CREATE_INFO,        // sType
    nullptr,                                // pNext
    0,                                      // flags
    VK_SHADER_STAGE_VERTEX_BIT,             // stage
    module,                                 // module
    "main",                                 // pName
    nullptr                                 // pSpecializationInfo
};

static const
VkPipelineVertexInputStateCreateInfo vertexInputStateCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_VERTEX_INPUT_STATE_CREATE_INFO, // sType
```

```
    nullptr,                                    // pNext
    0,                                          // flags
    0,                                          //
vertexBindingDescriptionCount
    nullptr,                                    // pVertexBindingDescriptions
    0,                                          //
vertexAttributeDescriptionCount
    nullptr                                     // pVertexAttributeDescriptions
};

static const
VkPipelineInputAssemblyStateCreateInfo inputAssemblyStateCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_INPUT_ASSEMBLY_STATE_CREATE_INFO,// sType
    nullptr,                                    // pNext
    0,                                          // flags
    VK_PRIMITIVE_TOPOLOGY_POINT_LIST,           // topology
    VK_FALSE                                    // primitiveRestartEnable
};

static const
VkViewport dummyViewport =
{
    0.0f, 0.0f,                                 // x, y
    1.0f, 1.0f,                                 // width, height
    0.1f, 1000.0f                               // minDepth, maxDepth
};

static const
VkRect2D dummyScissor =
{
    { 0, 0 },                                   // offset
    { 1, 1 }                                    // extent
};

static const
VkPipelineViewportStateCreateInfo viewportStateCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_VIEWPORT_STATE_CREATE_INFO,      // sType
    nullptr,                                    // pNext
    0,                                          // flags
    1,                                          // viewportCount
    &dummyViewport,                             // pViewports
    1,                                          // scissorCount
    &dummyScissor                               // pScissors
};

static const
VkPipelineRasterizationStateCreateInfo rasterizationStateCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_RASTERIZATION_STATE_CREATE_INFO,  // sType
    nullptr,                                    // pNext
    0,                                          // flags
```

```
        VK_FALSE,                                // depthClampEnable
        VK_TRUE,                                 // rasterizerDiscardEnable
        VK_POLYGON_MODE_FILL,                    // polygonMode
        VK_CULL_MODE_NONE,                       // cullMode
        VK_FRONT_FACE_COUNTER_CLOCKWISE,         // frontFace
        VK_FALSE,                                // depthBiasEnable
        0.0f,                                    // depthBiasConstantFactor
        0.0f,                                    // depthBiasClamp
        0.0f,                                    // depthBiasSlopeFactor
        0.0f                                     // lineWidth
    };

    static const
    VkGraphicsPipelineCreateInfo graphicsPipelineCreateInfo =
    {
        VK_STRUCTURE_TYPE_GRAPHICS_PIPELINE_CREATE_INFO,    // sType
        nullptr,                                 // pNext
        0,                                       // flags
        1,                                       // stageCount
        &shaderStageCreateInfo,                  // pStages
        &vertexInputStateCreateInfo,             // pVertexInputState
        &inputAssemblyStateCreateInfo,           // pInputAssemblyState
        nullptr,                                 // pTessellationState
        &viewportStateCreateInfo,                // pViewportState
        &rasterizationStateCreateInfo,           // pRasterizationState
        nullptr,                                 // pMultisampleState
        nullptr,                                 // pDepthStencilState
        nullptr,                                 // pColorBlendState
        nullptr,                                 // pDynamicState
        VK_NULL_HANDLE,                          // layout
        renderpass,                              // renderPass
        0,                                       // subpass
        VK_NULL_HANDLE,                          // basePipelineHandle
        0,                                       // basePipelineIndex
    };

    result = vkCreateGraphicsPipelines(device,
                                       VK_NULL_HANDLE,
                                       1,
                                       &graphicsPipelineCreateInfo,
                                       nullptr,
                                       &pipeline);
```

Of course, most of the time, you won't be using a graphics pipeline containing only a vertex shader. Up to five shader stages make up the graphics pipeline, as introduced earlier in this chapter. These stages include the following:

- The vertex shader, specified as `VK_SHADER_STAGE_VERTEX_BIT`, processes one vertex at a time and passes it to the next logical stage in the pipeline.

- The tessellation control shader, specified as `VK_SHADER_STAGE_TESSELLATION_CONTROL_BIT`, processes one control point at a time but has access to all of the data that makes up the patch. It can be considered to be a patch shader, and it produces the tessellation factors and per-patch data associated with the patch.

- The tessellation evaluation shader, specified using `VK_SHADER_STAGE_TESSELLATION_EVALUATION_BIT`, processes one tessellated vertex at a time. In many applications, it evaluates the patch function at each point—hence, the name. It also has access to the full patch data produced by the tessellation control shader.

- The geometry shader, specified using `VK_SHADER_STAGE_GEOMETRY_BIT`, executes once for each primitive that passes through the pipeline: points, lines, or triangles. It can produce new primitives or throw them away rather than passing them on. It can also change the type of a primitive as it passes by.

- The fragment shader, specified using `VK_SHADER_STAGE_FRAGMENT_BIT`, executes once per fragment, after rasterization. It is primarily responsible for computing the final color of each pixel.

Most straightforward rendering will include at least a vertex and a fragment shader. Each shader stage can consume data from the previous stage or pass data on to the next, forming a pipeline. In some cases, the inputs to a shader are supplied by fixed-function blocks, and sometimes the outputs from a shader are consumed by fixed-function blocks. Regardless of the source or destination of data, the means of declaring the inputs and outputs to shaders are the same.

To declare an input to a shader in SPIR-V, a variable must be decorated as `Input` when it is declared. Likewise, to create an output from the shader, decorate a variable as `Output` when it is declared. Unlike in GLSL, special-purpose inputs and outputs do not have predefined names in SPIR-V. Rather, they are decorated with their purpose. Then you write shaders in GLSL and compile them to SPIR-V using a GLSL compiler. The compiler will recognize access to built-in variables and translate them into appropriately declared and decorated input and output variables in the resulting SPIR-V shader.

Vertex Input State

To render real geometry, you need to feed data into the front of the Vulkan pipeline. You can use the vertex and instance indices that are provided by SPIR-V to programmatically generate geometry or explicitly fetch geometry data from a buffer. Alternatively, you can describe the layout of geometric data in memory and Vulkan can fetch it for you, supplying it directly to your shader.

To do this, we use the `pVertexInputState` member of `VkGraphicsPipelineCreateInfo`, which is a pointer to an instance of the `VkPipelineVertexInputStateCreateInfo` structure, the definition of which is

[Click here to view code image](#)

```
typedef struct VkPipelineVertexInputStateCreateInfo {
    VkStructureType                           sType;
    const void*                               pNext;
    VkPipelineVertexInputStateCreateFlags     flags;
    uint32_t                                  vertexBindingDescriptionCount;
    const VkVertexInputBindingDescription*    pVertexBindingDescriptions;
    uint32_t                                  vertexAttributeDescriptionCount;
    const
VkVertexInputAttributeDescription*  pVertexAttributeDescriptions;
} VkPipelineVertexInputStateCreateInfo;
```

The `VkPipelineVertexInputStateCreateInfo` structure begins with the familiar `sType` and `pNext` fields, which should be set to `VK_STRUCTURE_TYPE_PIPELINE_VERTEX_INPUT_STATE_CREATE_INFO` and `nullptr`, respectively. The `flags` field of `VkPipelineVertexInputStateCreateInfo` is reserved for future use and should be set to zero.

Vertex input state is divided into a set of vertex bindings to which you can bind buffers containing data and a set of vertex attributes that describe how vertex data is laid out in those buffers. Buffers bound to the vertex buffer binding points are sometimes referred to as vertex buffers. It should be noted, though, that there's not really any such thing as a "vertex buffer" in the sense that any buffer can store vertex data, and a single buffer can store vertex data and other kinds of data as well. The only requirement for a buffer to be used as storage for vertex data is that it must have been created with the `VK_BUFFER_USAGE_VERTEX_BUFFER_BIT` set.

`vertexBindingDescriptionCount` is the number of vertex bindings used by the pipeline, and `pVertexBindingDescriptions` is a pointer to an array of that many `VkVertexInputBindingDescription` structures, each describing one of the bindings. The definition of `VkVertexInputBindingDescription` is

Click here to view code image

```
typedef struct VkVertexInputBindingDescription {
    uint32_t              binding;
    uint32_t              stride;
    VkVertexInputRate     inputRate;
} VkVertexInputBindingDescription;
```

The `binding` field is the index of the binding described by this structure. Each pipeline can address a number of vertex buffer bindings, and their indices do not need to be contiguous. It is not necessary to describe every binding in a given pipeline so long as every binding that is used by that pipeline is described.

The last binding index addressed by the array of `VkVertexInputBindingDescription` structures must be less than the maximum number of bindings supported by the device. This limit is guaranteed to be at least 16, but for some devices, it could be higher. If you don't need more than 16 bindings, then there's no reason to check the limit. However, you can determine the highest binding index by checking the `maxVertexInputBindings` member of the device's `VkPhysicalDeviceLimits` structure, which is returned from a call to **vkGetPhysicalDeviceProperties()**.

Each binding can be seen as an array of structures located in a buffer object. The stride of the array—that is, the distance between the start of each structure, measured in bytes—is specified in `stride`. If the vertex data is specified as an array of structures, the `stride` parameter essentially contains the size of the structure, even if the shader doesn't use every member of it. The maximum value of `stride` for any particular binding is implementation-dependent but is guaranteed to be at least 2,048 bytes. If you want to use vertex data with a greater stride than this, you need to query the supported stride to make sure that the device can handle it.

To determine the maximum supported stride, check the `maxVertexInputBindingStride` field of the device's `VkPhysicalDeviceLimits` structure.

Further, Vulkan can iterate through the array either as a function of the vertex index or as a function of the instance index when instancing is in use. This is specified in the `inputRate` field, which should be either `VK_VERTEX_INPUT_RATE_VERTEX` or `VK_VERTEX_INPUT_RATE_INSTANCE`.

Each vertex attribute is essentially a member of one of the structures stored in the vertex buffer. Each vertex attribute sourced from the vertex buffer shares the step rate and stride of the array but has its own data type and offset within that structure. This is described using the `VkVertexInputAttributeDescription` structure. The address of an array of these structures is passed in the `pVertexAttributeDescriptions` field of `VkPipelineVertexInputStateCreateInfo`, and the number of elements in the array (which is the number of vertex attributes) is passed in

`vertexAttributeDescriptionCount`. The definition of `VkVertexInputAttributeDescription` is

Click here to view code image

```
typedef struct VkVertexInputAttributeDescription {
    uint32_t    location;
    uint32_t    binding;
    VkFormat    format;
    uint32_t    offset;
} VkVertexInputAttributeDescription;
```

Each attribute has a *location* that is used to refer to it in the vertex shader. Again, the vertex attribute locations don't need to be contiguous, and it's not necessary to describe every single vertex attribute location so long as all the attributes used by the pipeline are described. The attribute's location is specified though the `location` member of `VkVertexInputAttributeDescription`.

The binding to which the buffer is bound, and from which this attribute sources its data, is specified in `binding` and should match one of the bindings specified in the array of `VkVertexInputBindingDescription` structures described earlier. The format of the vertex data is specified in `format`, and the offset within each structure is specified in `offset`.

Just as the total size of the structure has an upper limit, there is an upper limit to the offset from the start of the structure for each attribute: the upper bound on `offset`. This is guaranteed to be at least 2,047 bytes, which is high enough to place a single byte right at the end of a structure of the maximum guaranteed size (2,048 bytes). If you need to use bigger structures than this, you need to check the capability of the device to handle it. The `maxVertexInputAttributeOffset` field of the device's `VkPhysicalDeviceLimits` structure contains the maximum value that can be used in `offset`. You can retrieve this structure by calling **vkGetPhysicalDeviceProperties()**.

Listing 7.3 shows how to create a structure in C++ and describe it using the `VkVertexInputBindingDescription` and `VkVertexInputAttributeDescription` such that you can use it to hand vertex data to Vulkan.

Listing 7.3: Describing Vertex Input Data

[**Click here to view code image**](#)

```
typedef struct vertex_t
{
    vmath::vec4 position;
    vmath::vec3 normal;
    vmath::vec2 texcoord;
} vertex;

static const
VkVertexInputBindingDescription vertexInputBindings[] =
{
    { 0, sizeof(vertex), VK_VERTEX_INPUT_RATE_VERTEX }    // Buffer
};

static const
VkVertexInputAttributeDescription vertexAttributes[] =
{
    { 0, 0, VK_FORMAT_R32G32B32A32_SFLOAT, 0 },                        //
Position
    { 1, 0, VK_FORMAT_R32G32B32_SFLOAT, offsetof(vertex, normal) }, //
Normal
    { 2, 0, VK_FORMAT_R32G32_SFLOAT, offsetof(vertex, texcoord) }    // Tex
Coord
};

static const
VkPipelineVertexInputStateCreateInfo vertexInputStateCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_VERTEX_INPUT_STATE_CREATE_INFO,  // sType
    nullptr,                                          // pNext
    0,                                                // flags
    vkcore::utils::arraysize(vertexInputBindings), //
vertexBindingDescriptionCount
    vertexInputBindings,                              //
pVertexBindingDescriptions
    vkcore::utils::arraysize(vertexAttributes),     //
vertexAttributeDescriptionCount
    vertexAttributes                                  //
pVertexAttributeDescriptions
};
```

The maximum number of input attributes that can be used in a single vertex shader is
implementation-dependent but is guaranteed to be at least 16. This is the upper limit on the number
of `VkVertexInputAttributeDescription` structures in the
`pVertexInputAttributeDescriptions` array. Some implementations may support more
inputs than this. To determine the maximum number of vertex shader inputs that you can use, check
the `maxVertexInputAttributes` field of the device's `VkPhysicalDeviceLimits`
structure.

Vertex data is read from the vertex buffers that you bind to the command buffer and then passed to the vertex shader. For the vertex shader to be able to interpret that vertex data, it must declare inputs corresponding to the vertex attributes you have defined. To do this, create a variable in your SPIR-V vertex shader with the `Input` storage class. In a GLSL shader, this can be expressed using an **in** variable.

Each input must have an assigned *location*. This is specified in GLSL using the `location` layout qualifier, which is then translated into a SPIR-V `Location` decoration applied to the input. Listing 7.4 shows a fragment of a GLSL vertex shader that declares a number of inputs. The resulting SPIR-V produced by `glslangvalidator` is shown in Listing 7.5.

The shader shown in Listing 7.5 is incomplete, as it has been edited to make the declared inputs clearer.

Listing 7.4: Declaring Inputs to a Vertex Shader (GLSL)

Click here to view code image

```
#version 450 core

layout (location = 0) in vec3 i_position;
layout (location = 1) in vec2 i_uv;

void main(void)
{
    gl_Position = vec4(i_position, 1.0f);
}
```

Listing 7.5: Declaring Inputs to a Vertex Shader (SPIR-V)

Click here to view code image

```
; SPIR-V
; Version: 1.0
; Generator: Khronos Glslang Reference Front End; 1
; Bound: 30
; Schema: 0
        OpCapability Shader
   %1 = OpExtInstImport "GLSL.std.450"
        OpMemoryModel Logical GLSL450
        OpEntryPoint Vertex %4 "main" %13 %18 %29
        OpSource GLSL 450
        OpName %18 "i_position"       ;; Name of i_position
        OpName %29 "i_uv"             ;; Name of i_uv
        OpDecorate %18 Location 0     ;; Location of i_position
        OpDecorate %29 Location 1     ;; Location of i_uv
...
   %6 = OpTypeFloat 32                ;; %6 is 32-bit floating-point type
  %16 = OpTypeVector %6 3             ;; %16 is a vector of 3 32-bit floats
  (vec3)
  %17 = OpTypePointer Input %16
  %18 = OpVariable %17 Input          ;; %18 is i _position - input pointer to
  vec3
```

```
%27 = OpTypeVector %6 2         ;; %27 is a vector of 2 32-bit floats
%28 = OpTypePointer Input %27
%29 = OpVariable %28 Input      ;; %29 is i _uv - input pointer to vec2
...
```

It is also possible to declare a vertex shader input that corresponds only to certain components of the vertex attribute. Again, the attribute is the data supplied by your application through vertex buffers, and the vertex shader input is the variable in the vertex shader corresponding to the data read by Vulkan on your behalf.

To create a vertex shader input that corresponds to a subset of the components of an input vector, use the GLSL `component` layout qualifier, which is translated into a SPIR-V `Component` decoration applied to the vertex shader input. Each vertex shader input can begin at a component numbered 0 through 3, corresponding to the `x`, `y`, `z`, and `w` channels of the source data. Each input consumes as many consecutive components as it requires. That is, a scalar consumes a single component, a **vec2** consumes 2, a **vec3** consumes 3, and so on.

Vertex shaders can also declare matrices as inputs. In GLSL, this is as simple as using the **in** storage qualifier on a variable in the vertex shader. In SPIR-V, a matrix is effectively declared as a special type of vector consisting of vector types. The matrix is considered to be column primary by default. Therefore, each set of contiguous data fills a single column of the matrix.

Input Assembly

The input assembly phase of the graphics pipeline takes the vertex data and groups it into primitives ready for processing by the rest of the pipeline. It is described by an instance of the `VkPipelineInputAssemblyStateCreateInfo` structure that is passed through the `pInputAssemblyState` member of the `VkGraphicsPipelineCreateInfo` structure. The definition of `VkPipelineInputAssemblyStateCreateInfo` is

[Click here to view code image](#)

```
typedef struct VkPipelineInputAssemblyStateCreateInfo {
    VkStructureType                            sType;
    const void*                                pNext;
    VkPipelineInputAssemblyStateCreateFlags    flags;
    VkPrimitiveTopology                        topology;
    VkBool32                                   primitiveRestartEnable;
} VkPipelineInputAssemblyStateCreateInfo;
```

The `sType` field should be set to `VK_STRUCTURE_TYPE_PIPELINE_VERTEX_INPUT_STATE_CREATE_INFO`, and `pNext` should be set to `nullptr`. The `flags` field is reserved for future use and should be set to zero.

The *primitive topology* is specified in `topology`, which should be one of the primitive topologies supported by Vulkan. These are members of the `VkPrimitiveTopology` enumeration. The simplest members of this enumeration are the *list* topologies, which are

- `VK_PRIMITIVE_TOPOLOGY_POINT_LIST`: Each vertex is used to construct an independent point.

- `VK_PRIMITIVE_TOPOLOGY_LINE_LIST`: Vertices are grouped into pairs, each pair forming a line segment from the first to the second vertex.

- `VK_PRIMITIVE_TOPOLOGY_TRIANGLE_LIST`: Vertices are grouped into triplets forming triangles.

Next are the strip and fan primitives. These are groupings of vertices into primitives (lines or triangles) in which each line or triangle shares one or two vertices with the previous one. The strip and fan primitives are as follows:

- `VK_PRIMITIVE_TOPOLOGY_LINE_STRIP`: The first two vertices in a draw form a single line segment. Each new vertex after them forms a new line segment from the last processed vertex. The result is a connected sequence of lines.

- `VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP`: The first three vertices in a draw form a single triangle. Each subsequent vertex forms a new triangle along with the last *two* vertices. The result is a connected row of triangles, each sharing an edge with the last.

- `VK_PRIMITIVE_TOPOLOGY_TRIANGLE_FAN`: The first three vertices in a draw form a single triangle. Each subsequent vertex forms a new triangle along with the last vertex and the *first* vertex in the draw.

Strip and fan topologies are not complex but can be difficult to visualize if you are not familiar with them. Figure 7.2 shows these topologies laid out graphically.
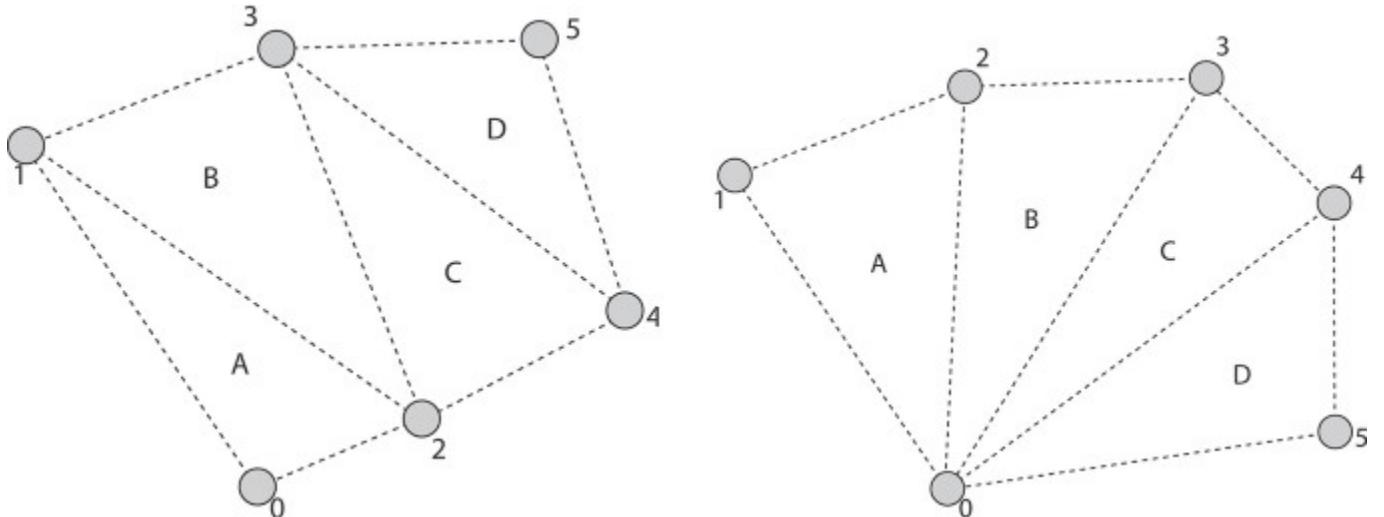
Figure 7.2: Strip (Left) and Fan (Right) Topologies

Next are the *adjacency primitives*, which are typically used only when a geometry shader is enabled and are able to convey additional information about primitives next to them in an original mesh. The adjacency primitive topologies are

- `VK_PRIMITIVE_TOPOLOGY_LINE_LIST_WITH_ADJACENCY`: Every four vertices in the draw form a single primitive, with the center two vertices forming a line and the first and last vertex in each group of four being presented to the geometry shader, when present.

- `VK_PRIMITIVE_TOPOLOGY_LINE_STRIP_WITH_ADJACENCY`: The first four vertices in the draw form a single primitive, with the center two vertices forming a line segment and the first and last being presented to the geometry shader as adjacency information. Each subsequent vertex essentially slides this window of four vertices along by one, forming a new line segment and presenting the new vertex as adjacency information.

- `VK_PRIMITIVE_TOPOLOGY_TRIANGLE_LIST_WITH_ADJACENCY`: Similar to lines with adjacency, each group of six vertices is formed into a single primitive, with the first, third, and

fifth in each group constructing a triangle and the second, fourth, and sixth being presented to the geometry shader as adjacency information.

- `VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP_WITH_ADJACENCY`: This is perhaps the most confusing primitive topology and certainly needs a diagram to visualize. Essentially, the strip begins with the first six vertices forming a triangle with adjacency information as in the list case. For every *two* new vertices, a new triangle is formed, with the odd-numbered vertices forming the triangle and the even-numbered vertices providing adjacency information.

Again, adjacency topologies can be quite difficult to visualize—especially the `VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP_WITH_ADJACENCY` topology. Figure 7.3 illustrates the layout of vertices within the `VK_PRIMITIVE_TOPOLOGY_TRIANGLE_LIST_WITH_ADJACENCY` topology. In the figure you can see that there are two triangles formed from a total of 12 vertices. The vertices wrap around the outside of each triangle, with the odd-numbered vertices forming the center triangles (A and B) and the even-numbered vertices forming virtual triangles that are not rendered, but carry adjacency information. This concept carries on to the triangle strip
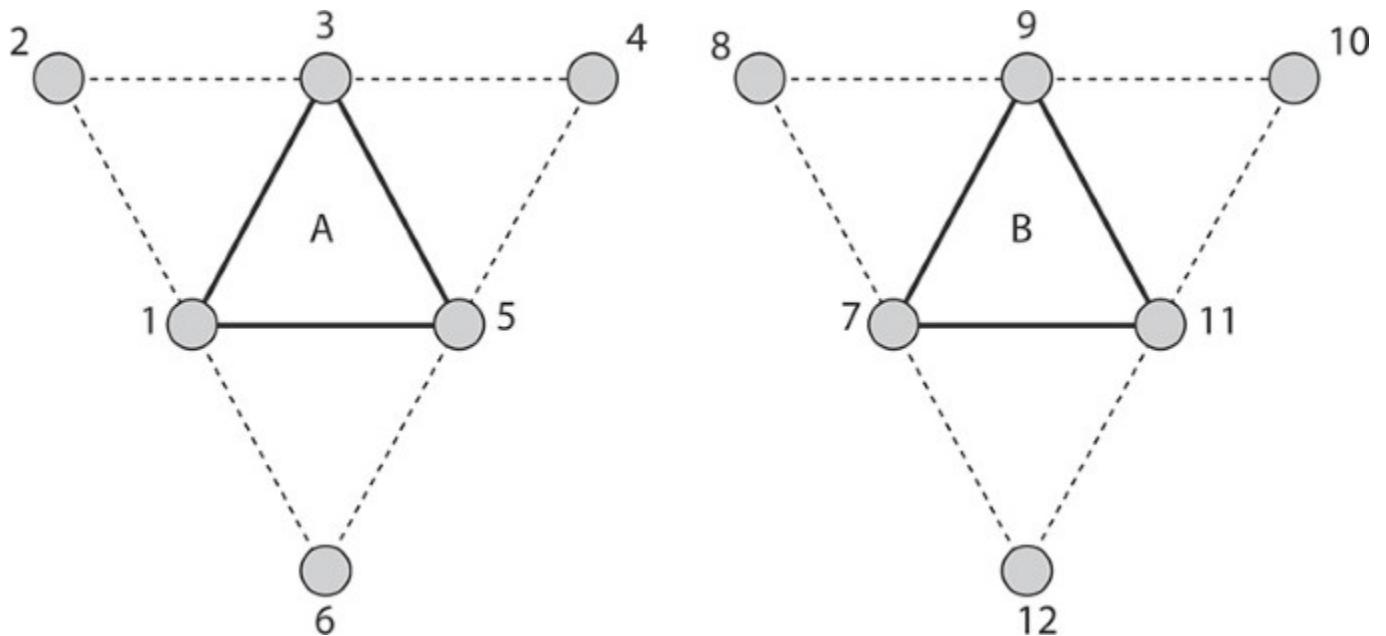
Figure 7.3: Triangles with Adjacency Topology

primitive. Figure 7.4 shows how it is applied to the `VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP_WITH_ADJACENCY`.
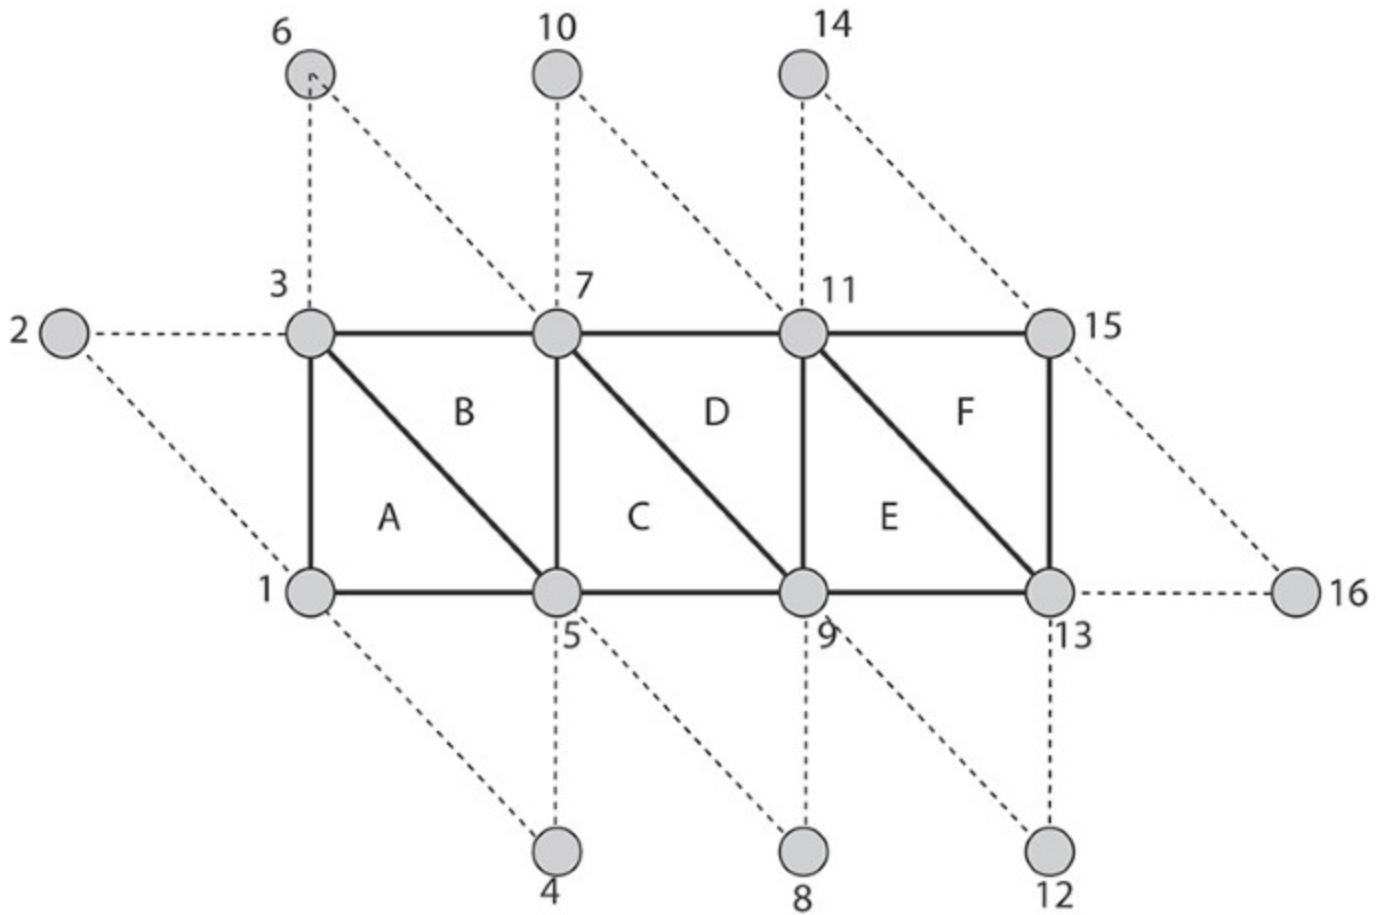
Figure 7.4: Triangle Strip with Adjacency Topology

Adjacency topologies are typically used only when a geometry shader is present, as the geometry shader is the only stage that really sees the adjacency vertices. However, it's possible to use adjacency primitives without a geometry shader; the adjacency vertices will simply be discarded.

The last primitive topology is `VK_PRIMITIVE_TOPOLOGY_PATCH_LIST`. This topology is used when tessellation is enabled, which requires additional information to be passed to pipeline construction.

The last field in `VkPipelineInputAssemblyStateCreateInfo` is `primitiveRestartEnable`. This is a flag that is used to allow strip and fan primitive topologies to be cut and restarted. Without this, each strip or fan would need to be a separate draw. When you use restarts, many strips or fans can be combined into a single draw. Restarts take effect only when *indexed* draws are used because the point at which to restart the strip is marked using a special, reserved value in the index buffer. This is covered in more detail in Chapter 8, "Drawing."

Tessellation State

Tessellation is the process of breaking a large, complex primitive into a large number of smaller primitives approximating the original. Vulkan can tessellate a patch primitive into many smaller point, line, or triangle primitives prior to geometry shading and rasterization. Most of the state related to tessellation is configured using the tessellation control shader and tessellation evaluation shader. However, because these shading stages don't run until vertex data has already been fetched and processed by the vertex shader, some information is needed up front to configure this stage of the pipeline.

This information is provided through an instance of the `VkPipelineTessellationStateCreateInfo` structure, pointed to by the `pTessellationState` member of `VkGraphicsPipelineCreateInfo`. The definition of `VkPipelineTessellationStateCreateInfo` is

[**Click here to view code image**](#)

```
typedef struct VkPipelineTessellationStateCreateInfo {
    VkStructureType                           sType;
    const void*                               pNext;
    VkPipelineTessellationStateCreateFlags    flags;
    uint32_t                                  patchControlPoints;
} VkPipelineTessellationStateCreateInfo;
```

When the `topology` field of the `VkPipelineInputAssemblyStateCreateInfo` structure is set to `VK_PRIMITIVE_TOPOLOGY_PATCH_LIST`, `pTessellationState` must be a pointer to a `VkPipelineTessellationStateCreateInfo` structure; otherwise, `pTessellationState` can be `nullptr`.

`sType` for `VkPipelineTessellationStateCreateInfo` is `VK_STRUCTURE_TYPE_PIPELINE_TESSELLATION_STATE_CREATE_INFO`. `pNext` should be set to `nullptr`, and `flags` is reserved for use in future versions of Vulkan and should be set to zero. The only field of significance in `VkPipelineTessellationStateCreateInfo` is `patchControlPoints`, which sets the number of control points that will be grouped into a single primitive (patch). Tessellation is a somewhat advanced topic and will be covered in more detail in [Chapter 9](#), "[Geometry Processing](#)."

Viewport State

Viewport transformation is the final coordinate transform in the Vulkan pipeline before rasterization occurs. It transforms vertices from normalized device coordinates into window coordinates. Multiple viewports can be in use simultaneously. The state of these viewports, including the number of active viewports and their parameters, is set through an instance of the `VkPipelineViewportStateCreateInfo` structure, the address of which is passed through the `pViewportState` member of `VkGraphicsPipelineCreateInfo`. The definition of `VkPipelineViewportStateCreateInfo` is

[**Click here to view code image**](#)

```
typedef struct VkPipelineViewportStateCreateInfo {
    VkStructureType                      sType;
    const void*                          pNext;
```

```
    VkPipelineViewportStateCreateFlags      flags;
    uint32_t                                viewportCount;
    const VkViewport*                       pViewports;
    uint32_t                                scissorCount;
    const VkRect2D*                         pScissors;
} VkPipelineViewportStateCreateInfo;
```

The `sType` field of `VkPipelineViewportStateCreateInfo` should be set to `VK_STRUCTURE_TYPE_PIPELINE_VIEWPORT_STATE_CREATE_INFO`, and `pNext` should be set to `nullptr`. The `flags` field is reserved for use in a future version of Vulkan and should be set to zero.

The number of viewports that will be available to the pipeline is set in `viewportCount`, and the dimensions of each viewport are passed in an array of `VkViewport` structures, the address of which is specified in `pViewports`. The definition of `VkViewport` is

```
typedef struct VkViewport {
    float      x;
    float      y;
    float      width;
    float      height;
    float      minDepth;
    float      maxDepth;
} VkViewport;
```

The `VkPipelineViewportStateCreateInfo` structure is also used to set the scissor rectangles for the pipeline. As with viewports, a single pipeline can define multiple scissor rectangles, and they are passed through an array of `VkRect2D` structures. The number of scissor rectangles is specified in `scissorCount`. Note that the index used for the viewport and scissor rectangles when drawing is the same, so you must set `scissorCount` to the same value as `viewportCount`. `VkRect2D` is a simple structure defining a rectangle in 2D and is used for many things in Vulkan. Its definition is

```
typedef struct VkRect2D {
    VkOffset2D    offset;
    VkExtent2D    extent;
} VkRect2D;
```

Support for multiple viewports is optional. When multiple viewports are supported, then at least 16 are available. The maximum number of viewports that can be enabled in a single graphics pipeline can be determined by inspecting the `maxViewports` member of the `VkPhysicalDeviceLimits` structure returned from a call to **`vkGetPhysicalDeviceProperties()`**. If multiple viewports are supported, then this limit will be at least 16. Otherwise, this field will contain the value 1.

More information about how the viewport transformation works and how to utilize multiple viewports in your application is given in Chapter 9, "Geometry Processing." Further information about scissor testing is contained in Chapter 10, "Fragment Processing." In order to simply render to the full framebuffer, disable the scissor test and create a single viewport with the same dimensions as the framebuffer's color attachments.

Rasterization State

Rasterization is the fundamental process whereby primitives represented by vertices are turned into streams of fragments ready to be shaded by your fragment shader. The state of the rasterizer controls how this process occurs and is set using an instance of the `VkPipelineRasterizationStateCreateInfo` passed through the `pRasterizationState` member of `VkGraphicsPipelineCreateInfo`. The definition of `VkPipelineRasterizationStateCreateInfo` is

[Click here to view code image](#)

```c
typedef struct VkPipelineRasterizationStateCreateInfo {
    VkStructureType                            sType;
    const void*                                pNext;
    VkPipelineRasterizationStateCreateFlags    flags;
    VkBool32                                   depthClampEnable;
    VkBool32                                   rasterizerDiscardEnable;
    VkPolygonMode                              polygonMode;
    VkCullModeFlags                            cullMode;
    VkFrontFace                                frontFace;
    VkBool32                                   depthBiasEnable;
    float                                      depthBiasConstantFactor;
    float                                      depthBiasClamp;
    float                                      depthBiasSlopeFactor;
    float                                      lineWidth;
} VkPipelineRasterizationStateCreateInfo;
```

The `sType` field of `VkPipelineRasterizationStateCreateInfo` should be set to `VK_STRUCTURE_TYPE_PIPELINE_RASTERIZATION_STATE_CREATE_INFO` and `pNext` should be set to `nullptr`. The `flags` field is reserved and should be set to zero.

The `depthClampEnable` field is used to turn depth clamping on or off. Depth clamping causes fragments that would have been clipped away by the near or far planes to instead be projected onto those planes and can be used to fill holes in geometry that would be caused by clipping.

`rasterizerDiscardEnable` is used to turn off rasterization altogether. When this flag is set, the rasterizer will not run, and no fragments will be produced.

The `polygonMode` field can be used to get Vulkan to turn triangles into points or lines automatically. The possible values for `polygonMode` are

- `VK_POLYGON_MODE_FILL`: This is the normal mode that is used to fill in triangles. Triangles will be drawn solid, and every point inside the triangle will create a fragment.
- `VK_POLYGON_MODE_LINE`: This mode turns the triangles into lines, with each edge of each triangle becoming a line. This is useful for drawing geometry in wireframe mode.
- `VK_POLYGON_MODE_POINT`: This mode simply draws each vertex as a point.

The advantage of using the polygon mode to turn geometry into wireframe or point clouds over simply drawing lines or points is that operations that operate only on complete triangles, such as back-face culling, are still performed. Thus, lines that would have encompassed a culled triangle are not drawn, whereas they would be if the geometry were simply drawn as lines.

Culling is controlled with `cullMode`, which can be zero or a bitwise combination of either of the following:

- `VK_CULL_MODE_FRONT_BIT`: Polygons (triangles) that are considered to face the viewer are discarded.
- `VK_CULL_MODE_BACK_BIT`: Polygons that are considered to face away from the viewer are discarded.

For convenience, Vulkan defines `VK_CULL_MODE_FRONT_AND_BACK` as the bitwise OR of both `VK_CULL_MODE_FRONT_BIT` and `VK_CULL_MODE_BACK_BIT`. Setting `cullMode` to this value will result in all triangles being dicarded. Note that culling doesn't affect lines or points because they don't have a facing direction.

Which direction a triangle is facing is determined from the winding order of its vertices—whether they proceed clockwise or counterclockwise in window space. Which of clockwise or counterclockwise is considered front-facing is determined by the `frontFace` field. This is a member of the `VkFrontFace` enumeration and can be either `VK_FRONT_FACE_COUNTER_CLOCKWISE` or `VK_FRONT_FACE_CLOCKWISE`.

The next four parameters—`depthBiasEnable`, `depthBiasConstantFactor`, `depthBiasClamp`, and `depthBiasSlopeFactor`—control the *depth bias* feature. This feature allows fragments to be offset in depth before the depth test and can be used to prevent depth fighting. This feature is discussed in some detail in [Chapter 10](#), "[Fragment Processing](#)."

Finally, `lineWidth` sets the width of line primitives, in pixels. This applies to all lines rasterized with the pipeline. This includes pipelines in which the primitive topology is one of the line primitives, the geometry or tessellation shaders turn the input primitives into lines, and the polygon mode (set by `polygonMode`) is `VK_POLYGON_MODE_LINE`. Note that some Vulkan implementations don't support wide lines and will ignore this field. Others may run very slowly when this field is not 1.0; still others may honor this field completely and throw away all your lines if you set `lineWidth` to 0.0. Therefore, you should always set this field to 1.0 unless you're sure you want something else.

Even when wide lines are supported, the maximum width of a line is device-dependent. It is guaranteed to be at least 8 pixels but could be much higher. To determine the maximum line width supported by a device, check the `lineWidthRange` field of its `VkPhysicalDeviceLimits` structure. This is an array of two floating-point values, the first being the minimum width of a line (which will be at most 1 pixel; its purpose is for drawing lines that are less than a pixel wide) and the second being the maximum width of a line. If variable line width is not supported, then both elements of the array will be 1.0.

Further, as line width is changed, a device may snap the width you specify into fixed-size increments. For example, it may support only whole-pixel size changes. This is the line width granularity, which can be determined by inspecting the `lineWidthGranularity` field of the `VkPhysicalDeviceLimits` structure.

Multisample State

Multisampling is the process of generating multiple samples for each pixel in an image. It is used to combat aliasing and can greatly improve image quality when used effectively. When you use multisampling, the color and depth-stencil attachments must be multisample images, and the multisample state of the pipeline should be set appropriately through the `pMultisampleState` member of `VkGraphicsPipelineCreateInfo`. This is a pointer to an instance of the `VkPipelineMultisampleStateCreateInfo` structure, the definition of which is

```
typedef struct VkPipelineMultisampleStateCreateInfo {
    VkStructureType                          sType;
    const void*                              pNext;
    VkPipelineMultisampleStateCreateFlags    flags;
    VkSampleCountFlagBits                    rasterizationSamples;
    VkBool32                                 sampleShadingEnable;
    float                                    minSampleShading;
    const VkSampleMask*                      pSampleMask;
    VkBool32                                 alphaToCoverageEnable;
    VkBool32                                 alphaToOneEnable;
} VkPipelineMultisampleStateCreateInfo;
```

The `sType` field of `VkPipelineMultisampleStateCreateInfo` should be set to `VK_STRUCTURE_TYPE_PIPELINE_MULTISAMPLE_STATE_CREATE_INFO` and `pNext` should be set to `nullptr`. The `flags` field is reserved and should be set to zero.

Depth and Stencil State

The depth-stencil state controls how the depth and stencil tests are conducted and what happens to a fragment should it pass or fail either of those tests. The depth and stencil tests can be performed either before or after the fragment shader runs. By default, the tests occur after the fragment shader.[1]

1. Most implementations will only keep up the appearance that the depth and stencil tests are running after the fragment shader and, if possible, run the tests before running the shader to avoid running shader code when the test would fail.

To run the fragment shader before the depth test, we can apply the SPIR-V `EarlyFragmentTests` execution mode to the entry point of our fragment shader.

The depth-stencil state is configured through the `pDepthStencilState` member of `VkGraphicsPipelineCreateInfo`, which is a pointer to an instance of the `VkPipelineDepthStencilStateCreateInfo` structure. The definition of `VkPipelineDepthStencilStateCreateInfo` is

```
typedef struct VkPipelineDepthStencilStateCreateInfo {
    VkStructureType                           sType;
    const void*                               pNext;
    VkPipelineDepthStencilStateCreateFlags    flags;
    VkBool32                                  depthTestEnable;
    VkBool32                                  depthWriteEnable;
    VkCompareOp                               depthCompareOp;
    VkBool32                                  depthBoundsTestEnable;
    VkBool32                                  stencilTestEnable;
    VkStencilOpState                          front;
    VkStencilOpState                          back;
    float                                     minDepthBounds;
    float                                     maxDepthBounds;
} VkPipelineDepthStencilStateCreateInfo;
```

The sType field for `VkPipelineDepthStencilStateCreateInfo` should be set to `VK_STRUCTURE_TYPE_PIPELINE_DEPTH_STENCIL_CREATE_INFO` and `pNext` should be set to `nullptr`. The `flags` field is reserved and should be set to zero.

The depth test is enabled if `depthTestEnable` is set to `VK_TRUE`. If the depth test is enabled, then the test to use is selected using `depthCompareOp`, which is one of the `VkCompareOp` enumerant values. The available depth testing operations are discussed in more detail in [Chapter 10](#), "[Fragment Processing](#)." If `depthTestEnable` is set to `VK_FALSE`, then the depth test is enabled. The value of `depthCompareOp` is enabled, and all fragments are considered to have passed the depth test. It should be noted, however, that when the depth test is disabled, no writes to the depth buffer occur.

If the depth test passes (or if the depth test is disabled), then the fragment passes on to the stencil test. The stencil test is enabled if the `stencilTestEnable` field of `VkPipelineDepthStencilCreateInfo` is set to `VK_TRUE` and disabled otherwise. When stencil testing is enabled, a separate state is provided for front- and back-facing primitives in the `front` and `back` members, respectively. If stencil test is disabled, all fragments are considered to have passed the stencil test.

The details of depth and stencil testing are covered in more depth in [Chapter 10](#), "[Fragment Processing](#)."

Color Blend State

The final stage in the Vulkan graphics pipeline is the color blend stage. This stage is responsible for writing fragments into the color attachments. In many cases, this is a simple operation that simply overwrites the existing content of the attachment with value(s) output from the fragment shader. However, the color blender is capable of mixing (blending) those values with the values already in the framebuffer and performing simple logical operations between the output of the fragment shader and the current content of the framebuffer.

The state of the color blender is specified using the `pColorBlendState` member of the `VkGraphicsPipelineCreateInfo` structure. This is a pointer to an instance of the `VkPipelineColorBlendStateCreateInfo` structure, the definition of which is

[Click here to view code image](#)

```
typedef struct VkPipelineColorBlendStateCreateInfo {
    VkStructureType                            sType;
    const void*                                pNext;
    VkPipelineColorBlendStateCreateFlags       flags;
    VkBool32                                   logicOpEnable;
    VkLogicOp                                  logicOp;
    uint32_t                                   attachmentCount;
    const VkPipelineColorBlendAttachmentState* pAttachments;
    float                                      blendConstants[4];
} VkPipelineColorBlendStateCreateInfo;
```

The sType field of `VkPipelineColorBlendStateCreateInfo` should be set to `VK_STRUCTURE_TYPE_PIPELINE_COLOR_BLEND_STATE_CREATE_INFO`, and `pNext` should be set to `nullptr`. The `flags` field is reserved for future use and should be set to zero.

The `logicOpEnable` field specifies whether to perform logical operations between the output of the fragment shader and the content of the color attachments. When `logicOpEnable` is `VK_FALSE`, then logical operations are disabled and values produced by the fragment shader are written into the color attachment unmodified. When `logicOpEnable` is `VK_TRUE`, logic operations are enabled for the attachments that support them. The logic operation to apply is the same for every attachment and is a member of the `VkLogicOp` numeration. The meaning of each of the enumerants and more information about logical operations is given in [Chapter 10](#), "[Fragment Processing](#)."

Each attachment can have a different format, and can support different blending operations. These are specified with an array of `VkPipelineColorBlendAttachmentState` structures, the address of which is passed through the `pAttachments` member of `VkPipelineColorBlendStateCreateInfo`. The number of attachments is set in `attachmentCount`. The definition of `VkPipelineColorBlendAttachmentState` is

[Click here to view code image](#)

```
typedef struct VkPipelineColorBlendAttachmentState {
    VkBool32                blendEnable;
    VkBlendFactor           srcColorBlendFactor;
    VkBlendFactor           dstColorBlendFactor;
    VkBlendOp               colorBlendOp;
    VkBlendFactor           srcAlphaBlendFactor;
    VkBlendFactor           dstAlphaBlendFactor;
    VkBlendOp               alphaBlendOp;
    VkColorComponentFlags   colorWriteMask;
} VkPipelineColorBlendAttachmentState;
```

For each color attachment, the members of `VkPipelineColorBlendAttachmentState` control whether blending is enabled, what the source and destination factors are, what the blending operation is (for both the color and alpha channels separately), and which channels in the output image are to be updated.

If the `colorBlendEnable` field of `VkPipelineColorBlendAttachmentState` is `VK_TRUE`, then the remaining parameters control the state of blending. Blending will be covered in more detail in [Chapter 10](#), "[Fragment Processing](#)." When `colorBlendEnable` is `VK_FALSE`, the blending parameters in `VkPipelineColorBlendAttachmentState` are ignored, and blending is disabled for that attachment.

Regardless of the state of `colorBlendEnable`, the final field, `colorWriteMask`, controls which channels of the output image are written in this attachment. It is a bitfield made up of bits from the `VkColorComponentFlagBits` enumeration. The four channels, represented by `VK_COLOR_COMPONENT_R_BIT`, `VK_COLOR_COMPONENT_G_BIT`, `VK_COLOR_COMPONENT_B_BIT`, and `VK_COLOR_COMPONENT_A_BIT`, can be individually masked out for writing. If the flag corresponding to a particular channel is not included in `colorWriteMask`, then that channel will not be modified. Only the channels included in `colorWriteMask` will be updated through rendering to the attachment.

Dynamic State

As you have seen, the graphics pipeline object is large and complex, and contains a lot of state. In many graphics applications, it is often desirable to be able to change some states at a relatively high frequency. If every change in every state required that you created a new graphics pipeline object, then the number of objects your application would have to manage would quickly become very large.

To make fine-grained state changes more manageable, Vulkan provides the ability to mark particular parts of the graphics pipeline as *dynamic*, which means that they can be updated on the fly using commands directly inside the command buffer rather than using an object. Because this reduces the opportunity for Vulkan to optimize or absorb parts of state, it's necessary to specify exactly what state you want to make dynamic. This is done through the `pDynamicState` member of the `VkGraphicsPipelineCreateInfo` structure, which is a pointer to an instance of the `VkPipelineDynamicStateCreateInfo` structure, the definition of which is

[**Click here to view code image**](#)

```
typedef struct VkPipelineDynamicStateCreateInfo {
    VkStructureType                    sType;
    const void*                        pNext;
    VkPipelineDynamicStateCreateFlags  flags;
    uint32_t                           dynamicStateCount;
    const VkDynamicState*              pDynamicStates;
} VkPipelineDynamicStateCreateInfo;
```

The `sType` for `VkPipelineDynamicStateCreateInfo` should be set to `VK_STRUCTURE_TYPE_PIPELINE_DYNAMIC_STATE_CREATE_INFO`, and `pNext` should be set to `nullptr`. The `flags` field is reserved and should be set to zero.

The number of states that you wish to be dynamic is specified in `dynamicStateCount`. This is the length of the array pointed to by `pDynamicStates`, which is an array of elements from the `VkDynamicState` enumeration. Including a member of this enumeration in the `pDynamicStates` array tells Vulkan that you want to be able to change that state using the corresponding dynamic state setting commands. The members of `VkDynamicState` and their meanings are as follows:

- `VK_DYNAMIC_STATE_VIEWPORT`: The viewport rectangle is dynamic and will be updated using **vkCmdSetViewport()**.

- `VK_DYNAMIC_STATE_SCISSOR`: The scissor rectangle is dynamic and will be updated using **vkCmdSetScissor()**.

- `VK_DYNAMIC_STATE_LINE_WIDTH`: The line width is dynamic and will be updated using **vkCmdSetLineWidth()**.

- `VK_DYNAMIC_STATE_DEPTH_BIAS`: The depth bias parameters are dynamic and will be updated using **vkCmdSetDepthBias()**.

- `VK_DYNAMIC_STATE_BLEND_CONSTANTS`: The color blend constants are dynamic and will be updated using **vkCmdSetBlendConstants()**.

- `VK_DYNAMIC_STATE_DEPTH_BOUNDS`: The depth bounds parameters are dynamic and will be updated using **vkCmdSetDepthBounds()**.

- `VK_DYNAMIC_STATE_STENCIL_COMPARE_MASK`,
  `VK_DYNAMIC_STATE_STENCIL_WRITE_MASK`, and
  `VK_DYNAMIC_STATE_STENCIL_REFERENCE`: The corresponding stencil parameters are
  dynamic and will be updated using **vkCmdSetStencilCompareMask()**,
  vkCmdSetStencilWriteMask(), and **vkCmdSetStencilReference()**,
  respectively.

If a state is specified as being dynamic, then it becomes your responsibility to set that state when
binding the pipeline. If the state is not marked as dynamic, then it is considered *static* and is set when
the pipeline is bound. Binding a pipeline with static state makes the dynamic state undefined. The
reason for this is that Vulkan implementations might optimize static states into the pipeline object and
not actually program them into hardware if they're not used or if it can otherwise be determined that
this is valid. When a pipeline with that state marked dynamic is subsequently bound, it's not defined
whether the dynamic state is consistent in hardware.

When you are switching between pipelines that mark the same state as dynamic, however, then the
state remains persistent across binds. Table 7.1 illustrates this.

Old Pipeline	New Pipeline	State Valid
Dynamic	Dynamic	Yes
Dynamic	Static	Yes
Static	Static	Yes
Static	Dynamic	No

Table 7.1: Dynamic and Static State Validity

As you can see in Table 7.1, the only case in which state becomes undefined occurs when switching
from a pipeline with that state marked as static to one in which the same state is marked as dynamic.
In all other cases, the state is well defined and comes from either the pipeline's state or the dynamic
state set with the appropriate command.

If you set a dynamic state when a pipeline with that state set as static is currently bound, the results
are undefined if you then draw with that pipeline. The effect could be to ignore the state-setting
command and continue to use the static version of the state from the pipeline, honor the state-setting
command and use the new dynamic state, or corrupt state altogether and break your application
entirely. The effect will be different across implementations and will likely depend on which state is
erroneously overridden.

Setting dynamic state and then binding a pipeline with that state marked as dynamic should cause the
dynamic state to be used. However, it's good practice to bind the pipeline first and then bind any
related state simply to avert the possibility of undefined behavior.

Summary

This chapter provided a whirlwind tour of the Vulkan graphics pipeline. The pipeline consists of multiple stages, some of which are configurable but fixed-function, and some of which are made up of extremely powerful shaders. Building on the concept of the pipeline object introduced in Chapter 6, "Shaders and Pipelines," the graphics pipeline object was introduced. This object includes a large amount of fixed function state. Although the pipelines built in this chapter were simple, a solid foundation was laid upon which to build more complex and expressive pipelines in later chapters.

Chapter 8. Drawing

What You'll Learn in This Chapter

- The details of the different drawing commands in Vulkan
- How to draw many copies of data through instancing
- How to pass drawing parameters through buffers

Drawing is the fundamental operation in Vulkan that triggers work to be performed by a graphics pipeline. Vulkan includes several drawing commands, each generating graphics work in slightly different ways. This chapter delves deep into the drawing commands supported by Vulkan. First, we reintroduce the basic drawing command first discussed in [Chapter 7](), "[Graphics Pipelines]()"; then we explore *indexed* and *instanced* drawing commands. Finally, we discuss a method to retrieve the parameters for a drawing command from device memory and even generate them on the device itself.

Back in [Chapter 7](), "[Graphics Pipelines]()," you were introduced to your first drawing command, **vkCmdDraw()**. This command simply pushes vertices into the Vulkan graphics pipeline. When we introduced the command, we glossed over some of its parameters. We also hinted at the existence of other drawing commands. For reference, here is the prototype for **vkCmdDraw()** again:

[**Click here to view code image**]()

```
void vkCmdDraw (
    VkCommandBuffer                     commandBuffer,
    uint32_t                            vertexCount,
    uint32_t                            instanceCount,
    uint32_t                            firstVertex,
    uint32_t                            firstInstance);
```

As with all commands that execute on the device, the first parameter is a `VkCommandBuffer` handle. The number of vertices in each draw is specified in `vertexCount`, and the vertex index from which the vertices start is specified in `firstVertex`. The vertices that are sent down the pipeline have indices starting from `firstVertex` and proceed through `vertexCount` contiguous vertices. If you're using vertex buffers and attributes to automatically feed data into your vertex shader, then the shader will see data fetched from that contiguous section of the arrays. If you're using the vertex index directly in your shader, you will see it count monotonically from `firstVertex` upward.

Getting Ready to Draw

As we mentioned back in [Chapter 7](), "[Graphics Pipelines]()," all drawing is contained inside a renderpass. Although renderpass objects can encapsulate many subpasses, even simple rendering that draws into a single output image must be part of a renderpass. The renderpass is created by calling **vkCreateRenderPass()** as described in [Chapter 7](). To prepare for rendering, we need to call **vkCmdBeginRenderPass()**, which sets the current renderpass object and, perhaps more important, configures the set of output images that will be drawn into. The prototype of **vkCmdBeginRenderPass()** is

```
void vkCmdBeginRenderPass (
    VkCommandBuffer                         commandBuffer,
    const VkRenderPassBeginInfo*            pRenderPassBegin,
    VkSubpassContents                       contents);
```

The command buffer that will contain the commands issued inside the renderpass is passed in `commandBuffer`. The bulk of the parameters describing the renderpass are passed through a pointer to an instance of the `VkRenderPassBeginInfo` structure in `pRenderPassBegin`. The definition of `VkRenderPassBeginInfo` is

```
typedef struct VkRenderPassBeginInfo {
    VkStructureType      sType;
    const void*          pNext;
    VkRenderPass         renderPass;
    VkFramebuffer        framebuffer;
    VkRect2D             renderArea;
    uint32_t             clearValueCount;
    const VkClearValue*  pClearValues;
} VkRenderPassBeginInfo;
```

The `sType` field of the `VkRenderPassBeginInfo` structure should be set to `VK_STRUCTURE_TYPE_RENDER_PASS_BEGIN_INFO`, and `pNext` should be set to `nullptr`. The renderpass that is begun is specified in `renderPass`, and the framebuffer that we're going to render into is specified in `framebuffer`. As discussed in Chapter 7, "Graphics Pipelines," the framebuffer is the collection of images that will be rendered to by graphics commands.

Within any particular use of a renderpass, we can choose to render only into a small area of the attached images. To do this, use the `renderArea` member of the `VkRenderPassBeginInfo` structure to specify the rectangle in which all rendering will be contained. Simply setting `renderArea.offset.x` and `renderArea.offset.y` to 0 and `renderArea.extent.width` and `renderArea.extent.height` to the width and height of the images in the framebuffer tells Vulkan that you're going to render into the whole rendering area of the framebuffer.

If any of the attachments in the renderpass have a load operation of `VK_ATTACHMENT_LOAD_OP_CLEAR`, then the colors or values that you want to clear them to are specified in an array of `VkClearValue` unions, a pointer to which is passed in `pClearValues`. The number of elements in `pClearValues` is passed in `clearValueCount`. The definition of `VkClearValue` is

```
typedef union VkClearValue {
    VkClearColorValue            color;
    VkClearDepthStencilValue     depthStencil;
} VkClearValue;
```

If the attachment is a color attachment, then the values stored in the `color` member of the `VkClearValue` union are used, and if the attachment is a depth, stencil, or depth-stencil

attachment, then the values stored in the `depthStencil` member are used. `color` and `depthStencil` are instances of the `VkClearColorValue` and `VkClearDepthStencilValue` structures, respectively, the definitions of which are

[Click here to view code image](#)

```
typedef union VkClearColorValue {
    float        float32[4];
    int32_t      int32[4];
    uint32_t     uint32[4];
} VkClearColorValue;
```

and

[Click here to view code image](#)

```
typedef struct VkClearDepthStencilValue {
    float        depth;
    uint32_t     stencil;
} VkClearDepthStencilValue;
```

The index of each attachment is used to index into the array of `VkClearValue` unions. This means that if only some of the attachments have a load operation of `VK_ATTACHMENT_LOAD_OP_CLEAR`, then there could be unused entries in the array. There must be at least as many entries in the `pClearValues` array as the highest-indexed attachment with a load operation of `VK_ATTACHMENT_LOAD_OP_CLEAR`.

For each attachment with a load operation of `VK_ATTACHMENT_LOAD_OP_CLEAR`, if it is a color attachment, then the values of the `float32`, `int32`, or `uint32` arrays are used to clear the attachment, depending on whether it is a floating-point or normalized format, a signed integer format, or an unsigned integer format, respectively. If the attachment is a depth, stencil, or depth-stencil attachment, then the values of the `depth` and `stencil` members of the `depthStencil` member of the `VkClearValue` union are used to clear the appropriate aspect of the attachment.

Once the renderpass has begun, you can place drawing commands (which are discussed in the next section) in the command buffer. All rendering will be directed into the framebuffer specified in the `VkRenderPassBeginInfo` structure passed to **vkCmdBeginRenderPass()**. To finalize rendering contained in the renderpass, you need to end it by calling **vkCmdEndRenderPass()**, the prototype of which is

[Click here to view code image](#)

```
void vkCmdEndRenderPass (
    VkCommandBuffer                        commandBuffer);
```

After **vkCmdEndRenderPass()** has executed, any rendering directed through the renderpass is completed, and the content of the framebuffer is updated. Until then, the framebuffer's content is undefined. Only attachments with a store operation of `VK_ATTACHMENT_STORE_OP_STORE` will reflect the new content produced by the rendering inside the renderpass. If an attachment has a store operation of `VK_ATTACHMENT_STORE_OP_DONT_CARE`, then its content is undefined after the renderpass has completed.

Vertex Data

If the graphics pipeline you're going to use requires vertex data, before performing any drawing commands, you need to bind buffers to source the data from. When buffers are in use as the sources of vertex data, they are sometimes known as *vertex buffers*. The command to buffers for use as vertex data is **vkCmdBindVertexBuffers()**, and its prototype is

[Click here to view code image](#)

```
void vkCmdBindVertexBuffers (
    VkCommandBuffer                    commandBuffer,
    uint32_t                           firstBinding,
    uint32_t                           bindingCount,
    const VkBuffer*                    pBuffers,
    const VkDeviceSize*                pOffsets);
```

The command buffer to which to bind the buffers is specified in `commandBuffer`. A given pipeline may reference many vertex buffers, and **vkCmdBindVertexBuffers()** is capable of updating a subset of the bindings on a particular command buffer. The index of the first binding to update is passed in `firstBinding`, and the number of contiguous bindings to update is passed in `bindingCount`. To update noncontiguous ranges of vertex buffer bindings, you need to call **vkCmdBindVertexBuffers()** multiple times.

The `pBuffers` parameter is a pointer to an array of `bindingCount VkBuffer` handles to the buffer objects to be bound, and `pOffsets` is a pointer to an array of `bindingCount` offsets into the buffer objects at which the data for each binding starts. The values in `pOffsets` are specified in bytes. It is perfectly reasonable to bind the same buffer object with different offsets (or even the same offset, if that's what's required) to a command buffer; simply include the same `VkBuffer` handle multiple times in the `pBuffers` array.

The layout and format of the data in the buffers are defined by the graphics pipeline that will consume the vertex data. Therefore, the format of the data is not specified here, but in the `VkPipelineVertexInputStateCreateInfo` structure passed via the `VkGraphicsPipelineCreateInfo` used to create the graphics pipeline. Back in Chapter 7, "Graphics Pipelines," we showed an example of setting up *interleaved* vertex data as a C++ structure in Listing 7.3. Listing 8.1 shows a slightly more advanced example that uses one buffer to store position data alone and a second buffer that stores a per-vertex normal and texture coordinate.

Listing 8.1: Separate Vertex Attribute Setup

[Click here to view code image](#)

```
typedef struct vertex_t
{
    vmath::vec3 normal;
    vmath::vec2 texcoord;
} vertex;

static const
VkVertexInputBindingDescription vertexInputBindings[] =
{
    { 0, sizeof(vmath::vec4), VK_VERTEX_INPUT_RATE_VERTEX }, // Buffer 1
```

```
    { 1, sizeof(vertex), VK_VERTEX_INPUT_RATE_VERTEX }        // Buffer 2
};

static const
VkVertexInputAttributeDescription vertexAttributes[] =
{
    { 0, 0, VK_FORMAT_R32G32B32A32_SFLOAT, 0 },              // Position
    { 1, 1, VK_FORMAT_R32G32B32_SFLOAT, 0 },                // Normal
    { 2, 1, VK_FORMAT_R32G32_SFLOAT, sizeof(vmath::vec3) }  // Tex Coord
};

static const
VkPipelineVertexInputStateCreateInfo vertexInputStateCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_VERTEX_INPUT_STATE_CREATE_INFO,  // sType
    nullptr,                                       // pNext
    0,                                             // flags
    vkcore::utils::arraysize(vertexInputBindings), //
vertexBindingDescription-
                                                   // Count
    vertexInputBindings,                           // pVertexBinding-
                                                   // Descriptions
    vkcore::utils::arraysize(vertexAttributes),    // vertexAttribute-
                                                   // DescriptionCount
    vertexAttributes                               // pVertexAttribute-
                                                   // Descriptions
};
```

In [Listing 8.1](#), we have defined three vertex attributes spread across two buffers. In the first buffer, only a single **vec4** variable is stored, and this is used for position. The stride for this buffer is therefore the size of a **vec4**, which is 16 bytes. In the second buffer, we store the interleaved normal and texture coordinates for the vertex. We represent this as the `vertex` structure, allowing the compiler to compute the stride for us.

Indexed Draws

Simply pushing contiguous runs of vertices into the pipeline isn't always what you want. In most geometric meshes, many vertices are used more than once. A fully connected mesh may share a single vertex among many triangles. Even a simple cube shares each vertex among three adjacent triangles. It is extremely wasteful to have to specify each vertex three times in your vertex buffers. Besides this, some Vulkan implementations are smart enough that if they see a vertex with the same input parameters more than once, they can skip processing it a second time and subsequent times, and instead reuse the results of the first vertex shader invocation.

To enable this, Vulkan allows *indexed draws*. The indexed equivalent of **vkCmdDraw()** is **vkCmdDrawIndexed()**, the prototype of which is

[Click here to view code image](#)

```
void vkCmdDrawIndexed (
    VkCommandBuffer                    commandBuffer,
    uint32_t                           indexCount,
    uint32_t                           instanceCount,
```

```
    uint32_t                                firstIndex,
    int32_t                                 vertexOffset,
    uint32_t                                firstInstance);
```

Again, the first parameter to **vkCmdDrawIndexed()** is the handle to the command buffer in which the draw will be executed. Rather than simply starting from zero and counting upward, however, **vkCmdDrawIndexed()** fetches indices from an *index buffer*. The index buffer is a regular buffer object that you bind to the command buffer by calling **vkCmdBindIndexBuffer()**, the prototype of which is

[Click here to view code image](#)

```
void vkCmdBindIndexBuffer (
    VkCommandBuffer                         commandBuffer,
    VkBuffer                                buffer,
    VkDeviceSize                            offset,
    VkIndexType                             indexType);
```

The command buffer to which to bind the index buffer is specified in `commandBuffer`, and the handle to the buffer object containing indexed data is specified in `buffer`. A section of a buffer object can be bound to the command buffer starting from `offset`. The bound section always extends to the end of the buffer object. There is no bounds checking on the index buffer; Vulkan will read as many indices from the buffer as you tell it to. However, it will never read past the end of the buffer object.

The data type of the indices in the buffer is specified in `indexType`. This is a member of the `VkIndexType` enumeration, the members of which are

• `VK_INDEX_TYPE_UINT16`: Unsigned 16-bit integers

• `VK_INDEX_TYPE_UINT32`: Unsigned 32-bit integers

When you call **vkCmdDrawIndexed()**, Vulkan will start fetching data from the currently bound index buffer at an offset of

[Click here to view code image](#)

```
            offset + firstIndex * sizeof(index)
```

where **sizeof**(index) is 2 for `VK_INDEX_TYPE_UINT16` and 4 for `VK_INDEX_TYPE_UINT32`. The code will fetch `indexCount` contiguous integers from the index buffer and then add `vertexOffset` to them. This addition is always performed in 32 bits, regardless of the index type for the currently bound index buffer. It is not defined what would happen if this addition overflowed the 32-bit unsigned integer range, so you should avoid that.

A schematic illustrating the data flow is shown in [Figure 8.1](#).
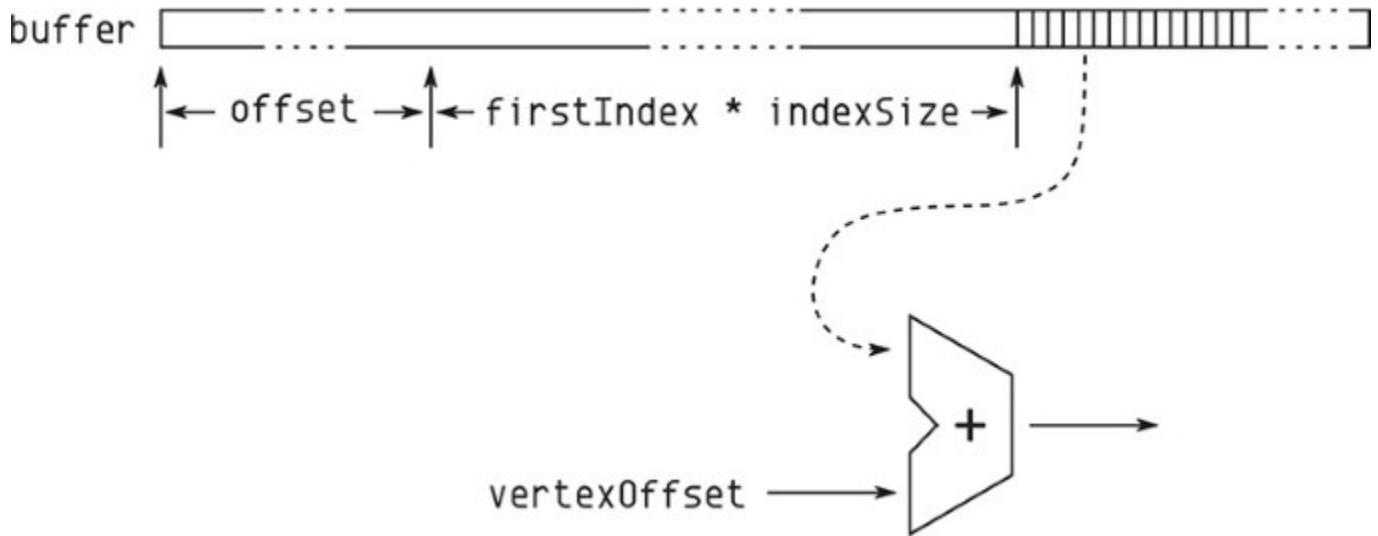
Figure 8.1: Index Data Flow

Note that when the index type is `VK_INDEX_TYPE_UINT32`, the maximum range of index values may not be supported. To check this, look at the `maxDrawIndexedIndexValue` field of the device's `VkPhysicalDeviceLimits` structure, which you can retrieve by calling

`vkGetPhysicalDeviceProperties()`. This value will always be at least 2^{24}-1 and may be as high as 2^{32}-1.

To demonstrate the effectiveness of the use of index data, Listing 8.2 shows the difference between the data required for drawing a simple cube using indexed and nonindexed data.

Listing 8.2: Indexed Cube Data

Click here to view code image

```
// Raw, non-indexed data
static const float vertex_positions[] =
{
    -0.25f, 0.25f, -0.25f,
    -0.25f, -0.25f, -0.25f,
     0.25f, -0.25f, -0.25f,


     0.25f, -0.25f, -0.25f,
     0.25f, 0.25f, -0.25f,
    -0.25f, 0.25f, -0.25f,

     0.25f, -0.25f, -0.25f,
     0.25f, -0.25f, 0.25f,
     0.25f, 0.25f, -0.25f,

     0.25f, -0.25f, 0.25f,
     0.25f, 0.25f, 0.25f,
     0.25f, 0.25f, -0.25f,
```

```
         0.25f, -0.25f,  0.25f,
        -0.25f, -0.25f,  0.25f,
         0.25f,  0.25f,  0.25f,

        -0.25f, -0.25f,  0.25f,
        -0.25f,  0.25f,  0.25f,
         0.25f,  0.25f,  0.25f,

        -0.25f, -0.25f,  0.25f,
        -0.25f, -0.25f, -0.25f,
        -0.25f,  0.25f,  0.25f,

        -0.25f, -0.25f, -0.25f,
        -0.25f,  0.25f, -0.25f,
        -0.25f,  0.25f,  0.25f,

        -0.25f, -0.25f,  0.25f,
         0.25f, -0.25f,  0.25f,
         0.25f, -0.25f, -0.25f,

         0.25f, -0.25f, -0.25f,
        -0.25f, -0.25f, -0.25f,
        -0.25f, -0.25f,  0.25f,

        -0.25f,  0.25f, -0.25f,
         0.25f,  0.25f, -0.25f,
         0.25f,  0.25f,  0.25f,

         0.25f,  0.25f,  0.25f,
        -0.25f,  0.25f,  0.25f,
        -0.25f,  0.25f, -0.25f
};

static const uint32_t vertex_count = sizeof(vertex_positions) /
                                     (3 * sizeof(float));

// Indexed vertex data
static const float indexed_vertex_positions[] =
{
        -0.25f, -0.25f, -0.25f,
        -0.25f,  0.25f, -0.25f,
         0.25f, -0.25f, -0.25f,
         0.25f,  0.25f, -0.25f,
         0.25f, -0.25f,  0.25f,
         0.25f,  0.25f,  0.25f,
        -0.25f, -0.25f,  0.25f,
        -0.25f,  0.25f,  0.25f,
};

// Index buffer
static const uint16_t vertex_indices[] =
{
        0, 1, 2,
```

```
        2, 1, 3,
        2, 3, 4,
        4, 3, 5,
        4, 5, 6,
        6, 5, 7,
        6, 7, 0,
        0, 7, 1,
        6, 0, 2,
        2, 4, 6,
        7, 5, 3,
        7, 3, 1
    };

    static const uint32_t index_count =
    vkcore::utils::arraysize(vertex_indices);
```

As you can see in Listing 8.2, the amount of data used to draw the cube is quite small. Only the vertex data for the 8 unique vertices is stored, along with 36 indices used to reference them. As geometry sizes go up with scene complexity, the savings can be quite large. In this simple example, the nonindexed vertex data is 36 vertices, each consisting of 3 elements of 4 bytes, which is a total of 432 bytes of data. Meanwhile, the indexed data is 12 vertices, again each of 3 elements of 4 bytes, plus 36 indices, each consuming 2 bytes of storage. This produces a total 168 bytes of data for the indexed cube.

In addition to the space savings provided by using indexed data, many Vulkan implementations include a vertex cache that can reuse the results of computations performed on vertex data. If the vertices are nonindexed, then the pipeline must assume that they are all unique. However, when the vertices are indexed, two vertices with the same index *are* the same. In any closed mesh, the same vertex will appear more than once, as it is shared among multiple primitives. This reuse can save quite a bit of work.

Index-Only Rendering

The raw index of the current vertex is available to your vertex shaders. This index appears in the variable decorated with VertexIndex in a SPIR-V shader, which is generated using the gl_VertexIndex built-in variable in GLSL. This contains the content of the index buffer (or the automatically generated vertex index) plus the value of vertexOffset passed to **vkCmdDrawIndexed()**.

You can use this index to fetch data from a buffer, for example. This allows you to pump geometry into the pipeline without worrying about vertex attributes. However, in some scenarios, a single 32-bit value might be all you need. In these cases, you can use the vertex index directly as data. Vulkan doesn't actually care what the values in the index buffer are so long as you don't use them to address into vertex buffers.

The object local vertex position for many pieces of geometry can be represented by 16-, 10-, or even 8-bit data with sufficient precision. Three 10-bit values can be packed inside a single 32-bit word. In fact, this is exactly what the VK_FORMAT_A2R10G10B10_SNORM_PACK32 format (and its unsigned counterpart) represent. Although the vertex data is not usable directly as an index buffer, it's possible to manually unpack the vertex data in the shader as though it had that format. As such, by simply unpacking the index in our shader, we can draw simple geometry without anything more than an index buffer.

Listing 8.3 shows the GLSL shader you use to do this unpacking operation.

Listing 8.3: Using the Vertex Index in a Shader

Click here to view code image

```
#version 450 core

vec3 unpackA2R10G10B10_snorm(uint value)
{
    int val_signed = int(value);
    vec3 result;
    const float scale = (1.0f / 512.0f);

    result.x = float(bitfieldExtract(val_signed, 20, 10));
    result.y = float(bitfieldExtract(val_signed, 10, 10));
    result.z = float(bitfieldExtract(val_signed, 0, 10));

    return result * scale;
}

void main(void)
{
    gl_Position = vec4(unpackA2R10G10B10_snorm(gl_VertexIndex), 1.0f);
}
```

The vertex shader shown in Listing 8.3 simply unpacks the incoming vertex index by using the `unpackA2R10G10B10_snorm` function. The resulting value is written to `gl_Position`. The 10 bits of precision in each of the `x`, `y`, and `z` coordinates effectively *snap* our vertices to a 1,024 × 1,024 × 1,024 grid of positions. This is sufficient in many cases. If an additional scale is applied, this can be passed to the shader via a push constant, and if the vertex is to undergo other transforms via matrix multiplications, for example, those transforms still proceed at full precision.

Reset Indices

Another subfeature of indexed draws allows you to use the primitive restart index. This special index value can be stored in the index buffer that can be used to signal the start of a new primitive. It is most useful when the primitive topology is one of the long, continuous primitives, including `VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP`, `VK_PRIMITIVE_TOPOLOGY_TRIANGLE_FAN`, and `VK_PRIMITIVE_TOPOLOGY_LINE_STRIP`, along with the adjacency versions of those topologies, `VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP_WITH_ADJACENCY`, and `VK_PRIMITIVE_TOPOLOGY_LINE_STRIP_WITH_ADJACENCY`.

The primitive restart feature is enabled using the `VkPipelineInputAssemblyStateCreateInfo` structure passed through the `pInputAssemblyState` member of the `VkGraphicsPipelineCreateInfo` structure used to create a graphics pipeline. Again, the definition of this structure is

```
typedef struct VkPipelineInputAssemblyStateCreateInfo {
    VkStructureType                            sType;
    const void*                                pNext;
    VkPipelineInputAssemblyStateCreateFlags    flags;
    VkPrimitiveTopology                        topology;
    VkBool32                                   primitiveRestartEnable;
} VkPipelineInputAssemblyStateCreateInfo;
```

The `topology` field has to be set to one of the primitive topologies that supports primitive restarts (the list and fan topologies mentioned earlier), and the `primitiveRestartEnable` field is set to `VK_TRUE`. When primitive restart is enabled, the special value of the maximum possible value for the index type (`0xFFFF` for `VK_INDEX_TYPE_UINT16` and `0xFFFFFFFF` for `VK_INDEX_TYPE_UINT32`) is used as the special restart marker.

If primitive restart is not enabled, the special reset marker is treated as a normal vertex index. While using 32-bit indices, it's unlikely that you'll ever need to use this value, because that would mean you had more than 4 billion vertices. However, the index value can still be passed to the vertex shader. It's not valid to enable the reset for primitive topologies other than the strip and fan topologies mentioned earlier.

When Vulkan encounters the reset value in the index buffer, it ends the current strip or fan and starts a new one beginning with the vertex addressed by the next index in the index buffer. If the reset value appears multiple times in a row, Vulkan simply skips them, looking for the next nonreset index value. If there aren't enough vertices to form a complete primitive (for example, if the reset index appears before three nonreset vertices are seen when the primitive topology is `VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP` or `VK_PRIMITIVE_TOPOLOGY_TRIANGLE_FAN`), then Vulkan will throw out all the vertices used so far and start a new primitive.

[Figure 8.2](#) shows how reset indices affect the `VK_PRIMITIVE_TOPOLOGY_TRIANGLE_STRIP` topology. In the top strip, contiguous indices between 0 and 12 are used to create a single long strip. When you enable primitive resets and replace index 6 with reset index value `0xFFFFFFFF`, the strip stops after the first four triangles and restarts with a triangle between vertices 7, 8, and 9.
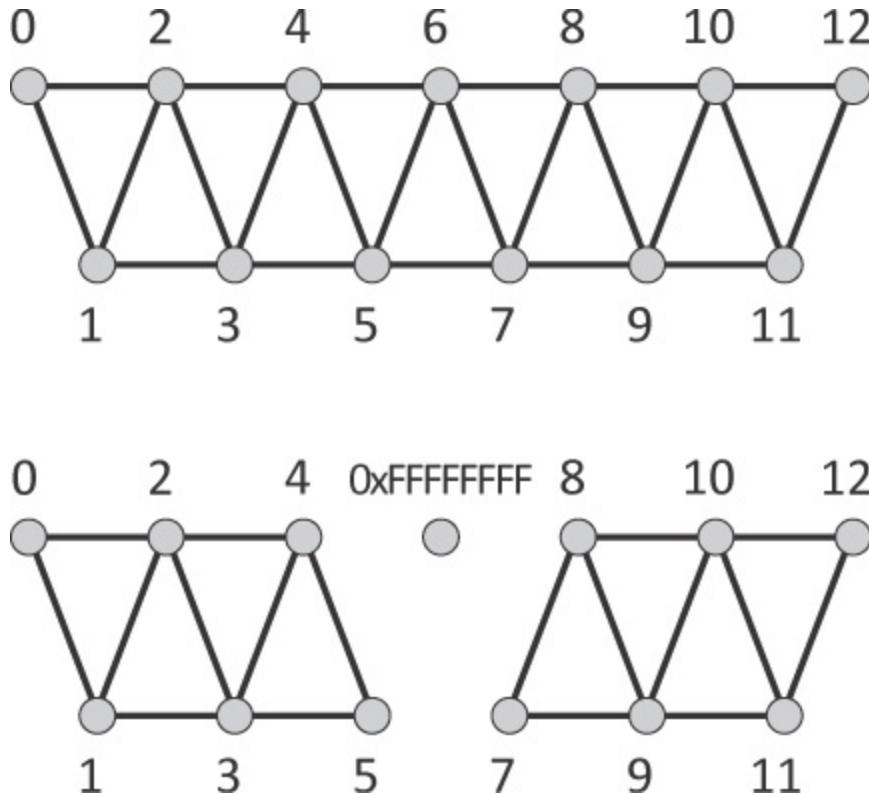
Figure 8.2: The Effect of Primitive Restart on Triangle Strips

The reset index is useful for cutting a very large draw using strips or fans into many smaller pieces. There comes a cut-off point in usefulness when the number of subdraws (individual strips or fans) decreases and their size increases, at which point it's probably best to simply produce two separate draws. This is especially true if it means switching pipelines between one with primitive restart enabled and one with it disabled.

If your model consists of hundreds or thousands of short strips, it might be a good idea to use primitive restart. If your model consists of a handful of very long strips, just make multiple drawing commands. Also, in some architectures, using the reset index can affect performance, and it may be better to simply use list topologies and unroll the index buffer rather than try to use strips.

Instancing

There are two parameters to **vkCmdDraw()** and **vkCmdDrawIndexed()** that we have thus far glossed over. These are the `firstInstance` and `instanceCount` parameters, and they are used to control instancing. This is a technique whereby many copies of the same geometry can be sent into the graphics pipeline. Each copy is known as an instance. At first, this seems like it wouldn't be much use, but there are two ways that your application can apply variation to each of the instances of the geometry:

- Use the `InstanceIndex` built-in decoration on a vertex shader input to receive the index of the current instance as an input to the shader. This input variable can then be used to fetch parameters from a uniform buffer or programmatically compute per-instance variation, for example.

- Use *instanced* vertex attributes to have Vulkan feed your vertex shader with unique data for each instance.

Listing 8.4 shows an example of using the instance index through the `gl_InstanceIndex` built-in variable in GLSL. The example draws many different cubes using instancing where each instance of the cube has a different color and transformation applied. The transformation matrix and color of each cube are placed in arrays that are stored in a pair of uniform buffers. The shader then indexes into these arrays with the `gl_InstanceIndex` built-in variable. The result of rendering with this shader is shown in Figure 8.3.

Listing 8.4: Using the Instance Index in a Shader

Click here to view code image

```glsl
#version 450 core

layout (set = 0, binding = 0) uniform matrix_uniforms_b
{
    mat4 mvp_matrix[1024];
};

layout (set = 0, binding = 1) uniform color_uniforms_b
{
    vec4 cube_colors[1024];
};
layout (location = 0) in vec3 i_position;

out vs_fs
{
    flat vec4 color;
};
void main(void)
{
    float f = float(gl_VertexIndex / 6) / 6.0f;
    vec4 color1 = cube_colors[gl_InstanceIndex];
    vec4 color2 = cube_colors[gl_InstanceIndex & 512];

    color = mix(color1, color2, f);
    gl_Position = mvp_matrix[gl_InstanceIndex] * vec4(i_position, 1.0f);
}
```
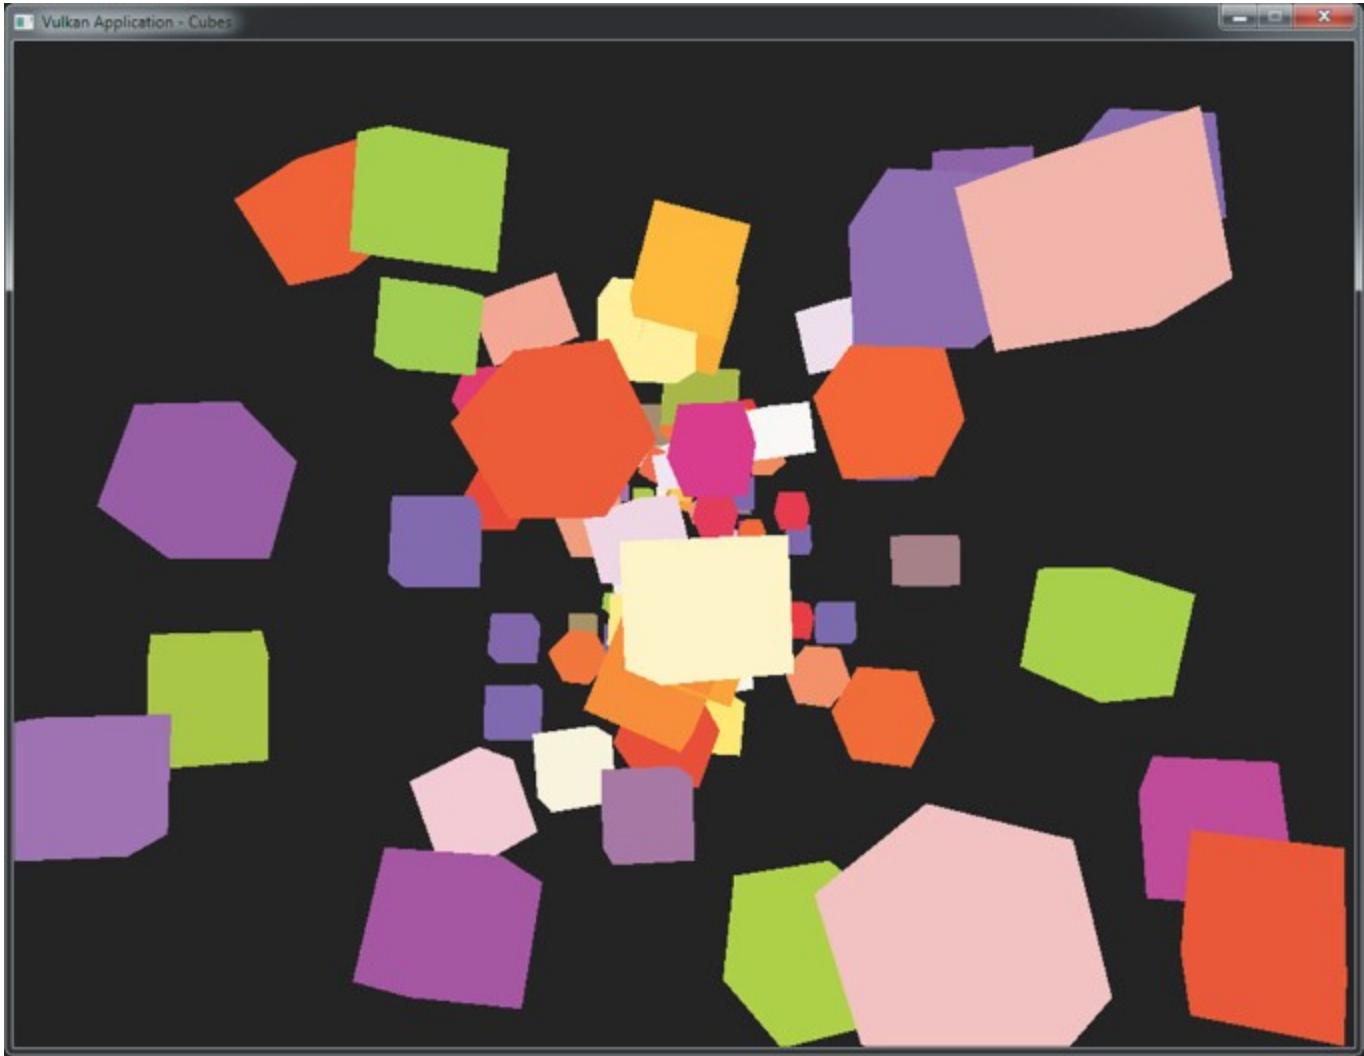
Figure 8.3: Many Instanced Cubes

Indirect Draws

In the **vkCmdDraw()** and **vkCmdDrawIndexed()** commands, the parameters to the command (vertexCount, vertexOffset, and so on) are passed as immediate parameters directly to the commands themselves. This means that you need to know the exact parameters of each draw call at the time that your application builds its command buffers. In most cases, having access to the parameters of drawing commands is a natural part of the application. However, in some situations, you don't know the exact parameters for each and every draw. Examples include the following:

- The overall structure of the geometry is known, but the exact number of vertices and locations of data in the vertex buffers is not known, such as when an object is always rendered the same way but its level of detail may change over time.

- The drawing commands are to be generated by the device, rather than the host. In this situation, the total number and layout of vertex data may *never* be known to the host.

In these cases, you can use an *indirect draw*, which is a drawing command that sources its parameters from device-accessible memory rather than embedding them in the command buffer along with the command. The first indirect draw command is **vkCmdDrawIndirect()**, which performs a nonindexed draw using parameters contained in a buffer. Its prototype is

```
void vkCmdDrawIndirect (
    VkCommandBuffer                         commandBuffer,
    VkBuffer                                buffer,
    VkDeviceSize                            offset,
    uint32_t                                drawCount,
    uint32_t                                stride);
```

The command itself is still placed into the command buffer, just as with **vkCmdDraw()**. commandBuffer is the command buffer into which the command is placed. However, the parameters of the command are sourced from the buffer specified in buffer at the offset specified in offset, which is measured in bytes. At this offset in the buffer, an instance of the VkDrawIndirectCommand structure should appear, containing the actual parameters of the command. Its definition is

```
typedef struct VkDrawIndirectCommand {
    uint32_t    vertexCount;
    uint32_t    instanceCount;
    uint32_t    firstVertex;
    uint32_t    firstInstance;
} VkDrawIndirectCommand;
```

The members of VkDrawIndirectCommand have the same meanings as the similarly named parameters of **vkCmdDraw()**. vertexCount and instanceCount are the numbers of vertices and indices to invoke, respectively, and firstVertex and firstInstance are the starting values for the vertex and instance indices, respectively.

vkCmdDrawIndirect() performs a nonindexed, indirect draw using parameters from a buffer object. It's also possible to perform an indexed indirect draw using **vkCmdDrawIndexedIndirect()**. The prototype of this function is

```
void vkCmdDrawIndexedIndirect (
    VkCommandBuffer                         commandBuffer,
    VkBuffer                                buffer,
    VkDeviceSize                            offset,
    uint32_t                                drawCount,
    uint32_t                                stride);
```

The parameters of **vkCmdDrawIndexedIndirect()** are identical to those of **vkCmdDrawIndirect()**. commandBuffer is the command buffer into which the command is written, buffer is the buffer containing the parameters; and offset is the offset, in bytes, at which the parameters are located in that buffer. However, the data structure containing the parameters of **vkCmdDrawIndexedIndirect()** is different. It is an instance of the VkDrawIndexedIndirectCommand structure, the definition of which is

```
typedef struct VkDrawIndexedIndirectCommand {
    uint32_t    indexCount;
```

```
    uint32_t    instanceCount;
    uint32_t    firstIndex;
    int32_t     vertexOffset;
    uint32_t    firstInstance;
} VkDrawIndexedIndirectCommand;
```

Again, the members of `VkDrawIndexedIndirectCommand` have the same meanings as the similarly named parameters of **`vkCmdDrawIndexed()`**. `indexCount` and `instanceCount` are the numbers of vertex indices and instances to push into the pipeline; the `firstIndex` member specifies where to start fetching indices from the index buffer; `vertexOffset` specifies the offset value to be added to the index data; and `firstInstance` specifies the value from which the instance counter should start counting.

What is important to remember about indirect drawing commands is that while the buffer object and the offset into it are baked into the command buffer, the parameters for the draw don't need to be in the sourced buffer object until the command buffer is executed by the device. As the device executes the command buffer, when it reaches the command, it will read whatever parameters are in the buffer and execute the drawing command as though those parameters had been specified directly to a regular drawing command. As far as the rest of the pipeline is concerned, there is no difference between a direct and an indirect draw.

This means several things:

- You can build command buffers with indirect draws long before they're needed, filling in the final parameters for the draw (in the buffer object rather than the command buffer) before the command buffer is submitted for execution.

- You can create a command buffer containing an indirect draw, submit it, overwrite the parameters in the buffer object, and submit the same command buffer again. This effectively *patches* new parameters into what could be a long, complex command buffer.

- You can write parameters into a buffer object by using stores from a shader object, or by using a command such as **`vkCmdFillBuffer()`** or **`vkCmdCopyBuffer()`** to generate drawing parameters on the device itself—either in the same command buffer or in another submitted just before the one containing the draw commands.

You may have noticed that both **`vkCmdDrawIndirect()`** and **`vkCmdDrawIndexedIndirect()`** take a `drawCount` and a `stride` parameter. These parameters allow you to pass *arrays* of drawing commands to Vulkan. A single call to **`vkCmdDrawIndirect()`** or **`vkCmdDrawIndexedIndirect()`** will kick off `drawCount` separate draws, each sourcing its parameters from a `VkDrawIndirectCommand` or `VkDrawIndexedIndirectCommand` structure, respectively.

The array of structures still begins at `offset` bytes into the buffer object, and each of these structures is separated from the previous by `stride` bytes. If `stride` is zero, then the same parameter structure will be used for every draw.[1]

1. Note that this behavior differs from OpenGL, in which a stride of zero causes the device to assume a tightly packed array, and it is impossible to source the same parameters over and over.

The number of draws is still baked into the command buffer, but draws whose `indexCount` or `instanceCount` parameters are zero will be skipped by the device. While this doesn't mean that you can truly produce a fully dynamic draw count, by using a fixed upper limit on the number of draws and ensuring that all unused entries in the array of parameters have at least one of

`vertexCoount`, `indexCount`, or `instanceCount`, set to zero, you can generate a variable number of draws by using a single command.

Note that support for counts other than one (and zero) is optional. To check whether the device supports a count greater than one, check the `multiDrawIndirect` field of the device's `VkPhysicalDeviceFeatures` structure as returned from a call to **`vkGetPhysicalDeviceFeatures()`**, and remember to enable the feature in the set of enabled features passed to **`vkCreateDevice()`** when creating the logical device.

When indirect counts are supported, the maximum number of draws that can be passed to a single call to **`vkCmdDrawIndirect()`** or **`vkCmdDrawIndexedIndirect()`** may still be limited. To check the supported count, inspect the `maxDrawIndirectCount` field of the device's `VkPhysicalDeviceLimits` structure. When `multiDrawIndirect` is not supported, this field will be 1. If it is supported, then it is guaranteed to be at least 65,535. If the number of draws you're pushing through each of these commands is less than this amount, then there's no need to directly check the limit.

Drawing many pieces of geometry back to back with the same pipeline and graphics state can sometimes be limiting. However, in many cases, all that is different between draws is parameters passed to shaders. This is especially true when applications use über shaders or physically based rendering techniques. There is no direct way to pass parameters to the individual draws that make up a single call to **`vkCmdDrawIndirect()`** or **`vkCmdDrawIndexedIndirect()`** with a `drawCount` greater than one. However, what is available in the shader is the SPIR-V decoration `DrawIndex` decoration on an input to the vertex shader. This is produced by using the `gl_DrawIDARB` input in GLSL.

When decorated with `DrawIndex`, the shader input will contain the index of the draw, starting from zero and counting upward as the draws are generated by the device. This can then be used to index into arrays of data stored in uniform or shader storage blocks. Listing 8.5 shows a GLSL shader that uses `gl_DrawIDARB` to retrieve per-draw parameters from a shader storage block.

Listing 8.5: Draw Index Used in a Shader

Click here to view code image

```
#version 450 core

// Enable the GL_ARB_shader_draw_parameters extensions.
#extension GL_ARB_shader_draw_parameters : require

layout (location = 0) in vec3 position_3;

layout (set = 0, binding = 0) uniform FRAME_DATA
{
    mat4 view_matrix;
    mat4 proj_matrix;
    mat4 viewproj_matrix;
};

layout (set = 0, binding = 1) readonly buffer OBJECT_TRANSFORMS
{
    mat4 model_matrix[];
```

```
    };

    void main(void)
    {
        // Extend input position to vec4.
        vec4 position = vec4(position_3, 1.0);

        // Compute per-object model-view matrix.
        mat4 mv_matrix = view_matrix * model_matrix[gl_DrawIDARB];

        // Output position using global projection matrix.
        gl_Position = proj_matrix * P;
    }
```

The shader in Listing 8.5 uses a single uniform block to store per-frame constants and a single shader storage block to store a large array of per-object transformation matrices.[2] The `gl_DrawIDARB` built-in variable is used to index into the `model_matrix` array stored in the shader storage block. The result is that each subdraw in the single **vkCmdDrawIndirect()** call uses its own model transformation matrix.

> 2. At the time of writing, the reference GLSL compiler does not contain support for the `GL_ARB_draw_parameters` extension that exposes `gl_DrawID`. This shader was developed in an OpenGL test environment and then edited to suit Vulkan. It is expected to work once support for `GL_ARB_draw_parameters` lands in the reference compiler.

Summary

This chapter covered the various drawing commands supported by Vulkan. You were reintroduced to **vkCmdDraw()**, which was first mentioned in Chapter 7, "Graphics Pipelines," and which produces nonindexed draws. Indexed draws were covered, and then we explored instancing, which is a technique for drawing many copies of the same geometry with varying parameters driven by the instance index. Finally, we looked at indirect draws, which allow the parameters for drawing commands to be sourced from device memory rather than specified at command-buffer construction time. Together, instancing and indirect draws are powerful tools that allow complex scenes to be built up with very few drawing commands.

Chapter 9. Geometry Processing

What You'll Learn in This Chapter

- Using tessellation to increase the geometric detail of your scene
- Using geometry shaders to process whole primitives
- Clipping geometry against user-specified planes

While many Vulkan programs will stick to vertex and fragment shaders, two optional pieces of functionality can be used to increase the geometric detail of the rendered images. These functions are tessellation and geometry shading. Although these concepts were briefly introduced earlier, this chapter digs deeper into the details of both tessellation and geometry shader functionality and discusses how to make effective use of these powerful sections of the geometry processing pipeline.

Tessellation

Tessellation is controlled by a collection of stages that appear near the front of the graphics pipeline, immediately after vertex shading. We briefly introduced tessellation in Chapter 7, "Graphics Pipelines." However, because tessellation is an optional stage in the pipeline, we mostly glossed over it in order to cover the remaining stages. This section covers it in more detail.

Tessellation takes as input patches, which are really just collections of control points represented as vertices, and breaks them down into many smaller, simpler primitives—such as points, lines, or triangles—that can be rendered by the rest of the pipeline in the normal manner. Tessellation is an optional feature in Vulkan. Presence of support can be determined by checking the `tessellationShader` member of the device's `VkPhysicalDeviceFeatures` structure. If this is `VK_FALSE`, then pipelines containing tessellation shaders cannot be created or used in your application.

Tessellation Configuration

From the application's perspective, the tessellation engine is a fixed-function, though highly configurable, block of functionality surrounded by two shader stages. The first stage, the tessellation control shader, is responsible for processing the control points of a patch, setting up some per-patch parameters, and handing control to the fixed-function tessellation block. This block takes the patch and breaks it up into the fundamental point, line, or triangle primitives, finally passing the resulting generated vertex data to a second shading stage: the tessellation evaluation shader. This shader appears much like a vertex shader except that it runs for each generated vertex.

Tessellation is controlled and configured through a combination of two sources of information. The first source is the `VkPipelineTessellationStateCreateInfo` structure passed through the `VkGraphicsPipelineCreateInfo` structure used to create the graphics pipeline. Introduced in Chapter 7, "Graphics Pipelines," `VkPipelineTessellationStateCreateInfo` is defined as

```
typedef struct VkPipelineTessellationStateCreateInfo {
    VkStructureType                          sType;
    const void*                              pNext;
    VkPipelineTessellationStateCreateFlags   flags;
    uint32_t                                 patchControlPoints;
} VkPipelineTessellationStateCreateInfo;
```

The only member that affects tessellation state in this structure is `patchControlPoints`, which sets the number of control points that make up a patch. The remaining state of the tessellation system is set by using the two shaders.

The maximum number of control points that can be used to construct a patch is implementation-dependent but is guaranteed to be at least 32. If tessellation is supported, then the Vulkan implementation will support at least 32 control points per patch, so if you never use patches larger than this, there is no reason to query the upper limit. If you need to use patches larger than 32 control points, you can determine the supported maximum by inspecting the `maxTessellationPatchSize` member of the device's `VkPhysicalDeviceLimits` structure as returned from a call to **vkGetPhysicalDeviceProperties()**.

Tessellation Modes

The fundamental operation of the tessellation engine is to take the patch and, given a set of tessellation *levels*, subdivide each edge according to its level. The distance of each subdivided point along each edge is assigned a value between 0.0 and 1.0. The two main modes of tessellation treat the patch as either a rectangle or a triangle. When the patch is tessellated as a rectangle, the subdivided coordinates form a 2D barycentric coordinate, and when the patch is tessellated as a triangle, the generated vertices have 3D barycentric coordinates.

Each patch has a set both of inner and of outer tessellation levels. The outer tessellation levels control the level of tessellation along the outer edge of the patch. If you set this level the same as that calculated for adjacent patches in a larger geometry, you can form seamless joins. The inner tessellation modes control the level of tessellation in the center of the patch. Figure 9.1 shows how the inner and outer levels are assigned to edges within quad patches and how barycentric coordinates are assigned to points within each patch.
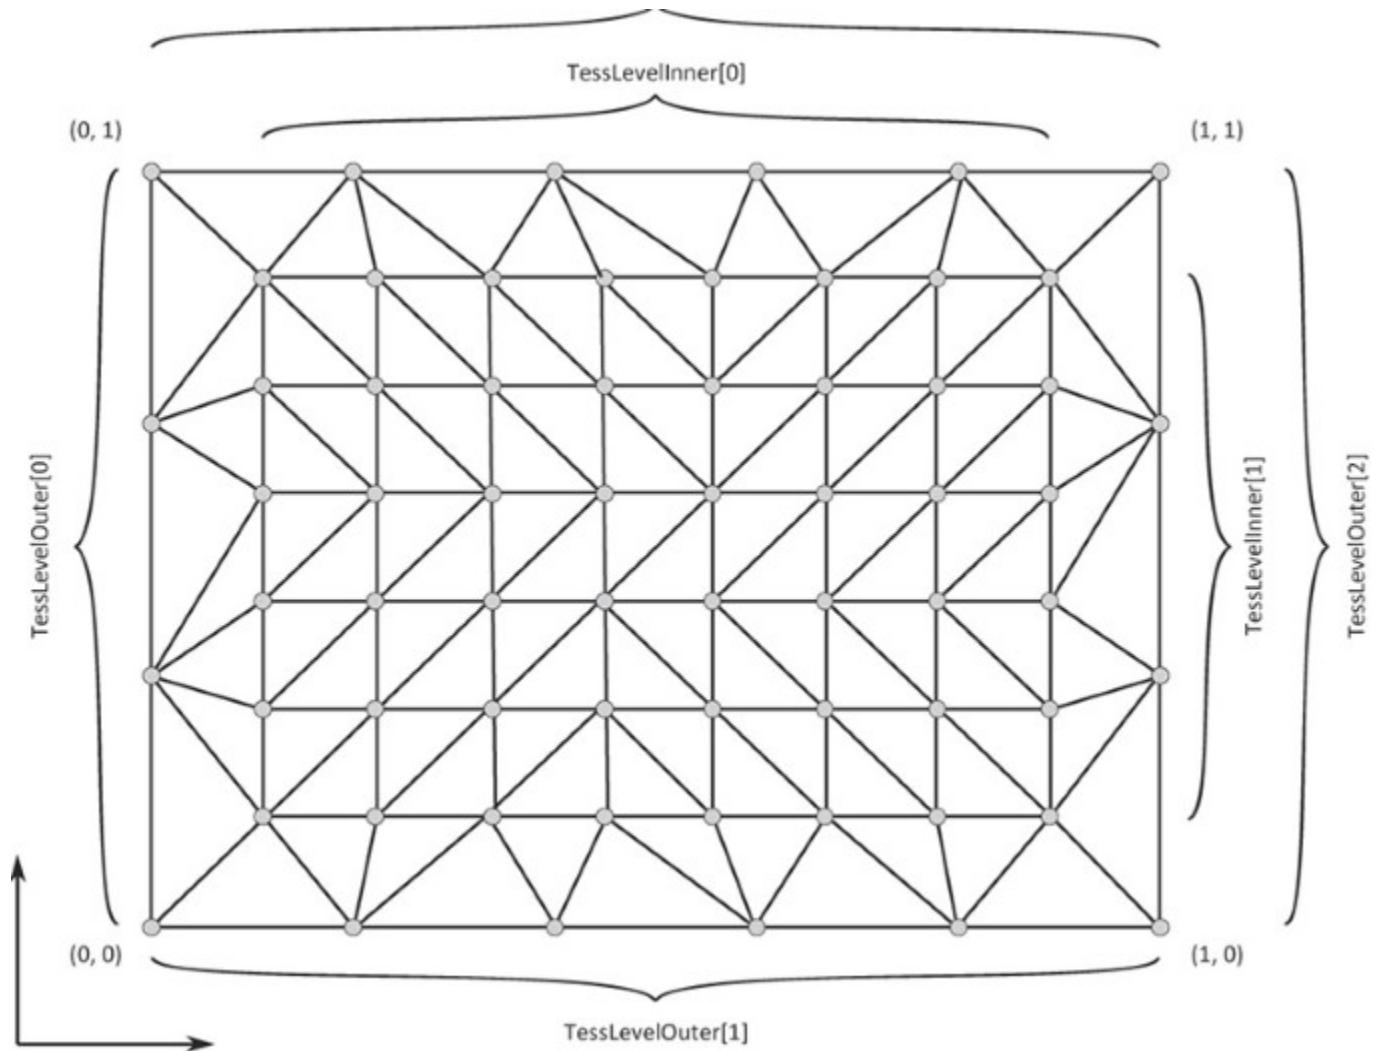
Figure 9.1: Quad Tessellation

As you can see from the figure, the four outer tessellation factors control the level of tessellation along each of the four outer edges of the quad. The *u* and *v* directions in barycentric coordinate space are marked in the figure. For triangle tessellation, the principle is similar, but the assignment of the 3D barycentric coordinate within the triangle is a little different. Figure 9.2 demonstrates.
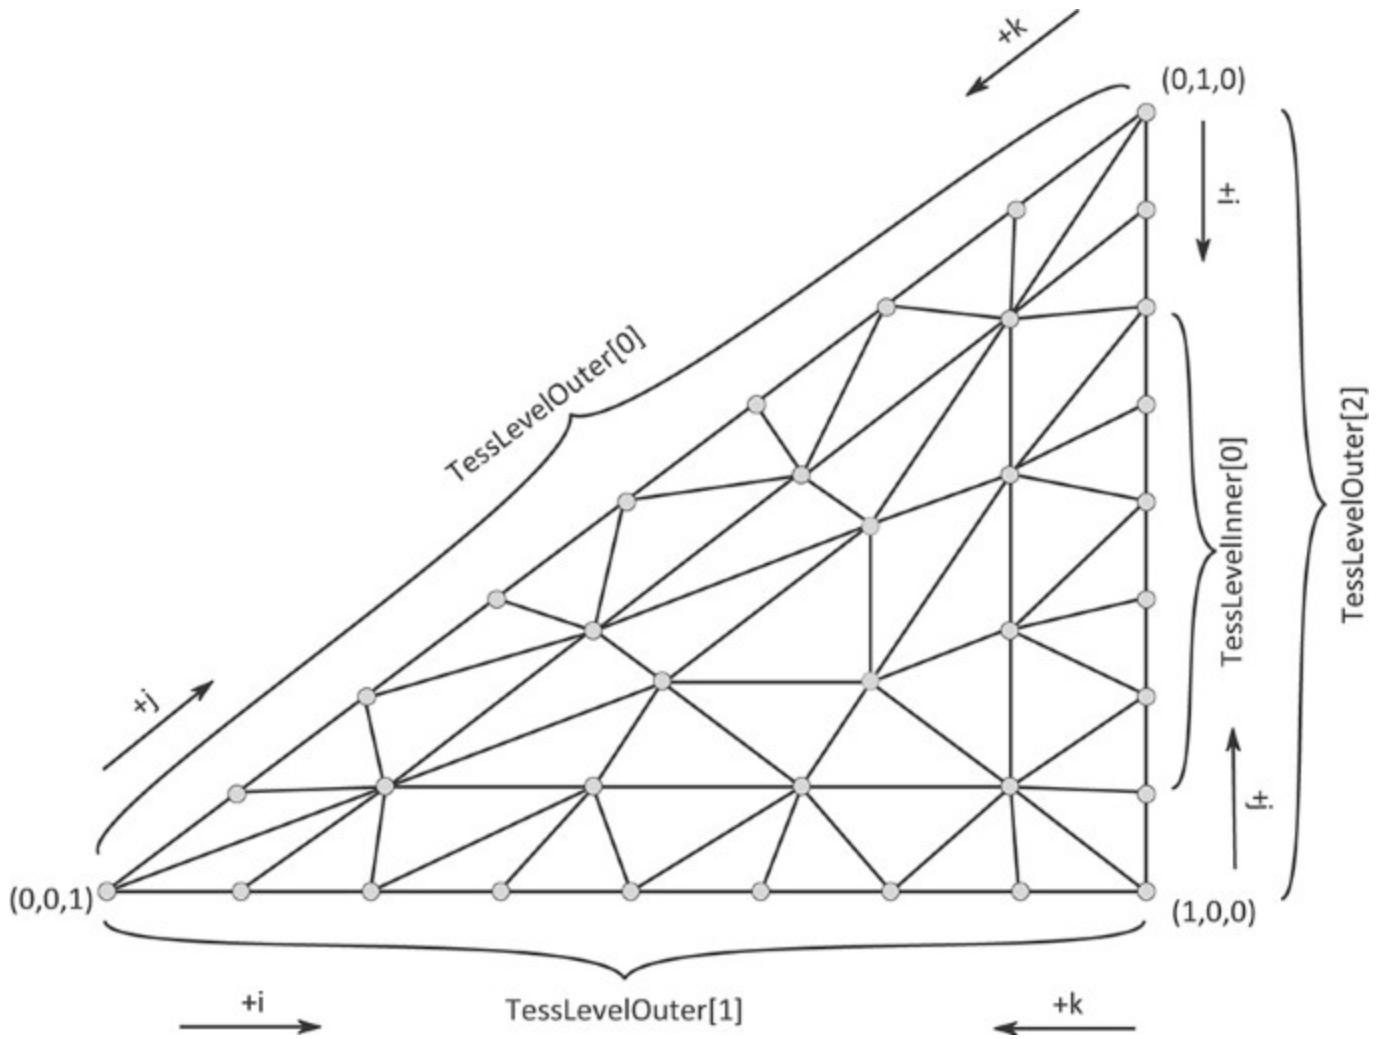
Figure 9.2: Triangle Tessellation.

As you can see in Figure 9.2, for triangle tessellation modes, the three outer tessellation factors control the level of tessellation along the outer edge of the triangular patch. Unlike with quad tessellation, triangle tessellation mode uses just a single tessellation factor, which is applied to the entire patch besides the outermost ring of triangles around its edge.

In addition to the quad and triangle tessellation modes, a special mode known as *isoline* mode allows a patch to be broken down into a series of straight lines. This can be considered to be a special case of quad tessellation mode. In isoline mode, the barycentric coordinates of generated points within the patch are still 2D, but there is no inner tessellation level, and there are only two outer tessellation levels. Figure 9.3 shows how this mode works.
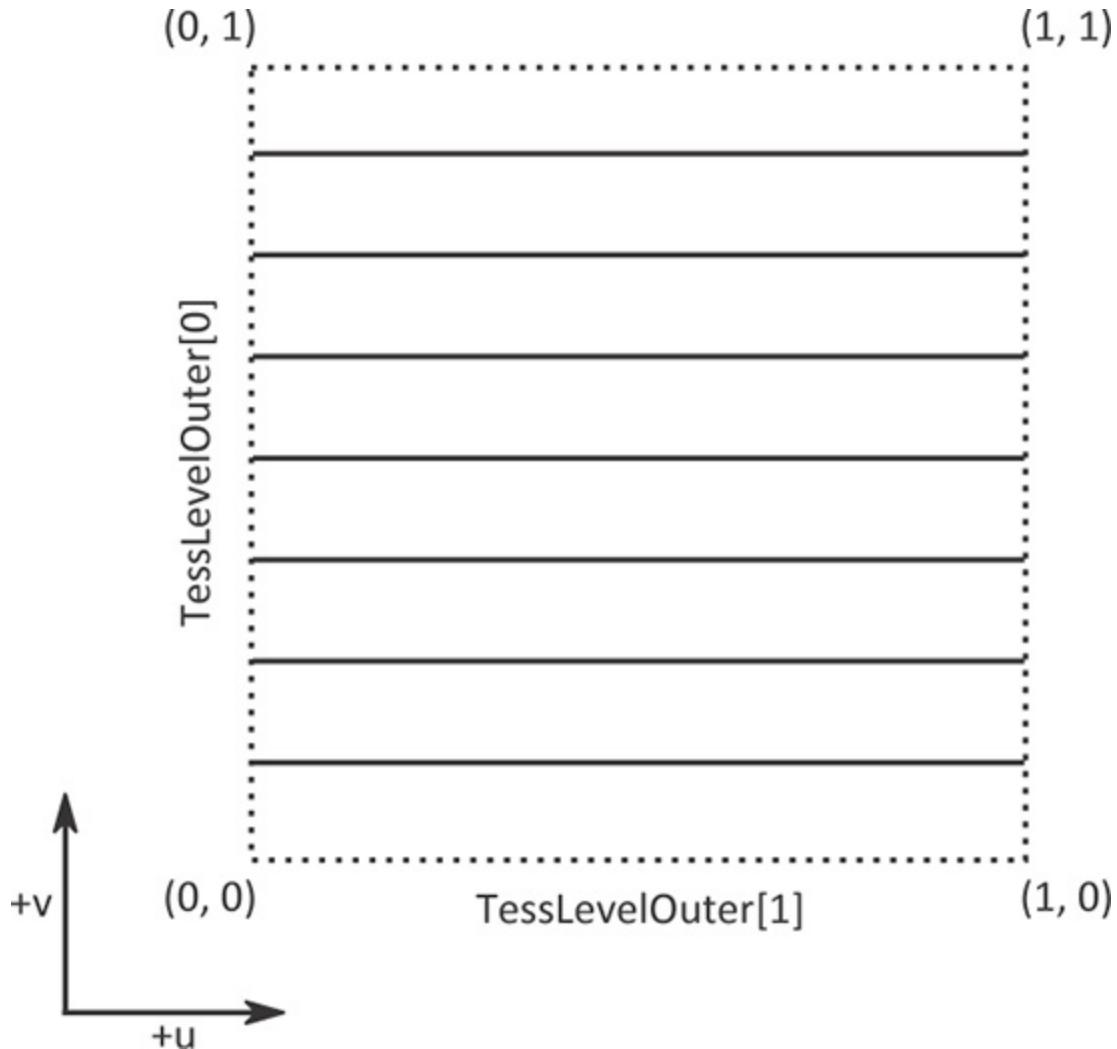
Figure 9.3: Isoline Tessellation

When you set the tessellation mode, one (or both) of the tessellation control or evaluation shaders must include the `OpExecutionMode` instruction with the `Triangles`, `Quads`, or `IsoLines` argument. To generate such a shader from GLSL, use an input layout qualifier in the tessellation evaluation shader, as shown in Table 9.1.

In the table, `%n` represents the index given to the main entry point. As SPIR-V modules can have multiple entry points, it's possible to create a tessellation evaluation shader with an entry point for each mode. Note, however, that the tessellation mode affects the definition of the barycentric coordinate, so care must be taken to interpret it correctly.

In SPIR-V, the tessellation mode instruction can appear in the tessellation control shader, the tessellation evaluation shader, or both, so long as both shaders agree.

In addition to the `Quads` and `Triangles` tessellation modes, which produce triangles, and the `IsoLines` tessellation mode, which produces lines, a special fourth mode, `PointMode`, is supported. As its name suggests, this mode allows the tessellation engine to produce points. To enable this mode, use the `OpExecutionMode` instruction with the `PointMode` argument. Again, this mode can appear in the tessellation evaluation shader, the tessellation control shader, or both, so long as they agree. In GLSL, this mode appears in the tessellation evaluation shader as an input layout qualifier, so that

```
layout (point_mode) in;
```

becomes

```
OpExecutionMode %n PointMode
```

`PointMode` applies *on top of* other tessellation modes such as `Quads` or `Triangles`. In this mode, the patch is tessellated as normal, but rather than being joined, the resulting vertices are sent into the remainder of the pipeline as though they were points. Note that this is subtly different from simply setting the `polygonMode` field of the `VkPipelineRasterizationStateCreateInfo` structure to `VK_POLYGON_MODE_POINT`. In particular, points produced by the tessellator in this mode appear to the geometry shader (if enabled) to be points and are rasterized exactly once, rather than once for each generated primitive in which they appear, as they would be otherwise.

GLSL	SPIR-V
`layout (triangles) in;`	`OpExecutionMode %n Triangles`
`layout (quads) in;`	`OpExecutionMode %n Quads`
`layout (isolines) in;`	`OpExecutionMode %n IsoLines`

Table 9.1: GLSL and SPIR-V Tessellation Modes

Controlling Subdivision

When subdividing the edges of the patches, the tessellator can use one of three strategies to place the split points, which eventually become vertices in the resulting tessellated mesh. This feature allows you to control the appearance of the tessellated patch and particularly to control how the edges of adjacent patches line up. The available modes are

- `SpacingEqual`: The tessellation level assigned to each edge is clamped to the range [1, *maxLevel*] and then rounded to the next higher integer n. The edge is then divided into n segments of equal length in barycentric space.

- `SpacingFractionalEven`: The tessellation level assigned to each edge is clamped to the range [2, *maxLevel*] and then rounded to the nearest *even* integer n. The edge is then subdivided into $n - 2$ segments of equal length, with two additional, shorter segments filling the center region of the edge.

- `SpacingFractionalOdd`: The tessellation level assigned to each edge is clamped to the range [1, *maxLevel* − 1] and then rounded to the nearest *odd* integer n. The edge is then subdivided into $n - 2$ segments of equal length, with two additional, shorter segments filling the center region of the edge.

For both `SpacingFractionalEven` and `SpacingFractionalOdd`, the edge is not tessellated at all if the clamped tessellation level is equal to 1. At levels over 1, the modes produce different visual effects. These effects are shown in Figure 9.4.
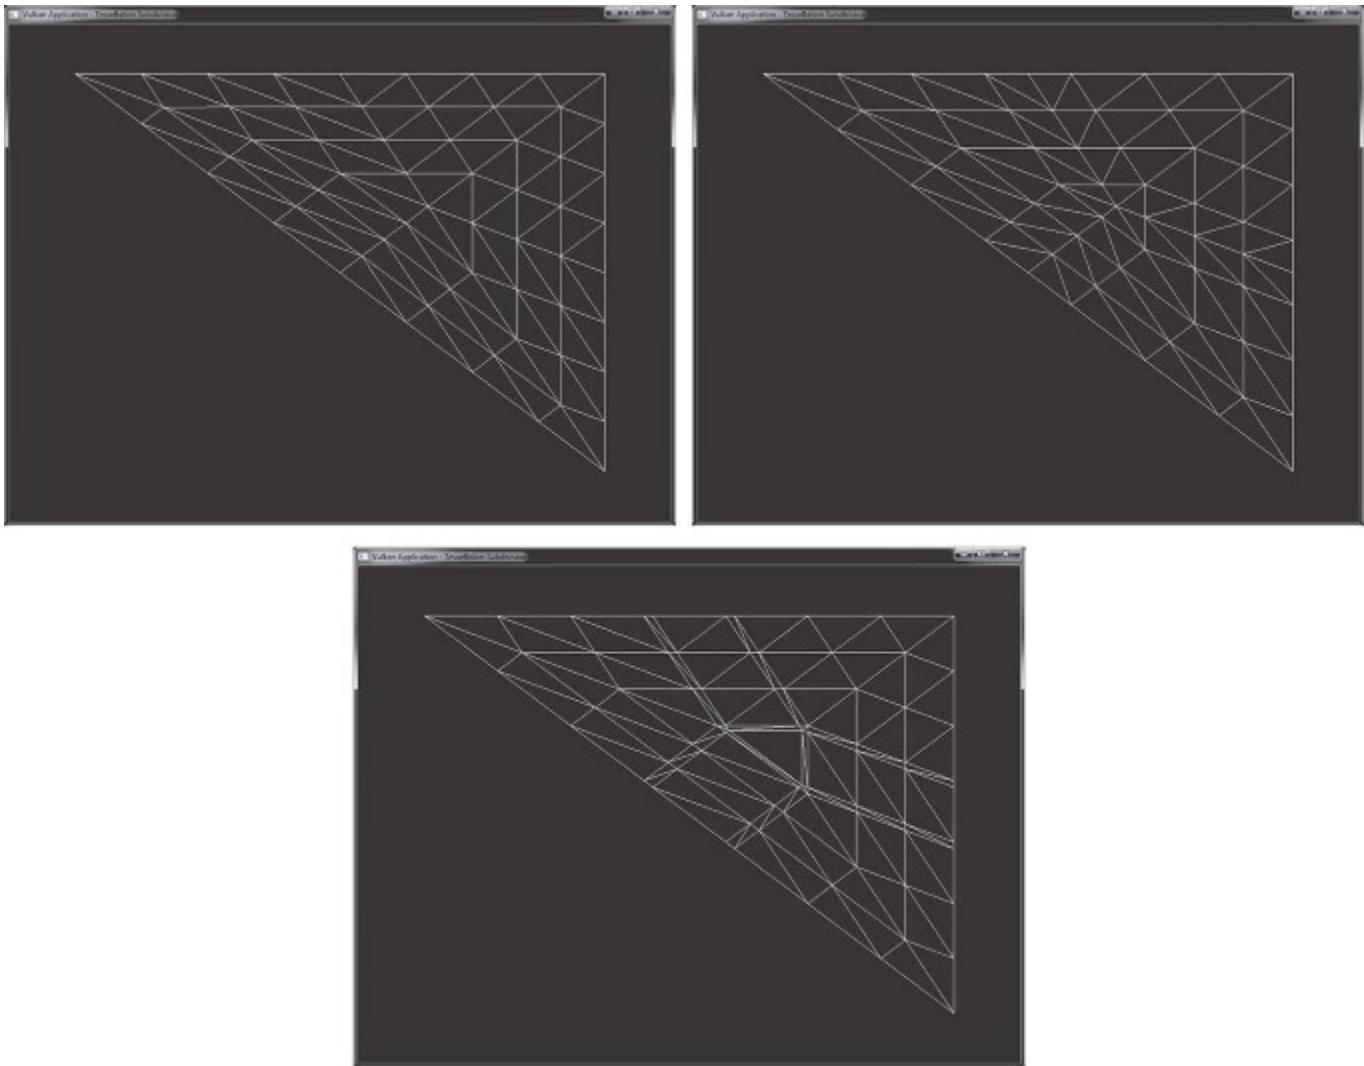
Figure 9.4: Tessellation Spacing Modes

In the top-left image of Figure 9.4, the tessellation mode is set to `SpacingEqual`. As you can see, each outer edge of the tessellated triangle is divided into a number of equal-length segments. The tessellation level in all three images is set to 7.3. This level has been rounded up to 8, and the edge has been divided into that number of segments.

In the top-right image of Figure 9.4, the tessellation mode is `SpacingFractionalEven`. The tessellation level (7.3) has been rounded down to the next-lower even integer (6), and the outer edges of the triangle have been divided into this many equal-length segments. The remaining section of the edge is then divided into two equal-size pieces, producing two shorter segments in the center of the edge. This is most easily seen on the long diagonal edge of the triangle.

Finally, in the bottom image of Figure 9.4, the tessellation mode is set to `SpacingFractionalOdd`. Here, the tessellation level has been rounded down to the next-lower odd number (7). The outer edges of the tessellated triangle are then divided into this many equal-length segments, with the remaining space made up from two smaller segments inserted on either side of the central large segment. Again, this is easiest to see on the long, diagonal edge of the triangle. However, here we see the effect of the two small segments as they produce lines of higher tessellation leading into the center of the tessellated region.

GLSL	SPIR-V
`layout (cw) in;`	`OpExecutionMode %n VertexOrderCw`
`layout (ccw) in;`	`OpExecutionMode %n VertexOrderCcw`

Table 9.2: GLSL and SPIR-V Tessellation Winding Order

Again, the tessellation spacing mode is set using the `OpExecutionMode` instruction in SPIR-V in the tessellation control shader, the evalaution shader, or both, so long as they agree. In GLSL, this instruction is also generated by using an input layout qualifier.

When the tessellation mode is either `Triangles` or `Quads`, the tessellation engine will produce triangles as outputs. The order in which the resulting vertices are processed by the remainder of the pipeline determines which way the triangles face relative to the original patch. The vertices of a triangle are said to appear either in clockwise or counterclockwise order, which is the order in which you would encounter them when traversing the triangle edges in the specified direction while viewing the triangle from the front.

Again, the tessellation winding order is set by using the `OpExecutionMode` instruction in the tessellation control shader, the evaluation shaders, or both, and the equivalent GLSL declaration is an input layout qualifier specified in the tessellation evaluation shader. The GLSL layout qualifier declarations and resulting SPIR-V `OpExecutionMode` instructions are shown in Table 9.2. Once again, in the table, the SPIR-V notation of `%n` indicates the index of the entry point to which the `OpExecutionMode` instruction applies.

Tessellation Variables

Each patch processed by the tessellation control shader has a fixed number of control points. This number is set by using the `patchControlPoints` member of the `VkPipelineTessellationStateCreateInfo` structure used to create the pipeline, as discussed earlier. Each control point is represented as a vertex passed into the pipeline by one of the drawing commands. The vertex shader processes the vertices one at a time before they are passed in groups to the tessellation control shader, which has access to all of the vertices making up a patch.

The tessellation evaluation shader also has access to all of the control points making up a patch, but the number of control points in the patch can be changed between the tessellation control shader and the tessellation evaluation shader. The number of control points that are passed from the tessellation control shader to the tessellation evaluation shader is set by using a SPIR-V `OutputVertices` argument to the `OpExecutionMode` instruction applied to the entry point. Again, this can appear in the tessellation control shader, the tessellation evaluation shader, or both, so long as they agree. This operation takes an integer constant (or specialization constant).

In GLSL, the number of control points passed from the tessellation control shader to the tessellation evaluation shader is specified by using an output layout qualifier in the tessellation control shader. For example, the GLSL declaration

```
layout (vertices = 9) out;
```

becomes

```
OpExecutionMode %n OutputVertices 9}
```

The inner and outer tessellation levels are set by using the tessellation control shader. This is accomplished in SPIR-V by decorating variables in the tessellation control shader with the `TessLevelInner` and `TessLevelOuter` decorations, respectively.

The variable representing the outer tessellation levels is an array of four floating-point values, all of which are used for `Quads` tessellation mode, the first three of which are used in `Triangles` tessellation mode, and only the first two of which are significant in `IsoLines` tessellation mode.

The variable representing the inner tessellation levels is an array of two floating-point values. In `Quads` tessellation mode, the two values control the inner tessellation level in the u and v domains. In `Triangles` tessellation mode, the first element of the array sets the tessellation mode for the center patch, and the second is ignored. In `IsoLines` mode, there is no inner tessellation level.

The inner and outer tessellation levels appear in GLSL as the `gl_TessLevelInner` and `gl_TessLevelOuter` built-in variables, respectively. When you are using these variables in a GLSL tessellation control shader, the compiler will generate the appropriate SPIR-V variable declaration and decorate it accordingly.

The maximum tessellation level that can be used in a Vulkan pipeline is device-dependent. You can determine the maximum tessellation level that the device supports by checking the `maxTessellationGenerationLevel` field of the device's `VkPhysicalDeviceLimits` structure, which can be retrieved by calling **vkGetPhysicalDeviceProperties()**. The minimum guaranteed limit for `maxTessellationGenerationLevel` is 64, but some devices may support higher levels. However, most applications will not need higher levels of tessellation than this, and in this case, there is no reason to query the limit.

Consider the GLSL tessellation control shader shown in , which simply sets the inner and outer tessellation levels of a patch to some hard-coded constants. This is not a complete tessellation control shader, but it is sufficient to demonstrate how tessellation assignments are translated from GLSL to SPIR-V.

Listing 9.1: Trivial Tessellation Control Shader (GLSL)

```
#version 450 core

layout (vertices = 1) out;

void main(void)
{
    gl_TessLevelInner[0] = 7.0f;
    gl_TessLevelInner[1] = 8.0f;

    gl_TessLevelOuter[0] = 3.0f;
    gl_TessLevelOuter[1] = 4.0f;
    gl_TessLevelOuter[2] = 5.0f;
    gl_TessLevelOuter[3] = 6.0f;
}
```

After compilation into SPIR-V, the shader shown in [Listing 9.1](#) becomes the (substantially longer) SPIR-V shader shown in [Listing 9.2](#). This listing is the raw output of the SPIR-V disassembler with comments added by hand.

Listing 9.2: Trivial Tessellation Control Shader (SPIR-V)

Click here to view code image

```
;; Require tessellation capability; import GLSL450 constructs.
          OpCapability Tessellation
     %1 = OpExtInstImport "GLSL.std.450"
          OpMemoryModel Logical GLSL450
;; Define "main" as the entry point for a tessellation control shader.
          OpEntryPoint TessellationControl %5663 "main" %3290 %5448
;; Number of patch output vertices = 1
          OpExecutionMode %5663 OutputVertices 1
;; Decorate the tessellation level variables appropriately.
          OpDecorate %3290 Patch
          OpDecorate %3290 BuiltIn TessLevelInner
          OpDecorate %5448 Patch
          OpDecorate %5448 BuiltIn TessLevelOuter
;; Declare types used in this shader.
      %8 = OpTypeVoid
   %1282 = OpTypeFunction %8
     %13 = OpTypeFloat 32
     %11 = OpTypeInt 32 0
;; This is the declaration of the gl_TessLevelInner[2] variable.
   %2576 = OpConstant %11 2
    %549 = OpTypeArray %13 %2576
   %1186 = OpTypePointer Output %549
   %3290 = OpVariable %1186 Output
     %12 = OpTypeInt 32 1
   %2571 = OpConstant %12 0
   %1330 = OpConstant %13 7
    %650 = OpTypePointer Output %13
   %2574 = OpConstant %12 1
   %2807 = OpConstant %13 8
;; Declare the gl_TessLevelOuter[4] variable.
   %2582 = OpConstant %11 4
    %609 = OpTypeArray %13 %2582
   %1246 = OpTypePointer Output %609
   %5448 = OpVariable %1246 Output
;; Declare constants used for indexing into our output arrays and the
;; values written into those arrays.
   %2978 = OpConstant %13 3
   %2921 = OpConstant %13 4
   %2577 = OpConstant %12 2
   %1387 = OpConstant %13 5
   %2580 = OpConstant %12 3
   %2864 = OpConstant %13 6
;; Start of the main function
   %5663 = OpFunction %8 None %1282
  %23934 = OpLabel
```

```
;; Declare references to elements of the output arrays and write constants
;; into them.
   %6956 = OpAccessChain %650 %3290 %2571
            OpStore %6956 %1330
  %19732 = OpAccessChain %650 %3290 %2574
            OpStore %19732 %2807
  %19733 = OpAccessChain %650 %5448 %2571
            OpStore %19733 %2978
  %19734 = OpAccessChain %650 %5448 %2574
            OpStore %19734 %2921
  %19735 = OpAccessChain %650 %5448 %2577
            OpStore %19735 %1387
  %23304 = OpAccessChain %650 %5448 %2580
            OpStore %23304 %2864
;; End of main
            OpReturn
            OpFunctionEnd
```

Tessellation control shaders execute a single invocation for each *output* control point defined in the patch. All of those invocations have access to all of the data associated with the input control points for the patch. As a result, the input variables to the tessellation control shader are defined as arrays. As discussed, the number of input and output control points in a patch do not have to be equal. In addition to the tessellation level outputs, the tessellation shader can define more outputs to be used for per-control point and per-patch data.

The per-control-point output from the tessellation control shader is declared as arrays whose sizes match the number of output control points in the patch. There is one tessellation control shader invocation per output control point, so there is one entry in each of the output arrays that corresponds to each of those invocations. The per-control-point outputs can only be written by the corresponding invocation. The index of the shader invocation within the patch is available as a built-in variable, which can be accessed by declaring an integer variable decorated by the SPIR-V InvocationId built-in. In GLSL, this variable is declared as the gl_InvocationID built-in variable. This variable must be used to index into the output arrays.

[Listing 9.3](#) shows how to declare output variables in a GLSL tessellation control shader, and [Listing 9.4](#) shows how that shader is translated into SPIR-V. Again, [Listing 9.3](#) is not a complete tessellation control shader and, while legal, will not produce any useful output. Also, the SPIR-V shader shown in [Listing 9.4](#) has been commented by hand.

Listing 9.3: Declaring Outputs in Tessellation Control Shaders (GLSL)

Click here to view code image

```glsl
#version 450 core

layout (vertices = 4) out;
out float o_outputData[4];

void main(void)
{
    o_outputData[gl_InvocationId] = 19.0f;
}
```

Listing 9.4: Declaring Outputs in Tessellation Control Shaders (SPIR-V)

```
;; Declare a tessellation control shader.
        OpCapability Tessellation
   %1 = OpExtInstImport "GLSL.std.450"
        OpMemoryModel Logical GLSL450
        OpEntryPoint TessellationControl %5663 "main" %3227 %4585
;; 4 output vertices per patch, declare InvocationId built-in
        OpExecutionMode %5663 OutputVertices 4
        OpDecorate %4585 BuiltIn InvocationId
   %8 = OpTypeVoid
%1282 = OpTypeFunction %8
  %13 = OpTypeFloat 32
  %11 = OpTypeInt 32 0
%2582 = OpConstant %11 4
 %549 = OpTypeArray %13 %2582
%1186 = OpTypePointer Output %549
%3227 = OpVariable %1186 Output
  %12 = OpTypeInt 32 1
;; This declares the InvocationId input.
 %649 = OpTypePointer Input %12
%4585 = OpVariable %649 Input
  %37 = OpConstant %13 19
 %650 = OpTypePointer Output %13
;; Beginning of main
%5663 = OpFunction %8 None %1282
%24683 = OpLabel
;; Load the invocation ID.
%20081 = OpLoad %12 %4585
;; Define a reference to output variable.
%13546 = OpAccessChain %650 %3227 %20081
;; Store into it at invocation ID.
        OpStore %13546 %37
;; End of main
        OpReturn
        OpFunctionEnd
```

The output variables declared in the tessellation control shader are made available in the tessellation evaluation shader as inputs. The total number of components that can be produced per output vertex by a tessellation control shader is device-dependent and is determined by inspecting the `maxTessellationControlTotalOutputComponents` member of the device's `VkPhysicalDeviceLimits` structure, which you can obtain by calling **`vkGetPhysicalDeviceProperties()`**. This is guaranteed to be at least 2,048 components. Some of these components may be used per vertex, and some may be applied to the patch, as discussed in the next section.

Of course, the per-vertex limit is the limit that applies to variables passed for each vertex from tessellation control to tessellation evaluation shader. The total number of components that a tessellation evaluation shader can receive as input from the tessellation control shader is contained in the `maxTessellationEvaluationInputComponents` field of the

`VkPhysicalDeviceLimits` structure. This limit is at least 64 but could be higher, depending on the device.

Patch Variables

Although normal output variables in tessellation control shaders are instantiated as arrays corresponding to the output control points, sometimes specific pieces of data are needed that apply everywhere in the patch. These pieces of data can be declared as *patch* outputs. Patch outputs serve two purposes:

- They store per-patch data and pass it from tessellation control shader to tessellation evaluation shader.
- They allow data sharing between tessellation control shader invocations within a single patch.

Within the group of tessellation control shader invocations corresponding to a single patch, patch outputs are actually readable as well as writable. If other invocations in the same patch have written to a patch output, it is possible to read the data they have placed there.

To declare a patch output variable in GLSL, use the **patch** qualifier on the variable declaration. This declaration is translated into the `Patch` decoration on the variable subsequently declared in the SPIR-V shader. For example,

```
patch out myVariable;
```

becomes

```
OpName %n "myVariable"
OpDecorate %n Patch
```

where `%n` is the identifier assigned to the `myVariable` variable.

Because all of the tessellation control shader invocations corresponding to a single patch may be running in parallel, perhaps at different rates, simply writing to one patch variable and then reading from another patch variable won't produce well-defined results. If some invocations get ahead of others while processing a patch, the invocations that are behind won't "see" the results of writes performed by the other invocations that haven't reached the write yet.

To synchronize the invocations within a patch and ensure that they all reach the same place at the same time, we can use the `OpControlBarrier` instruction, which can synchronize control flow of invocations within the tessellation control shader. Further, to ensure that the write to the patch variable is made visible to other invocations within the same patch, we also need to include either an `OpMemoryBarrier` instruction or further memory semantics in the `OpControlBarrier` instruction.

In GLSL, these instructions can be generated by calling the `barrier()` built-in function inside the tessellation control shader. When called, the GLSL compiler generates an `OpMemoryBarrier` instruction to force memory coherency across the set of tessellation control shader invocations, and then it generates an `OpControlBarrier` instruction to synchronize their control flow. After these instructions have executed, tessellation control shader invocations can read data written to patch variables by other invocations in the same patch.

If any one of the tessellation levels written by the tessellation control shader is 0.0 or a floating-point NaN, then the entire patch is discarded. This provides a mechanism for the tessellation control shader

to programmatically throw out patches that it determines will not contribute to the output image. For example, if the maximum deviation from a plane for a displacement map is known, then the tessellation control shader can inspect a patch, determine whether all of the geomery that would result from tessellating that patch would face away from the viewer, and cull the patch. If such a patch were to be passed through by the tessellation control shader, then it would be tessellated, the tessellation evaluation shader would run, and all the resulting triangles would individually be culled by subsequent stages in the pipeline. This would be very inefficient.

Tessellation Evaluation Shaders

After the tessellation control shader runs and passes tessellation factors to the fixed-function tessellation unit, it generates new vertices inside the patch and assigns them barycentric coordinates in patch space. Each new vertex generates an invocation of the tessellation evaluation shader, to which the barycentric coordinate of the vertex is passed.

For `IsoLines` and `Quads` tessellation mode, these are 2D coordinates. For `Triangles` tessellation mode, these are 3D coordinates. Regardless of the mode, they are delivered to the tessellation evaluation shader via built-in variables.

In GLSL, this declaration corresponds to the `gl_TessCoord` built-in variable. When this variable is used, the GLSL compiler automatically generates the appropriate variable declarations and decorations in SPIR-V.

In SPIR-V, the result is the declaration of a three-element vector of floating-point values, decorated with the `TessCoord` decoration. Note that this is always a three-element array, even when the tessellation mode calls for a 2D barycentric coordinate (`IsoLines` or `Quads`). In those modes, the third component of the array is simply zero.

For example, the minimal tessellation evaluation shader in [Listing 9.5](#) becomes the SPIR-V shader in [Listing 9.6](#).

Listing 9.5: Accessing `gl_TessCoord` in Evaluation Shader (GLSL)

Click here to view code image

```
#version 450 core

layout (quads) in;

void main(void)
{
    gl_Position = vec4(gl_TessCoord, 1.0);
}
```

Listing 9.6: Accessing `gl_TessCoord` in Evaluation Shader (SPIR-V)

Click here to view code image

```
;; This is a GLSL 450 tessellation shader; enable capabilities.
         OpCapability Tessellation
         OpCapability TessellationPointSize
         OpCapability ClipDistance
         OpCapability CullDistance
```

```
           %1 = OpExtInstImport "GLSL.std.450"
                OpMemoryModel Logical GLSL450
;; Declare the entry point and decorate it appropriately.
                OpEntryPoint TessellationEvaluation %5663 "main" %4930 %3944
                OpExecutionMode %5663 Quads
                OpExecutionMode %5663 SpacingEqual
                OpExecutionMode %5663 VertexOrderCcw
;; Declare GLSL built-in outputs.
                OpMemberDecorate %2935 0 BuiltIn Position
                OpMemberDecorate %2935 1 BuiltIn PointSize
                OpMemberDecorate %2935 2 BuiltIn ClipDistance
                OpMemberDecorate %2935 3 BuiltIn CullDistance
                OpDecorate %2935 Block
;; This is the decoration of gl_TessCoord.
                OpDecorate %3944 BuiltIn TessCoord
           %8 = OpTypeVoid
        %1282 = OpTypeFunction %8
          %13 = OpTypeFloat 32
          %29 = OpTypeVector %13 4
          %11 = OpTypeInt 32 0
        %2573 = OpConstant %11 1
         %554 = OpTypeArray %13 %2573
        %2935 = OpTypeStruct %29 %13 %554 %554
         %561 = OpTypePointer Output %2935
        %4930 = OpVariable %561 Output
          %12 = OpTypeInt 32 1
        %2571 = OpConstant %12 0
;; Vector of 3 components, as input, decorated with BuiltIn TessCoord
          %24 = OpTypeVector %13 3
         %661 = OpTypePointer Input %24
        %3944 = OpVariable %661 Input
         %138 = OpConstant %13 1
         %666 = OpTypePointer Output %29
        %5663 = OpFunction %8 None %1282
       %24987 = OpLabel
;; Read from gl_TessCoord.
       %17674 = OpLoad %24 %3944
;; Extract the three elements.
       %22014 = OpCompositeExtract %13 %17674 0
       %23496 = OpCompositeExtract %13 %17674 1
        %7529 = OpCompositeExtract %13 %17674 2
;; Construct the new vec4.
       %18260 = OpCompositeConstruct %29 %22014 %23496 %7529 %138
;; Write to gl_Postion.
       %12055 = OpAccessChain %666 %4930 %2571
                OpStore %12055 %18260
;; End of main
                OpReturn
                OpFunctionEnd
```

The outputs declared in the tessellation evaluation shader are fed to the next stage of the pipeline.
When a geometry shader is enabled, it receives its input from the tessellation evaluation shader;
otherwise, the tessellation evaluation shader outputs are used to feed interpolation and are

subsequently passed to the fragment shader. The total amount of data that a tessellation evaluation shader can produce is device-dependent and can be determined by checking the `maxTessellationEvaluationOutputComponents` field of the device's `VkPhysicalDeviceLimits` structure.

Tessellation Example: Displacement Mapping

To tie all of this together, we'll walk through a simple but complete example of using tessellation to implement *displacement mapping*. Displacement mapping is a common technique to add detail to a surface by offsetting along its normal by using a texture. To implement this, we'll take a patch with four control points, each having a position and a normal. We will transform their positions and normals into world space in the vertex shader; then, in the tessellation control shader, we will compute their view space positions and set their tessellation levels. The tessellation control shader then passes the world-space positions to the tessellation evaluation shader.

In the tessellelation evaluation shader, we take the barycentric coordinates generated by the tessellator and use them to compute an interpolated normal and world space position, fetch from a texture, and then displace the computed vertex coordinate along the computed normal by a value derived from the texture. This final world-space coordinate is transformed into view space by using the world-to-view matrix, which is stored in the same buffer as our object-to-world matrix used in our vertex shader.

To set this up, we'll use a texture to store our displacement map. In addition to our texture resources, we'll use push constants to communicate with our shaders. For the purposes of this example, we'll use a simple push constant to send a transformation matrix into the tessellation control shader.

The next two push constants are floating-point values. The first is used in the tessellation evaluation shader to scale the level of tessellation applied to the patch. The second is used in the tessellation evaluation shader to scale the amount of displacement applied to each vertex. The values read from the texture are normalized to the range [0.0, 1.0], so an external tessellation scale allows us to use this range optimally.

The description of the `VkDescriptorSetLayoutCreateInfo` used to create this descriptor set layout is shown in [Listing 9.7](#).

Listing 9.7: Descriptor Setup for Displacement Mapping

Click here to view code image

```
struct PushConstantData
{
    vmath::mat4 mvp_matrix;
    float       displacement_scale;
};

static const VkDescriptorSetLayoutBinding descriptorBindings[] =
{
    // Only binding is a sampled image with combined sampler.
    {
        0,
        VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER,
        1,
        VK_SHADER_STAGE_TESSELLATION_EVALUATION_BIT,
```

```
            nullptr
        }
};

static const VkDescriptorSetLayoutCreateInfo descriptorSetLayoutCreateInfo
=
{
    VK_STRUCTURE_TYPE_DESCRIPTOR_SET_LAYOUT_CREATE_INFO, nullptr,
    0,
    vkcore::utils::arraysize(descriptorBindings),
    descriptorBindings
};
vkCreateDescriptorSetLayout(getDevice(),
                            &descriptorSetLayoutCreateInfo,
                            nullptr,
                            &m_descriptorSetLayout);

// Define a push constant range.
static const VkPushConstantRange pushConstantRange[] =
{
    // Push data into the evaluation shader.
    {
        VK_SHADER_STAGE_TESSELLATION_EVALUATION_BIT,
        0,
        sizeof (PushConstantData)
    }
};
VkPipelineLayoutCreateInfo pipelineLayoutCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_LAYOUT_CREATE_INFO, nullptr,
    0,
    1,
    &m_descriptorSetLayout,
    vkcore::utils::arraysize(pushConstantRange),
    pushConstantRange
};

result = vkCreatePipelineLayout(getDevice(),
                                &pipelineLayoutCreateInfo,
                                &nullptr,
                                &m_pipelineLayout);
```

As the geometry for the tessellated quad is simply four vertices arranged in a square, we're not going to use vertex buffers in this example, but instead we'll programmatically generate the object space positions in the vertex shader. The GLSL vertex shader used in this example is shown in Listing 9.8.

Listing 9.8: Vertex Shader for Displacement Mapping

Click here to view code image

```
#version 450 core

void main(void)
```

```
{
    float x = float(gl_VertexIndex & 1) - 0.5f;
    float y = float(gl_VertexIndex & 2) * 0.5f - 0.5f;

    gl_Position = vec4(x, y, 0.0f, 1.0f);
}
```

The output of the vertex shader is simply a quad of side length 1, centered on the origin. This is passed into the tessellation control shader shown in <u>Listing 9.9</u>, which sets the tessellation factors and passes the position of the vertices through otherwise unchanged.

Listing 9.9: Tessellation Control Shader for Displacement Mapping

<u>**Click here to view code image**</u>

```
#version 450 core

layout (vertices = 4) out;

void main(void)
{
    if (gl_InvocationID == 0)
    {
        gl_TessLevelInner[0] = 64.0f;
        gl_TessLevelInner[1] = 64.0f;

        gl_TessLevelOuter[0] = 64.0f;
        gl_TessLevelOuter[1] = 64.0f;
        gl_TessLevelOuter[2] = 64.0f;
        gl_TessLevelOuter[3] = 64.0f;
    }
    gl_out[gl_InvocationID].gl_Position =
gl_in[gl_InvocationID].gl_Position;
}
```

Note that the tessellation control shader shown in <u>Listing 9.9</u> unconditionally sets the tessellation factors to 64, which is the minimum required tessellation level that must be supported by Vulkan implementations. This is a very high amount of tessellation, and most applications will not need this level. Rather, we would divide the patch into several smaller patches and then use a smaller, probably different tessellation level for each of them. The output of the tessellation control shader is then passed to the tessellation evaluation shader shown in <u>Listing 9.10</u>.

Listing 9.10: Tessellation Evaluation Shader for Displacement Mapping

<u>**Click here to view code image**</u>

```
#version 450 core

layout (quads, fractional_odd_spacing) in;

layout (push_constant) uniform push_constants_b
{
    mat4 mvp_matrix;
```

```
    float displacement_scale;
} push_constants;

layout (set = 0, binding = 0) uniform sampler2D texDisplacement;

void main(void)
{
    vec4 mid1 = mix(gl_in[0].gl_Position, gl_in[1].gl_Position,
gl_TessCoord.x);
    vec4 mid2 = mix(gl_in[2].gl_Position, gl_in[3].gl_Position,
gl_TessCoord.x);
    vec4 pos = mix(mid1, mid2, gl_TessCoord.y);
    float displacement = texture(texDisplacement, gl_TessCoord.xy).x;

    pos.z = displacement * push_constants.displacement_scale;

    gl_Position = push_constants.mvp_matrix * pos;
}
```

The shader in uses the content of `gl_TessCoord.xy` to interpolate the position of the vertices that land at the four corners of the quad patch produced by the tessellation control shader. The final position, `pos`, is simply a weighted average of the four corners. Additionally, `gl_TessCoord.xy` is used as a texture coordinate to sample from the displacement map texture, `texDisplacement`.

This displacement value is scaled by the `displacement_scale` push constant, which allows our application to vary the amount of displacement applied to the mesh. We know that our vertex and tessellation control shaders set up our patch in the *x-y* plane and that the *z* value for all of our tessellated points will therefore be zero. Overwriting the *z* component with displacement produces a patch that is like a landscape with the *z* direction being up. We transform this through our final transformation matrix to produce a vertex in view space.

The rest of the graphics pipeline setup for the displacement mapping example is similar to other applications shown thus far. However, because tessellation is enabled, we must set the `pTessellationState` member of our `VkGraphicsPipelineCreateInfo` to the address of a valid `VkPipelineTessellationStateCreateInfo` structure. This structure is extremely simple and is used only to set the number of control points in our patch. The `VkPipelineTessellationStateCreateInfo` for this application is shown in .

Listing 9.11: Tessellation State Creation Information

Click here to view code image

```
VkPipelineTessellationStateCreateInfo tessellationStateCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_TESSELLATION_STATE_CREATE_INFO,  // sType
    nullptr,                                                    // pNext
    0,                                                          // flags
    4                                                // patchControlPoints
};
```

The result of running this application is shown in Figure 9.5. As you can see in the figure, we have set the polygon mode to `VK_POLYGON_MODE_LINE`, which allows us to see the structure of the tessellated patch.
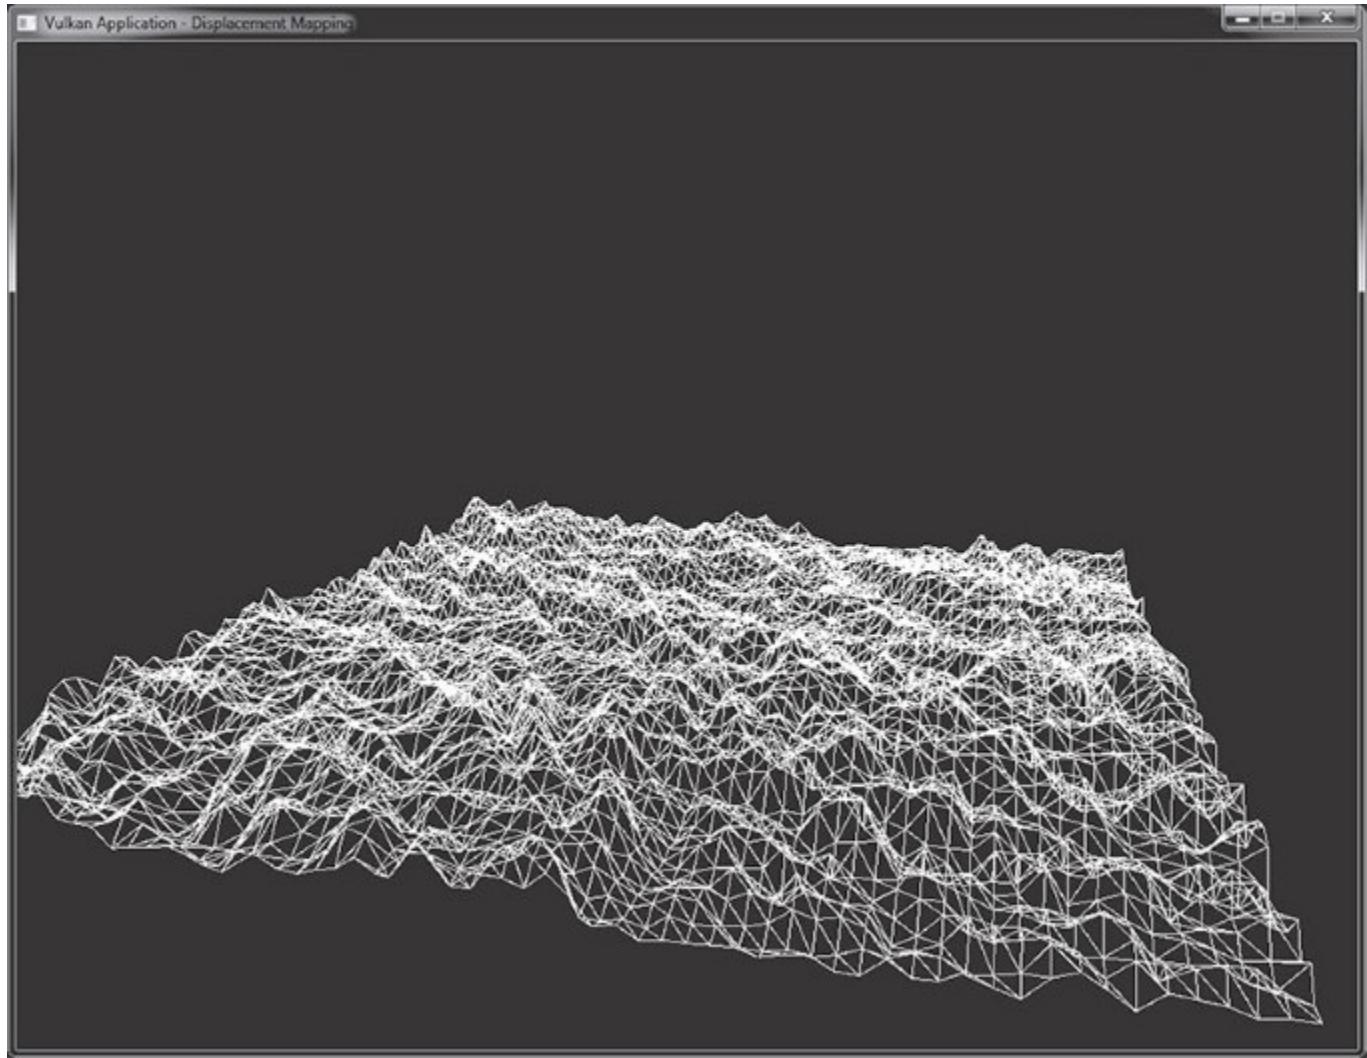
Figure 9.5: Result of Tessellated Displacement Mapping

For all of the geometric density shown in Figure 9.5, remember that no real geometry was passed into the graphics pipeline. The four corners of the patch are programmatically generated by the vertex shader, and no real vertex attributes are used. All of the triangles in the patch are generated by the tessellation engine, and their positions are computed by using a texture. By applying shading and detail textures to the surface of the patch, you can simulate even more apparent detail. Depending on the performance of the system, the tessellation levels chosen by the tessellation control shader can be scaled to balance performance and visual quality. Tessellation is therefore a very effective way to introduce geometric density to a scene without needing to regenerate mesh data or produce multiple versions of 3D art assets.

Geometry Shaders

The geometry shader is the final stage in the front-end, geometry-processing part of the pipeline. When enabled, it runs immediately after the tessellation evaluation shader if tessellation is enabled and after the vertex shader if tessellation is not in use. A geometry shader is unique in that a single invocation of the geometry shader processes an entire primitive. Further, the geometry shader is the only shader stage that can see adjacency primitives.[1] Finally, the geometry shader is special in that it can both destroy and programmatically create new geometry.

> 1. The vertex shader can see the additional vertices that make up the adjacency primitive, but it has no knowledge of whether any particular vertex is part of the main primitive or the adjacency information.

To enable the geometry shader stage, include a geometry shader in the pipeline by including an instance of the `VkPipelineShaderStageCreateInfo` structure describing a geometry shader in the array of stages passed through `pStages` in the `VkGraphicsPipelineCreateInfo` structure. Support for geometry shaders is optional, so before creating a pipeline containing a geometry shader, your application should check for support by inspecting the `geometryShader` member of the `VkPhysicalDeviceFeatures` structure returned from a call to **`vkGetPhysicalDeviceFeatures()`**, and then set that same member to `VK_TRUE` in the `VkDeviceCreateInfo` structure passed to **`vkCreateDevice()`**.

A geometry shader is created in SPIR-V by using the `Geometry` execution model with the `OpEntryPoint` instruction. A geometry shader must include the following pieces of information:

- The input primitive type, which must be one of points, lines, triangles, lines with adjacency, or triangles with adjacency

- The output primitive type, which must be one of points, line strips, or triangle strips

- The maximum number of vertices expected to be produced by a single invocation of the geometry shader

All of these attributes are specified as arguments to `OpExecutionMode` instructions in the SPIR-V shader. In GLSL, the first is specified by using an input layout qualifier, and the second two are specified by using output layout qualifiers. The shader shown in Listing 9.12 is a minimal GLSL geometry shader, which is legal, but it will throw away all of the geometry that passes down the pipe.

Listing 9.12: Minimal Geometry Shader (GLSL)

Click here to view code image

```
#version 450 core

layout (triangles) in;
layout (triangle_strip) out;
layout (max_vertices = 3) out;

void main(void)
{
    // Do nothing.
}
```

As you can see, Listing 9.12 contains only the required input and output layout definitions, while its `main` function is empty. When compiled to SPIR-V, the shader shown in Listing 9.13 is produced.

Listing 9.13: Minimal Geometry Shader (SPIR-V)

```
;; This is a geometry shader written in GLSL 450.
        OpCapability Geometry
   %1 = OpExtInstImport "GLSL.std.450"
        OpMemoryModel Logical GLSL450
;; Declare the main entry point.
        OpEntryPoint Geometry %5663 "main"
;; Triangles input, triangle strip output, maximum vertex count is 3.
        OpExecutionMode %5663 Triangles
        OpExecutionMode %5663 Invocations 1
        OpExecutionMode %5663 OutputTriangleStrip
        OpExecutionMode %5663 OutputVertices 3
;; Start of main
   %8 = OpTypeVoid
%1282 = OpTypeFunction %8
%5663 = OpFunction %8 None %1282
%16103 = OpLabel
;; End of empty main
        OpReturn
        OpFunctionEnd
```

In Listing 9.13, the `main` entry point is decorated using three `OpExecutionMode` instructions. The first uses `Triangles` to indicate that the shader expects triangles as input. The second, `OutputTriangleStrip`, specifies that the shader produces triangle strips. Finally, `OutputVertices` is used to specify that each shader invocation will produce a maximum of 3 vertices—a single triangle. The `Invocations` decoration is also specified, but it is set to 1, which is the default, and could therefore have been omitted.

The maximum number of vertices that a single geometry shader can produce is device-dependent. To determine the limit for your device, check the `maxGeometryOutputVertices` field of the device's `VkPhysicalDeviceLimits` structure, which you can retrieve by calling **vkGetPhysicalDeviceProperties()**. The `OutputVertices` argument to the execution mode instruction must be less than or equal to this value. `maxGeometryOutputVertices` is guaranteed to be at least 256.

The maximum number of components that a geometry shader can produce for each vertex it outputs is also device-dependent and is determined by checking the `maxGeometryOutputComponents` field of the device's `VkPhysicalDeviceLimits` structure. This is guaranteed to be at least 64.

In addition to the limit on the number of vertices that the shader can produce, the number of *components* that the shader can produce is subject to a device-dependent limit. This is stored in the `maxGeometryTotalOutputComponents` field of the `VkPhysicalDeviceLimits` structure. This is guaranteed to be at least 1,024—enough for 256 vertices (the minimum guarantee for `maxGeometryOutputVertices`), each consisting of a single **vec4**. Some devices may advertise a higher value for one or both limits. Note that the minimum guarantees aren't a simple product of one another. That is, the gaurantee for `maxGeometryTotalOutputComponents` (1,024) is not equal to the product of `maxGeometryOutputComponents` (64) and

`maxGeometryOutputVertices` (256). Many devices will support producing a large number of small vertices or a small number of larger vertices.

The inputs to the geometry shader come from two places:

- Built-in inputs that are declared inside a `gl_PerVertex` input block declaration
- User-defined inputs that match the corresponding output declaration in the next-earlier shader stage (tessellation evaluation or vertex, depending on whether tessellation is enabled)

To declare the `gl_PerVertex` input block, declare an input block called `gl_PerVertex` and include in it all of the built-in per-vertex input variables that your shader needs to use. For example, [Listing 9.14](#) shows the input block declared containing `gl_Position`, which corresponds to `gl_Position` as written by the previous stage. This block is declared as an array of instances called `gl_in[]`, which is implicitly sized given the primitive type.

Listing 9.14: Declaring `gl_PerVertex` in a GLSL Geometry Shader

```
in gl_PerVertex
{
    vec4 gl_Position;
} gl_in[];
```

When a shader containing the declaration shown in [Listing 9.14](#) and performing a read from `gl_Position` is compiled to SPIR-V, the disassembly shown in [Listing 9.15](#) is produced.

Listing 9.15: Reading `gl_PerVertex` in a SPIR-V Geometry Shader

[**Click here to view code image**](#)

```
...
;; Decorate the first member of our block with the BuiltIn Position.
          OpMemberDecorate %1017 0 BuiltIn Position
          OpDecorate %1017 Block
...
;; Declare an array of structures, with a pointer to this array as an
input.
   %1017 = OpTypeStruct %29
    %557 = OpTypeArray %1017 %2573
   %1194 = OpTypePointer Input %557
   %5305 = OpVariable %1194 Input
    %666 = OpTypePointer Input %29
...
;; Access the input using OpLoad.
   %7129 = OpAccessChain %666 %5305 %2571 %2571
  %15646 = OpLoad %29 %7129
```

As with other shader stages, declaring an input in the geometry shader that does not include a `BuiltIn` decoration means that it will read its value from those produced in the vertex shader that precedes it. The total number of components across all inputs to the geometry shader is subject to a device-dependent limit. This can be determined by checking the `maxGeometryInputComponents` member of the device's `VkPhysicalDeviceLimits` structure. This is guaranteed to be at least 64 components.

Unless a geometry shader explicitly produces output data, it effectively does nothing. The shader can produce individual vertices one at a time, and the vertices will then be assembled into primitives after the shader has executed. To produce a vertex, the shader should execute the `OpEmitVertex` instruction, which is produced by the GLSL built-in function `EmitVertex()`.

When the `OpEmitVertex` instruction is executed, the current values of all output variables are used to create a new vertex and pushed down the pipeline. The values of the output variables become undefined at this point, so for each output vertex, it is necessary to rewrite all of the outputs of the shader. Before `OpEmitVertex` is useful, then, we must declare some outputs in our shader.

Outputs are declared in GLSL using an output-block declaration. For example, to produce a **vec4** and pass it to the subsequent fragment shader, declare a block as shown in Listing 9.16.

Listing 9.16: Declaring an Output Block in GLSL

```
out gs_out
{
    vec4 color;
};
```

Again, compiling this fragment to SPIR-V produces a declaration of a block with a single member containing a vector of four floating-point values. When this variable is written to in the shader, an `OpStore` operation is performed to write into the block.

A complete, pass-through geometry shader is shown in Listing 9.17, and the resulting SPIR-V is shown in Listing 9.18.

Listing 9.17: Pass-Through GLSL Geometry Shader

Click here to view code image

```
#version 450 core

layout (points) in;
layout (points) out;
layout (max_vertices = 1) out;

in gl_PerVertex
{
    vec4 gl_Position;
} gl_in[];

out gs_out
{
    vec4 color;
};

void main(void)
{
    gl_Position = gl_in[0].gl_Position;
    color = vec4(0.5, 0.1, 0.9, 1.0);
```

```
    EmitVertex();
}
```

Listing 9.18: Pass-Through SPIR-V Geometry Shader

[Click here to view code image](#)

```
               OpCapability Geometry
               OpCapability GeometryPointSize
               OpCapability ClipDistance
               OpCapability CullDistance
          %1 = OpExtInstImport "GLSL.std.450"
               OpMemoryModel Logical GLSL450
               OpEntryPoint Geometry %5663 "main" %22044 %5305 %4930
               OpExecutionMode %5663 InputPoints
               OpExecutionMode %5663 Invocations 1
               OpExecutionMode %5663 OutputPoints
               OpExecutionMode %5663 OutputVertices 1
               OpMemberDecorate %2935 0 BuiltIn Position
               OpMemberDecorate %2935 1 BuiltIn PointSize
               OpMemberDecorate %2935 2 BuiltIn ClipDistance
               OpMemberDecorate %2935 3 BuiltIn CullDistance
               OpDecorate %2935 Block
               OpMemberDecorate %1017 0 BuiltIn Position
               OpDecorate %1017 Block
               OpDecorate %1018 Block
          %8 = OpTypeVoid
       %1282 = OpTypeFunction %8
         %13 = OpTypeFloat 32
         %29 = OpTypeVector %13 4
         %11 = OpTypeInt 32 0
       %2573 = OpConstant %11 1
        %554 = OpTypeArray %13 %2573
       %2935 = OpTypeStruct %29 %13 %554 %554
        %561 = OpTypePointer Output %2935
      %22044 = OpVariable %561 Output
         %12 = OpTypeInt 32 1
       %2571 = OpConstant %12 0
       %1017 = OpTypeStruct %29
        %557 = OpTypeArray %1017 %2573
       %1194 = OpTypePointer Input %557
       %5305 = OpVariable %1194 Input
        %666 = OpTypePointer Input %29
        %667 = OpTypePointer Output %29
       %1018 = OpTypeStruct %29
       %1654 = OpTypePointer Output %1018
       %4930 = OpVariable %1654 Output
        %252 = OpConstant %13 0.5
       %2936 = OpConstant %13 0.1
       %1364 = OpConstant %13 0.9
        %138 = OpConstant %13 1
        %878 = OpConstantComposite %29 %252 %2936 %1364 %138
       %5663 = OpFunction %8 None %1282
```

```
   %23915 = OpLabel
    %7129 = OpAccessChain %666 %5305 %2571 %2571
   %15646 = OpLoad %29 %7129
   %19981 = OpAccessChain %667 %22044 %2571
            OpStore %19981 %15646
   %22639 = OpAccessChain %667 %4930 %2571
            OpStore %22639 %878
            OpEmitVertex
            OpReturn
            OpFunctionEnd
```

As you can see in Listing 9.18, the shader enables the `Geometry` capability and then simply copies its input to its output. You can see the call to the `OpEmitVertex` instruction toward the end of the shader, immediately before the `main` function ends.

Cutting Primitives

You may have noticed that the only output primitive types available in a geometry shader are points, line strips, and triangle strips. It is not possible to directly output individual lines or triangles. A geometry shader can output any number of vertices, up to an implementation-defined limit, so long as the maximum output vertex count is properly declared. However, just calling `EmitVertex()` many times will produce a single long strip.

In order to output several smaller strips and (within the limit) individual lines or triangles, the `EndPrimitive()` function is available. This function ends the current strip and starts a new one when the next vertex is emitted. The current strip is automatically ended when the shader exits, so if your maximum vertex count is 3, and the output primitive type is **triangle_strip**, then it's not necessary to explicitly call `EndPrimitive()`. However, if you want to produce several independent lines or triangles from a single geometry shader invocation, call `EndPrimitive()` between each strip.

The geometry shaders shown in previous listings use point primitives as both input and output to sidestep this issue. Listing 9.19 shows a shader that uses the **triangle_strip** output primitive type to output triangle strips. However, this shader produces six output vertices, each group of three representing a separate triangle. The `EndPrimitive()` function is used to cut the triangle strip after each triangle in order to produce two separate strips of one triangle each.

Listing 9.19: Cutting Strips in a Geometry Shader

Click here to view code image

```
#version 450 core

layout (triangles) in;
layout (triangle_strip, max_vertices = 6) out;

void main(void)
{
    int i, j;
    vec4 scale = vec4(1.0f, 1.0f, 1.0f, 1.0f);

    for (j = 0; j < 2; j++)
```

```
    {
        for (i = 0; i < 3; i++)
        {
            gl_Position = gl_in[i].gl_Position * scale;

            EmitVertex();
        }

        EndPrimitive();
        scale.xy = -scale.xy;
    }
}
```

In the shader shown in [Listing 9.19](), the first iteration of the outer loop produces a first triangle by simply copying all three vertices' positions (in the inner loop) into the output position, multiplying by a scale factor, and then calling `EmitVertex()`. After the inner loop has completed, the shader calls `EndPrimitive()` to cut the strip at that point. It then inverts the x and y axes of the scale, rotating the triangle 180 degress around the z axis and iterates a second time. This produces a second triangle.

Geometry Shader Instancing

[Chapter 8](), "[Drawing]()," covered instancing, which is a technique to quickly draw many copies of the same geometry using a single drawing command. The number of instances is passed as a parameter to **vkCmdDraw()** or **vkCmdDrawIndexed()**, or through a structure in memory to **vkCmdDrawIndirect()** or **vkCmdDrawIndexedIndirect()**. With draw-level instancing, the entire draw is effectively executed many times. This includes fetching data from index and vertex buffers that does not change, checking for primitive restart indices (if enabled), and so on.

When a geometry shader is enabled, a special mode of instancing is available that runs the pipeline from the geometry shader onward multiple times, leaving all of the stages behind the geometry shader (including tessellation, for example) running only once. While it may not be worth introducing a geometry shader simply to use this form of instancing, when a geometry shader is required anyway, this can be a very efficient mechanism to quickly render multiple copies of a piece of geometry, with the properties of each instance controlled by the geometry shader.

You may have noticed that some of the SPIR-V listings shown earlier in this chapter include a declaration of the following form at the front of all of the geometry shaders:

[Click here to view code image]()

```
OpExecutionMode %n Invocations 1
```

This `Invocations` execution mode tells Vulkan how many times to run the geometry shader—which is the number of instances[2] to run.

> [2.]() In references to an instanced geometry shader, the term *invocation* is often used to mean *instance* because an invocation of the geometry shader is run for each instance. This disambiguates the *instance* used in instanced draws from the *invocation* being the execution of the geometry shader. In fact, both can be used at the same time to have an instanced draw using a geometry shader with multiple invocations.

Setting the invocation count to 1 simply means that the shader will run once as expected. This is the default and is inserted by the GLSL compiler if you don't instruct it otherwise. To control the number of invocations of the geoemtry, use the **invocations** input layout qualifier in the GLSL shader. For example, the following declaration sets the invocation count for the shader to 8:

```
    layout (invocations = 8) in;
```

When this layout qualifier is included in a GLSL shader, the compiler will insert the appropriate `OpExecutionMode` instruction to set the invocation count for the entry point to the value you specified. While the invocation count is hard-coded into the shader rather than passed as a parameter, as it is in draw instancing, you can use a specialization constant to set its value. This allows you to customize a single geometry shader to run in a number of different scenarios.

As the shader executes, the invocation number is made available to it through the GLSL `gl_InvocationID` built-in variable. In SPIR-V, this is translated into a decoration attaching the built-in `InvocationId` to an input variable.

The GLSL shader shown in [Listing 9.20](#) is a complete example of running two instances with each sourcing a separate object-to-world matrix in order to draw two copies of an object, each with a different transform. By using the `gl_InvocationID` (which turns into the `InvocationId` decoration), we can index into an array of matrices in order to get a different transform for each invocation.

Listing 9.20: Instanced GLSL Geometry Shader

Click here to view code image

```
#version 450 core

layout (triangles, invocations = 2) in;
layout (triangle_strip, max_vertices = 3) out;

layout (set = 0, binding = 0) uniform
{
    mat4 projection;
    mat4 objectToWorld[2];
} transforms;

void main(void)
{
    int i;
    mat4 objectToWorld = transforms.objectToWorld[gl_InvocationID];
    mat4 objectToClip = objectToWorld * transforms.projection;

    for (i = 0; i < 3; i++)
    {
        gl_Position = gl_in[i].gl_Position * objectToClip;

        EmitVertex();
    }
}
```

The maximum number of geometry shader invocations is implementation-dependent but is guaranteed to be at least 32. Some implementations support more invocations than this, and the limit for a particular implementation can be determined by checking the `maxGeometryShaderInvocations` field of the device's `VkPhysicalDeviceLimits` structure as returned from a call to **`vkGetPhysicalDeviceProperties()`**.

Programmable Point Size

When you are rendering points, by default, a single vertex produces a point that is exactly 1 pixel wide. On modern high-resolution displays, a single pixel is extremely small, so it is often the case that you wish to render points that are much larger than this. When the primitives are rasterized as points, it is possible to set the point size by using the last stage of the geometry-processing pipeline.

Points are rasterized in one of three ways:

- Rendering with only a vertex and fragment shader, and setting the primitive topology to `VK_PRIMITIVE_TOPOLOGY_POINT_LIST`
- Enabling tessellation and setting the tessellation mode to points by decorating the tessellation shader's entry point with the SPIR-V `PointMode` execution mode
- Using a geometry shader that produces points

The last stage in the geometry pipeline (vertex, tessellation evaluation, or geometry) can specify the size of points by decorating a floating-point output with the `PointSize` parameter to the `BuiltIn` decoration. The value written to this output will be used as the diameter of the rasterized points.

In GLSL, such a decorated output can be produced by writing to the `gl_PointSize` built-in output. An example vertex shader that writes to `gl_PointSize` is shown in Listing 9.21, and the resulting SPIR-V output is shown in Listing 9.22.

Listing 9.21: Use of `gl_PointSize` in GLSL

Click here to view code image

```
#version 450 core

layout (location = 0) in vec3 i_position;
layout (location = 1) in float i_pointSize;

void main(void)
{
    gl_Position = vec4(i_position, 1.0f);
    gl_PointSize = i_pointSize;
}
```

The shader in Listing 9.21 declares two inputs: `i_position` and `i_pointSize`. Both are passed through to their respective outputs. Simply writing to `gl_PointSize` causes the GLSL compiler to automatically declare an output variable in the SPIR-V shader and decorate it with the `PointSize` decoration, as you can see in Listing 9.22, which has been manually edited and commented.

Listing 9.22: Decorating an Output with `PointSize`

Click here to view code image

```
        ...
        ;; GLSL compiler automatically declares per-vertex output block.
        OpName %11 "gl_PerVertex"
        OpMemberName %11 0 "gl_Position"
        OpMemberName %11 1 "gl_PointSize"
        OpMemberName %11 2 "gl_ClipDistance"
```

```
        OpMemberName %11 3 "gl_CullDistance"
        OpName %13 ""
        ;; Naming inputs
        OpName %18 "i_position"
        OpName %29 "i_pointSize"
        ;; Decorating members of the default output block
        OpMemberDecorate %11 0 BuiltIn Position
        OpMemberDecorate %11 1 BuiltIn PointSize ;; gl_PointSize
        OpMemberDecorate %11 2 BuiltIn ClipDistance
        OpMemberDecorate %11 3 BuiltIn CullDistance
        OpDecorate %11 Block
        OpDecorate %18 Location 0
        OpDecorate %29 Location 1
        ...
  %4 = OpFunction %2 None %3
        ;; Start of "main"
  %5 = OpLabel
        ;; Load i_position.
 %19 = OpLoad %16 %18
 %21 = OpCompositeExtract %6 %19 0
 %22 = OpCompositeExtract %6 %19 1
 %23 = OpCompositeExtract %6 %19 2
        ;; Construct vec4 and write to gl_Position.
 %24 = OpCompositeConstruct %7 %21 %22 %23 %20
 %26 = OpAccessChain %25 %13 %15
        OpStore %26 %24
        ;; Load from i_pointSize (%29).
 %30 = OpLoad %6 %29
        ;; Use access chain to dereference built-in output.
 %32 = OpAccessChain %31 %13 %27
        ;; Store to output decorated with PointSize.
        OpStore %32 %30
        OpReturn
        OpFunctionEnd
```

The size produced by the shader and written to the output decorated as PointSize must fall within the range of point sizes supported by the device. This can be determined by inspecting the pointSizeRange member of the device's VkPhysicalDeviceLimits structure, which you can obtain through a call to **vkGetPhysicalDeviceProperties()**. This is an array of two floating-point values, the first being the smallest point that can be rasterized and the second being the diameter of the largest point that can be rasterized.

The pixel size of a point will be quantized to a device-dependent scale. The delta between supported point sizes can be determined from the pointSizeGranularity field of the VkPhysicalDeviceLimits structure. For example, if an implementation can render points at any size in quarter-pixel increments, then pointSizeGranularity will be 0.25. All devices must be able to render any supported point size with at least single-pixel accuracy, though many devices will provide much more precision than this.

The largest guaranteed value for the minimum point size is a single pixel. That is, some implementations may be able to accurately rasterize points smaller than a single pixel, but all implementations must be able to rasterize single-pixel points. The smallest guaranteed point size, in

pixels, is 64.0 minus one unit of the device's granularity as given in `pointSizeGranularity`. So if the device renders points in quarter-pixel increments, the maximum point size will be 63.75 pixels.

If none of the geometry processing shaders writes a variable decorated as `PointSize`, then the default point size of 1 pixel will be assumed. Many implementations may be more efficient when a compile-time constant is written to `PointSize`, often removing the executable code from the shader and instead programming the point size as the state of the rasterizer. It is possible to set the point size through a specialization constant to produce a shader that writes a compile-time constant to the point size but is still configurable at pipeline build time, much as the line width is.

Line Width and Rasterization

Optionally, Vulkan is able to rasterize lines that are greater than a single pixel wide. The width of lines to be rasterized can be specified when a graphics pipeline is created by setting the `lineWidth` of the `VkPipelineRasterizationStateCreateInfo` structure used to create the pipeline. Lines greater than a single pixel wide are known as *wide lines*. Wide lines are supported by the implementation if the `wideLines` field of its `VkPhysicalDeviceFeatures` is `VK_TRUE`. In this case, the range of supported line widths is contained in the `lineWidthRange` member of the device's `VkPhysicalDeviceLimits` structure, which can be retrieved with a call to **`vkGetPhysicalDeviceProperties()`**.

A given Vulkan implementation may render lines in one of two ways: *strict* or *nonstrict*. Which method the device uses is reported in the `strictLines` field of its `VkPhysicalDeviceLimits` structure. If this is `VK_TRUE`, then strict line rasterization is implemented; otherwise, only nonstrict rasterization is supported by the device. There is no way to choose which method is used; a device will implement only one method.

In general, strict or nonstrict line rasterization does not affect single-pixel wide lines, so it really applies only to wide lines. When wide lines are in use, a strict line is essentially rasterized as though it were a rectangle centered on the line running from the starting point to the ending point of each segment. As a result, the square end caps of the line are perpendicular to the direction of the line. This is illustrated in Figure 9.6, which shows a line segment from $\{x_a, y_a, z_a\}$ to $\{x_b, y_b, z_b\}$ rotated onto the diagonal. The original line is shown as a dotted line, and the outline of the rasterized line is solid. You should be able to see from Figure 9.6 that as the line rotates away from horizontal or vertical, the width of its cross-section does not change. It is effectively rasterized as a rectangle whose long edge is the length of the line and whose short edge is the width of the line.
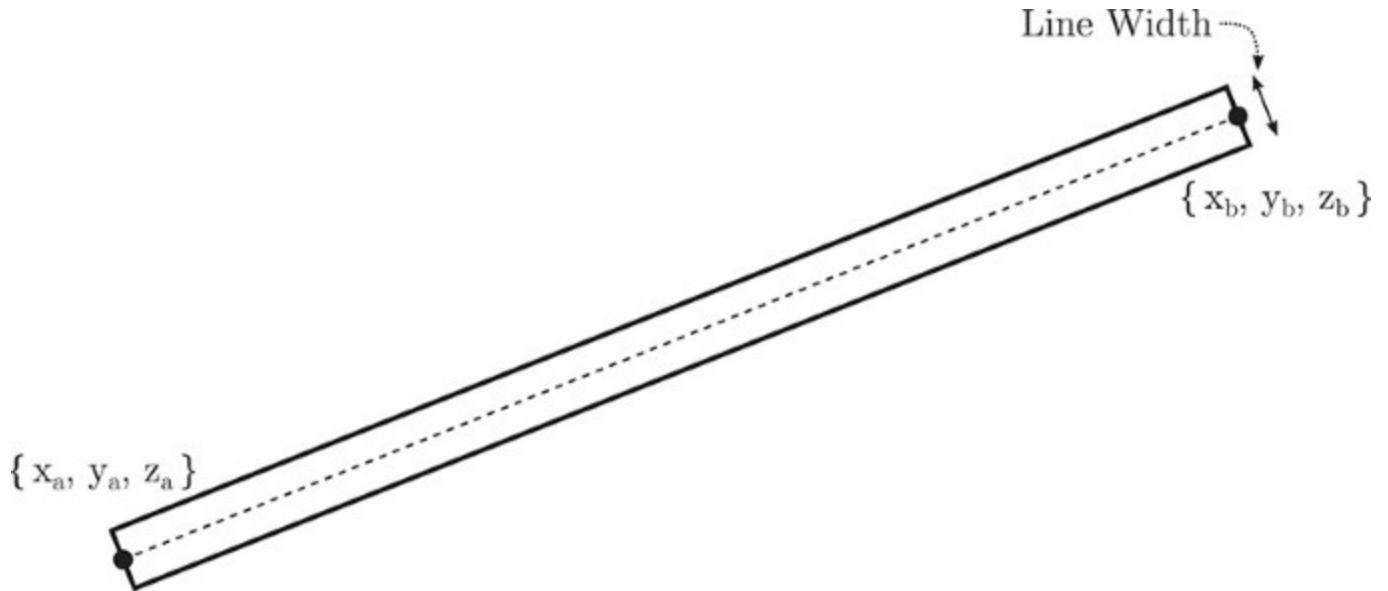
Figure 9.6: Rasterization of Strict Lines

Meanwhile, nonstrict lines are rasterized as a series of horizontal or vertical columns of fragments, depending on the major axis of the line. The major axis of a line is the axis (either x or y) in which it moves the farthest. If the change in x coordinate along the line is greatest, then it is an x-major line and will be rasterized as a sequence of vertical columns of fragments. Conversely, if it moves farthest in y, then it is a y-major line and will be rasterized as a sequence of horizontal rows of fragments.

As a consequence, the end of the line is no longer square unless the line is perfectly horizontal or vertical. Instead, it is wedge-shaped, with the endcap of the line perpendicular to the line's major axis rather than to the line itself. For small line widths up to a few pixels, this may not be noticeable. For larger line widths, however, this can be visually distracting. Figure 9.7 illustrates nonstrict line rasterization. As you can see in the figure, as the line rotates away from the horizontal or vertical, it becomes a parallelogram, with the short side length being the line width. The farther the line becomes from the horizontal or vertical, the narrower it will appear as its perpendicular width decreases.
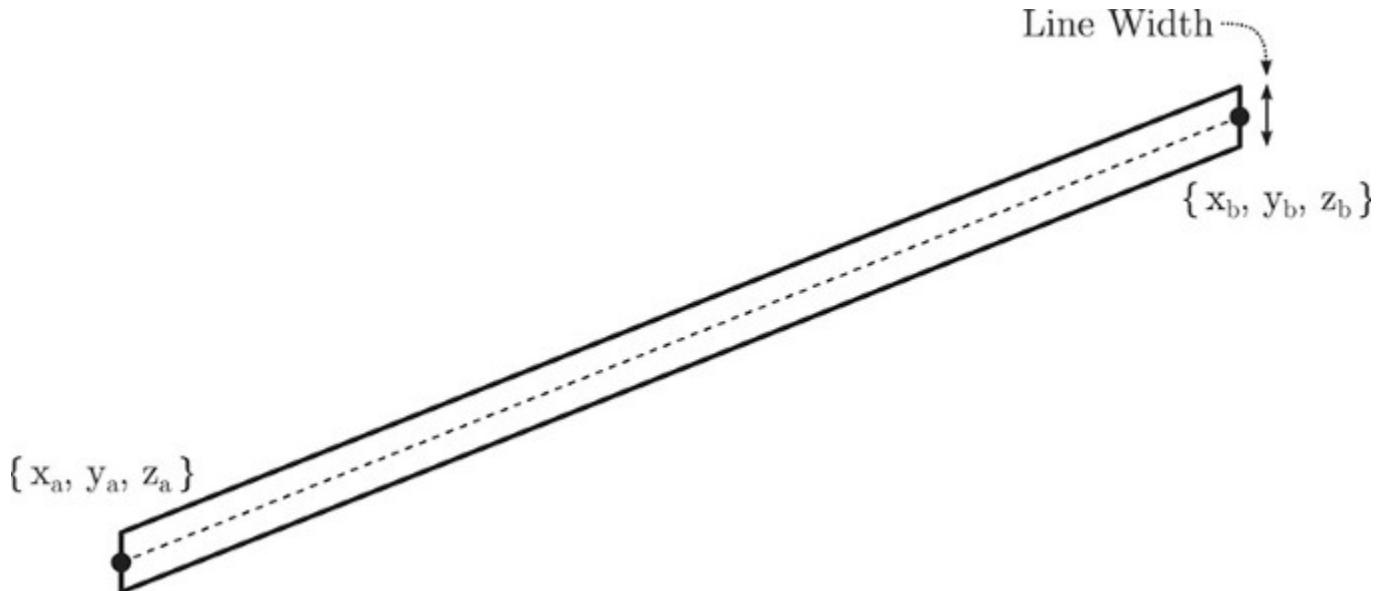
Figure 9.7: Rasterization of Nonstrict Lines

The width of lines used for rasterization can also be marked as a dynamic state. To do this, include the VK_DYNAMIC_STATE_LINE_WIDTH token in the list of dynamic states passed through the pDynamicStates member of the VkPipelineDynamicStateCreateInfo structure used to create the graphics pipeline. Once the line width is marked as dynamic, the lineWidth field of VkPipelineRasterizationStateCreateInfo is ignored, and the line width is instead set by calling **vkCmdSetLineWidth()**, whose prototype is

[Click here to view code image](#)

```
void vkCmdSetLineWidth (
    VkCommandBuffer                         commandBuffer,
    float                                   lineWidth);
```

The lineWidth parameter sets the width of lines, in pixels. Any line primitive rendered will take this thickness, whether it is the result of drawing with one of the line topologies or whether a tessellation or geometry shader turns another type of primitive into lines. The new line-width parameter must be between the minimum and maximum line widths supported by the Vulkan implementation. This can be determined by checking the lineWidthRange member of the device's VkPhysicalDeviceLimits structure. The first element of lineWidthRange is the minimum width, and the second element is the maximum. Support for lines is optional, so some implementations will return 1.0 for both elements. However, if wide lines are supported, the Vulkan implementation will support a range of line widths that includes 1.0 to 8.0 pixels.

User Clipping and Culling

In order to ensure that no geometry is rendered outside the viewport, Vulkan performs clipping of geometry against the viewport bounds. A typical method for doing this is to determine the distance of each vertex to the planes, defining the viewport as a signed quantity. Positive distances are on the "inside" of the volume, and negative distances are "outside" the volume. The distance to each plane is computed separately for each vertex. If all of the vertices belonging to a primitive are on the outside of a single plane, then the whole primitive can be discarded. Conversely, if all of the vertices have a positive distance to all planes, that means that they are all inside the volume, so the entire primitive can be safely rendered.

If there is a mix of "inside" and "outside" that make up the primitive, then it must be clipped, which normally means breaking it down into a set of smaller primitives. Figure 9.8 illustrates this. As you can see in the figure, four triangles have been sent into the pipeline. Triangle A is entirely contained within the viewport and is rendered directly. Triangle B, on the other hand, lies entirely outside the viewport and can be trivially discarded. Triangle C penetrates a single edge of the viewport and is therefore clipped against it. The smaller triangle is produced by the clipper, and this is the one that is rasterized. Finally, triangle D presents a more complex scenario. It lies partially inside the viewport but penetrates two of the viewport's edges. In this case, the clipper breaks the triangle into several smaller triangles and rasterizes those. The resulting polygon is shown in bold, and the generated interior edges are shown as dotted lines.
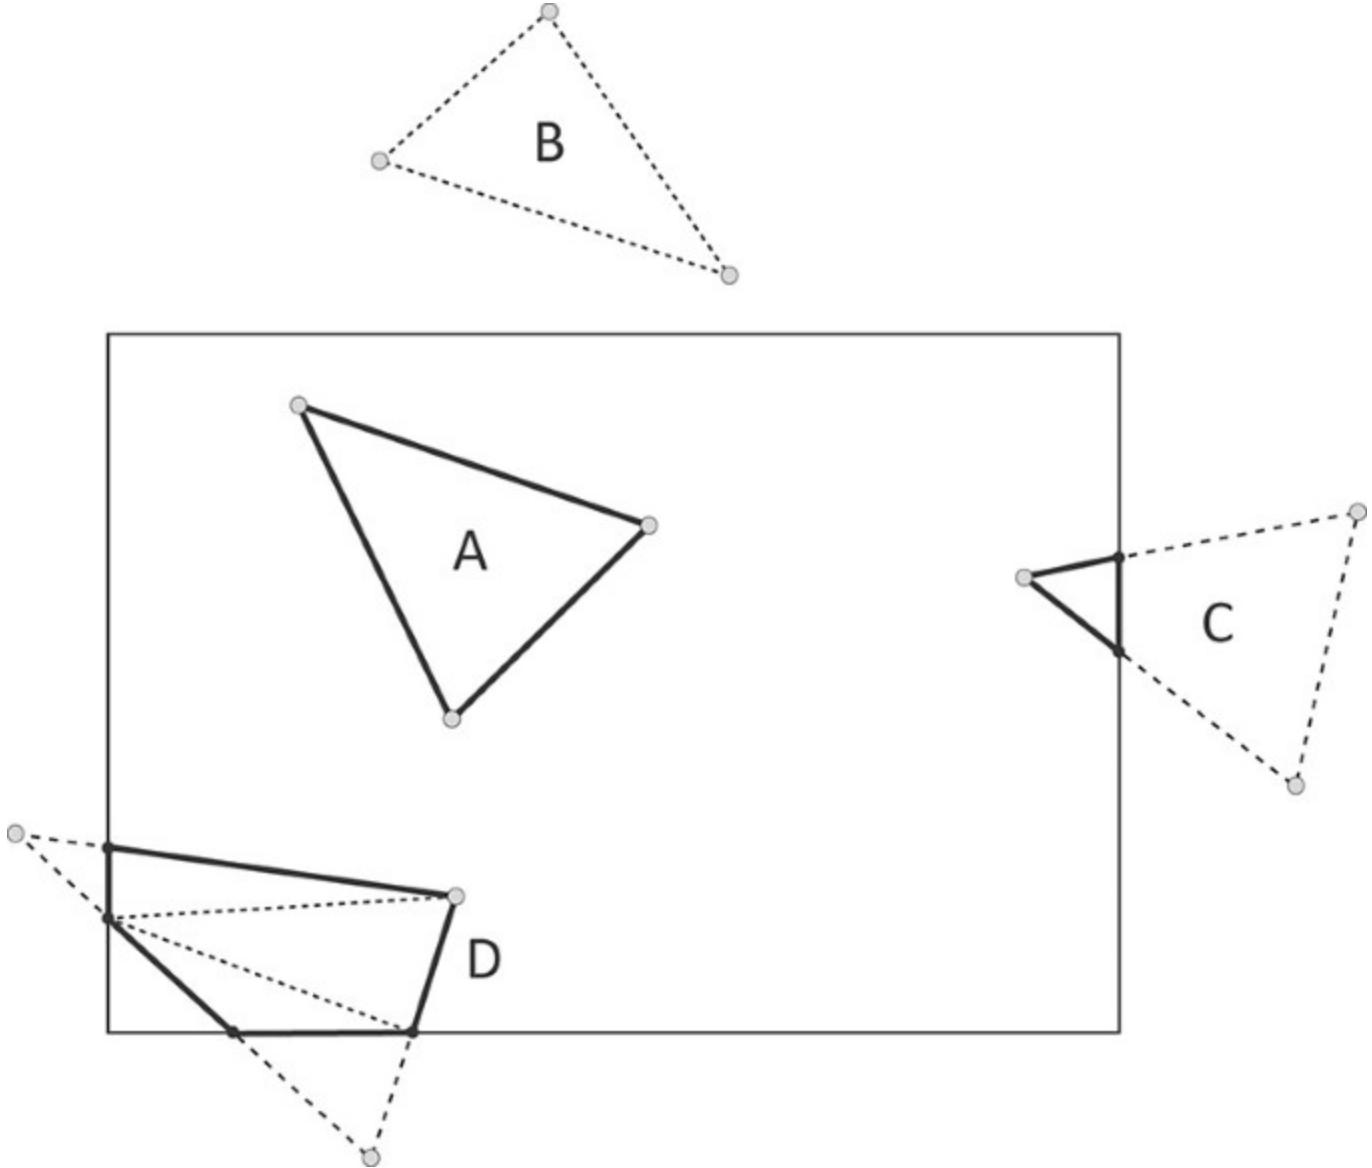
Figure 9.8: Clipping Against a Viewport

In addition to being able to clip a primitive against the edges of the viewport, it is possible to supply other distance values generated in your shaders that contribute to this clipping procedure. These are known as *clip distances* and, once assigned, are treated exactly as the distances computed to the viewport planes. Positive values are treated as being inside the view volume, and negative values are treated as being outside the view volume.

To generate a clip distance in your shader, decorate an output variable by using the `ClipDistance` decoration. This variable must be declared as an array of 32-bit floating-point values, each element of which is a separate clip distance. Vulkan will aggregate all of the clip distances written by your shaders when performing clipping on the generated primitives.

Not all devices support clip distances. If clip distance is supported by a device, then the maximum number of clip distances supported by that device is guaranteed to be at least 8. Some devices may support more distances than this. You can check how many distances are supported by the device by inspecting the `maxClipDistances` field of its `VkPhysicalDeviceLimits` structure. You

can retrieve this by calling **vkGetPhysicalDeviceProperties()**. If clip distances are not supported or are not enabled at device-creation time, then this field will be 0.

The last geometry-processing stage (vertex, tessellation evaluation, or geometry shader) produces the clip distances that will be used by the clipping stage to clip the generated primitives. In any stage after the clip distances have been declared as outputs, the produced distances can be made available as inputs. Therefore, in the tessellation control, tessellation evaluation, or geometry shader stages, you can read (and rewrite) the values produced in the previous stage if you need to.

The clip distance is also available as an input to the fragment shader. While clipping will be performed at the primitive level by most implementations, another way to implement clipping by using the clip distance is to interpolate the clip distances applied to the vertices across the primitive. Fragments to which are assigned any negative clip distance values are discarded before fragment processing completes. Even if an implementation doesn't implement clip distances this way, the interpolated value of each clip distance is made available to the fragment shader should it decorate an input variable with the `ClipDistance` decoration.

Listing 9.24 shows an example of decorating an output variable with `ClipDistance` in SPIR-V, and Listing 9.23 shows the GLSL fragment that was used to generate this SPIR-V. As you can see, in GLSL, the built-in variable `gl_ClipDistance` is used to write clip distances and is translated to an output variable with the appropriate decorations by the GLSL compiler.

Listing 9.23: Use of `gl_ClipDistance` in GLSL

Click here to view code image

```
#version 450 core

// Redeclare gl_ClipDistance to explicitly size it.
out float gl_ClipDistance[1];

layout (location = 0) in vec3 i_position;

// Push constant from which to assign clip distance
layout (push_constant) uniform push_constants_b
{
    float clip_distance[4];
} push_constant;

void main(void)
{
    gl_ClipDistance[0] = push_constant.clip_distance[0];
    gl_Position = vec4(i_position, 1.0f);
}
```

The shader in Listing 9.23 simply assigns the value of a push constant directly to the `gl_ClipDistance` output. A more practical use of `gl_ClipDistance` would compute a distance to a plane for each vertex and assign that to the clip-distance output. This rather simple shader serves to demonstrate how the SPIR-V result is generated. This is shown in Listing 9.24.

Listing 9.24: Decorating Outputs with `ClipDistance`

[Click here to view code image](#)

```
            OpCapability Shader
;; The shader requires the ClipDistance capability.
            OpCapability ClipDistance
       %1 = OpExtInstImport "GLSL.std.450"
            OpMemoryModel Logical GLSL450
            OpEntryPoint Vertex %4 "main" %13 %29
            OpSource GLSL 450
            OpName %4 "main"
            OpName %11 "gl_PerVertex"
            OpMemberName %11 0 "gl_Position"
            OpMemberName %11 1 "gl_PointSize"
;; Redeclaration of gl_ClipDistance built-in
            OpMemberName %11 2 "gl_ClipDistance"
            OpMemberName %11 3 "gl_CullDistance"
            OpName %13 ""
            OpName %19 "push_constants_b"
            OpMemberName %19 0 "clip_distance"
            OpName %21 "push_constant"
            OpName %29 "i_position"
            OpMemberDecorate %11 0 BuiltIn Position
            OpMemberDecorate %11 1 BuiltIn PointSize
;; Decorate the built-in variable as ClipDistance.
            OpMemberDecorate %11 2 BuiltIn ClipDistance
            OpMemberDecorate %11 3 BuiltIn CullDistance
            OpDecorate %11 Block
            OpDecorate %18 ArrayStride 4
            OpMemberDecorate %19 0 Offset 0
            OpDecorate %19 Block
            OpDecorate %29 Location 0
       %2 = OpTypeVoid
       %3 = OpTypeFunction %2
       %6 = OpTypeFloat 32
       %7 = OpTypeVector %6 4
       %8 = OpTypeInt 32 0
       %9 = OpConstant %8 1
      %10 = OpTypeArray %6 %9
;; This creates the built-in output structure containing gl_ClipDistance.
      %11 = OpTypeStruct %7 %6 %10 %10
      %12 = OpTypePointer Output %11
;; Instantiate the built-in outputs.
      %13 = OpVariable %12 Output
      ...
;; Beginning of main()
       %5 = OpLabel
;; Load from the push-constant array.
      %23 = OpAccessChain %22 %21 %16 %16
      %24 = OpLoad %6 %23
;; Store to clip distance.
      %26 = OpAccessChain %25 %13 %15 %16
```

```
            OpStore %26 %24
        ...
;; End of main()
            OpReturn
            OpFunctionEnd
```

There is no reason that your shader should assign values to the `ClipDistance` outputs by computing distance to a plane. You could assign the distance to an analytic function, to a higher-order surface, or even to a displacement map read from a texture, for example. However, bear in mind that because the primitives are clipped as though the distance was computed from a flat plane, the resulting edges will be a piecewise linear approximation of the function used to compute those distances. If your function represents tight curves or detailed surfaces, you will need quite a bit of geometry to make the resulting edges appear smooth.

While the tessellation control and geometry shading stages have access to entire primitives, the vertex and tessellation evaluation shaders do not. Therefore, if you wish to discard an entire primitive from one of these stages, it is hard to coordinate the shader invocations corresponding to that primitive in order to assign a negative clip distance to all of them. For this purpose, you can instead use the *cull distance* for each vertex. The cull distance works very similarly to the clip distance. The difference is that the entire primitive is discarded if *any* vertex it contains has a negative cull distance, regardless of the values of the cull distances for the other vertices.

To use cull distance, decorate an output variable with the `CullDistance` decoration. The cull distance is also available as in input to subsequent shader stages, just as `ClipDistance` is. Also, if `CullDistance` is used as an input in the fragment shader, its content will be the interpolated value of the distance assigned in the geometry-processing stages. Listings 9.25 and 9.26 show minor modifications to Listings 9.23 and 9.24, respectively, to show assignment to variables decorated with `CullDistance` rather than `ClipDistance`.

Listing 9.25: Use of `gl_CullDistance` in GLSL

Click here to view code image

```glsl
#version 450 core

out float gl_CullDistance[1];

layout (location = 0) in vec3 i_position;

layout (push_constant) uniform push_constants_b
{
    float cull_distance[4];
} push_constant;

void main(void)
{
    gl_CullDistance[0] = push_constant.cull_distance[0];
    gl_Position = vec4(i_position, 1.0f);
}
```

As you can see, Listing 9.25 is almost identical to Listing 9.23 except that we have used `gl_CullDistance` in place of `gl_ClipDistance`. As you might expect, the resulting SPIR-V

shader shown in is also almost identical to that in , which uses the `ClipDistance` decoration.

Listing 9.26: Decorating Outputs with `CullDistance`

[Click here to view code image](#)

```
           OpCapability Shader
;; The shader requires the CullDistance capability.
           OpCapability CullDistance
      %1 = OpExtInstImport "GLSL.std.450"
           OpMemoryModel Logical GLSL450
           OpEntryPoint Vertex %4 "main" %13 %29
           OpSource GLSL 450
           OpName %4 "main"
           OpName %11 "gl_PerVertex"
           OpMemberName %11 0 "gl_Position"
           OpMemberName %11 1 "gl_PointSize"
           OpMemberName %11 2 "gl_ClipDistance"
;; Redeclaration of gl_CullDistance built-in
           OpMemberName %11 3 "gl_CullDistance"
           OpName %13 ""
           OpName %19 "push_constants_b"
           OpMemberName %19 0 "cull_distance"
           OpName %21 "push_constant"
           OpName %29 "i_position"
           OpMemberDecorate %11 0 BuiltIn Position
           OpMemberDecorate %11 1 BuiltIn PointSize
           OpMemberDecorate %11 2 BuiltIn ClipDistance
;; Decorate the built-in variable as CullDistance.
           OpMemberDecorate %11 3 BuiltIn CullDistance
           OpDecorate %11 Block
           OpDecorate %18 ArrayStride 4
           OpMemberDecorate %19 0 Offset 0
           OpDecorate %19 Block
           OpDecorate %29 Location 0
      %2 = OpTypeVoid
      %3 = OpTypeFunction %2
      %6 = OpTypeFloat 32
      %7 = OpTypeVector %6 4
      %8 = OpTypeInt 32 0
      %9 = OpConstant %8 1
     %10 = OpTypeArray %6 %9
;; This creates the built-in output structure containing gl_CullDistance.
     %11 = OpTypeStruct %7 %6 %10 %10
     %12 = OpTypePointer Output %11
;; Instantiate the built-in outputs.
     %13 = OpVariable %12 Output
           ...
;; Beginning of main()
      %5 = OpLabel
;; Load from the push-constant array.
     %23 = OpAccessChain %22 %21 %16 %16
```

```
      %24 = OpLoad %6 %23
;; Store to cull distance.
      %26 = OpAccessChain %25 %13 %15 %16

      ...
;; End of main()
            OpReturn
            OpFunctionEnd
```

As you can see in [Listing 9.25](), outputs decorated with `CullDistance` are produced by writing to the GLSL built-in variable `gl_CullDistance`.

Whenever `ClipDistance` or `CullDistance` is used in the fragment shader, you should only see *positive* values for those inputs. This is because any fragment that would have had a negative clip or cull distance should have been discarded. The only possible exception to this is *helper invocations*, which are invocations of the fragment shader that are used to generate deltas during the computation of gradients. If your fragment shader is sensitive to negative values of `ClipDistance`, then this may be a case you care about.

As with clip distances, cull distances are declared in your shaders as arrays of floating-point values, and the number of elements in those arrays is device-dependent. Some devices do not support cull distances, but if they do, they are guaranteed to support at least 8. To determine the number of distances supported, check the `maxCullDistances` field of the device's `VkPhysicalDeviceLimits` structure. If cull distances are not supported, then this field will be 0.

Before cull distances can be used in a SPIR-V shader, the shader must enable the `CullDistance` capability using an `OpCapability` instruction. You can check whether the device supports the `CullDistance` capability in SPIR-V shaders by inspecting the `shaderCullDistance` member of the device's `VkPhysicalDeviceFeatures` structure obtained through a call to **`vkGetPhysicalDeviceProperties()`**. You must also enable this feature when the device is created by setting `shaderCullDistance` to `VK_TRUE` in the `VkPhysicalDeviceFeatures` structure used to create the device. Likewise, for clip distances, the capability is enabled by executing `OpCapability` with the `ClipDistance` argument, and checked and enabled through the `shaderClipDistance` field of `VkPhysicalDeviceFeatures`.

Because clip and cull distances may consume similar resources, in addition to the `maxClipDistances` and `maxCullDistances` limits, many devices have a *combined* limit for the number of distances that you can use at once. This often precludes you from using the maximum number of both at the same time. To check the combined limit, look at the `maxCombinedClipAndCullDistances` field of the device's `VkPhysicalDeviceLimits` structure.

The Viewport Transformation

The final transformation in the pipeline before rasterization is the viewport transformation. The coordinates produced by the last stage in the geometry pipeline (or produced by the clipper) are in homogeneous clip coordinates. First, the homogeneous coordinates are all divided through by their own *w* components, producing normalized device coordinates.

$$
\begin{pmatrix} x_c \\ y_c \\ z_c \\ w_c \end{pmatrix} \times \frac{1}{w_c} = \begin{pmatrix} x_d \\ y_d \\ z_d \\ 1.0 \end{pmatrix}
\tag{9.1}
$$

Here, the vertex's clip coordinates are represented by $\{x_c, y_c, z_c, w_c\}$, and the normalized device coordinates are represented by $\{x_d, y_d, z_d\}$. Because the vertex's w component is divided by itself, it always becomes 1.0 and so can be discarded at this point, and the normalized device coordinate is considered to be a 3D coordinate.

Before the primitive can be rasterized, it needs to be transformed into framebuffer coordinates, which are coordinates relative to the origin of the framebuffer. This is performed by scaling and biasing the vertex normalized device coordinates by the selected viewport transform.

One or more viewports are configured as part of the graphics pipeline. This was briefly introduced in Chapter 7, "Graphics Pipelines." Most applications will use a single viewport at a time. When there are multiple viewports configured, each behaves the same way; only the selection of viewport is controlled by the geometry shader if it is present. If there is no geometry shader, then only the first configured viewport is accessible.

Each viewport is defined by an instance of the `VkViewport` structure, the definition of which is

```
typedef struct VkViewport {
    float     x;
    float     y;
    float     width;
    float     height;
    float     minDepth;
    float     maxDepth;
} VkViewport;
```

The transformation from normalized device coordinates to framebuffer coordinates is performed as

$$
\begin{pmatrix} x_f \\ y_f \\ z_f \end{pmatrix} = \begin{pmatrix} \frac{p_x}{2} \times x_d + o_x \\ \frac{p_y}{2} \times y_d + o_y \\ p_z \times z_d + o_z \end{pmatrix}
\tag{9.2}
$$

Here, x_f, y_f, and z_f are the coordinates of each vertex in framebuffer coordinates. The x and y fields of `VkViewport` are o_x and o_y, respectively, and `minDepth` is o_z. The `width` and `height` fields are used for p_x and p_y. p_z is formed from the expression (`maxDepth` - `minDepth`).

The maximum size of a viewport is a device-dependent quantity, although it is guaranteed to be at least 4,096 × 4,096 pixels in size. If all of your color attachments are smaller than this (or if the width and height of all of your viewports are smaller), then there is no reason to query the upper limit of viewport size. If you want to render into an image that is larger than this, you can determine the upper limit on viewport size by inspecting the `maxViewportDimensions` member of the device's `VkPhysicalDeviceLimits` structure, which you can retrieve by calling **vkGetPhysicalDeviceFeatures()**.

`maxViewportDimensions` is an array of two floating-point values. The first element is the maximum supported viewport width, and the second is the maximum supported viewport height. The `width` and `height` members of `VkViewport` must be less than or equal to these values.

Although the width and height of each viewport must be within the limits reported in `maxViewportDimensions`, it is possible to offset the viewport within a larger set of attachments. The maximum extent to which the viewport can be offset is determined by checking the `viewportBoundsRange` field of the device's `VkPhysicalDeviceLimits` structure. So long as the left, right, top, and bottom of the viewport lie within the range of `viewportBoundsRange[0]` and `viewportBoundsRange[1]`, then the viewport can be used.

Although the outputs of the vertex shader and the parameters of the viewport are all floating-point quantities, the resulting framebuffer coordinates are generally converted to fixed-point representation before rasterization. The range supported by the viewport coordinates must obviously be large enough to represent the maximum viewport size. The number of fractional bits determines the accuracy at which vertices are snapped to pixel coordinates.

This precision is device-dependent, and some devices might snap directly to pixel centers. However, this is uncommon, and most devices employ some subpixel precision. The amount of precision a device uses in its viewport coordinate representation is contained in the `viewportSubPixelBits` field of its `VkPhysicalDeviceLimits` structure.

As well as specifying the bounds of the viewports when the graphics pipeline is created, the viewport state can be made dynamic. To do this, include `VK_DYNAMIC_STATE_VIEWPORT` in the list of dynamic states when the pipeline is created. When this token is included, the values of the viewport bounds specified at pipeline-creation time are ignored. Only the number of viewports is relevant. To set the viewport bounds dynamically, call **`vkCmdSetViewport()`**, the prototype of which is

Click here to view code image

```
void vkCmdSetViewport (
    VkCommandBuffer                    commandBuffer,
    uint32_t                           firstViewport,
    uint32_t                           viewportCount,
    const VkViewport*                  pViewports);
```

Any subset of the active viewports can be updated with **`vkCmdSetViewport()`**. The `firstViewport` parameter specifies the first viewport to update, and the `viewportCount` parameter specifies the number of viewports (starting from `firstViewport`) to update. The dimensions of the viewports is specified in `pViewports`, which is a pointer to an array of `viewportCount VkViewport` structures.

The number of viewports supported by the current pipeline is specified in the `viewportCount` field of `VkPipelineViewportStateCreateInfo` and is always considered to be static. Unless the pipeline disables rasterization, there must be at least one viewport in the pipeline. You can set viewports outside the number supported by the current pipeline and then switch to a pipeline with more viewports, and it will use the state you specified.

Usually, you will use a single viewport that covers all of the framebuffer attachments. However, in some cases you might want to render to a smaller window within the framebuffer. Further, you may want to render to multiple windows. For example, in a CAD-type application, you may have a top,

side, front, and perspective view of an object being modeled. It is possible to render to multiple viewports simultaneously with a single graphics pipeline.

Support for multiple viewports is optional. The total number of viewports supported by a device can be determined by checking the `maxViewports` member of its `VkPhysicalDeviceLimits` structure, as returned from a call to **vkGetPhysicalDeviceProperties()**. If multiple viewports are supported, then this will be at least 16 and may be higher. If multiple viewports are not supported, then this field will be 1.

When you are creating the graphics pipeline, the number of viewports the pipeline will use is specified in the `viewportCount` field of the `VkPipelineViewportStateCreateInfo` structure passed through `pViewportState` in `VkGraphicsPipelineCreateInfo`. When the viewports are configured as static state, their parameters are passed through the `pViewports` member of the same structure.

Once a pipeline is in use, its geometry shader can select the viewport index by decorating one of its outputs with the `ViewportIndex` decoration. You can generate this code by writing to the `gl_ViewportIndex` built-in output in a GLSL shader.

All of the vertices in a primitive produced by the geometry shader should have the same viewport index. When the primitive is rasterized, it will use the viewport parameters from the selected viewport. Because geometry shaders can run instanced, a simple way to broadcast geometry to multiple viewports is to run the geometry shader with as many instances as there are viewports in the pipeline, and then, for each viewport, assign the invocation index to the viewport index, causing the version of the geometry to be rendered into the appropriate viewport. Listing 9.27 shows an example of how to do this.

Listing 9.27: Using Multiple Viewports in a Geometry Shader (GLSL)

Click here to view code image

```
#version 450 core

layout (triangles, invocations = 4) in;
layout (triangle_strip, max_vertices = 3) out;

layout (set = 0, binding = 0) uniform transform_block
{
    mat4 mvp_matrix[4];
};

in VS_OUT
{
    vec4 color;
} gs_in[];

out GS_OUT
{
    vec4 color;
} gs_out;

void main(void)
{
```

```
        for (int i = 0; i < gl_in.length(); i++)
    {
        gs_out.color = gs_in[i].color;
        gl_Position = mvp_matrix[gl_InvocationID]*
                        gl_in[i].gl_Position;
        gl_ViewportIndex = gl_InvocationID;
        EmitVertex();
    }
    EndPrimitive();
}
```

In [Listing 9.27](#), the transformation matrix to be applied to the geometry for each shader invocation is stored in a uniform block called `transform_block`. The number of invocations is set to 4, using an input layout qualifier, and then the `gl_InvocationID` built-in variable is used to index into the array of matrices. The invocation index is also used to specify the viewport index.

Summary

This chapter discussed the optional geometry processing stages of the Vulkan graphics pipeline: tessellation and geometry. You saw that the tessellation stage is made up of a pair of shaders surrounding a fixed-function, configurable block that breaks large patches into many smaller points, lines, or triangles. After the tessellation stage is the geometry shading stage, which receives primitives from the previous stage and can process entire primitives, discard them, or create new ones, decimating or amplifying geometry as it passes down the pipeline.

You also saw how geometry can be clipped and culled using per-vertex controls and how selection of the viewport index in a geometry shader can be used to confine geometry to user-specified regions.

Chapter 10. Fragment Processing

What You'll Learn in This Chapter

- What happens after a primitive is rasterized
- How your fragment shaders determine the color of fragments
- How the results of your fragment shaders are merged into the final picture

---

The previous chapters covered everything that occurs in the Vulkan graphics pipeline up until the point of rasterization. The rasterizer takes your primitives and breaks them into many fragments, which ultimately come together to form the final pixels that will be shown to your users. In this chapter, you will see what happens to those fragments as they undergo per-fragment tests, shading and then blending into color attachments used by your application.

Scissor Testing

The scissor test is a stage in fragment processing that operates before any other testing is performed. This test simply ensures that fragments are within a specified rectangle on the framebuffer. While this is somewhat similar to viewport transformations, there are two important differences between using a viewport that is not the full size of the framebuffer and the scissor test:

- The viewport transformation changes the locations of primitives in the framebuffer; as the viewport rectangle moves, so do the primitives inside it. The scissor rectangles have no effect on the the position of primitives and operate after they have been rasterized.
- While the viewport rectangles affect clipping and in some cases may produce new primitives, the scissor rectangle operates directly on rasterized fragments, discarding them before fragment shading.

The scissor test always runs before the fragment shader, so if your fragment shader produces side effects, those side effects will not be seen for fragments that are scissored away. The scissor test is always enabled. However, in most implementations, setting the scissor rectangle to the full size of the framebuffer effectively disables it.

The number of scissor rectangles is specified in the `scissorCount` field of the `VkPipelineViewportStateCreateInfo` structure used to create the graphics pipeline. This must be the same as the `viewportCount` field. As discussed in [Chapter 9](#), "[Geometry Processing](#)," the viewport to use for the viewport transformation is chosen by writing to an output in the geometry shader decorated with the `ViewportIndex` decoration. This same output is used to select the scissor rectangle used for the scissor test. As such, it's not possible to select an arbitrary combination of scissor rectangle and viewport. If you want to use multiple viewports and disable the scissor test, then you need to specify as many scissor rectangles as you do viewports, but set all of the scissor rectangles to the full size of the framebuffer.

Each scissor rectangle is represented by an instance of the `VkRect2D` structure, the definition of which is

```
typedef struct VkRect2D {
    VkOffset2D    offset;
```

```
        VkExtent2D      extent;
    } VkRect2D;
```

A rectangle is made up of an origin and a size, stored in the `offset` and `extent` fields of `VkRect2D`. These are `VkOffset2D` and `VkExtent2D` structures, the definitions of which are

```
    typedef struct VkOffset2D {
        int32_t    x;
        int32_t    y;
    } VkOffset2D;
```

and

```
    typedef struct VkExtent2D {
        uint32_t    width;
        uint32_t    height;
    } VkExtent2D;
```

The `x` and `y` fields of `offset` contain the coordinates, in pixels, of the origin of each scissor rectangle, and the `width` and `height` fields of `extent` contain its size.

As with the viewport rectangles, the number of rectangles accessible to a pipeline is always considered to be static state, but the sizes of the rectangles may be made dynamic. If `VK_DYNAMIC_STATE_SCISSOR` is included in the list of dynamic states when the pipeline is created, then the scissor rectangle state becomes dynamic and can be modified using **vkCmdSetScissor()**.

[**Click here to view code image**](#)

```
    void vkCmdSetScissor (
        VkCommandBuffer                     commandBuffer,
        uint32_t                            firstScissor,
        uint32_t                            scissorCount,
        const VkRect2D*                     pScissors);
```

The index of the first scissor rectangle to update is passed in `firstScissor`, and the number of scissor rectangles to update is passed in `scissorCount`. The range of scissor rectangles can be a subset of the scissor rectangles supported by Vulkan. However, it's important to set all the scissor rectangles that the current pipeline might use before rendering into them. The extent of the scissor rectangles is contained in an array of `VkRect2D` structures, the address of which is passed in `pScissors`.

The scissor test is essentially always enabled. Setting one or more of the scissor rectangles to cover the entire renderable area effectively disables the test for that rectangle. The number of scissor rectangles used by a pipeline is always considered to be part of the pipeline's static state. This number is set through the `scissorCount` member of the `VkPipelineViewportStateCreateInfo` structure used to create the pipeline. The range of scissor rectangles specified in **vkCmdSetScissor()** can extend outside the number of scissor rectangles supported by the currently bound pipeline. If you switch to another pipeline with more scissor rectangles, the rectangles you set will be used.

Depth and Stencil Operations

The depth and stencil buffers are special attachments that allow fragments to be evaluated against information already contained in them either before or after the fragment shader runs. Chapter 7, "Graphics Pipelines," introduced the depth state. Additional state in the rasterization state also controls how rasterized fragments interact with the depth and stencil buffers.

Logically, the depth and stencil operations run after the fragment shader has executed. In practice, most implementations will execute the depth and stencil operations before[1] running the fragment shader wherever they can prove that execution of the fragment shader has no visible side effects, so the result of running tests (and potentially rejecting fragments) will have no visible effect on the rendered scene. In the following sections, we will discuss operation in logical terms—that is, assuming that tests occur after shading—but mention explicitly where this is not the case or where some caveat may preclude an implementation from running tests early.

[1]. An entire class of rendering hardware—deferred shading hardware—will attempt to run the depth test for every fragment in the scene before running any fragment shading. Configuring state to preclude this can have serious performance consequences on this type of hardware.

During testing against the depth and stencil buffers, those buffers can be optionally updated with new data. In fact, writing to depth or stencil buffers is so much considered to be part of the test that the tests must be enabled in order to see writes to depth or stencil buffers occur. This is a common omission for developers new to graphics programming, so it's important to mention it here.

As introduced in Chapter 7, "Graphics Pipelines," depth and stencil operation state is configured using the `VkPipelineDepthStencilStateCreateInfo` structure passed through the `VkGraphicsPipelineCreateInfo` structure used to create the graphics pipeline. For reference, the definition of `VkPipelineDepthStencilStateCreateInfo` is

Click here to view code image

```
typedef struct VkPipelineDepthStencilStateCreateInfo {
    VkStructureType                            sType;
    const void*                                pNext;
    VkPipelineDepthStencilStateCreateFlags     flags;
    VkBool32                                   depthTestEnable;
    VkBool32                                   depthWriteEnable;
    VkCompareOp                                depthCompareOp;
    VkBool32                                   depthBoundsTestEnable;
    VkBool32                                   stencilTestEnable;
    VkStencilOpState                           front;
    VkStencilOpState                           back;
    float                                      minDepthBounds;
    float                                      maxDepthBounds;
} VkPipelineDepthStencilStateCreateInfo;
```

As you can see, this structure is rather large and contains several substructures that define properties of the depth and stencil tests. We necessarily glossed over much of this detail in Chapter 7, "Graphics Pipelines," but dig into more of it here. Each of the substructures is discussed in the following sections.

Depth Testing

The first operation performed after rasterization is the depth test. The depth test compares the depth value for the fragment with the value stored in the current depth-stencil attachment using an operation chosen when the pipeline is created. The fragment's depth value can come from one of two places:

- It may be interpolated as part of rasterization using the depth values at each of the primitive's vertices.
- It may be generated in the fragment shader and output along with other color attachments.

Of course, generating the depth value in the fragment shader is one of the cases in which an implementation would be forced to run the shader before performing the depth test.

The depth value interpolated by rasterization is taken from the result of the viewport transformation. Each vertex has a depth value in the range 0.0, 1.0 as produced from the last stage in the geometry-processing pipeline. This is then scaled and biased using the parameters of the selected viewport to yield a *framebuffer* depth.

The depth test must be enabled by setting the `depthTestEnable` field of the `VkPipelineDepthStencilStateCreateInfo` structure to `VK_TRUE`.

When enabled, the operation used to compare the fragment's computed depth value (either interpolated during rasterization or produced by your fragment shader) is specified in the `depthCompareOp` field of the pipeline's `VkPipelineDepthStencilStateCreateInfo`. This is one of the standard `VkCompareOp` tokens, and their meanings when applied to depth testing are shown in .

If the depth test passes, the resulting depth (whether interpolated or produced by the fragment shader) may be written to the depth buffer. To do this, set `depthWriteEnable` to `VK_TRUE`. If `depthWriteEnable` is `VK_FALSE`, then the depth buffer is not written to, regardless of the outcome of the depth test.

Function VK_COMPARE_OP_...	Meaning
ALWAYS	The depth test always passes; all fragments are considered to have passed the depth test.
NEVER	The depth test never passes; all fragments are considered to have failed the depth test.
LESS	The depth test passes if the new fragment's depth value is less than the old fragment's depth value.
LESS_OR_EQUAL	The depth test passes if the new fragment's depth value is less than or equal to the old fragment's depth value.
EQUAL	The depth test passes if the new fragment's depth value is equal to the old fragment's depth value.
NOT_EQUAL	The depth test passes if the new fragment's depth value is not equal to the old fragment's depth value.
GREATER	The depth test passes if the new fragment's depth value is greater than the old fragment's depth value.
GREATER_OR_EQUAL	The depth test passes if the new fragment's depth value is greater than or equal to the old fragment's depth value.

Table 10.1: Depth Comparison Functions

It is important to note that the depth buffer is *not* updated unless the depth test is enabled, regardless of the state of depthWriteEnable. Therefore, if you want to unconditionally write the fragment's depth value into the depth buffer, you need to enable the test (set depthTestEnable to VK_TRUE), enable depth writes (set depthWriteEnable to VK_TRUE), and configure the depth test to always pass (set depthCompareEnable to VK_COMPARE_OP_ALWAYS).

Depth-Bounds Testing

The depth-bounds test is a special, additional test that can be performed as part of depth testing. The value stored in the depth buffer for the current fragment is compared with a specified *range* of values that is part of the pipeline. If the value in the depth buffer falls within the specified range, then the test passes; otherwise, it fails. What's interesting about the depth-bounds test is that it is not in any way dependent on the depth value of the fragment being tested. This means that it can be quickly evaluated at the same time as or even before depth interpolation or fragment shading executes.

A use case for the depth-bounds test is to intersect volumetric geometry with an existing depth buffer. For example, if we prerender the depth buffer for a scene, we can project a light's sphere of influence into the scene. We then set the depth bounds to the minimum and maximum distance from the light's center and enable the test. When we render the light geometry (using a fragment shader that will perform deferred shading computations), the depth-bounds test will quickly reject fragments that will not be influenced by the light.

To enable the depth-bounds test, set the `depthBoundsTestEnable` member of `VkPipelineDepthStencilStateCreateInfo` to `VK_TRUE`, and configure the minimum and maximum depth-bounds values in `minDepthBounds` and `maxDepthBounds`, respectively. Perhaps more usefully, though, the minimum and maximum extents for the depth-bounds test can be configured as dynamic state.

To do this, include `VK_DYNAMIC_STATE_DEPTH_BOUNDS` as one of the dynamic states in the pipeline's `VkPipelineDynamicStateCreateInfo` structure. Once the depth-bounds test is set to dynamic, the minimum and maximum extents for the test are set with **`vkCmdSetDepthBounds()`**, the prototype of which is

[**Click here to view code image**](#)

```
void vkCmdSetDepthBounds (
    VkCommandBuffer                         commandBuffer,
    float                                   minDepthBounds,
    float                                   maxDepthBounds);
```

Again, for the depth-bounds test to take effect, it must be enabled by setting the `depthBoundsTestEnable` member of the `VkPipelineDepthStencilStateCreateInfo` structure used to create the pipeline to `VK_TRUE`. The `minDepthBounds` and `maxDepthBounds` parameters take similar meaning to the similarly named parameters in the `VkPipelineDepthStencilStateCreateInfo` structure. Whether the depth-bounds test is enabled or not is *always* static state, even if the depth-bounds values themselves are marked as dynamic.

Note that the depth-bounds test is an optional feature, and not all Vulkan implementations support it. To determine whether the Vulkan implementation supports the depth-bounds test, check the `depthBounds` member of the device's `VkPhysicalDeviceFeatures` structure. In order to use the depth-bounds test, you should also set the `depthBounds` field of the `VkPhysicalDeviceFeatures` structure used to create the device to `VK_TRUE`.

Depth Bias

When two primitives are rendered on top of each other or very close to it, their interpolated depth values may be the same or very close. If the primitives are exactly coplanar, then their interpolated depth values should be identical. If there is any deviation, then their depth values will differ by a very small amount. As the interpolated depth values are subject to floating-point imprecision, this can result in depth testing producing inconsistent and implementation-dependent results. The visual artifact that results is known as *depth fighting*.

To counteract this, and to enable predictable results from depth testing generally, a programmable bias value can be applied to the interpolated depth values to force them to be offset toward or away from the viewer. This is known as depth bias and is actually part of rasterization state because it is the rasterizer that produces the interoplated depth values.

To enable depth bias, set the `depthBiasEnable` field of the `VkPipelineRasterizationStateCreateInfo` structure used to create the rasterizer state in the graphics pipeline to `VK_TRUE`. For reference, the definition of `VkPipelineRasterizationStateCreateInfo` is

[**Click here to view code image**](#)

```
typedef struct VkPipelineRasterizationStateCreateInfo {
    VkStructureType                            sType;
    const void*                                pNext;
    VkPipelineRasterizationStateCreateFlags    flags;
    VkBool32                                   depthClampEnable;
    VkBool32                                   rasterizerDiscardEnable;
    VkPolygonMode                              polygonMode;
    VkCullModeFlags                            cullMode;
    VkFrontFace                                frontFace;
    VkBool32                                   depthBiasEnable;
    float                                      depthBiasConstantFactor;
    float                                      depthBiasClamp;
    float                                      depthBiasSlopeFactor;
    float                                      lineWidth;
} VkPipelineRasterizationStateCreateInfo;
```

When `depthBiasEnable` is `VK_TRUE`, the next three fields—`depthBiasConstantFactor`, `depthBiasClamp`, and `depthBiasSlopeFactor`—configure the depth-bias equations used to compute an offset that is applied to the interpolated depth values.

First, the maximum depth slope of the polygon m is computed as

$$m = \sqrt{\left(\frac{\delta z_f}{\delta x_f}\right)^2 + \left(\frac{\delta z_f}{\delta y_f}\right)^2} \qquad (10.1)$$

This is often approximated as

$$m = \max\left\{\left|\frac{\delta z_f}{\delta x_f}\right|, \left|\frac{\delta z_f}{\delta y_f}\right|\right\} \qquad (10.2)$$

In both equations, x_f, y_f, and z_f represent a point on the triangle.

Once m is computed, the offset o is computed as

$$o = m \times \mathtt{depthBiasSlopeFactor} + r \times \mathtt{depthBiasConstantFactor} \quad (10.3)$$

The term r is the minimum resolvable delta in depth values for the primitive. For fixed-point depth buffers, this is a constant that depends only on the number of bits in the depth buffer. For floating-point depth buffers, this depends on the range of depth values spanned by the primitive.

If `depthBiasClamp` in `VkPipelineRasterizationStateCreateInfo` is 0.0 or NaN, then the computed value of o is used to directly bias the interpolated depth values. However, if `depthBiasClamp` is positive, then it forms an upper bound on o, and if it is negative, it forms a lower bound on o.

By biasing the depth value toward or away from the viewer slightly, it's possible to ensure that for a given pair of coplanar (or almost coplanar) primitives, one always "wins" the depth test. This can essentially eliminate depth fighting.

The depth-bias parameters can either be static state contained in the rasterization state of the graphics pipeline or configured to be dynamic state. To enable the depth-bias state to be dynamic, include `VK_DYNAMIC_STATE_DEPTH_BIAS` in the list of dynamic states passed in the `pDynamicStates` member of the `VkPipelineDynamicStateCreateInfo` structure used to create the graphics pipeline.

Once depth bias is configured to be dynamic, the parameters for depth-bias equations can be set by calling **vkCmdSetDepthBias()**, the prototype of which is

[Click here to view code image](#)

```
void vkCmdSetDepthBias (
    VkCommandBuffer                         commandBuffer,
    float                                   depthBiasConstantFactor,
    float                                   depthBiasClamp,
    float                                   depthBiasSlopeFactor);
```

The command buffer to set the depth-bias state for is specified in `commandBuffer`. The `depthBiasConstantFactor`, `depthBiasClampFactor`, and `depthBiasSlopeFactor` have the same meanings as the similarly named members of `VkPipelineRasterizationStateCreateInfo`. Note that while the depth bias parameters can be made part of a pipeline's dynamic state, the depth bias must still be enabled by setting the `depthBiasEnable` flag to `VK_TRUE` in the `VkPipelineRasterizationStateCreateInfo` structure used to create the pipeline. The `depthBiasEnable` flag is always considered to be static state, although setting the depth-bias factors to 0.0 effectively disables it.

Stencil Testing

The stencil test is enabled by setting the `stencilTestEnable` field of `VkPipelineDepthStencilStateCreateInfo` to `VK_TRUE`.

The stencil test can actually be different for front- and back-facing primitives. The state of the stencil test is represented by an instance of the `VkStencilOpState` structure, and there is one for each of the front- and back-facing states. The definition of `VkStencilOpState` is

```
typedef struct VkStencilOpState {
    VkStencilOp    failOp;
    VkStencilOp    passOp;
    VkStencilOp    depthFailOp;
    VkCompareOp    compareOp;
    uint32_t       compareMask;
    uint32_t       writeMask;
    uint32_t       reference;
} VkStencilOpState;
```

There are three possible outcomes between the depth and stencil tests. If the depth test fails, then the operation specified in `depthFailOp` is performed, and the stencil test is skipped. If the depth test passes, however, then the stencil test is performed, producing one of two further outcomes: If the stencil test fails, then the operation specified in `failOp` is performed, and if the stencil test passes, the operation specified in `passOp` is performed. Each of `depthFailOp`, `failOp`, and `passOp` is a member of the `VkStencilOp` enumeration. The meaning of each operation is shown in Table 10.2.

If enabled, the stencil test compares the stencil reference value with that of the current content of the stencil buffer. The operator used to compare the two values is specified in the `compareOp` field of the `VkStencilOpState` structure. This is the same set of values used for depth testing, only comparing the stencil reference with the content of the stencil buffer.

Function	Result
KEEP	Do not modify the stencil buffer.
ZERO	Set stencil buffer value to 0.
REPLACE	Replace stencil value with reference value.
INCR_AND_CLAMP	Increment stencil with saturation.
DECR_AND_CLAMP	Decrement stencil with saturation.
INVERT	Bitwise-invert stencil value.
INCR_AND_WRAP	Increment stencil without saturation.
DECR_AND_WRAP	Decrement stencil without saturation.

Table 10.2: Stencil Operations

The reference value, compare mask value, and write mask value can also be made dynamic. To mark these states as dynamic, include `VK_DYNAMIC_STATE_STENCIL_REFERENCE`, `VK_DYNAMIC_STATE_STENCIL_COMPARE_MASK`, and `VK_DYNAMIC_STATE_STENCIL_WRITE_MASK`, respectively, in the list of dynamic states passed in the `pDynamicStates` member of the `VkPipelineDynamicStateCreateInfo` structure used to create the graphics pipeline.

Parameters of the stencil test can be set using three functions: **vkCmdSetStencilReference()**, **vkCmdSetStencilCompareMask()**, and **vkCmdSetStencilWriteMask()**. They set the stencil reference value and the compare and write masks, respectively. Their prototypes are

[**Click here to view code image**](#)

```
void vkCmdSetStencilReference (
    VkCommandBuffer                       commandBuffer,
    VkStencilFaceFlags                    faceMask,
    uint32_t                              reference);

void vkCmdSetStencilCompareMask (
    VkCommandBuffer                       commandBuffer,
    VkStencilFaceFlags                    faceMask,
    uint32_t                              compareMask);

void vkCmdSetStencilWriteMask (
    VkCommandBuffer                       commandBuffer,
    VkStencilFaceFlags                    faceMask,
    uint32_t                              writeMask);
```

Each of these three states can be independently made either static or dynamic. All three functions take a `faceMask` parameter, which determines whether the new state applies to front-facing primitives, back-facing primitives, or both. To specify each of these, set `faceMask` to `VK_STENCIL_FACE_FRONT_BIT`, `VK_STENCIL_FACE_BACK_BIT`, or `VK_STENCIL_FRONT_AND_BACK`, respectively. Note that `VK_STENCIL_FRONT_AND_BACK` has the same numeric value as `VK_STENCIL_FACE_FRONT_BIT` and `VK_STENCIL_FACE_BACK_BIT` added together.

Early Fragment Tests

Normally, the depth and stencil tests are configured to run *after* the fragment shader has run. Because updating the depth buffer occurs as part of the test—the test is considered to be an indivisible read-modify-write operation—this has three primary consequences that are visible to the application:

- The fragment shader can update the depth value of the fragment by writing to a floating-point output decorated with the `BuiltIn FragDepth` decoration. In this case, the depth test executes after the fragment shader and uses the value produced by the shader rather than the value generated by interpolation.
- The fragment shader has side effects such as storing into an image. In this case, the side effects from shader execution will be visible, even for fragments that later fail the depth test.
- If the fragment shader throws away the fragment by executing the `OpKill` SPIR-V instruction (produced by calling the **discard()** function in GLSL), then the depth buffer is not updated.

If none of these conditions can occur, Vulkan implementations are permitted to detect this case and reorder testing such that the depth and stencil tests occur before the fragment shader executes. This saves the work of executing the fragment shader for fragments that would fail the tests.

Because the default logical order of the pipeline specifies that the depth and stencil tests occur after the fragment shader has executed, when any of the conditions hold, the implementation *must* run operations in that order. You can, however, force the shader to run after the depth and stencil tests even when one of these conditions is true. There are two ways to do this.

First, you can force the fragment shader to run early regardless of any other condition detected by the Vulkan implementation. To do this, decorate the SPIR-V entry point for the fragment shader with the `EarlyFragmentTests` decoration. This will cause execution of the fragment shader to be delayed (or at least appear to be delayed) until after the depth and stencil tests have been executed. Should the depth test fail, then the fragment shader will not be executed, and any side effects that it may have produced, such as updates to an image via image stores, will not be visible.

This first mode is useful, for example, to run depth tests against a prerendered depth image produced, say, in an earlier subpass of a renderpass, but to produce output from the fragment shader using image stores rather than normal fragment shader outputs. Because the number of fixed-function outputs from the fragment shader is limited but the number of stores it may perform to images is effectively unlimited, this is a good way to increase the amount of data that a fragment shader can produce for a single fragment.

The second mechanism for running the fragment shader after depth testing is more subtle. It interacts specifically with the case where the fragment shader writes to the fragment's depth value. Under normal circumstances, the implementation must assume that the fragment shader may write any arbitrary value to the fragment's depth unless it can definitively prove otherwise.

In cases where you know that the fragment shader will move the depth value in only one direction (toward or away from the viewer), you can apply the SPIR-V `DepthGreater` or `DepthLess` execution modes to the fragment shader's entry point with the `OpExecutionMode` instruction.

When `DepthGreater` is applied, then Vulkan knows that no matter what your shader does, the resulting depth values produced by the fragment shader will only be *greater* than the values produced by interpolation. Therefore, if the depth test is `VK_COMPARE_OP_GREATER` or `VK_COMPARE_OP_GREATER_OR_EQUAL`, then the fragment shader cannot negate the result of a depth test that's already passed.

Likewise, when `DepthLess` is applied, then Vulkan knows that the fragment shader will only make the resulting depth value *less* than it would have been and therefore cannot negate the result of a passing `VK_COMPARE_OP_LESS` or `VK_COMPARE_OP_LESS_OR_EQUAL` test.

Finally, the SPIR-V `DepthUnchanged` execution mode tells Vulkan that no matter what the fragment shader appears to do to the depth value of the fragment shader, it should treat it as though it did not update the value at all. Therefore, any optimizations it might make if the shader did not write to the fragment's depth remain valid and can be enabled where appropriate.

Multisample Rendering

Multisample rendering is a method to improve image quality by storing multiple depth, stencil, or color values for each pixel. As the image is rendered, multiple samples within each pixel are generated (hence the term *multisample* rendering) and when the image is to be displayed to the user, the samples are merged using a filter to produce a final, single value per pixel.

There are two general ways to produce a multisample image:

- Multisampling: Determining which samples within a pixel are covered, computing a single color value for that pixel, and then broadcasting that value to all covered samples within the pixel
- Supersampling: Computing unique color values for each and every sample within a pixel

Obviously, supersampling comes at significantly higher cost than multisampling because potentially many more color values are computed and stored.

To create a multisample image, set the `samples` field of the `VkImageCreateInfo` structure used to create the image to one of the values of the `VkSampleCountFlagBits` enumeration. The `samples` field is a *one hot* encoding of the sample count to be used in the image. The numeric value for the enumeration is simply *n*, where *n* is the number of samples in the image. Only power-of-two sample counts are supported.

While the current Vulkan header defines enumerants from `VK_SAMPLE_COUNT_1_BIT` to `VK_SAMPLE_COUNT_64_BIT`, most Vulkan implementations will support sample counts between 1 and 8 or maybe 16. Further, not every sample count between 1 and the maximum sample count supported by an implementation will be supported for every image format. In fact, support for different sample counts varies per format and even per tiling mode and should be queried using **vkGetPhysicalDeviceFormatProperties()**. As discussed in Chapter 2, "Memory and Resources," **vkGetPhysicalDeviceImageFormatProperties()** returns information about a particular format. For reference, its prototype is

Click here to view code image

```
VkResult vkGetPhysicalDeviceImageFormatProperties (
    VkPhysicalDevice                        physicalDevice,
    VkFormat                                format,
    VkImageType                             type,
    VkImageTiling                           tiling,
    VkImageUsageFlags                       usage,
    VkImageCreateFlags                      flags,
    VkImageFormatProperties*                pImageFormatProperties);
```

The format to be queried is passed along with the image type, tiling mode, usage flags and other information, and **vkGetPhysicalDeviceImageFormatProperties()** writes the support information about the format into the `VkImageFormatProperties` structure pointed to by `pImageFormatProperties`. The `sampleCounts` field of this structure then contains the supported sample counts for an image in the specified format with the configuration described by the remaining parameters.

As a primitive is rasterized, the rasterizer computes coverage information for the pixels hit by it. When multisample rendering is off, rasterization works by simply determining whether the center of each pixel is inside the primitive and, if so, considering the pixel hit. The depth and stencil tests are then performed at the center of the pixel, and the fragment shader is executed to determine the resulting shading information for that pixel.

When multisample rendering is enabled, coverage is determined for each sample in each pixel. For each sample that is considered to be inside the pixel, the depth and stencil tests are computed individually. If any sample is determined to be visible, then the fragment shader is run *once* to determine the output color for the fragment. That output color is then written into every visible sample in the image.

By default, the samples are distributed evenly within a pixel at a set of standard locations. These locations are illustrated in Figure 10.1. The standard sample locations shown in the figure are supported if the `standardSampleLocations` field of the device's `VkPhysicalDeviceLimits` structure is `VK_TRUE`. In this case, the sample locations shown in Figure 10.1 are used for 1-, 2-, 4-, 8-, and 16-sample images (if supported). Even if the device

supports 32, 64, or higher sample count images, there are no defined standard sample locations for those sample counts.
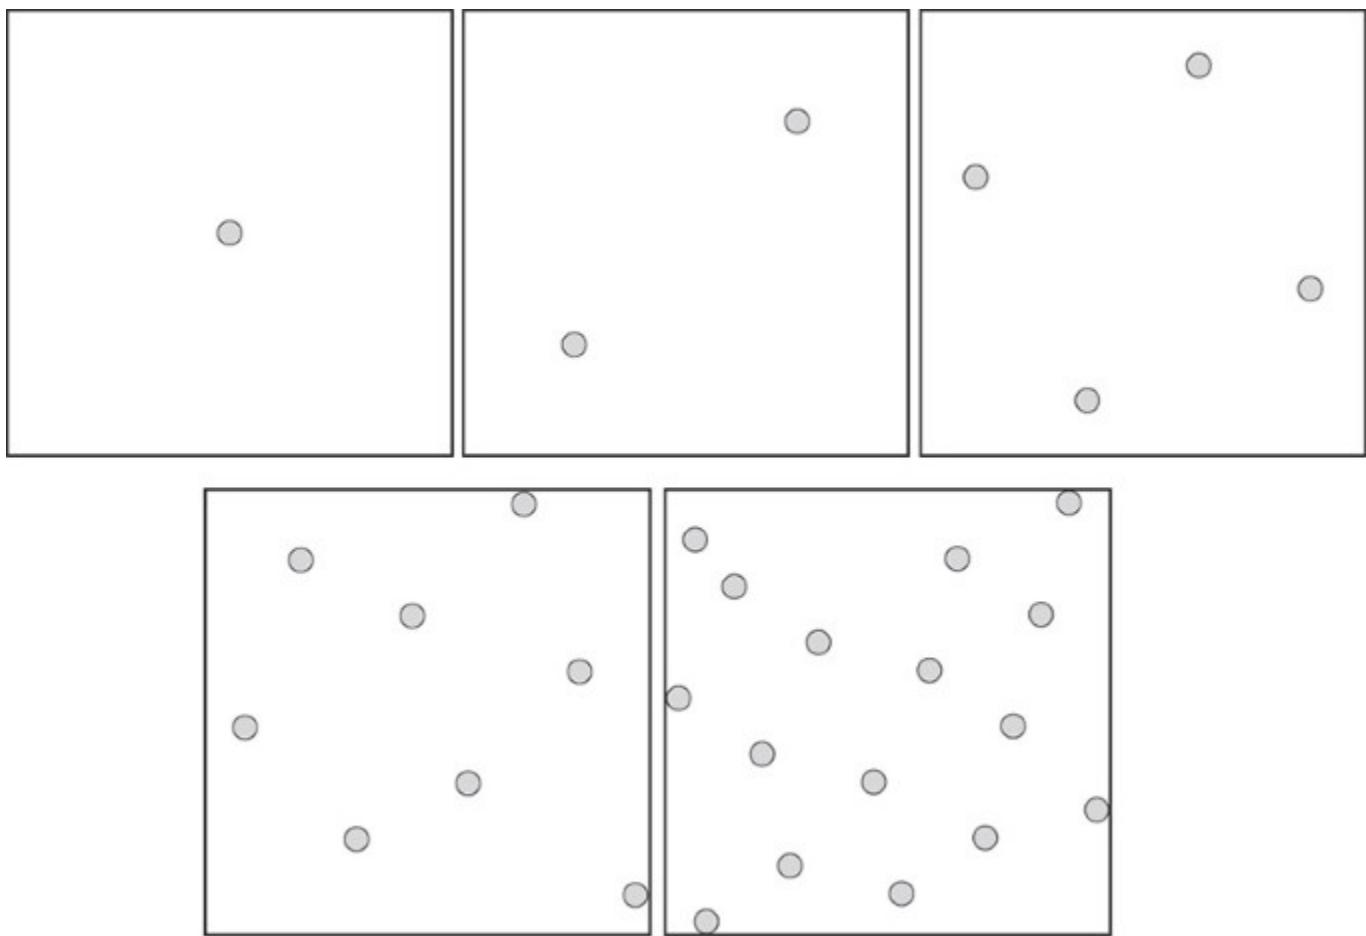
Figure 10.1: Standard Sample Locations

If `standardSampleLocations` is `VK_FALSE` in the `VkPhysicalDeviceLimits` structure, then those sample counts may be supported with nonstandard sample locations. Determining the sample locations used by a device in this scenario would likely require a device extension.

Sample Rate Shading

Usually, when a multisample image is bound to the framebuffer, multisampling is used to render into it. As described earlier, coverage is determined at each of the sample locations in each pixel, and if any of them is covered by the primitive, the fragment shader is run once to determine the values of its outputs. Those values are then broadcast to all covered samples in the pixel.

For even higher image quality, we can use supersampling, which causes the fragment shader to be executed for each covered sample and the unique output values it produces to be written directly to the pixel's samples.

While multisampling can improve the appearance of the edges of primitives, it can do nothing to prevent aliasing artifacts resulting from high spacial frequencies produced by the fragment shader. However, by running a fragment shader at sample rate, high frequencies produced by the fragment shader can be antialiased along with primitive edges.

To enable sample rate shading, set the `sampleShadingEnable` member of the `VkPipelineMultisampleStateCreateInfo` structure used to create the graphics pipeline to `VK_TRUE`. When `sampleShadingEnable` is enabled, the `minSampleShading` field further controls the frequency at which the shader will be executed. This is a floating-point value between 0.0 and 1.0. *At least* this fraction of the samples in the pixel will receive a unique set of values as produced by a separate invocation of the fragment shader.

If `minSampleShading` is 1.0, then every sample in the pixel is guaranteed to receive its own data from a separate invocation of the fragment shader. Any value less than 1.0 causes the device to compute colors for at least that many samples and then distribute the values to the samples in a device-dependent manner. For example, if `minSampleShading` is 0.5, then a fully covered 8-sample pixel would cause the fragment shader to execute *at least* 4 times and for the 4 sets of outputs to be distributed to the 8 samples in the pixel in a device-dependent manner.

It is possible to configure the pipeline to update only a subset of the samples in the framebuffer. You can specify a sample mask, which is a bitmask with a single bit corresponding to each of the samples in the framebuffer. This mask is specified by passing an array of unsigned 32-bit integers through the `pSampleMask` member of `VkPipelineMultisampleStateCreateInfo`. The array divides the arbitrarily long bitmask into 32-bit chunks, each chunk being an element of the array. If the framebuffer has 32 or fewer samples, then the array becomes a single element long.

The N^{th} sample is represented by the $(N\%32)^{th}$ bit of the $(N/32)^{th}$ element of the array. When the bit is set, the sample may be updated by the pipeline. If the bit is cleared, then the pipeline will not modify the content of samples at that index. In effect, the sample mask is logically ANDed with the coverage computed during rasterization. If `pSampleMask` is `nullptr`, then all samples are enabled for writing.

In addition to generating coverage during rasterization, it is possible for a fragment shader to produce a pseudo-coverage value by writing into the alpha channel of its output. By setting the `alphaToCoverage` field of `VkPipelineMultisampleStateCreateInfo` to `VK_TRUE`, the alpha value written into the fragment shader's first output is used to create a new coverage value for the fragment. A fraction of the samples in the temporarily created sample mask are enabled in an implementation-defined order, according to the value written by the shader into the alpha channel. This is then logically ANDed with the coverage mask computed during rasterization. This allows simple transparency and coverage effects to be implemented by simply exporting alpha values from the shader.

The question then remains what to do with the actual alpha channel in the framebuffer. If the shader writes a fraction of coverage into the alpha channel, it may not make sense to write this value directly into the framebuffer along with the RGB color data. Instead, you can ask Vulkan to subsitute a value of 1.0 for alpha for the other alpha-related operations, as though the shader had not produced an alpha value at all. To do this, set the `alphaToOneEnable` field to `VK_TRUE`.

Multisample Resolves

When rendering to a multisample image is completed, it can be *resolved* to a single sample image. This process aggregates all of the values stored in the multiple samples for each pixel into a single value, producing a nonmultisampled image. There are two ways to resolve an image.

The first is to include a nonmultisampled image in the `pResolveAttachments` array passed to through the `VkSubpassDescription` used to create a subpass that would render to the original

multisample color attachments. The subpass will produce a multisample image into the corresponding color attachment, and when the subpass ends (or at least by the time the renderpass ends), Vulkan will automatically resolve the multisample image into the corresponding nonmultisampled image in `pResolveAttachments`.

You can then set the `storeOp` for the multisample image to `VK_ATTACHMENT_STORE_OP_DONT_CARE` to discard the original multisampled data, assuming that it is not needed for something else. This is likely the most efficient way of resolving a multisample image, as many implementations will be able to fold the resolve operation into some other internal operations that are already part of the renderpass, or at least perform the resolve operation while the original multisample data is in caches rather than writing it all into memory and reading it back later.

It's also possible to explicitly resolve a multisample image into a single sample image by calling **`vkCmdResolveImage()`**, the prototype of which is

[**Click here to view code image**](#)

```
void vkCmdResolveImage (
    VkCommandBuffer                 commandBuffer,
    VkImage                         srcImage,
    VkImageLayout                   srcImageLayout,
    VkImage                         dstImage,
    VkImageLayout                   dstImageLayout,
    uint32_t                        regionCount,
    const VkImageResolve*           pRegions);
```

The command buffer that will perform the resolve operation is passed in `commandBuffer`. The source image and its expected layout at the time of the resolve operation are passed in `srcImage` and `srcImageLayout`, respectively. Likewise, the destination image and its expected layout are passed in `dstImage` and `dstImageLayout`. **`vkCmdResolveImage()`** behaves much like the blit and copy operations discussed in [Chapter 4](#), "[Moving Data](#)." As such, the layout of the source image should be either `VK_IMAGE_LAYOUT_GENERAL` or `VK_IMAGE_LAYOUT_TRANSFER_SRC_OPTIMAL`, and the layout of the destination image should be `VK_IMAGE_LAYOUT_GENERAL` or `VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL`.

As with the other blit and copy commands, **`vkCmdResolveImage()`** can resolve only parts of an image if needed. The number of regions to resolve is passed in `regionCount`, and `pRegions` points to an array of `VkImageResolve` structures, each defining one of the regions. The definition of `VkImageResolve` is

[**Click here to view code image**](#)

```
typedef struct VkImageResolve {
    VkImageSubresourceLayers    srcSubresource;
    VkOffset3D                  srcOffset;
    VkImageSubresourceLayers    dstSubresource;
    VkOffset3D                  dstOffset;
    VkExtent3D                  extent;
} VkImageResolve;
```

Each region to resolve is represented by one instance of the `VkImageResolve` structure. Each structure contains a description of the regions in the source and destination regions. The subresource

from which source data is to be taken is in the `srcSubresource` field, and the subresource to which the resolved image data is to be written is in the `dstSubresource` field. **vkCmdResolveImage()** cannot rescale the multisample image, so the sizes of the source and destination regions are the same. This is stored in the `extent` member of the `VkImageResolve` structure. However, the region in the source and destination image does not need to be at the same location in the two images. The origin of the region in the source image is stored in the `srcOffset` field, and the origin of the output region is stored in the `dstOffset` field.

When an image is resolved by specifying it as one of the resolve attachments in the renderpass, the entire region is resolved at the end of the subpass that references it—at least the region that is included in the `renderArea` passed to **vkCmdBeginRenderPass()**. With **vkCmdResolveImage()**, however, it's possible to resolve parts of an image. Although explicitly calling **vkCmdResolveImage()** is likely to be less efficient than resolving attachments at the end of a renderpass, it may be that you need to resolve only part of the image, so **vkCmdResolveImage()** is the more appropriate choice in that case.

Logic Operations

Logic operations allow logical operations such as `AND` and `XOR` to be applied between the output of the fragment shader and the content of a color attachment. Logical operations are supported on most integer formats that can be used as color attachments. Logical operations are enabled by setting the `logicOpEnable` field of the `VkPipelineColorBlendStateCreateInfo` used to create a color blend state object to `VK_TRUE`. Introduced in Chapter 7, "Graphics Pipelines," the definition of `VkPipelineColorBlendStateCreateInfo` is

Click here to view code image

```
typedef struct VkPipelineColorBlendStateCreateInfo {
    VkStructureType                               sType;
    const void*                                   pNext;
    VkPipelineColorBlendStateCreateFlags          flags;
    VkBool32                                      logicOpEnable;
    VkLogicOp                                     logicOp;
    uint32_t                                      attachmentCount;
    const VkPipelineColorBlendAttachmentState*    pAttachments;
    float                                         blendConstants[4];
} VkPipelineColorBlendStateCreateInfo;
```

When logical operations are enabled, the logical operation specified by `logicOp`, is performed between each of the fragment shader's outputs and the corresponding color attachment. The same logical operation is used for every attachment—that is, it is not possible to use a different logical operation for different attachments. The available logical operations, represented by members of the `VkLogicOp`, are shown in Table 10.3.

Operation (VK_LOGIC_OP_...)	Result
CLEAR	Set all values to 0
AND	Source & Destination
AND_REVERSE	Source & ~Destination
COPY	Source
AND_INVERTED	~Source & Destination
NOOP	Destination
XOR	Source ^Destination
OR	Source | Destination
NOR	~(Source | Destination)
EQUIV	~(Source ^Destination)
INVERT	~Destination
OR_REVERSE	Source | ~Destination
COPY_INVERTED	~Source
OR_INVERTED	~Source | Destination
NAND	~(Source & Destination)
SET	Set all values to 1

Table 10.3: Logic Operations

In Table 10.3, *Source* refers to the value(s) produced by the fragment shader, and *Destination* refers to the values already in the color attachment.

Although the selected logical operation is global (if enabled) and applies to all color outputs, logical operations are applied only to color attachments with formats that support them. If the format used for one of the color attachments does not support logical operations, then the selected logic operation is ignored for that attachment, and the value is written directly to the attachment.

Fragment Shader Outputs

Each fragment shader may have one or more outputs. It is possible to construct a fragment shader that has no outputs and produces visible side effects by performing image store operations. However, in the majority of cases, fragment shaders produce outputs by writing them to special variables declared as outputs.

To declare an output in a GLSL fragment, simply create a variable at global scope with the **out** storage qualifier. This produces a SPIR-V declaration of an output variable, which in turn is used by Vulkan to connect the fragment shader's output to subsequent processing. Listing 10.1 shows a simple GLSL fragment shader that declares a single output and writes an opaque red color into it. The resulting SPIR-V shader is shown in Listing 10.2.

Listing 10.1: Declaring an Output in a Fragment Shader (GLSL)

```
#version 450 core
out vec4 o_color;
void main(void)
{
    o_color = vec4(1.0f, 0.0f, 0.0f, 1.0f);
}
```

When the code in Listing 10.1 is used to create the SPIR-V shader in Listing 10.2, you can see that the o_color output is translated into an output variable (%9) of type vector of four floating-point values (%7). The OpStore instruction is then used to write to it.

Listing 10.2: Declaring an Output in a Fragment Shader (SPIR-V)

```
; SPIR-V
; Version: 1.0
; Generator: Khronos Glslang Reference Front End; 1
; Bound: 13
; Schema: 0
                OpCapability Shader
           %1 = OpExtInstImport "GLSL.std.450"
                OpMemoryModel Logical GLSL450
                OpEntryPoint Fragment %4 "main" %9
                OpExecutionMode %4 OriginUpperLeft
                OpSource GLSL 450
                OpName %4 "main"
                OpName %9 "o_color"
           %2 = OpTypeVoid
           %3 = OpTypeFunction %2
           %6 = OpTypeFloat 32
           %7 = OpTypeVector %6 4
           %8 = OpTypePointer Output %7
           %9 = OpVariable %8 Output
          %10 = OpConstant %6 1
          %11 = OpConstant %6 0
          %12 = OpConstantComposite %7 %10 %11 %11 %10
           %4 = OpFunction %2 None %3
           %5 = OpLabel
                OpStore %9 %12
                OpReturn
                OpFunctionEnd
```

When the framebuffer color attachment is floating-point, or signed- or unsigned-normalized format, the fragment shader output should be declared using a floating-point variable. When the color attachment is a signed or unsigned integer format, then the fragment shader output should be declared as a signed or unsigned integer type too. The number of components in the fragment shader output should be at least as many as the number of components in the corresponding color attachment.

As you saw in [Chapter 7](#), "[Graphics Pipelines](#)," a single graphics pipeline can have many color attachments and access these attachments through attachment references in each subpass. Each subpass can reference several output attachments, which means that a fragment shader can write to several attachments. To do this, we declare multiple outputs in our fragment shader and, in GLSL, specify an output *location* using a `location` layout qualifier.

[Listing 10.3](#) shows a GLSL fragment shader that declares multiple outputs, writing a different constant color to each one. Each is assigned a different location using a `location` layout qualifier.

Listing 10.3: Several Outputs in a Fragment Shader (GLSL)

Click here to view code image

```
#version 450 core

layout (location = 0) out vec4 o_color1;
layout (location = 1) out vec4 o_color2;
layout (location = 5) out vec4 o_color3;

void main(void)
{
    o_color1 = vec4(1.0f, 0.0f, 0.0f, 1.0f);
    o_color2 = vec4(0.0f, 1.0f, 0.0f, 1.0f);
    o_color3 = vec4(0.0f, 0.0f, 1.0f, 1.0f);
}
```

The result of compiling the shader in [Listing 10.3](#) to SPIR-V is shown in [Listing 10.4](#).

Listing 10.4: Several Outputs in a Fragment Shader (SPIR-V)

Click here to view code image

```
; SPIR-V
; Version: 1.0
; Generator: Khronos Glslang Reference Front End; 1
; Bound: 17
; Schema: 0
             OpCapability Shader
        %1 = OpExtInstImport "GLSL.std.450"
             OpMemoryModel Logical GLSL450
             OpEntryPoint Fragment %4 "main" %9 %13 %15
             OpExecutionMode %4 OriginUpperLeft
             OpSource GLSL 450
             OpName %4 "main"
             OpName %9 "o_color1"
             OpName %13 "o_color2"
             OpName %15 "o_color3"
             OpDecorate %9 Location 0
             OpDecorate %13 Location 1
             OpDecorate %15 Location 5
        %2 = OpTypeVoid
        %3 = OpTypeFunction %2
        %6 = OpTypeFloat 32
```

```
 %7 = OpTypeVector %6 4
 %8 = OpTypePointer Output %7
 %9 = OpVariable %8 Output
%10 = OpConstant %6 1
%11 = OpConstant %6 0
%12 = OpConstantComposite %7 %10 %11 %11 %10
%13 = OpVariable %8 Output
%14 = OpConstantComposite %7 %11 %10 %11 %10
%15 = OpVariable %8 Output
%16 = OpConstantComposite %7 %11 %11 %10 %10
 %4 = OpFunction %2 None %3
 %5 = OpLabel
      OpStore %9 %12
      OpStore %13 %14
      OpStore %15 %16
      OpReturn
      OpFunctionEnd
```

As you can see in Listing 10.4, the outputs are declared as normal but are decorated using OpDecorate instructions with the locations assigned in the original GLSL shader of Listing 10.3.

The locations assigned in the fragment shader do not have to be contiguous. That is, you can leave gaps. The shader in Listing 10.3 assigns to outputs at locations 0, 1, and 5, leaving locations 2, 3, and 4 (and locations beyond 5) unwritten. However, it is best practice to tightly assign locations to fragment shader outputs.

Further, the maximum location that can be assigned to a fragment shader output is device-dependent. All Vulkan devices support writing to at least four color attachments from a fragment shader. The actual maximum supported by a device can be determined by checking the maxFragmentOutputAttachments member of the device's VkPhysicalDeviceLimits structure as returned from a call to **vkGetPhysicalDeviceProperties()**. Most desktop-class hardware will support writing to eight or possibly more color attachments.

Fragment shader outputs may also be aggregated into arrays. Of course, all elements in the array have the same type, so this method is suitable only for writing to multiple color attachments with the same types. When a fragment shader declares an output as an array, the first element in that array consumes the location assigned by the location layout qualifier, and each subsequent element in the array consumes a consecutive location. In the resulting SPIR-V shader, this produces a single output variable declared as an array with a single OpDecorate instruction. When writing to the outputs, an OpAccessChain instruction is used to dereference the appropriate element of the array, the result of which is passed to the OpStore instruction as normal.

Color Blending

Blending is the process of merging the outputs of the fragment shader into the corresponding color attachments. When blending is disabled, the outputs of the fragment shader are simply written into the color attachment unmodified, and the original content of the attachment is overwritten. However, when blending is enabled, the value written into the color attachment becomes a function of the value produced by the shader, some configurable constants, and the value already in the color attachment. This allows effects such as transparency, accumulation, translucency, and so on to be implemented using fast, often fixed-function hardware acceleration rather than shader code. Blending also happens entirely in order, whereas read-modify-write operations in a fragment shader may execute out of order with respect to the primitives that they are part of.

Blending, whether it's enabled, and the parameters of the equation are controlled on a per-attachment basis using the `VkPipelineColorBlendAttachmentState` structure passed for the attachment in the `VkPipelineColorBlendStateCreateInfo` structure used to create a color blend state object. Introduced in [Chapter 7](#), "[Graphics Pipelines](#)," the definition of `VkPipelineColorBlendAttachmentState` is

[**Click here to view code image**](#)

```
typedef struct VkPipelineColorBlendAttachmentState {
    VkBool32                blendEnable;
    VkBlendFactor           srcColorBlendFactor;
    VkBlendFactor           dstColorBlendFactor;
    VkBlendOp               colorBlendOp;
    VkBlendFactor           srcAlphaBlendFactor;
    VkBlendFactor           dstAlphaBlendFactor;
    VkBlendOp               alphaBlendOp;
    VkColorComponentFlags   colorWriteMask;
} VkPipelineColorBlendAttachmentState;
```

If the `colorBlendEnable` member of each `VkPipelineColorBlendAttachmentState` structure is `VK_TRUE`, then blending is applied to the corresponding color attachment. The outputs of the fragment shader are considered to be the *source*, and the existing content of the color attachment is considered to be the *destination*. First, the source color (R, G, and B) channels are multiplied by a factor specified in the `srcColorBlendFactor`. Likewise, the destination color channels are multiplied by factors specified by `dstColorBlendFactor`. The results of these multiplications are then combined using an operation specified by `colorBlendOp`.

Operation (VK_BLEND_OP_...)	RGB	Alpha
ADD	$S_{rgb} * RGB_s + $ $D_{rgb} * RGB_d$	$S_a * A_s + $ $D_a * A_d$
SUBTRACT	$S_{rgb} * RGB_s - $ $D_{rgb} * RGB_d$	$S_a * A_s - $ $D_a * A_d$
REVERSE_SUBTRACT	$D_{rgb} * RGB_d - $ $S_{rgb} * RGB_s$	$D_a * A_d - $ $S_a * A_s$
MIN	$\min(RGB_s, RGB_d)$	$\min(A_s, A_d)$
MAX	$\max(RGB_s, RGB_d)$	$\min(A_s, A_d)$

Table 10.4: Blend Equations

The blend operations are specified with members of the `VkBlendOp` enumeration, and their effects are listed in Table 10.4. In the table, S_{rgb} and D_{rgb} are the the source and destination color factors, S_a and D_a are the source and destination alpha factors, RGB_s and RGB_d are the source and destination RGB values, and A_s and A_d are the source and destination alpha values.

Note that the `VK_BLEND_FACTOR_MIN` and `VK_BLEND_FACTOR_MAX` modes do not include the source or destination factors (S_{rgb}, S_a, D_{rgb} or D_a), only the source and destination colors or alpha values. These modes can be used to find the maximum or minumum values produced at a particular pixel for the attachment.

The available blend factors for use in `srcColorBlendFactor`, `dstColorBlendFactor`, `srcAlphaBlendFactor`, and `dstAlphaBlendFactor` are represented by members of the `VkBlendFactor` enumeration. Table 10.5 shows its members and their meanings.

In the table, R_{s0}, G_{s0}, B_{s0}, and A_{s0} are the R, G, B, and A channels of the first color output from the fragment shader, and R_{s1}, G_{s1}, B_{s1}, and A_{s1} are the R, G, B, and A channels of the secondary color output from the fragment shader. These are used for *dual-source* blending, which is explained in more detail in the next section.

Blend Factor (VK_BLEND_FACTOR_ ...)	RGB	Alpha
ZERO	$(0, 0, 0)$	0
ONE	$(1, 1, 1)$	1
SRC_COLOR	(R_{s0}, G_{s0}, B_{s0})	A_{s0}
ONE_MINUS_SRC_COLOR	$(1, 1, 1) - (R_{s0}, G_{s0}, B_{s0})$	$1 - A_{s0}$
DST_COLOR	(R_d, G_d, B_d)	A_d
ONE_MINUS_DST_COLOR	$(1, 1, 1) - (R_d, G_d, B_d)$	$1 - A_d$
SRC_ALPHA	(A_{s0}, A_{s0}, A_{s0})	A_{s0}
ONE_MINUS_SRC_ALPHA	$(1, 1, 1) - (A_{s0}, A_{s0}, A_{s0})$	$1 - A_{s0}$
DST_ALPHA	(A_d, A_d, A_d)	A_d
ONE_MINUS_DST_ALPHA	$(1, 1, 1) - (A_d, A_d, A_d)$	$1 - A_d$
CONSTANT_COLOR	(R_c, G_c, B_c)	A_c
ONE_MINUS_CONSTANT_COLOR	$(1, 1, 1) - (R_c, G_c, B_c)$	$1 - A_c$
CONSTANT_ALPHA	(A_c, A_c, A_c)	A_c
ONE_MINUS_CONSTANT_ALPHA	$(1, 1, 1) - (A_c, A_c, A_c)$	$1 - A_c$
ALPHA_SATURATE	(f, f, f) $f = \min(A_{s0}, 1 - A_d)$	1
SRC1_COLOR	(R_{s1}, G_{s1}, B_{s1})	A_{s1}
ONE_MINUS_SRC1_COLOR	$(1, 1, 1) - (R_{s1}, G_{s1}, B_{s1})$	$1 - A_{s1}$
SRC1_ALPHA	(A_{s1}, A_{s1}, A_{s1})	A_{s1}
ONE_MINUS_SRC1_ALPHA	$(1, 1, 1) - (A_{s1}, A_{s1}, A_{s1})$	$1 - A_{s1}$

Table 10.5: Blend Factors

In the entries `VK_BLEND_FACTOR_CONSTANT_COLOR` and `VK_BLEND_FACTOR_ONE_MINUS_CONSTANT_COLOR` in Table 10.5, the term `CONSTANT_COLOR` refers to the (R_c, G_c, B_c, A_c) *constant color*, which is a constant that is part of the pipeline. This can be any arbitrary set of floating-point values and can be used, for example, to scale the content of the framebuffer by a fixed amount.

When this state is static, it is specified in the `blendConstants` field of the `VkPipelineColorBlendStateCreateInfo` structure used to create the color blend state object. When the color blend constant state is configured as dynamic, the **vkCmdSetBlendConstants()** function is used to change the blend constants. The prototype of **vkCmdSetBlendConstants()** is

```
void vkCmdSetBlendConstants (
    VkCommandBuffer                         commandBuffer,
    const float                             blendConstants[4]);
```

The command buffer to set the blend constant state for is specified in `commandBuffer`, and the new blend constants are specified in `blendConstants`. This is an array of four floating-point values that take the place of the `blendConstants` member of the `VkPipelineColorBlendStateCreateInfo` structure used to initialize the graphics pipeline's blend state.

As discussed in [Chapter 7](#), "[Graphics Pipelines](#)," to set the constant blend color as part of the dynamic state for a pipeline, include `VK_DYNAMIC_STATE_BLEND_CONSTANTS` in the list of dynamic states passed through the `VkPipelineDynamicStateCreateInfo` structure when creating the graphics pipeline.

In [Table 10.5](#), you notice that the last few factor tokens include the terms `SRC1_COLOR` or `SRC1_ALPHA`, which correspond to RGB_{S1} and A_{S1}, respectively. These are the second source-blending values taken from the fragment shader. The fragment shader can output *two* terms for a single color attachment, and these terms can be used in the blending stage to implement slightly more advanced blending modes than are achievable with a single color output.

For example, it is possible to multiply the destination color (the values already in the color attachment) by one set of values produced by the fragment shader and then add a second set of values produced by the same shader invocation. To implement this, set `srcColorBlendFactor` to `VK_BLEND_FACTOR_ONE` and `dstColorBlendFactor` to `VK_BLEND_FACTOR_SRC1_COLOR`.

When dual-source blending is in use, the two source colors produced by the fragment shader are both directed to the same attachment location, but with different color indices. To specify the color index in SPIR-V, decorate the output from the fragment shader with the `Index` decoration using the `OpDecorate` or `OpMemberDecorate` instruction. Index 0 (which is the default if the decoration is missing) corresponds to the first color output, and index 1 corresponds to the second color output.

Dual-source blending may not be supported by all Vulkan implementations. To determine whether your Vulkan drivers and hardware supports dual-source blending, check the `dualSrcBlend` member of the device's `VkPhysicalDeviceFeatures` structure, which you can retrieve by calling **vkGetPhysicalDeviceFeatures()**. When dual-source blending is supported and enabled, the total number of color attachments that can be referenced by a subpass in the pipeline using that mode may be limited. This limit is stored in `maxFragmentDualSrcAttachments`. If there is any support at all, then at least one attachment can be used in dual-source blending mode.

Summary

This chapter covered the operations that occur as part of fragment processing, after rasterization. These operations include depth and stencil tests, scissor testing, fragment shading, blending, and logic operations. Together, these operations compute the final colors of the pixels produced by your application. They determine visibility and ultimately produce the image that will be shown to your users.

Chapter 11. Synchronization

What You'll Learn in This Chapter

- • How to synchronize the host and the device
- • How to synchronize work on different queues on the same device
- • How to synchronize work conducted at different points in the pipeline

---

Vulkan is designed to run work asynchronously, in parallel, with multiple queues on a device working together with the host to keep physical resources busy and in use. At various points in your application, you will need to keep the host and the various parts of the device in sync. In this chapter, we discuss the multiple synchronization primitives that are available to Vulkan applications for this purpose.

Synchronization in Vulkan is accomplished through the use of various synchronization primitives. There are several types of synchronization primitives and they are intended for different uses in an application. The three main types of synchronization primitives are

- • Fences: Used when the host needs to wait for the device to complete execution of large pieces of work represented by a submission, usually with the assistance of the operating system.
- • Events: Represent a fine-grained synchronization primitive that can be signaled either by the host or the device. It can be signaled mid-command buffer when signaled by the device, and it can be waited on by the device at specific points in the pipeline.
- • Semaphores: Synchronization primitives that are used to control ownership of resources across different queues on a single device. They can be used to synchronize work executing on different queues that would otherwise operate asynchronously.

We'll cover each of these three synchronization primitives in the following few sections.

Fences

A fence is a medium-weight synchronization primitive that is generally implemented with the help of the operating system. A fence is given to commands that interact with the operating system, such as **vkQueueSubmit()**, and when the work that these commands provoke is completed, the fence is signaled.

Because the fence often corresponds to a native synchronization primitive provided by the operating system, it is generally possible to put threads to sleep while they wait on fences, which saves power. However, this is intended for operations in which waiting may take some time, such as waiting for the completion of the execution of a number of command buffers or the presentation of a completed frame to the user.

To create a new fence object, call **vkCreateFence()**, the prototype of which is

[Click here to view code image](#)

```
VkResult vkCreateFence (
    VkDevice                        device,
```

```
    const VkFenceCreateInfo*              pCreateInfo,
    const VkAllocationCallbacks*          pAllocator,
    VkFence*                              pFence);
```

The device that will create the fence object is specified in `device`, and the remaining parameters of the fence are passed through a pointer to an instance of the `VkFenceCreateInfo` structure, the definition of which is

```
typedef struct VkFenceCreateInfo {
    VkStructureType      sType;
    const void*          pNext;
    VkFenceCreateFlags   flags;
} VkFenceCreateInfo;
```

The `sType` field of the `VkFenceCreateInfo` structure should be set to `VK_STRUCTURE_TYPE_FENCE_CREATE_INFO`, and `pNext` should be set to `nullptr`. The only remaining field, `flags`, specifies a set of flags that controls the behavior of the fence. The only flag defined for use here is `VK_FENCE_CREATE_SIGNALED_BIT`. If this bit is set in `flags`, then the initial state of the fence is signaled; otherwise, it is unsignaled.

If **vkCreateFence()** is successful, a handle to the new fence object is placed in the variable pointed to by `pFence`. If `pAllocator` is not `nullptr`, it should point to a host memory allocation structure that will be used to allocate any host memory required by the fence.

As with most other Vulkan objects, when you are done with the fence, you should destroy it in order to free its resources. To do this, call **vkDestroyFence()**, the prototype of which is

```
void vkDestroyFence (
    VkDevice                         device,
    VkFence                          fence,
    const VkAllocationCallbacks*     pAllocator);
```

The device that owns the fence object is specified in `device`, and the handle to the fence to be destroyed is passed in `fence`. If a host memory allocator was used with **vkCreateFence()**, then `pAllocator` should point to a host memory allocation structure compatible with the one used to allocate the object; otherwise, `pAllocator` should be `nullptr`.

The fence may be used in any command that takes a fence parameter. These commands usually operate on queues and provoke work to be executed on that queue. For example, here is the prototype of **vkQueueSubmit()**:

```
VkResult vkQueueSubmit (
    VkQueue                          queue,
    uint32_t                         submitCount,
    const VkSubmitInfo*              pSubmits,
    VkFence                          fence);
```

Notice that the last parameter to **vkQueueSubmit()** is a `VkFence` handle. When all of the work provoked on `queue` is complete, the fence specified in `fence` will be set to signaled. In some cases,

the device can signal the fence directly. In other cases, the device will signal the operating system through an interrupt or other hardware mechanism, and the operating system will change the state of the fence.

The application can determine the state of the fence at any time by calling **vkGetFenceStatus()**, the prototype of which is

[**Click here to view code image**](#)

```
VkResult vkGetFenceStatus (
    VkDevice                              device,
    VkFence                               fence);
```

The device that owns the fence is specified in `device`, and the fence whose status to query is passed in `fence`. The value returned from **vkGetFenceStatus()** indicates the state of the fence. On success, **vkGetFenceStatus()** can be one of the following values:

- `VK_SUCCESS`: The fence is currently signaled.
- `VK_NOT_READY`: The fence is currently unsignaled.

If there was some problem retrieving the status of the fence, **vkGetFenceStatus()** may return an error code. It can be tempting to *poll* the state of a fence by spinning in a loop until **vkGetFenceStatus()** returns `VK_SUCCESS`. However, this is extremely inefficient and bad for performance, especially if there is any chance that the application could wait for a long time. Rather than spin, your application should call **vkWaitForFences()**, which allows the Vulkan implementation to provide a more optimal mechanism for waiting on one or more fences.

The prototype of **vkWaitForFences()** is

[**Click here to view code image**](#)

```
VkResult vkWaitForFences (
    VkDevice                              device,
    uint32_t                              fenceCount,
    const VkFence*                        pFences,
    VkBool32                              waitAll,
    uint64_t                              timeout);
```

The device that owns the fences that will be waited on is passed in `device`. **vkWaitForFences()** can wait on any number of fences. The number of fences upon which it should wait is passed in `fenceCount`, and `pFences` should point to an array of this many `VkFence` handles to the fences to wait on.

vkWaitForFences() can either wait for all the fences in the `pFences` array to become signaled or can return as soon as any fence in the array becomes signaled. If `waitAll` is `VK_TRUE`, then **vkWaitForFences()** will wait for all of the fences in `pFences` to become signaled; otherwise, it will return as soon as any of them becomes signaled. In this case, you can use **vkGetFenceStatus()** to determine which fences are signaled and which are unsignaled.

Because **vkWaitForFences()** could end up waiting for a very long time, it is possible to ask it to time out. The `timeout` parameter specifies the time, in nanoseconds, that **vkWaitForFences()** should wait for the termination condition (at least one of the fences becomes signaled) to be met. If

timeout is zero, then **vkWaitForFences()** will simply check the status of the fences and return immediately. This is similar to **vkGetFenceStatus()** except for two differences:

- **vkWaitForFences()** can check the status of multiple fences, possibly more efficiently than multiple calls to **vkGetFenceStatus()**.

- If the fences in the pFences array are not signaled, then **vkWaitForFences()** indicates a timeout rather than a not-ready condition.

There is no way to ask **vkWaitForFences()** to wait forever. However, because timeout is a 64-bit value measured in nanoseconds and $2^{64} - 1$ nanoseconds is a little over 584 years, the value ~0ull is probably a sufficient proxy for infinity.

The result of the wait operation is represented by the value returned from **vkWaitForFences()**. It may be one of the following values:

- VK_SUCCESS: The condition being waited on was met. If waitAll was VK_FALSE, then at least one fence in pFences was signaled. If waitAll was VK_TRUE, then all of the fences in pFences became signaled.

- VK_TIMEOUT: The condition being waited on was not met before the timeout period occurred. If timeout was zero, this represents an instantaneous polling of the fences.

If an error occurs, such as one of the handles in the pFences array being invalid, then **vkWaitForFences()** will return an appropriate error code.

Once a fence has become signaled, it will remain in that state until it is reset to an unsignaled state explicitly. When a fence is no longer signaled, it can be reused in commands such as **vkQueueSubmit()**. There are no commands that cause the device to reset a fence, and there are no commands that can wait for a fence to become unsignaled. To reset one or more fences to an unsignaled state, call **vkResetFences()**, the prototype of which is

[**Click here to view code image**](#)

```
VkResult vkResetFences (
    VkDevice                        device,
    uint32_t                        fenceCount,
    const VkFence*                  pFences);
```

The device that owns the fences is passed in device; the number of fences to reset is passed in fenceCount; and a pointer to an array of VkFence handles is passed in pFences. Take care that you don't try waiting on a fence that has been reset without calling another function that will eventually signal that fence.

A primary use case for fences is to prevent the host from overwriting data that might be in use by the device—or, more accurately, data that might be *about to be used* by the device. When a command buffer is submitted, it does not execute immediately, but is placed in queue and executed by the device in turn. It also does not execute infinitely quickly, but takes some time to execute. Therefore, if you place data in memory and then submit a command buffer that references that data, you must be careful to ensure that the data remains valid and in place until the command buffer has executed. The fence supplied to **vkQueueSubmit()** is the means to determine this.

In this example, we create a buffer, map it, and place some data in it using the host; then we submit a command buffer that references that data. Three mechanisms are provided to perform

synchronization between the host and device and to ensure that the host does not overwrite the data in the buffer before it's used by the device.

In the first method, we call **`vkQueueWaitIdle()`** to ensure that all work submitted to the queue (including our work that consumes the data in the buffer) has completed. In the second method, we use a single fence, associated with the submission that consumes the data, and we wait on that fence before overwriting the content of the buffer. In the third method, we subdivide the buffer into four quarters, associate a fence with each quarter, and wait for the fence associated with that section of the buffer before overwriting it.

In addition, a method with no synchronization at all is provided. This method simply overwrites the data in the buffer without waiting at all. In this method, we invalidate the data in the buffer before overwriting it with new, valid data. This should demonstrate the kind of corruption that can occur in the absence of proper synchronization.

The code for the first two methods—brute-force idle of the queue and waiting for a single fence (which is essentially equivalent in this scenario)—is somewhat trivial and is therefore not shown here. Listing 11.1 shows the initialization sequence for the method using four separate fences.

Listing 11.1: Setup for Four-Fence Synchronization

Click here to view code image

```
// kFrameDataSize is the size of the data consumed in a single frame.
// kRingBufferSegments is the number of frames' worth of data to keep.
// Create a buffer large enough to hold kRingBufferSegments copies
// of kFrameDataSize.
static const VkBufferCreateInfo bufferCreateInfo =
{
    VK_STRUCTURE_TYPE_BUFFER_CREATE_INFO, nullptr,  // sType,pNext
    0,                                              // flags
    kFrameDataSize*kRingBufferSegments,             // size
    VK_BUFFER_USAGE_TRANSFER_SRC_BIT,               // usage
    VK_SHARING_MODE_EXCLUSIVE,                      // sharingMode
    0,                                              //
queueFamilyIndexCount
    nullptr                                         // pQueueFamilyIndices
};

result = vkCreateBuffer(device,
                        &bufferCreateInfo,
                        nullptr,
                        &m_buffer);

// Create kRingBufferSegments fences, all initially in signaled state.
static const VkFenceCreateInfo fenceCreateInfo =
{
    VK_STRUCTURE_TYPE_FENCE_CREATE_INFO,nullptr,

    VK_FENCE_CREATE_SIGNALED_BIT
};

for (int i = 0;i < kRingBufferSegments; ++i)
```

```
{
    result = vkCreateFence(device,
                           &fenceCreateInfo,
                           nullptr,
                           &m_fence[i]);
}
```

As you can see, in [Listing 11.1](#) we create a buffer that is four times the size of the data to be used by the device; then we create four fences, each guarding one quarter of the total buffer size. The fences are created in the signaled state. This is so we can simply enter our loop that waits on the fence guarding the current section of the buffer to be signaled, fills the buffer, and generates a new command buffer referencing that section of the buffer. The first time this loop executes, the fence is already signaled because it was created in that state, and no special case is required for the first entry into the loop.

[Listing 11.2](#) shows the inner loop that waits for each fence to be signaled, fills the appropriate section of the buffer, resets the fence, generates a command buffer to consume the data, submits it to the queue, and specifies the fence.

Listing 11.2: Loop Waiting on Fences for Synchronization

[Click here to view code image](#)

```
// Beginning of frame - compute the segment index.
static int framesRendered = 0;
const int segmentIndex = framesRendered % kRingBufferSegments;

// Use a ring of command buffers indexed by segment.
const VkCommandBuffer cmdBuffer = m_cmdBuffer[segmentIndex];

// Wait for the fence associated with this segment.
result = vkWaitForFences(device,
                         1,
                         &m_fence[segmentIndex],
                         VK_TRUE,
                         UINT64_MAX);

// It's now safe to overwrite the data. m_mappedData is an array of
// kRingBufferSegments pointers to persistently mapped backing store for
// the source buffer.
fillBufferWithData(m_mappedData[segmentIndex]);

// Reset the command buffer. We always use the same command buffer to copy
// from a given segment of the the staging buffer, so it's safe to reset
it
// here because we've already waited for the associated fence.
vkResetCommandBuffer(cmdBuffer, 0);

// Rerecord the command buffer.
static const VkCommandBufferBeginInfo beginInfo =

{
    VK_STRUCTURE_TYPE_COMMAND_BUFFER_BEGIN_INFO,      // sType
```

```cpp
    nullptr,                                          // pNext
    VK_COMMAND_BUFFER_USAGE_ONE_TIME_SUBMIT_BIT,      // flags
    nullptr                                           // pInheritanceInfo
};

vkBeginCommandBuffer(cmdBuffer, &beginInfo);

// Copy from the staging buffer at the appropriate segment index into
// the final destination buffer.
VkBufferCopy copyRegion =
{
    segmentIndex * kFrameDataSize,                    // srcOffset
    0,                                                // dstOffset
    kFrameDataSize                                    // size
};

vkCmdCopyBuffer(cmdBuffer,
                m_stagingBuffer,
                m_targetBuffer,
                1,
                &copyRegion);

vkEndCommandBuffer(cmdBuffer);

// Reset the fence for this segment before submitting this chunk of work
to
// the queue.
vkResetFences(device, 1, &m_fence[segmentIndex]);

// Note that this example doesn't use any submission semaphores. In a real
// application, you would submit many command buffers in a single
submission
// and protect that submission using wait and signal semaphores.
VkSubmitInfo submitInfo =
{
    VK_STRUCTURE_TYPE_SUBMIT_INFO, nullptr,           // sType, pNext
    0,                                                // waitSemaphoreCount
    nullptr,                                          // pWaitSemaphores
    nullptr,                                          // pWaitDstStageMask
    1,                                                // commandBufferCount
    &cmdBuffer,                                       // pCommandBuffers
    0,                                                //
signalSemaphoreCount
    nullptr                                           // pSignalSemaphores
};

vkQueueSubmit(m_queue,
              1,
              &submitInfo,
              m_fence[segmentIndex]);

framesRendered++;
```

Note that the code in [Listing 11.2](#) is incomplete and serves only to demonstrate the use of fences to protect shared data. In particular, the example omits several things that a real application will need:

- It does not include any pipeline barriers to protect the target buffer (`m_targetBuffer`) from being overwritten or to move the source buffer from host writable to the source for transfers.

- It does not use any semaphores to protect the submission. Semaphores will be needed if this is part of a larger frame that includes presentation or submission to other queues.

- It does not include any calls to **`vkFlushMappedMemoryRanges()`**, which will be required unless the memory backing the buffer was allocated with the `VK_MEMORY_PROPERTY_HOST_COHERENT_BIT` set.

Events

Event objects represent a fine-grained synchronization primitive that can be used to accurately delimit operations occurring within a pipeline. An event exists in one of two states: signaled and unsignaled. Unlike a fence, which is typically moved from state to state by the operating system, events can be signaled or reset explicitly by the device or host. Not only can the device directly manipulate the state of an event, but it can do this at specific points in the pipeline.

Unlike a fence, where only the host can wait on the object, the device can wait on event objects. When waiting, the device can wait at a specific point in the pipeline. Because the device can signal an event at one part of the pipeline and wait for it at another, events provide a way to synchronize execution occurring at different parts of the same pipeline.

To create an event object, call **`vkCreateEvent()`**, the prototype of which is

[**Click here to view code image**](#)

```
VkResult vkCreateEvent (
    VkDevice                        device,
    const VkEventCreateInfo*        pCreateInfo,
    const VkAllocationCallbacks*    pAllocator,
    VkEvent*                        pEvent);
```

The device which is to create the event is specified in `device`, and the remaining parameters describing the event are passed through a pointer to an instance of the `VkEventCreateInfo` structure, the definition of which is

[**Click here to view code image**](#)

```
typedef struct VkEventCreateInfo {
    VkStructureType     sType;
    const void*         pNext;
    VkEventCreateFlags  flags;
} VkEventCreateInfo;
```

The `sType` field of `VkEventCreateInfo` should be set to `VK_STRUCTURE_TYPE_EVENT_CREATE_INFO`, and `pNext` should be set to `nullptr`. The `flags` field of `VkEventCreateInfo` specifies additional behavior of the event. However, there are currently no event creation flags defined, and `flags` should be treated as reserved and set to zero.

If **vkCreateEvent()** is successful, it will place the VkEvent handle to the newly created object in the variable pointed to by pEvent. If **vkCreateEvent()** requires host memory, then it will use the allocator specified in pAllocator. A compatible allocator should be used when the event is destroyed.

When you are done with the event, you should free its resources by destroying the object. To destroy an event object, call **vkDestroyEvent()**, the prototype of which is

[**Click here to view code image**](#)

```
void vkDestroyEvent (
    VkDevice                          device,
    VkEvent                           event,
    const VkAllocationCallbacks*      pAllocator);
```

The device that owns the event object should be passed in device, and the event object to be destroyed should be passed in event. You should make sure that the event is not accessed again after it has been destroyed. This includes direct access to the event on the host by passing its handle to a function call, and also implied access to the event through the execution of a command buffer that may contain a reference to it.

The initial state of an event is unsignaled or reset. Either the host or the device can change the state of the event. To change the state of an event to set on the host, call **vkSetEvent()**, the prototype of which is

[**Click here to view code image**](#)

```
VkResult vkSetEvent (
    VkDevice                         device,
    VkEvent                          event);
```

The device that owns the event object should be passed in device, and the event object to set is specified in event. Access to the event object should be externally synchronized. Attempting to set or reset the event from multiple threads concurrently will result in a race condition and produce an undefined result. When an event object is set by the host, its state immediately changes to set. If another thread is waiting on the event through a call to **vkCmdWaitEvents()**, then that thread will immediately be unblocked, assuming that other waiting conditions are satisfied.

Events can also be moved from set to reset state on the host by calling **vkResetEvent()**. The prototype of vkResetEvent is

[**Click here to view code image**](#)

```
VkResult vkResetEvent (
    VkDevice                        device,
    VkEvent                         event);
```

Again, the device that owns the event is passed in device, and the event object to reset is passed in event. Again, access to the event must be externally synchronized. The specified event is immediately moved to the reset state. The **vkSetEvent()** and **vkResetEvent()** commands have no effect if the specified event is already in the set or reset state, respectively. That is, it is not an error to call **vkSetEvent()** on an event that is already set or **vkResetEvent()** on an event that is already reset, although it might not be what you intended to do.

You can determine the immediate state of an event with **vkGetEventStatus()**, the prototype of which is

```
VkResult vkGetEventStatus (
    VkDevice                            device,
    VkEvent                             event);
```

The device that owns the event is passed in device, and the event whose state to query is passed in event. Unlike with **vkSetEvent()** and **vkResetEvent()**, **vkGetEventStatus()** can be called on an event object without synchronization. In fact, this is the intended use case.

The return value of **vkGetEventStatus()** reports the state of the event. The possible return values include

- VK_EVENT_SET: The specified event is in the set or signaled state.
- VK_EVENT_RESET: The specified event is in the reset or unsignaled state.

If anything goes wrong, such as the event parameter not being a handle to a valid event, **vkGetEventStatus()** will return an appropriate error code.

There is no way for the host to wait on an event besides spinning in a loop waiting for **vkGetEventStatus()** to return VK_EVENT_SET. This is not particularly efficient, and you should be sure to cooperate with the system in the case that you need to do this by, for example, sleeping the current thread or doing other useful work in between queries of the event object status.

Event objects may also be manipulated by the device. As with almost any other work executed by the device, this is done by placing commands in a command buffer and then submitting the command buffer to one of the device's queues for execution. The command for setting an event is **vkCmdSetEvent()**, the prototype of which is

```
void vkCmdSetEvent (
    VkCommandBuffer                 commandBuffer,
    VkEvent                         event,
    VkPipelineStageFlags            stageMask);
```

The command buffer that will eventually set the event is specified in commandBuffer, and a handle to the event object that will be signaled is passed in event. Unlike **vkSetEvent()**, **vkCmdSetEvent()** takes a pipeline stage at which to signal the event. This is passed in stageMask, which is a bitfield made up of members of the VkPipelineStageFlagBits enumeration. If multiple bits are set in stageMask, then the event will be set when execution of the command passes each of the specified stages. This may seem redundant, but if there are calls to **vkCmdResetEvent()**, it may be possible to observe the state of the event object toggle between set and reset as the commands pass down the pipeline.

The corresponding command to reset an event is **vkCmdResetEvent()**, the prototype of which is

```
void vkCmdResetEvent (
    VkCommandBuffer                         commandBuffer,
```

```
    VkEvent                                 event,
    VkPipelineStageFlags                    stageMask);
```

Again, the command buffer that is to reset the event is passed in `commandBuffer`; the event to reset is passed in `event`; and as with **vkCmdSetEvent()**, the **vkCmdResetEvent()** command takes a `stageMask` parameter that contains a bitfield representing the stages at which the event will assume the reset state.

While it is possible on the host to retrieve the immediate state of an event object using **vkGetEventStatus()** but it is not possible to directly wait on the event, the converse is true on the device; there is no command to directly retrieve the state of an event on the device, but it *is* possible to wait on one or more events. To do this, call **vkCmdWaitEvents()**, the prototype of which is

[**Click here to view code image**](#)

```
void vkCmdWaitEvents (
    VkCommandBuffer                         commandBuffer,
    uint32_t                                eventCount,
    const VkEvent*                          pEvents,
    VkPipelineStageFlags                    srcStageMask,
    VkPipelineStageFlags                    dstStageMask,
    uint32_t                                memoryBarrierCount,
    const VkMemoryBarrier*                  pMemoryBarriers,
    uint32_t                                bufferMemoryBarrierCount,
    const VkBufferMemoryBarrier*            pBufferMemoryBarriers,
    uint32_t                                imageMemoryBarrierCount,
    const VkImageMemoryBarrier*             pImageMemoryBarriers);
```

The **vkCmdWaitEvents()** command takes many parameters and behaves much as a barrier does. The command buffer whose execution will be stalled waiting for the events is specified in `commandBuffer`; the number of events to wait on is specified in `eventCount`; and a pointer to an array of `VkEvent` handles to the events to be waited on is passed in `pEvents`.

The `srcStageMask` and `dstStageMask` specify the pipeline stages in which the events may have been signaled and the stages that will wait for the events to become signaled, respectively. Waiting will always occur at the stages specified in `dstStageMask`. The stages specified in `srcStageMask` allow a Vulkan implementation to synchronize pairs of stages rather than all work in the pipeline, which can be somewhat more efficient.

Once all of the events contained in the `pEvents` array have become signaled, the device will execute all of the memory, buffer, and image barriers specified in `pMemoryBarriers`, `pBufferMemoryBarriers`, and `pImageMemoryBarriers` before resuming execution of the work in the pipeline. The number of elements in the array pointed to by `pMemoryBarriers` is `memoryBarrierCount`; the number of elements in the array pointed to by `pBufferMemoryBarriers` is `bufferMemoryBarrierCount`; and the number of elements in the array pointed to by `pImageMemoryBarriers` is `imageMemoryBarrierCount`.

Although the flags contained in `VkPipelineStageFlagBits` are quite fine-grained, most Vulkan implementations will be able to perform a wait operation at only some of them. It is always legal for an implementation to wait at an earlier point in the pipeline. It is still guaranteed that work

executed at later points in the pipeline will happen *after* the events in `pEvents` have become signaled.

Semaphores

The final type of synchronization primitive supported in Vulkan is the semaphore. Semaphores represent flags that can be atomically set or reset by the hardware, the views of which are coherent across queues. When you are setting the semaphore, the device will wait for it to be unset, set it, and then return control to the caller. Likewise, when resetting the semaphore, the device waits for the semaphore to be set, resets it, and then returns to the caller. This all happens atomically. If multiple agents are waiting on a semaphore, only one of them will "see" the semaphore as unset or set and receive control. The rest will continue to wait.

Semaphores cannot be explicitly signaled or waited on by the device. Rather, they are signaled and waited on by queue operations such as **`vkQueueSubmit()`**.

To create a semaphore object, call **`vkCreateSemaphore()`**, the prototype of which is

[**Click here to view code image**](#)

```
VkResult vkCreateSemaphore (
    VkDevice                        device,
    const VkSemaphoreCreateInfo*    pCreateInfo,
    const VkAllocationCallbacks*    pAllocator,
    VkSemaphore*                    pSemaphore);
```

The device that will create the semaphore is passed in `device`, and additional parameters of the semaphore are passed through `pCreateInfo`, which is a pointer to an instance of the `VkSemaphoreCreateInfo` structure. The definition of `VkSemaphoreCreateInfo` is

[**Click here to view code image**](#)

```
typedef struct VkSemaphoreCreateInfo {
    VkStructureType         sType;
    const void*             pNext;
    VkSemaphoreCreateFlags  flags;
} VkSemaphoreCreateInfo;
```

The `sType` field of `VkSemaphoreCreateInfo` should be set to `VK_STRUCTURE_TYPE_SEMAPHORE_CREATE_INFO`, and `pNext` should be set to `nullptr`. The `flags` field is reserved and should be set to zero.

If **`vkCreateSemaphore()`** is successful, the resulting semaphore object is written into the variable pointed to by `pSemaphore`.

When you are done with the semaphore object, you should destroy it in order to free any resources associated with it. To destroy a semaphore, call **`vkDestroySemaphore()`**, the prototype of which is

[**Click here to view code image**](#)

```
void vkDestroySemaphore (
    VkDevice                        device,
    VkSemaphore                     semaphore,
    const VkAllocationCallbacks*    pAllocator);
```

The device that owns the semaphore object should be passed in `device`, and the semaphore that is to be destroyed should be passed in `semaphore`. Access to the semaphore must be externally synchronized. In particular, a semaphore must not be destroyed while it may be accessed from another thread. If a host memory allocator was used to create the semaphore, a pointer to a compatible allocator should be passed in `pAllocator`.

Unlike the other synchronization primitives, events and fences, semaphore objects do not allow you to explicitly set, reset, or wait on them. Instead, you use these objects to synchronize access to resources across queues and form an integral part of submission of work to the device. Recall that the prototype of **vkQueueSubmit()** is

[**Click here to view code image**](#)

```
VkResult vkQueueSubmit (
    VkQueue                          queue,
    uint32_t                         submitCount,
    const VkSubmitInfo*              pSubmits,
    VkFence                          fence);
```

vkQueueSubmit() takes an array of `VkSubmitInfo` structures, the definition of which is

[**Click here to view code image**](#)

```
typedef struct VkSubmitInfo {
    VkStructureType                  sType;
    const void*                      pNext;
    uint32_t                         waitSemaphoreCount;
    const VkSemaphore*               pWaitSemaphores;
    const VkPipelineStageFlags*      pWaitDstStageMask;
    uint32_t                         commandBufferCount;
    const VkCommandBuffer*           pCommandBuffers;
    uint32_t                         signalSemaphoreCount;
    const VkSemaphore*               pSignalSemaphores;
} VkSubmitInfo;
```

Each entry of the `pSubmits` array passed to **vkQueueSubmit()** contains a `pWaitSemaphores` and a `pSignalSemaphores` member, which are pointers to arrays of semaphore objects. Before executing the commands in `pCommandBuffers`, the queue will wait for all of the semaphores in `pWaitSemaphores`. In doing so, it will take "ownership" of the semaphores.

It will then execute the commands contained in each of the command buffers in the `pCommandBuffers` array, and when it is done, it will signal each of the semaphores contained in `pSignalSemaphores`.

Access to the semaphores referenced by the `pWaitSemaphores` and `pSignalSemaphores` arrays must be externally synchronized. In practice, this means that if you submit command buffers to two different queues from two different threads (which is perfectly legal), you need to be careful that the same semaphore doesn't end up in the wait list for one submission and the signal list for the other.

This mechanism can be used to synchronize access to resources that is performed by work submitted to two or more different queues. For example, we can submit a number of command buffers to a compute-only queue, which then signals a semaphore upon completion. That same semaphore

appears in the wait list for a second submission to a graphics queue. <u>Listing 11.3</u> shows an example of this.

Listing 11.3: Cross-Queue Submission with Semaphores

<u>Click here to view code image</u>

```cpp
// First, perform a submission to the compute queue. The compute
submission
// includes the compute -> graphics semaphore (m_computeToGfxSemaphore) in
// the pSignalSemaphores list. This sempaphore will become signaled once
the
// compute queue finishes processing the command buffers in the
submission.
VkSubmitInfo computeSubmitInfo =
{
    VK_STRUCTURE_TYPE_SUBMIT_INFO, nullptr,        // sType, pNext
    0,                                             // waitSemaphoreCount
    nullptr,                                       // pWaitSemaphores
    nullptr,                                       // pWaitDstStageMask
    1,                                             // commandBufferCount
    &computeCmdBuffer,                             // pCommandBuffers
    1,                                             // signalSemaphoreCount
    &m_computeToGfxSemaphore                       // pSignalSemaphores
};

vkQueueSubmit(m_computeQueue,
              1,
              &computeSubmitInfo,
              VK_NULL_HANDLE);

// Now perform a submission to the graphics queue. This submission
includes
// m_computeToGfxSemaphore in the pWaitSemaphores list passed to the
// submission. We need to wait at a specific stage. Here, we just use
// VK_PIPELINE_STAGE_TOP_OF_PIPE_BIT, but if you know that the data
produced
// on the source queue will be consumed at a later stage of the pipeline
in
// the destination queue, you could place the wait point later.
static const VkFlags waitStages = VK_PIPELINE_STAGE_TOP_OF_PIPE_BIT;
VkSubmitInfo graphicsSubmitInfo =
{
    VK_STRUCTURE_TYPE_SUBMIT_INFO, nullptr,        // sType, pNext
    1,                                             // waitSemaphoreCount
    &m_computeToGfxSemaphore,                      // pWaitSemaphores
    &waitStages,                                   // pWaitDstStageMask
    1,                                             // commandBufferCount
    &graphicsCmdBuffer,                            // pCommandBuffers
    0,                                             //
signalSemaphoreCount
    nullptr                                        // pSignalSemaphores
};
```

```
vkQueueSubmit(m_graphicsQueue,
              1,
              &graphicsSubmitInfo,
              VK_NULL_HANDLE);
```

In [Listing 11.3](#), you can see that a single semaphore object is used in both submissions. The first time, it appears in the signal list for work performed on the compute queue. When the work is complete, the compute queue signals the semaphore. The same semaphore then appears in the wait list for the submission on the graphics queue. The graphics queue will wait for this semaphore to become signaled before proceeding to execute the command buffers submitted to it. This ensures that work executed on the compute queue to produce data is completed before work executing on the graphics queue that consumes that data begins.

The same synchonization mechanism using semaphores is employed by the sparse memory binding commands that also operate at the queue level. The **vkQueueBindSparse()** function was introduced in the "[Sparse Resources](#)" section of [Chapter 2](#), "[Memory and Resources](#)." To recap, the prototype of **vkQueueBindSparse()** is

[Click here to view code image](#)

```
VkResult vkQueueBindSparse (
    VkQueue                         queue,
    uint32_t                        bindInfoCount,
    const VkBindSparseInfo*         pBindInfo,
    VkFence                         fence);
```

Each binding operation is represented as an instance of the VkBindSparseInfo structure, and each of those operations also has a pWaitSemaphores and pSignalSemaphores array that behaves similarly to the parameters passed to **vkQueueSubmit()**. Again, access to those semaphores must be externally synchronized, meaning that you must ensure that no two threads attempt to access the same semaphores at the same time.

Summary

This chapter taught you about the synchronization primitives available in Vulkan: fences, events, and semaphores. Fences provide a mechanism for the operating system to signal your application when it has completed operations requested of it, such as submission of command buffers or presentation of images via the window system. Events provide a fine-grained signaling mechanism that can be used to control flow of data through the pipeline and allow different points within the pipeline to synchronize. Finally, semaphores provide a primitive that can be signaled and waited on different queues on the same device, allowing synchronization and transfer of ownership of resources across queues.

Together, these primitives provide a powerful toolbox. As Vulkan is an asynchronous API with work occurring in parallel across the host and device, and across multiple queues on a single device, synchronization primitives and their correct use are key to the operation of an efficient application.

Chapter 12. Getting Data Back

What You'll Learn in This Chapter

- Gather information about the execution of your application on the device
- Time operations performed by the device
- Read data produced by the device on the host

For the most part, graphics and compute operations supported by Vulkan are "fire and forget," in that you build a command buffer and submit it, and eventually, data is displayed to the user. Your application has very little feedback or input from Vulkan. However, there are reasons to want to retrieve data *from* Vulkan. This chapter covers topics related to reading data and information back from Vulkan. This data includes statistics about the operations your application performs, timing information, and reading data produced directly by your application.

Queries

The primary mechanism for reading statistical data back from Vulkan is the *query object*. Query objects are created and managed in *pools*, and each object is effectively one slot in a pool rather than a discrete object that is managed alone. There are several types of query objects, each measuring a different aspect of the device's operation as it executes work you submit to it. All types of queries are managed in pools, so the first step in using queries is creating a pool object to store them in.

Most queries execute by wrapping commands contained inside a command buffer with a pair of commands to start and stop the query. The exception to this is the timestamp query, which takes an instantaneous snapshot of the device time and therefore doesn't really have a duration. For other query types, while they execute, statistics about the operation of the device are gathered, and when the query is stopped, the results are written into device memory represented by the pool. At some later time, your application can gather the results of any number of queries contained in the pool and read them back.

A query pool is created by calling**vkCreateQueryPool()**, the prototype of which is

[Click here to view code image](#)

```
VkResult vkCreateQueryPool (
    VkDevice                            device,
    const VkQueryPoolCreateInfo*        pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkQueryPool*                        pQueryPool);
```

The device that will create the pool is specified in `device`. The remaining parameters of the pool are passed through a pointer to an instance of the `VkQueryPoolCreateInfo` structure, the definition of which is

[Click here to view code image](#)

```
typedef struct VkQueryPoolCreateInfo {
    VkStructureType                     sType;
```

```
    const void*                    pNext;
    VkQueryPoolCreateFlags         flags;
    VkQueryType                    queryType;
    uint32_t                       queryCount;
    VkQueryPipelineStatisticFlags  pipelineStatistics;
} VkQueryPoolCreateInfo;
```

The `sType` field of `VkQueryPoolCreateInfo` should be set to
`VK_STRUCTURE_TYPE_QUERY_POOL_CREATE_INFO`, and `pNext` should be set to `nullptr`.
The `flags` field is reserved for future use and should be set to zero. The type of query to be stored
in the pool is specified in the `queryType` field, which is a member of the `VkQueryType`
enumeration.

Each pool can contain only one type of query, although it is possible to create as many pools as you
wish and to run several queries at the same time.

The types of query are

- `VK_QUERY_TYPE_OCCLUSION`: Occlusion queries count the number of samples that pass
  the depth and stencil tests.

- `VK_QUERY_TYPE_PIPELINE_STATISTICS`: Pipeline statistics queries count various
  statistics generated by operations of the device.

- `VK_QUERY_TYPE_TIMESTAMP`: Timestamp queries measure the amount of time the
  execution of commands in a command buffer takes.

Each type of query is discussed in more detail later in this chapter.

The number of queries that can be stored in the pool is specified in `queryCount`. When the pool is
used to execute queries, the individual queries are referenced by their index within the pool.

Finally, when the query type is `VK_QUERY_TYPE_PIPELINE_STATISTICS`, some additional
flags controlling how those statistics are gathered is specified in `pipelineStatistics`.

If **vkCreateQueryPool()** is successful, a handle to the new query pool object will be written
into the variable pointed to by `pQueryPool`. If `pAllocator` is a pointer to a valid host memory
allocator, then that allocator will be used to allocate any host memory needed by the pool object.
Otherwise, `pAllocator` should be `nullptr`.

As with any other object in Vulkan, when you are done with it, you should destroy the query pool to
free its resources. To do this, call **vkDestroyQueryPool()**, the prototype of which is

[**Click here to view code image**](#)

```
void vkDestroyQueryPool (
    VkDevice                       device,
    VkQueryPool                    queryPool,
    const VkAllocationCallbacks*   pAllocator);
```

The device that owns the pool should be passed in `device`, and the pool to be destroyed should be
passed in `queryPool`. If a host memory allocator was used when the pool was created, a pointer to
a compatible allocator should be passed in `pAllocator`; otherwise, `pAllocator` should be
`nullptr`.

Each query in the query pool is marked as available or unavailable. Initially, all queries in the pool
are in an undefined state, and before any query can be used, you need to reset the pool. As the only

agent that can write into the pool is the device, you need to execute a command on the device to reset the pool. The command to do this is **vkCmdResetQueryPool()**, the prototype of which is

```
void vkCmdResetQueryPool (
    VkCommandBuffer                         commandBuffer,
    VkQueryPool                             queryPool,
    uint32_t                                firstQuery,
    uint32_t                                queryCount);
```

The command buffer that will execute the reset command is specified in `commandBuffer`, and the pool containing the queries is specified in `queryPool`. It is possible to reset only a selection of queries within a pool. The `firstQuery` parameter specifies the index of the first query to reset, and `queryCount` is the number of queries to reset. You must submit the command buffer containing**vkCmdResetQueryPool()** to an appropriate queue before the pool can be used for anything else.

Executing Queries

Queries are executed by wrapping commands contained inside a command buffer in a pair of additional commands to start and stop the query:**vkCmdBeginQuery()** and **vkCmdEndQuery()**, respectively.

The prototype for **vkCmdBeginQuery()** is

```
void vkCmdBeginQuery (
    VkCommandBuffer                 commandBuffer,
    VkQueryPool                     queryPool,
    uint32_t                        query,
    VkQueryControlFlags             flags);
```

The command buffer that contains the commands to gather statistics about and that will execute the query is specified in `commandBuffer`. Queries are referred to by their index within the pool that contains them. The pool is specified in `queryPool`, and the index of the query within the pool is specified in `query`.

Additional flags that control the execution of the query can be specified in `flags`. The only defined flag is `VK_QUERY_CONTROL_PRECISE_BIT`. If this flag is set, then the results gathered by the query will be precise (the meaning of *precise* varies by query type); otherwise, Vulkan might produce approximate results. In some cases, gathering precise results may reduce performance, so you should set this flag only when you require exact results.

Be aware, however, that if there is no performance penalty to running an exact query, an implementation might ignore this flag and always return exact results. You should be sure to test your application on multiple Vulkan implementations to make sure that it really *is* tolerant of inexact results.

Once the query has begun, place the commands to gather statistics about in the command buffer, and after you have executed those commands, end the query by calling**vkCmdEndQuery()**. The prototype of**vkCmdEndQuery()** is

```
void vkCmdEndQuery (
    VkCommandBuffer                    commandBuffer,
    VkQueryPool                       queryPool,
    uint32_t                          query);
```

The command buffer containing the currently executing query is specified in commandBuffer. The pool containing the query is passed in queryPool, and the index of the query within the pool is specified in query.

Before you begin a query, the query must be reset. The queries in a pool are reset by the **vkCmdResetQueryPool()** command, as discussed earlier. This command must be executed when the queue is created and also between each use of a query object.

Calls to **vkCmdBeginQuery()** and **vkCmdEndQuery()** must appear in pairs. If you begin a query and forget to end it, the result will never become available to your application. If you end a query more than once or without beginning it first, the result of the query will be undefined.

To retrieve the results of one or more queries from a pool, call **vkGetQueryPoolResults()**. Its prototype is

```
VkResult vkGetQueryPoolResults (
    VkDevice                          device,
    VkQueryPool                       queryPool,
    uint32_t                          firstQuery,
    uint32_t                          queryCount,
    size_t                            dataSize,
    void*                             pData,
    VkDeviceSize                      stride,
    VkQueryResultFlags                flags);
```

The device that owns the pool from which to retrieve results is passed in device, and the pool is passed in queryPool. The index of the first query to retrieve results for is passed in firstQuery, and the number of queries is passed in queryCount.

vkGetQueryPoolResults() places the results of the requested queries in host memory pointed to by pData. The size of the memory region is passed in dataSize. Vulkan will not write more than this amount of data into memory. The result of each query is written to memory stride bytes apart. If stride is not at least as large as the amount of data produced by each query, the results may overwrite one another, and the result is undefined.

What is written to memory depends on the query type. flags provides additional information to Vulkan about how the queries should be reported. The flags available for use in flags are

- VK_QUERY_RESULT_64_BIT: If this bit is set, results are returned as 64-bit quantities; otherwise, results are returned as 32-bit quantities.

- VK_QUERY_RESULT_WAIT_BIT: If this bit is set, then **vkGetQueryPoolResults()** will wait until results of the queries are available. Otherwise, **vkGetQueryPoolResults()** returns a status code to report whether the commands contributing to the results of the queries were ready.

- `VK_QUERY_RESULT_WITH_AVAILABILITY_BIT`: If this bit is set, then Vulkan will write a zero result for queries that were not ready when **vkGetQueryPoolResults()** was called. Any query that was ready will have a nonzero result.

- `VK_QUERY_RESULT_PARTIAL_BIT`: If this bit is set, Vulkan might write the current value of a query into the result buffer even if the commands surrounded by the query have not finished executing.

It is also possible to write the results of queries directly into a buffer object. This allows results to be gathered asynchronously by the device, depositing results into a buffer for later use. The buffer can then either be mapped and accessed by the host or used as the source of data in subsequent graphics or compute operations.

To write the results of queries into a buffer object, call **vkCmdCopyQueryPoolResults()**, the prototype of which is

[**Click here to view code image**](#)

```
void vkCmdCopyQueryPoolResults (
    VkCommandBuffer                          commandBuffer,
    VkQueryPool                              queryPool,
    uint32_t                                 firstQuery,
    uint32_t                                 queryCount,
    VkBuffer                                 dstBuffer,
    VkDeviceSize                             dstOffset,
    VkDeviceSize                             stride,
    VkQueryResultFlags                       flags);
```

The command buffer that will execute the copy operation is specified in `commandBuffer`. This does not need to be the same command buffer that executed the queries. The `queryPool` parameter specifies the pool containing the queries that will be summarized into the buffer, and `firstQuery` and `queryCount` specify the index of the first query and the number of queries to copy, respectively. They have the same meanings as the similarly named parameters to **vkGetQueryPoolResults()**.

Rather than taking a pointer to host memory as **vkGetQueryPoolResults()** does, **vkCmdCopyQueryPoolResults()** takes a buffer object handle in `dstBuffer` and an offset into that buffer, measured in bytes, where the results will be written in `dstOffset`. The `stride` parameter is the number of bytes between each result in the buffer, and `flags` is a bitfield made up of the same flags as the `flags` parameter to **vkGetQueryPoolResults()**.

After **vkCmdCopyQueryPoolResults()** executes, access to the resulting values written to the buffer object must be synchronized using a barrier where the source is `VK_PIPELINE_STAGE_TRANSFER_BIT` and the access is `VK_ACCESS_TRANSFER_WRITE_BIT`.

Occlusion Queries

If the query type of the pool is `VK_QUERY_TYPE_OCCLUSION`, then the count is the number of fragments that pass the depth and stencil tests. This can be used to determine visibility or even to measure the area of geometry, in pixels. If the depth and stencil tests are disabled, then the result of the occlusion query is simply the area of the rasterized primitives.

A common use case is to render a view of part of a scene—for example, buildings or terrain—only to the depth buffer. Then render a simplified version of characters or other high-detail geometry—such as trees and vegetation, objects, or building details—with an occlusion query surrounding each. Such low-detail stand-ins are often referred to as *proxies*. Finally, make the decision as to whether to render the full-detail version of the object based on the result of each query. Because the query also tells you the approximate area of the object, you can perhaps render different versions of the final geometry based on its expected size onscreen. As an object gets farther away, you can substitute a version with lower geometric detail, substitute simplified shaders, or reduce tessellation levels, for example.

If you don't care too much about the visible area of an object but only about *whether* the object is visible, be sure not to set the `VK_QUERY_CONTROL_PRECISE_BIT` flag in the `flags` parameter when you create the query pool. If this flag is not set (indicating that you're interested only in approximate results), then the results of the queries should be treated as Boolean values. That is, they will be zero if the object was not visible and nonzero if it was. The actual value is not defined.

Pipeline Statistics Queries

Pipeline statistics queries allow your application to measure various aspects of the operation of the graphics pipeline. Each query can measure a number of different counters that are updated by the device as it executes the command buffer. The set of counters to enable is a property of the query pool and is specified at pool-creation time in the `pipelineStatistics` parameter.

The counters available and the flags that need to be set in `pipelineStatistics` are `VK_QUERY_PIPELINE_STATISTIC_...`

- `...INPUT_ASSEMBLY_VERTICES_BIT`: When enabled, the pipeline statistics query will count the number of vertices assembled by the vertex assembly stage of the graphics pipeline.

- `...INPUT_ASSEMBLY_PRIMITIVES_BIT`: When enabled, the pipeline statistics query will count the number of complete primitives assembled by the primitive assembly stage of the graphics pipeline.

- `...VERTEX_SHADER_INVOCATIONS_BIT`: When enabled, the pipeline statistics query will count the total number of invocations of the vertex shader produced in the graphics pipeline. Note that this may not be the same as the number of vertices assembled because Vulkan can sometimes skip vertex shader execution if a vertex is determined not to be part of a primitive or if it is part of multiple primitives and its result can be reused.

- `...GEOMETRY_SHADER_INVOCATIONS_BIT`: When enabled, the pipeline statistics query will count the total number of invocations of the geometry shader produced by the graphics pipeline.

- `...GEOMETRY_SHADER_PRIMITIVES_BIT`: When enabled, the pipeline statistics query will count the total number of primitives produced by the geometry shader.

- `...CLIPPING_INVOCATIONS_BIT`: When enabled, the pipeline statistics query will count the number of primitives that enter the clipping stage of the graphics pipeline. If a primitive can be trivially discarded without clipping, this counter does not increment.

- `...CLIPPING_PRIMITIVES_BIT`: When enabled, the pipeline statistics query will count the number of primitives produced when clipping. If the clipping stage of a Vulkan

implementation breaks primitives that clip against the viewport or a user-defined plane into multiple smaller primitives, this query will count those smaller primitives.

- `...FRAGMENT_SHADER_INVOCATIONS_BIT`: When enabled, the pipeline statistics query will count the total number of invocations of the fragment shader. This includes helper invocations and invocations of a fragment shader that produce a fragment that is ultimately discarded due to late depth or stencil tests.

- `...TESSELLATION_CONTROL_SHADER_PATCHES_BIT`: When enabled, the pipeline statistics query will count the total number of patches processed by the tessellation control shader. This is not the same as the number of tessellation control shader invocations because the tessellation control shader runs an invocation for each output control point in each patch, whereas this counter increments once for each *patch*.

- `...TESSELLATION_EVALUATION_SHADER_INVOCATIONS_BIT`: When enabled, the pipeline statistics query will increment each time the tessellation evaluation shader is invoked during tessellation processing. This is not necessarily the same as the number of vertices produced by the tessellator as, depending on the implementation, the tessellation evaluation shader may be invoked more than once for some tessellated vertices.

- `...COMPUTE_SHADER_INVOCATIONS_BIT`: When enabled, the pipeline statistics query will count the total number of compute shader invocations. This is the only counter that will count when dispatches are passed down the pipeline.

When reading the results of a pipeline statistics query, the number of counters written into memory (either host memory using **`vkGetQueryPoolResults()`** or buffer memory using **`vkCmdCopyQueryPoolResults()`**) depends on the number of enabled counters—i.e., the number of set bits in `pipelineStatistics`. Each result is written into consecutive 32- or 64-bit unsigned integers, and the start of each block of results is separated by `stride` bytes, as specified to the relevant command.

The counters are written into memory in the order of the lowest-valued member of the `VkQueryPipelineStatisticFlagBits` enumeration to the highest.

For a given set of enabled pipeline statistics queries and bitness of results, it is therefore possible to construct a C structure that represents the results. For example, the C structure shown in Listing 12.1 represents the full set of counters available when every defined bit is set in `pipelineStatistics` for 64-bit queries.

Listing 12.1: C Structure for All Pipeline Statistics

Click here to view code image

```
// Example structure containing all available pipeline statistics counters
typedef struct VkAllPipelineStatistics {
    uint64_t inputAssemblyVertices;
    uint64_t inputAssemblyPrimitives;
    uint64_t vertexShaderInvocations;
    uint64_t geometryShaderInvocations;
    uint64_t geometryShaderPrimitives;
    uint64_t clipperInvocations;
    uint64_t clipperOutputPrimitives;
    uint64_t fragmentShaderInvocations;
```

```
        uint64_t tessControlShaderPatches;
        uint64_t tessEvaluationShaderInvocations;
        uint64_t computeShaderEvaluations;
    } VkAllPipelineStatistics;
```

Note that there may be a performance penalty for accumulating the statistics queries or gathering their results. You should enable only the counters you actually need.

Because the results of some of the counters may be approximate, and because their exact values are generally implementation-dependent—for example, how many output primitives the clipping stage produces depends on how the clipper is implemented—you shouldn't use the results of these queries to compare Vulkan implementations. However, you can use these queries to get a measure of the relative complexity of different parts of your application, which can help you find bottlenecks while performance tuning.

Also, comparing different counters can give you insight into the operation of Vulkan. For example, by comparing the number of primitives produced by the primitive assembler with the number of clipper invocations and the number of clipper output primitives, you can determine some details of how clipping is implemented by a particular Vulkan device and how that device handles the geometry rendered by your application.

Timing Queries

A timing query measures the amount of time taken to execute commands in a command buffer. If the type of query in the query pool is `VK_QUERY_TYPE_TIMESTAMP`, then the values written into the output buffer are the number of nanoseconds taken to execute the commands in the command buffer between **vkCmdBeginQuery()** and **vkCmdEndQuery()**.

You can also retrieve an instantaneous measure of time from the pipeline by using the **vkCmdWriteTimestamp()** command to write the current device time into a slot in a query pool. The prototype of **vkCmdWriteTimestamp()** is

[Click here to view code image](#)

```
void vkCmdWriteTimestamp (
    VkCommandBuffer                     commandBuffer,
    VkPipelineStageFlagBits             pipelineStage,
    VkQueryPool                         queryPool,
    uint32_t                            query);
```

The command buffer that will write a timestamp into the pool is specified in `commandBuffer`, and the pool and the index within it where the timestamp should be written are specified in `queryPool` and `query`, respectively. When the device executes the **vkCmdWriteTimestamp()** command, it writes the current device time into the specified query at the pipeline stage specified in `pipelineStage`. This is a single member of the `VkPipelineStageFlagBits` enumeration, each bit representing a single logical pipeline stage.

In theory, it's possible to ask the device to write multiple timestamps from different stages of the pipeline. However, not all devices are capable of writing a timestamp from every stage of the pipeline. If a device cannot write a timestamp from a particular stage, it will write it at the next logical stage of the pipeline after the one requested. Therefore, the results might not reflect the actual time taken to process the commands between timestamps if they are not written from the same stage.

For example, if you perform a draw with tessellation enabled, request a timestamp after vertex processing by using `VK_PIPELINE_STAGE_VERTEX_SHADER_BIT`, and then request another after the tessellation evaluation shader has executed by using `VK_PIPELINE_STAGE_TESSELLATION_EVALUATION_SHADER_BIT`, you might expect that the measured time delta would be the amount of time taken to execute the tessellation control shader, the fixed-function tessellation, and the tessellation evaluation shader for that draw. However, if the implementation cannot write timestamps in the middle of geometry processing, it may honor your timestamps at some later stage in the pipeline, possibly even after fragment processing is complete. Hence, the time measured will be very small.

The units of measure used for timestamps are device-dependent. The timestamp will always be monotonically increasing, but each increment of the timestamp value represents a device-dependent number of nanoseconds. You can determine the number of nanoseconds represented by a single increment of the device timestamp by inspecting the `timestampPeriod` member of the device's `VkPhysicalDeviceProperties` structure, which you can retrieve by calling**vkGetPhysicalDeviceProperties()**.

To determine the absolute time delta between two timestamps, therefore, you should take the two integer timestamp values, subtract the first from the second to compute a delta in "ticks," and then multiply this integer delta by the floating-point `timestampPeriod` value in order to get a time in nanoseconds.

Reading Data with the Host

In some cases, it may be necessary to read data produced by the device into your application. Example uses of this include reading image data to take screen shots or, in compute applications, to save the results of compute shaders to disk. The primary mechanism for this is to issue commands to copy data into a buffer that is mapped and then read the data from the mapping using the host.

Memory allocations are mapped into host address space by calling**vkMapMemory()**. This function was introduced in [Chapter 2](#), "[Memory and Resources](#)." You can map memory and leave it mapped indefinitely. This is known as a *persistent mapping*. When you are performing a read from a mapped memory region, data is typically written to the memory by the device with a command such as**vkCmdCopyBuffer()** or **vkCmdCopyImageToBuffer()**. Before the host reads the memory, it must do two things:

- Ensure that the device has executed the command. This is normally accomplished by waiting on a fence that is signaled when the command buffer containing the copy command completes execution.

- Ensure that the host memory system's view of the data is the same as the device's.

If the memory mapping was made using a memory object allocated from a memory type that exposes the `VK_MEMORY_PROPERTY_HOST_COHERENT_BIT` property, then the mapping between the host and device is always coherent. That is, the host and device communicate in order to ensure that their respective caches are synchronized and that any reads or writes to shared memory are seen by the other peer.

If memory is not allocated from a memory type exposing the `VK_MEMORY_PROPERTY_HOST_COHERENT_BIT` property, then you must execute a pipeline barrier to move the resource into a host-readable state. To do this, make sure that the destination access mask includes `VK_ACCESS_HOST_READ_BIT`. [Listing 12.2](#) shows an example of how to

construct a pipeline barrier that moves a buffer resource (and therefore the memory backing it) into a host-readable state after being written to by a **vkCmdCopyImageToBuffer()** command.

Listing 12.2: Moving a Buffer to Host-Readable State

[Click here to view code image](#)

```
VkBufferMemoryBarrier bufferMemoryBarrier =
{
    VK_STRUCTURE_TYPE_BUFFER_MEMORY_BARRIER,  // sType
    nullptr,                                  // pNext
    VK_ACCESS_TRANSFER_WRITE_BIT,             // srcAccessMask
    VK_ACCESS_HOST_READ_BIT,                  // dstAccessMask
    VK_QUEUE_FAMILY_IGNORED,                  // srcQueueFamilyIndex
    VK_QUEUE_FAMILY_IGNORED,                  // dstQueueFamilyIndex
    buffer,                                   // buffer
    0,                                        // offset
    VK_WHOLE_SIZE                             // size
};

vkCmdPipelineBarrier(
    cmdBuffer,                                // commandBuffer
    VK_PIPELINE_STAGE_TRANSFER_BIT,           // srcStageMask
    VK_PIPELINE_STAGE_HOST_BIT,               // dstStageMask
    0,                                        // dependencyFlags
    0,                                        // memoryBarrierCount
    nullptr,                                  // pMemoryBarriers
    1,                                        // bufferMemoryBarrierCount
    &bufferMemoryBarrier,                     // pBufferMemoryBarriers
    0,                                        // imageMemoryBarrierCount
    nullptr);                                 // pImageMemoryBarriers
```

In [Listing 12.2](#), we insert a pipeline barrier after the **vkCmdCopyImageToBuffer()** command (not shown in the listing), which is considered to be a transfer command. Therefore, the source pipeline stage is VK_PIPELINE_STAGE_TRANSFER_BIT, and as the host will read the data, the destination pipeline stage is VK_PIPELINE_STAGE_HOST_BIT. These stages are virtual pipeline stages that do not participate in normal graphics operations but potentially represent points in internally created pipelines that perform copy operations.

In addition to the pipeline stages, which is what the barrier synchronizes, we specify the access mask for each end of the barrier. The source of data is writes by the transfer operation, so we specify VK_ACCESS_TRANSFER_WRITE_BIT, and the destination of the data is reads performed by the host, so we specify VK_ACCESS_HOST_READ_BIT. These bits are used to ensure that any caches that need to be synchronized are correctly made coherent between device and host.

Summary

In this chapter, you read about the two ways that Vulkan can produce information that your application can consume: queries and explicit data reads.

Queries provide a series of counters that can be enabled and will increment when events occur inside the graphics pipeline. You learned about occlusion queries, which count fragments that pass the depth

and stencil tests. You learned about pipeline statistics queries, which can provide insight into the inner operation of Vulkan graphics and compute pipelines. You saw how timing queries allow you to measure the amount of time taken to execute commands inside a command buffer, as well as how to ask Vulkan to write immediate timestamps into buffers that you can read.

Finally, you learned about reading data from mapped buffers and correctly synchronizing access to buffers from the device and the host.

Chapter 13. Multipass Rendering

What You'll Learn in This Chapter

- How to use renderpass objects to accelerate multipass rendering
- How to fold clears and barriers into renderpass objects
- How to control how and when data gets saved to memory

Many graphics applications make multiple passes over each frame or are otherwise able to subdivide rendering into multiple logical phases. Vulkan brings this into the core of its operation, exposing the concept of multipass rendering within a single object. This object was briefly introduced in Chapter 7, "Graphics Pipelines," but we skimmed over many of the details, instead going only into enough depth to enable basic single-pass rendering to be achieved. In this chapter, we dig deeper into the topic to explain how multipass rendering algorithms can be implemented in a few renderpass objects or even a single one.

When we introduced the renderpass object back in Chapter 7, "Graphics Pipelines," we covered it in only enough detail to explain how a framebuffer can be attached to a command buffer at the beginning of a renderpass and how the renderpass could be configured to allow drawing into a single set of color attachments. A renderpass object, however, can contain many subpasses, each performing some of the operations required to render the final scene. Dependency information can be introduced, allowing a Vulkan implementation to build a directed acyclic graph (DAG) and determine where data flows, who produces it and who consumes it, what needs to be ready by when, and so on.

Input Attachments

Recall the `VkRenderPassCreateInfo` structure, the definition of which is

Click here to view code image

```
typedef struct VkRenderPassCreateInfo {
    VkStructureType                 sType;
    const void*                     pNext;
    VkRenderPassCreateFlags         flags;
    uint32_t                        attachmentCount;
    const VkAttachmentDescription*  pAttachments;
    uint32_t                        subpassCount;
    const VkSubpassDescription*     pSubpasses;
    uint32_t                        dependencyCount;
    const VkSubpassDependency*      pDependencies;
} VkRenderPassCreateInfo;
```

Within this structure, we have pointers to arrays of attachments, subpasses, and dependency information. Each subpass is defined by a `VkSubpassDescription` structure, the definition of which is

```
typedef struct VkSubpassDescription {
    VkSubpassDescriptionFlags        flags;
    VkPipelineBindPoint              pipelineBindPoint;
    uint32_t                         inputAttachmentCount;
    const VkAttachmentReference *    pInputAttachments;
    uint32_t                         colorAttachmentCount;
    const VkAttachmentReference *    pColorAttachments;
    const VkAttachmentReference *    pResolveAttachments;
    const VkAttachmentReference *    pDepthStencilAttachment;
    uint32_t                         preserveAttachmentCount;
    const uint32_t*                  pPreserveAttachments;
} VkSubpassDescription;
```

In the example renderpass we set up in Chapter 7, we used a single subpass with no dependencies and a single set of outputs. However, each subpass can have one or more input attachments, which are attachments from which you can *read* in your fragment shaders. The primary difference between an input attachment and a normal texture bound into a descriptor set is that when you read from an input attachment, you read from the current fragment.

Each subpass may write to one or more output attachments. These are either the color attachments or the depth-stencil attachment (of which there is only one). By inspecting the subpasses, which output attachments they write to and input attachments they read from, Vulkan can build up a graph of where data flows within a renderpass.

In order to demonstrate this, we will construct a simple three-pass renderpass object that performs deferred shading. In the first pass, we render only to a depth attachment in order to produce what is known as a depth prepass.

In a second pass, we render all the geometry with a special shader that produces a g-buffer, which is a color attachment (or set of color attachments) that stores a normal diffuse color, specular power, and other parameters needed for shading. In this second pass, we test against the depth buffer we just generated, so we do not write out the large amounts of data for geometry that will not be visible in the final output. Even when the geometry is visible, we do not run complex lighting shaders; we simply write out data.

In our third pass, we perform all of our shading calculations. We read from the depth buffer in order to reconstruct the view-space position, which allows us to create our eye and view vectors. We also read from our normal, specular, and diffuse buffers, which supply parameters for our lighting computation. Note that in the third pass, we don't actually need the real geometry, and we instead render a single triangle that, after clipping, covers the entire viewport.

Figure 13.1 shows this schematically. As you can see from the figure, the first subpass has no inputs and only a depth attachment. The second subpass uses the same depth attachment for testing but also has no inputs, producing only outputs. The third and final pass uses the depth buffer produced by the first pass and the g-buffer attachments produced by the second pass as input attachments. It can do this because the lighting calculations at each pixel require only the data computed by previous shader invocations *at the same location*.
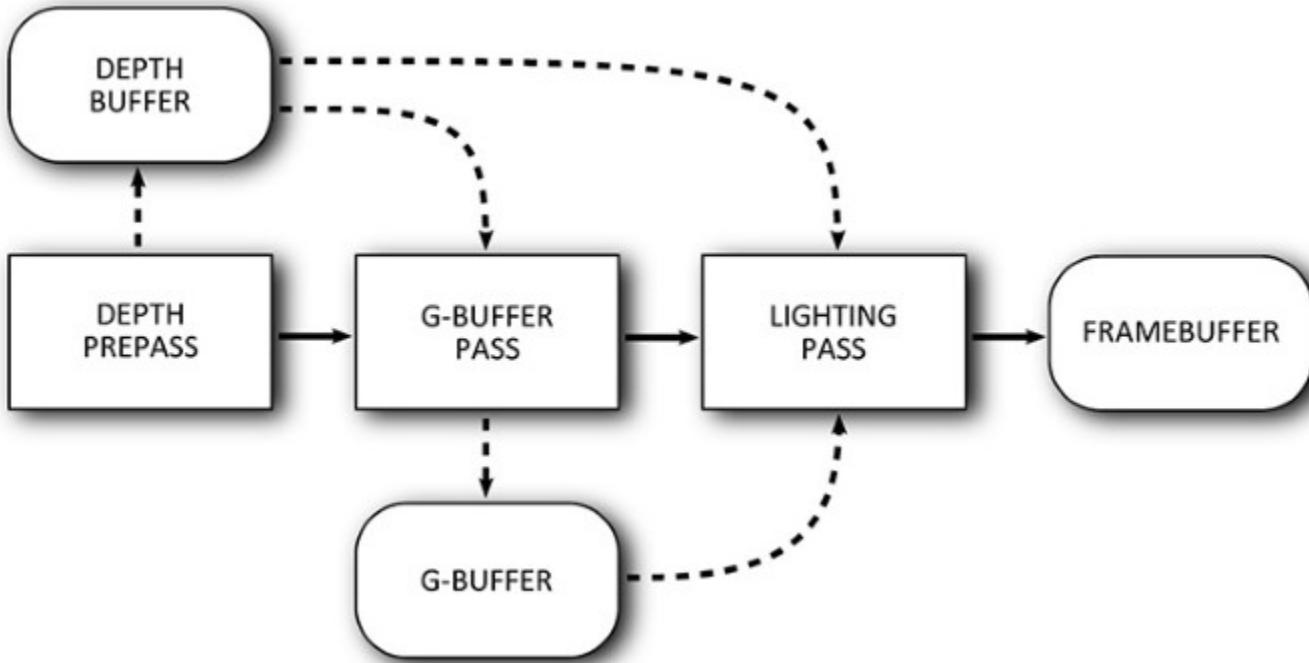
Figure 13.1: Data Flow for a Simple Deferred Renderer

Listing 13.1 shows the code required to construct a renderpass that represents these three subpasses and their attachments.

Listing 13.1: Deferred Shading Renderpass Setup

Click here to view code image

```
enum
{
    kAttachment_BACK            = 0,
    kAttachment_DEPTH           = 1,
    kAttachment_GBUFFER         = 2
};
enum
{
    kSubpass_DEPTH              = 0,
    kSubpass_GBUFFER            = 1,
    kSubpass_LIGHTING           = 2
};

static const VkAttachmentDescription attachments[] =
{
    // Back buffer
    {
        0,                                      // flags
        VK_FORMAT_R8G8B8A8_UNORM,              // format
        VK_SAMPLE_COUNT_1_BIT,                 // samples
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,       // loadOp
        VK_ATTACHMENT_STORE_OP_STORE,          // storeOp
```

```
            VK_ATTACHMENT_LOAD_OP_DONT_CARE,    // stencilLoadOp
            VK_ATTACHMENT_STORE_OP_DONT_CARE,   // stencilStoreOp
            VK_IMAGE_LAYOUT_UNDEFINED,          // initialLayout
            VK_IMAGE_LAYOUT_PRESENT_SRC_KHR     // finalLayout
        },
        // Depth buffer
        {
            0,                                  // flags
            VK_FORMAT_D32_SFLOAT,               // format
            VK_SAMPLE_COUNT_1_BIT,              // samples
            VK_ATTACHMENT_LOAD_OP_DONT_CARE,    // loadOp
            VK_ATTACHMENT_STORE_OP_DONT_CARE,   // storeOp
            VK_ATTACHMENT_LOAD_OP_DONT_CARE,    // stencilLoadOp
            VK_ATTACHMENT_STORE_OP_DONT_CARE,   // stencilStoreOp
            VK_IMAGE_LAYOUT_UNDEFINED,          // initialLayout
            VK_IMAGE_LAYOUT_UNDEFINED           // finalLayout
        },
        // G-buffer 1
        {
            0,                                  // flags
            VK_FORMAT_R32G32B32A32_UINT,        // format
            VK_SAMPLE_COUNT_1_BIT,              // samples
            VK_ATTACHMENT_LOAD_OP_DONT_CARE,    // loadOp
            VK_ATTACHMENT_STORE_OP_DONT_CARE,   // storeOp
            VK_ATTACHMENT_LOAD_OP_DONT_CARE,    // stencilLoadOp
            VK_ATTACHMENT_STORE_OP_DONT_CARE,   // stencilStoreOp
            VK_IMAGE_LAYOUT_UNDEFINED,          // initialLayout
            VK_IMAGE_LAYOUT_UNDEFINED           // finalLayout
        }
    };
    // Depth prepass depth buffer reference (read/write)
    static const VkAttachmentReference depthAttachmentReference =
    {
        kAttachment_DEPTH,                                  // attachment
        VK_IMAGE_LAYOUT_DEPTH_STENCIL_ATTACHMENT_OPTIMAL    // layout
    };

    // G-buffer attachment references (render)
    static const VkAttachmentReference gBufferOutputs[] =
    {
        {
            kAttachment_GBUFFER,                            // attachment
            VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL        // layout
        }
    };

    // Lighting input attachment references
    static const VkAttachmentReference gBufferReadRef[] =
    {
        // Read from g-buffer.
        {
            kAttachment_GBUFFER,                            // attachment
            VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL        // layout
```

```cpp
        },
        // Read depth as texture.
        {
            kAttachment_DEPTH,                          // attachment
            VK_IMAGE_LAYOUT_DEPTH_STENCIL_READ_ONLY_OPTIMAL // layout
        }
    };


    // Final pass-back buffer render reference
    static const VkAttachmentReference backBufferRenderRef[] =
    {
        {
            kAttachment_BACK,                          // attachment
            VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL        // layout
        }
    };


    static const VkSubpassDescription subpasses[] =
    {
        // Subpass 1 - depth prepass
        {
            0,                                      // flags
            VK_PIPELINE_BIND_POINT_GRAPHICS,      // pipelineBindPoint
            0,                                      // inputAttachmentCount
            nullptr,                                // pInputAttachments
            0,                                      // colorAttachmentCount
            nullptr,                                // pColorAttachments
            nullptr,                                // pResolveAttachments
            &depthAttachmentReference,              // pDepthStencilAttachment
            0,                                      // preserveAttachmentCount
            nullptr                                 // pPreserveAttachments
        },
        // Subpass 2 - g-buffer generation
        {
            0,                                      // flags
            VK_PIPELINE_BIND_POINT_GRAPHICS,      // pipelineBindPoint
            0,                                      // inputAttachmentCount
            nullptr,                                // pInputAttachments
            vkcore::utils::arraysize(gBufferOutputs), // colorAttachmentCount
            gBufferOutputs,                         // pColorAttachments
            nullptr,                                // pResolveAttachments
            &depthAttachmentReference,              // pDepthStencilAttachment
            0,                                      // preserveAttachmentCount
            nullptr                                 // pPreserveAttachments
        },
        // Subpass 3 - lighting
        {
            0,                                      // flags
            VK_PIPELINE_BIND_POINT_GRAPHICS,      // pipelineBindPoint
            vkcore::utils::arraysize(gBufferReadRef),  // inputAttachmentCount
            gBufferReadRef,                         // pInputAttachments
            vkcore::utils::arraysize(backBufferRenderRef),//
colorAttachmentCount
```

```
            backBufferRenderRef,                     // pColorAttachments
            nullptr,                                 // pResolveAttachments
            nullptr,                                 // pDepthStencilAttachment
            0,                                       // preserveAttachmentCount
            nullptr                                  // pPreserveAttachments
        },
    };

    static const VkSubpassDependency dependencies[] =
    {
        // G-buffer pass depends on depth prepass.
        {
            kSubpass_DEPTH,                                         // srcSubpass
            kSubpass_GBUFFER,                                       // dstSubpass
            VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,  // srcStageMask
            VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,          // dstStageMask
            VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,           // srcAccessMask
            VK_ACCESS_SHADER_READ_BIT,                      // dstAccessMask
            VK_DEPENDENCY_BY_REGION_BIT                     // dependencyFlags
        },
        // Lighting pass depends on g-buffer.
        {
            kSubpass_GBUFFER,                                       // srcSubpass
            kSubpass_LIGHTING,                                      // dstSubpass
            VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,  // srcStageMask
            VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,          // dstStageMask
            VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,           // srcAccessMask
            VK_ACCESS_SHADER_READ_BIT,                      // dstAccessMask
            VK_DEPENDENCY_BY_REGION_BIT                     // dependencyFlags
        },
    };

    static const VkRenderPassCreateInfo renderPassCreateInfo =
    {
        VK_STRUCTURE_TYPE_RENDER_PASS_CREATE_INFO, nullptr,
        0,                                              // flags
        vkcore::utils::arraysize(attachments),          // attachmentCount
        attachments,                                    // pAttachments
        vkcore::utils::arraysize(subpasses),            // subpassCount
        subpasses,                                      // pSubpasses
        vkcore::utils::arraysize(dependencies),         // dependencyCount
        dependencies                                    // pDependencies
    };

    result = vkCreateRenderPass(device,
                                &renderPassCreateInfo,
                                nullptr,
                                &m_renderPass);
```

As you can see, Listing 13.1 is quite long. However, the code complexity is relatively low; most of the listing is simply definitions of static data structures describing the renderpass.

The `attachments[]` array contains a list of all of the attachments used in the renderpass. This is referenced by index by the arrays of `VkAttachmentReference` structures,

`depthAttachmentReference`, `gBufferOutputs`, `gBufferReadRef`, and `backBufferRenderRef`. These reference the depth buffer, the g-buffer as an output, the g-buffer as an input, and the back buffer, respectively.

The `subpasses[]` array is a description of the subpasses contained in the renderpass. Each is desribed by an instance of the `VkSubpassDescription` structure, and you can see there is one for each of the depth prepass, g-buffer generation, and lighting passes.

Note that for the lighting pass, we include the g-buffer read references and the depth buffer as input attachments to the pass. This is so that the lighting computations performed in the shader can read the g-buffer content and the pixels' depth value.

Finally, we see the dependencies between the passes described in the `dependencies[]` array. There are two entries in the array, the first describing the dependency of the g-buffer pass on the depth prepass, and the second describing the dependency of the lighting pass on the g-buffer pass. Note that there is no reason to have a dependency between the lighting pass and the depth prepass even though one technically exists, because there is already an implicit depencency through the g-buffer generation pass.

The subpasses inside a renderpass are logically executed in the order in which they are declared in the array of subpasses referenced by the `VkRenderPassCreateInfo` structure used to create the renderpass object. When **vkCmdBeginRenderPass()** is called, the first subpass in the array is automatically begun. In simple renderpasses with a single subpass, this is sufficient to execute the entire renderpass. However, once we have multiple subpasses, we need to be able to tell Vulkan when to move from subpass to subpass.

To do this, we call **vkCmdNextSubpass()**, the prototype of which is

[**Click here to view code image**](#)

```
void vkCmdNextSubpass (
    VkCommandBuffer                         commandBuffer,
    VkSubpassContents                       contents);
```

The command buffer in which to place the command is specified in `commandBuffer`. The `contents` parameter specifies where the commands for the subpass will come from. For now, we'll set this to `VK_SUBPASS_CONTENTS_INLINE`, which indicates that you will continue to add commands to the same command buffer. It's also possible to call other command buffers, in which case we'd set this to `VK_SUBPASS_CONTENTS_SECONDARY_COMMAND_BUFFERS`. We'll cover this scenario later in this chapter.

When **vkCmdNextSubpass()** is called, the current command buffer moves to the next subpass in the current renderpass. Correspondingly, you can call **vkCmdNextSubpass()** only between calls to **vkCmdBeginRenderPass()** and **vkCmdEndRenderPass()**, and only until you have exhausted the subpasses within the renderpass.

With renderpasses containing multiple subpasses, we still must call **vkCmdEndRenderPass()** to terminate the current renderpass and finalize rendering.

Attachment Contents

Each color and depth-stencil attachment associated with a renderpass has a load operation and a store operation that determine how its contents are loaded from and stored to memory as the renderpass is begun and ended.

Attachment Initialization

When the renderpass is begun, the operations that should be performed on each of the attachments are specified in the `loadOp` field of the `VkAttachmentDescription` structure describing the attachment. There are two possible values for this field.

`VK_ATTACHMENT_LOAD_OP_DONT_CARE` means that you don't care about the initial contents of the attachment. This means that Vulkan can do whatever it needs to do to get the attachment ready to render into (including doing nothing), without worrying about the actual values in the attachment. For example, if it has a super-fast clear that clears only to purple, then purple it shall be.

Setting `loadOp` for the attachment to `VK_ATTACHMENT_LOAD_OP_CLEAR` means that the attachment will be cleared to a value you specify at **vkCmdBeginRenderPass()** time. While logically, this operation happens at the very start of the renderpass, in practice, implementations may delay the actual clear operation to the beginning of the first pass that uses the attachment. This is the preferred method of clearing color attachments.

It's also possible to explicitly clear one or more color or depth-stencil attachments inside a renderpass by calling **vkCmdClearAttachments()**, the prototype of which is

[**Click here to view code image**](#)

```
void vkCmdClearAttachments (
    VkCommandBuffer                          commandBuffer,
    uint32_t                                 attachmentCount,
    const VkClearAttachment*                 pAttachments,
    uint32_t                                 rectCount,
    const VkClearRect*                       pRects);
```

The command buffer that will execute the command is specified in `commandBuffer`. **vkCmdClearAttachments()** will clear the contents of several attachments. The number of attachments to clear is specified in `attachmentCount`, and `pAttachments` should be a pointer to an array of `attachmentCount` `VkClearAttachment` structures, each defining one of the attachments to clear. The definition of `VkClearAttachment` is

[**Click here to view code image**](#)

```
typedef struct VkClearAttachment {
    VkImageAspectFlags    aspectMask;
    uint32_t              colorAttachment;
    VkClearValue          clearValue;
} VkClearAttachment;
```

The `aspectMask` field of `VkClearAttachment` specifies the aspect or aspects of the attachment to be cleared. If `aspectMask` contains `VK_IMAGE_ASPECT_DEPTH_BIT`, `VK_IMAGE_ASPECT_STENCIL_BIT`, or both, then the clear operation will be applied to the depth-stencil attachment for the current subpass. Each subpass can have at most one depth-stencil

attachment. If `aspectMask` contains `VK_IMAGE_ASPECT_COLOR_BIT`, then the clear operation will be applied to the color attachment at index `colorAttachment` in the current subpass. It is not possible to clear a color attachment and a depth-stencil attachment with a single `VkClearAttachment` structure, so `aspectMask` should not contain `VK_IMAGE_ASPECT_COLOR_BIT` along with `VK_IMAGE_ASPECT_DEPTH_BIT` or `VK_IMAGE_ASPECT_STENCIL_BIT`.

The values to clear the attachment with are specified in the `clearValue` field, which is an instance of the `VkClearValue` union. This was introduced in [Chapter 8](), "[Drawing]()," and its definition is

Click here to view code image

```
typedef union VkClearValue {
    VkClearColorValue              color;
    VkClearDepthStencilValue       depthStencil;
} VkClearValue;
```

If the referenced attachment is a color attachment, then the values from the `color` field of the `VkClearAttachment` structure will be used. Otherwise, the values contained in the `depthStencil` field of the structure will be used.

In addition to clearing multiple attachments, a single call to **vkCmdClearAttachments()** can clear rectangular regions of each attachment. This provides additional functionality over setting the `loadOp` for the attachment to `VK_ATTACHMENT_LOAD_OP_CLEAR`. If an attachment is cleared with `VK_ATTACHMENT_LOAD_OP_CLEAR` (which is what you want in the majority of cases), the whole attachment is cleared, and there is no opportunity to clear only part of the attachment. However, when you use **vkCmdClearAttachments()**, multiple smaller regions can be cleared.

The number of regions to clear is specified in the `rectCount` parameter to **vkCmdClearAttachments()**, and the `pRects` parameter is a pointer to an array of `rectCount` `VkClearRect` structures, each defining one of the rectangles to clear. The definition of `VkClearRect` is

Click here to view code image

```
typedef struct VkClearRect {
    VkRect2D    rect;
    uint32_t    baseArrayLayer;
    uint32_t    layerCount;
} VkClearRect;
```

The `VkClearRect` structure defines more than a rectangle. The `rect` field contains the actual rectangle to clear. If the attachment is an array image, then some or all of its layers can be cleared by specifying the range of layers in `baseArrayLayer` and `layerCount`, which contain, respectively, the index of the first layer to clear and the number of layers to clear.

In addition to containing more information and providing more functionality than the attachments load operation, **vkCmdClearAttachments()** also potentially provides more convenience than **vkCmdClearColorImage()** or **vkCmdClearDepthStencilImage()**. Both of these commands require that the image be in `VK_IMAGE_LAYOUT_GENERAL` or `VK_IMAGE_LAYOUT_TRANSFER_DST_OPTIMAL` layout and therefore may require a pipeline barrier around them to ensure this. Also, these commands cannot be called inside a renderpass. On

the other hand, **vkCmdClearAttachments()** takes advantage of the fact that the attachments are already bound for rendering and are considered to be part of the renderpass content, almost like a special kind of draw. Therefore, no barrier or special handling is needed beyond ensuring that the command is executed inside a renderpass.

That said, it is still recommended that you use the VK_ATTACHMENT_LOAD_OP_CLEAR load operation when you need the entirety of an attachment to be cleared as part of a renderpass, and the VK_ATTACHMENT_LOAD_OP_DONT_CARE operation when you will guarantee that you will overwrite every pixel in the attachment by the time the renderpass has completed.

Render Areas

When a renderpass instance is executed, it is possible to tell Vulkan that you're going to update only a small area of the attachments. This area is known as the *render area* and can be specified when **vkCmdBeginRenderPass()** is called. We introduced this briefly in Chapter 8, "Drawing," and the renderArea member of the VkRenderPassBeginInfo structure is passed to **vkCmdBeginRenderPass()**.

If you are rendering to the entire framebuffer, set the renderArea field to cover the entire area of the framebuffer. However, if you want to update only a small part of the framebuffer, you can set the renderArea rectangle accordingly. Any part of the renderpass's attachments that are not contained inside this render area are not affected by any of the operations in the renderpass including the renderpass's load and store operations for those attachments.

When you are using a render area smaller than the entire attachment, it is the application's responsibility to ensure that it doesn't render outside this area. Some implementations may ignore the render area entirely and trust your application to stick within it; some may round the render area up to some multiple of an internal rectangular region; and some may strictly adhere to the area you've specified. The only way to get well-defined behavior is to make sure you render only inside this area, using a scissor test if needed.

Rendering to a smaller render area than the entire attachment may also come at some performance cost unless the area matches the *granularity* of the supported render areas for the implementation. Consider the framebuffer to be tiles in a grid that are potentially rendered one at a time. Completely covering and redefining the content of a single tile should be fast, but updating only part of a tile may cause Vulkan to do extra work to keep the untouched parts of the tile well defined.

To acheive maximum performance, you should ensure that the render areas you use match the render area granularity supported by the implementation. This can be queried using **vkGetRenderAreaGranularity()**, the prototype of which is

[Click here to view code image](#)

```
void vkGetRenderAreaGranularity (
    VkDevice                              device,
    VkRenderPass                          renderPass,
    VkExtent2D*                           pGranularity);
```

For a renderpass specified in renderPass, **vkGetRenderAreaGranularity()** returns, in the variable pointed to by pGranularity, the dimensions of a tile used for rendering. The device that owns the renderpass should be passed in device.

To ensure that the render area you pass to **vkCmdBeginRenderPass()** performs optimally, you should ensure two things: First, that the x and y components of its origin are integer multiples of the width and height of the render-area granularity; and second, that the width and height of the render area are either integer multiples of that granularity or extend to the edge of the framebuffer. Obviously, a render area that completely covers the attachments trivially meets these requirements.

Preserving Attachment Content

In order to preserve the contents of the attachment, we need to set the attachment's store operation (contained in the storeOp field of the VkAttachmentDescription structure used to create the renderpass) to VK_ATTACHMENT_STORE_OP_STORE. This causes Vulkan to ensure that after the renderpass has completed, the contents of the image used as the attachment accurately reflect what was rendered during the renderpass.

The only other choice for this field is VK_ATTACHMENT_STORE_OP_DONT_CARE, which tells Vulkan that you don't need the content of the attachment after the renderpass has completed. This is used, for example, when an attachment is used to store intermediate data that will be consumed by some later subpass in the same renderpass. In this case, the content doesn't need to live longer than the renderpass itself.

In some cases, you need to produce content in one subpass, execute an unrelated subpass, and then consume the content created more than one subpass ago. In this case, you should tell Vulkan that it cannot discard the content of an attachment over the course of rendering another subpass. In practice, a Vulkan implementation should be able to tell by inspecting the input and output attachments for the subpasses in a renderpass which ones produce and which ones consume data and will do the right thing. However, to be fully correct, every live attachment should appear as an input, an output, or a preserve attachment in each subpass. Furthermore, by including an attachment in the preserve attachment array for a subpass, you are telling Vulkan that you are about to use the attachment content in an upcoming subpass. This may enable it to keep some of the data in cache or some other high-speed memory.

The list of attachments to preserve across a subpass is specified using the pPreserveAttachments field of the VkSubpassDescription structure describing each subpass. This is a pointer to an array of uint32_t indices into the renderpass's attachment list, and the number of integers in this array is contained in the preserveAttachmentCount field of VkSubpassDescription.

To demonstrate this, we extend our example further to render transparent objects. We render a depth buffer, then render a *g-buffer* containing per-pixel information, and finally we render a shading pass that calculates lighting information. Because we have only 1 pixel's worth of information, this deferred shading approach cannot render transparent or translucent objects. Therefore, these objects must be rendered separately. The traditional approach is to simply render all the opaque geometry in one pass (or passes) and then composite the translucent geometry on top at the end. This introduces a serial dependency, causing rendering of the translucent geometry to wait for opaque geometry to complete rendering.

This serial dependency is shown in the DAG illustrated in .
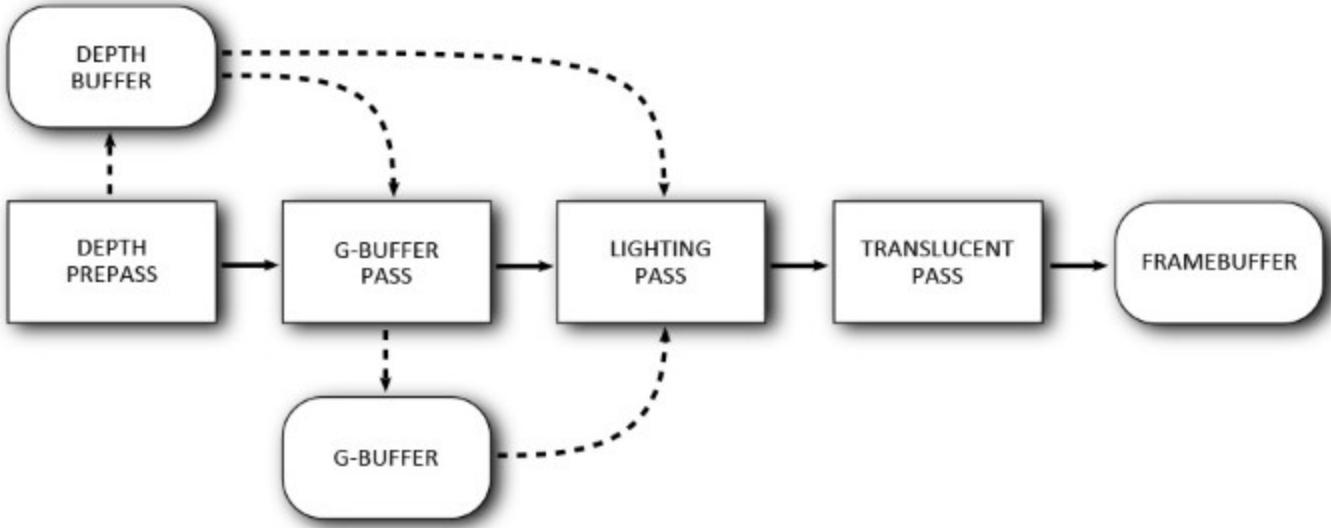
Figure 13.2: Serial Dependency of Translucent on Opaque Geometry

Rather than introduce a serial dependency, which would preclude Vulkan from rendering any of the translucent geometry in parallel with opaque geometry, we take another approach: Render the translucent geometry to another color attachment (using the same depth prepass information for depth rejection) and the opaque geometry into a second temporary attachment. After both the opaque and transparent geometry have been rendered, we perform a composition pass that blends the translucent geometry on top of the opaque geometry. This pass can also perform other per-pixel operations, such as color grading, vignetting, film-grain application, and so on. The new DAG for this is shown in Figure 13.3.
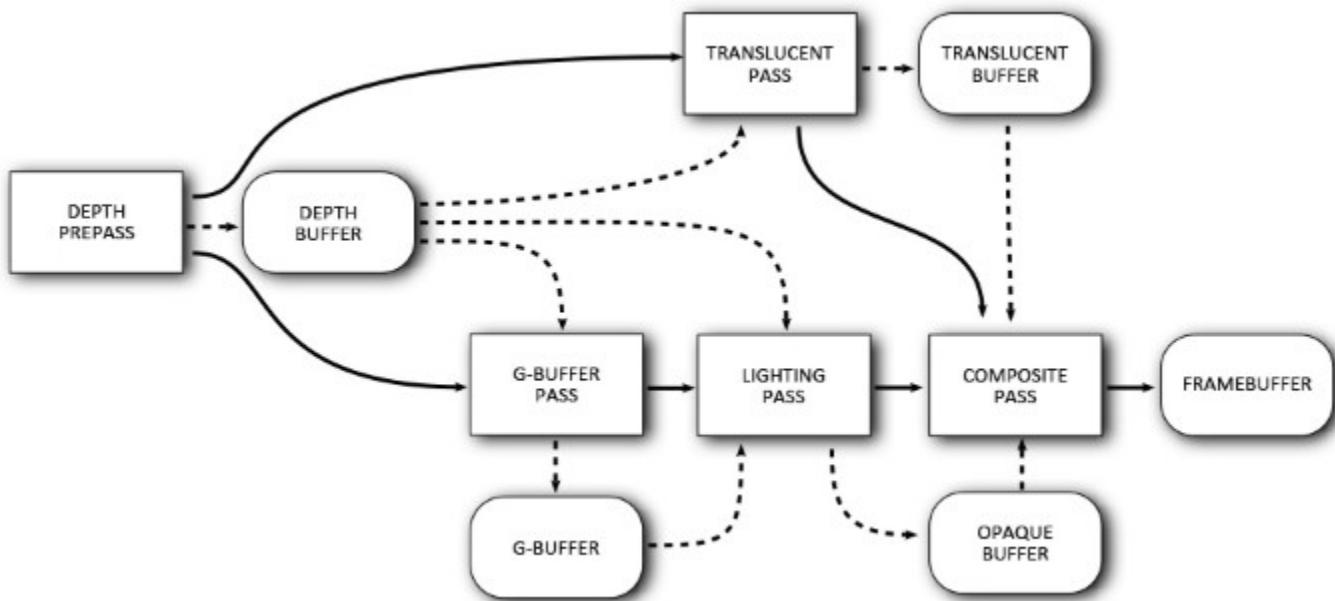
Figure 13.3: Parallel Rendering of Translucent and Opaque Geometry

As you can see from the updated DAG in Figure 13.3, first the depth information is rendered and then the g-buffer generation pass executes, followed by the lighting pass. The translucency buffer generation pass has no dependency on the g-buffer or the result of the lighting pass, so it is able to

run in parallel, depending only on the depth information from the depth prepass. The new composite pass now depends on the result of the lighting pass and the translucent pass.

As the subpasses in the renderpass are expressed serially, regardless of the serial ordering of the opaque g-buffer pass and the translucency pass in the renderpass, we need to preserve the content of the first pass's outputs until the shading pass is able to execute. Because there is less data to store, we render the translucent objects first and preserve the translucency buffer across the g-buffer generation pass. The g-buffer and translucency buffer are then used as input attachments to the shading pass.

The code to set all this up is shown in .

Listing 13.2: Translucency and Deferred Shading Setup

Click here to view code image

```
enum
{
    kAttachment_BACK         = 0,
    kAttachment_DEPTH        = 1,
    kAttachment_GBUFFER      = 2,
    kAttachment_TRANSLUCENCY = 3,
    kAttachment_OPAQUE       = 4
};

enum
{
    kSubpass_DEPTH           = 0,
    kSubpass_GBUFFER         = 1,
    kSubpass_LIGHTING        = 2,
    kSubpass_TRANSLUCENTS    = 3,
    kSubpass_COMPOSITE       = 4
};

static const VkAttachmentDescription attachments[] =
{
    // Back buffer
    {
        0,                                  // flags
        VK_FORMAT_R8G8B8A8_UNORM,           // format
        VK_SAMPLE_COUNT_1_BIT,              // samples
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,    // loadOp
        VK_ATTACHMENT_STORE_OP_STORE,       // storeOp
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,    // stencilLoadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,   // stencilStoreOp
        VK_IMAGE_LAYOUT_UNDEFINED,          // initialLayout
        VK_IMAGE_LAYOUT_PRESENT_SRC_KHR     // finalLayout
    },
    // Depth buffer
    {
        0,                                  // flags
        VK_FORMAT_D32_SFLOAT,               // format
        VK_SAMPLE_COUNT_1_BIT,              // samples
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,    // loadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,   // storeOp
```

```
            VK_ATTACHMENT_LOAD_OP_DONT_CARE,     // stencilLoadOp
            VK_ATTACHMENT_STORE_OP_DONT_CARE,    // stencilStoreOp
            VK_IMAGE_LAYOUT_UNDEFINED,           // initialLayout
            VK_IMAGE_LAYOUT_UNDEFINED            // finalLayout
        },
        // G-buffer 1
        {
            0,                                   // flags
            VK_FORMAT_R32G32B32A32_UINT,         // format
            VK_SAMPLE_COUNT_1_BIT,               // samples
            VK_ATTACHMENT_LOAD_OP_DONT_CARE,     // loadOp
            VK_ATTACHMENT_STORE_OP_DONT_CARE,    // storeOp
            VK_ATTACHMENT_LOAD_OP_DONT_CARE,     // stencilLoadOp
            VK_ATTACHMENT_STORE_OP_DONT_CARE,    // stencilStoreOp
            VK_IMAGE_LAYOUT_UNDEFINED,           // initialLayout
            VK_IMAGE_LAYOUT_UNDEFINED            // finalLayout
        },
        // Translucency buffer
        {
            0,                                   // flags
            VK_FORMAT_R8G8B8A8_UNORM,            // format
            VK_SAMPLE_COUNT_1_BIT,               // samples
            VK_ATTACHMENT_LOAD_OP_DONT_CARE,     // loadOp
            VK_ATTACHMENT_STORE_OP_DONT_CARE,    // storeOp
            VK_ATTACHMENT_LOAD_OP_DONT_CARE,     // stencilLoadOp
            VK_ATTACHMENT_STORE_OP_DONT_CARE,    // stencilStoreOp
            VK_IMAGE_LAYOUT_UNDEFINED,           // initialLayout
            VK_IMAGE_LAYOUT_UNDEFINED            // finalLayout
        }
    };


    // Depth prepass depth buffer reference (read/write)
    static const VkAttachmentReference depthAttachmentReference =
    {
        kAttachment_DEPTH,                               // attachment
        VK_IMAGE_LAYOUT_DEPTH_STENCIL_ATTACHMENT_OPTIMAL // layout
    };
    // G-buffer attachment references (render)
    static const VkAttachmentReference gBufferOutputs[] =
    {
        {
            kAttachment_GBUFFER,                         // attachment
            VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL     // layout
        }
    };


    // Lighting input attachment references
    static const VkAttachmentReference gBufferReadRef[] =
    {
        // Read from g-buffer.
        {
            kAttachment_GBUFFER,                         // attachment
            VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL     // layout
```

```
    },
    // Read depth as texture.
    {
        kAttachment_DEPTH,                                // attachment
        VK_IMAGE_LAYOUT_DEPTH_STENCIL_READ_ONLY_OPTIMAL // layout
    }
};


// Lighting pass - write to opaque buffer.
static const VkAttachmentReference opaqueWrite[] =
{
    // Write to opaque buffer.
    {
        kAttachment_OPAQUE,                               // attachment
        VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL        // layout
    }
};


// Translucency rendering pass - translucency buffer write
static const VkAttachmentReference translucentWrite[] =
{
    // Write to translucency buffer.
    {
        kAttachment_TRANSLUCENCY,                         // attachment
        VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL        // layout
    }
};


static const VkAttachmentReference compositeInputs[] =
{
    // Read from translucency buffer.
    {
        kAttachment_TRANSLUCENCY,                         // attachment
        VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL        // layout
    },
    // Read from opaque buffer.
    {
        kAttachment_OPAQUE,                               // attachment
        VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL        // layout
    }
};


// Final pass - back buffer render reference
static const VkAttachmentReference backBufferRenderRef[] =
{
    {
        kAttachment_BACK,                                 // attachment
        VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL        // layout
    }
};


static const VkSubpassDescription subpasses[] =
```

```cpp
{
    // Subpass 1 - depth prepass
    {
        0,                                          // flags
        VK_PIPELINE_BIND_POINT_GRAPHICS,            // pipelineBindPoint
        0,                                          // inputAttachmentCount
        nullptr,                                    // pInputAttachments
        0,                                          // colorAttachmentCount
        nullptr,                                    // pColorAttachments
        nullptr,                                    // pResolveAttachments
        &depthAttachmentReference,                  // pDepthStencilAttachment
        0,                                          // preserveAttachmentCount
        nullptr                                     // pPreserveAttachments
    },
    // Subpass 2 - g-buffer generation
    {
        0,                                          // flags
        VK_PIPELINE_BIND_POINT_GRAPHICS,            // pipelineBindPoint
        0,                                          // inputAttachmentCount
        nullptr,                                    // pInputAttachments
        vkcore::utils::arraysize(gBufferOutputs),   // colorAttachmentCount
        gBufferOutputs,                             // pColorAttachments
        nullptr,                                    // pResolveAttachments
        &depthAttachmentReference,                  // pDepthStencilAttachment
        0,                                          // preserveAttachmentCount
        nullptr                                     // pPreserveAttachments
    },
    // Subpass 3 - lighting
    {
        0,                                          // flags
        VK_PIPELINE_BIND_POINT_GRAPHICS,            // pipelineBindPoint
        vkcore::utils::arraysize(gBufferReadRef),   // inputAttachmentCount
        gBufferReadRef,                             // pInputAttachments
        vkcore::utils::arraysize(opaqueWrite),      // colorAttachmentCount
        opaqueWrite,                                // pColorAttachments
        nullptr,                                    // pResolveAttachments
        nullptr,                                    // pDepthStencilAttachment
        0,                                          // preserveAttachmentCount
        nullptr                                     // pPreserveAttachments
    },
    // Subpass 4 - translucent objects
    {
        0,                                          // flags
        VK_PIPELINE_BIND_POINT_GRAPHICS,            // pipelineBindPoint
        0,                                          // inputAttachmentCount
        nullptr,                                    // pInputAttachments
        vkcore::utils::arraysize(translucentWrite), //
colorAttachmentCount
        translucentWrite,                           // pColorAttachments
        nullptr,                                    // pResolveAttachments
        nullptr,                                    // pDepthStencilAttachment
        0,                                          // preserveAttachmentCount
        nullptr                                     // pPreserveAttachments
```

```
    },
    // Subpass 5 - composite
    {
        0,                                       // flags
        VK_PIPELINE_BIND_POINT_GRAPHICS,     // pipelineBindPoint
        0,                                       // inputAttachmentCount
        nullptr,                                 // pInputAttachments
        vkcore::utils::arraysize(backBufferRenderRef), //
colorAttachmentCount
        backBufferRenderRef,                 // pColorAttachments
        nullptr,                             // pResolveAttachments
        nullptr,                             // pDepthStencilAttachment
        0,                                   // preserveAttachmentCount
        nullptr                              // pPreserveAttachments
    }
};
static const VkSubpassDependency dependencies[] =
{
    // G-buffer pass depends on depth prepass.
    {

        kSubpass_DEPTH,                                         // srcSubpass
        kSubpass_GBUFFER,                                       // dstSubpass
        VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,  // srcStageMask
        VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,          // dstStageMask
        VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,           // srcAccessMask
        VK_ACCESS_SHADER_READ_BIT,                      // dstAccessMask
        VK_DEPENDENCY_BY_REGION_BIT                     // dependencyFlags
    },
    // Lighting pass depends on g-buffer.
    {

        kSubpass_GBUFFER,                                       // srcSubpass
        kSubpass_LIGHTING,                                      // dstSubpass
        VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,  // srcStageMask
        VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,          // dstStageMask
        VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,           // srcAccessMask
        VK_ACCESS_SHADER_READ_BIT,                      // dstAccessMask
        VK_DEPENDENCY_BY_REGION_BIT                     // dependencyFlags
    },
    // Composite pass depends on translucent pass.
    {
        kSubpass_TRANSLUCENTS,                                  // srcSubpass
        kSubpass_COMPOSITE,                                     // dstSubpass
        VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,  // srcStageMask
        VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,          // dstStageMask
        VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,           // srcAccessMask
        VK_ACCESS_SHADER_READ_BIT,                      // dstAccessMask
        VK_DEPENDENCY_BY_REGION_BIT                     // dependencyFlags
    },
    // Composite pass also depends on lighting.
    {
        kSubpass_LIGHTING,                                      // srcSubpass
        kSubpass_COMPOSITE,                                     // dstSubpass
        VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,  // srcStageMask
```

```
            VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,              // dstStageMask
            VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,               // srcAccessMask
            VK_ACCESS_SHADER_READ_BIT,                          // dstAccessMask
            VK_DEPENDENCY_BY_REGION_BIT                         // dependencyFlags
        }
    };

    static const VkRenderPassCreateInfo renderPassCreateInfo =
    {
        VK_STRUCTURE_TYPE_RENDER_PASS_CREATE_INFO, nullptr,
        0,                                                      // flags
        vkcore::utils::arraysize(attachments),                 // attachmentCount
        attachments,                                           // pAttachments
        vkcore::utils::arraysize(subpasses),                   // subpassCount
        subpasses,                                             // pSubpasses
        vkcore::utils::arraysize(dependencies),                // dependencyCount
        dependencies                                           // pDependencies
    };

    result = vkCreateRenderPass(device,
                                &renderPassCreateInfo,
                                nullptr,
                                &m_renderPass);
```

Again, the code in [Listing 13.2](#) is extremely long but is mostly a set of constant data structures. Here, we've added the translucent pass and the composite pass to the list of passes in subpasses[]. The final pass is now the composite pass, so it is the one that references the back buffer in its pColorAttachments array. The result of the lighting pass is now written to the temporary opaque buffer, indexed by kAttachment_OPAQUE.

Although this appears to consume a significant amount of memory, we can note several redeeming points about this configuration:

- There are likely to be at least two or three back buffers, while you can use the same buffer every frame for intermediate results. The additional overhead of one extra buffer is not that large.

- The lighting pass consumes the g-buffer, after which it is no longer needed. You can either write the opaque result back into the g-buffer or mark the attachment as transient and hope that the Vulkan implementation does this for you.

- If you are rendering high dynamic range, then you may want your rendering results to be in a higher-precision format than the back buffer anyway and then perform tone mapping or other processing during the composition pass. In this case, you'll need the intermediate buffers.

Secondary Command Buffers

Secondary command buffers are command buffers that can be *called* from primary command buffers. Although not directly related to multipass rendering, they are used primarily to allow the commands contributing to a render consisting of many subpasses to be built up in multiple command buffers. As you know, a renderpass must begin and end in the same command buffer. That is, the call to **vkCmdEndRenderPass()** must appear in the same command buffer as its corresponding **vkCmdBeginRenderPass()**.

Given this requirement, it's very difficult to render a large amount of the scene in a single renderpass and still build command buffers in parallel. In the ideal case (from an implementation point of view), the entire scene will be rendered in a single large renderpass, with potentially many subpasses. Without secondary command buffers, this would require most, if not all, of a scene to be rendered using a single long command buffer, precluding parallelized command generation.

To create a secondary command buffer, create a command pool and then from it allocate one or more command buffers. In the `VkCommandBufferAllocateInfo` structure passed to **`vkAllocateCommandBuffers()`**, set the `level` field to `VK_COMMAND_BUFFER_LEVEL_SECONDARY`. We then record commands into the command buffer as usual, but with certain restrictions as to which commands can be executed. A table listing which commands may and may not be recorded in secondary command buffers is shown in the Appendix, "Vulkan Functions."

When the secondary command buffer is ready to execute from another primary command buffer, call **`vkCmdExecuteCommands()`**, the prototype of which is

Click here to view code image

```
void vkCmdExecuteCommands (
    VkCommandBuffer                  commandBuffer,
    uint32_t                         commandBufferCount,
    const VkCommandBuffer*           pCommandBuffers);
```

The command buffer from which to call the secondary command buffers is passed in `commandBuffer`. A single call to **`vkCmdExecuteCommands()`** can execute many secondary-level command buffers. The number of command buffers to execute is passed in `commandBufferCount`, and `pCommandBuffers` should point to an array of this many `VkCommandBuffer` handles to the command buffers to execute.

Vulkan command buffers contain a certain amount of state. In particular, the currently bound pipeline, the various dynamic states, and the currently bound descriptor sets are effectively properties of each command buffer. When multiple command buffers are executed back to back, even when sent to the same call to **`vkQueueSubmit()`**, no state is inherited from one to the next. That is, the initial state of each command buffer is undefined, even if the previously executed command buffer left everything as it needs to be for the next.

When you are executing a large, complex command buffer, it's probably fine to begin with an undefined state because the first thing you'll do in the command buffer is set up everything required for the first few drawing commands. That cost is likely to be small relative to the cost of the whole command buffer, even if it's partially redundant with respect to previously executed command buffers.

When a primary command buffer calls a secondary command buffer, and especially when a primary command buffer calls many short secondary command buffers back to back, it can be costly to reset the complete state of the pipeline in each and every secondary command buffer. To compensate for this, some state can be inherited from primary to secondary command buffers. This is done using the `VkCommandBufferInheritanceInfo` structure, which is passed to **`vkBeginCommandBuffer()`**. The definition of this structure is

```
typedef struct VkCommandBufferInheritanceInfo {
    VkStructureType                    sType;
    const void*                        pNext;
    VkRenderPass                       renderPass;
    uint32_t                           subpass;
    VkFramebuffer                      framebuffer;
    VkBool32                           occlusionQueryEnable;
    VkQueryControlFlags                queryFlags;
    VkQueryPipelineStatisticFlags      pipelineStatistics;
} VkCommandBufferInheritanceInfo;
```

The `VkCommandBufferInheritanceInfo` provides a mechanism for your application to tell Vulkan that you know what the state will be when the secondary command buffer is executed. This allows the properties of the command buffer to begin in a well-defined state.

The `sType` field of `VkCommandBufferInheritanceInfo` should be set to `VK_STRUCTURE_TYPE_COMMAND_BUFFER_INHERITANCE_INFO`, and `pNext` should be set to `nullptr`.

The `renderpass` and `subpass` fields specify the renderpass and the subpass of the render that the command buffer will be called inside, respectively. If the framebuffer to which the renderpass will be rendering is known, then it can be specified in the `framebuffer` field. This can sometimes result in better performance when the command buffer is executed. However, if you don't know which framebuffer will be used, then you should set this field to `VK_NULL_HANDLE`.

The `occlusionQueryEnable` field should be set to `VK_TRUE` if the secondary command buffer will be executed while the primary command buffer is executing an occlusion query. This tells Vulkan to keep any counters associated with occlusion queries consistent during execution of the secondary command buffer. If this flag is `VK_FALSE`, then the secondary command buffer should not be executed while occlusion queries are active in the calling command buffer. While the behavior is technically undefined, the most likely outcome is that the results of the occlusion queries are garbage.

You can execute occlusion queries inside secondary command buffers regardless of the value of `occlusionQueryEnable`. You will need to begin and end the query inside the same secondary command buffer if you don't inherit the state from a calling primary.

If occlusion query inheritance is enabled, then the `queryFlags` field contains additional flags that control the behavior of occlusion queries. The only flag defined for use here is `VK_QUERY_CONTROL_PRECISE_BIT`, which, if set, indicates that precise occlusion-query results are needed.

The `pipelineStatistics` field includes flags that tell Vulkan which pipeline statistics are being gathered by calling the primary command buffer. Again, you can gather pipeline statistics during the execution of a secondary command buffer, but if you want the operation of the pipeline invoked by a secondary command buffer to contribute counters accumulated by the primary command buffer, you need to set these flags accurately. The available bits are members of the `VkQueryPipelineStatisticFlagBits` enumeration.

Summary

This chapter delved deeper into the renderpass, a fundamental feature of Vulkan that enables efficient multipass rendering. You saw how to construct a nontrivial renderpass that contains many subpasses and how to build the contents of those subpasses as separate command buffers that can be called from your main command buffer. We discussed some potential optimizations that Vulkan implementations could make to improve the performance of rendering when they are given all the information about what *will* come in the frame. You also saw how many of the functions performed by barriers and clears can be folded into the renderpass, sometimes making them close to free. The renderpass is a powerful feature and, if possible, you should endeavor to make use of it in your applications.

Appendix. Vulkan Functions

This appendix contains a table of the Vulkan command buffer building commands. It provides a quick reference to what can and cannot be used inside and outside a renderpass, as well as to what is legal inside a primary or secondary command buffer.

Function	Renderpass		Command Buffer Level	
	Inside	Outside	Primary	Secondary
vkCmdBeginQuery	✓	✓	✓	✓
vkCmdBeginRenderPass		✓	✓	
vkCmdBindDescriptorSets	✓	✓	✓	✓
vkCmdBindIndexBuffer	✓	✓	✓	✓
vkCmdBindPipeline	✓	✓	✓	✓
vkCmdBindVertexBuffers	✓	✓	✓	✓
vkCmdBlitImage		✓	✓	✓
vkCmdClearAttachments	✓		✓	✓
vkCmdClearColorImage		✓	✓	✓
vkCmdClearDepthStencilImage		✓	✓	✓
vkCmdCopyBuffer		✓	✓	✓
vkCmdCopyBufferToImage		✓	✓	✓
vkCmdCopyImage		✓	✓	✓
vkCmdCopyImageToBuffer		✓	✓	✓
vkCmdCopyQueryPoolResults		✓	✓	✓
vkCmdDispatch		✓	✓	✓
vkCmdDispatchIndirect		✓	✓	✓
vkCmdDraw	✓		✓	✓
vkCmdDrawIndexed	✓		✓	✓
vkCmdDrawIndexedIndirect	✓		✓	✓
vkCmdDrawIndirect	✓		✓	✓
vkCmdEndQuery	✓	✓	✓	✓
vkCmdEndRenderPass	✓		✓	
vkCmdExecuteCommands	✓	✓	✓	

vkCmdFillBuffer		✓	✓	✓
vkCmdNextSubpass	✓		✓	
vkCmdPipelineBarrier	✓	✓	✓	✓
vkCmdPushConstants	✓	✓	✓	✓
vkCmdResetEvent		✓	✓	✓
vkCmdResetQueryPool		✓	✓	✓
vkCmdResolveImage		✓	✓	✓
vkCmdSetBlendConstants	✓	✓	✓	✓
vkCmdSetDepthBias	✓	✓	✓	✓
vkCmdSetDepthBounds	✓	✓	✓	✓
vkCmdSetEvent		✓	✓	✓
vkCmdSetLineWidth	✓	✓	✓	✓
vkCmdSetScissor	✓	✓	✓	✓
vkCmdSetStencilCompareMask	✓	✓	✓	✓
vkCmdSetStencilReference	✓	✓	✓	✓
vkCmdSetStencilWriteMask	✓	✓	✓	✓
vkCmdSetViewport	✓	✓	✓	✓
vkCmdUpdateBuffer		✓	✓	✓
vkCmdWaitEvents	✓	✓	✓	✓
vkCmdWriteTimestamp	✓	✓	✓	✓

Command Buffer Building Functions

Glossary

adjacency primitive One of the primitive topologies that includes, for each primitive, additional vertex data representing adjacent primitives in the original geometry. Examples include triangles and lines with adjacency.

aliasing Technically, the loss of signal information in an image reproduced at some finite resolution. It is most often characterized by the appearance of sharp, jagged edges along points, lines, or polygons due to the nature of having a limited number of fixed-size pixels.

alpha A fourth color value added to provide a degree of transparency to the color of an object. An alpha value of 0.0 means complete transparency; 1.0 denotes no transparency (opaque).

ambient light Light in a scene that doesn't come from any specific point source or direction. Ambient light illuminates all surfaces evenly and on all sides.

antialiasing A rendering method used to smooth lines, curves, and polygon edges. This technique averages the color of pixels adjacent to the line. It has the visual effect of softening the transition from the pixels on the line and those adjacent to the line, thus providing a smoother appearance.

Apple An apple is a piece of fruit. Fruit does not support Vulkan.

ARB Acronym for the Architecture Review Board, the committee body consisting of 3D graphics hardware vendors, previously charged with maintaining the OpenGL Specification. This function has since been assumed by the Khronos Group. See *Khronos Group*.

aspect When applied to an image, a logical part of that image, such as the depth or stencil component of a combined depth-stencil image.

aspect ratio The ratio of the width of a window to the height of the window—specifically, the width of the window in pixels divided by the height of the window in pixels.

associativity A sequence of operations in which changing the order of the operations (but not the order of the arguments) does not affect the result. For example, addition is associative because $a + (b + c) = (a + b) + c$.

atomic operation A sequence of operations that must be indivisible for correct operation. The term usually refers to a read-modify-write sequence on a single memory location.

attachment An image associated with a renderpass that can be used as the input or output of one or more of its subpasses.

barrier A point in a computer program that serves as a marker across which operations may not be reordered. Between barriers, certain operations may be exchanged if their movement does not logically change the operation of the program. Barriers operate on resources or memory and can also be used to change the layout of images.

Bézier curve A curve whose shape is defined by control points near the curve rather than by the precise set of points that define the curve itself.

bitplane An array of bits mapped directly to screen pixels.

blending The process of merging a new color value into an existing color attachment using an equation and parameters that are configured as part of a graphics pipeline.

blit Short for *bl*ock *i*mage *t*ransfer—an operation that copies image data from one place to another, potentially processing it further as it is copied. In Vulkan, a blit is used to scale image data and to perform basic operations such as format conversion.

branch prediction An optimization strategy used in processor design whereby the processor tries to guess (or predict) the outcome of some conditional code and start executing the more likely branch before it is certain that it is required. If the processor is right, it gets ahead by a few instructions. If the processor is wrong, it needs to throw away the work and start again with the other branch.

buffer An area of memory used to store image information. This information can be color, depth, or blending information. The red, green, blue, and alpha buffers are often collectively referred to as the color buffers.

Cartesian A coordinate system based on three directional axes placed at a 90° orientation to one another. These coordinates are labeled *x*, *y*, and *z*.

clip coordinates The 2D geometric coordinates that result from the model-view and projection transformation.

clip distance A distance value assigned by a shader that is used by fixed-function clipping to allow primitives to be clipped against an arbitrary set of planes before rasterization.

clipping The elimination of a portion of a single primitive or group of primitives. The points that would be rendered outside the clipping region or volume are not drawn. The clipping volume is generally specified by the projection matrix. Clipped primitives are reconstructed such that the edges of the primitive do not lie outside the clipping region.

command buffer A list of commands that can be executed by the device.

commutative An operation in which changing the order of its operands does not change its result. For example, addition is commutative, whereas subtraction is not.

compute pipeline A Vulkan pipeline object consisting of a *compute shader* and related state that is used to execute computational work on a Vulkan device.

compute shader A shader that executes a work item per invocation as part of a local work group, a number of which may be grouped into a global work group.

concave A reference to the shape of a polygon. A polygon is said to be concave if a straight line through it will enter and subsequently exit the polygon more than once.

contention A term used to describe the condition in which two or more threads of execution attempt to use a single shared resource.

convex A reference to the shape of a polygon. A convex polygon has no indentations, and no straight line can be drawn through the polygon that intersects it more than twice (once entering, once leaving).

CRT Cathode ray tube.

cull distance A value applied to a vertex within a primitive that will cause the entire primitive to be discarded if *any* of its vertices are assigned a negative value.

culling The elimination of graphics primitives that would not be seen if rendered. Back-face culling eliminates the front or back face of a primitive so that the face isn't drawn. Frustum culling eliminates whole objects that would fall outside the viewing frustum.

depth fighting A visual artifact caused by primitives with very close depth values being rendered on top of one another and producing inconsistent depth-test results.

depth test A test performed between a depth value computed for a fragment and a value stored in a depth buffer.

descriptor A data structure containing an implementation-specific description of a resource such as a buffer or an image.

destination color The stored color at a particular location in the color buffer. This terminology is usually used when describing blending operations to distinguish between the color already present in the color buffer and the color coming into the color buffer (source color).

device memory Memory that is accessible to a Vulkan device. This may be dedicated memory that's physically attached to the device, or some or all of the host's memory that is accessible to the device.

dispatch A command that begins the execution of compute shaders.

displacement mapping The act of taking a flat surface and displacing it along a normal by an amount determined by a texture or other source of data.

dithering A method used to simulate a wider range of color depth by placing different-colored pixels together in patterns that create the illusion of shading between the two colors.

double buffering A drawing technique used to present stable images. The image to be displayed is assembled in memory and then placed on the screen in a single update operation rather than built primitive by primitive onscreen. Double buffering is a much faster and smoother update operation and can produce animations.

event A synchronization primitive used within a Vulkan command buffer to synchronize the execution of different parts of the pipeline.

extruding The process of taking a 2D image or shape and adding a third dimension uniformly across the surface. This process can transform 2D fonts into 3D lettering.

eye coordinates The coordinate system based on the position of the viewer. The viewer's position is placed along the positive z axis, looking down the negative z axis.

fence An object used to synchronize execution on the host with completion of command buffers executed by the device.

FMA Acronym for fused multiply add, an operation commonly implemented in a single piece of hardware that multiplies two numbers together and adds a third, with the intermediate result generally computed at higher precision than a stand-alone multiplication or addition operation.

fragment A single piece of data that may eventually contribute to the color of a pixel in an image.

fragment shader A shader that executes once per fragment and generally computes the final color of that fragment.

framebuffer An object containing references to images that will be used as the attachments during rendering with graphics pipelines.

frustum A pyramid-shaped viewing volume that creates a perspective view. (Near objects are large; far objects are small.)

gamma correction The process of transforming a linear value through a gamma curve to produce nonlinear output. Gamma (γ) can be greater or less than one to move from linear to gamma space and back again, respectively.

garbage Uninitialized data that is read and consumed by a computer program, often resulting in corruption, crashes, or other undesired behavior.

geometry shader A shader that executes once per primitive, having access to all vertices making up that primitive.

gimbal lock A state in which a sequence of rotations can essentially become stuck on a single axis. This occurs when one of the rotations early in the sequence rotates one Cartesian axis onto another. After this, rotation around either of the axes results in the same rotation, making it impossible to escape from the locked position.

GLSL Acronym for OpenGL Shading Language, a high-level C-like shading language.

GPU Acronym for graphics processing unit—a specialized processor that does most of the heavyweight lifting for Vulkan.

graphics pipeline A Vulkan pipeline object constructed using one or more graphics shaders (vertex, tessellation, geometry, and fragment) and related state that is capable of processing drawing commands and rendering graphics primitives on a Vulkan device. See also *pipeline* and *compute pipeline*.

handle An opaque variable used to refer to another object. In Vulkan, all handles are 64-bit integers.

hazard In reference to memory operations, a situation in which undefined order of transactions in memory may lead to undefined or undesired results. Typical examples include read-after-write (RAW), write-after-write (WAW), and write-after-read (WAR) hazards.

helper invocation An invocation of a fragment shader that does not represent a fragment that lies inside a primitive, but is executed in order to produce derivatives or other information required to correctly process shader invocations for fragments that *do* lie inside the primitive.

host memory Memory that is dedicated for use by the host and inaccessible to the device. This type of memory is generally used to store data for use by the CPU.

implementation A software- or hardware-based device that performs Vulkan rendering operations.

index buffer A buffer bound for use as the source of indices for an indexed draw.

indexed draw A drawing command that chooses vertices by using indices read from a buffer rather than simply counting them monotonically.

indirect draw A drawing command that sources its parameters from device memory accessed through a buffer object rather than reading them directly from a command buffer.

instance Applied to drawing commands, the process of drawing the same set of data many times, potentially with different parameters applied.

instancing See *instance*.

invocation A single execution of a shader. The term is most commonly used to describe compute shaders but is applicable to any shader stage.

Khronos Group The industry consortium that manages the maintenance and promotion of the Vulkan specification.

linked list A list of data elements joined by pointers or indices stored or associated with each element in the list.

literal A value, not a variable name. A literal is a specific string or numeric constant embedded directly in source code.

mapping See *memory mapping*.

matrix A 2D array of numbers. Matrices can be operated on mathematically and are used to perform coordinate transformations.

memory mapping Obtaining a pointer to a region of device memory that is usable by the host.

mipmapping A technique that uses multiple levels of detail for a texture. This technique selects among the different sizes of an image available or possibly combines the two nearest matches to produce the final fragments used for texturing.

model-view matrix The matrix that transforms position vectors from model (or object) space to view (or eye) space

multisample An image in which each pixel has more than one sample. Multisample images are typically used in antialiasing. See *antialiasing*.

multisampling The act of rendering to a multisample image.

normal A directional vector that points perpendicularly to a plane or surface. When used, normals must be specified for each vertex in a primitive.

normalize The reduction of a normal to a unit normal. A unit normal is a vector that has a length of exactly 1.0.

normalized device coordinate The coordinate space produced by taking a homogeneous position and dividing it through by its own w component.

occlusion query A graphics operation whereby visible (or more accurately potentially visible) pixels are counted and the count is returned to the application.

one hot A method of encoding a number or state using a single bit of a binary number. For example, a state machine or enumeration with 32 states would require only 5 bits using regular encoding but an entire 32-bit value to encode as one hot. In hardware, one hot is often simpler to decode, and multiple one-hot flags can be included in a single integer value.

orthographic A drawing mode (also called *parallel projection*) in which no perspective or foreshortening takes place. The lengths and dimensions of all primitives are undistorted regardless of orientation or distance from the viewer.

out-of-order execution The ability of a processor to determine interinstruction dependencies and start executing those instructions whose inputs are ready *before* other instructions that may have preceded them in program order.

overloading In computer languages, the practice of creating two or more functions that share a name but have different function signatures.

perspective A drawing mode in which objects farther from the viewer appear smaller than nearby objects.

perspective divide The transformation applied to homogeneous vectors to move them from clip space into normalized device coordinates by dividing them through by their w components.

pipeline An object representing the state of a large part of the device that is used to execute work.

pixel Condensed from the words *picture element*, the smallest visual division available on the computer screen. Pixels are arranged in rows and columns and are individually set to the appropriate color to render any given image.

pixmap A 2D array of color values that comprise a color image. Pixmaps are so called because each picture element corresponds to a pixel onscreen.

polygon A 2D shape drawn with any number of sides (but must have at least three sides).

presentation The act of taking an image and displaying it to the user.

primitive A group of one or more vertices formed into a geometric shape by Vulkan, such as a line, point, or triangle. All objects and scenes are composed of various combinations of primitives.

primitive topology The arrangement of vertices into a single primitive or group of primitives. This arrangement includes points, lines, strips, and triangles, as well as adjoined primives such as line strips, triangle strips, and adjacency primitives.

projection The transformation of lines, points, and polygons from eye coordinates to clipping coordinates onscreen.

push constant A uniform constant accessible to a shader that can be updated directly from a command buffer by calling `vkCmdPushConstants()` without synchronization, avoiding the need for memory backing or barriers associated with the update of small uniform buffers. In some implementations, these constants may be accelerated in hardware.

quadrilateral A polygon with exactly four sides.

query object An object used to retrieve statistical data about the operation of your application as it runs on the device.

queue A construct within a device that can execute work, potentially in parallel with work submitted to other queues on the same device.

queue family Queues with identical properties grouped into families within a device. Queues within a family are considered to be compatible, and command buffers built for a particular family may be submitted to any member of that family.

race condition A state in which multiple parallel processes such as threads in a program or invocations of a shader attempt to communicate or otherwise depend on one another in some way, but no ensurance of ordering is performed.

rasterization The process of converting projected primitives and bitmaps into pixel fragments in the framebuffer.

render The conversion of primitives in object coordinates to an image in the framebuffer. The rendering pipeline is the process by which Vulkan commands and statements become pixels onscreen.

renderpass An object representing one or more subpasses over a common set of framebuffer attachments.

sample A component of a pixel or texel that, when resolved, will contribute to the final output color. Sampling an image is the process of gathering one or more samples from it and resolving that group of samples to produce a final color.

scintillation A sparkling or flashing effect produced on objects when a nonmipmapped texture map is applied to a polygon that is significantly smaller than the size of the texture being applied.

scissor A fragment ownership test that rejects fragments that lie outside a window-aligned rectangle.

semaphore A synchronization primitive that can be used to synchronize work executed by different queues on a single device.

shader A small program that is executed by the graphics hardware, often in parallel, to operate on individual vertices or pixels.

single static assignment A form of writing a program in which any variable, often represented as a register for a virtual machine, is written to or assigned only once. This makes data-flow

resolution simple and facilitates a number of common optimizations in compilers and other code generation applications.

source color The color of the incoming fragment, as opposed to the color already present in the color buffer (destination color). This term is usually used to describe how the source and destination colors are combined during a blending operation.

specification The design document that specifies Vulkan operation and fully describes how an implementation must work.

spline A general term used to describe any curve created by placing control points near the curve that have a pulling effect on the curve's shape. This effect is similar to the reaction of a piece of flexible material when pressure is applied at various points along its length.

sRGB Nonlinear encoding used for image data using a gamma curve that matches phosphor-based CRTs of the late 1990s.

SSA See *single static assignment*.

stipple A binary bit pattern used to mask out pixel generation in the framebuffer. This is similar to a monochrome bitmap, but 1D patterns are used for lines and 2D patterns are used for polygons.

submission See *submit*.

submit When applied to work executed by devices, the act of sending the work in the form of command buffers to one of the device's queues.

subpass A single pass of rendering contained within a *renderpass* object.

super scalar A processor architecture that is capable of executing two or more independent instructions at the same time on multiple processor pipelines, which may or may not have the same capabilities.

supersampling The process of computing multiple color, depth, and stencil values for every sample within a multisample image.

surface An object representing an output for presentation. This type of object is generally platform-specific and is provided using the `VK_KHR_surface` extension.

swap chain An object representing a sequence of images that can be presented into a surface for viewing by the user.

tessellation The process of breaking down a complex polygon or analytic surface into a mesh of convex polygons. This process can also be applied to separate a complex curve into a series of less-complex lines.

tessellation control shader A shader that runs before fixed-function tessellation occurs. The shader executes once per control point in a patch primitive, producing tessellation factors and a new set of control points as an output primitive.

tessellation evaluation shader A shader that runs after fixed-function tessellation occurs. The shader executes once per vertex generated by the tessellator.

tessellation shader A term used to describe either a tessellation control shader or a tessellation evaluation shader.

texel A texture element. A texel represents a color from a texture that is applied to a pixel fragment in the framebuffer.

texture An image pattern of colors applied to the surface of a primitive.

texture mapping The process of applying a texture image to a surface. The surface does not have to be planar (flat). Texture mapping is often used to wrap an image around a curved object or to produce a patterned surface, such as wood or marble.

transformation The manipulation of a coordinate system. This can include rotation, translation, scaling (both uniform and nonuniform), and perspective division.

translucence A degree of transparency of an object. In Vulkan, this is represented by an alpha value ranging from 1.0 (opaque) to 0.0 (transparent).

unified memory architecture A memory architecture in which multiple devices in the system, such as a Vulkan device and the host processor, have equal access to a common pool of memory rather than accessing physically segmented memory.

vector A directional quantity usually represented by x, y, and z components.

vertex A single point in space. Except when used for point and line primitives, the term also defines the point at which two edges of a polygon meet.

vertex buffer A buffer being used as the source of data for vertex shader inputs.

vertex shader A shader that executes once per incoming vertex.

view Applied to a resource, another object that references the same underlying data in a different manner from the object of which it is a view, allowing that data to be reinterpreted or sectioned as required by an application.

viewing volume The area in 3D space that can be viewed in the window. Objects and points outside the viewing volume are clipped (cannot be seen).

viewport The area within a window that is used to display a Vulkan image. Usually, this area encompasses the entire client area. Stretched viewports can produce enlarged or shrunken output within the physical window.

wireframe The representation of a solid object by a mesh of lines rather than solid shaded polygons. Wireframe models are usually rendered faster and can be used to view both the front and back of an object at the same time.

Index

Enumerations

Code Snippets

```
VkResult vkCreateInstance (
    const VkInstanceCreateInfo*             pCreateInfo,
    const VkAllocationCallbacks*            pAllocator,
    VkInstance*                             pInstance);
```

```c
typedef struct VkInstanceCreateInfo {
    VkStructureType             sType;
    const void*                 pNext;
    VkInstanceCreateFlags       flags;
    const VkApplicationInfo*    pApplicationInfo;
    uint32_t                    enabledLayerCount;
    const char* const*          ppEnabledLayerNames;
    uint32_t                    enabledExtensionCount;
    const char* const*          ppEnabledExtensionNames;
} VkInstanceCreateInfo;
```

```c
typedef struct VkApplicationInfo {
    VkStructureType      sType;
    const void*          pNext;
    const char*          pApplicationName;
    uint32_t             applicationVersion;
    const char*          pEngineName;
    uint32_t             engineVersion;
    uint32_t             apiVersion;
} VkApplicationInfo;
```

```c
VkResult vkEnumeratePhysicalDevices (
    VkInstance                              instance,
    uint32_t*                               pPhysicalDeviceCount,
    VkPhysicalDevice*                       pPhysicalDevices);
```

```cpp
VkResult vkapp::init()
{
    VkResult result = VK_SUCCESS;
    VkApplicationInfo appInfo = { };
    VkInstanceCreateInfo instanceCreateInfo = { };

    // A generic application info structure
    appInfo.sType = VK_STRUCTURE_TYPE_APPLICATION_INFO;
    appInfo.pApplicationName = "Application";
    appInfo.applicationVersion = 1;
    appInfo.apiVersion = VK_MAKE_VERSION(1, 0, 0);

    // Create the instance.
    instanceCreateInfo.sType = VK_STRUCTURE_TYPE_INSTANCE_CREATE_INFO;
    instanceCreateInfo.pApplicationInfo = &appInfo;

    result = vkCreateInstance(&instanceCreateInfo, nullptr, &m_instance);

    if (result == VK_SUCCESS)
    {
        // First figure out how many devices are in the system.
        uint32_t physicalDeviceCount = 0;
        vkEnumeratePhysicalDevices(m_instance, &physicalDeviceCount, nullptr);

        if (result == VK_SUCCESS)
        {
            // Size the device array appropriately and get the physical
            // device handles.
            m_physicalDevices.resize(physicalDeviceCount);
            vkEnumeratePhysicalDevices(m_instance,
                                       &physicalDeviceCount,
                                       &m_physicalDevices[0]);
        }
    }

    return result;
}
```

```c
void vkGetPhysicalDeviceProperties (
    VkPhysicalDevice                    physicalDevice,
    VkPhysicalDeviceProperties*         pProperties);
```

```
typedef struct VkPhysicalDeviceProperties {
    uint32_t                            apiVersion;
    uint32_t                            driverVersion;
    uint32_t                            vendorID;
    uint32_t                            deviceID;
    VkPhysicalDeviceType                deviceType;
    char                                deviceName
                                            [VK_MAX_PHYSICAL_DEVICE_NAME_SIZE];
    uint8_t                             pipelineCacheUUID[VK_UUID_SIZE];
    VkPhysicalDeviceLimits              limits;
    VkPhysicalDeviceSparseProperties    sparseProperties;
} VkPhysicalDeviceProperties;
```

```
void vkGetPhysicalDeviceFeatures (
    VkPhysicalDevice                        physicalDevice,
    VkPhysicalDeviceFeatures*               pFeatures);
```

```
void vkGetPhysicalDeviceMemoryProperties (
    VkPhysicalDevice                    physicalDevice,
    VkPhysicalDeviceMemoryProperties*   pMemoryProperties);
```

```c
typedef struct VkPhysicalDeviceMemoryProperties {
    uint32_t            memoryTypeCount;
    VkMemoryType        memoryTypes[VK_MAX_MEMORY_TYPES];
    uint32_t            memoryHeapCount;
    VkMemoryHeap        memoryHeaps[VK_MAX_MEMORY_HEAPS];
} VkPhysicalDeviceMemoryProperties;
```

```c
typedef struct VkMemoryType {
    VkMemoryPropertyFlags      propertyFlags;
    uint32_t                   heapIndex;
} VkMemoryType;
```

```c
typedef struct VkMemoryHeap {
    VkDeviceSize            size;
    VkMemoryHeapFlags       flags;
} VkMemoryHeap;
```

```c
void vkGetPhysicalDeviceQueueFamilyProperties (
    VkPhysicalDevice                        physicalDevice,
    uint32_t*                               pQueueFamilyPropertyCount,
    VkQueueFamilyProperties*                pQueueFamilyProperties);
```

```c
typedef struct VkQueueFamilyProperties {
    VkQueueFlags    queueFlags;
    uint32_t        queueCount;
    uint32_t        timestampValidBits;
    VkExtent3D      minImageTransferGranularity;
} VkQueueFamilyProperties;
```

```cpp
uint32_t queueFamilyPropertyCount;
std::vector<VkQueueFamilyProperties> queueFamilyProperties;
VkPhysicalDeviceMemoryProperties physicalDeviceMemoryProperties;

// Get the memory properties of the physical device.
vkGetPhysicalDeviceMemoryProperties(m_physicalDevices[deviceIndex],
                                    &physicalDeviceMemoryProperties);

// First determine the number of queue families supported by the physical
// device.
vkGetPhysicalDeviceQueueFamilyProperties(m_physicalDevices[0],
                                         &queueFamilyPropertyCount,
                                         nullptr);

// Allocate enough space for the queue property structures.
queueFamilyProperties.resize(queueFamilyPropertyCount);

// Now query the actual properties of the queue families.
vkGetPhysicalDeviceQueueFamilyProperties(m_physicalDevices[0],
                                         &queueFamilyPropertyCount,
                                         queueFamilyProperties.data());
```

```
VkResult vkCreateDevice (
    VkPhysicalDevice                    physicalDevice,
    const VkDeviceCreateInfo*           pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkDevice*                           pDevice);
```

```c
typedef struct VkDeviceCreateInfo {
    VkStructureType                    sType;
    const void*                        pNext;
    VkDeviceCreateFlags                flags;
    uint32_t                           queueCreateInfoCount;
    const VkDeviceQueueCreateInfo*     pQueueCreateInfos;
    uint32_t                           enabledLayerCount;
    const char* const*                 ppEnabledLayerNames;
    uint32_t                           enabledExtensionCount;
    const char* const*                 ppEnabledExtensionNames;
    const VkPhysicalDeviceFeatures*    pEnabledFeatures;
} VkDeviceCreateInfo;
```

```c
typedef struct VkDeviceQueueCreateInfo {
    VkStructureType                 sType;
    const void*                     pNext;
    VkDeviceQueueCreateFlags        flags;
    uint32_t                        queueFamilyIndex;
    uint32_t                        queueCount;
    const float*                    pQueuePriorities;
} VkDeviceQueueCreateInfo;
```

```cpp
VkResult result;
VkPhysicalDeviceFeatures supportedFeatures;
VkPhysicalDeviceFeatures requiredFeatures = {};

vkGetPhysicalDeviceFeatures(m_physicalDevices[0],
                            &supportedFeatures);

requiredFeatures.multiDrawIndirect      = supportedFeatures.multiDrawIndirect;
requiredFeatures.tessellationShader     = VK_TRUE;
requiredFeatures.geometryShader         = VK_TRUE;

const VkDeviceQueueCreateInfo deviceQueueCreateInfo =
{
    VK_STRUCTURE_TYPE_DEVICE_QUEUE_CREATE_INFO,     // sType
    nullptr,                                        // pNext
    0,                                              // flags
    0,                                              // queueFamilyIndex
    1,                                              // queueCount
    nullptr                                         // pQueuePriorities
};

const VkDeviceCreateInfo deviceCreateInfo =
{
    VK_STRUCTURE_TYPE_DEVICE_CREATE_INFO,           // sType
    nullptr,                                        // pNext

    0,                                              // flags
    1,                                              // queueCreateInfoCount
    &deviceQueueCreateInfo,                         // pQueueCreateInfos
    0,                                              // enabledLayerCount
    nullptr,                                        // ppEnabledLayerNames
    0,                                              // enabledExtensionCount
    nullptr,                                        // ppEnabledExtensionNames
    &requiredFeatures                               // pEnabledFeatures
};

result = vkCreateDevice(m_physicalDevices[0],
                        &deviceCreateInfo,
                        nullptr,
                        &m_logicalDevice);
```

```
VkResult vkEnumerateInstanceLayerProperties (
    uint32_t*                           pPropertyCount,
    VkLayerProperties*                  pProperties);
```

```
typedef struct VkLayerProperties {
    char        layerName[VK_MAX_EXTENSION_NAME_SIZE];
    uint32_t    specVersion;
    uint32_t    implementationVersion;
    char        description[VK_MAX_DESCRIPTION_SIZE];
} VkLayerProperties;
```

```cpp
uint32_t numInstanceLayers = 0;
std::vector<VkLayerProperties> instanceLayerProperties;

// Query the instance layers.
vkEnumerateInstanceLayerProperties(&numInstanceExtensions,
                                   nullptr);

// If there are any layers, query their properties.
if (numInstanceLayers != 0)
{
    instanceLayerProperties.resize(numInstanceLayers);
    vkEnumerateInstanceLayerProperties(nullptr,
                                       &numInstanceLayers,
                                       instanceLayerProperties.data());
}
```

```
VkResult vkEnumerateDeviceLayerProperties (
    VkPhysicalDevice                    physicalDevice,
    uint32_t*                           pPropertyCount,
    VkLayerProperties*                  pProperties);
```

```
VkResult vkEnumerateInstanceExtensionProperties (
    const char*                         pLayerName,
    uint32_t*                           pPropertyCount,
    VkExtensionProperties*              pProperties);
```

```cpp
uint32_t numInstanceExtensions = 0;
std::vector<VkExtensionProperties> instanceExtensionProperties;

// Query the instance extensions.
vkEnumerateInstanceExtensionProperties(nullptr,
                                       &numInstanceExtensions,
                                       nullptr);

// If there are any extensions, query their properties.
if (numInstanceExtensions != 0)
{
    instanceExtensionProperties.resize(numInstanceExtensions);
    vkEnumerateInstanceExtensionProperties(nullptr,
                                           &numInstanceExtensions,
                                           instanceExtensionProperties.data());
}
```

```
typedef struct VkExtensionProperties {
    char        extensionName[VK_MAX_EXTENSION_NAME_SIZE];
    uint32_t    specVersion;
} VkExtensionProperties;
```

```
VkResult vkEnumerateDeviceExtensionProperties (
    VkPhysicalDevice                        physicalDevice,
    const char*                             pLayerName,
    uint32_t*                               pPropertyCount,
    VkExtensionProperties*                  pProperties);
```

```
PFN_vkVoidFunction vkGetInstanceProcAddr (
    VkInstance                              instance,
    const char*                             pName);
```

```
VKAPI_ATTR void VKAPI_CALL vkVoidFunction(void);
```

```
PFN_vkVoidFunction vkGetDeviceProcAddr (
    VkDevice                              device,
    const char*                           pName);
```

```
VkResult vkDeviceWaitIdle (
    VkDevice                                    device);
```

```c
void vkDestroyDevice (
    VkDevice                                    device,
    const VkAllocationCallbacks*                pAllocator);
```

```
void vkDestroyInstance (
    VkInstance                          instance,
    const VkAllocationCallbacks*        pAllocator);
```

```
VkResult vkCreateInstance (
    const VkInstanceCreateInfo*          pCreateInfo,
    const VkAllocationCallbacks*         pAllocator,
    VkInstance*                          pInstance);
```

```c
typedef struct VkAllocationCallbacks {
    void*                                    pUserData;
    PFN_vkAllocationFunction                 pfnAllocation;
    PFN_vkReallocationFunction               pfnReallocation;
    PFN_vkFreeFunction                       pfnFree;
    PFN_vkInternalAllocationNotification     pfnInternalAllocation;
    PFN_vkInternalFreeNotification           pfnInternalFree;
} VkAllocationCallbacks;
```

```
void* VKAPI_CALL Allocation(
    void*                                       pUserData,
    size_t                                      size,
    size_t                                      alignment,
    VkSystemAllocationScope                     allocationScope);

void* VKAPI_CALL Reallocation(
    void*                                       pUserData,
    void*                                       pOriginal,
    size_t                                      size,
    size_t                                      alignment,
    VkSystemAllocationScope                     allocationScope);

void VKAPI_CALL Free(
    void*                                       pUserData,
    void*                                       pMemory);
```

```c
void VKAPI_CALL InternalAllocationNotification(
    void*                               pUserData,
    size_t                              size,
    VkInternalAllocationType            allocationType,
    VkSystemAllocationScope             allocationScope);

void VKAPI_CALL InternalFreeNotification(
    void*                               pUserData,
    size_t                              size,
    VkInternalAllocationType            allocationType,
    VkSystemAllocationScope             allocationScope);
```

```cpp
class allocator
{
public:
    // Operator that allows an instance of this class to be used as a
    // VkAllocationCallbacks structure
    inline operator VkAllocationCallbacks() const
    {
        VkAllocationCallbacks result;

        result.pUserData = (void*)this;
        result.pfnAllocation = &Allocation;
        result.pfnReallocation = &Reallocation;
        result.pfnFree = &Free;
        result.pfnInternalAllocation = nullptr;
        result.pfnInternalFree = nullptr;

        return result;
    };

private:
    // Declare the allocator callbacks as static member functions.
    static void* VKAPI_CALL Allocation(
        void*                           pUserData,
        size_t                          size,
        size_t                          alignment,
        VkSystemAllocationScope         allocationScope);

    static void* VKAPI_CALL Reallocation(
        void*                           pUserData,
        void*                           pOriginal,
        size_t                          size,
        size_t                          alignment,
        VkSystemAllocationScope         allocationScope);

    static void VKAPI_CALL Free(
        void*                           pUserData,
        void*                           pMemory);

    // Now declare the nonstatic member functions that will actually perform
    // the allocations.
    void* Allocation(
        size_t                          size,
        size_t                          alignment,
        VkSystemAllocationScope         allocationScope);

    void* Reallocation(
        void*                           pOriginal,
        size_t                          size,
        size_t                          alignment,
        VkSystemAllocationScope         allocationScope);

    void Free(
        void*                           pMemory);
};
```

```cpp
void* allocator::Allocation(
    size_t                                          size,
    size_t                                          alignment,
    VkSystemAllocationScope                         allocationScope)
{
    return aligned_malloc(size, alignment);
}

void* VKAPI_CALL allocator::Allocation(
    void*                                           pUserData,
    size_t                                          size,
    size_t                                          alignment,
    VkSystemAllocationScope                         allocationScope)
{
    return static_cast<allocator*>(pUserData)->Allocation(size,
                                                          alignment,
                                                          allocationScope);
}

void* allocator::Reallocation(
    void*                                           pOriginal,
    size_t                                          size,
    size_t                                          alignment,
    VkSystemAllocationScope                         allocationScope)
{
    return aligned_realloc(pOriginal, size, alignment);
}

void* VKAPI_CALL allocator::Reallocation(
    void*                                           pUserData,
    void*                                           pOriginal,
    size_t                                          size,
    size_t                                          alignment,
    VkSystemAllocationScope                         allocationScope)
{

    return static_cast<allocator*>(pUserData)->Reallocation(pOriginal,
                                                            size,
                                                            alignment,
                                                            allocationScope);
}

void allocator::Free(
    void*                                           pMemory)
{
    aligned_free(pMemory);
}

void VKAPI_CALL allocator::Free(
    void*                                           pUserData,
    void*                                           pMemory)
{
    return static_cast<allocator*>(pUserData)->Free(pMemory);
}
```

```
VkResult vkCreateBuffer (
    VkDevice                            device,
    const VkBufferCreateInfo*           pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkBuffer*                           pBuffer);
```

```c
typedef struct VkBufferCreateInfo {
    VkStructureType         sType;
    const void*             pNext;
    VkBufferCreateFlags     flags;
    VkDeviceSize            size;
    VkBufferUsageFlags      usage;
    VkSharingMode           sharingMode;
    uint32_t                queueFamilyIndexCount;
    const uint32_t*         pQueueFamilyIndices;
} VkBufferCreateInfo;
```

```cpp
static const VkBufferCreateInfo bufferCreateInfo =
{
    VK_STRUCTURE_TYPE_BUFFER_CREATE_INFO, nullptr,
    0,
    1024 * 1024,
    VK_BUFFER_USAGE_TRANSFER_SRC_BIT | VK_BUFFER_USAGE_TRANSFER_DST_BIT,
    VK_SHARING_MODE_EXCLUSIVE,
    0, nullptr
};

VkBuffer buffer = VK_NULL_HANDLE;

vkCreateBuffer(device, &bufferCreateInfo, &buffer);
```

```
void vkGetPhysicalDeviceFormatProperties (
    VkPhysicalDevice                        physicalDevice,
    VkFormat                                format,
    VkFormatProperties*                     pFormatProperties);
```

```
typedef struct VkFormatProperties {
    VkFormatFeatureFlags    linearTilingFeatures;
    VkFormatFeatureFlags    optimalTilingFeatures;
    VkFormatFeatureFlags    bufferFeatures;
} VkFormatProperties;
```

```
VkResult vkGetPhysicalDeviceImageFormatProperties (
    VkPhysicalDevice                        physicalDevice,
    VkFormat                                format,
    VkImageType                             type,
    VkImageTiling                           tiling,
    VkImageUsageFlags                       usage,
    VkImageCreateFlags                      flags,
    VkImageFormatProperties*                pImageFormatProperties);
```

```
typedef struct VkImageFormatProperties {
    VkExtent3D              maxExtent;
    uint32_t               maxMipLevels;
    uint32_t               maxArrayLayers;
    VkSampleCountFlags      sampleCounts;
    VkDeviceSize            maxResourceSize;
} VkImageFormatProperties;
```

```c
typedef struct VkExtent3D {
    uint32_t    width;
    uint32_t    height;
    uint32_t    depth;
} VkExtent3D;
```

```
VkResult vkCreateImage (
    VkDevice                              device,
    const VkImageCreateInfo*              pCreateInfo,
    const VkAllocationCallbacks*          pAllocator,
    VkImage*                              pImage);
```

```c
typedef struct VkImageCreateInfo {
    VkStructureType          sType;
    const void*              pNext;
    VkImageCreateFlags       flags;
    VkImageType              imageType;
    VkFormat                 format;
    VkExtent3D               extent;
    uint32_t                 mipLevels;
    uint32_t                 arrayLayers;
    VkSampleCountFlagBits    samples;
    VkImageTiling            tiling;
    VkImageUsageFlags        usage;
    VkSharingMode            sharingMode;
    uint32_t                 queueFamilyIndexCount;
    const uint32_t*          pQueueFamilyIndices;
    VkImageLayout            initialLayout;
} VkImageCreateInfo;
```

```cpp
VkImage image = VK_NULL_HANDLE;
VkResult result = VK_SUCCESS;

static const
VkImageCreateInfo imageCreateInfo =
{
    VK_STRUCTURE_TYPE_IMAGE_CREATE_INFO,        // sType
    nullptr,                                    // pNext
    0,                                          // flags
    VK_IMAGE_TYPE_2D,                           // imageType
    VK_FORMAT_R8G8B8A8_UNORM,                   // format
    { 1024, 1024, 1 },                          // extent
    10,                                         // mipLevels
    1,                                          // arrayLayers
    VK_SAMPLE_COUNT_1_BIT,                      // samples
    VK_IMAGE_TILING_OPTIMAL,                    // tiling
    VK_IMAGE_USAGE_SAMPLED_BIT,                 // usage
    VK_SHARING_MODE_EXCLUSIVE,                  // sharingMode
    0,                                          // queueFamilyIndexCount
    nullptr,                                    // pQueueFamilyIndices
    VK_IMAGE_LAYOUT_UNDEFINED                   // initialLayout
};

result = vkCreateImage(device, &imageCreateInfo, nullptr, &image);
```

```c
void vkGetImageSubresourceLayout (
    VkDevice                          device,
    VkImage                           image,
    const VkImageSubresource*         pSubresource,
    VkSubresourceLayout*              pLayout);
```

```c
typedef struct VkImageSubresource {
    VkImageAspectFlags    aspectMask;
    uint32_t              mipLevel;
    uint32_t              arrayLayer;
} VkImageSubresource;
```

```c
typedef struct VkSubresourceLayout {
    VkDeviceSize    offset;
    VkDeviceSize    size;
    VkDeviceSize    rowPitch;
    VkDeviceSize    arrayPitch;
    VkDeviceSize    depthPitch;
} VkSubresourceLayout;
```

```
if (cl >= 1.0)
{
    cs = 1.0;
}
else if (cl <= 0.0)
{
    cs = 0.0;
}
else if (cl < 0.0031308)
{
    cs = 12.92 * cl;
}
else
{
    cs = 1.055 * pow(cl, 0.41666) - 0.055;
}
```

```
if (cs >= 1.0)
{
    cl = 1.0;
}
else if (cs <= 0.0)
{
    cl = 0.0;
}
else if (cs <= 0.04045)
{
    cl = cs / 12.92;
}
else
{
    cl = pow((cs + 0.0555) / 1.055), 2.4)
}
```

```
VkResult vkCreateBufferView (
    VkDevice                            device,
    const VkBufferViewCreateInfo*       pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkBufferView*                       pView);
```

```c
typedef struct VkBufferViewCreateInfo {
    VkStructureType              sType;
    const void*                  pNext;
    VkBufferViewCreateFlags      flags;
    VkBuffer                     buffer;
    VkFormat                     format;
    VkDeviceSize                 offset;
    VkDeviceSize                 range;
} VkBufferViewCreateInfo;
```

```
VkResult vkCreateImageView (
    VkDevice                            device,
    const VkImageViewCreateInfo*        pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkImageView*                        pView);
```

```c
typedef struct VkImageViewCreateInfo {
    VkStructureType               sType;
    const void*                   pNext;
    VkImageViewCreateFlags        flags;
    VkImage                       image;
    VkImageViewType               viewType;
    VkFormat                      format;
    VkComponentMapping            components;
    VkImageSubresourceRange       subresourceRange;
} VkImageViewCreateInfo;
```

```c
typedef struct VkComponentMapping {
    VkComponentSwizzle    r;
    VkComponentSwizzle    g;
    VkComponentSwizzle    b;
    VkComponentSwizzle    a;
} VkComponentMapping;
```

```c
typedef struct VkImageSubresourceRange {
    VkImageAspectFlags      aspectMask;
    uint32_t                baseMipLevel;
    uint32_t                levelCount;
    uint32_t                baseArrayLayer;
    uint32_t                layerCount;
} VkImageSubresourceRange;
```

```c
void vkDestroyBuffer (
    VkDevice                          device,
    VkBuffer                          buffer,
    const VkAllocationCallbacks*      pAllocator);
```

```
void vkDestroyBufferView (
    VkDevice                        device,
    VkBufferView                    bufferView,
    const VkAllocationCallbacks*    pAllocator);
```

```
void vkDestroyImage (
    VkDevice                        device,
    VkImage                         image,
    const VkAllocationCallbacks*    pAllocator);
```

```
void vkDestroyImageView (
    VkDevice                        device,
    VkImageView                     imageView,
    const VkAllocationCallbacks*    pAllocator);
```

```
VkResult vkAllocateMemory (
    VkDevice                            device,
    const VkMemoryAllocateInfo*         pAllocateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkDeviceMemory*                     pMemory);
```

```
typedef struct VkMemoryAllocateInfo {
    VkStructureType        sType;
    const void*            pNext;
    VkDeviceSize           allocationSize;
    uint32_t               memoryTypeIndex;
} VkMemoryAllocateInfo;
```

```c
void vkFreeMemory (
    VkDevice                          device,
    VkDeviceMemory                    memory,
    const VkAllocationCallbacks*      pAllocator);
```

```
void vkGetDeviceMemoryCommitment (
    VkDevice                            device,
    VkDeviceMemory                      memory,
    VkDeviceSize*                       pCommittedMemoryInBytes);
```

```
VkResult vkMapMemory (
    VkDevice                        device,
    VkDeviceMemory                  memory,
    VkDeviceSize                    offset,
    VkDeviceSize                    size,
    VkMemoryMapFlags                flags,
    void**                          ppData);
```

```
void vkUnmapMemory (
    VkDevice                                    device,
    VkDeviceMemory                              memory);
```

```
VkResult vkFlushMappedMemoryRanges (
    VkDevice                        device,
    uint32_t                        memoryRangeCount,
    const VkMappedMemoryRange*      pMemoryRanges);
```

```c
typedef struct VkMappedMemoryRange {
    VkStructureType     sType;
    const void*         pNext;
    VkDeviceMemory      memory;
    VkDeviceSize        offset;
    VkDeviceSize        size;
} VkMappedMemoryRange;
```

```
VkResult vkInvalidateMappedMemoryRanges (
    VkDevice                                    device,
    uint32_t                                    memoryRangeCount,
    const VkMappedMemoryRange*                  pMemoryRanges);
```

```c
void vkGetBufferMemoryRequirements (
    VkDevice                                device,
    VkBuffer                                buffer,
    VkMemoryRequirements*                   pMemoryRequirements);
```

```
void vkGetImageMemoryRequirements (
    VkDevice                                device,
    VkImage                                 image,
    VkMemoryRequirements*                   pMemoryRequirements);
```

```c
typedef struct VkMemoryRequirements {
    VkDeviceSize     size;
    VkDeviceSize     alignment;
    uint32_t         memoryTypeBits;
} VkMemoryRequirements;
```

```cpp
uint32_t application::chooseHeapFromFlags(
    const VkMemoryRequirements& memoryRequirements,
    VkMemoryPropertyFlags requiredFlags,
    VkMemoryPropertyFlags preferredFlags)
{
    VkPhysicalDeviceMemoryProperties deviceMemoryProperties;

    vkGetPhysicalDeviceMemoryProperties(m_physicalDevices[0],
                                        &deviceMemoryProperties);

    uint32_t selectedType = ~0u;
    uint32_t memoryType;

    for (memoryType = 0; memoryType < 32; ++memoryType)
    {
        if (memoryRequirements.memoryTypeBits & (1 << memoryType))
        {
            const VkMemoryType& type =
                deviceMemoryProperties.memoryTypes[memoryType];

            // If it exactly matches my preferred properties, grab it.
            if ((type.propertyFlags & preferredFlags) == preferredFlags)
            {
                selectedType = memoryType;
                break;
            }
        }
    }

    if (selectedType != ~0u)
    {
        for (memoryType = 0; memoryType < 32; ++memoryType)
        {
            if (memoryRequirements.memoryTypeBits & (1 << memoryType))
            {
                const VkMemoryType& type =
                    deviceMemoryProperties.memoryTypes[memoryType];

                // If it has all my required properties, it'll do.
                if ((type.propertyFlags & requiredFlags) == requiredFlags)
                {
                    selectedType = memoryType;
                    break;
                }
            }
        }
    }

    return selectedType;
}
```

```
VkResult vkBindBufferMemory (
    VkDevice                                    device,
    VkBuffer                                    buffer,
    VkDeviceMemory                              memory,
    VkDeviceSize                               memoryOffset);
```

```
VkResult vkBindImageMemory (
    VkDevice                            device,
    VkImage                             image,
    VkDeviceMemory                      memory,
    VkDeviceSize                        memoryOffset);
```

```
void vkGetImageSparseMemoryRequirements (
    VkDevice                                device,
    VkImage                                 image,
    uint32_t*                               pSparseMemoryRequirementCount,
    VkSparseImageMemoryRequirements*        pSparseMemoryRequirements);
```

```c
typedef struct VkSparseImageMemoryRequirements {
    VkSparseImageFormatProperties    formatProperties;
    uint32_t                         imageMipTailFirstLod;
    VkDeviceSize                     imageMipTailSize;
    VkDeviceSize                     imageMipTailOffset;
    VkDeviceSize                     imageMipTailStride;
} VkSparseImageMemoryRequirements;
```

```
typedef struct VkSparseImageFormatProperties {
    VkImageAspectFlags          aspectMask;
    VkExtent3D                  imageGranularity;
    VkSparseImageFormatFlags    flags;
} VkSparseImageFormatProperties;
```

```c
void vkGetPhysicalDeviceSparseImageFormatProperties (
    VkPhysicalDevice                        physicalDevice,
    VkFormat                                format,
    VkImageType                             type,
    VkSampleCountFlagBits                   samples,
    VkImageUsageFlags                       usage,
    VkImageTiling                           tiling,
    uint32_t*                               pPropertyCount,
    VkSparseImageFormatProperties*          pProperties);
```

```
VkResult vkQueueBindSparse (
    VkQueue                         queue,
    uint32_t                        bindInfoCount,
    const VkBindSparseInfo*         pBindInfo,
    VkFence                         fence);
```

```c
typedef struct VkBindSparseInfo {
    VkStructureType                      sType;
    const void*                          pNext;
    uint32_t                             waitSemaphoreCount;
    const VkSemaphore*                   pWaitSemaphores;
    uint32_t                             bufferBindCount;
    const VkSparseBufferMemoryBindInfo*  pBufferBinds;
    uint32_t                             imageOpaqueBindCount;
    const VkSparseImageOpaqueMemoryBindInfo* pImageOpaqueBinds;
    uint32_t                             imageBindCount;
    const VkSparseImageMemoryBindInfo*   pImageBinds;
    uint32_t                             signalSemaphoreCount;
    const VkSemaphore*                   pSignalSemaphores;
} VkBindSparseInfo;
```

```c
typedef struct VkSparseBufferMemoryBindInfo {
    VkBuffer                        buffer;
    uint32_t                        bindCount;
    const VkSparseMemoryBind*       pBinds;
} VkSparseBufferMemoryBindInfo;
```

```
typedef struct VkSparseMemoryBind {
    VkDeviceSize                    resourceOffset;
    VkDeviceSize                    size;
    VkDeviceMemory          memory;
    VkDeviceSize                    memoryOffset;
    VkSparseMemoryBindFlags    flags;
} VkSparseMemoryBind;
```

```
typedef struct VkSparseImageOpaqueMemoryBindInfo {
    VkImage                      image;
    uint32_t                     bindCount;
    const VkSparseMemoryBind*    pBinds;
} VkSparseImageOpaqueMemoryBindInfo;
```

```c
typedef struct VkSparseImageMemoryBindInfo {
    VkImage                         image;
    uint32_t                        bindCount;
    const VkSparseImageMemoryBind*  pBinds;
} VkSparseImageMemoryBindInfo;
```

```c
typedef struct VkSparseImageMemoryBind {
    VkImageSubresource          subresource;
    VkOffset3D                  offset;
    VkExtent3D                  extent;
    VkDeviceMemory              memory;
    VkDeviceSize                memoryOffset;
    VkSparseMemoryBindFlags     flags;
} VkSparseImageMemoryBind;
```

```c
typedef struct VkImageSubresource {
    VkImageAspectFlags      aspectMask;
    uint32_t                mipLevel;
    uint32_t                arrayLayer;
} VkImageSubresource;
```

```c
typedef struct VkDeviceQueueCreateInfo {
    VkStructureType                 sType;
    const void*                     pNext;
    VkDeviceQueueCreateFlags        flags;
    uint32_t                        queueFamilyIndex;
    uint32_t                        queueCount;
    const float*                    pQueuePriorities;
} VkDeviceQueueCreateInfo;
```

```
void vkGetDeviceQueue (
    VkDevice                                device,
    uint32_t                                queueFamilyIndex,
    uint32_t                                queueIndex,
    VkQueue*                                pQueue);
```

```
VkResult vkCreateCommandPool (
    VkDevice                            device,
    const VkCommandPoolCreateInfo*      pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkCommandPool*                      pCommandPool);
```

```c
typedef struct VkCommandPoolCreateInfo {
    VkStructureType              sType;
    const void*                  pNext;
    VkCommandPoolCreateFlags     flags;
    uint32_t                     queueFamilyIndex;
} VkCommandPoolCreateInfo;
```

```c
VkResult vkAllocateCommandBuffers (
    VkDevice                            device,
    const VkCommandBufferAllocateInfo*  pAllocateInfo,
    VkCommandBuffer*                    pCommandBuffers);
```

```c
typedef struct VkCommandBufferAllocateInfo {
    VkStructureType        sType;
    const void*            pNext;
    VkCommandPool          commandPool;
    VkCommandBufferLevel   level;
    uint32_t               commandBufferCount;
} VkCommandBufferAllocateInfo;
```

```c
void vkFreeCommandBuffers (
    VkDevice                            device,
    VkCommandPool                       commandPool,
    uint32_t                            commandBufferCount,
    const VkCommandBuffer*              pCommandBuffers);
```

```c
void vkDestroyCommandPool (
    VkDevice                             device,
    VkCommandPool                        commandPool,
    const VkAllocationCallbacks*         pAllocator);
```

```
VkResult vkBeginCommandBuffer (
    VkCommandBuffer                         commandBuffer,
    const VkCommandBufferBeginInfo*         pBeginInfo);
```

```c
typedef struct VkCommandBufferBeginInfo {
    VkStructureType                          sType;
    const void*                              pNext;
    VkCommandBufferUsageFlags                flags;
    const VkCommandBufferInheritanceInfo*    pInheritanceInfo;
} VkCommandBufferBeginInfo;
```

```c
void vkCmdCopyBuffer (
    VkCommandBuffer                      commandBuffer,
    VkBuffer                             srcBuffer,
    VkBuffer                             dstBuffer,
    uint32_t                             regionCount,
    const VkBufferCopy*                  pRegions);
```

```c
typedef struct VkBufferCopy {
    VkDeviceSize    srcOffset;
    VkDeviceSize    dstOffset;
    VkDeviceSize    size;
} VkBufferCopy;
```

```
void CopyDataBetweenBuffers(VkCmdBuffer cmdBuffer,
                            VkBuffer srcBuffer, VkDeviceSize srcOffset,
                            VkBuffer dstBuffer, VkDeviceSize dstOffset,
                            VkDeviceSize size)
{
    const VkBufferCopy copyRegion =
    {
        srcOffset, dstOffset, size
    };

    vkCmdCopyBuffer(cmdBuffer, srcBuffer, dstBuffer, 1, &copyRegion);
}
```

```
VkResult vkEndCommandBuffer (
     VkCommandBuffer                              commandBuffer);
```

```
VkResult vkResetCommandBuffer (
    VkCommandBuffer                         commandBuffer,
    VkCommandBufferResetFlags               flags);
```

```
VkResult vkResetCommandPool (
    VkDevice                            device,
    VkCommandPool                       commandPool,
    VkCommandPoolResetFlags             flags);
```

```
VkResult vkQueueSubmit (
    VkQueue                              queue,
    uint32_t                             submitCount,
    const VkSubmitInfo*                  pSubmits,
    VkFence                              fence);
```

```
typedef struct VkSubmitInfo {
    VkStructureType                 sType;
    const void*                     pNext;
    uint32_t                        waitSemaphoreCount;
    const VkSemaphore*              pWaitSemaphores;
    const VkPipelineStageFlags*     pWaitDstStageMask;
    uint32_t                        commandBufferCount;
    const VkCommandBuffer*          pCommandBuffers;
    uint32_t                        signalSemaphoreCount;
    const VkSemaphore*              pSignalSemaphores;
} VkSubmitInfo;
```

```
VkResult vkQueueWaitIdle (
    VkQueue                                    queue);
```

```
VkResult vkDeviceWaitIdle (
    VkDevice                                    device);
```

```
void vkCmdPipelineBarrier (
    VkCommandBuffer                      commandBuffer,
    VkPipelineStageFlags                 srcStageMask,
    VkPipelineStageFlags                 dstStageMask,
    VkDependencyFlags                    dependencyFlags,
    uint32_t                             memoryBarrierCount,
    const VkMemoryBarrier*               pMemoryBarriers,
    uint32_t                             bufferMemoryBarrierCount,
    const VkBufferMemoryBarrier*         pBufferMemoryBarriers,
    uint32_t                             imageMemoryBarrierCount,
    const VkImageMemoryBarrier*          pImageMemoryBarriers);
```

```c
typedef struct VkMemoryBarrier {
    VkStructureType     sType;
    const void*         pNext;
    VkAccessFlags       srcAccessMask;
    VkAccessFlags       dstAccessMask;
} VkMemoryBarrier;
```

```c
typedef struct VkBufferMemoryBarrier {
    VkStructureType      sType;
    const void*          pNext;
    VkAccessFlags        srcAccessMask;
    VkAccessFlags        dstAccessMask;
    uint32_t             srcQueueFamilyIndex;
    uint32_t             dstQueueFamilyIndex;
    VkBuffer             buffer;
    VkDeviceSize         offset;
    VkDeviceSize         size;
} VkBufferMemoryBarrier;
```

```c
typedef struct VkImageMemoryBarrier {
    VkStructureType             sType;
    const void*                 pNext;
    VkAccessFlags               srcAccessMask;
    VkAccessFlags               dstAccessMask;
    VkImageLayout               oldLayout;
    VkImageLayout               newLayout;
    uint32_t                    srcQueueFamilyIndex;
    uint32_t                    dstQueueFamilyIndex;
    VkImage                     image;
    VkImageSubresourceRange     subresourceRange;
} VkImageMemoryBarrier;
```

```
typedef struct VkImageSubresourceRange {
    VkImageAspectFlags      aspectMask;
    uint32_t                baseMipLevel;
    uint32_t                levelCount;
    uint32_t                baseArrayLayer;
    uint32_t                layerCount;
} VkImageSubresourceRange;
```

```cpp
const VkImageMemoryBarrier imageMemoryBarriers =
{
    VK_STRUCTURE_TYPE_IMAGE_MEMORY_BARRIER,        // sType
    nullptr,                                        // pNext
    VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,           // srcAccessMask
    VK_ACCESS_SHADER_READ_BIT,                      // dstAccessMask
    VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL,       // oldLayout
    VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL,       // newLayout
    VK_QUEUE_FAMILY_IGNORED,                        // srcQueueFamilyIndex
    VK_QUEUE_FAMILY_IGNORED,                        // dstQueueFamilyIndex
    image,                                          // image
    {                                               // subresourceRange
        VK_IMAGE_ASPECT_COLOR_BIT,                  // aspectMask
        0,                                          // baseMipLevel
        VK_REMAINING_MIP_LEVELS,                    // levelCount
        0,                                          // baseArrayLayer
        VK_REMAINING_ARRAY_LAYERS                   // layerCount
    }
};

vkCmdPipelineBarrier(m_currentCommandBuffer,
                     VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,
                     VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,
                     0,
                     0, nullptr,
                     0, nullptr,
                     1, &imageMemoryBarrier);
```

```c
void vkCmdFillBuffer (
    VkCommandBuffer                              commandBuffer,
    VkBuffer                                     dstBuffer,
    VkDeviceSize                                 dstOffset,
    VkDeviceSize                                 size,
    uint32_t                                     data);
```

```
void FillBufferWithFloats(VkCommandBuffer cmdBuffer,
                          VkBuffer dstBuffer,
                          VkDeviceSize offset,
                          VkDeviceSize length,
                          const float value)
{
    vkCmdFillBuffer(cmdBuffer,
                    dstBuffer,
                    0,
                    1024,
                    *(const uint32_t*)&value);
}
```

```c
void vkCmdUpdateBuffer (
    VkCommandBuffer                    commandBuffer,
    VkBuffer                           dstBuffer,
    VkDeviceSize                       dstOffset,
    VkDeviceSize                       dataSize,
    const uint32_t*                    pData);
```

```c
void vkCmdClearColorImage (
    VkCommandBuffer                 commandBuffer,
    VkImage                         image,
    VkImageLayout                   imageLayout,
    const VkClearColorValue*        pColor,
    uint32_t                        rangeCount,
    const VkImageSubresourceRange*  pRanges);
```

```
typedef union VkClearColorValue {
    float           float32[4];
    int32_t         int32[4];
    uint32_t        uint32[4];
} VkClearColorValue;
```

```c
typedef struct VkImageSubresourceRange {
    VkImageAspectFlags      aspectMask;
    uint32_t                baseMipLevel;
    uint32_t                levelCount;
    uint32_t                baseArrayLayer;
    uint32_t                layerCount;
} VkImageSubresourceRange;
```

```
void vkCmdClearDepthStencilImage (
    VkCommandBuffer                     commandBuffer,
    VkImage                             image,
    VkImageLayout                       imageLayout,
    const VkClearDepthStencilValue*     pDepthStencil,
    uint32_t                            rangeCount,
    const VkImageSubresourceRange*      pRanges);
```

```
typedef struct VkClearDepthStencilValue {
    float           depth;
    uint32_t        stencil;
} VkClearDepthStencilValue;
```

```
void vkCmdCopyBufferToImage (
    VkCommandBuffer                         commandBuffer,
    VkBuffer                                srcBuffer,
    VkImage                                 dstImage,
    VkImageLayout                           dstImageLayout,
    uint32_t                                regionCount,
    const VkBufferImageCopy*                pRegions);
```

```c
typedef struct VkBufferImageCopy {
    VkDeviceSize                    bufferOffset;
    uint32_t                        bufferRowLength;
    uint32_t                        bufferImageHeight;
    VkImageSubresourceLayers        imageSubresource;
    VkOffset3D                      imageOffset;
    VkExtent3D                      imageExtent;
} VkBufferImageCopy;
```

```c
typedef struct VkImageSubresourceLayers {
    VkImageAspectFlags      aspectMask;
    uint32_t                mipLevel;
    uint32_t                baseArrayLayer;
    uint32_t                layerCount;
} VkImageSubresourceLayers;
```

```
void vkCmdCopyImageToBuffer (
    VkCommandBuffer                      commandBuffer,
    VkImage                              srcImage,
    VkImageLayout                        srcImageLayout,
    VkBuffer                             dstBuffer,
    uint32_t                             regionCount,
    const VkBufferImageCopy*             pRegions);
```

```c
void vkCmdCopyImage (
    VkCommandBuffer                     commandBuffer,
    VkImage                             srcImage,
    VkImageLayout                       srcImageLayout,
    VkImage                             dstImage,
    VkImageLayout                       dstImageLayout,
    uint32_t                            regionCount,
    const VkImageCopy*                  pRegions);
```

```
typedef struct VkImageCopy {
    VkImageSubresourceLayers        srcSubresource;
    VkOffset3D                      srcOffset;
    VkImageSubresourceLayers        dstSubresource;
    VkOffset3D                      dstOffset;
    VkExtent3D                      extent;
} VkImageCopy;
```

```
void vkCmdBlitImage (
    VkCommandBuffer                 commandBuffer,
    VkImage                         srcImage,
    VkImageLayout                   srcImageLayout,
    VkImage                         dstImage,
    VkImageLayout                   dstImageLayout,
    uint32_t                        regionCount,
    const VkImageBlit*              pRegions,
    VkFilter                        filter);
```

```
typedef struct VkImageBlit {
    VkImageSubresourceLayers    srcSubresource;
    VkOffset3D                  srcOffsets[2];
    VkImageSubresourceLayers    dstSubresource;
    VkOffset3D                  dstOffsets[2];
} VkImageBlit;
```

```
VkBool32 vkGetPhysicalDeviceWin32PresentationSupportKHR(
    VkPhysicalDevice                            physicalDevice,
    uint32_t                                    queueFamilyIndex);
```

```
VkResult vkCreateWin32SurfaceKHR(
    VkInstance                              instance,
    const VkWin32SurfaceCreateInfoKHR*      pCreateInfo,
    const VkAllocationCallbacks*            pAllocator,
    VkSurfaceKHR*                           pSurface);
```

```c
typedef struct VkWin32SurfaceCreateInfoKHR {
    VkStructureType                 sType;
    const void*                     pNext;
    VkWin32SurfaceCreateFlagsKHR    flags;
    HINSTANCE                       hinstance;
    HWND                            hwnd;
} VkWin32SurfaceCreateInfoKHR;
```

```
VkBool32 vkGetPhysicalDeviceXlibPresentationSupportKHR(
    VkPhysicalDevice                            physicalDevice,
    uint32_t                                    queueFamilyIndex,
    Display*                                    dpy,
    VisualID                                    visualID);
```

```c
VkResult vkCreateXlibSurfaceKHR(
    VkInstance                                  instance,
    const VkXlibSurfaceCreateInfoKHR*           pCreateInfo,
    const VkAllocationCallbacks*                pAllocator,
    VkSurfaceKHR*                               pSurface);
```

```
typedef struct VkXlibSurfaceCreateInfoKHR {
    VkStructureType                 sType;
    const void*                     pNext;
    VkXlibSurfaceCreateFlagsKHR     flags;
    Display*                        dpy;
    Window                          window;
} VkXlibSurfaceCreateInfoKHR;
```

```
VkBool32 vkGetPhysicalDeviceXcbPresentationSupportKHR(
    VkPhysicalDevice                            physicalDevice,
    uint32_t                                    queueFamilyIndex,
    xcb_connection_t*                           connection,
    xcb_visualid_t                              visual_id);
```

```c
VkResult vkCreateXcbSurfaceKHR(
    VkInstance                                  instance,
    const VkXcbSurfaceCreateInfoKHR*            pCreateInfo,
    const VkAllocationCallbacks*                pAllocator,
    VkSurfaceKHR*                               pSurface);
```

```c
typedef struct VkXcbSurfaceCreateInfoKHR {
    VkStructureType                 sType;
    const void*                     pNext;
    VkXcbSurfaceCreateFlagsKHR      flags;
    xcb_connection_t*               connection;
    xcb_window_t                    window;
} VkXcbSurfaceCreateInfoKHR;
```

```
VkResult vkCreateSwapchainKHR(
    VkDevice                            device,
    const VkSwapchainCreateInfoKHR*     pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkSwapchainKHR*                     pSwapchain);
```

```c
typedef struct VkSwapchainCreateInfoKHR {
    VkStructureType                  sType;
    const void*                      pNext;
    VkSwapchainCreateFlagsKHR        flags;
    VkSurfaceKHR                     surface;
    uint32_t                         minImageCount;
    VkFormat                         imageFormat;
    VkColorSpaceKHR                  imageColorSpace;
    VkExtent2D                       imageExtent;
    uint32_t                         imageArrayLayers;
    VkImageUsageFlags                imageUsage;
    VkSharingMode                    imageSharingMode;
    uint32_t                         queueFamilyIndexCount;
    const uint32_t*                  pQueueFamilyIndices;
    VkSurfaceTransformFlagBitsKHR    preTransform;
    VkCompositeAlphaFlagBitsKHR      compositeAlpha;
    VkPresentModeKHR                 presentMode;
    VkBool32                         clipped;
    VkSwapchainKHR                   oldSwapchain;
} VkSwapchainCreateInfoKHR;
```

```
VkResult vkGetPhysicalDeviceSurfaceCapabilitiesKHR(
    VkPhysicalDevice                            physicalDevice,
    VkSurfaceKHR                                surface,
    VkSurfaceCapabilitiesKHR*                   pSurfaceCapabilities);
```

```c
typedef struct VkSurfaceCapabilitiesKHR {
    uint32_t                        minImageCount;
    uint32_t                        maxImageCount;
    VkExtent2D                      currentExtent;
    VkExtent2D                      minImageExtent;
    VkExtent2D                      maxImageExtent;
    uint32_t                        maxImageArrayLayers;
    VkSurfaceTransformFlagsKHR      supportedTransforms;
    VkSurfaceTransformFlagBitsKHR   currentTransform;
    VkCompositeAlphaFlagsKHR        supportedCompositeAlpha;
    VkImageUsageFlags               supportedUsageFlags;
} VkSurfaceCapabilitiesKHR;
```

```
VkResult vkGetSwapchainImagesKHR(
    VkDevice                                    device,
    VkSwapchainKHR                              swapchain,
    uint32_t*                                   pSwapchainImageCount,
    VkImage*                                    pSwapchainImages);
```

```cpp
VkResult result;

// First, we create the swap chain.
VkSwapchainCreateInfoKHR swapChainCreateInfo =
{
    VK_STRUCTURE_TYPE_SWAPCHAIN_CREATE_INFO_KHR, // sType
    nullptr,                                     // pNext
    0,                                           // flags
    m_mainSurface,                               // surface
    2,                                           // minImageCount
    VK_FORMAT_R8G8B8A8_UNORM,                    // imageFormat
    VK_COLORSPACE_SRGB_NONLINEAR_KHR,            // imageColorSpace
    { 1024, 768 },                               // imageExtent
    1,                                           // imageArrayLayers
    VK_IMAGE_USAGE_COLOR_ATTACHMENT_BIT,         // imageUsage
    VK_SHARING_MODE_EXCLUSIVE,                   // imageSharingMode
    0,                                           // queueFamilyIndexCount
    nullptr,                                     // pQueueFamilyIndices
    VK_SURFACE_TRANSFORM_INHERIT_BIT_KHR,        // preTransform
    VK_COMPOSITE_ALPHA_OPAQUE_BIT_KHR,           // compositeAlpha
    VK_PRESENT_MODE_FIFO_KHR,                    // presentMode
    VK_TRUE,                                     // clipped
    m_swapChain                                  // oldSwapchain
};

result = vkCreateSwapchainKHR(m_logicalDevice,
                              &swapChainCreateInfo,
                              nullptr,
                              &m_swapChain);

// Next, we query the swap chain for the number of images it actually contains.
uint32_t swapChainImageCount = 0;

if (result == VK_SUCCESS)
{
    result = vkGetSwapchainImagesKHR(m_logicalDevice,
                                     m_swapChain,
                                     &swapChainImageCount,
                                     nullptr);
}

if (result == VK_SUCCESS)
{
    // Now we resize our image array and retrieve the image handles from the
    // swap chain.
    m_swapChainImages.resize(swapChainImageCount);

    result = vkGetSwapchainImagesKHR(m_logicalDevice,
                                     m_swapChain,
                                     &swapChainImageCount,
                                     m_swapChainImages.data());
}

return result == VK_SUCCESS ? m_swapChain : VK_NULL_HANDLE;
```

```c
VKAPI_ATTR VkResult VKAPI_CALL vkGetPhysicalDeviceSurfaceFormatsKHR(
    VkPhysicalDevice                            physicalDevice,
    VkSurfaceKHR                                surface,
    uint32_t*                                   pSurfaceFormatCount,
    VkSurfaceFormatKHR*                         pSurfaceFormats);
```

```c
typedef struct VkSurfaceFormatKHR {
    VkFormat            format;
    VkColorSpaceKHR     colorSpace;
} VkSurfaceFormatKHR;
```

```
VkResult vkAcquireNextImageKHR(
    VkDevice                                device,
    VkSwapchainKHR                          swapchain,
    uint64_t                                timeout,
    VkSemaphore                             semaphore,
    VkFence                                 fence,
    uint32_t*                               pImageIndex);
```

```
VkResult vkGetPhysicalDeviceDisplayPropertiesKHR(
    VkPhysicalDevice                                physicalDevice,
    uint32_t*                                       pPropertyCount,
    VkDisplayPropertiesKHR*                         pProperties);
```

```c
typedef struct VkDisplayPropertiesKHR {
    VkDisplayKHR                    display;
    const char*                     displayName;
    VkExtent2D                      physicalDimensions;
    VkExtent2D                      physicalResolution;
    VkSurfaceTransformFlagsKHR      supportedTransforms;
    VkBool32                        planeReorderPossible;
    VkBool32                        persistentContent;
} VkDisplayPropertiesKHR;
```

```
VkResult vkGetPhysicalDeviceDisplayPlanePropertiesKHR(
    VkPhysicalDevice                            physicalDevice,
    uint32_t*                                   pPropertyCount,
    VkDisplayPlanePropertiesKHR*                pProperties);
```

```c
typedef struct VkDisplayPlanePropertiesKHR {
    VkDisplayKHR        currentDisplay;
    uint32_t            currentStackIndex;
} VkDisplayPlanePropertiesKHR;
```

```
VkResult vkGetDisplayPlaneSupportedDisplaysKHR(
    VkPhysicalDevice                                physicalDevice,
    uint32_t                                        planeIndex,
    uint32_t*                                       pDisplayCount,
    VkDisplayKHR*                                   pDisplays);
```

```
VkResult vkGetDisplayPlaneCapabilitiesKHR(
    VkPhysicalDevice                            physicalDevice,
    VkDisplayModeKHR                            mode,
    uint32_t                                    planeIndex,
    VkDisplayPlaneCapabilitiesKHR*              pCapabilities);
```

```c
typedef struct VkDisplayPlaneCapabilitiesKHR {
    VkDisplayPlaneAlphaFlagsKHR        supportedAlpha;
    VkOffset2D                         minSrcPosition;
    VkOffset2D                         maxSrcPosition;
    VkExtent2D                         minSrcExtent;
    VkExtent2D                         maxSrcExtent;
    VkOffset2D                         minDstPosition;
    VkOffset2D                         maxDstPosition;
    VkExtent2D                         minDstExtent;
    VkExtent2D                         maxDstExtent;
} VkDisplayPlaneCapabilitiesKHR;
```

```
VkResult vkGetDisplayModePropertiesKHR(
    VkPhysicalDevice                            physicalDevice,
    VkDisplayKHR                                display,
    uint32_t*                                   pPropertyCount,
    VkDisplayModePropertiesKHR*                 pProperties);
```

```c
typedef struct VkDisplayModePropertiesKHR {
    VkDisplayModeKHR                displayMode;
    VkDisplayModeParametersKHR      parameters;
} VkDisplayModePropertiesKHR;
```

```c
typedef struct VkDisplayModeParametersKHR {
    VkExtent2D      visibleRegion;
    uint32_t        refreshRate;
} VkDisplayModeParametersKHR;
```

```
VkResult vkCreateDisplayModeKHR(
    VkPhysicalDevice                        physicalDevice,
    VkDisplayKHR                            display,
    const VkDisplayModeCreateInfoKHR*       pCreateInfo,
    const VkAllocationCallbacks*            pAllocator,
    VkDisplayModeKHR*                       pMode);
```

```
typedef struct VkDisplayModeCreateInfoKHR {
    VkStructureType                 sType;
    const void*                     pNext;
    VkDisplayModeCreateFlagsKHR     flags;
    VkDisplayModeParametersKHR      parameters;
} VkDisplayModeCreateInfoKHR;
```

```
VkResult vkCreateDisplayPlaneSurfaceKHR(
    VkInstance                                  instance,
    const VkDisplaySurfaceCreateInfoKHR*        pCreateInfo,
    const VkAllocationCallbacks*                pAllocator,
    VkSurfaceKHR*                               pSurface);
```

```c
typedef struct VkDisplaySurfaceCreateInfoKHR {
    VkStructureType                     sType;
    const void*                         pNext;
    VkDisplaySurfaceCreateFlagsKHR      flags;
    VkDisplayModeKHR                    displayMode;
    uint32_t                            planeIndex;
    uint32_t                            planeStackIndex;
    VkSurfaceTransformFlagBitsKHR       transform;
    float                               globalAlpha;
    VkDisplayPlaneAlphaFlagBitsKHR      alphaMode;
    VkExtent2D                          imageExtent;
} VkDisplaySurfaceCreateInfoKHR;
```

```
VkResult vkGetPhysicalDeviceSurfaceSupportKHR(
    VkPhysicalDevice                              physicalDevice,
    uint32_t                                      queueFamilyIndex,
    VkSurfaceKHR                                  surface,
    VkBool32*                                     pSupported);
```

```cpp
const VkImageMemoryBarrier barrier =
{
    VK_STRUCTURE_TYPE_IMAGE_MEMORY_BARRIER,         // sType
    nullptr,                                        // pNext
    VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,           // srcAccessMask
    VK_ACCESS_MEMORY_READ_BIT,                      // dstAccessMask
    VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL,       // oldLayout
    VK_IMAGE_LAYOUT_PRESENT_SRC_KHR,                // newLayout
    0,                                              // srcQueueFamilyIndex
    0,                                              // dstQueueFamilyIndex
    sourceImage,                                    // image
    {                                               // subresourceRange
        VK_IMAGE_ASPECT_COLOR_BIT,                  // aspectMask
        0,                                          // baseMipLevel
        1,                                          // levelCount
        0,                                          // baseArrayLayer
        1                                           // layerCount
    }
};

vkCmdPipelineBarrier(cmdBuffer,
                     VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,
                     VK_PIPELINE_STAGE_BOTTOM_OF_PIPE_BIT,
                     0,
                     0, nullptr,
                     0, nullptr,
                     1, &barrier);
```

```
VkResult vkQueuePresentKHR(
    VkQueue                                     queue,
    const VkPresentInfoKHR*                     pPresentInfo);
```

```c
typedef struct VkPresentInfoKHR {
    VkStructureType              sType;
    const void*                  pNext;
    uint32_t                     waitSemaphoreCount;
    const VkSemaphore*           pWaitSemaphores;
    uint32_t                     swapchainCount;
    const VkSwapchainKHR*        pSwapchains;
    const uint32_t*              pImageIndices;
    VkResult*                    pResults;
} VkPresentInfoKHR;
```

```c
void vkDestroySwapchainKHR(
    VkDevice                                    device,
    VkSwapchainKHR                              swapchain,
    const VkAllocationCallbacks*                pAllocator);
```

```
glslangvalidator simple.comp -o simple.spv
```

```
; SPIR-V
; Version: 1.0
; Generator: Khronos Glslang Reference Front End; 1
; Bound: 6
; Schema: 0
               OpCapability Shader
          %1 = OpExtInstImport "GLSL.std.450"
               OpMemoryModel Logical GLSL450
               OpEntryPoint GLCompute %4 "main"
               OpExecutionMode %4 LocalSize 1 1 1
               OpSource GLSL 450
               OpName %4 "main"
          %2 = OpTypeVoid
          %3 = OpTypeFunction %2
          %4 = OpFunction %2 None %3
          %5 = OpLabel
               OpReturn
               OpFunctionEnd
```

```
VkResult vkCreateShaderModule (
    VkDevice                        device,
    const VkShaderModuleCreateInfo* pCreateInfo,
    const VkAllocationCallbacks*    pAllocator,
    VkShaderModule*                 pShaderModule);
```

```c
typedef struct VkShaderModuleCreateInfo {
    VkStructureType                sType;
    const void*                    pNext;
    VkShaderModuleCreateFlags      flags;
    size_t                         codeSize;
    const uint32_t*                pCode;
} VkShaderModuleCreateInfo;
```

```c
void vkDestroyShaderModule (
    VkDevice                            device,
    VkShaderModule                      shaderModule,
    const VkAllocationCallbacks*        pAllocator);
```

```glsl
#version 450 core

layout (local_size_x = 4, local_size_y = 5, local_size_z 6) in;

void main(void)
{
    // Do nothing.
}
```

```
...
        OpCapability Shader
   %1 = OpExtInstImport "GLSL.std.450"
        OpMemoryModel Logical GLSL450
        OpEntryPoint GLCompute %4 "main"
        OpExecutionMode %4 LocalSize 4 5 6
        OpSource GLSL 450
...
```

```
VkResult vkCreateComputePipelines (
    VkDevice                                   device,
    VkPipelineCache                            pipelineCache,
    uint32_t                                   createInfoCount,
    const VkComputePipelineCreateInfo*         pCreateInfos,
    const VkAllocationCallbacks*               pAllocator,
    VkPipeline*                                pPipelines);
```

```c
typedef struct VkComputePipelineCreateInfo {
    VkStructureType                        sType;
    const void*                            pNext;
    VkPipelineCreateFlags                  flags;
    VkPipelineShaderStageCreateInfo        stage;
    VkPipelineLayout                       layout;
    VkPipeline                             basePipelineHandle;
    int32_t                                basePipelineIndex;
} VkComputePipelineCreateInfo;
```

```
typedef struct VkPipelineShaderStageCreateInfo {
    VkStructureType                     sType;
    const void*                         pNext;
    VkPipelineShaderStageCreateFlags    flags;
    VkShaderStageFlagBits               stage;
    VkShaderModule                      module;
    const char*                         pName;
    const VkSpecializationInfo*         pSpecializationInfo;
} VkPipelineShaderStageCreateInfo;
```

```glsl
layout (constant_id = 0) const int numThings = 42;
layout (constant_id = 1) const float thingScale = 4.2f;
layout (constant_id = 2) const bool doThat = false;
```

```
...
            OpDecorate %7 SpecId 0
            OpDecorate %9 SpecId 1
            OpDecorate %11 SpecId 2
     %6 = OpTypeInt 32 1
     %7 = OpSpecConstant %6 42
     %8 = OpTypeFloat 32
     %9 = OpSpecConstant %8 4.2
    %10 = OpTypeBool
    %11 = OpSpecConstantFalse %10
...
```

```c
typedef struct VkSpecializationInfo {
    uint32_t                       mapEntryCount;
    const VkSpecializationMapEntry* pMapEntries;
    size_t                         dataSize;
    const void*                    pData;
} VkSpecializationInfo;
```

```c
typedef struct VkSpecializationMapEntry {
    uint32_t    constantID;
    uint32_t    offset;
    size_t      size;
} VkSpecializationMapEntry;
```

```
void vkDestroyPipeline (
    VkDevice                            device,
    VkPipeline                          pipeline,
    const VkAllocationCallbacks*        pAllocator);
```

```
VkResult vkCreatePipelineCache (
    VkDevice                            device,
    const VkPipelineCacheCreateInfo*    pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkPipelineCache*                    pPipelineCache);
```

```
typedef struct VkPipelineCacheCreateInfo {
    VkStructureType                  sType;
    const void*                      pNext;
    VkPipelineCacheCreateFlags       flags;
    size_t                           initialDataSize;
    const void*                      pInitialData;
} VkPipelineCacheCreateInfo;
```

```
VkResult vkGetPipelineCacheData (
    VkDevice                                    device,
    VkPipelineCache                             pipelineCache,
    size_t*                                     pDataSize,
    void*                                       pData);
```

```cpp
VkResult SaveCacheToFile(VkDevice device, VkPipelineCache cache,
          const char* fileName)
{
    size_t cacheDataSize;
    VkResult result = VK_SUCCESS;

    // Determine the size of the cache data.
    result = vkGetPipelineCacheData(device,
                                    cache,
                                    &cacheDataSize,
                                    nullptr);

    if (result == VK_SUCCESS && cacheDataSize != 0)
    {
        FILE* pOutputFile;
        void* pData;

        // Allocate a temporary store for the cache data.
        result = VK_ERROR_OUT_OF_HOST_MEMORY;
        pData = malloc(cacheDataSize);

        if (pData != nullptr)
        {
            // Retrieve the actual data from the cache.
            result = vkGetPipelineCacheData(device,
                                            cache,
                                            &cacheDataSize,
                                            pData);
```

```cpp
        if (result == VK_SUCCESS)
        {
            // Open the file and write the data to it.
            pOutputFile = fopen(fileName, "wb");

            if (pOutputFile != nullptr)
            {
                fwrite(pData, 1, cacheDataSize, pOutputFile);

                fclose(pOutputFile);
            }

            free(pData);
        }
    }
}

return result;
}
```

```c
// This structure does not exist in official headers but is included here
// for illustration.
typedef struct VkPipelineCacheHeader {
    uint32_t        length;
    uint32_t        version;
    uint32_t        vendorID;
    uint32_t        deviceID;
    uint8_t         uuid[16];
} VkPipelineCacheHeader;
```

```
VkResult vkMergePipelineCaches (
    VkDevice                              device,
    VkPipelineCache                       dstCache,
    uint32_t                              srcCacheCount,
    const VkPipelineCache*                pSrcCaches);
```

```
void vkDestroyPipelineCache (
    VkDevice                            device,
    VkPipelineCache                     pipelineCache,
    const VkAllocationCallbacks*        pAllocator);
```

```
void vkCmdBindPipeline (
    VkCommandBuffer                             commandBuffer,
    VkPipelineBindPoint                         pipelineBindPoint,
    VkPipeline                                  pipeline);
```

```
void vkCmdDispatch (
    VkCommandBuffer                         commandBuffer,
    uint32_t                                x,
    uint32_t                                y,
    uint32_t                                z);
```

```
void vkCmdDispatchIndirect (
    VkCommandBuffer                         commandBuffer,
    VkBuffer                                buffer,
    VkDeviceSize                            offset);
```

```
typedef struct VkDispatchIndirectCommand {
    uint32_t     x;
    uint32_t     y;
    uint32_t     z;
} VkDispatchIndirectCommand;
```

```
VkResult vkCreateDescriptorSetLayout (
    VkDevice                                device,
    const VkDescriptorSetLayoutCreateInfo*  pCreateInfo,
    const VkAllocationCallbacks*            pAllocator,
    VkDescriptorSetLayout*                  pSetLayout);
```

```c
typedef struct VkDescriptorSetLayoutCreateInfo {
    VkStructureType                     sType;
    const void*                         pNext;
    VkDescriptorSetLayoutCreateFlags    flags;
    uint32_t                            bindingCount;
    const VkDescriptorSetLayoutBinding* pBindings;
} VkDescriptorSetLayoutCreateInfo;
```

```c
typedef struct VkDescriptorSetLayoutBinding {
    uint32_t                binding;
    VkDescriptorType        descriptorType;
    uint32_t                descriptorCount;
    VkShaderStageFlags      stageFlags;
    const VkSampler*        pImmutableSamplers;
} VkDescriptorSetLayoutBinding;
```

```glsl
#version 450 core

layout (set = 0, binding = 0) uniform sampler2D myTexture;
layout (set = 0, binding = 2) uniform sampler3D myLut;
layout (set = 1, binding = 0) uniform myTransforms
{
    mat4 transform1;
    mat3 transform2;
};

void main(void)
{
    // Do nothing!
}
```

```
; SPIR-V
; Version: 1.0
; Generator: Khronos Glslang Reference Front End; 1
; Bound: 22
; Schema: 0
               OpCapability Shader
          %1 = OpExtInstImport "GLSL.std.450"
               OpMemoryModel Logical GLSL450
               OpEntryPoint GLCompute %4 "main"
               OpExecutionMode %4 LocalSize 1 1 1
               OpSource GLSL 450
               OpName %4 "main"
               OpName %10 "myTexture"
               OpName %14 "myLut"
               OpName %19 "myTransforms"
               OpMemberName %19 0 "transform1"
               OpMemberName %19 1 "transform2"
               OpName %21 ""
               OpDecorate %10 DescriptorSet 0
               OpDecorate %10 Binding 0
               OpDecorate %14 DescriptorSet 0
               OpDecorate %14 Binding 2
               OpMemberDecorate %19 0 ColMajor
               OpMemberDecorate %19 0 Offset 0
               OpMemberDecorate %19 0 MatrixStride 16
               OpMemberDecorate %19 1 ColMajor
               OpMemberDecorate %19 1 Offset 64
               OpMemberDecorate %19 1 MatrixStride 16
               OpDecorate %19 Block
               OpDecorate %21 DescriptorSet 1
               OpDecorate %21 Binding 0
          %2 = OpTypeVoid
          %3 = OpTypeFunction %2
          %6 = OpTypeFloat 32
          %7 = OpTypeImage %6 2D 0 0 0 1 Unknown
          %8 = OpTypeSampledImage %7
          %9 = OpTypePointer UniformConstant %8
         %10 = OpVariable %9 UniformConstant
         %11 = OpTypeImage %6 3D 0 0 0 1 Unknown
         %12 = OpTypeSampledImage %11
         %13 = OpTypePointer UniformConstant %12
         %14 = OpVariable %13 UniformConstant
         %15 = OpTypeVector %6 4
         %16 = OpTypeMatrix %15 4
         %17 = OpTypeVector %6 3
         %18 = OpTypeMatrix %17 3
         %19 = OpTypeStruct %16 %18
         %20 = OpTypePointer Uniform %19
         %21 = OpVariable %20 Uniform
          %4 = OpFunction %2 None %3
          %5 = OpLabel
               OpReturn
               OpFunctionEnd
```

```c
VkResult vkCreatePipelineLayout (
    VkDevice                            device,
    const VkPipelineLayoutCreateInfo*   pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkPipelineLayout*                   pPipelineLayout);
```

```c
typedef struct VkPipelineLayoutCreateInfo {
    VkStructureType                 sType;
    const void*                     pNext;
    VkPipelineLayoutCreateFlags     flags;
    uint32_t                        setLayoutCount;
    const VkDescriptorSetLayout*    pSetLayouts;
    uint32_t                        pushConstantRangeCount;
    const VkPushConstantRange*      pPushConstantRanges;
} VkPipelineLayoutCreateInfo;
```

```cpp
// This describes our combined image-samplers. One set, two disjoint bindings
static const VkDescriptorSetLayoutBinding Samplers[] =
{
    {
        0,                                                // Start from binding 0
        VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER,        // Combined image-sampler
        1,                                                // Create one binding
        VK_SHADER_STAGE_ALL,                              // Usable in all stages
        nullptr                                           // No static samplers
    },
    {
        2,                                                // Start from binding 2

        VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER,        // Combined image-sampler
        1,                                                // Create one binding
        VK_SHADER_STAGE_ALL,                              // Usable in all stages
        nullptr                                           // No static samplers
    }
};

// This is our uniform block. One set, one binding.
static const VkDescriptorSetLayoutBinding UniformBlock =
{
    0,                                                    // Start from binding 0
    VK_DESCRIPTOR_TYPE_UNIFORM_BUFFER,                    // Uniform block
    1,                                                    // One binding
    VK_SHADER_STAGE_ALL,                                  // All stages
    nullptr                                               // No static samplers
};

// Now create the two descriptor set layouts.
static const VkDescriptorSetLayoutCreateInfo createInfoSamplers =
{
    VK_STRUCTURE_TYPE_DESCRIPTOR_SET_LAYOUT_CREATE_INFO,
    nullptr,
    0,
    2,
    &Samplers[0]
};

static const VkDescriptorSetLayoutCreateInfo createInfoUniforms =
{
    VK_STRUCTURE_TYPE_DESCRIPTOR_SET_LAYOUT_CREATE_INFO,
    nullptr,
    0,
    1,
    &UniformBlock
};

// This array holds the two set layouts.
VkDescriptorSetLayout setLayouts[2];

vkCreateDescriptorSetLayout(device, &createInfoSamplers,
                            nullptr, &setLaouts[0]);
vkCreateDescriptorSetLayout(device, &createInfoUniforms,
                            nullptr, &setLayouts[1]);

// Now create the pipeline layout.
const VkPipelineLayoutCreateInfo pipelineLayoutCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_LAYOUT_CREATE_INFO, nullptr,
    0,
    2, setLayouts,
    0, nullptr
};

VkPipelineLayout pipelineLayout;

vkCreatePipelineLayout(device, &pipelineLayoutCreateInfo,
                       nullptr, pipelineLayout);
```

```
void vkDestroyPipelineLayout (
    VkDevice                            device,
    VkPipelineLayout                    pipelineLayout,
    const VkAllocationCallbacks*        pAllocator);
```

```
void vkDestroyDescriptorSetLayout (
    VkDevice                            device,
    VkDescriptorSetLayout               descriptorSetLayout,
    const VkAllocationCallbacks*        pAllocator);
```

```
VkResult vkCreateDescriptorPool (
    VkDevice                            device,
    const VkDescriptorPoolCreateInfo*   pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkDescriptorPool*                   pDescriptorPool);
```

```c
typedef struct VkDescriptorPoolCreateInfo {
    VkStructureType                sType;
    const void*                    pNext;
    VkDescriptorPoolCreateFlags    flags;
    uint32_t                       maxSets;
    uint32_t                       poolSizeCount;
    const VkDescriptorPoolSize*    pPoolSizes;
} VkDescriptorPoolCreateInfo;
```

```c
typedef struct VkDescriptorPoolSize {
    VkDescriptorType    type;
    uint32_t            descriptorCount;
} VkDescriptorPoolSize;
```

```
VkResult vkAllocateDescriptorSets (
    VkDevice                                    device,
    const VkDescriptorSetAllocateInfo*          pAllocateInfo,
    VkDescriptorSet*                            pDescriptorSets);
```

```c
typedef struct VkDescriptorSetAllocateInfo {
    VkStructureType                  sType;
    const void*                      pNext;
    VkDescriptorPool                 descriptorPool;
    uint32_t                         descriptorSetCount;
    const VkDescriptorSetLayout*     pSetLayouts;
} VkDescriptorSetAllocateInfo;
```

```
VkResult vkFreeDescriptorSets (
    VkDevice                                    device,
    VkDescriptorPool                            descriptorPool,
    uint32_t                                    descriptorSetCount,
    const VkDescriptorSet*                      pDescriptorSets);
```

```
VkResult vkResetDescriptorPool (
    VkDevice                          device,
    VkDescriptorPool                  descriptorPool,
    VkDescriptorPoolResetFlags        flags);
```

```
void vkDestroyDescriptorPool (
    VkDevice                        device,
    VkDescriptorPool                descriptorPool,
    const VkAllocationCallbacks*    pAllocator);
```

```
void vkUpdateDescriptorSets (
    VkDevice                        device,
    uint32_t                        descriptorWriteCount,
    const VkWriteDescriptorSet*     pDescriptorWrites,
    uint32_t                        descriptorCopyCount,
    const VkCopyDescriptorSet*      pDescriptorCopies);
```

```
typedef struct VkWriteDescriptorSet {
    VkStructureType                   sType;
    const void*                       pNext;
    VkDescriptorSet                   dstSet;
    uint32_t                          dstBinding;
    uint32_t                          dstArrayElement;
    uint32_t                          descriptorCount;
    VkDescriptorType                  descriptorType;
    const VkDescriptorImageInfo*      pImageInfo;
    const VkDescriptorBufferInfo*     pBufferInfo;
    const VkBufferView*               pTexelBufferView;
} VkWriteDescriptorSet;
```

```
typedef struct VkDescriptorImageInfo {
    VkSampler          sampler;
    VkImageView        imageView;
    VkImageLayout      imageLayout;
} VkDescriptorImageInfo;
```

```c
typedef struct VkDescriptorBufferInfo {
    VkBuffer            buffer;
    VkDeviceSize        offset;
    VkDeviceSize        range;
} VkDescriptorBufferInfo;
```

```
typedef struct VkCopyDescriptorSet {
    VkStructureType      sType;
    const void*          pNext;
    VkDescriptorSet      srcSet;
    uint32_t             srcBinding;
    uint32_t             srcArrayElement;
    VkDescriptorSet      dstSet;
    uint32_t             dstBinding;
    uint32_t             dstArrayElement;
    uint32_t             descriptorCount;
} VkCopyDescriptorSet;
```

```c
void vkCmdBindDescriptorSets (
    VkCommandBuffer                 commandBuffer,
    VkPipelineBindPoint             pipelineBindPoint,
    VkPipelineLayout                layout,
    uint32_t                        firstSet,
    uint32_t                        descriptorSetCount,
    const VkDescriptorSet*          pDescriptorSets,
    uint32_t                        dynamicOffsetCount,
    const uint32_t*                 pDynamicOffsets);
```

```glsl
layout (set = 0, binding = 1) uniform my_uniform_buffer_t
{
    float foo;
    vec4 bar;
    int baz[42];
} my_uniform_buffer;

layout (set = 0, binding = 2) buffer my_storage_buffer_t
{
    int peas;
    float carrots;
    vec3 potatoes[99];
} my_storage_buffer;
```

```
       OpCapability Shader
%1 =   OpExtInstImport "GLSL.std.450"
       OpMemoryModel Logical GLSL450
       OpEntryPoint Vertex %4 "main"
       OpSource GLSL 450
       OpName %4 "main"
       ;; Name the my_uniform_buffer_t block and its members.
       OpName %12 "my_uniform_buffer_t"
       OpMemberName %12 0 "foo"
       OpMemberName %12 1 "bar"
       OpMemberName %12 2 "baz"
       OpName %14 "my_uniform_buffer"
       ;; Name the my_storage_buffer_t block and its members.
       OpName %18 "my_storage_buffer_t"
       OpMemberName %18 0 "peas"
       OpMemberName %18 1 "carrots"
       OpMemberName %18 2 "potatoes"
       OpName %20 "my_storage_buffer"
       OpDecorate %11 ArrayStride 16
       ;; Assign offsets to the members of my_uniform_buffer_t.
       OpMemberDecorate %12 0 Offset 0
       OpMemberDecorate %12 1 Offset 16
       OpMemberDecorate %12 2 Offset 32
       OpDecorate %12 Block
       OpDecorate %14 DescriptorSet 0
       OpDecorate %14 Binding 1
       OpDecorate %17 ArrayStride 16
       ;; Assign offsets to the members of my_storage_buffer_t.
       OpMemberDecorate %18 0 Offset 0
       OpMemberDecorate %18 1 Offset 4
       OpMemberDecorate %18 2 Offset 16
       OpDecorate %18 BufferBlock
       OpDecorate %20 DescriptorSet 0
       OpDecorate %20 Binding 2
       ...
```

```glsl
layout (set = 0, binding = 3) uniform samplerBuffer my_float_texel_buffer;
layout (set = 0, binding = 4) uniform isamplerBuffer my_signed_texel_buffer;
layout (set = 0, binding = 5) uniform usamplerBuffer my_unsigned_texel_buffer;
```

```
         OpCapability Shader
         OpCapability SampledBuffer
    %1 = OpExtInstImport "GLSL.std.450"
         OpMemoryModel Logical GLSL450
         OpEntryPoint Vertex %4 "main"
         OpSource GLSL 450
         OpName %4 "main"
         ;; Name our texel buffers.
         OpName %10 "my_float_texel_buffer"
         OpName %15 "my_signed_texel_buffer"
         OpName %20 "my_unsigned_texel_buffer"
         ;; Assign set and binding decorations.
         OpDecorate %10 DescriptorSet 0
         OpDecorate %10 Binding 3
         OpDecorate %15 DescriptorSet 0
         OpDecorate %15 Binding 4
         OpDecorate %20 DescriptorSet 0
         OpDecorate %20 Binding 5
    %2 = OpTypeVoid
    %3 = OpTypeFunction %2
         ;; Declare the three texel buffer variables.
    %6 = OpTypeFloat 32
    %7 = OpTypeImage %6 Buffer 0 0 0 1 Unknown
    %8 = OpTypeSampledImage %7
    %9 = OpTypePointer UniformConstant %8
   %10 = OpVariable %9 UniformConstant
   %11 = OpTypeInt 32 1
   %12 = OpTypeImage %11 Buffer 0 0 0 1 Unknown
   %13 = OpTypeSampledImage %12
   %14 = OpTypePointer UniformConstant %13
   %15 = OpVariable %14 UniformConstant
   %16 = OpTypeInt 32 0
   %17 = OpTypeImage %16 Buffer 0 0 0 1 Unknown
   %18 = OpTypeSampledImage %17
   %19 = OpTypePointer UniformConstant %18
   %20 = OpVariable %19 UniformConstant

         ...
```

```c
typedef struct VkPushConstantRange {
    VkShaderStageFlags      stageFlags;
    uint32_t                offset;
    uint32_t                size;
} VkPushConstantRange;
```

```glsl
layout (push_constant) uniform my_push_constants_t
{
    int bourbon;
    int scotch;
    int beer;
} my_push_constants;
```

```
        OpCapability Shader
%1 = OpExtInstImport "GLSL.std.450"
        OpMemoryModel Logical GLSL450
        OpEntryPoint GLCompute %4 "main"
        OpExecutionMode %4 LocalSize 1 1 1
        OpSource GLSL 450
        OpName %4 "main"
        ;; Name the push constant block and its members.
        OpName %7 "my_push_constants_t"
        OpMemberName %7 0 "bourbon"
        OpMemberName %7 1 "scotch"
        OpMemberName %7 2 "beer"
        OpName %9 "my_push_constants"
        ;; Assign offsets to the members of the push constant block.
        OpMemberDecorate %7 0 Offset 0
        OpMemberDecorate %7 1 Offset 4
        OpMemberDecorate %7 2 Offset 8
        OpDecorate %7 Block
%2 = OpTypeVoid
%3 = OpTypeFunction %2
%6 = OpTypeInt 32 1
%7 = OpTypeStruct %6 %6 %6
        ;; Declare the push constant block itself.
%8 = OpTypePointer PushConstant %7
%9 = OpVariable %8 PushConstant
        ...
```

```
void vkCmdPushConstants (
    VkCommandBuffer                     commandBuffer,
    VkPipelineLayout                    layout,
    VkShaderStageFlags                  stageFlags,
    uint32_t                            offset,
    uint32_t                            size,
    const void*                         pValues);
```

```c
VkResult vkCreateSampler (
    VkDevice                            device,
    const VkSamplerCreateInfo*          pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkSampler*                          pSampler);
```

```c
typedef struct VkSamplerCreateInfo {
    VkStructureType             sType;
    const void*                 pNext;
    VkSamplerCreateFlags        flags;
    VkFilter                    magFilter;
    VkFilter                    minFilter;
    VkSamplerMipmapMode         mipmapMode;
    VkSamplerAddressMode        addressModeU;
    VkSamplerAddressMode        addressModeV;
    VkSamplerAddressMode        addressModeW;
    float                       mipLodBias;
    VkBool32                    anisotropyEnable;
    float                       maxAnisotropy;
    VkBool32                    compareEnable;
    VkCompareOp                 compareOp;
    float                       minLod;
    float                       maxLod;
    VkBorderColor               borderColor;
    VkBool32                    unnormalizedCoordinates;
} VkSamplerCreateInfo;
```

```
void vkDestroySampler (
    VkDevice                           device,
    VkSampler                          sampler,
    const VkAllocationCallbacks*       pAllocator);
```

```
void vkCmdDraw (
    VkCommandBuffer                          commandBuffer,
    uint32_t                                 vertexCount,
    uint32_t                                 instanceCount,
    uint32_t                                 firstVertex,
    uint32_t                                 firstInstance);
```

```
VkResult vkCreateRenderPass (
    VkDevice                             device,
    const VkRenderPassCreateInfo*        pCreateInfo,
    const VkAllocationCallbacks*         pAllocator,
    VkRenderPass*                        pRenderPass);
```

```c
typedef struct VkRenderPassCreateInfo {
    VkStructureType                     sType;
    const void*                         pNext;
    VkRenderPassCreateFlags             flags;
    uint32_t                            attachmentCount;
    const VkAttachmentDescription*      pAttachments;
    uint32_t                            subpassCount;
    const VkSubpassDescription*         pSubpasses;
    uint32_t                            dependencyCount;
    const VkSubpassDependency*          pDependencies;
} VkRenderPassCreateInfo;
```

```c
typedef struct VkAttachmentDescription {
    VkAttachmentDescriptionFlags    flags;
    VkFormat                        format;
    VkSampleCountFlagBits           samples;
    VkAttachmentLoadOp              loadOp;
    VkAttachmentStoreOp             storeOp;
    VkAttachmentLoadOp              stencilLoadOp;
    VkAttachmentStoreOp             stencilStoreOp;
    VkImageLayout                   initialLayout;
    VkImageLayout                   finalLayout;
} VkAttachmentDescription;
```

```c
typedef struct VkSubpassDescription {
    VkSubpassDescriptionFlags       flags;
    VkPipelineBindPoint             pipelineBindPoint;
    uint32_t                        inputAttachmentCount;
    const VkAttachmentReference*    pInputAttachments;
    uint32_t                        colorAttachmentCount;
    const VkAttachmentReference*    pColorAttachments;
    const VkAttachmentReference*    pResolveAttachments;
    const VkAttachmentReference*    pDepthStencilAttachment;
    uint32_t                        preserveAttachmentCount;
    const uint32_t*                 pPreserveAttachments;
} VkSubpassDescription;
```

```
typedef struct VkAttachmentReference {
    uint32_t            attachment;
    VkImageLayout       layout;
} VkAttachmentReference;
```

```c
typedef struct VkSubpassDependency {
    uint32_t                    srcSubpass;
    uint32_t                    dstSubpass;
    VkPipelineStageFlags        srcStageMask;
    VkPipelineStageFlags        dstStageMask;
    VkAccessFlags               srcAccessMask;
    VkAccessFlags               dstAccessMask;
    VkDependencyFlags           dependencyFlags;
} VkSubpassDependency;
```

```
// This is our color attachment. It's an R8G8B8A8_UNORM single sample image.
// We want to clear it at the start of the renderpass and save the contents
// when we're done. It starts in UNDEFINED layout, which is a key to
// Vulkan that it's allowed to throw the old content away, and we want to
// leave it in COLOR_ATTACHMENT_OPTIMAL state when we're done.
static const VkAttachmentDescription attachments[] =
{
    {
        0,                                          // flags
        VK_FORMAT_R8G8B8A8_UNORM,                   // format
        VK_SAMPLE_COUNT_1_BIT,                      // samples
        VK_ATTACHMENT_LOAD_OP_CLEAR,                // loadOp
        VK_ATTACHMENT_STORE_OP_STORE,               // storeOp
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,            // stencilLoadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,           // stencilStoreOp
        VK_IMAGE_LAYOUT_UNDEFINED,                  // initialLayout
        VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL    // finalLayout
    }
};

// This is the single reference to our single attachment.
static const VkAttachmentReference attachmentReferences[] =
{
    {
        0,                                          // attachment
        VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL    // layout
    }
};

// There is one subpass in this renderpass, with only a reference to the
// single output attachment.
static const VkSubpassDescription subpasses[] =
{
    {
```

```
            0,                                            // flags
            VK_PIPELINE_BIND_POINT_GRAPHICS,              // pipelineBindPoint
            0,                                            // inputAttachmentCount
            nullptr,                                      // pInputAttachments
            1,                                            // colorAttachmentCount
            &attachmentReferences[0],                     // pColorAttachments
            nullptr,                                      // pResolveAttachments
            nullptr,                                      // pDepthStencilAttachment
            0,                                            // preserveAttachmentCount
            nullptr                                       // pPreserveAttachments
    }
};

// Finally, this is the information that Vulkan needs to create the renderpass
// object.
static VkRenderPassCreateInfo renderpassCreateInfo =
{
    VK_STRUCTURE_TYPE_RENDER_PASS_CREATE_INFO,         // sType
    nullptr,                                           // pNext
    0,                                                 // flags
    1,                                                 // attachmentCount
    &attachments[0],                                   // pAttachments
    1,                                                 // subpassCount
    &subpasses[0],                                     // pSubpasses
    0,                                                 // dependencyCount
    nullptr                                            // pDependencies
};

VkRenderPass renderpass = VK_NULL_HANDLE;

// The only code that actually executes is this single call, which creates
// the renderpass object.
vkCreateRenderPass(device,
                   &renderpassCreateInfo,
                   nullptr,
                   &renderpass);
```

```c
void vkDestroyRenderPass (
    VkDevice                        device,
    VkRenderPass                    renderPass,
    const VkAllocationCallbacks*    pAllocator);
```

```
VkResult vkCreateFramebuffer (
    VkDevice                             device,
    const VkFramebufferCreateInfo*       pCreateInfo,
    const VkAllocationCallbacks*         pAllocator,
    VkFramebuffer*                       pFramebuffer);
```

```c
typedef struct VkFramebufferCreateInfo {
    VkStructureType                 sType;
    const void*                     pNext;
    VkFramebufferCreateFlags        flags;
    VkRenderPass                    renderPass;
    uint32_t                        attachmentCount;
    const VkImageView*              pAttachments;
    uint32_t                        width;
    uint32_t                        height;
    uint32_t                        layers;
} VkFramebufferCreateInfo;
```

```c
void vkDestroyFramebuffer (
    VkDevice                        device,
    VkFramebuffer                   framebuffer,
    const VkAllocationCallbacks*    pAllocator);
```

```
VkResult vkCreateGraphicsPipelines (
    VkDevice                              device,
    VkPipelineCache                       pipelineCache,
    uint32_t                              createInfoCount,
    const VkGraphicsPipelineCreateInfo*   pCreateInfos,
    const VkAllocationCallbacks*          pAllocator,
    VkPipeline*                           pPipelines);
```

```c
typedef struct VkGraphicsPipelineCreateInfo {
    VkStructureType                                 sType;
    const void*                                     pNext;
    VkPipelineCreateFlags                           flags;
    uint32_t                                        stageCount;
    const VkPipelineShaderStageCreateInfo*          pStages;
    const VkPipelineVertexInputStateCreateInfo*     pVertexInputState;
    const VkPipelineInputAssemblyStateCreateInfo*   pInputAssemblyState;
    const VkPipelineTessellationStateCreateInfo*    pTessellationState;
    const VkPipelineViewportStateCreateInfo*        pViewportState;
    const VkPipelineRasterizationStateCreateInfo*   pRasterizationState;
    const VkPipelineMultisampleStateCreateInfo*     pMultisampleState;
    const VkPipelineDepthStencilStateCreateInfo*    pDepthStencilState;
    const VkPipelineColorBlendStateCreateInfo*      pColorBlendState;
    const VkPipelineDynamicStateCreateInfo*         pDynamicState;
    VkPipelineLayout                                layout;
    VkRenderPass                                    renderPass;
    uint32_t                                        subpass;
    VkPipeline                                      basePipelineHandle;
    int32_t                                         basePipelineIndex;
} VkGraphicsPipelineCreateInfo;
```

```c
typedef struct VkPipelineShaderStageCreateInfo {
    VkStructureType                     sType;
    const void*                         pNext;
    VkPipelineShaderStageCreateFlags    flags;
    VkShaderStageFlagBits               stage;
    VkShaderModule                      module;
    const char*                         pName;
    const VkSpecializationInfo*         pSpecializationInfo;
} VkPipelineShaderStageCreateInfo;
```

```cpp
VkPipelineShaderStageCreateInfo shaderStageCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_SHADER_STAGE_CREATE_INFO,        // sType
    nullptr,                                    // pNext
    0,                                          // flags
    VK_SHADER_STAGE_VERTEX_BIT,                 // stage
    module,                                     // module
    "main",                                     // pName
    nullptr                                     // pSpecializationInfo
};

static const
VkPipelineVertexInputStateCreateInfo vertexInputStateCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_VERTEX_INPUT_STATE_CREATE_INFO,  // sType
    nullptr,                                    // pNext
    0,                                          // flags
    0,                                          // vertexBindingDescriptionCount
    nullptr,                                    // pVertexBindingDescriptions
    0,                                          // vertexAttributeDescriptionCount
    nullptr                                     // pVertexAttributeDescriptions
};

static const
VkPipelineInputAssemblyStateCreateInfo inputAssemblyStateCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_INPUT_ASSEMBLY_STATE_CREATE_INFO,// sType
    nullptr,                                    // pNext
    0,                                          // flags
    VK_PRIMITIVE_TOPOLOGY_POINT_LIST,           // topology
    VK_FALSE                                    // primitiveRestartEnable
};
```

```
static const
VkViewport dummyViewport =
{
    0.0f, 0.0f,                                 // x, y
    1.0f, 1.0f,                                 // width, height
    0.1f, 1000.0f                               // minDepth, maxDepth
};

static const
VkRect2D dummyScissor =
{
    { 0, 0 },                                   // offset
    { 1, 1 }                                    // extent
};

static const
VkPipelineViewportStateCreateInfo viewportStateCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_VIEWPORT_STATE_CREATE_INFO,       // sType
    nullptr,                                    // pNext
    0,                                          // flags
    1,                                          // viewportCount
    &dummyViewport,                             // pViewports
    1,                                          // scissorCount
    &dummyScissor                               // pScissors
};

static const
VkPipelineRasterizationStateCreateInfo rasterizationStateCreateInfo =
{
```

```
    VK_STRUCTURE_TYPE_PIPELINE_RASTERIZATION_STATE_CREATE_INFO, // sType
    nullptr,                                    // pNext
    0,                                          // flags
    VK_FALSE,                                   // depthClampEnable
    VK_TRUE,                                    // rasterizerDiscardEnable
    VK_POLYGON_MODE_FILL,                       // polygonMode
    VK_CULL_MODE_NONE,                          // cullMode
    VK_FRONT_FACE_COUNTER_CLOCKWISE,            // frontFace
    VK_FALSE,                                   // depthBiasEnable
    0.0f,                                       // depthBiasConstantFactor
    0.0f,                                       // depthBiasClamp
    0.0f,                                       // depthBiasSlopeFactor
    0.0f                                        // lineWidth
};

static const
VkGraphicsPipelineCreateInfo graphicsPipelineCreateInfo =
{
    VK_STRUCTURE_TYPE_GRAPHICS_PIPELINE_CREATE_INFO,    // sType
    nullptr,                                    // pNext
    0,                                          // flags
    1,                                          // stageCount
    &shaderStageCreateInfo,                     // pStages
    &vertexInputStateCreateInfo,                // pVertexInputState
    &inputAssemblyStateCreateInfo,              // pInputAssemblyState
    nullptr,                                    // pTessellationState
    &viewportStateCreateInfo,                   // pViewportState
    &rasterizationStateCreateInfo,              // pRasterizationState
    nullptr,                                    // pMultisampleState
    nullptr,                                    // pDepthStencilState
    nullptr,                                    // pColorBlendState
    nullptr,                                    // pDynamicState
    VK_NULL_HANDLE,                             // layout
    renderpass,                                 // renderPass
    0,                                          // subpass
    VK_NULL_HANDLE,                             // basePipelineHandle
    0,                                          // basePipelineIndex
};

result = vkCreateGraphicsPipelines(device,
                                   VK_NULL_HANDLE,
                                   1,
                                   &graphicsPipelineCreateInfo,
                                   nullptr,
                                   &pipeline);
```

```c
typedef struct VkPipelineVertexInputStateCreateInfo {
    VkStructureType                             sType;
    const void*                                 pNext;
    VkPipelineVertexInputStateCreateFlags       flags;
    uint32_t                                    vertexBindingDescriptionCount;
    const VkVertexInputBindingDescription*      pVertexBindingDescriptions;
    uint32_t                                    vertexAttributeDescriptionCount;
    const VkVertexInputAttributeDescription*    pVertexAttributeDescriptions;
} VkPipelineVertexInputStateCreateInfo;
```

```c
typedef struct VkVertexInputBindingDescription {
    uint32_t                binding;
    uint32_t                stride;
    VkVertexInputRate       inputRate;
} VkVertexInputBindingDescription;
```

```c
typedef struct VkVertexInputAttributeDescription {
    uint32_t    location;
    uint32_t    binding;
    VkFormat    format;
    uint32_t    offset;
} VkVertexInputAttributeDescription;
```

```cpp
typedef struct vertex_t
{
    vmath::vec4 position;
    vmath::vec3 normal;
    vmath::vec2 texcoord;
} vertex;

static const
VkVertexInputBindingDescription vertexInputBindings[] =
{
    { 0, sizeof(vertex), VK_VERTEX_INPUT_RATE_VERTEX }        // Buffer
};

static const
VkVertexInputAttributeDescription vertexAttributes[] =
{
    { 0, 0, VK_FORMAT_R32G32B32A32_SFLOAT, 0 },                       // Position
    { 1, 0, VK_FORMAT_R32G32B32_SFLOAT, offsetof(vertex, normal) },   // Normal
    { 2, 0, VK_FORMAT_R32G32_SFLOAT, offsetof(vertex, texcoord) }     // Tex Coord
};

static const
VkPipelineVertexInputStateCreateInfo vertexInputStateCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_VERTEX_INPUT_STATE_CREATE_INFO,  // sType
    nullptr,                                                    // pNext
    0,                                                          // flags
    vkcore::utils::arraysize(vertexInputBindings),   // vertexBindingDescriptionCount
    vertexInputBindings,                             // pVertexBindingDescriptions
    vkcore::utils::arraysize(vertexAttributes),      // vertexAttributeDescriptionCount
    vertexAttributes                                 // pVertexAttributeDescriptions
};
```

```glsl
#version 450 core

layout (location = 0) in vec3 i_position;
layout (location = 1) in vec2 i_uv;

void main(void)
{
    gl_Position = vec4(i_position, 1.0f);
}
```

```
; SPIR-V
; Version: 1.0
; Generator: Khronos Glslang Reference Front End; 1
; Bound: 30
; Schema: 0
        OpCapability Shader
  %1 = OpExtInstImport "GLSL.std.450"
        OpMemoryModel Logical GLSL450
        OpEntryPoint Vertex %4 "main" %13 %18 %29
        OpSource GLSL 450
        OpName %18 "i_position"        ;; Name of i_position
        OpName %29 "i_uv"              ;; Name of i_uv
        OpDecorate %18 Location 0      ;; Location of i_position
        OpDecorate %29 Location 1      ;; Location of i_uv
...
  %6 = OpTypeFloat 32                  ;; %6 is 32-bit floating-point type
 %16 = OpTypeVector %6 3              ;; %16 is a vector of 3 32-bit floats (vec3)
 %17 = OpTypePointer Input %16
 %18 = OpVariable %17 Input           ;; %18 is i_position - input pointer to vec3
 %27 = OpTypeVector %6 2              ;; %27 is a vector of 2 32-bit floats
 %28 = OpTypePointer Input %27
 %29 = OpVariable %28 Input           ;; %29 is i_uv - input pointer to vec2
```

```c
typedef struct VkPipelineInputAssemblyStateCreateInfo {
    VkStructureType                             sType;
    const void*                                 pNext;
    VkPipelineInputAssemblyStateCreateFlags     flags;
    VkPrimitiveTopology                         topology;
    VkBool32                                    primitiveRestartEnable;
} VkPipelineInputAssemblyStateCreateInfo;
```

```c
typedef struct VkPipelineTessellationStateCreateInfo {
    VkStructureType                        sType;
    const void*                            pNext;
    VkPipelineTessellationStateCreateFlags flags;
    uint32_t                               patchControlPoints;
} VkPipelineTessellationStateCreateInfo;
```

```c
typedef struct VkPipelineViewportStateCreateInfo {
    VkStructureType                        sType;
    const void*                            pNext;
    VkPipelineViewportStateCreateFlags     flags;
    uint32_t                               viewportCount;
    const VkViewport*                      pViewports;
    uint32_t                               scissorCount;
    const VkRect2D*                        pScissors;
} VkPipelineViewportStateCreateInfo;
```

```c
typedef struct VkPipelineRasterizationStateCreateInfo {
    VkStructureType                            sType;
    const void*                                pNext;
    VkPipelineRasterizationStateCreateFlags    flags;
    VkBool32                                   depthClampEnable;
    VkBool32                                   rasterizerDiscardEnable;
    VkPolygonMode                              polygonMode;
    VkCullModeFlags                            cullMode;
    VkFrontFace                                frontFace;
    VkBool32                                   depthBiasEnable;
    float                                      depthBiasConstantFactor;
    float                                      depthBiasClamp;
    float                                      depthBiasSlopeFactor;
    float                                      lineWidth;
} VkPipelineRasterizationStateCreateInfo;
```

```c
typedef struct VkPipelineMultisampleStateCreateInfo {
    VkStructureType                           sType;
    const void*                               pNext;
    VkPipelineMultisampleStateCreateFlags     flags;
    VkSampleCountFlagBits                     rasterizationSamples;
    VkBool32                                  sampleShadingEnable;
    float                                     minSampleShading;
    const VkSampleMask*                       pSampleMask;
    VkBool32                                  alphaToCoverageEnable;
    VkBool32                                  alphaToOneEnable;
} VkPipelineMultisampleStateCreateInfo;
```

```c
typedef struct VkPipelineDepthStencilStateCreateInfo {
    VkStructureType                        sType;
    const void*                            pNext;
    VkPipelineDepthStencilStateCreateFlags flags;
    VkBool32                               depthTestEnable;
    VkBool32                               depthWriteEnable;
    VkCompareOp                            depthCompareOp;
    VkBool32                               depthBoundsTestEnable;
    VkBool32                               stencilTestEnable;
    VkStencilOpState                       front;
    VkStencilOpState                       back;
    float                                  minDepthBounds;
    float                                  maxDepthBounds;
} VkPipelineDepthStencilStateCreateInfo;
```

```c
typedef struct VkPipelineColorBlendStateCreateInfo {
    VkStructureType                              sType;
    const void*                                  pNext;
    VkPipelineColorBlendStateCreateFlags         flags;
    VkBool32                                     logicOpEnable;
    VkLogicOp                                    logicOp;
    uint32_t                                     attachmentCount;
    const VkPipelineColorBlendAttachmentState*   pAttachments;
    float                                        blendConstants[4];
} VkPipelineColorBlendStateCreateInfo;
```

```c
typedef struct VkPipelineColorBlendAttachmentState {
    VkBool32                    blendEnable;
    VkBlendFactor               srcColorBlendFactor;
    VkBlendFactor               dstColorBlendFactor;
    VkBlendOp                   colorBlendOp;
    VkBlendFactor               srcAlphaBlendFactor;
    VkBlendFactor               dstAlphaBlendFactor;
    VkBlendOp                   alphaBlendOp;
    VkColorComponentFlags       colorWriteMask;
} VkPipelineColorBlendAttachmentState;
```

```c
typedef struct VkPipelineDynamicStateCreateInfo {
    VkStructureType                     sType;
    const void*                         pNext;
    VkPipelineDynamicStateCreateFlags   flags;
    uint32_t                            dynamicStateCount;
    const VkDynamicState*               pDynamicStates;
} VkPipelineDynamicStateCreateInfo;
```

```
void vkCmdDraw (
    VkCommandBuffer                     commandBuffer,
    uint32_t                           vertexCount,
    uint32_t                           instanceCount,
    uint32_t                           firstVertex,
    uint32_t                           firstInstance);
```

```
void vkCmdBeginRenderPass (
    VkCommandBuffer                              commandBuffer,
    const VkRenderPassBeginInfo*                 pRenderPassBegin,
    VkSubpassContents                            contents);
```

```c
typedef struct VkRenderPassBeginInfo {
    VkStructureType         sType;
    const void*             pNext;
    VkRenderPass            renderPass;
    VkFramebuffer           framebuffer;
    VkRect2D                renderArea;
    uint32_t                clearValueCount;
    const VkClearValue*     pClearValues;
} VkRenderPassBeginInfo;
```

```
typedef union VkClearValue {
    VkClearColorValue           color;
    VkClearDepthStencilValue    depthStencil;
} VkClearValue;
```

```c
typedef union VkClearColorValue {
    float           float32[4];
    int32_t         int32[4];
    uint32_t        uint32[4];
} VkClearColorValue;
```

```
typedef struct VkClearDepthStencilValue {
    float          depth;
    uint32_t       stencil;
} VkClearDepthStencilValue;
```

```c
void vkCmdEndRenderPass (
    VkCommandBuffer                                    commandBuffer);
```

```c
void vkCmdBindVertexBuffers (
    VkCommandBuffer                     commandBuffer,
    uint32_t                            firstBinding,
    uint32_t                            bindingCount,
    const VkBuffer*                     pBuffers,
    const VkDeviceSize*                 pOffsets);
```

```cpp
typedef struct vertex_t
{
    vmath::vec3 normal;
    vmath::vec2 texcoord;
} vertex;

static const
VkVertexInputBindingDescription vertexInputBindings[] =
{
    { 0, sizeof(vmath::vec4), VK_VERTEX_INPUT_RATE_VERTEX }, // Buffer 1
    { 1, sizeof(vertex), VK_VERTEX_INPUT_RATE_VERTEX }      // Buffer 2
};

static const
VkVertexInputAttributeDescription vertexAttributes[] =
{
    { 0, 0, VK_FORMAT_R32G32B32A32_SFLOAT, 0 },             // Position
    { 1, 1, VK_FORMAT_R32G32B32_SFLOAT, 0 },                // Normal
    { 2, 1, VK_FORMAT_R32G32_SFLOAT, sizeof(vmath::vec3) }  // Tex Coord
};

static const
VkPipelineVertexInputStateCreateInfo vertexInputStateCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_VERTEX_INPUT_STATE_CREATE_INFO,  // sType
    nullptr,                                            // pNext
    0,                                                  // flags
    vkcore::utils::arraysize(vertexInputBindings),      // vertexBindingDescription-
                                                        // Count
    vertexInputBindings,                                // pVertexBinding-
                                                        // Descriptions
    vkcore::utils::arraysize(vertexAttributes),         // vertexAttribute-
                                                        // DescriptionCount
    vertexAttributes                                    // pVertexAttribute-
                                                        // Descriptions
};
```

```
void vkCmdDrawIndexed (
    VkCommandBuffer                         commandBuffer,
    uint32_t                                indexCount,
    uint32_t                                instanceCount,
    uint32_t                                firstIndex,
    int32_t                                 vertexOffset,
    uint32_t                                firstInstance);
```

```
void vkCmdBindIndexBuffer (
    VkCommandBuffer                           commandBuffer,
    VkBuffer                                  buffer,
    VkDeviceSize                              offset,
    VkIndexType                               indexType);
```

```
offset + firstIndex * sizeof(index)
```

```c
// Raw, non-indexed data
static const float vertex_positions[] =
{
    -0.25f,  0.25f, -0.25f,
    -0.25f, -0.25f, -0.25f,
     0.25f, -0.25f, -0.25f,

     0.25f, -0.25f, -0.25f,
     0.25f,  0.25f, -0.25f,
    -0.25f,  0.25f, -0.25f,

     0.25f, -0.25f, -0.25f,
     0.25f, -0.25f,  0.25f,
     0.25f,  0.25f, -0.25f,

     0.25f, -0.25f,  0.25f,
     0.25f,  0.25f,  0.25f,
     0.25f,  0.25f, -0.25f,

     0.25f, -0.25f,  0.25f,
    -0.25f, -0.25f,  0.25f,
     0.25f,  0.25f,  0.25f,

    -0.25f, -0.25f,  0.25f,
    -0.25f,  0.25f,  0.25f,
     0.25f,  0.25f,  0.25f,

    -0.25f, -0.25f,  0.25f,
    -0.25f, -0.25f, -0.25f,
    -0.25f,  0.25f,  0.25f,

    -0.25f, -0.25f, -0.25f,
    -0.25f,  0.25f, -0.25f,
    -0.25f,  0.25f,  0.25f,

    -0.25f, -0.25f,  0.25f,
     0.25f, -0.25f,  0.25f,
     0.25f, -0.25f, -0.25f,
```

```cpp
     0.25f, -0.25f, -0.25f,
    -0.25f, -0.25f, -0.25f,
    -0.25f, -0.25f,  0.25f,

    -0.25f,  0.25f, -0.25f,
     0.25f,  0.25f, -0.25f,
     0.25f,  0.25f,  0.25f,

     0.25f,  0.25f,  0.25f,
    -0.25f,  0.25f,  0.25f,
    -0.25f,  0.25f, -0.25f
};

static const uint32_t vertex_count = sizeof(vertex_positions) /
                                      (3 * sizeof(float));

// Indexed vertex data
static const float indexed_vertex_positions[] =
{
    -0.25f, -0.25f, -0.25f,
    -0.25f,  0.25f, -0.25f,
     0.25f, -0.25f, -0.25f,
     0.25f,  0.25f, -0.25f,
     0.25f, -0.25f,  0.25f,
     0.25f,  0.25f,  0.25f,
    -0.25f, -0.25f,  0.25f,
    -0.25f,  0.25f,  0.25f,
};

// Index buffer
static const uint16_t vertex_indices[] =

    {
        0, 1, 2,
        2, 1, 3,
        2, 3, 4,
        4, 3, 5,
        4, 5, 6,
        6, 5, 7,
        6, 7, 0,
        0, 7, 1,
        6, 0, 2,
        2, 4, 6,
        7, 5, 3,
        7, 3, 1
    };

    static const uint32_t index_count = vkcore::utils::arraysize(vertex_indices);
```

```glsl
#version 450 core

vec3 unpackA2R10G10B10_snorm(uint value)
{
    int val_signed = int(value);
    vec3 result;
    const float scale = (1.0f / 512.0f);

    result.x = float(bitfieldExtract(val_signed, 20, 10));
    result.y = float(bitfieldExtract(val_signed, 10, 10));
    result.z = float(bitfieldExtract(val_signed, 0, 10));

    return result * scale;
}

void main(void)
{
    gl_Position = vec4(unpackA2R10G10B10_snorm(gl_VertexIndex), 1.0f);
}
```

```c
typedef struct VkPipelineInputAssemblyStateCreateInfo {
    VkStructureType                             sType;
    const void*                                 pNext;
    VkPipelineInputAssemblyStateCreateFlags     flags;
    VkPrimitiveTopology                         topology;
    VkBool32                                    primitiveRestartEnable;
} VkPipelineInputAssemblyStateCreateInfo;
```

```glsl
#version 450 core

layout (set = 0, binding = 0) uniform matrix_uniforms_b
{
    mat4 mvp_matrix[1024];
};

layout (set = 0, binding = 1) uniform color_uniforms_b
{

    vec4 cube_colors[1024];
};

layout (location = 0) in vec3 i_position;

out vs_fs
{
    flat vec4 color;
};

void main(void)
{
    float f = float(gl_VertexIndex / 6) / 6.0f;
    vec4 color1 = cube_colors[gl_InstanceIndex];
    vec4 color2 = cube_colors[gl_InstanceIndex & 512];

    color = mix(color1, color2, f);
    gl_Position = mvp_matrix[gl_InstanceIndex] * vec4(i_position, 1.0f);
}
```

```
void vkCmdDrawIndirect (
    VkCommandBuffer                     commandBuffer,
    VkBuffer                            buffer,
    VkDeviceSize                        offset,
    uint32_t                            drawCount,
    uint32_t                            stride);
```

```c
typedef struct VkDrawIndirectCommand {
    uint32_t    vertexCount;
    uint32_t    instanceCount;
    uint32_t    firstVertex;
    uint32_t    firstInstance;
} VkDrawIndirectCommand;
```

```
void vkCmdDrawIndexedIndirect (
    VkCommandBuffer                 commandBuffer,
    VkBuffer                        buffer,
    VkDeviceSize                    offset,
    uint32_t                        drawCount,
    uint32_t                        stride);
```

```c
typedef struct VkDrawIndexedIndirectCommand {
    uint32_t      indexCount;
    uint32_t      instanceCount;
    uint32_t      firstIndex;
    int32_t       vertexOffset;
    uint32_t      firstInstance;
} VkDrawIndexedIndirectCommand;
```

```glsl
#version 450 core

// Enable the GL_ARB_shader_draw_parameters extensions.
#extension GL_ARB_shader_draw_parameters : require

layout (location = 0) in vec3 position_3;

layout (set = 0, binding = 0) uniform FRAME_DATA
{
    mat4 view_matrix;
    mat4 proj_matrix;
    mat4 viewproj_matrix;
};

layout (set = 0, binding = 1) readonly buffer OBJECT_TRANSFORMS
{
    mat4 model_matrix[];
};

void main(void)
{
    // Extend input position to vec4.
    vec4 position = vec4(position_3, 1.0);

    // Compute per-object model-view matrix.
    mat4 mv_matrix = view_matrix * model_matrix[gl_DrawIDARB];

    // Output position using global projection matrix.
    gl_Position = proj_matrix * P;
}
```

```
typedef struct VkPipelineTessellationStateCreateInfo {
    VkStructureType                         sType;
    const void*                             pNext;
    VkPipelineTessellationStateCreateFlags  flags;
    uint32_t                                patchControlPoints;
} VkPipelineTessellationStateCreateInfo;
```

```
OpExecutionMode %n OutputVertices 9}
```

```glsl
#version 450 core

layout (vertices = 1) out;

void main(void)
{
    gl_TessLevelInner[0] = 7.0f;
    gl_TessLevelInner[1] = 8.0f;

    gl_TessLevelOuter[0] = 3.0f;
    gl_TessLevelOuter[1] = 4.0f;
    gl_TessLevelOuter[2] = 5.0f;
    gl_TessLevelOuter[3] = 6.0f;
}
```

```
;; Require tessellation capability; import GLSL450 constructs.
        OpCapability Tessellation
   %1 = OpExtInstImport "GLSL.std.450"
        OpMemoryModel Logical GLSL450
;; Define "main" as the entry point for a tessellation control shader.
        OpEntryPoint TessellationControl %5663 "main" %3290 %5448
;; Number of patch output vertices = 1
        OpExecutionMode %5663 OutputVertices 1
;; Decorate the tessellation level variables appropriately.
        OpDecorate %3290 Patch
        OpDecorate %3290 BuiltIn TessLevelInner
        OpDecorate %5448 Patch
        OpDecorate %5448 BuiltIn TessLevelOuter
;; Declare types used in this shader.
    %8 = OpTypeVoid
 %1282 = OpTypeFunction %8
   %13 = OpTypeFloat 32
   %11 = OpTypeInt 32 0
;; This is the declaration of the gl_TessLevelInner[2] variable.
 %2576 = OpConstant %11 2
  %549 = OpTypeArray %13 %2576
 %1186 = OpTypePointer Output %549
 %3290 = OpVariable %1186 Output
   %12 = OpTypeInt 32 1
 %2571 = OpConstant %12 0
 %1330 = OpConstant %13 7
  %650 = OpTypePointer Output %13
 %2574 = OpConstant %12 1
 %2807 = OpConstant %13 8
;; Declare the gl_TessLevelOuter[4] variable.
 %2582 = OpConstant %11 4
  %609 = OpTypeArray %13 %2582
```

```
    %1246 = OpTypePointer Output %609
    %5448 = OpVariable %1246 Output
;; Declare constants used for indexing into our output arrays and the
;; values written into those arrays.
    %2978 = OpConstant %13 3
    %2921 = OpConstant %13 4
    %2577 = OpConstant %12 2
    %1387 = OpConstant %13 5
    %2580 = OpConstant %12 3
    %2864 = OpConstant %13 6
;; Start of the main function
    %5663 = OpFunction %8 None %1282
   %23934 = OpLabel
;; Declare references to elements of the output arrays and write constants
;; into them.
    %6956 = OpAccessChain %650 %3290 %2571
            OpStore %6956 %1330
   %19732 = OpAccessChain %650 %3290 %2574
            OpStore %19732 %2807
   %19733 = OpAccessChain %650 %5448 %2571
            OpStore %19733 %2978
   %19734 = OpAccessChain %650 %5448 %2574
            OpStore %19734 %2921
   %19735 = OpAccessChain %650 %5448 %2577
            OpStore %19735 %1387
   %23304 = OpAccessChain %650 %5448 %2580
            OpStore %23304 %2864
;; End of main
            OpReturn
            OpFunctionEnd
```

```glsl
#version 450 core

layout (vertices = 4) out;
out float o_outputData[4];

void main(void)
{
    o_outputData[gl_InvocationId] = 19.0f;
}
```

```
;; Declare a tessellation control shader.
          OpCapability Tessellation
    %1 = OpExtInstImport "GLSL.std.450"
          OpMemoryModel Logical GLSL450
          OpEntryPoint TessellationControl %5663 "main" %3227 %4585
;; 4 output vertices per patch, declare InvocationId built-in
          OpExecutionMode %5663 OutputVertices 4
          OpDecorate %4585 BuiltIn InvocationId
    %8 = OpTypeVoid
 %1282 = OpTypeFunction %8
   %13 = OpTypeFloat 32
   %11 = OpTypeInt 32 0
 %2582 = OpConstant %11 4
  %549 = OpTypeArray %13 %2582
 %1186 = OpTypePointer Output %549
 %3227 = OpVariable %1186 Output
   %12 = OpTypeInt 32 1
;; This declares the InvocationId input.
  %649 = OpTypePointer Input %12
 %4585 = OpVariable %649 Input
   %37 = OpConstant %13 19
  %650 = OpTypePointer Output %13
;; Beginning of main
 %5663 = OpFunction %8 None %1282
%24683 = OpLabel
;; Load the invocation ID.
%20081 = OpLoad %12 %4585
;; Define a reference to output variable.
%13546 = OpAccessChain %650 %3227 %20081
;; Store into it at invocation ID.
          OpStore %13546 %37
;; End of main
          OpReturn
          OpFunctionEnd
```

```
#version 450 core

layout (quads) in;

void main(void)
{
    gl_Position = vec4(gl_TessCoord, 1.0);
}
```

```
;; This is a GLSL 450 tessellation shader; enable capabilities.
            OpCapability Tessellation
            OpCapability TessellationPointSize
            OpCapability ClipDistance
            OpCapability CullDistance
     %1 = OpExtInstImport "GLSL.std.450"
            OpMemoryModel Logical GLSL450
;; Declare the entry point and decorate it appropriately.
            OpEntryPoint TessellationEvaluation %5663 "main" %4930 %3944
            OpExecutionMode %5663 Quads
            OpExecutionMode %5663 SpacingEqual
            OpExecutionMode %5663 VertexOrderCcw
;; Declare GLSL built-in outputs.
            OpMemberDecorate %2935 0 BuiltIn Position
            OpMemberDecorate %2935 1 BuiltIn PointSize
            OpMemberDecorate %2935 2 BuiltIn ClipDistance
            OpMemberDecorate %2935 3 BuiltIn CullDistance
            OpDecorate %2935 Block
;; This is the decoration of gl_TessCoord.
            OpDecorate %3944 BuiltIn TessCoord
     %8 = OpTypeVoid
  %1282 = OpTypeFunction %8
    %13 = OpTypeFloat 32
    %29 = OpTypeVector %13 4
    %11 = OpTypeInt 32 0
  %2573 = OpConstant %11 1
   %554 = OpTypeArray %13 %2573
  %2935 = OpTypeStruct %29 %13 %554 %554
   %561 = OpTypePointer Output %2935
  %4930 = OpVariable %561 Output
    %12 = OpTypeInt 32 1
  %2571 = OpConstant %12 0
;; Vector of 3 components, as input, decorated with BuiltIn TessCoord
    %24 = OpTypeVector %13 3
   %661 = OpTypePointer Input %24
  %3944 = OpVariable %661 Input
```

```
    %138 = OpConstant %13 1
    %666 = OpTypePointer Output %29
   %5663 = OpFunction %8 None %1282
  %24987 = OpLabel
;; Read from gl_TessCoord.
  %17674 = OpLoad %24 %3944
;; Extract the three elements.
  %22014 = OpCompositeExtract %13 %17674 0
  %23496 = OpCompositeExtract %13 %17674 1
   %7529 = OpCompositeExtract %13 %17674 2
;; Construct the new vec4.
  %18260 = OpCompositeConstruct %29 %22014 %23496 %7529 %138
;; Write to gl_Postion.
  %12055 = OpAccessChain %666 %4930 %2571
           OpStore %12055 %18260

  ;; End of main
           OpReturn
           OpFunctionEnd
```

```cpp
struct PushConstantData
{
    vmath::mat4 mvp_matrix;
    float       displacement_scale;
};

static const VkDescriptorSetLayoutBinding descriptorBindings[] =
{
    // Only binding is a sampled image with combined sampler.
    {
        0,
        VK_DESCRIPTOR_TYPE_COMBINED_IMAGE_SAMPLER,
        1,
        VK_SHADER_STAGE_TESSELLATION_EVALUATION_BIT,
        nullptr
    }
};

static const VkDescriptorSetLayoutCreateInfo descriptorSetLayoutCreateInfo =
{
    VK_STRUCTURE_TYPE_DESCRIPTOR_SET_LAYOUT_CREATE_INFO, nullptr,
    0,
    vkcore::utils::arraysize(descriptorBindings),
    descriptorBindings
};

vkCreateDescriptorSetLayout(getDevice(),
                            &descriptorSetLayoutCreateInfo,
                            nullptr,
                            &m_descriptorSetLayout);
```

```cpp
// Define a push constant range.
static const VkPushConstantRange pushConstantRange[] =
{
    // Push data into the evaluation shader.
    {
        VK_SHADER_STAGE_TESSELLATION_EVALUATION_BIT,
        0,
        sizeof(PushConstantData)
    }
};

VkPipelineLayoutCreateInfo pipelineLayoutCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_LAYOUT_CREATE_INFO, nullptr,
    0,
    1,

    &m_descriptorSetLayout,
    vkcore::utils::arraysize(pushConstantRange),
    pushConstantRange
};

result = vkCreatePipelineLayout(getDevice(),
                                &pipelineLayoutCreateInfo,
                                nullptr,
                                &m_pipelineLayout);
```

```glsl
#version 450 core

void main(void)
{
    float x = float(gl_VertexIndex & 1) - 0.5f;
    float y = float(gl_VertexIndex & 2) * 0.5f - 0.5f;

    gl_Position = vec4(x, y, 0.0f, 1.0f);
}
```

```glsl
#version 450 core

layout (vertices = 4) out;

void main(void)
{
    if (gl_InvocationID == 0)
    {
        gl_TessLevelInner[0] = 64.0f;
        gl_TessLevelInner[1] = 64.0f;

        gl_TessLevelOuter[0] = 64.0f;
        gl_TessLevelOuter[1] = 64.0f;
        gl_TessLevelOuter[2] = 64.0f;
        gl_TessLevelOuter[3] = 64.0f;
    }
    gl_out[gl_InvocationID].gl_Position = gl_in[gl_InvocationID].gl_Position;
}
```

```glsl
#version 450 core

layout (quads, fractional_odd_spacing) in;

layout (push_constant) uniform push_constants_b
{
    mat4 mvp_matrix;
    float displacement_scale;
} push_constants;

layout (set = 0, binding = 0) uniform sampler2D texDisplacement;

void main(void)
{
    vec4 mid1 = mix(gl_in[0].gl_Position, gl_in[1].gl_Position, gl_TessCoord.x);
    vec4 mid2 = mix(gl_in[2].gl_Position, gl_in[3].gl_Position, gl_TessCoord.x);
    vec4 pos = mix(mid1, mid2, gl_TessCoord.y);
    float displacement = texture(texDisplacement, gl_TessCoord.xy).x;

    pos.z = displacement * push_constants.displacement_scale;

    gl_Position = push_constants.mvp_matrix * pos;
}
```

```cpp
VkPipelineTessellationStateCreateInfo tessellationStateCreateInfo =
{
    VK_STRUCTURE_TYPE_PIPELINE_TESSELLATION_STATE_CREATE_INFO,  // sType
    nullptr,                                                     // pNext
    0,                                                           // flags
    4                                                            // patchControlPoints
};
```

```glsl
#version 450 core

layout (triangles) in;
layout (triangle_strip) out;
layout (max_vertices = 3) out;

void main(void)
{
    // Do nothing.
}
```

```
;; This is a geometry shader written in GLSL 450.
        OpCapability Geometry
   %1 = OpExtInstImport "GLSL.std.450"
        OpMemoryModel Logical GLSL450
;; Declare the main entry point.
        OpEntryPoint Geometry %5663 "main"
;; Triangles input, triangle strip output, maximum vertex count is 3.
        OpExecutionMode %5663 Triangles
        OpExecutionMode %5663 Invocations 1
        OpExecutionMode %5663 OutputTriangleStrip
        OpExecutionMode %5663 OutputVertices 3
;; Start of main
   %8 = OpTypeVoid
%1282 = OpTypeFunction %8
%5663 = OpFunction %8 None %1282
%16103 = OpLabel
;; End of empty main
        OpReturn
        OpFunctionEnd
```

```
...
;; Decorate the first member of our block with the BuiltIn Position.
          OpMemberDecorate %1017 0 BuiltIn Position
          OpDecorate %1017 Block

...
;; Declare an array of structures, with a pointer to this array as an input.
    %1017 = OpTypeStruct %29
     %557 = OpTypeArray %1017 %2573
    %1194 = OpTypePointer Input %557
    %5305 = OpVariable %1194 Input
     %666 = OpTypePointer Input %29

...
;; Access the input using OpLoad.
    %7129 = OpAccessChain %666 %5305 %2571 %2571
   %15646 = OpLoad %29 %7129
```

```glsl
#version 450 core

layout (points) in;
layout (points) out;
layout (max_vertices = 1) out;

in gl_PerVertex
{
    vec4 gl_Position;
} gl_in[];

out gs_out
{
    vec4 color;
};

void main(void)
{
    gl_Position = gl_in[0].gl_Position;
    color = vec4(0.5, 0.1, 0.9, 1.0);

    EmitVertex();
}
```

```
              OpCapability Geometry
              OpCapability GeometryPointSize
              OpCapability ClipDistance
              OpCapability CullDistance
       %1 = OpExtInstImport "GLSL.std.450"
              OpMemoryModel Logical GLSL450
              OpEntryPoint Geometry %5663 "main" %22044 %5305 %4930
              OpExecutionMode %5663 InputPoints
              OpExecutionMode %5663 Invocations 1
              OpExecutionMode %5663 OutputPoints
              OpExecutionMode %5663 OutputVertices 1
              OpMemberDecorate %2935 0 BuiltIn Position
              OpMemberDecorate %2935 1 BuiltIn PointSize
              OpMemberDecorate %2935 2 BuiltIn ClipDistance
              OpMemberDecorate %2935 3 BuiltIn CullDistance
              OpDecorate %2935 Block
              OpMemberDecorate %1017 0 BuiltIn Position
              OpDecorate %1017 Block
              OpDecorate %1018 Block
       %8 = OpTypeVoid
    %1282 = OpTypeFunction %8
      %13 = OpTypeFloat 32
      %29 = OpTypeVector %13 4
      %11 = OpTypeInt 32 0
    %2573 = OpConstant %11 1
     %554 = OpTypeArray %13 %2573
    %2935 = OpTypeStruct %29 %13 %554 %554
     %561 = OpTypePointer Output %2935
   %22044 = OpVariable %561 Output
      %12 = OpTypeInt 32 1
    %2571 = OpConstant %12 0
    %1017 = OpTypeStruct %29
     %557 = OpTypeArray %1017 %2573
    %1194 = OpTypePointer Input %557
    %5305 = OpVariable %1194 Input
     %666 = OpTypePointer Input %29
```

```
  %667 = OpTypePointer Output %29
 %1018 = OpTypeStruct %29
 %1654 = OpTypePointer Output %1018
 %4930 = OpVariable %1654 Output
  %252 = OpConstant %13 0.5
 %2936 = OpConstant %13 0.1
 %1364 = OpConstant %13 0.9
  %138 = OpConstant %13 1
  %878 = OpConstantComposite %29 %252 %2936 %1364 %138
 %5663 = OpFunction %8 None %1282
%23915 = OpLabel
 %7129 = OpAccessChain %666 %5305 %2571 %2571
%15646 = OpLoad %29 %7129
%19981 = OpAccessChain %667 %22044 %2571
         OpStore %19981 %15646
%22639 = OpAccessChain %667 %4930 %2571
         OpStore %22639 %878
         OpEmitVertex
         OpReturn
         OpFunctionEnd
```

```glsl
#version 450 core

layout (triangles) in;
layout (triangle_strip, max_vertices = 6) out;

void main(void)
{
    int i, j;
    vec4 scale = vec4(1.0f, 1.0f, 1.0f, 1.0f);

    for (j = 0; j < 2; j++)
    {
        for (i = 0; i < 3; i++)
        {
            gl_Position = gl_in[i].gl_Position * scale;

            EmitVertex();
        }

        EndPrimitive();
        scale.xy = -scale.xy;
    }
}
```

```
OpExecutionMode %n Invocations 1
```

```
#version 450 core

layout (triangles, invocations = 2) in;
layout (triangle_strip, max_vertices = 3) out;

layout (set = 0, binding = 0) uniform
{
    mat4 projection;
    mat4 objectToWorld[2];
} transforms;

void main(void)
{
    int i;
    mat4 objectToWorld = transforms.objectToWorld[gl_InvocationID];
    mat4 objectToClip = objectToWorld * transforms.projection;

    for (i = 0; i < 3; i++)
    {
        gl_Position = gl_in[i].gl_Position * objectToClip;

        EmitVertex();
    }
}
```

```glsl
#version 450 core

layout (location = 0) in vec3 i_position;
layout (location = 1) in float i_pointSize;

void main(void)
{
    gl_Position = vec4(i_position, 1.0f);
    gl_PointSize = i_pointSize;
}
```

```
        ...
        ;; GLSL compiler automatically declares per-vertex output block.
        OpName %11 "gl_PerVertex"
        OpMemberName %11 0 "gl_Position"
        OpMemberName %11 1 "gl_PointSize"
        OpMemberName %11 2 "gl_ClipDistance"
        OpMemberName %11 3 "gl_CullDistance"
        OpName %13 ""
        ;; Naming inputs
        OpName %18 "i_position"
        OpName %29 "i_pointSize"
        ;; Decorating members of the default output block
        OpMemberDecorate %11 0 BuiltIn Position
        OpMemberDecorate %11 1 BuiltIn PointSize ;; gl_PointSize
        OpMemberDecorate %11 2 BuiltIn ClipDistance
        OpMemberDecorate %11 3 BuiltIn CullDistance
        OpDecorate %11 Block
        OpDecorate %18 Location 0
        OpDecorate %29 Location 1

        ...

 %4 = OpFunction %2 None %3
        ;; Start of "main"
 %5 = OpLabel
        ;; Load i_position.
%19 = OpLoad %16 %18
%21 = OpCompositeExtract %6 %19 0
%22 = OpCompositeExtract %6 %19 1
%23 = OpCompositeExtract %6 %19 2
        ;; Construct vec4 and write to gl_Position.
%24 = OpCompositeConstruct %7 %21 %22 %23 %20
%26 = OpAccessChain %25 %13 %15
        OpStore %26 %24
        ;; Load from i_pointSize (%29).
%30 = OpLoad %6 %29
        ;; Use access chain to dereference built-in output.
%32 = OpAccessChain %31 %13 %27
        ;; Store to output decorated with PointSize.
        OpStore %32 %30
        OpReturn
        OpFunctionEnd
```

```
void vkCmdSetLineWidth (
    VkCommandBuffer                              commandBuffer,
    float                                        lineWidth);
```

```glsl
#version 450 core

// Redeclare gl_ClipDistance to explicitly size it.
out float gl_ClipDistance[1];

layout (location = 0) in vec3 i_position;

// Push constant from which to assign clip distance
layout (push_constant) uniform push_constants_b
{
    float clip_distance[4];
} push_constant;

void main(void)
{
    gl_ClipDistance[0] = push_constant.clip_distance[0];
    gl_Position = vec4(i_position, 1.0f);
}
```

```
            OpCapability Shader
;; The shader requires the ClipDistance capability.
            OpCapability ClipDistance
       %1 = OpExtInstImport "GLSL.std.450"
            OpMemoryModel Logical GLSL450
            OpEntryPoint Vertex %4 "main" %13 %29
            OpSource GLSL 450
            OpName %4 "main"
            OpName %11 "gl_PerVertex"
            OpMemberName %11 0 "gl_Position"
            OpMemberName %11 1 "gl_PointSize"
;; Redeclaration of gl_ClipDistance built-in
            OpMemberName %11 2 "gl_ClipDistance"
            OpMemberName %11 3 "gl_CullDistance"
            OpName %13 ""
            OpName %19 "push_constants_b"
            OpMemberName %19 0 "clip_distance"
            OpName %21 "push_constant"
            OpName %29 "i_position"
            OpMemberDecorate %11 0 BuiltIn Position
            OpMemberDecorate %11 1 BuiltIn PointSize
;; Decorate the built-in variable as ClipDistance.
            OpMemberDecorate %11 2 BuiltIn ClipDistance
            OpMemberDecorate %11 3 BuiltIn CullDistance
            OpDecorate %11 Block
            OpDecorate %18 ArrayStride 4
            OpMemberDecorate %19 0 Offset 0
            OpDecorate %19 Block
            OpDecorate %29 Location 0
       %2 = OpTypeVoid
       %3 = OpTypeFunction %2
       %6 = OpTypeFloat 32
       %7 = OpTypeVector %6 4
       %8 = OpTypeInt 32 0
       %9 = OpConstant %8 1
      %10 = OpTypeArray %6 %9
;; This creates the built-in output structure containing gl_ClipDistance.
      %11 = OpTypeStruct %7 %6 %10 %10
      %12 = OpTypePointer Output %11
;; Instantiate the built-in outputs.
      %13 = OpVariable %12 Output
      ...
;; Beginning of main()
       %5 = OpLabel
;; Load from the push-constant array.
      %23 = OpAccessChain %22 %21 %16 %16
      %24 = OpLoad %6 %23
;; Store to clip distance.
      %26 = OpAccessChain %25 %13 %15 %16
            OpStore %26 %24
      ...
;; End of main()
            OpReturn
            OpFunctionEnd
```

```glsl
#version 450 core

out float gl_CullDistance[1];

layout (location = 0) in vec3 i_position;

layout (push_constant) uniform push_constants_b
{
    float cull_distance[4];
} push_constant;

void main(void)
{
    gl_CullDistance[0] = push_constant.cull_distance[0];
    gl_Position = vec4(i_position, 1.0f);
}
```

```
          OpCapability Shader
;; The shader requires the CullDistance capability.
          OpCapability CullDistance
     %1 = OpExtInstImport "GLSL.std.450"
          OpMemoryModel Logical GLSL450
          OpEntryPoint Vertex %4 "main" %13 %29
          OpSource GLSL 450
          OpName %4 "main"
          OpName %11 "gl_PerVertex"
          OpMemberName %11 0 "gl_Position"
          OpMemberName %11 1 "gl_PointSize"
          OpMemberName %11 2 "gl_ClipDistance"
;; Redeclaration of gl_CullDistance built-in
          OpMemberName %11 3 "gl_CullDistance"
          OpName %13 ""
          OpName %19 "push_constants_b"
          OpMemberName %19 0 "cull_distance"
          OpName %21 "push_constant"
          OpName %29 "i_position"
          OpMemberDecorate %11 0 BuiltIn Position
          OpMemberDecorate %11 1 BuiltIn PointSize
          OpMemberDecorate %11 2 BuiltIn ClipDistance
;; Decorate the built-in variable as CullDistance.
          OpMemberDecorate %11 3 BuiltIn CullDistance
          OpDecorate %11 Block
          OpDecorate %18 ArrayStride 4
          OpMemberDecorate %19 0 Offset 0
          OpDecorate %19 Block
          OpDecorate %29 Location 0
     %2 = OpTypeVoid
     %3 = OpTypeFunction %2
     %6 = OpTypeFloat 32
     %7 = OpTypeVector %6 4
     %8 = OpTypeInt 32 0
     %9 = OpConstant %8 1
    %10 = OpTypeArray %6 %9
;; This creates the built-in output structure containing gl_CullDistance.
    %11 = OpTypeStruct %7 %6 %10 %10
    %12 = OpTypePointer Output %11
;; Instantiate the built-in outputs.
    %13 = OpVariable %12 Output
    ...
;; Beginning of main()
     %5 = OpLabel
;; Load from the push-constant array.
    %23 = OpAccessChain %22 %21 %16 %16
    %24 = OpLoad %6 %23
;; Store to cull distance.
    %26 = OpAccessChain %25 %13 %15 %16
    ...
;; End of main()
          OpReturn
          OpFunctionEnd
```

```
void vkCmdSetViewport (
    VkCommandBuffer                         commandBuffer,
    uint32_t                                firstViewport,
    uint32_t                                viewportCount,
    const VkViewport*                       pViewports);
```

```glsl
#version 450 core

layout (triangles, invocations = 4) in;
layout (triangle_strip, max_vertices = 3) out;

layout (set = 0, binding = 0) uniform transform_block
{
    mat4 mvp_matrix[4];
};

in VS_OUT
{
    vec4 color;
} gs_in[];

out GS_OUT
{
    vec4 color;
} gs_out;

void main(void)
{
    for (int i = 0; i < gl_in.length(); i++)

    {
        gs_out.color = gs_in[i].color;
        gl_Position = mvp_matrix[gl_InvocationID] *
                        gl_in[i].gl_Position;
        gl_ViewportIndex = gl_InvocationID;
        EmitVertex();
    }
    EndPrimitive();
}
```

```c
void vkCmdSetScissor (
    VkCommandBuffer                             commandBuffer,
    uint32_t                                    firstScissor,
    uint32_t                                    scissorCount,
    const VkRect2D*                             pScissors);
```

```c
typedef struct VkPipelineDepthStencilStateCreateInfo {
    VkStructureType                           sType;
    const void*                               pNext;
    VkPipelineDepthStencilStateCreateFlags    flags;
    VkBool32                                  depthTestEnable;
    VkBool32                                  depthWriteEnable;
    VkCompareOp                               depthCompareOp;
    VkBool32                                  depthBoundsTestEnable;
    VkBool32                                  stencilTestEnable;
    VkStencilOpState                          front;
    VkStencilOpState                          back;
    float                                     minDepthBounds;
    float                                     maxDepthBounds;
} VkPipelineDepthStencilStateCreateInfo;
```

```
void vkCmdSetDepthBounds (
    VkCommandBuffer                                  commandBuffer,
    float                                            minDepthBounds,
    float                                            maxDepthBounds);
```

```c
typedef struct VkPipelineRasterizationStateCreateInfo {
    VkStructureType                            sType;
    const void*                                pNext;
    VkPipelineRasterizationStateCreateFlags    flags;
    VkBool32                                   depthClampEnable;
    VkBool32                                   rasterizerDiscardEnable;
    VkPolygonMode                              polygonMode;
    VkCullModeFlags                            cullMode;
    VkFrontFace                                frontFace;
    VkBool32                                   depthBiasEnable;
    float                                      depthBiasConstantFactor;
    float                                      depthBiasClamp;
    float                                      depthBiasSlopeFactor;
    float                                      lineWidth;
} VkPipelineRasterizationStateCreateInfo;
```

```
void vkCmdSetDepthBias (
    VkCommandBuffer                              commandBuffer,
    float                                        depthBiasConstantFactor,
    float                                        depthBiasClamp,
    float                                        depthBiasSlopeFactor);
```

```c
typedef struct VkStencilOpState {
    VkStencilOp     failOp;
    VkStencilOp     passOp;
    VkStencilOp     depthFailOp;
    VkCompareOp     compareOp;
    uint32_t        compareMask;
    uint32_t        writeMask;
    uint32_t        reference;
} VkStencilOpState;
```

```
void vkCmdSetStencilReference (
    VkCommandBuffer                              commandBuffer,
    VkStencilFaceFlags                           faceMask,
    uint32_t                                     reference);


void vkCmdSetStencilCompareMask (
    VkCommandBuffer                              commandBuffer,
    VkStencilFaceFlags                           faceMask,
    uint32_t                                     compareMask);


void vkCmdSetStencilWriteMask (
    VkCommandBuffer                              commandBuffer,
    VkStencilFaceFlags                           faceMask,
    uint32_t                                     writeMask);
```

```
VkResult vkGetPhysicalDeviceImageFormatProperties (
    VkPhysicalDevice                        physicalDevice,
    VkFormat                                format,
    VkImageType                             type,
    VkImageTiling                           tiling,
    VkImageUsageFlags                       usage,
    VkImageCreateFlags                      flags,
    VkImageFormatProperties*                pImageFormatProperties);
```

```
void vkCmdResolveImage (
    VkCommandBuffer                          commandBuffer,
    VkImage                                  srcImage,
    VkImageLayout                            srcImageLayout,
    VkImage                                  dstImage,
    VkImageLayout                            dstImageLayout,
    uint32_t                                 regionCount,
    const VkImageResolve*                    pRegions);
```

```c
typedef struct VkImageResolve {
    VkImageSubresourceLayers    srcSubresource;
    VkOffset3D                  srcOffset;
    VkImageSubresourceLayers    dstSubresource;
    VkOffset3D                  dstOffset;
    VkExtent3D                  extent;
} VkImageResolve;
```

```c
typedef struct VkPipelineColorBlendStateCreateInfo {
    VkStructureType                             sType;
    const void*                                 pNext;
    VkPipelineColorBlendStateCreateFlags        flags;
    VkBool32                                    logicOpEnable;
    VkLogicOp                                   logicOp;
    uint32_t                                    attachmentCount;
    const VkPipelineColorBlendAttachmentState*  pAttachments;
    float                                       blendConstants[4];
} VkPipelineColorBlendStateCreateInfo;
```

```glsl
#version 450 core

out vec4 o_color;
void main(void)
{
    o_color = vec4(1.0f, 0.0f, 0.0f, 1.0f);
}
```

```
; SPIR-V
; Version: 1.0
; Generator: Khronos Glslang Reference Front End; 1
; Bound: 13
; Schema: 0
               OpCapability Shader
          %1 = OpExtInstImport "GLSL.std.450"
               OpMemoryModel Logical GLSL450
               OpEntryPoint Fragment %4 "main" %9
               OpExecutionMode %4 OriginUpperLeft
               OpSource GLSL 450
               OpName %4 "main"
               OpName %9 "o_color"
          %2 = OpTypeVoid
          %3 = OpTypeFunction %2
          %6 = OpTypeFloat 32
          %7 = OpTypeVector %6 4
          %8 = OpTypePointer Output %7
          %9 = OpVariable %8 Output
         %10 = OpConstant %6 1
         %11 = OpConstant %6 0
         %12 = OpConstantComposite %7 %10 %11 %11 %10
          %4 = OpFunction %2 None %3
          %5 = OpLabel
               OpStore %9 %12
               OpReturn
               OpFunctionEnd
```

```glsl
#version 450 core

layout (location = 0) out vec4 o_color1;
layout (location = 1) out vec4 o_color2;
layout (location = 5) out vec4 o_color3;

void main(void)
{
    o_color1 = vec4(1.0f, 0.0f, 0.0f, 1.0f);
    o_color2 = vec4(0.0f, 1.0f, 0.0f, 1.0f);
    o_color3 = vec4(0.0f, 0.0f, 1.0f, 1.0f);
}
```

```
; SPIR-V
; Version: 1.0
; Generator: Khronos Glslang Reference Front End; 1
; Bound: 17
; Schema: 0
                OpCapability Shader
        %1 =    OpExtInstImport "GLSL.std.450"
                OpMemoryModel Logical GLSL450
                OpEntryPoint Fragment %4 "main" %9 %13 %15
                OpExecutionMode %4 OriginUpperLeft
                OpSource GLSL 450
                OpName %4 "main"
                OpName %9 "o_color1"
                OpName %13 "o_color2"
                OpName %15 "o_color3"
                OpDecorate %9 Location 0
                OpDecorate %13 Location 1
                OpDecorate %15 Location 5
        %2 =    OpTypeVoid
        %3 =    OpTypeFunction %2
        %6 =    OpTypeFloat 32
        %7 =    OpTypeVector %6 4
        %8 =    OpTypePointer Output %7
        %9 =    OpVariable %8 Output
       %10 =    OpConstant %6 1
       %11 =    OpConstant %6 0
       %12 =    OpConstantComposite %7 %10 %11 %11 %10
       %13 =    OpVariable %8 Output
       %14 =    OpConstantComposite %7 %11 %10 %11 %10
       %15 =    OpVariable %8 Output
       %16 =    OpConstantComposite %7 %11 %11 %10 %10
        %4 =    OpFunction %2 None %3
        %5 =    OpLabel
                OpStore %9 %12
                OpStore %13 %14
                OpStore %15 %16
                OpReturn
                OpFunctionEnd
```

```c
typedef struct VkPipelineColorBlendAttachmentState {
    VkBool32                    blendEnable;
    VkBlendFactor               srcColorBlendFactor;
    VkBlendFactor               dstColorBlendFactor;
    VkBlendOp                   colorBlendOp;
    VkBlendFactor               srcAlphaBlendFactor;
    VkBlendFactor               dstAlphaBlendFactor;
    VkBlendOp                   alphaBlendOp;
    VkColorComponentFlags       colorWriteMask;
} VkPipelineColorBlendAttachmentState;
```

```
void vkCmdSetBlendConstants (
    VkCommandBuffer                         commandBuffer,
    const float                             blendConstants[4]);
```

```c
VkResult vkCreateFence (
        VkDevice                                device,
const VkFenceCreateInfo*                        pCreateInfo,
const VkAllocationCallbacks*                    pAllocator,
VkFence*                                        pFence);
```

```c
typedef struct VkFenceCreateInfo {
    VkStructureType        sType;
    const void*            pNext;
    VkFenceCreateFlags     flags;
} VkFenceCreateInfo;
```

```c
void vkDestroyFence (
    VkDevice                          device,
    VkFence                           fence,
    const VkAllocationCallbacks*      pAllocator);
```

```
VkResult vkQueueSubmit (
    VkQueue                             queue,
    uint32_t                            submitCount,
    const VkSubmitInfo*                 pSubmits,
    VkFence                             fence);
```

```
VkResult vkGetFenceStatus (
    VkDevice                                    device,
    VkFence                                     fence);
```

```
VkResult vkWaitForFences (
    VkDevice                              device,
    uint32_t                              fenceCount,
    const VkFence*                        pFences,
    VkBool32                              waitAll,
    uint64_t                              timeout);
```

```
VkResult vkResetFences (
    VkDevice                          device,
    uint32_t                          fenceCount,
    const VkFence*                    pFences);
```

```cpp
// kFrameDataSize is the size of the data consumed in a single frame.
// kRingBufferSegments is the number of frames' worth of data to keep.
// Create a buffer large enough to hold kRingBufferSegments copies
// of kFrameDataSize.
static const VkBufferCreateInfo bufferCreateInfo =
{
    VK_STRUCTURE_TYPE_BUFFER_CREATE_INFO, nullptr,   // sType, pNext
    0,                                               // flags
    kFrameDataSize * kRingBufferSegments,            // size
    VK_BUFFER_USAGE_TRANSFER_SRC_BIT,                // usage
    VK_SHARING_MODE_EXCLUSIVE,                       // sharingMode
    0,                                               // queueFamilyIndexCount
    nullptr                                          // pQueueFamilyIndices
};

result = vkCreateBuffer(device,
                        &bufferCreateInfo,
                        nullptr,
                        &m_buffer);

// Create kRingBufferSegments fences, all initially in signaled state.
static const VkFenceCreateInfo fenceCreateInfo =
{
    VK_STRUCTURE_TYPE_FENCE_CREATE_INFO, nullptr,

    VK_FENCE_CREATE_SIGNALED_BIT
};

for (int i = 0; i < kRingBufferSegments; ++i)
{
    result = vkCreateFence(device,
                           &fenceCreateInfo,
                           nullptr,
                           &m_fence[i]);
}
```

```cpp
// Beginning of frame - compute the segment index.
static int framesRendered = 0;
const int segmentIndex = framesRendered % kRingBufferSegments;

// Use a ring of command buffers indexed by segment.
const VkCommandBuffer cmdBuffer = m_cmdBuffer[segmentIndex];

// Wait for the fence associated with this segment.
result = vkWaitForFences(device,
                         1,
                         &m_fence[segmentIndex],
                         VK_TRUE,
                         UINT64_MAX);

// It's now safe to overwrite the data. m_mappedData is an array of
// kRingBufferSegments pointers to persistently mapped backing store for
// the source buffer.
fillBufferWithData(m_mappedData[segmentIndex]);

// Reset the command buffer. We always use the same command buffer to copy
// from a given segment of the the staging buffer, so it's safe to reset it
// here because we've already waited for the associated fence.
vkResetCommandBuffer(cmdBuffer, 0);

// Rerecord the command buffer.
static const VkCommandBufferBeginInfo beginInfo =

{
    VK_STRUCTURE_TYPE_COMMAND_BUFFER_BEGIN_INFO,    // sType
    nullptr,                                        // pNext
    VK_COMMAND_BUFFER_USAGE_ONE_TIME_SUBMIT_BIT,    // flags
    nullptr                                         // pInheritanceInfo
};
```

```
vkBeginCommandBuffer(cmdBuffer, &beginInfo);

// Copy from the staging buffer at the appropriate segment index into
// the final destination buffer.
VkBufferCopy copyRegion =
{
    segmentIndex * kFrameDataSize,                  // srcOffset
    0,                                              // dstOffset
    kFrameDataSize                                  // size
};

vkCmdCopyBuffer(cmdBuffer,
                m_stagingBuffer,
                m_targetBuffer,
                1,
                &copyRegion);

vkEndCommandBuffer(cmdBuffer);

// Reset the fence for this segment before submitting this chunk of work to
// the queue.
vkResetFences(device, 1, &m_fence[segmentIndex]);

// Note that this example doesn't use any submission semaphores. In a real
// application, you would submit many command buffers in a single submission
// and protect that submission using wait and signal semaphores.
VkSubmitInfo submitInfo =
{
    VK_STRUCTURE_TYPE_SUBMIT_INFO, nullptr,         // sType, pNext
    0,                                              // waitSemaphoreCount
    nullptr,                                        // pWaitSemaphores
    nullptr,                                        // pWaitDstStageMask
    1,                                              // commandBufferCount
    &cmdBuffer,                                     // pCommandBuffers
    0,                                              // signalSemaphoreCount
    nullptr                                         // pSignalSemaphores
};

vkQueueSubmit(m_queue,
              1,
              &submitInfo,
              m_fence[segmentIndex]);

framesRendered++;
```

```
VkResult vkCreateEvent (
    VkDevice                            device,
    const VkEventCreateInfo*            pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkEvent*                            pEvent);
```

```
typedef struct VkEventCreateInfo {
    VkStructureType         sType;
    const void*             pNext;
    VkEventCreateFlags      flags;
} VkEventCreateInfo;
```

```c
void vkDestroyEvent (
    VkDevice                        device,
    VkEvent                         event,
    const VkAllocationCallbacks*    pAllocator);
```

```
VkResult vkSetEvent (
    VkDevice                                    device,
    VkEvent                                     event);
```

```
VkResult vkResetEvent (
    VkDevice                                   device,
    VkEvent                                    event);
```

```
VkResult vkGetEventStatus (
    VkDevice                                device,
    VkEvent                                 event);
```

```c
void vkCmdSetEvent (
    VkCommandBuffer                         commandBuffer,
    VkEvent                                 event,
    VkPipelineStageFlags                    stageMask);
```

```
void vkCmdResetEvent (
    VkCommandBuffer                         commandBuffer,
    VkEvent                                 event,
    VkPipelineStageFlags                    stageMask);
```

```
void vkCmdWaitEvents (
    VkCommandBuffer                 commandBuffer,
    uint32_t                        eventCount,
    const VkEvent*                  pEvents,
    VkPipelineStageFlags            srcStageMask,
    VkPipelineStageFlags            dstStageMask,
    uint32_t                        memoryBarrierCount,
    const VkMemoryBarrier*          pMemoryBarriers,
    uint32_t                        bufferMemoryBarrierCount,
    const VkBufferMemoryBarrier*    pBufferMemoryBarriers,
    uint32_t                        imageMemoryBarrierCount,
    const VkImageMemoryBarrier*     pImageMemoryBarriers);
```

```
VkResult vkCreateSemaphore (
    VkDevice                        device,
    const VkSemaphoreCreateInfo*    pCreateInfo,
    const VkAllocationCallbacks*    pAllocator,
    VkSemaphore*                    pSemaphore);
```

```
typedef struct VkSemaphoreCreateInfo {
    VkStructureType             sType;
    const void*                 pNext;
    VkSemaphoreCreateFlags      flags;
} VkSemaphoreCreateInfo;
```

```
void vkDestroySemaphore (
    VkDevice                        device,
    VkSemaphore                     semaphore,
    const VkAllocationCallbacks*    pAllocator);
```

```c
VkResult vkQueueSubmit (
    VkQueue                             queue,
    uint32_t                            submitCount,
    const VkSubmitInfo*                 pSubmits,
    VkFence                             fence);
```

```c
typedef struct VkSubmitInfo {
    VkStructureType              sType;
    const void*                  pNext;
    uint32_t                     waitSemaphoreCount;
    const VkSemaphore*           pWaitSemaphores;
    const VkPipelineStageFlags*  pWaitDstStageMask;
    uint32_t                     commandBufferCount;
    const VkCommandBuffer*       pCommandBuffers;
    uint32_t                     signalSemaphoreCount;
    const VkSemaphore*           pSignalSemaphores;
} VkSubmitInfo;
```

```cpp
// First, perform a submission to the compute queue. The compute submission
// includes the compute -> graphics semaphore (m_computeToGfxSemaphore) in
// the pSignalSemaphores list. This sempaphore will become signaled once the
// compute queue finishes processing the command buffers in the submission.
VkSubmitInfo computeSubmitInfo =
{
    VK_STRUCTURE_TYPE_SUBMIT_INFO, nullptr,        // sType, pNext
    0,                                             // waitSemaphoreCount
    nullptr,                                       // pWaitSemaphores
    nullptr,                                       // pWaitDstStageMask
    1,                                             // commandBufferCount
    &computeCmdBuffer,                             // pCommandBuffers
    1,                                             // signalSemaphoreCount
    &m_computeToGfxSemaphore                       // pSignalSemaphores
};

vkQueueSubmit(m_computeQueue,
              1,
              &computeSubmitInfo,
              VK_NULL_HANDLE);

// Now perform a submission to the graphics queue. This submission includes
// m_computeToGfxSemaphore in the pWaitSemaphores list passed to the
// submission. We need to wait at a specific stage. Here, we just use
// VK_PIPELINE_STAGE_TOP_OF_PIPE_BIT, but if you know that the data produced
// on the source queue will be consumed at a later stage of the pipeline in
// the destination queue, you could place the wait point later.
static const VkFlags waitStages = VK_PIPELINE_STAGE_TOP_OF_PIPE_BIT;
VkSubmitInfo graphicsSubmitInfo =
{
    VK_STRUCTURE_TYPE_SUBMIT_INFO, nullptr,        // sType, pNext
    1,                                             // waitSemaphoreCount
    &m_computeToGfxSemaphore,                      // pWaitSemaphores
    &waitStages,                                   // pWaitDstStageMask
    1,                                             // commandBufferCount
    &graphicsCmdBuffer,                            // pCommandBuffers
    0,                                             // signalSemaphoreCount
    nullptr                                        // pSignalSemaphores
};

vkQueueSubmit(m_graphicsQueue,
              1,
              &graphicsSubmitInfo,
              VK_NULL_HANDLE);
```

```
VkResult vkQueueBindSparse (
    VkQueue                             queue,
    uint32_t                            bindInfoCount,
    const VkBindSparseInfo*             pBindInfo,
    VkFence                             fence);
```

```c
VkResult vkCreateQueryPool (
    VkDevice                            device,
    const VkQueryPoolCreateInfo*        pCreateInfo,
    const VkAllocationCallbacks*        pAllocator,
    VkQueryPool*                        pQueryPool);
```

```
typedef struct VkQueryPoolCreateInfo {
    VkStructureType                    sType;
    const void*                        pNext;
    VkQueryPoolCreateFlags             flags;
    VkQueryType                        queryType;
    uint32_t                           queryCount;
    VkQueryPipelineStatisticFlags      pipelineStatistics;
} VkQueryPoolCreateInfo;
```

```
void vkDestroyQueryPool (
    VkDevice                            device,
    VkQueryPool                         queryPool,
    const VkAllocationCallbacks*        pAllocator);
```

```
void vkCmdResetQueryPool (
    VkCommandBuffer                          commandBuffer,
    VkQueryPool                              queryPool,
    uint32_t                                 firstQuery,
    uint32_t                                 queryCount);
```

```c
void vkCmdBeginQuery (
    VkCommandBuffer                             commandBuffer,
    VkQueryPool                                 queryPool,
    uint32_t                                    query,
    VkQueryControlFlags                         flags);
```

```
void vkCmdEndQuery (
    VkCommandBuffer                          commandBuffer,
    VkQueryPool                              queryPool,
    uint32_t                                 query);
```

```
VkResult vkGetQueryPoolResults (
    VkDevice                             device,
    VkQueryPool                          queryPool,
    uint32_t                             firstQuery,
    uint32_t                             queryCount,
    size_t                               dataSize,
    void*                                pData,
    VkDeviceSize                         stride,
    VkQueryResultFlags                   flags);
```

```
void vkCmdCopyQueryPoolResults (
    VkCommandBuffer                    commandBuffer,
    VkQueryPool                        queryPool,
    uint32_t                           firstQuery,
    uint32_t                           queryCount,
    VkBuffer                           dstBuffer,
    VkDeviceSize                       dstOffset,
    VkDeviceSize                       stride,
    VkQueryResultFlags                 flags);
```

```
void vkCmdWriteTimestamp (
    VkCommandBuffer                    commandBuffer,
    VkPipelineStageFlagBits            pipelineStage,
    VkQueryPool                        queryPool,
    uint32_t                           query);
```

```c
// Example structure containing all available pipeline statistics counters
typedef struct VkAllPipelineStatistics {
    uint64_t inputAssemblyVertices;
    uint64_t inputAssemblyPrimitives;
    uint64_t vertexShaderInvocations;
    uint64_t geometryShaderInvocations;
    uint64_t geometryShaderPrimitives;
    uint64_t clipperInvocations;
    uint64_t clipperOutputPrimitives;
    uint64_t fragmentShaderInvocations;
    uint64_t tessControlShaderPatches;
    uint64_t tessEvaluationShaderInvocations;
    uint64_t computeShaderEvaluations;
} VkAllPipelineStatistics;
```

```
VkBufferMemoryBarrier bufferMemoryBarrier =
{
    VK_STRUCTURE_TYPE_BUFFER_MEMORY_BARRIER,  // sType
    nullptr,                                   // pNext
    VK_ACCESS_TRANSFER_WRITE_BIT,              // srcAccessMask
    VK_ACCESS_HOST_READ_BIT,                   // dstAccessMask
    VK_QUEUE_FAMILY_IGNORED,                   // srcQueueFamilyIndex
    VK_QUEUE_FAMILY_IGNORED,                   // dstQueueFamilyIndex
    buffer,                                    // buffer
    0,                                         // offset
    VK_WHOLE_SIZE                              // size
};

vkCmdPipelineBarrier(
    cmdBuffer,                                 // commandBuffer
    VK_PIPELINE_STAGE_TRANSFER_BIT,            // srcStageMask

    VK_PIPELINE_STAGE_HOST_BIT,                // dstStageMask
    0,                                         // dependencyFlags
    0,                                         // memoryBarrierCount
    nullptr,                                   // pMemoryBarriers
    1,                                         // bufferMemoryBarrierCount
    &bufferMemoryBarrier,                      // pBufferMemoryBarriers
    0,                                         // imageMemoryBarrierCount
    nullptr);                                  // pImageMemoryBarriers
```

```c
typedef struct VkRenderPassCreateInfo {
    VkStructureType                    sType;
    const void*                        pNext;
    VkRenderPassCreateFlags            flags;
    uint32_t                           attachmentCount;
    const VkAttachmentDescription*     pAttachments;
    uint32_t                           subpassCount;
    const VkSubpassDescription*        pSubpasses;
    uint32_t                           dependencyCount;
    const VkSubpassDependency*         pDependencies;
} VkRenderPassCreateInfo;
```

```c
typedef struct VkSubpassDescription {
    VkSubpassDescriptionFlags       flags;
    VkPipelineBindPoint             pipelineBindPoint;
    uint32_t                        inputAttachmentCount;
    const VkAttachmentReference*    pInputAttachments;
    uint32_t                        colorAttachmentCount;
    const VkAttachmentReference*    pColorAttachments;
    const VkAttachmentReference*    pResolveAttachments;
    const VkAttachmentReference*    pDepthStencilAttachment;
    uint32_t                        preserveAttachmentCount;
    const uint32_t*                 pPreserveAttachments;
} VkSubpassDescription;
```

```
enum
{
    kAttachment_BACK                    = 0,
    kAttachment_DEPTH                   = 1,
    kAttachment_GBUFFER                 = 2
};

enum
{
    kSubpass_DEPTH                      = 0,
    kSubpass_GBUFFER                    = 1,
    kSubpass_LIGHTING                   = 2
};

static const VkAttachmentDescription attachments[] =
{
    // Back buffer
    {
        0,                                      // flags
        VK_FORMAT_R8G8B8A8_UNORM,               // format
        VK_SAMPLE_COUNT_1_BIT,                  // samples
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,        // loadOp
        VK_ATTACHMENT_STORE_OP_STORE,           // storeOp
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,        // stencilLoadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,       // stencilStoreOp
        VK_IMAGE_LAYOUT_UNDEFINED,              // initialLayout
        VK_IMAGE_LAYOUT_PRESENT_SRC_KHR         // finalLayout
    },
    // Depth buffer
    {
        0,                                      // flags
        VK_FORMAT_D32_SFLOAT,                   // format
        VK_SAMPLE_COUNT_1_BIT,                  // samples
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,        // loadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,       // storeOp
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,        // stencilLoadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,       // stencilStoreOp
        VK_IMAGE_LAYOUT_UNDEFINED,              // initialLayout
        VK_IMAGE_LAYOUT_UNDEFINED               // finalLayout
    },
```

```cpp
    // G-buffer 1
    {
        0,                                          // flags
        VK_FORMAT_R32G32B32A32_UINT,                // format
        VK_SAMPLE_COUNT_1_BIT,                      // samples
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,            // loadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,           // storeOp
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,            // stencilLoadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,           // stencilStoreOp
        VK_IMAGE_LAYOUT_UNDEFINED,                  // initialLayout
        VK_IMAGE_LAYOUT_UNDEFINED                   // finalLayout
    }
};

// Depth prepass depth buffer reference (read/write)
static const VkAttachmentReference depthAttachmentReference =
{
    kAttachment_DEPTH,                                  // attachment
    VK_IMAGE_LAYOUT_DEPTH_STENCIL_ATTACHMENT_OPTIMAL    // layout
};

// G-buffer attachment references (render)
static const VkAttachmentReference gBufferOutputs[] =
{
    {
        kAttachment_GBUFFER,                        // attachment
        VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL    // layout
    }
};

// Lighting input attachment references
static const VkAttachmentReference gBufferReadRef[] =
{
    // Read from g-buffer.
    {
        kAttachment_GBUFFER,                        // attachment
        VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL    // layout
    },
    // Read depth as texture.
    {
        kAttachment_DEPTH,                              // attachment
        VK_IMAGE_LAYOUT_DEPTH_STENCIL_READ_ONLY_OPTIMAL // layout
    }
};
```

```cpp
// Final pass-back buffer render reference
static const VkAttachmentReference backBufferRenderRef[] =
{
    {
        kAttachment_BACK,                              // attachment
        VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL       // layout
    }
};

static const VkSubpassDescription subpasses[] =
{
    // Subpass 1 - depth prepass
    {
        0,                                      // flags
        VK_PIPELINE_BIND_POINT_GRAPHICS,        // pipelineBindPoint
        0,                                      // inputAttachmentCount
        nullptr,                                // pInputAttachments
        0,                                      // colorAttachmentCount
        nullptr,                                // pColorAttachments
        nullptr,                                // pResolveAttachments
        &depthAttachmentReference,              // pDepthStencilAttachment
        0,                                      // preserveAttachmentCount
        nullptr                                 // pPreserveAttachments
    },
    // Subpass 2 - g-buffer generation
    {
        0,                                      // flags
        VK_PIPELINE_BIND_POINT_GRAPHICS,        // pipelineBindPoint
        0,                                      // inputAttachmentCount
        nullptr,                                // pInputAttachments
        vkcore::utils::arraysize(gBufferOutputs), // colorAttachmentCount
        gBufferOutputs,                         // pColorAttachments
        nullptr,                                // pResolveAttachments
        &depthAttachmentReference,              // pDepthStencilAttachment
        0,                                      // preserveAttachmentCount
        nullptr                                 // pPreserveAttachments
    },
```

```cpp
    // Subpass 2 - g-buffer generation
    {
        0,                                              // flags
        VK_PIPELINE_BIND_POINT_GRAPHICS,    // pipelineBindPoint
        0,                                              // inputAttachmentCount
        nullptr,                                        // pInputAttachments
        vkcore::utils::arraysize(gBufferOutputs), // colorAttachmentCount
        gBufferOutputs,                             // pColorAttachments
        nullptr,                                        // pResolveAttachments
        &depthAttachmentReference,              // pDepthStencilAttachment
        0,                                              // preserveAttachmentCount
        nullptr                                         // pPreserveAttachments
    },
    // Subpass 3 - lighting
    {
        0,                                              // flags
        VK_PIPELINE_BIND_POINT_GRAPHICS,    // pipelineBindPoint
        vkcore::utils::arraysize(gBufferReadRef),    // inputAttachmentCount
        gBufferReadRef,                             // pInputAttachments
        vkcore::utils::arraysize(backBufferRenderRef), // colorAttachmentCount
        backBufferRenderRef,                        // pColorAttachments
        nullptr,                                        // pResolveAttachments
        nullptr,                                        // pDepthStencilAttachment
        0,                                              // preserveAttachmentCount
        nullptr                                         // pPreserveAttachments
    },
};

static const VkSubpassDependency dependencies[] =
{
    // G-buffer pass depends on depth prepass.
    {
        kSubpass_DEPTH,                                     // srcSubpass
        kSubpass_GBUFFER,                                   // dstSubpass
        VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,  // srcStageMask
        VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,          // dstStageMask
        VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,           // srcAccessMask
        VK_ACCESS_SHADER_READ_BIT,                      // dstAccessMask
        VK_DEPENDENCY_BY_REGION_BIT                     // dependencyFlags
    },
```

```cpp
    // Lighting pass depends on g-buffer.
    {
        kSubpass_GBUFFER,                                 // srcSubpass
        kSubpass_LIGHTING,                                // dstSubpass
        VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,    // srcStageMask
        VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,            // dstStageMask
        VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,             // srcAccessMask
        VK_ACCESS_SHADER_READ_BIT,                        // dstAccessMask
        VK_DEPENDENCY_BY_REGION_BIT                       // dependencyFlags
    },
};

static const VkRenderPassCreateInfo renderPassCreateInfo =
{
    VK_STRUCTURE_TYPE_RENDER_PASS_CREATE_INFO, nullptr,

    0,                                           // flags
    vkcore::utils::arraysize(attachments),       // attachmentCount
    attachments,                                 // pAttachments
    vkcore::utils::arraysize(subpasses),         // subpassCount
    subpasses,                                   // pSubpasses
    vkcore::utils::arraysize(dependencies),      // dependencyCount
    dependencies                                 // pDependencies
};

result = vkCreateRenderPass(device,
                            &renderPassCreateInfo,
                            nullptr,
                            &m_renderPass);
```

```
void vkCmdNextSubpass (
    VkCommandBuffer                                      commandBuffer,
    VkSubpassContents                                    contents);
```

```
void vkCmdClearAttachments (
    VkCommandBuffer                 commandBuffer,
    uint32_t                        attachmentCount,
    const VkClearAttachment*        pAttachments,
    uint32_t                        rectCount,
    const VkClearRect*              pRects);
```

```
typedef struct VkClearAttachment {
    VkImageAspectFlags      aspectMask;
    uint32_t                colorAttachment;
    VkClearValue            clearValue;
} VkClearAttachment;
```

```c
typedef union VkClearValue {
    VkClearColorValue          color;
    VkClearDepthStencilValue   depthStencil;
} VkClearValue;
```

```c
typedef struct VkClearRect {
    VkRect2D        rect;
    uint32_t        baseArrayLayer;
    uint32_t        layerCount;
} VkClearRect;
```

```c
void vkGetRenderAreaGranularity (
    VkDevice                                device,
    VkRenderPass                            renderPass,
    VkExtent2D*                             pGranularity);
```

```cpp
enum
{
    kAttachment_BACK            = 0,
    kAttachment_DEPTH           = 1,
    kAttachment_GBUFFER         = 2,
    kAttachment_TRANSLUCENCY    = 3,
    kAttachment_OPAQUE          = 4
};

enum
{
    kSubpass_DEPTH              = 0,
    kSubpass_GBUFFER            = 1,
    kSubpass_LIGHTING           = 2,
    kSubpass_TRANSLUCENTS       = 3,
    kSubpass_COMPOSITE          = 4
};

static const VkAttachmentDescription attachments[] =
{
    // Back buffer
    {
        0,                                  // flags
        VK_FORMAT_R8G8B8A8_UNORM,           // format
        VK_SAMPLE_COUNT_1_BIT,              // samples
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,    // loadOp
        VK_ATTACHMENT_STORE_OP_STORE,       // storeOp
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,    // stencilLoadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,   // stencilStoreOp
        VK_IMAGE_LAYOUT_UNDEFINED,          // initialLayout
        VK_IMAGE_LAYOUT_PRESENT_SRC_KHR     // finalLayout
```

```
    },
    // Depth buffer
    {
        0,                                          // flags
        VK_FORMAT_D32_SFLOAT,                       // format
        VK_SAMPLE_COUNT_1_BIT,                      // samples
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,            // loadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,           // storeOp
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,            // stencilLoadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,           // stencilStoreOp
        VK_IMAGE_LAYOUT_UNDEFINED,                  // initialLayout
        VK_IMAGE_LAYOUT_UNDEFINED                   // finalLayout
    },
    // G-buffer 1
    {
        0,                                          // flags
        VK_FORMAT_R32G32B32A32_UINT,                // format
        VK_SAMPLE_COUNT_1_BIT,                      // samples
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,            // loadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,           // storeOp
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,            // stencilLoadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,           // stencilStoreOp
        VK_IMAGE_LAYOUT_UNDEFINED,                  // initialLayout
        VK_IMAGE_LAYOUT_UNDEFINED                   // finalLayout
    },
    // Translucency buffer
    {
        0,                                          // flags
        VK_FORMAT_R8G8B8A8_UNORM,                   // format
        VK_SAMPLE_COUNT_1_BIT,                      // samples
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,            // loadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,           // storeOp
        VK_ATTACHMENT_LOAD_OP_DONT_CARE,            // stencilLoadOp
        VK_ATTACHMENT_STORE_OP_DONT_CARE,           // stencilStoreOp
        VK_IMAGE_LAYOUT_UNDEFINED,                  // initialLayout
        VK_IMAGE_LAYOUT_UNDEFINED                   // finalLayout
    }
};
```

```cpp
// Depth prepass depth buffer reference (read/write)
static const VkAttachmentReference depthAttachmentReference =
{
    kAttachment_DEPTH,                              // attachment
    VK_IMAGE_LAYOUT_DEPTH_STENCIL_ATTACHMENT_OPTIMAL    // layout
};

// G-buffer attachment references (render)
static const VkAttachmentReference gBufferOutputs[] =
{
    {
        kAttachment_GBUFFER,                        // attachment
        VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL    // layout
    }
};

// Lighting input attachment references
static const VkAttachmentReference gBufferReadRef[] =
{
    // Read from g-buffer.
    {
        kAttachment_GBUFFER,                        // attachment
        VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL    // layout
    },
    // Read depth as texture.
    {
        kAttachment_DEPTH,                          // attachment
        VK_IMAGE_LAYOUT_DEPTH_STENCIL_READ_ONLY_OPTIMAL // layout
    }
};

// Lighting pass - write to opaque buffer.
static const VkAttachmentReference opaqueWrite[] =
{
```

```cpp
    // Write to opaque buffer.
    {
        kAttachment_OPAQUE,                          // attachment
        VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL     // layout
    }
};

// Translucency rendering pass - translucency buffer write
static const VkAttachmentReference translucentWrite[] =
{
    // Write to translucency buffer.
    {
        kAttachment_TRANSLUCENCY,                    // attachment
        VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL     // layout
    }
};

static const VkAttachmentReference compositeInputs[] =
{
    // Read from translucency buffer.
    {
        kAttachment_TRANSLUCENCY,                    // attachment
        VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL     // layout
    },
    // Read from opaque buffer.
    {
        kAttachment_OPAQUE,                          // attachment
        VK_IMAGE_LAYOUT_SHADER_READ_ONLY_OPTIMAL     // layout
    }
};

// Final pass - back buffer render reference
static const VkAttachmentReference backBufferRenderRef[] =
{
    {
        kAttachment_BACK,                            // attachment
        VK_IMAGE_LAYOUT_COLOR_ATTACHMENT_OPTIMAL     // layout
    }
};
```

```cpp
static const VkSubpassDescription subpasses[] =
{
    // Subpass 1 - depth prepass
    {
        0,                                          // flags
        VK_PIPELINE_BIND_POINT_GRAPHICS,            // pipelineBindPoint
        0,                                          // inputAttachmentCount
        nullptr,                                    // pInputAttachments
        0,                                          // colorAttachmentCount
        nullptr,                                    // pColorAttachments
        nullptr,                                    // pResolveAttachments
        &depthAttachmentReference,                  // pDepthStencilAttachment
        0,                                          // preserveAttachmentCount
        nullptr                                     // pPreserveAttachments
    },
    // Subpass 2 - g-buffer generation
    {
        0,                                          // flags
        VK_PIPELINE_BIND_POINT_GRAPHICS,            // pipelineBindPoint
        0,                                          // inputAttachmentCount
        nullptr,                                    // pInputAttachments
        vkcore::utils::arraysize(gBufferOutputs),   // colorAttachmentCount
        gBufferOutputs,                             // pColorAttachments
        nullptr,                                    // pResolveAttachments
        &depthAttachmentReference,                  // pDepthStencilAttachment
        0,                                          // preserveAttachmentCount
        nullptr                                     // pPreserveAttachments
    },
    // Subpass 3 - lighting
    {
        0,                                          // flags
        VK_PIPELINE_BIND_POINT_GRAPHICS,            // pipelineBindPoint
        vkcore::utils::arraysize(gBufferReadRef),   // inputAttachmentCount
        gBufferReadRef,                             // pInputAttachments
        vkcore::utils::arraysize(opaqueWrite),      // colorAttachmentCount
        opaqueWrite,                                // pColorAttachments
        nullptr,                                    // pResolveAttachments
        nullptr,                                    // pDepthStencilAttachment
        0,                                          // preserveAttachmentCount
        nullptr                                     // pPreserveAttachments
```

```cpp
    },
    // Subpass 4 - translucent objects
    {
        0,                                          // flags
        VK_PIPELINE_BIND_POINT_GRAPHICS,            // pipelineBindPoint
        0,                                          // inputAttachmentCount
        nullptr,                                    // pInputAttachments
        vkcore::utils::arraysize(translucentWrite), // colorAttachmentCount
        translucentWrite,                           // pColorAttachments
        nullptr,                                    // pResolveAttachments
        nullptr,                                    // pDepthStencilAttachment
        0,                                          // preserveAttachmentCount
        nullptr                                     // pPreserveAttachments
    },
    // Subpass 5 - composite
    {
        0,                                          // flags
        VK_PIPELINE_BIND_POINT_GRAPHICS,            // pipelineBindPoint
        0,                                          // inputAttachmentCount
        nullptr,                                    // pInputAttachments
        vkcore::utils::arraysize(backBufferRenderRef), // colorAttachmentCount
        backBufferRenderRef,                        // pColorAttachments
        nullptr,                                    // pResolveAttachments
        nullptr,                                    // pDepthStencilAttachment
        0,                                          // preserveAttachmentCount
        nullptr                                     // pPreserveAttachments
    }
};

static const VkSubpassDependency dependencies[] =
{
```

```cpp
    // G-buffer pass depends on depth prepass.
    {
        kSubpass_DEPTH,                                 // srcSubpass
        kSubpass_GBUFFER,                               // dstSubpass
        VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,  // srcStageMask
        VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,          // dstStageMask
        VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,           // srcAccessMask
        VK_ACCESS_SHADER_READ_BIT,                      // dstAccessMask
        VK_DEPENDENCY_BY_REGION_BIT                     // dependencyFlags
    },
    // Lighting pass depends on g-buffer.
    {
        kSubpass_GBUFFER,                               // srcSubpass
        kSubpass_LIGHTING,                              // dstSubpass
        VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,  // srcStageMask
        VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,          // dstStageMask
        VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,           // srcAccessMask
        VK_ACCESS_SHADER_READ_BIT,                      // dstAccessMask
        VK_DEPENDENCY_BY_REGION_BIT                     // dependencyFlags
    },
    // Composite pass depends on translucent pass.
    {
        kSubpass_TRANSLUCENTS,                          // srcSubpass
        kSubpass_COMPOSITE,                             // dstSubpass
        VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,  // srcStageMask
        VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,          // dstStageMask
        VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,           // srcAccessMask
        VK_ACCESS_SHADER_READ_BIT,                      // dstAccessMask
        VK_DEPENDENCY_BY_REGION_BIT                     // dependencyFlags
    },
    // Composite pass also depends on lighting.
    {
        kSubpass_LIGHTING,                              // srcSubpass
        kSubpass_COMPOSITE,                             // dstSubpass
        VK_PIPELINE_STAGE_COLOR_ATTACHMENT_OUTPUT_BIT,  // srcStageMask
        VK_PIPELINE_STAGE_FRAGMENT_SHADER_BIT,          // dstStageMask
        VK_ACCESS_COLOR_ATTACHMENT_WRITE_BIT,           // srcAccessMask
        VK_ACCESS_SHADER_READ_BIT,                      // dstAccessMask
        VK_DEPENDENCY_BY_REGION_BIT                     // dependencyFlags
    }
};

static const VkRenderPassCreateInfo renderPassCreateInfo =
{
    VK_STRUCTURE_TYPE_RENDER_PASS_CREATE_INFO, nullptr,
    0,                                          // flags
    vkcore::utils::arraysize(attachments),      // attachmentCount
    attachments,                                // pAttachments
    vkcore::utils::arraysize(subpasses),        // subpassCount
    subpasses,                                  // pSubpasses
    vkcore::utils::arraysize(dependencies),     // dependencyCount
    dependencies                                // pDependencies
};

result = vkCreateRenderPass(device,
                            &renderPassCreateInfo,
                            nullptr,
                            &m_renderPass);
```

```c
void vkCmdExecuteCommands (
    VkCommandBuffer                     commandBuffer,
    uint32_t                            commandBufferCount,
    const VkCommandBuffer*              pCommandBuffers);
```

```c
typedef struct VkCommandBufferInheritanceInfo {
    VkStructureType                  sType;
    const void*                      pNext;
    VkRenderPass                     renderPass;
    uint32_t                         subpass;
    VkFramebuffer                    framebuffer;
    VkBool32                         occlusionQueryEnable;
    VkQueryControlFlags              queryFlags;
    VkQueryPipelineStatisticFlags    pipelineStatistics;
} VkCommandBufferInheritanceInfo;
```